
The Verilog PLI Handbook

A User's Guide
and
Comprehensive Reference
on the
Verilog Programming Language Interface

The Verilog PLI Handbook

A User's Guide
and
Comprehensive Reference
on the
Verilog Programming Language Interface

Stuart Sutherland

Sutherland HDL, Inc.



SPRINGER SCIENCE+BUSINESS MEDIA, LLC

Library of Congress Cataloging-in-Publication Data

A C.I.P. Catalogue record for this book is available
from the Library of Congress.

ISBN 978-1-4613-7279-0
DOI 10.1007/978-1-4615-5017-4

ISBN 978-1-4615-5017-4 (eBook)

Copyright © 1999 Springer Science+Business Media New York
Originally published by Kluwer Academic Publishers in 1999

All rights reserved. No part of this publication may be reproduced, stored in a retrieval system or transmitted in any form or by any means, mechanical, photocopying, recording, or otherwise, without the prior written permission of the publisher, Springer Science+Business Media, LLC

Printed on acid-free paper.

Dedication

*To my wonderful wife,
LeeAnn, and my children:
Ammon, Tamara, Hannah,
Seth and Samuel*

About the Author

Mr. Stuart Sutherland is a member of the IEEE Verilog standards committee, where he is co-chair of the PLI standards task force and technical editor for the PLI sections of the IEEE 1364 Verilog Language Reference Manual.

Mr. Sutherland has more than 14 years of experience in hardware design and over ten years of experience with Verilog. He is the founder of *Sutherland HDL Inc.*, located in Portland Oregon. Sutherland HDL provides expert Verilog HDL and Verilog PLI design services, including training, modeling, design verification and software tool evaluation. Verilog training is one of the specialties of Sutherland HDL. Prior to founding Sutherland HDL in 1992, Mr. Sutherland was as an engineer at Sanders Display Products Division in New Hampshire, where he worked on high speed graphics systems for the defense industry. In 1988, he became a senior applications engineer for Gateway Design Automation, the founding company of Verilog. At Gateway, which was acquired by Cadence Design Systems in 1989, Mr. Sutherland specialized in training and support for logic simulation, timing analysis, fault simulation, and the Verilog PLI. Mr. Sutherland has also worked closely with several EDA vendors to specify, test and bring to market Verilog simulation products.

Mr. Sutherland holds a Bachelor of Science in Computer Science, with an emphasis in Electronic Engineering Technology, from Weber State University (Ogden, Utah) and Franklin Pierce College (Nashua, New Hampshire). He has taught Verilog engineering courses at the University of California, Santa Cruz (Santa Clara extension), and has authored the popular “*Verilog HDL Quick Reference Guide*” and “*Verilog PLI Quick Reference Guide*”. He has presented tutorials and papers at the International Verilog Conference and at the International Cadence User’s Group Conference, and has won awards for best speaker and best tutorial.

Table of Contents

Table of Contents.....	ix
List of Examples.....	xvii
Foreword.....	xxi
Acknowledgments	xxiii
Introduction: The VPI and TF/ACC Parts of the Verilog PLI Standard	1
Intended Audience:	1
Generations of the PLI standard	2
TF routines	4
ACC routines	4
OVI PLI 1.0	4
OVI PLI 2.0	4
IEEE 1364-1995 PLI standard and VPI routines.....	5
IEEE 1364-1999 PLI (proposed)	5
Organization of this book	6
About the PLI examples in this book.....	7
Other sources of information	7

Part One: The VPI Portion of the Verilog PLI Standard

CHAPTER 1: Creating PLI Applications Using VPI Routines	13
1.1 The capabilities of the Verilog PLI	13
1.2 General steps to create a PLI application	15
1.3 User-defined system tasks and system functions	16
1.4 The \$hello PLI application example.....	17
1.5 The \$show_value PLI application example.....	21
1.6 Summary.....	28
CHAPTER 2: Interfacing VPI PLI Applications to Verilog Simulators	29
2.1 General PLI terms as used in this book	29
2.2 System tasks and system functions.....	31
2.3 Instantiated Verilog designs	33
2.4 How PLI applications work.....	36
2.5 calltf routines	37
2.6 compiletf routines	38
2.7 sizetf routines.....	40
2.8 VPI Simulation callback routines	41
2.9 PLI routine inputs and outputs.....	42
2.10 A complete system function example — \$pow.....	43
2.11 Interfacing PLI applications to Verilog simulators	47
2.12 Using the VPI user_data field.....	52
2.13 Compiling and linking PLI applications.....	54
2.14 Summary.....	54
CHAPTER 3: How to Use the VPI Routines	55
3.1 Specification of \$show_all_nets and \$show_all_signals.....	55
3.2 The VPI routine library.....	56
3.3 Advantages of the VPI library	57
3.4 Verilog HDL objects	59
3.5 Obtaining object handles	62
3.6 Printing messages from VPI applications.....	66
3.7 Accessing object properties	67
3.8 Reading the logic values of Verilog objects	69
3.9 Reading the current simulation time.....	71
3.10 Controlling simulation from PLI applications.....	72
3.11 A complete PLI application using VPI routines.....	73
3.12 Obtaining handles for reg and variable data types	77
3.13 Obtaining handles to the current hierarchy scope	85
3.14 Obtaining handles to multiple task/function arguments	91
3.15 Summary.....	96

CHAPTER 4: Details about the VPI Routine Library	97
4.1 PLI application performance considerations	98
4.2 The VPI string buffer.....	99
4.3 VPI error handling	100
4.4 VPI object diagrams	102
4.5 Obtaining handles for objects	105
4.6 System task and system function objects	111
4.7 Traversing Verilog hierarchy using object relationships.....	117
4.8 Writing messages to files.....	127
4.9 Reading and using simulation times	128
4.10 User-defined invocation options.....	133
4.11 Comparing VPI handles.....	137
4.12 Summary.....	138
CHAPTER 5: Reading and Modifying Values Using VPI Routines	139
5.1 Accessing objects which have logic values.....	139
5.2 Reading object logic values.....	140
5.3 Writing values to Verilog objects	161
5.4 Returning logic values of system functions.....	165
5.5 Using specparam constants as model attributes	168
5.6 Accessing Verilog memory arrays and variable arrays	171
5.7 Reading and modifying delay values.....	172
5.8 Summary.....	182
CHAPTER 6: Synchronizing to Verilog Simulations Using VPI Callbacks ...	183
6.1 PLI application callbacks.....	183
6.2 Creating an instance specific work area	184
6.3 Registering simulation callback routines.....	192
6.4 Simulation action and simulation feature callbacks	196
6.5 Simulation time callbacks.....	204
6.6 Application-scheduled callbacks at a future simulation time.....	211
6.7 Simulation event callbacks	220
6.8 Summary.....	226
CHAPTER 7: Interfacing to C Models Using VPI Routines	227
7.1 How to interface C models with Verilog simulations	228
7.2 Creating the C language model	230
7.3 A C model example	231
7.4 Creating a Verilog shell module.....	233
7.5 Creating a combinational logic interface to a C model	235
7.6 Creating a sequential logic interface to a C model.....	242
7.7 Synchronizing with the end of a simulation time step.....	245

7.8	Synchronizing with a future simulation time step	250
7.9	Allocating storage within a C model	250
7.10	Representing propagation delays in a C model	255
7.11	Summary.....	257

Part Two: The TF/ACC Portion of the Verilog PLI Standard

CHAPTER 8: Creating PLI Applications Using TF and ACC Routines261

8.1	The capabilities of the Verilog PLI	261
8.2	General steps to create a PLI application	263
8.3	User-defined system tasks and system functions	264
8.4	The \$hello PLI application example.....	265
8.5	The \$show_value PLI application example.....	269
8.6	Summary.....	274

CHAPTER 9: Interfacing TF/ACC PLI Applications to Verilog Simulators 275

9.1	General PLI terms as used in this book	275
9.2	System tasks and system functions.....	277
9.3	Instantiated Verilog designs	279
9.4	How PLI applications work.....	282
9.5	calltf routines	283
9.6	checktf routines.....	284
9.7	sizetf routines.....	286
9.8	misctf routines	287
9.9	PLI routine inputs and outputs.....	289
9.10	A complete system function example — \$pow	290
9.11	Interfacing PLI applications to Verilog simulators	292
9.12	Using the user_data field	297
9.13	Compiling and linking PLI applications.....	298
9.14	Summary.....	299

CHAPTER 10: How to Use the TF Routines301

10.1	The TF Library	301
10.2	System tasks and system functions.....	303
10.3	The PLI string buffer	306
10.4	Controlling simulation.....	307
10.5	Printing messages	307
10.6	Checking system task/function arguments	310
10.7	The TF work area	313
10.8	Reading and using simulation times.....	316
10.9	Reading simulation invocation options	322

10.10 Utility TF routines	323
10.11 A complete PLI application using TF Routines	325
10.12 Summary.....	327
CHAPTER 11: Reading and Writing Values Using TF Routines	329
11.1 Working with a 4-logic value, multiple strength level system.....	329
11.2 How the PLI writes values into Verilog	330
11.3 Reading and writing 2-state values.....	332
11.4 Reading and writing 4-state logic values using C strings.....	337
11.5 Reading Verilog strings.....	345
11.6 Reading and writing 4-state values using aval/bval encoding.....	346
11.7 Reading 4-state logic values with strengths	359
11.8 Reading from and writing into Verilog memory arrays	367
11.9 Summary.....	388
CHAPTER 12: Synchronizing to Verilog Simulations Using Miscf Routines 389	
12.1 The purpose of the miscf routine.....	389
12.2 Automatic callbacks for simulation events.....	393
12.3 Application scheduled callbacks at the end of a time step	395
12.4 Application-scheduled callbacks at a future simulation time	400
12.5 System task/function argument value change callbacks.....	403
12.6 Simulation save and restart callbacks	407
12.7 A complete example of using miscf routine callbacks	409
12.8 Summary.....	415
CHAPTER 13: Interfacing to C Models Using TF Routines	417
13.1 How to interface C models with Verilog simulations	418
13.2 Creating the C language model	420
13.3 A C model example	421
13.4 Creating a Verilog shell module	423
13.5 Creating a combinational logic interface to a C model	425
13.6 Creating a sequential logic interface to a C model	431
13.7 Synchronizing with the end of a simulation time step.....	434
13.8 Synchronizing with a future simulation time step	437
13.9 Allocating storage within a C model	437
13.10 Representing propagation delays in a C model	441
13.11 Summary.....	443
CHAPTER 14: How to Use the ACC Routines	445
14.1 Specification of \$show_all_nets and \$show_all_signals.....	445
14.2 The ACC routine library	446
14.3 Advantages of the ACC library	448
14.4 Verilog HDL objects	449

14.5	ACC handle routines	451
14.6	ACC next routines	452
14.7	Accessing object types and fulltypes.....	454
14.8	Accessing the names of objects	456
14.9	Reading the logic values of Verilog objects	457
14.10	A complete PLI application using ACC routines	459
14.11	Accessing handles for reg and variable data types	461
14.12	Obtaining handles to the current hierarchy scope	466
14.13	Obtaining handles to multiple task/function arguments	470
14.14	Summary.....	472
CHAPTER 15: Details on the ACC Routine Library		473
15.1	PLI application performance considerations	473
15.2	Initializing and configuring ACC routines	475
15.3	ACC routine error handling.....	478
15.4	Using ACC object diagrams	480
15.5	Using ACC handle routines	484
15.6	Using ACC next routines.....	485
15.7	Traversing Verilog hierarchy using object relationships.....	487
15.8	Traversing hierarchy across module ports.....	493
15.9	Identifying modules and library cells	499
15.10	Accessing loads and drivers	502
15.11	Accessing model timing	503
15.12	Counting the number of objects.....	511
15.13	Collecting and maintaining lists of object handles	513
15.14	Obtaining object handles using an object's name	516
15.15	Comparing ACC handles.....	520
15.16	Summary.....	521
CHAPTER 16: Reading and Modifying Values Using ACC Routines		523
16.1	Using ACC fetch routines.....	524
16.2	Reading object type properties	526
16.3	Accessing an object's source code location	530
16.4	Reading the simulation invocation commands	532
16.5	Accessing objects in simulation which have logic values	535
16.6	Reading the values of system task/function arguments.....	536
16.7	Reading object logic values	540
16.8	Writing values into Verilog objects.....	553
16.9	Returning logic values of system functions.....	559
16.10	Verilog objects which can have delay values	561
16.11	Reading module time scale information	563
16.12	Reading delay values	567

16.13 Writing delay values into an object	575
16.14 Reading parameter constant values	581
16.15 Using specparam constants as model attributes	583
16.16 Reading and modifying path pulse controls	585
16.17 Summary.....	589

CHAPTER 17: Synchronizing to Simulations Using the Value Change Link 591

17.1 An overview of the VCL routines	591
17.2 Adding and removing VCL flags on Verilog objects.....	592
17.3 Using the VCL consumer routine.....	595
17.4 An example of using Value Change Link routines.....	598
17.5 Summary.....	600

CHAPTER 18: Interfacing to C Models Using ACC Routines601

18.1 How to interface C models with Verilog simulations	602
18.2 Creating the C language model	604
18.3 A C model example	605
18.4 Creating a Verilog shell module.....	607
18.5 Creating a combinational logic interface to a C model	609
18.6 Creating a sequential logic interface to a C model	614
18.7 Synchronizing with the end of a simulation time step.....	617
18.8 Synchronizing with a future simulation time step	621
18.9 Allocating storage within a C model	621
18.10 Representing propagation delays in a C model	626
18.11 Summary.....	627

Appendices

APPENDIX A: Linking PLI Applications to Verilog Simulators631

A.1 The PLI Interface Mechanism.....	632
A.2 Linking PLI app's to the Cadence Verilog-XL simulator.....	641
A.3 Linking to NC-Verilog from Cadence Design Systems, Inc.	649
A.4 Linking PLI applications to the Model Technology ModelSim simulator	654
A.5 Linking PLI applications to the Avant! Polaris simulator	658
A.6 Linking to Silos III from Simucad, Inc.	661
A.7 Linking to VCS from Synopsys, Inc.	667
A.8 Linking to VeriBest from VeriBest, Inc.....	674
A.9 Summary	680

APPENDIX B: The IEEE 1364-1995 VPI Routine Library	681
B.1 VPI Objects Relationships	682
B.2 VPI routine definitions (listed alphabetically).....	707
APPENDIX C: The IEEE 1364-1995 TF Routine Library.....	717
C.1 TF routine definitions (listed alphabetically).....	718
APPENDIX D: The IEEE 1364-1995 ACC Routine Library	735
D.1 ACC Objects Relationships	736
D.2 ACC routine definitions (listed alphabetically)	749
Index	773

List of Examples

The source code files for the major examples used in this book are provided on a CD located inside the back of this book. The examples can also be downloaded from the Author's web page. Go to <http://www.sutherland.com>, and navigate the links to: "*The Verilog PLI Handbook*" — "*Download Book Examples*".

CHAPTER 1: Creating PLI Applications Using VPI Routines

Example 1-1:	\$hello — a VPI calltf routine for a PLI application	18
Example 1-2:	\$hello — VPI register function for a PLI application	19
Example 1-3:	\$show_value — a VPI compiletf routine for a PLI application	23
Example 1-4:	\$show_value — a VPI calltf routine for a PLI application	25
Example 1-5:	\$show_value — a VPI register function for a PLI application	26
Example 1-6:	\$show_value — a Verilog HDL test case using a PLI application	26

CHAPTER 2: Interfacing VPI based PLI Applications to Verilog Simulators

Example 2-1:	\$pow — a system function using VPI routines	44
Example 2-2:	\$pow — a VPI register function for a system function	50
Example 2-3:	a vlog_startup_routines array with 2 register functions	51

CHAPTER 3: How to Use the VPI Routines

Example 3-1:	\$show_all_nets — using VPI routines to access simulation objects	73
Example 3-2:	\$show_all_nets — a Verilog HDL test using a PLI application	76
Example 3-3:	\$show_all_signals, version 1 — printing values of all Verilog types	80
Example 3-4:	\$show_all_signals — a Verilog HDL test for the PLI application	83
Example 3-5:	\$show_all_signals, version 2 — obtaining the local scope handle	87
Example 3-6:	\$show_all_signals, version 3 — obtaining handles for multiple tfargs	92

CHAPTER 4: Details about the VPI Routine Library

Example 4-1:	PLIbook_numargs_vpi() — accessing system task/function arguments	114
Example 4-2:	PLIbook_getarg_handle_vpi() — accessing system task/function args	115
Example 4-3:	\$list_pathout_ports — traversing Verilog hierarchy	121
Example 4-4:	\$port_info — traversing hierarchy across module ports	123
Example 4-5:	\$count_all_prims — traversing multiple levels of hierarchy	126
Example 4-6:	\$test_invoke_options — testing for user-defined invocation options	135

CHAPTER 5: Reading and Modifying Values Using VPI Routines

Example 5-1:	PLIbook_getarg_intval_vpi() — reading Verilog values as C integers	146
Example 5-2:	PLIbook_getarg_realval_vpi() — reading Verilog values as C doubles	147
Example 5-3:	PLIbook_getarg_stringval_vpi() — reading Verilog values as C strings	148
Example 5-4:	using vpi_get_value() to read 4-state values as a C string	149
Example 5-5:	\$read_vecval — reading 4-state vector values as aval/bval pairs	154
Example 5-6:	\$read_strengthval — reading logic and strength values	157
Example 5-7:	\$read_timeval — reading time variable values	159
Example 5-8:	\$realpow — returning values of system functions	167
Example 5-9:	using vpi_get_value() to read specparam attribute values	169
Example 5-10:	\$read_delays — reading module path delays	177
Example 5-11:	\$mipd_delays — using vpi_put_delays()	180

CHAPTER 6: Synchronizing to Verilog Simulations Using VPI Callbacks

Example 6-1:	PLIbook_set_vpiworkarea(), PLIbook_get_vpiworkarea()—instance storage	187
Example 6-2:	\$read_test_vector — using VPI simulation action callbacks	200
Example 6-3:	\$my_strobe — simulation callbacks at the end of the current time	209
Example 6-4:	\$read_stimulus_ba — scheduling callbacks at a future time	214
Example 6-5:	\$my_monitor — scheduling simulation callbacks at a future time	223

CHAPTER 7: Interfacing to C Models Using VPI Routines

Example 7-1:	scientific ALU C model — combinational logic version	233
Example 7-2:	Verilog shell module for the scientific ALU C model	234
Example 7-3:	combinational logic C model interface using VPI routines	237
Example 7-4:	sequential logic C model interface using VPI routines	242
Example 7-5:	C model interface synchronized to the end of a time step	246
Example 7-6:	scientific ALU C model with latched outputs	251
Example 7-7:	combinational logic interface to the latched scientific ALU C model	252
Example 7-8:	scientific ALU Verilog shell module with pin-to-pin path delays	256

CHAPTER 8: Creating PLI Applications Using TF and ACC Routines

Example 8-1:	\$hello — a TF/ACC calltf routine for the PLI application	266
Example 8-2:	\$show_value — a TF/ACC checktf routine for a PLI application	270
Example 8-3:	\$show_value — a TF/ACC calltf routine for a PLI application	271
Example 8-4:	\$show_value — a Verilog HDL test case using a PLI application	272

CHAPTER 9: Interfacing TF/ACC based PLI Applications to Verilog Simulators

Example 9-1:	\$pow — a system function using TF/ACC routines	291
Example 9-2:	example veriusertfs array for registering TF/ACC PLI applications	296

CHAPTER 10: How to Use the TF Routines

Example 10-1:	using the tf_setworkarea() and tf_getworkarea() routines	314
Example 10-2:	storing multiple values in the TF work area	316
Example 10-3:	\$read_test_vector — using the TF routines	326

CHAPTER 11: Reading and Writing Values Using TF Routines

Example 11-1: \$realpow — using TF routines to read/write 2-state values	336
Example 11-2: \$exprinfo_test — using <code>tf_exprinfo()</code> to read 4-state values	352
Example 11-3: \$read_4state_value — using the <code>tf_evaluatep()</code> routine	354
Example 11-4: \$propagatep_test — using <code>tf_propagatep()</code> to write 4-state values	357
Example 11-5: \$nodeinfo_test — using the <code>tf_nodeinfo()</code> routine	365
Example 11-6: \$dump_mem_?? — using <code>tf_nodeinfo()</code> with Verilog memories	375
Example 11-7: \$fill_mem — using <code>tf_nodeinfo()</code> to modify Verilog memories	380
Example 11-8: \$dump_mem_bin — accessing a memory word bit; obvious method	384
Example 11-9: \$dump_mem_bin — accessing a memory word bit; efficient method	386

CHAPTER 12: Synchronizing to Verilog Simulations Using Misctf Routines

Example 12-1: using the misctf reason input	392
Example 12-2: pseudo-code for using automatic misctf routine callbacks	394
Example 12-3: \$my_strobe — using the <code>tf_rosynchronize()</code> routine	399
Example 12-4: \$read_stimulus — using the <code>tf_setdelay()</code> routine	402
Example 12-5: \$my_monitor1 — asynchronous argument value change callbacks	404
Example 12-6: \$my_monitor2 — synchronized analysis of argument value changes	406
Example 12-7: \$read_stimulus_ba — using several misctf routine callbacks	411

CHAPTER 13: Interfacing to C Models Using TF Routines

Example 13-1: scientific ALU C model — combinational logic version	423
Example 13-2: Verilog shell module for the scientific ALU C model	424
Example 13-3: combinational logic C model interface using a calltf routine	427
Example 13-4: combinational logic C model interface using a misctf routine	430
Example 13-5: sequential logic C model interface using TF routines	432
Example 13-6: C model interface synchronized to the end of a time step	435
Example 13-7: scientific ALU C model with latched outputs	438
Example 13-8: combinational logic interface to latched scientific ALU C model	440
Example 13-9: scientific ALU Verilog shell module with pin-to-pin path delays	442

CHAPTER 14: How to Use the ACC Routines

Example 14-1: \$show_all_nets — using ACC routines in a PLI application	459
Example 14-2: \$show_all_nets — Verilog HDL test case for the PLI application	460
Example 14-3: \$show_all_signals, version 1 — using the <code>acc_next()</code> routine	464
Example 14-4: \$show_all_signals1 — Verilog HDL test case for the PLI application	465
Example 14-5: \$show_all_signals, version 2—obtaining a handle for the local scope	469
Example 14-6: \$show_all_signals, version 3—obtaining handles for multiple tfargs	471

CHAPTER 15: Details on the ACC Routine Library

Example 15-1: \$list_pathout_ports — traversing Verilog hierarchy	492
Example 15-2: \$port_info — traversing Verilog hierarchy across module ports	497
Example 15-3: \$list_cells — working with module cells	501
Example 15-4: \$count_loads — using the <code>acc_count()</code> routine	511
Example 15-5: \$display_all_nets — using the <code>acc_collect()</code> routine	514

CHAPTER 16: Reading and Modifying Values Using ACC Routines

Example 16-1: \$print_invoke_commands — printing invocation commands	534
Example 16-2: \$list_nets — using acc_fetch_value() to read values as strings	543
Example 16-3: \$read_vector_value — reading vector values as aval/bval pairs	552
Example 16-4: \$func_72bit — returning 4-state vector values with system functions	560
Example 16-5: \$list_prim_delays — reading typical delays	570
Example 16-6: \$list_path_delays — reading min:typ:max delay values	574
Example 16-7: \$mipd_delays — modifying min:typ:max delays	580
Example 16-8: \$list_params — reading parameter and specparam values	582

CHAPTER 17: Synchronizing to Simulations Using the Value Change Link

Example 17-1: \$my_monitor — using the acc_vcl_add() routine	599
--------------------------------------------------------------------	-----

CHAPTER 18: Interfacing to C Models Using ACC Routines

Example 18-1: scientific ALU C model — combinational logic version	607
Example 18-2: Verilog shell module for the Scientific ALU C model	608
Example 18-3: combinational logic C model interface using ACC routines	611
Example 18-4: sequential logic C model interface using ACC routines	615
Example 18-5: C model interface synchronized to the end of a time step	618
Example 18-6: scientific ALU C model with latched outputs	622
Example 18-7: combinational logic interface to latched scientific ALU C model	624

APPENDIX A: Linking PLI Applications to Verilog Simulators

Example A-1: VPI register function for the \$hello system function	635
Example A-2: sample vlog_startup_routines array	635
Example A-3: sample veriusertfs array, as used by many Verilog simulators	639
Example A-4: Sample Verilog-XL vpi_user.c file (partial listing)	642
Example A-5: Sample Verilog-XL veriuser.c file (partial listing)	643
Example A-6: Sample Verilog-XL vconfig program output	645
Example A-7: Sample Cadence libvpi.mak file (partial listing)	647
Example A-8: Sample Cadence makfile.nc file (partial)	652
Example A-9: Sample file to specify PLI applications for the ModelSim simulator	655
Example A-10: Sample Polaris baseuser.c file	659
Example A-11: Sample Polaris fdagenmak program output	660
Example A-12: Sample file to specify PLI applications for the Silos III simulator	662
Example A-13: Sample PLI definitions file for the Silos III simulator	664
Example A-14: Sample Visual C++ make file for the Silos III simulator	664
Example A-15: PLI table file to specify PLI applications for the VCS simulator	672
Example A-16: Sample VeriBest veriuser.c file (partial listing)	675

Foreword

by Maq Mannan

*President and CEO, DSM Technologies
Chairman of the IEEE 1364 Verilog Standards Group
Past Chairman of Open Verilog International*

One of the major strengths of the Verilog language is the Programming Language Interface (PLI), which allows users and Verilog application developers to infinitely extend the capabilities of the Verilog language and the Verilog simulator. In fact, the overwhelming success of the Verilog language can be partly attributed to the existence of its PLI.

Using the PLI, add-on products, such as graphical waveform displays or pre and post simulation analysis tools, can be easily developed. These products can then be used with any Verilog simulator that supports the Verilog PLI. This ability to create third-party add-on products for Verilog simulators has created new markets and provided the Verilog user base with multiple sources of software tools.

Hardware design engineers can, and should, use the Verilog PLI to customize their Verilog simulation environment. A Company that designs graphics chips, for example, may wish to see the simulation results of a new design in some custom graphical display. The Verilog PLI makes it possible, and even trivial, to integrate custom software, such as a graphical display program, into a Verilog simulator. The simulation results can then dynamically be displayed in the custom format during simulation. And, if the company uses Verilog simulators from multiple simulator vendors, this integrated graphical display will work with all the simulators.

As another example, a company designing communication products can benefit from simulating more than just the hardware that makes up one piece of equipment. An internet router, for instance, can be better verified if the simulation models of the router have real data flowing into the router, and pass the outputs to a real environment. The Verilog PLI makes it possible to place an entire Verilog simulation into a

larger infrastructure, thus allowing real information to dynamically flow into and out of the simulation.

The possibilities are endless for how the Verilog Programming Language Interface can be utilized in design verification. If an engineer can conceive of a design verification task, and can implement the idea as a program, then the PLI will provide a way to integrate the program into the Verilog simulation environment.

Every engineer working with Verilog logic simulators can benefit from knowing the capabilities of the Verilog PLI. Knowing how to use the PLI for design verification tasks will enable an engineer to be more productive, and for designs to be of the highest quality.

Maq Mannan

Acknowledgments

I would like to express my gratitude for the many people who helped with this book.

Thank you, all those who reviewed drafts of this book and provided valuable feedback on how to make the book usable and accurate. The reviewers include:

Drew Lynch of *SureFire Verification*, **Charles Dawson** and **David Roberts** of *Cadence Design Systems*, **Steve Meyer** of *Pragmatic C Software*, **Chris Spear** of *Synopsys*, and **Marcelo Siero**, an independent software development contractor. Several of these reviewers are also members of the PLI Task Force of the IEEE 1364 Verilog Standards Committee. Thank you again, for patiently answering countless technical questions throughout the writing of this book. The task force members are **Drew Lynch**, **Charles Dawson**, **David Roberts** and **Steve Meyer**.

Thanks to all the technical and marketing representatives from various Verilog simulator vendors, for providing software for testing the examples in this book. These include: **Dave Kelf** and **Steve Gendron** of *Cadence Design Systems*, **Chris Spear** of *Synopsys*, **Tom Floodeen** of *VeriBest*, **Louie Nguyen** and **Tarak Parikh** of *Avant!*, and **Richard Jones** of *Simucad*.

A special thank you to my wife, **LeeAnn Sutherland**, for painstakingly wading through all the hardware engineering jargon and C programming terms in order to review the grammar, punctuation and spelling of the book.

Finally, thank you all others who have helped with the complex process of bringing this book from concept to print. There are too many of you to mention by name, but your efforts are all recognized and appreciated.

Stuart Sutherland

Introduction: The VPI and TF/ACC Parts of the Verilog PLI Standard

The Verilog Programming Language Interface, commonly called the Verilog PLI, is one of the more powerful features of Verilog. The PLI provides a means for both hardware designers and software engineers to interface their own programs to commercial Verilog simulators. Through this interface, a Verilog simulator can be customized to perform virtually any engineering task desired. Just a few of the common uses of the PLI include interfacing Verilog simulations to C language models, adding custom graphical tools to a simulator, reading and writing proprietary file formats from within a simulation, performing test coverage analysis during simulation, and so forth. The applications possible with the Verilog PLI are endless.

Intended Audience:

This book is written for digital design engineers with a background in the Verilog Hardware Description Language and a fundamental knowledge of the C programming language. It is expected that the reader:

- Has a basic knowledge of hardware engineering, specifically digital design of ASIC and FPGA technologies.
- Is familiar with the Verilog Hardware Description Language (HDL), and can write models of hardware circuits in Verilog, can write simulation test fixtures in Verilog, and can run at least one Verilog logic simulator.
- Knows basic C-language programming, including the use of functions, pointers, structures and file I/O.

Explanations of the concepts and terminology of digital engineering, the Verilog language or C programming are not included in this book.

Generations of the PLI standard

The Institute of Electrical and Electronics Engineers, Inc. (IEEE) Verilog standard encapsulates two major generations of the Verilog PLI. In the *IEEE 1364 Language Reference Manual* (the Verilog “LRM”), these two generations are referred to as:

- ***The TF/ACC routines***

(referred to as **PLI 1.0** prior to the IEEE standardization of Verilog)

- ***The VPI routines***

(referred to as **PLI 2.0** prior to the IEEE standardization of Verilog)

The VPI routines make up the newest generation of the PLI. These routines offer much greater capabilities and opportunities than the older TF/ACC generation of the PLI. The TF/ACC routines comprise an older generation of the Verilog PLI. These routines have been in use for many years, and are used by hundreds—possibly thousands—of PLI applications.

NOTE As the Verilog Hardware Description Language evolves to meet new engineering methods and requirements, the IEEE Verilog Standards Group will enhance the VPI routines to meet these new requirements. The IEEE 1364 standard includes the TF/ACC routines in order to provide backward compatibility and portability for older PLI applications. However, the IEEE 1364 Verilog Standards Group is only maintaining the older TF/ACC routines—the routines are not being enhanced as the Verilog language and engineering methods evolve.

To understand some of the terms used in the PLI standard, it is helpful to know a little about the origin and evolution of the PLI.

The Verilog PLI was first introduced in 1985 as part of a digital simulator called Verilog-XL, which was developed by Gateway Design Automation (Gateway later merged into Cadence Design Systems). From 1985 through 1989, both the Verilog language (the HDL) and the PLI were proprietary to Gateway—they could only be used with Gateway’s simulator products. During this time, Verilog-XL grew from an obscure product to one of the most popular digital simulators used by hardware engineers. The powerful capabilities of the PLI were one of the reasons for the growing popularity of Verilog-XL. Gateway recognized the potential of the PLI, and each year from 1985 to 1989, Gateway added new features to the PLI to meet user needs.

The widespread acceptance of Verilog-XL, along with a number of other market factors, encouraged Gateway to make the Verilog language public domain. This would

allow companies to share Verilog models and expertise, and for other software products to use the Verilog language. Gateway began the process of separating the Verilog HDL and PLI from their proprietary Verilog-XL simulator in 1989. In late 1989, Gateway Design Automation merged with Cadence Design Systems. Just a few months later, early in 1990, Cadence released the Verilog HDL and PLI to the public domain.

In 1990, a Verilog user organization was formed, called Open Verilog International (OVI), to manage the public domain Verilog language. In early 1993, OVI released a new version of the OVI Verilog HDL and PLI standard, which contained many enhancements to Verilog. In late 1993, OVI submitted a request to the IEEE to standardize both the Verilog HDL and PLI. The IEEE labeled the proposed standard as IEEE 1364, and formed a standards working group to document the Verilog standard. The IEEE 1364 Verilog Standards Group submitted a Verilog Language Reference Manual for vote in 1995, which was approved and adopted by the IEEE general membership as the official Verilog HDL and PLI standard.

The IEEE 1364 Verilog Standards Group has continued to meet and to propose enhancements to the Verilog HDL and PLI, so that the Verilog standard can evolve with hardware design technologies. In 1998, at the time this book was written, the 1364 Verilog Standards Group was preparing a new Verilog Language Reference Manual. The proposed 1364-1999 PLI standard contains many clarifications to the 1995 standard, and several enhancements to both the Verilog HDL and PLI. This book includes the clarifications of the proposed 1364-1999 standard. This book does not include the proposed enhancements because these changes were not finalized and had not been approved by the IEEE at the time this book was prepared.

This 13 year history of the PLI has led to several generations of the PLI standard. As a whole, the evolving generations of the PLI have been good for Verilog users, as each generation has added significant new features and capabilities to the PLI. However, each major generation of the PLI was developed by different engineers and sometimes different organizations. This has resulted in a number of inconsistent conventions within the standard, as well as a number of redundancies. In order to maintain an accurate and backward compatible PLI standard, the IEEE 1364 Verilog Standards Group chose to allow these inconsistencies and redundancies to remain in the IEEE standard. At times, this may make learning the PLI standard a little more difficult, but it is necessary for backward compatibility.

The version names for the major generations of the PLI are defined in the following paragraphs.

TF routines

From 1985 to 1988, the earliest generation of the PLI provided a procedural interface to Verilog simulation through constructs known as a *system tasks* and *system functions*. The library of C routines in this generation of the PLI mostly begin with the letters “**tf_**”, where the “tf” stands for “task/function”. This first generation of the PLI routines is sometimes referred to as “utility” routines.

ACC routines

In 1988 and 1989, Gateway Design Automation added several major new features to the Verilog-XL product, and the accompanying Verilog HDL and PLI. For the PLI, the enhancements provided users a procedural interface that could directly access many kinds of information within a simulation, instead of just the system tasks and functions that the TF routines could access. The ACC library of C functions all begin with the letters “**acc_**”, which stands for “access” routines.

The combined TF and ACC libraries are treated as the first generation of the PLI in this book.

OVI PLI 1.0

In 1990, Open Verilog International (OVI) was formed to manage the public domain version of the Verilog HDL and PLI. OVI labeled the TF and ACC libraries of PLI routines which Cadence had released to the public domain as the “**PLI 1.0**” standard. This term is still prevalently used, and almost every commercial Verilog simulator available today supports the 1990 OVI PLI 1.0 standard.

OVI PLI 2.0

OVI identified three major shortcomings in the PLI 1.0 standard. First, the standard had evolved over nearly 10 years with no controlling specification. This evolution had led to a very large library of C routines, with many inconsistencies and redundancies. Second, Cadence was continuing to evolve the PLI in the proprietary Verilog-XL simulator, but these new features were not part of the public domain standard. And third, while the TF and ACC routines were very useful, they did not provide a complete procedural interface to everything within a simulation.

In 1993, OVI released a “**PLI 2.0**” standard. The PLI 2.0 standard was well thought out, very simple to use, and corrected many of the shortcomings of the older PLI 1.0 standard. *OVI intended that the PLI 2.0 standard would completely replace the aging PLI 1.0 standard, making the older TF and ACC libraries obsolete.*

IEEE 1364-1995 PLI standard and VPI routines

In 1993, OVI submitted the Verilog HDL and the Verilog PLI to the Institute of Electrical and Electronics Engineers, Inc. (the IEEE) for standardization. The standardization process was completed in 1995.

OVI had intended that its PLI 2.0 standard would totally replace the older PLI 1.0 standard. The IEEE 1364 Verilog Standards Group, however, felt that maintaining backward compatibility was important in a standard, and chose to include both the older OVI PLI 1.0 and the newer OVI PLI 2.0 in the IEEE 1364 Verilog standard.

In the IEEE 1364 standard, the terms PLI 1.0 and PLI 2.0 do not exist. The 1364 Verilog standard instead refers to the three primary generations of the PLI as “*TF routines*”, “*ACC routines*” and “*VPI routines*”. The IEEE TF and ACC routines are derived from the OVI PLI 1.0 standard, and the VPI routines are derived from the OVI PLI 2.0 standard.

The VPI routines are a super set of the older TF and ACC routines, making the VPI redundant with the older TF/ACC routines. This means that most applications can be written with either TF and ACC routines, or with just the VPI routines. Note, however, that because the VPI routines are a super set of the older TF/ACC routines, there are a number of powerful applications that can only be written using the VPI routines. Many of the PLI application examples in this book are provided in two forms, using the newer VPI routines and also using older TF and ACC routines.

IEEE 1364-1999 PLI (proposed)

Hardware design technology and engineering methods continuously change, and the IEEE 1364 Verilog Standards Group is dynamically evolving the Verilog HDL and PLI standard to meet the changing needs of hardware engineering. In 1999, the IEEE Verilog Standards Group is submitting a proposed 1364-1999 Verilog standard for vote by the IEEE membership.

All enhancements and new features in the Verilog PLI standard will only be made in the VPI routines. The older ACC and TF routines will continue to remain in the standard to provide backward compatibility, but no new ACC and TF routines will be added. It should be noted, however, that the proposed 1999 standard does contain several corrections and clarifications to the 1364-1995 ACC and TF routines.

This book includes many of the clarifications to the PLI that are in the proposed 1364-1999 standard. The proposed enhancements to the PLI are not included, however, because these enhancements were not finalized and approved by the IEEE at the time this book was being prepared.

Organization of this book

The PLI is not a trivial topic to explain or to comprehend. It combines in a very unique way digital electronic engineering, digital modeling and simulation, and advanced C programming. To add to the complexity, the PLI has a long and complex history, having been in use—and evolving—since 1985.

This book is structured to serve two specific needs:

1. A tutorial on how to write PLI applications.

Towards this goal, the book provides explanations of the extensive PLI library from a usage point of view, with many examples showing how the PLI routines are applied.

2. A reference book on the IEEE 1364 Verilog PLI standard.

The appendices of this book contain a full description of the IEEE standard for the PLI. These appendices are intended to serve as a quick and convenient reference whenever you need to see the syntax or semantics of a specific PLI routine. These reference appendices do not explain the usage of each routine. That is covered in the more generalized “how-to” chapters.

This book is divided into two major parts, in order to effectively cover both generations of the PLI. Each part of the book is completely self contained, allowing a reader to focus on either or both of the generations of the Verilog PLI.

Part One of this book, comprising chapters 1 through 7, presents the **VPI** portion of the PLI standard, which is sometimes referred to as “**PLI 2.0**”. The VPI is the newest of generations of the PLI, and is by far the most versatile of all the PLI generations. The first chapters in Part One explain how to use the VPI routines. A large number of small, but useful examples illustrate the concepts presented. The later chapters in Part One contain a more comprehensive description of the VPI library of routines, and how each routine is used. More extensive examples of PLI applications are presented in these chapters. Appendix B contains a quick reference of the IEEE 1364 Verilog PLI standard for the VPI portion of the PLI.

Part Two of the book, comprising chapters 8 through 18, presents the **TF** portions and the **ACC** portions of the PLI standard. This is the older generation of the PLI, and is still very widely used in PLI applications. The TF and ACC portions of the standard are often referred to as “**PLI 1.0**”. The first chapters in Part Two explain how to use the TF and ACC routines, and include many useful examples. The later chapters present the TF library of routines in much greater detail, followed by a detailed presentation of the ACC library of routines. Appendices C and D contain a quick reference of the IEEE 1364 Verilog PLI standard for the TF and ACC portion of the standard.

About the PLI examples in this book

This book covers using the PLI for simple to intermediate applications. A large number of examples are provided, which are realistic and useful programs for engineers designing with Verilog. The examples are necessarily kept reasonably short and simplified in capability. It is expected that readers will build larger and more advanced PLI applications based on the examples provided in this book.

Portability is a primary focus in this book. The IEEE has standardized the Verilog PLI, but some Verilog simulators do not strictly adhere to this standard, and sometimes ambiguities in the complex IEEE standard are interpreted differently. Many Verilog simulators do not support the entire IEEE standard, and other simulators may add proprietary extensions to the IEEE standard. The examples presented within this book will work with any Verilog simulator that is fully IEEE compliant. Many of the examples will also work with simulators that support a subset of the IEEE standard. This book does not use any PLI extensions that are proprietary to a specific simulator.

The PLI examples in this book were tested using the **Verilog-XL™** simulator, from *Cadence Design Systems, Inc.* Verilog-XL runs on most Unix operating systems and on the Windows-NT operating system. For this book, all examples were tested using Verilog-XL version 2.6.9, on a Pentium II system with the Windows-NT operating system. The examples were compiled with the Microsoft Visual C++ compiler.

Many of the examples were also tested using the **NC-Verilog™** simulator from *Cadence Design Systems, Inc.*, the **VCS™** simulator from *Synopsys®, Inc.*, the **ModelSim™** simulator from *Model Technology, Inc.*, the **Polaris™** simulator from *Avant! Corporation.*, the **Silos III™** simulator from *Simucad®, Inc.*, and the **VeriBest™** simulator from *VeriBest Inc.*.

The source code files for the major examples used in this book are provided on a CD located inside the back of this book. The examples can also be downloaded from the Author's web page. Go to <http://www.sutherland.com>, and navigate the links to: "The Verilog PLI Handbook" — "Download Book Examples".

Other sources of information

There are many excellent books on the C programming language. Readers of this book who are novice C programmers will want to have their personal favorite C reference book handy.

There are also number of books on the Verilog HDL, any of which can serve as an excellent companion to this book. Some suggested books are:

IEEE Std 1364-1995, Language Reference Manual (LRM)—IEEE Standard Hardware Description Language based on the Verilog Hardware Description Language.

Copyright 1996, IEEE, Inc., New York, NY. ISBN 1-55937-727-5.

This is the official Verilog HDL and PLI standard. The book is a syntax and semantics reference, not a tutorial for learning Verilog. Softcover, 665 pages. For information on ordering, call 1-800-678-4333 (US and Canada), 1-908-981-9667 (elsewhere).



NOTE At the time this book was being prepared, in mid 1998, the IEEE 1364 Verilog Standards Group was preparing a proposed 1364-1999 version of Verilog standard. This book includes many of the clarifications to the PLI from the proposed 1364-1999 Language Reference Manual, but does not include any enhancements to the PLI.

Verilog Quickstart by James M. Lee

Copyright 1997, Kluwer Academic Publishers, Norwell MA.
ISBN: 0-7923-9927-7.

An excellent book for learning the Verilog HDL. The book teaches the basics of Verilog modeling, without getting into the more obscure or advanced aspects of the Verilog language. Hardcover, 328 pages. For more information, refer to the web site www.wkap.nl/book.htm/0-7923-9927-7.

The Verilog Hardware Description Language, 4th Edition by Donald E. Thomas and Philip R. Moorby

Copyright 1998, Kluwer Academic Publishers, Norwell MA.
ISBN: 0-7923-8166-1

A complete book on Verilog, covering RTL modeling, behavioral modeling and gate level modeling. The book has more detail on the gate, switch and strength level aspects of Verilog than many other books. Hardcover, 336 pages. For more information, refer to the web site www.wkap.nl/book.htm/0-7923-8166-1.

The Complete Verilog Book by Vivek Sagdeo

Copyright 1998, Kluwer Academic Publishers, Norwell MA.
ISBN: 0-7923-8188-2

A comprehensive reference book on the Verilog HDL, with discussion on event-driven and cycle-based simulation, analog modeling, and a brief introduction of the Verilog PLI. Hardcover, 496 pages. For more information, refer to the web site www.wkap.nl/book.htm/0-7923-8188-2.

Verilog HDL Quick Reference Guide by Stuart Sutherland

Copyright, 1992-1998, Sutherland HDL, Inc., Portland, OR.

A small, pocket-sized quick reference on the syntax and semantics of the complete Verilog language. Softcover, 32 pages. For more information, refer to the web site www.sutherland.com.

Verilog PLI Quick Reference Guide by Stuart Sutherland

Copyright, 1992-1998, Sutherland HDL, Inc., Portland, OR.

A small, pocket-sized quick reference on the syntax and semantics of the complete Verilog PLI standard. The main contents of this quick reference are reprinted with permission as the Appendices in this book. Softcover, 64 pages. For more information, refer to the web site www.sutherland.com.

Part One:

The VPI Portion

of the Verilog PLI Standard

CHAPTER 1

Creating PLI Applications Using VPI Routines

This chapter uses two short examples to introduce how PLI applications are created using the VPI routines. The first example will allow a Verilog simulator to execute the ubiquitous *Hello World* C program. The second example uses the PLI to access specific activity within a Verilog simulation, by listing the name and current logic value of a Verilog signal. The purpose of these examples is to introduce how to write a PLI application using the library of VPI routines in the PLI standard. Subsequent chapters in this part of the book build on the concepts presented in this chapter and show how to write more complex PLI applications.

1.1 The capabilities of the Verilog PLI

The Verilog **Programming Language Interface**, commonly referred to as the Verilog **PLI**, is a user-programmable procedural interface for Verilog digital logic simulators. The PLI gives Verilog users a way to extend the capabilities of a Verilog simulator by providing a means for the simulator to invoke other software programs, and then exchange information and synchronize activity with those programs.

The creators of the Verilog Hardware Description Language (HDL), which was first introduced in 1985, intended that the Verilog language should be focused on modeling digital logic. They felt that design verification tasks not directly related to modeling the design itself, such as reading and writing disk files, were already capably handled in other languages and should not be replicated in the Verilog HDL. Instead, the creators of Verilog added a procedural interface—the PLI—that provides a way for end-users of Verilog to write design verification tasks in the C programming language, and then have Verilog logic simulators execute those programs.

The capabilities of the PLI extend far beyond simple file I/O operations. The Verilog PLI allows Verilog designers to interface virtually any program to Verilog simulators. Some common engineering tasks that are required in the design of hardware, and for which the PLI is aptly fitted, include:

- C-language models

Abstract hardware models are sometimes represented in the C language instead of the Verilog HDL, perhaps to protect intellectual property or to optimize the performance of simulation. The PLI provides a number of ways to pass data from Verilog simulation to a C-language model and from the C-language model back into Verilog simulation.

- Access to programming language libraries

The PLI allows Verilog models and test programs to access the libraries of the C programming language, such as the C math library. Through the PLI, Verilog can pass arguments to a C math function, and pass the return of the math function back to the Verilog simulator.

- Reading test vector files

Test vectors are a common method of representing stimulus for simulation. A file reader program in C is easy to write, and, using the PLI, it is easy to pass values read by the C program to a Verilog simulation.

- Delay calculation

Through the PLI, a Verilog simulator can call an ASIC or FPGA vendor's delay calculator program. The delay calculator can estimate the effects that fanout, voltage, temperature and other variances will have on the delays of ASIC cells or gates. Through the PLI, the delay calculator program can modify the Verilog simulation data structure so that simulation is using the more accurate delays.

- Custom output displays

In order to verify design functionality, a logic simulator must generate outputs. The Verilog HDL provides formatted text output, and most simulators provide waveform displays and other graphical output displays. The PLI can be used to extend these output format capabilities. A video controller, for example, might generate video output for a CRT or other type of display panel. Using the PLI, an engineer can take the outputs from simulation of the video controller, and, while simulation is running, pass the data to a program which can display the data in the same form as the real display output. The PLI allows custom output programs to dynamically read simulation values while simulation is running.

- Co-simulation

Complex designs often require several types of logic simulation, such as a mix of analog and digital simulations, or a mix of Verilog and VHDL simulations. Each simulator type could be run independently on its own regions of the design, but it is far more effective to simulate all regions of the design at the same time, using multiple simulators together. The PLI can be used as a communication channel for Verilog simulators to transfer data to and from other types of simulators.

- Design debug utilities

The Verilog logic simulators on the market vary a great deal in what they provide for debugging design problems. There may be some type of information about a design that is needed, which cannot be accessed by the simulator's debugger. The Verilog PLI can access information deep inside a simulation data structure. It is often a very trivial task to write a short C program that uses the PLI to extract some specific information needed for debugging a design.

- Simulation analysis

Most Verilog logic simulators are intended to simply apply input stimulus to a model and show the resulting model outputs. The PLI can be used to greatly enhance simulation by analyzing what had to happen inside a design to generate the output values. The PLI can be used to generate toggle check reports (how many times each node in a design changed value), code coverage reports (what Verilog HDL statements were exercised by the input tests), power usage (how much energy the chip consumed), etc.

The preceding examples are just a few of the tasks that can be performed using the Verilog PLI. The PLI provides full access to anything that is happening in a Verilog simulation, and allows external programs to modify the simulation. This open access enables truly unlimited possibilities. If an engineer can conceive of an application that requires interaction with Verilog simulations, and can write a program to perform the task, the PLI can be used to interface that application to a Verilog simulator.

1.2 General steps to create a PLI application

A PLI application is a user-defined C language application which can be executed by a Verilog simulator. The PLI application can interact with the simulation by both reading and modifying the simulation logic and delay values.

The general steps to create a PLI application are:

1. Define a **system task** or **system function** name for the application.
2. Write a C language **callif routine** which will be executed by the simulator whenever simulation encounters the system task name or the system function name.
Optionally, additional C language routines can be written which will be executed by the simulator for special conditions, such as when the simulator compiler encounters the system task/function name.
3. **Register** the system task or system function name and the associated C language routines with the Verilog simulator. This registration tells the simulator about the new system task or system function name, and the name of the **callif routine** associated with that task or system function (along with any other routines).
4. **Compile** the C source files which contain the PLI application routines, and link the object files into the Verilog simulator.

1.3 User-defined system tasks and system functions

In the Verilog language, a **system task** or a **system function** is a command which is executed by a Verilog simulator. The name of a system task or a system function begins with a *dollar sign* (\$).

A **system task** is used like a programming statement that is executed from a Verilog *initial procedure* or *always procedure*. For example, to print a message at every positive edge of a clock, the **\$display** system task can be used.

```
always @(posedge clock)
$display("chip_out = %h", chip_out);
```

A **system function** is used like a programming function which returns a value. A function can be called anywhere that a logic value can be used. For example, to assign a random number to a vector at every positive edge of a clock, the **\$random** system function can be used.

```
always @(posedge clock)
vector <= $random();
```

The IEEE 1364 Verilog standard allows system tasks and system functions to be defined in three ways:

- A standard set of built-in system tasks and system functions.

These are defined as part of the IEEE standard, such as `$display`, `$random`, and `$finish`. All IEEE compliant Verilog simulators will have these standard system tasks and system functions.

- Simulator specific system tasks and system functions.

These are proprietary commands which are defined as part of a simulator, and may not exist in other simulators. Examples are `$save` to create a simulation check point file or `$db_settrace` to enable debug tracing.

- User-defined system tasks and system functions.

These are created through the Programming Language Interface. A PLI application developer specifies the name and functionality of the system task/function.

User-defined system task or system function names begin with a dollar sign (\$), just as the built-in system tasks and system functions. The user-defined system task/function name is then associated with a user-defined C application. When a Verilog simulator encounters the system task/function name, it will execute the C application associated with the name. This simple association allows Verilog simulation users to extend the capability of a Verilog simulator in any manner desired. Chapter 2 presents the concept of system tasks and system functions in more detail.

1.4 The \$hello PLI application example

The well known “Hello world” C program is a quick way to show how PLI applications are created. Though the C program itself is simple, it illustrates one of the most powerful features of the Verilog PLI—the ability for a Verilog designer to extend the Verilog language by having a Verilog simulation dynamically execute a user-defined program while simulation is running.

1.4.1 Step One: defining a \$hello system task

The first step in creating a PLI application is create a new system task or system function name. A `$hello` user-defined system task will be created for this PLI application. An example of using `$hello` is:

```
module test;  
initial  
  $hello();  
endmodule
```

1.4.2 Step Two: writing a calltf routine for \$hello

The second step in developing a PLI application is to write a C language routine which will be called when a Verilog simulator executes the `$hello` system task. In this example, the simulator will call a C language routine which prints the message:

```
Hello World!
```

The C language routine which will be called by the simulator is referred to as a *calltf routine*. This routine is a user-defined C function and can be any name. Example 1-1 lists the C source code for the *calltf routine* for `$hello`.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.01/hello_vpi.c
- Verilog test bench: Chapter.01/hello_test.v
- Verilog-XL results log: Chapter.01/hello_test.log

Example 1-1: `$hello` — a VPI *calltf routine* for a PLI application

```
#include <stdlib.h>      /* ANSI C standard library */  
#include <stdio.h>        /* ANSI C standard input/output library */  
#include "vpi_user.h"     /* IEEE 1364 PLI VPI routine library */  
  
int PLIBbook_hello_calltf(char *user_data)  
{  
    vpi_printf("\nHello World!\n\n");  
    return(0);  
}
```

In the above example:

- The header file `vpi_user.h` contains the VPI library of functions provided by the PLI standard. The library is a collection of C functions which can be used by a PLI application to interact with a Verilog simulator. The `vpi_user.h` header file is part of the IEEE 1364 Verilog standard.

- The `vpi_printf()` function used in the preceding example is very similar to the `C printf()` function. The difference is that `vpi_printf()` will write its message to the simulator's output window and the simulator's output log file, whereas the `C printf()` writes its message to the operating system's standard out channel.

1.4.3 Step Three: Registering the \$hello system task

The third step in creating a new PLI application is to tell the Verilog simulator about the new system task or system function name and the C routines which are associated with the application. This process is referred to as *registering* the PLI application.

The PLI IEEE 1364 standard provides an *interface mechanism* for this process. The interface mechanism defines:

- The type of application, which is a system task or system function.
- The system task or system function name.
- The name of the *calltf routine* and other C routines associated with the system task or system function.

The VPI interface mechanism for PLI applications is part of the IEEE standard, and is implemented in the same way with all IEEE compliant Verilog simulators. The interface mechanism requires two steps:

1. Create a *register function* which specifies the PLI application information in a `s_vpi_register_systf` structure and calls the `vpi_register_systf()` VPI routine.
2. List the name of the register function in a C array called `vlog_startup_routines`.

Example 1-2 lists a register function for the `$hello` PLI application. Chapter 2 presents more details of registering PLI applications using the VPI interface mechanism.

Example 1-2: \$hello — VPI register function for a PLI application

```
void PLIbook_hello_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname   = "$hello";
    tf_data.calltf   = PLIbook_hello_calltf;
    tf_data.compiletf = NULL;
    tf_data.sizetf   = NULL;
    vpi_register_systf(&tf_data);
}
```

1.4.4 Step Four: Compiling and linking the \$hello system task

The final step in creating a new PLI application is to compile the C source code containing the application and linking the compiled files to the Verilog simulator. Once the application has been linked to the simulator, the simulator can invoke the *calltf routine* when simulation executes the \$hello system task name.

Compiling and linking is not part of the IEEE standard. This process is both specific to the simulator and the operating system on which simulations are run. Appendix A presents the steps for compiling and linking PLI applications for several major Verilog simulators.

1.4.5 Running simulations with the \$hello system task

Once a PLI application has been registered with a Verilog simulator, the new system task or system function can be used in a Verilog model, just as with built-in system tasks and system functions. When the Verilog simulator encounters the new user-defined system task name, it will call the C routine which has been associated with that application.

The following Verilog model can be used to test the \$hello PLI application:

```
module test;

initial
begin
    $hello();
    #10 $stop;
    $finish;
end
endmodule
```

Figure 1-1 shows the results of simulating this Verilog model using the Cadence Verilog-XL simulator with the Cadence SimVision graphical user interface.

Figure 1-1: \$hello — simulation results using a PLI application

The screenshot shows the Cadence Verilog-XL Turbo NT simulation interface. The top menu bar includes File, Control, Set, Show, Select, Tools, Options, and Help. Below the menu is a toolbar with icons for play, stop, and other simulation controls. The main window displays a Verilog testbench code and its simulation results.

```
08 ****
09 `timescale 1ns / 1ns
10 module test;
11
12 initial
13 begin
14     $hello;
15     #10 $stop;
16     $finish;
17 end
18 endmodule
```

Scope: test Subscopes: |

Compiling source file "hello_test.v"
Highest level modules:
test

Hello World!

L15 "hello_test.v": \$stop at simulation time 10
Type ? for help
C1 > |

Ready

1.5 The \$show_value PLI application example

The *\$hello* PLI application, though trivial from a C programming aspect, effectively illustrates the powerful capability provided by the Programming Language Interface—the ability to have a Verilog simulator call a user-defined C routine during simulation. The PLI provides far more capabilities, however. The C routine which is called by the simulator can use the VPI library to read information from the simulation data structure, and to dynamically modify logic values and delay values in the data structure.

The system task *\$show_value* illustrates using the PLI to allow a C routine to read current logic values within a Verilog simulation. The *\$show_value* system task requires one argument, which is the name of a net or reg in the Verilog design. The *calltf* routine will then use the library of VPI routines to read the current logic value of that net or reg, and print its name and logic value to the simulator's output screen. An example of using *\$show_value* is:

```
module test;
    reg a, b, ci;
    wire sum, co;
    ...
    initial
        begin
            ...
            $show_value(sum);
        end
    endmodule
```

Two user-defined C routines will be associated with `$show_value`:

- A C routine to verify that `$show_value` has the correct type of argument.
- A C routine to print the name and logic value of the signal.

These programs are presented in the following sections.

1.5.1 Writing a compiletf routine for `$show_value`

The PLI allows users to provide a C routine to verify that a system task or system function is being used correctly and has the correct types of arguments. This C routine is referred to as a *compiletf routine*.

The C functions in the VPI library can access the arguments of a system task or system function and determine the Verilog data types of the objects listed in the arguments. This example uses some of these C functions. Only a brief explanation of each function is provided in this chapter. Subsequent chapters will present more details about each function in the VPI library.

The C source code for the *compiletf routine* associated with `$show_value` is listed below. The C functions and constants from the PLI library are shown in bold text.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.01/show_value_vpi.c`
- Verilog test bench: `Chapter.01/show_value_test.v`
- Verilog-XL results log: `Chapter.01/show_value_test.log`

Example 1-3: \$show_value — a VPI compiletf routine for a PLI application

```
int PLIbook_ShowVal_compiletf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, arg_handle;
    int         arg_type;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);
    if (systf_handle == NULL) {
        vpi_printf("ERROR: $show_value failed to obtain systf handle\n");
        tf_dofinish(); /* abort simulation */
        return(0);
    }

    /* obtain handles to system task arguments */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    if (arg_iterator == NULL) {
        vpi_printf("ERROR: $show_value requires 1 argument\n");
        tf_dofinish(); /* abort simulation */
        return(0);
    }

    /* check the type of object in system task arguments */
    arg_handle = vpi_scan(arg_iterator);
    arg_type = vpi_get(vpiType, arg_handle);
    if (arg_type != vpiNet && arg_type != vpiReg) {
        vpi_printf("ERROR: $show_value arg must be a net or reg\n");
        vpi_free_object(arg_iterator); /* free iterator memory */
        tf_dofinish(); /* abort simulation */
        return(0);
    }

    /* check that there are no more system task arguments */
    arg_handle = vpi_scan(arg_iterator);
    if (arg_handle != NULL) {
        vpi_printf("ERROR: $show_value can only have 1 argument\n");
        vpi_free_object(arg_iterator); /* free iterator memory */
        tf_dofinish(); /* abort simulation */
        return(0);
    }
    return(0);
}
```

Observe the following key points in the above example:

- **vpiHandle** is a special data type defined in the VPI library. This data type is used to store pointers to objects in a Verilog simulation data structure.
- The **vpi_handle()** function returns a **handle**, which is a form of pointer, to a specific object in a Verilog simulation data structure. In this example, handles are

obtained for the instance of the `$show_value` system task which called the PLI application, and for the object listed as the first argument of the system task.

- The `vpi_iterate()` function obtains an *iterator* for all of a specific type of object. An iterator is essentially a pointer to the next object in a series of objects. The iterator object is stored in a `vpiHandle` data type. In the above example, an iterator is obtained for all of the arguments to the `$show_value` system task.
- The `vpi_scan()` obtains a handle for the next object that is referenced by an iterator. In the above example, each object returned by `vpi_scan()` is the next argument of the `$show_value` system task.
- The `vpi_get()` function returns the value of integer properties of a specific object in the simulation data structure. The first input to this function is a constant that defines the property to be obtained. The `vpiType` property used in this example identifies the type of object passed as the argument to `$show_value`. In this example, the test is checking that the argument is a Verilog net or reg data type.
- The `vpi_free_object()` function releases memory allocated by the `vpi_iterate()` routine. Refer to section 4.5.3 on page 108 in Chapter 4 for more details on when `vpi_free_object()` needs to be used in a PLI application.
- The `vpi_printf()` function writes a message to the simulator's output window.
- The `tf_dofinish()` function causes the simulator to exit. In this compiletf routine, this function is used to abort simulation if there is a usage error with the `$show_value` system task.



NOTE In the IEEE 1364-1995 standard, there is no VPI routine that will cause a simulator to exit if a PLI application detects an error. The `tf_dofinish()` routine is part of the older TF library. The proposed IEEE 1364-1999 standard will add a new VPI routine that can be used to abort simulation.

1.5.2 Writing the calltf routine for `$show_value`

The *calltf routine* is the C routine which will be executed when simulation encounters the `$show_value` system task during simulation.

The following example lists the C source code for the *calltf routine* associated with `$show_value`. Later chapters define the routines used in this example in full detail. The C functions, constants and data types from the PLI library are shown in bold text.

Example 1-4: \$show_value — a VPI calltf routine for a PLI application

```
int PLIBook_ShowVal_calltf(char *user_data)
{
    vpiHandle      systf_handle, arg_iterator, arg_handle, net_handle;
    s_vpi_value   current_value;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument
       compiletf has already verified only 1 arg with correct type */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    net_handle = vpi_scan(arg_iterator);
    vpi_free_object(arg_iterator); /* free iterator memory */

    /* read current value */
    current_value.format = vpiBinStrVal; /* read value as a string */
    vpi_get_value(net_handle, &current_value);
    vpi_printf("Signal %s has the value %s\n",
               vpi_get_str(vpiFullName, net_handle),
               current_value.value.str);
    return(0);
}
```

Most of the VPI routines used in the above example were also used in the *compiletf routine*, and are described in section 1.5.1 on page 22. Additional routines used in the *calltf routine* are:

- The **vpi_get_str()** function returns the value of string properties of a specific object in the simulation data structure. The first input to this function is a constant that defines the property to be obtained. The **vpiFullName** property used in this example is the Verilog hierarchical path name of the net listed as an argument to *\$show_value*.
- The **vpi_get_value()** function obtains the logic value of a Verilog object. The value is returned into an **s_vpi_value** structure, which is defined as part of the VPI standard. The **vpi_get_value()** function allows the value to be obtained in a variety of formats. In this example, the value is obtained as a string, with a binary representation of the value.

1.5.3 Registering the \$show_value PLI application

The *register function* is used to inform the Verilog simulator about the PLI application. The information about the application is specified in a **s_vpi_register_systf** structure.

The following example lists the C source code for the register function for `$show_value`. The next chapter explains the process of registering PLI applications in greater detail.

Example 1-5: `$show_value` — a VPI register function for a PLI application

```
void PLIbook_ShowVal_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$show_value";
    tf_data.calltf   = PLIbook_ShowVal_calltf;
    tf_data.compiletf = PLIbook_ShowVal_compiletf;
    tf_data.sizetf    = NULL;
    vpi_register_systf(&tf_data);
    return;
}
```

1.5.4 A test case for `$show_value`

The following Verilog HDL source code is a small test case for the `$show_value` application. The example lists the values of a simple 1-bit adder after a few input test values have been applied. Observe how `$show_value` is used as a Verilog programming statement within the test bench.

Example 1-6: `$show_value` — a Verilog HDL test case using a PLI application


```
'timescale 1ns / 1ns
module test;
    reg a, b, ci, clk;
    wire sum, co;
    addbit i1 (a, b, ci, sum, co);
    initial
        begin
            clk = 0;
            a = 0;
            b = 0;
            ci = 0;
            #10 a = 1;
            #10 b = 1;
            $show_value(sum);
            $show_value(co);
            $show_value(i1.n3);
            #10 $stop;
            $finish;
        end
endmodule
```

```
/** A gate level 1 bit adder model ***/
`timescale 1ns / 1ns
module addbit (a, b, ci, sum, co);
    input a, b, ci;
    output sum, co;
    wire a, b, ci, sum, co,
          n1, n2, n3;
    xor (n1, a, b);
    xor #2 (sum, n1, ci);
    and (n2, a, b);
    and (n3, n1, ci);
    or #2 (co, n2, n3);
endmodule
```

1.5.5 Output from running the \$show_value test case

Following is the output from simulating the test case for the `$show_value` application, using the Cadence Verilog-XL simulator with its SimVision user interface.

Figure 1-2: `$show_value` — sample output using Verilog-XL



1.6 Summary

This chapter has presented two short examples which illustrate how PLI applications are created using the VPI routine library in the PLI standard. The *\$hello* PLI application showed how a Verilog simulation can execute another C program. The *\$show_value* application illustrated how arguments can be passed from a Verilog simulation to a PLI application, and how the application can access information about the arguments.

The principles illustrated by the *\$show_value* example can be readily expanded to provide much more powerful and useful capabilities. In Chapter 3, the functionality of *\$show_value* will be extended to create an application called *\$list_nets*, which automatically finds all nets in a module and prints the values of each net. Another example in the same chapter extends *\$show_value* to create *\$list_signals*, which automatically finds all nets, regs and variables in a module and prints the values of each one.

The next chapter presents much greater detail on how PLI applications interact with Verilog simulations. Subsequent chapters in Part One of this book will discuss the complete VPI library, and include many examples of how the VPI routines are used to create PLI applications which interact with Verilog simulators.

CHAPTER 2

Interfacing VPI based PLI Applications to Verilog Simulators

D*on't skip this chapter!* This chapter defines the terminology used by the Verilog PLI and how PLI applications which use the VPI library are interfaced to Verilog simulators. All remaining chapters in Part One of this book assume the principles covered in this chapter are understood. The general concepts presented are:

- PLI terms as used in this book
- Generations of the PLI standard
- System tasks and system functions
- How VPI routines work
- A Complete PLI application example
- Interfacing PLI applications to Verilog simulators

2.1 General PLI terms as used in this book

The Verilog PLI has been in use since 1985, and, over the years, many reference manuals, technical papers and training courses have been written about the PLI. Unfortunately, these documents have not used consistent terms for describing the PLI. This inconsistency can make it difficult to cross reference information about the PLI in different documents.

This book strives to clearly define the PLI terms and consistently use this terminology. The general PLI terms used in this book are.

C program (often abbreviated to *program*)

A complete software program written in the C or C++ programming language. A C program must include a C *main* function.

C function

A function written in the C or C++ programming language. A C function does not include a C *main* function.

Verilog function

A Verilog HDL function, written in the Verilog Hardware Description Language. Verilog functions can only be called from Verilog source code.

User-defined system task or system function (abbreviated to *system task/function*)

A user-defined system task or system function is a construct that is used in Verilog HDL source code. The name of a user-defined system task or system function must begin with a dollar sign (\$), and it is used in the Verilog language in the same way as a Verilog HDL standard system task or system function. When simulation encounters the user-defined system task or system function, the simulator will execute a *PLI routine* associated with the system task/function.

PLI routine

A C function which is part of a *PLI application*. The PLI routine is executed by the simulator when simulation encounters a user-defined system task or system function. The VPI portion of the PLI standard defines several PLI routine types: *calltf routines*, *completf routines*, *sizetf routines* and *simulation callback routines*. These terms are defined in more detail later in this chapter.

PLI application (sometimes abbreviated to *application*)

A *user-defined system task or system function* with an associated with a set of one or more *PLI routines*.

PLI library

A library of C functions which are defined in the Verilog PLI standard. PLI library functions are called from *PLI routines*, and enable the *PLI routines* to interact with a Verilog simulation. The IEEE 1364 standard provides three PLI libraries, referred to as the *VPI library*, the *TF library* and the *ACC library*.

VPI routines, TF routines and ACC routines

C functions contained in the *PLI library*. The term *routine* is used for these functions, to avoid confusion with *Verilog functions* and *C functions*.

2.2 System tasks and system functions

The Verilog Hardware Description Language provides constructs called “**system tasks**” and “**system functions**”. A system task or system function is not a hardware modeling construct. It is a command that is executed by a Verilog simulator.

The names of system tasks and system functions always begin with a dollar sign (\$). A **system task** is analogous to a subroutine. When the task is called, the simulator branches to a program that executes the functionality of the task. When the task has completed, the simulation returns to the next statement following the task call. A **system function** returns a value. When the function is called, a simulator executes the program associated with the function, which returns a value into the simulation. The return value becomes part of the statement that called the function. The return value of a function may be defined to be an integer, a vector (with a specific bit size) or a floating point number.

2.2.1 Built-in system tasks and system functions

The IEEE 1364 Verilog standard defines a number of system tasks and system functions that are built into all Verilog simulators. Examples of the standard system tasks are \$display, \$stop, and \$finish. Some of the standard system functions are \$time and \$random. These system tasks and system functions are part of the Verilog language, and the syntax and usage of these routines are covered in Verilog language books.

The IEEE 1364 standard also allows simulation vendors to add proprietary system tasks and system functions which are specific to a simulator product. For example, the Cadence Verilog-XL simulator adds the command \$db_settrace, which is a system task that turns on Verilog statement execution tracing to aid in design debugging.

2.2.2 User-defined system tasks and system functions

The Verilog PLI provides a means for Verilog users to add additional system tasks and system functions. When simulation encounters a user-defined system task or system function, the simulator branches to a user-supplied C function, executes that function, and then returns back to executing the Verilog source at the point where the system task/function was called.

User-defined system task and system function names must begin with a dollar sign (\$), and may only use the characters that are legal in Verilog names, which are:

a—z A—Z 1—9 _ \$

Following are examples of legal user-defined system task and system function names:

`$rand64`

`$cell_count`

`$GetVector`

2.2.3 Overloading built-in system tasks and system functions

A user-defined system task or system function can be given the same name as a system task/function that is built into the simulator. This allows a PLI application developer to overload the operation of a built-in system task/function with new functionality. For example, the built-in system function `$random` will return a 32-bit signed random number, with the random sequence defined by the simulator. If the verification of a design requires a special random number sequence, a random number generator could be written in the C language and then associated with a user-defined system function with the same `$random` name. The simulator would then call the user PLI application instead of the built-in random number generator.

2.2.4 System tasks are used as procedural programming statements

System tasks are procedural programming statements in the Verilog HDL. This means a system task may only be called from a Verilog `initial` procedure, an `always` procedure, a Verilog HDL `task` or a Verilog HDL `function`. A system task may not be called outside of a procedure, such as from a continuous assignment statement.

The following example calls a user-defined system task from an `always` procedure:

```
always @(posedge clock)
$read_test_vector("vectors.pat", input_vector);
```

2.2.5 System functions are used as expressions

System functions are expressions in the Verilog HDL. This means a system function may be called anywhere a logic value may be used. The return value of the system function is considered the result of the expression. System functions may be called from a Verilog `initial` procedure, an `always` procedure, a Verilog HDL `task`, a Verilog HDL `function`, an `assign` continuous assignment statement, or as an operand in a compound expression.

The following examples call a user-defined system function several different ways:

```

always @(posedge clock)
  if (chip_out !== $pow(base,exponent))
    ...
initial
  $monitor("output = %f", $pow(base,exponent));
assign temp = i + $pow(base,exponent);

```

2.2.6 System task/function arguments

A system task or system function may have any number of arguments, including none. The user-supplied PLI application can read the values of the arguments of a system task/function. If the argument is a Verilog `reg` or variable (`integer`, `real` and `time`), then the PLI application can also write values into the task/function argument.

System task/function arguments are typically numbered from 1 to N, with the left-most argument being argument number 1 (this numbering scheme is only a convention when working with VPI routines, but it is part of the PLI standard when using the older TF/ACC routines). For example:

```

module top (...);
  ...
  reg [15:0] in1;
  ...
  my_chip u1 (in1, out1);
  ...
  initial
    $cell_count(u1);
  ...
  always @(posedge clock)
    $read_vector_file("vectors.pat", in1);
  ...
endmodule

```

2.3 Instantiated Verilog designs

In a Verilog HDL module, a reference to another module is referred to as a **module instance**. In the following Verilog source code, module `top` contains two *instances* of module `bottom`.

```
module top;
    reg [7:0] in1;
    wire [7:0] out1;
    bottom b1 (in1[7:4], out1[7:4]);
    bottom b2 (in2[3:0], out2[3:0]);
endmodule

module bottom (in, out);
    ...
endmodule
```

The Verilog PLI requires an instantiated Verilog data structure. This means that a Verilog simulator has created a data structure which contains the complete hierarchy tree that has been represented by one module containing instances of other modules. There is no limit to the number of levels of hierarchy in the Verilog language.

The PLI does not work directly with Verilog HDL source code—The PLI is simply a procedural interface provided by a simulator so that PLI applications can access the data structure of a simulation.

Multiple instances of a system task or system function

The system task or system function which invokes a PLI application can be instantiated any number of times in the Verilog HDL source code. Consider the following Verilog HDL source code fragment:

```
module top;
    ...
    middle m1 (...);           //module instance
    middle m2 (...);           //module instance

    initial
        $my_app_1(in1, out1);   //system task instance

    always @(posedge clock)
        $my_app_1(in2, out2);   //system task instance

endmodule

module middle (...);
    ...
    bottom b1 (...);           //module instance
    bottom b2 (...);           //module instance

endmodule
```

```

module bottom (...);
  ...
  initial
    $my_app_2(in3, out3);      //system task instance
  always @(posedge clock)
    $my_app_2(in4, out4);      //system task instance
  endmodule

```

In the preceding example, four different conditions are shown that can exist in a Verilog simulation.

- A single instance of a system task that is invoked one time

In module **top**, the first *instance* of **\$my_app_1** is in a Verilog **initial** procedure, which will be called once at the beginning of a simulation.

- A single instance of a system task that is invoked many times

In module **top**, the second *instance* of **\$my_app_1** is in a Verilog **always** procedure, which will be called every positive edge of the clock signal.

- Multiple instances of a system task that are each invoked one time

In module **bottom**, the first *instance* of **\$my_app_2** is in a Verilog **initial** procedure, which will be called once at the beginning of a simulation.

- Multiple instances of a system task that are each invoked many times


In module **bottom**, the second *instance* of **\$my_app_2** is in a Verilog **always** procedure, which will be called every positive edge of the clock signal.

Note that in the preceding Verilog source code example, module **bottom** is instantiated two times in module **middle**. The two instances of **bottom**, with two instances of **\$my_app_2** inside **bottom**, creates four unique instances of **\$my_app_2**. Module **middle** is instantiated twice in module **top**, resulting in eight unique instances of **\$my_app_2**.

Thus, in the instantiated simulation data structure for the preceding Verilog source code, there are *two instances* of the **\$my_app_1** system task, and *eight instances* of the **\$my_app_2** system task. Each of these system task instances will have unique logic values for the inputs and outputs. One of the requirements of a PLI application is to allow for multiple unique instances of the application. This is an important consideration when an application allocates static variables or allocates system memory. Section 6.2 on page 184 in Chapter 6 discusses these issues in more depth.

2.4 How PLI applications work

The PLI standard allows a Verilog user to create a user-defined system task/function name and to associate one or more user-defined C routines with the system task/function name. When a Verilog simulator encounters the user-defined system task/function, it will execute the user-defined C routine associated with the name.



2.4.1 The types of PLI routines

The VPI portion of PLI standard defines several types of PLI routines which can be associated with a system task or system function. The type of the routine determines *when* the simulator will execute the routine. Some types of routines are run-time routines, which are invoked during simulation, and some types are compile time routines, which are invoked prior to simulation. The types of PLI routines are:

- *calltf routines*
- *compiletf routines*
- *sizetf routines*
- *simulation callback routines*

The purpose of these routines is defined in the following sections of this chapter.

2.4.2 Associating routine types with system task/functions

A system task/function can have multiple PLI routines associated with it, as long as each routine is a different type. A `$read_vector_file` application, for example, might have three routines associated with it—one which is called at compile time to perform syntax checking, one which is called when simulation first starts running to open the test vector file, and one which is called during simulation at every positive edge of clock to read vectors from the file.

Up to 3 different C routines may be associated with a system task/function in the VPI portion of the PLI standard: the *calltf routine*, *compiletf routine* and *sizetf routine*. The

simulation callback routine is not directly associated with a system task/function name. Instead, the *simulation callback routine* is called for various types of activity that can occur during a simulation, such as logic value changes and start or completion of simulation.

The next section of this chapter defines in detail the purpose and usage of the *calltf routine*, *compiletf routine*, *sizetf routine* and the *simulation callback routine*.


NOTE → PLI routines are sometimes referred to as PLI programs or C programs. In correct C terminology, PLI routines are actually C *functions*. A *program* in the C language terminology requires a *main* function. PLI applications do not contain a main function. Instead, the routines which make up the PLI application are linked into a Verilog simulator, and the simulator contains the C main function. This allows the simulator to call the appropriate PLI routine whenever its associated system task/function name is encountered.

2.4.3 C versus C++

Most Verilog simulators expect that PLI applications are written in ANSI C. The PLI libraries are compliant with the ANSI C standard, but also include the appropriate prototypes and other information required for C++. This allows PLI applications to be developed in either C or C++. However, many Verilog simulators were written in the C language, and may or may not support linking C++ applications to the simulator. PLI application developers who wish to work with C++ should first check the limitations of the simulator product to which the applications will be linked. For maximum portability, PLI applications should be written using ANSI C.


2.5 *calltf routines*

The *calltf routine* is executed when simulation is running. For the \$pow example that follows, at every positive edge of clock the *calltf routine* associated with \$pow will be executed by the simulator. The routine is user-defined, and, for this example, will calculate a 32-bit number representing x to the power of y.



2.6 compiletf routines

The **compiletf routine** is called by the simulator before simulation starts running—in other words, before simulation time 0. The routine may be called by the simulator’s compiler, or when the simulator loads and prepares its simulation data structure. The purpose of **compiletf routine** is to verify that a system task/function is being used correctly (e.g.: to check that the call to the PLI application has the correct number of arguments, and that the arguments are the correct Verilog data types). For example, if the \$pow system function were passed a module name as its second argument instead of a value, then the C math library would probably get an error when trying to use the name as an exponent value. Bad data can result in errors ranging from incorrect results to program crashes. Therefore, it is important to verify that each usage of a system task/function has valid arguments.



TIP Use a *completf routine* to improve the performance of PLI programs. Since this routine is only called one time prior to simulation time 0, the code to check the correctness of arguments is only executed once. The *calltf routine*, which may be invoked millions of times during a simulation, does not need to repeat the syntax checking, and therefore can execute more efficiently.



NOTE The *completf routine* will be called one time for each instance of a user-defined system task/function. If a design used the *\$pow* user-defined system function in three different places, the *completf routine* for *\$pow* would be called three times. If a simulator allows system tasks and system functions to be invoked from the simulator's debug command line, then the *completf routine* will be called prior to execution of the *calltf routine* for each interactive usage.

Limitations on completf callbacks

The intent of the *completf routine* is to verify the correctness of the arguments of an instance of a system task/function. The IEEE 1364 standard does not impose explicit

restrictions on what VPI routines will, or will not work in a *compiletf routine*. There are, however, considerations which should be observed. Since the *compiletf routine* is called prior to simulation time 0, variables have not yet been assigned any values, and nets which are assigned values as soon as simulation starts may not yet reflect those values.

NOTE → *The compiletf routine should only be used for syntax checking!* Do not use this routine to perform run-time duties such as allocating memory or opening files. The PLI standard does ensure that activity performed at compile time will remain in effect at simulation time. A *simulation callback routine* should be used instead of the *compiletf routine* to perform run-time work at the very beginning of a simulation.

To ensure that a PLI application will be portable, only the following activities should be performed in a *compiletf routine*:

- Accessing the arguments of an instance of a system task/function. Handles for the arguments can be obtained, and the properties of the arguments can be accessed to verify correctness. The values of literal numbers can be read, but variables and nets will not have been initialized. Note that Verilog parameter constants can be redefined for each instance of a Verilog module. The IEEE 1364 standard does not state whether the *compiletf routine* will be called before or after parameter redefinitions have occurred. Reading parameter constant values from a *compiletf routine* may not yield the same results on every simulator.
- Using VPI routines such as `vpi_printf()`, which do not access objects in simulation. Attempting to obtain handles and access properties beyond the system task arguments may not work correctly in all simulators. Attempting to write logic or delay values onto objects may result in PLI application errors.
- Registering *simulation callback routines* using `vpi_register_cb()` for the end of compilation or the start of simulation. Using *simulation callback routines* is presented in Chapter 6).


2.7 *sizetf routines*

The *sizetf routine* is only used with system functions which return scalar or vector values (system functions can also return integer and real number values). A function can return a user-specified number of bits, and the simulator compiler may need to know how many bits to expect from the return, in order to correctly compile the statement from which the system function is called. A *sizetf routine* is called one time, before simulation time 0. The *sizetf routine* returns to the simulator how many bits wide the return value of system function will be.

NOTE

Even if a system function is used multiple times in the Verilog source code, the *sizetf routine* is only called once. The return value from the first call is used to determine the return size for all instances of the system function.

In the `$pow` example, the *calltf routine* will return a 32-bit value. Therefore, the *sizetf routine* associated with `$pow` needs to return a value of 32 to the simulator.



Limitations on sizetf callbacks

The intent of the *sizetf routine* is to notify the simulator compiler of the return size for system functions. The simulator may invoke the *sizetf routine* very early in the compilation phase of a Verilog design, and, at this early stage, the Verilog hierarchy may not have been generated. In addition, the *sizetf routine* is only called one time for a system function name, and the return size applied to all instances of the system function. For these reasons, only standard C language statements and functions should be used in a *sizetf routine*. An error may result if any VPI routines are called from a *sizetf routine*. Any memory or static variables allocated by a *sizetf routine* may not remain in effect for when simulation starts running.

2.8 VPI Simulation callback routines


The VPI provides a means for PLI applications to be called for specific simulation events. The VPI portion of the PLI standard refers to these types of routines as *simulation callback routines*. The IEEE 1364 standard defines a number of simulation events for which a PLI routine can be called. The full list of callback reasons is presented in Chapter 6. Some examples of simulation related callbacks are:

- The beginning of Verilog simulation (just before the start of simulation time 0).

- Entering debug mode (such as when the `$stop` built-in system task is executed).
- End of simulation (such as when the `$finish` built-in system task is executed).
- Change of value of a signal.
- Execution of a Verilog procedural statement.
- End of a the current simulation time step.
- Simulation has reached a specified simulation time step.

A common usage of *simulation callback routines* is to perform tasks at the very beginning and the very end of a simulation. A PLI application to read test vectors might need to open a test vector file at the start of simulation, and close the file at the end of simulation.

```
always @(posedge clock)
  $read_test_vector("vectors.pat", in1);
```



2.9 PLI routine inputs and outputs

All types of PLI routines are C functions, and a Verilog simulator will utilize the inputs and return values of the functions when the simulator calls the functions. The

PLI standard defines what the inputs should be, and how the simulator should use the return value.

The input that is passed to the *compileif routine*, *calltf routine* and *sizetf routine* is a character pointer, which points to the user_data value that was specified when the system task/function was registered. Refer to section 2.12 on page 52 for more details on defining and using the user_data value.

The input that is passed to the *simulation callback routine* is a pointer to a structure containing information about the simulation callback. Refer to Chapter 6 for more details on using *simulation callback routines*.

The *compileif routine*, *calltf routine*, *sizetf routine* and *simulation callback routine* are expected to be integer functions in the PLI standard. However, the only return value which is used by the simulator is the return of the *sizetf routine*. This return represents the bit-width of the system function return value. The return value from the *compileif routine*, *calltf routine* and *simulation callback routine* are not used, and are ignored by the simulator.

It is a common practice to declare the *compileif routine*, *calltf routine* and *simulation callback routine* as void functions, since their return value is ignored, even though all of these routines are expected to be integer functions. However, declaring the functions as a different type from the PLI standard prototype may result in a C compiler warning message when the routine is compiled. To prevent this warning message, the pointer to the void function can be cast to a pointer to an int function. For example:

```
(int(*)())PLIBbook_PowCalltf
```

Note, however, that some C compilers might not accept this cast.

2.10 A complete system function example — \$pow

The following example illustrates all the parts of a complete system function, with the user-defined name of *\$pow*.

- The system function will return a 32-bit value.
- The system function requires two arguments, a base value and an exponent value. Both arguments must be numeric integer values.

To implement the `$pow` functionality, four user-defined PLI routines are used:

- A *sizetf routine* to establish the return size of `$pow`.
- A VPI *compiletf routine* to verify that the `$pow` arguments are valid values.
- A *calltf routine* to calculate the base to the power of the exponent each time `$pow` is executed by the simulator.
- A VPI *simulation callback routine* to print a message when simulation firsts starts running (immediately prior to simulation time 0).

The `$pow` system function does not really need the VPI *simulation callback routine*. This example PLI application includes the routine, in order to illustrate its usage.



TIP The *compiletf routine*, *sizetf routine*, *calltf routine*, and *simulation callback routine* are C functions. These functions may be located in separate files, or they can all be in the same file. Typically, smaller PLI applications place all the routines in a single file, while larger, more complex applications might break them into multiple files.



TIP The *compiletf routine*, *sizetf routine*, *calltf routine*, and *simulation callback routine* names can be any legal C function name. However, a typical Verilog HDL design may include several PLI applications. Using unique function name conventions can avoid possible name conflicts when multiple PLI applications are linked together.

For this chapter, the `$pow` application has been simplified to only work with integer numbers. In practice, this application would probably accept both integer and floating point numbers for inputs, and return a floating point number for the result. Example 5-8 on page 167 in Chapter 5 presents another version of `$pow` that works with floating point numbers.



CD The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.02/pow_vpi.c`
- Verilog test bench: `Chapter.02/pow_test.v`
- Verilog-XL results log: `Chapter.02/pow_test.log`

Example 2-1: `$pow` — a system function using VPI routines

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include "vpi_user.h"     /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"      /* IEEE 1364 PLI TF routine library */
                           /* TF library is used for aborting on error */
```

```
*****
 * Sizetf application
 *****/
int PLIbook_PowSizetf(char *user_data)
{
    return(32); /* $pow returns 32-bit values */
}

*****
 * compiletf application to verify valid systf args.
 *****/
int PLIbook_PowCompiletf(char *user_data)
{
    vpiHandle systf_handle, arg_itr, arg_handle;
    int tfarg_type;

    systf_handle = vpi_handle(vpiSysTfCall, NULL);
    arg_itr = vpi_iterate(vpiArgument, systf_handle);
    if (arg_itr == NULL) {
        vpi_printf("ERROR: $pow requires 2 arguments\n");
        tf_dofinish();
        return(0);
    }
    arg_handle = vpi_scan(arg_itr);
    tfarg_type = vpi_get(vpiType, arg_handle);
    if ( (tfarg_type != vpiReg) &&
        (tfarg_type != vpiIntegerVar) &&
        (tfarg_type != vpiConstant) ) {
        vpi_printf("ERROR: $pow arg1 must be number, variable or net\n");
        tf_dofinish();
        return(0);
    }
    arg_handle = vpi_scan(arg_itr);
    if (arg_handle == NULL) {
        vpi_printf("ERROR: $pow requires 2nd argument\n");
        tf_dofinish();
        return(0);
    }
    tfarg_type = vpi_get(vpiType, arg_handle);
    if ( (tfarg_type != vpiReg) &&
        (tfarg_type != vpiIntegerVar) &&
        (tfarg_type != vpiConstant) ) {
        vpi_printf("ERROR: $pow arg2 must be number, variable or net\n");
        tf_dofinish();
        return(0);
    }
    if (vpi_scan(arg_itr) != NULL) {
        vpi_printf("ERROR: $pow requires only 2 arguments\n");
        vpi_free_object(arg_itr);
        tf_dofinish();
        return(0);
    }
}
```

```

/*****
 * calltf to calculate base to power of exponent and return result.
 *****/
#include <math.h>
int PLIbook_PowCalltf(char *user_data)
{
    s_vpi_value value_s;
    vpiHandle systf_handle, arg_itr, arg_handle;
    int base, exp, result;

    systf_handle = vpi_handle(vpiSysTfCall, NULL);
    arg_itr = vpi_iterate(vpiArgument, systf_handle);
    if (arg_itr == NULL) {
        vpi_printf("ERROR: $pow failed to obtain systf arg handles\n");
        return(0);
    }

    /* read base from systf arg 1 (compiletf has already verified) */
    arg_handle = vpi_scan(arg_itr);
    value_s.format = vpiIntVal;
    vpi_get_value(arg_handle, &value_s);
    base = value_s.value.integer;

    /* read exponent from systf arg 2 (compiletf has already verified) */
    arg_handle = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* not calling scan until returns null */
    vpi_get_value(arg_handle, &value_s);
    exp = value_s.value.integer;

    /* calculate result of base to power of exponent */
    result = (int)pow( (double)base, (double)exp );

    /* write result to simulation as return value $pow */
    value_s.value.integer = result;
    vpi_put_value(systf_handle, &value_s, NULL, vpiNoDelay);
    return(0);
}

/*****
 * Start-of-simulation application
 *****/
int PLIbook_PowStartOfSim()
{
    vpi_printf("\n$pow PLI application is being used.\n\n");
    return(0);
}


```

NOTE → In the IEEE 1364-1995 standard, there is no VPI routine to cause a simulator to exit if a PLI application detects an error. The above example uses the `tf_dofinish()` routine from the older TF library to abort compilation. The proposed IEEE 1364-1999 standard will add a new VPI routine that can be used to abort simulation.

Simulation results for \$pow

Following is the output from running a simulating which uses the *\$pow* application, using the Cadence Verilog-XL simulator with its SimVision user interface.

Figure 2-1: *\$pow* — sample simulation results using Verilog-XL



The screenshot shows the Cadence Verilog-XL Turbo NT window. The menu bar includes File, Control, Set, Show, Select, Tools, Options, and Help. The toolbar contains icons for play, stop, zoom, and other simulation controls. The main window displays a Verilog testbench for the \$pow application. The code is as follows:

```

30
31     /* Test $pow with valid values */
32     #1 $display("$pow[2,3] returns %d", $pow[2,3]);
33     #1 result = $pow(a,b);
34     #1 $display("$pow(a,b) returns %d [a=%d b=%d]", result, a, b);
35     #1 $stop;
36     $finish;
37 end
38
39 endmodule

```

The scope dropdown shows "test". The message area indicates:

Compiling source file "pow_test.v"
Highest level modules:
test

\$pow PLI application is being used.

\$pow(2,3) returns 8
\$pow(a,b) returns 1 (a=1 b=0)
L35 "pow_test.v": \$stop at simulation time 4
Type ? for help

Ready

2.11 Interfacing PLI applications to Verilog simulators

As the previous sections in this chapter have shown, a PLI application may comprise:

- A system task or system function name
- A *calltf* routine
- A *completetf* routine
- A *sizetf* routine
- Any number of *simulation callback routines*

After these items have been defined, the PLI application must be interfaced with a Verilog simulator. The PLI standard provides an *interface mechanism* to make the associations between the system task/function name and the application routines. There are two generations of the interface mechanism, one that was created for the older TF and ACC libraries, and a newer mechanism created for the VPI library. This chapter presents the VPI interface mechanism. The older TF/ACC interface mechanism is presented in Part Two of this book.

The VPI interface mechanism involves three basic steps:

1. Create a register function, which associates the system task/function name with the application routines.
2. Notify the Verilog simulator about the registration function.
3. Link the applications into a Verilog simulator, so the simulator can call the appropriate routine when the system task/function name is encountered.

NOTE → The PLI standard does not provide any guidelines on how PLI applications should be linked into a Verilog simulator. There are many different C compilers and operating systems available to Verilog users, and each compiler and operating system has unique methods for compiling and linking programs. Since the commands for compiling and linking are not part of the IEEE 1364 standard, this process is not covered in this chapter. Appendix A presents the steps involved to compile and link PLI applications for several Verilog simulators.

The VPI Interface Mechanism is defined as part of the IEEE 1364 standard, providing a consistent method for all Verilog simulators to use. The VPI interface is used to specify:

1. A system task/function *name*.
2. The application *type*, which is a *task*, *sized function*, *integer function*, *time function* or *real function*.
3. Pointers to the C functions for a *calltf routine*, *compiletf routine* and *size tf routine*, if the routines exist (it is not required—and often not necessary—to provide each class of routine).
4. A character pointer *user_data* value, which the simulator will pass to the *calltf routine*, *compiletf routine* and *size tf routine* each time they are called.

The process of specifying the PLI application information is referred to as *registering* the application. To register a PLI application, the information about the application is specified in an **s_vpi_systf_data** structure. This structure is defined as part of the VPI standard, in the PLI *vpi_user.h* file. The definition is:

```

typedef struct t_vpi_systf_data {
    int type;
    int subtype;
    char *tfname;
    int (*calltf)();
    int (*completf)();
    int (*sizetf)();
    char *user_data;
} s_vpi_systf_data, *p_vpi_systf_data;

```

Table 49 explains the fields of the `s_vpi_systf_data` structure:

s_vpi_systf_data Field	Definition
type	Defines the type of a PLI application as being either a system task or a system function. This field must be set to the C constant: vpiSysTask or vpiSysFunc .
subtype	Defines the return type of a system function. This field is only used if the type field is vpiSysFunc , in which case the subtype must be set to the C constant: vpiSysFuncInt , vpiSysFuncReal , vpiSysFuncTime or vpiSysFuncSized .
tfname	Specifies the name of the system task/function; must be a quoted literal string.
calltf	Specifies a pointer to the C function that will be called by a simulator for the application's <i>calltf routine</i> .
completf	Specifies a pointer to the C function that will be called by a simulator for the application's <i>completf routine</i> .
sizetf	Specifies a pointer to the C function that will be called by a simulator for the application's <i>sizetf routine</i> .
user_data	Specifies a character pointer—the value of the pointer will be passed to the PLI application routines each time a routine is called.

Table 2-1: VPI interface mechanism `s_vpi_systf_data` structure fields

NOTE

The proposed IEEE 1364-1999 standard will change the names of the constants involved with system task/functions:

- **vpiSysFuncInt** changes to **vpiIntFunc**
- **vpiSysFuncTime** changes to **vpiTimeFunc**
- **vpiSysFuncReal** changes to **vpiRealFunc**
- **vpiSysFuncSized** changes to **vpiSizedFunc**

The constant names from the 1364-1995 standard will be aliased to the new constant names, to provide backward compatibility.

2.11.1 The steps required to register a PLI application

To register a system task or system function using the VPI interface mechanism involves the following steps:

1. Create a C function to register the system task/function. The C function name is user-defined and can be any legal C name.
2. Allocate an **s_vpi_systf_data** C structure.
3. Fill in the fields of the structure with the information about the system task or system function.
4. Register the system task/function by calling the VPI routine **vpi_register_systf()**.
5. Add the name of the C function created in step 1 to a C language array called **vlog_startup_routines**. This array is typically contained in a C source file provided with the simulator called **vpi_user.c**, though the IEEE 1364 standard does not require that the file be called that name.

The following example registers the **\$pow** application presented in this chapter.



TIP The register function can be located in any C source file, but is typically located in the same file as the PLI application C functions.

Example 2-2: **\$pow** — a VPI register function for a system function

```
/* prototypes of PLI application routine names */  
int PLIBbook_PowSizetf(), PLIBbook_PowCalltf(), PLIBbook_PowCompletf(),  
PLIBbook_PowStartOfSim();
```

```

void PLIbook_pow_register()
{
    s_vpi_systf_data tf_data;
    s_cb_data    cb_data_s;
    s_vpi_time   time_s;

    tf_data.type      = vpiSysFunc;
    tf_data.sysfuncype = vpiSysFuncSized;
    tf_data.tfname    = "$pow";
    tf_data.calltf    = PLIbook_PowCalltf;
    tf_data.compiletf = PLIbook_PowCompiletf;
    tf_data.sizetf    = PLIbook_PowSizetf;
    vpi_register_systf(&tf_data);
}

```

2.11.2 Notifying Verilog simulators about PLI applications

Once the register function has been defined, a Verilog simulator must be notified of the name of the register function, so that the simulator can call the functions and register the PLI applications. The VPI standard requires that all Verilog simulators provide a special array, called *vlog_startup_routines*, in order to notify the simulators about the register functions. All PLI applications which will be called by the simulator should have register function listed in this array.

An example *vlog_startup_routines* array is listed below, with entries for two register functions.

Example 2-3: a *vlog_startup_routines* array with 2 register functions

```

/* prototypes of the PLI application routines */
extern void PLIbook_pow_register(), PLIbook_ShowVal_register();

void (*vlog_startup_routines[])() =
{
    /*** add user entries here ***/
    PLIbook_pow_register,
    PLIbook_ShowVal_register,
    0 /*** final entry must be 0 ***/
};

```

NOTE → The IEEE 1364 standard does not define where the *vlog_startup_routines* array should be located. The only requirement is that a simulator must provide the array in a place that the simulator user can access and edit. Consult the reference manual of the simulator for the location of the start-up array used by that simulator.



TIP *Do not place the vlog_startup_routines array in the same file as the PLI application!* The `vlog_startup_routines` is a global array, and the C programming language does not permit multiple global arrays with the same name. In a typical design environment, PLI applications will come from several sources, such as internally developed applications, 3rd party applications, and ASIC vendor applications. If the `vlog_startup_routines` array and a PLI application are in the same file, then the source code for the application must be available whenever another PLI application needs to be added to the start-up array. If two PLI applications were to both include the `vlog_startup_routines` array in the application, then the object files for both applications could not be used together.

Limitations on functions listed in `vlog_startup_routines`

The functions listed in the `vlog_startup_routines` array are executed automatically by the simulator before simulation time 0, before the creation of the simulation data structure is complete. This limits what operations can be performed within functions that are executed from the `vlog_startup_routines`. The VPI routines which can be used in a function that is listed in `vlog_startup_routines` are:

- `vpi_register_systf()`
- `vpi_register_cb()` (this routine has limitations when used in a start-up callback, which are presented in Chapter 6, section 6.4.3 on page 198).
- VPI routines which do not access objects in simulation, such as `vpi_printf()`. It is an error to attempt obtain a handle for an object or any object properties.

2.12 Using the VPI user_data field

When a system task or system function is registered, a `user_data` value can be specified. The `user_data` is a character pointer, which can store a 32-bit value, or to point to a character string or to any other block of data. This `user_data` value is passed to the *calltf routine*, *compiletf routine* and *sizetf routine* as a C function input each time the simulator calls one of these routines.

In the following example, two different system tasks, `$read_test_vector_bin` and `$read_test_vector_hex`, are associated with the same *calltf routine*. The registration function for these system tasks is:

```

void ReadVector_register()
{
    s_vpi_systf_data  tf_data_bin, tf_data_hex;

    tf_data_bin.type      = vpiSysTask;
    tf_data_bin.tfname   = "$read_test_vector_bin";
    tf_data_bin.calltf   = ReadVectorCalltf;
    tf_data_bin.compiletf = NULL;
    tf_data_bin.sizetf    = NULL;
    tf_data_bin.user_data = "1";
    vpi_register_systf(&tf_data_bin);

    tf_data_hex.type      = vpiSysTask;
    tf_data_hex.tfname   = "$read_test_vector_hex";
    tf_data_hex.calltf   = ReadVectorCalltf;
    tf_data_hex.compiletf = NULL;
    tf_data_hex.sizetf    = NULL;
    tf_data_hex.user_data = "2";
    vpi_register_systf(&tf_data_hex);
}

```

Both system tasks invoke the same *calltf routine*, but the *user_data* value for the two system task names is different. Therefore, the *calltf routine* can check the *user_data* value to determine which system task name was used to call the routine.

The *user_data* value is passed to the callback routine as a C function input of type *char **. The following example illustrates reading the *user_data* value for the two system tasks registered in the previous example:

```

void ReadVectorCalltf(char *user_data)
{
    if (strcmp(user_data, "1") == 0)
        /* read test vectors as binary values */
    else if (strcmp(user_data, "2") == 0)
        /* read test vectors as hex values */
}

```

NOTE → Functions listed in the *vlog_startup_routines* array are not instance specific. The registration functions are only executed once per simulation. Therefore, the *user_data* value listed in the register routine will be common to all instances of the registered system task. If, for example, the *\$read_test_vector_hex* system task were used in two places in the Verilog design, each task will be passed the same *user_data* value. Chapter 6, section 6.2 on page 184 discusses how to allocate data that is unique to each instance of a system task/function.

2.13 Compiling and linking PLI applications

After the C source files for the PLI applications have been defined, they must be compiled and linked into a Verilog simulator. This allows the simulator to call the appropriate PLI routine when the PLI application system task or system function is encountered by the simulator.

The compiling and linking process is not part of the IEEE standard for the PLI. This process is defined by the simulator vendor, and is specific to both the simulator and the operating system on which the application is compiled. Appendix A presents the instructions for compiling and linking PLI applications for several of the Verilog simulators that were available at the time this book was written.

2.14 Summary

This chapter has shown the major parts of PLI applications. For PLI applications which are developed using the VPI library, the main parts of an application are: the name of a system task or system function, a *completf routine*, a *calltf routine*, a *sizetf routine*, and *simulation callback routines*. A system function called `$pow`, which calculates x^y and returns a 32-bit result, was used in this chapter to illustrate the major parts of a PLI application. The VPI interface mechanism is used to interface a PLI application to a Verilog simulator. The interface mechanism involves registering a PLI application using an application-defined register function, and notifying the simulator about the register function by editing a *vlog_startup_routines* array provided by the simulator. The next four chapters present how to use each of the routines in the VPI library.

CHAPTER 3

How to Use the VPI Routines

This chapter introduces the VPI portion of the PLI standard, and shows how to use the VPI routines to access information within a simulation data structure. Two complete PLI applications, `$show_all_nets` and `$show_all_signals`, will be created in this chapter to illustrate how the VPI routines work. The remaining chapters in this part of the book then build on the principles presented in this chapter by explaining the VPI library in much more detail.

The concepts presented in this chapter are:

- An overview of how VPI routines work
- Advantages of the VPI library
- Creating a complete PLI application using the VPI library
- Obtaining handles to Verilog HDL objects
- Accessing properties of Verilog HDL objects
- Reading values of Verilog HDL objects

3.1 Specification of `$show_all_nets` and `$show_all_signals`

To show how the VPI routines are used, two PLI applications will be created. The first example presented is an application called `$show_all_nets`. The usage of this application is:

```
$show_all_nets(<module_instance_name>) ;
```

This PLI application will:

1. Access the first argument of the system task, which is the name of a module instance.
2. Print the hierarchical path and name of that module.
3. Print the current simulation time.
4. Search for all net signals in the module.
5. Print the current logic value of each net.

This chapter will first illustrate a *completetf routine* for `$show_all_nets`, which verifies that the argument provided as an input is a valid module instance name. Then a *calltf routine* will be created to perform the functionality of the system task.

The second example is a PLI application called `$show_all_signals`. This application prints the current value of all net, reg and variable data types in a module. The usage of this application is:

```
$show_all_signals(<module_instance_name>);
```

To illustrate some additional ways to use the VPI routines, two enhancements to the `$show_all_signals` example will be presented. These are:

- Use no argument or a null argument to `$show_all_signals` to represent the module instance containing the `$show_all_signals` system task.
- Allow multiple arguments to `$show_all_signals`, so the values of signals in several modules can be printed with one call to `$show_all_signals`.

3.2 The VPI routine library

“VPI” stands for “Verilog Procedural Interface”. The VPI routines are the third of the three generations of the Verilog PLI routines (the TF routines were the first generation, and the ACC routines were the second). The primary purpose of the VPI routines is to provide a PLI application access to the internal data structures of a simulation. The VPI interface provides a consistent layer between a user’s PLI application and the underlying data structures of a simulation. By using this procedural interface layer, the PLI application does not need to know the specifics about how the simulator stores its data, and the same application will work with many different simulators.

The VPI routines treat Verilog HDL constructs as *objects*, and several of the VPI routines provide ways to locate specific objects within a simulation data structure. Other VPI routines can then read and modify information about the object.

The VPI library can be divided into five basic groups of routines:

- A *handle* routine obtains a handle for one specific Verilog HDL object.
- An *iterate* routine and a *scan* routine obtain handles for all of a specific type of Verilog object.
- *get* routines access information about an object.
- *set* routines modify information about an object.
- A few *miscellaneous* routines perform a variety of operations.

The library of VPI routines is defined in a C header file called **vpi_user.h**, which is part of the IEEE 1364 standard. This header file also defines a number of C constants and C structures used by the VPI routines. All PLI applications which use VPI routines must include the vpi_user.h file.

The VPI library is designed to work with the standard ANSI C libraries, such as **stdlib.h** and **stdio.h**. An example of including the header files for these libraries is:

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard I/O library */
#include "vpi_user.h"      /* IEEE 1364 PLI VPI library */
```

3.3 Advantages of the VPI library

The VPI routines provide direct access to virtually everything within a Verilog simulation data structure. This direct and all-inclusive access allows a PLI application to fully analyze and interact with a Verilog simulation in any way desired.

The VPI library is a concise set of 27 routines. (The proposed IEEE 1364-1999 Verilog standard will add a few additional routines). These routines have a very simple and consistent syntax, and are easy to learn and use. These 27 VPI routines were designed to completely replace the older TF and ACC generations of the PLI standard, which together total more than 200 routines. The older PLI generations had evolved over many years without the guidance of a standards body such as the IEEE. The TF and ACC libraries are bloated, full of redundancies, and very inconsistent in their syntax and semantics. The concise and consistent VPI library is much easier to work with.

As much as possible, the VPI library has been designed that so it can be very efficiently implemented by simulators. Potentially, this can improve the run-time performance of PLI applications compared to applications written with the older ACC generation of the standard. It should be noted, however, that the TF library only provides very limited access to the internals of a simulation data structure. Some simulators, such as the Synopsys VCS simulator, take advantage of this limited access to highly optimize their data structures when only the TF library is used. This restricted access and additional optimization may provide the best run-time performance of all the PLI libraries, but the restricted access also limits the capabilities of a PLI application.



The VPI library was designed to replace the TF and ACC libraries with a more concise, more robust, and more versatile procedural interface. The IEEE 1364 standard includes the TF and ACC libraries, in order to provide backward compatibility for older PLI applications. The official policy of the IEEE 1364 standards committee is that, as enhancements are added to the Verilog language, only the VPI library will be expanded to support the new features. The TF and ACC libraries will not be enhanced in future versions of the IEEE 1364 standard.

Some disadvantages of VPI routines

At the time this book was written, the TF and ACC libraries were supported by virtually every major Verilog simulator, but the VPI library was only supported by a few simulators. The more widespread support of TF and ACC routines makes a PLI application portable to many more simulators and engineering environments. This disadvantage of the VPI routines will disappear as more simulators companies implement the full IEEE PLI standard.

Many new users of the PLI find that it is sometimes much easier and faster to develop PLI applications using the TF and ACC libraries as compared to the VPI library. The TF and ACC libraries are much larger, and often have built-in routines which take care of a lot of the work that a PLI application needs to accomplish. In the smaller VPI library, the PLI application developer must code much of the corresponding functionality by hand. However, when using the larger TF and ACC libraries, PLI application developers may tend to implement poorly structured C code, which can impact simulation performance and becomes difficult to maintain.

3.4 Verilog HDL objects

The VPI routines treat Verilog HDL constructs as *objects*, and several of the VPI routines provide ways to locate any specific object or type of object within a simulation data structure. Other VPI routines can then read and modify information about each object. The simple Verilog HDL example which follows has several objects which can be accessed by the library of VPI routines.

```
module test;
    reg [1:0] test_in;
    wire [1:0] test_out;
    buf2 u1 (test_in, test_out);
    initial
        begin
            test_in = 3;
            #50 $display("in=%d, out=%d", test_in, test_out);
        end
    endmodule

module buf2 (in, out);
    input [1:0] in;
    output [1:0] out;
    wire [1:0] in, out;
    buf #5 n0 (out[0], in[0]);
    buf #7 n1 (out[1], in[1]);
endmodule
```

In this Verilog HDL example, the objects that a PLI application can access include:

- A top-level module, with the definition name “*test*”. Within this module are:
 - A *reg* signal, with a vector size of 2 and the name “*test_in*”. The signal will have a logic value which can be read and modified by the PLI application.
 - A *wire* net, with a vector size of 2 and the name “*test_out*”. The net reflects a resolved logic value which can be read by the PLI application.
- A module instance, with the definition name “*buf2*” and the instance name “*u1*”. Within this module are:
 - Two ports, with the names “*in*” and “*out*”. Each port has a vector size and direction.
 - Two *wire* nets, with vector sizes and names. The nets reflect a resolved logic value which can be read by the PLI application.
 - Two primitive instances, with the definition name “*buf*” and the instance names “*n0*” and “*n1*”. Each primitive has a delay value which can be read and modified by the PLI.

- Terminals on each primitive instance, with “*out[0]*” (a bit-select of “*out*”) and “*in[0]*” connected to one instance, and “*out[1]*” and “*in[1]*” connected to the other instance.
- An initial procedure. Within this procedure are:
 - A begin—end statement group.
 - An assignment statement, with expressions on the right-hand side and left-hand side of the assignment.
 - A time control (‘#’).
 - An instance of a system task;

In addition to the Verilog HDL objects, there are simulation objects which do not exist in the Verilog language, but which can be accessed by the PLI. Examples of simulation objects are the simulation event queue and propagation delays between two modules (referred to as an inter-connect delay).

3.4.1 The vpiHandle data type

The VPI routines use a special data type, called a **handle**, to access Verilog HDL and simulation objects. The declaration type for variables to store a handle is **vpiHandle**. The **vpiHandle** data type is defined in the VPI library (the **vpi_user.h** file). An example declaration for two handle variables is:

```
vpiHandle primitive_handle, net_handle;
```

There are several VPI routines that locate an object within a simulation data structure and return the handle for the object. Other VPI routines are used to access information about the object, using the object’s handle as a reference point. The information that can be accessed depends on the type of the object, but might include the object’s name and current logic value.

NOTE → ***Do not share handles between VPI routines and ACC routines!*** The ACC routines in the PLI standard also use the concept of a handle for referencing Verilog objects. The IEEE 1364 standard does not guarantee that a handle which is obtained with the ACC library will be the same as a handle which is obtained with the VPI library.

3.4.2 Object relationships

Each object in a Verilog design is related to other objects. In the preceding Verilog HDL example, module “*buf2*” contains instance “*n0*” of the buffer primitive. This is a parent–child relationship, where “*buf2*” is the parent which contains the child “*n0*”. Similarly, instance “*n0*” of the buffer has two terminals, which is another parent–child relationship.


The VPI standard documents three types of object relationships:

- **One-to-one relationships** occur when an object is related to only one other object of a certain type. In the preceding Verilog HDL example, the buffer instance “*n0*” has a single parent, which is module “*buf2*”, so there is a one-to-one relationship from this child to its parent.
- **One-to-many relationships** occur when an object is related to several other objects of a certain type. In the preceding Verilog HDL example, the buffer instance “*n0*” has multiple terminals, so there is a one-to-many relationship from this parent to its children.
- **Many-to-one relationships** occur when many objects are related to a single other object of a certain type. For example, an input port of a module might be driven by several sources, so there are many drivers for the one port.

All of the possible relationships that an object might have are documented in the VPI standard. This documentation is made in the form of **object diagrams**, which use enclosures to represent objects and arrows to represent the type of relationship. The full details of these data diagrams is not presented in this chapter, which is oriented towards getting started with using the VPI library. These details are presented in the next chapter.

Figure 3-1 serves as quick example of how the VPI diagrams document Verilog object relationships.

Figure 3-1: VPI object relationship for a module to a net



In this example:

- The double arrow from a module to a net indicates a *one-to-many relationship*. A

module can contain any number of nets.

- The single arrow from a net to a module indicates a *many-to-one relationship*. A specific net can only occur within one module.
- The names **vpiModule** and **vpiNet** are referred to as *tags*. These tags are constants which represent the type of object in the relationship. The constants are defined in the VPI library, and the VPI routines will use these tag constants to traverse from one object to another, such as from a child object to its parent.

3.5 Obtaining object handles

The usage of the `$show_all_nets` PLI application to be created in this chapter is:

```
$show_all_nets(<module_instance_name>);
```

`$show_all_nets` PLI will need to obtain a handle for the first system task/function argument, in order to access all signals within that module. The application will then need to obtain handles for all the signals in that module. The VPI library provides special routines that return handles for objects.

There are two important terms used with the VPI routines that obtain object handles:

- *target objects* are the type of objects for which the VPI routine will obtain handles.
- A *reference object* is where the VPI routine will search for the target objects. For example, to find all nets within a module, the reference object is the module.

3.5.1 Obtaining a handle for a one-to-one relationship

vpiHandle() obtains a handle for a target object with a one-to-one relationship from the reference object. The syntax for this routine is:

vpiHandle vpi_handle (type, reference)

int **type** constant representing an object type.

vpiHandle **reference** handle for an object.

As an example, to obtain the parent module that contains a net, the type constant would be **vpiModule**, and the reference object would be a handle for the net. The C code that would obtain the module handle is:

```
vpiHandle module_handle, net_handle;  
/* assume a net handle has already been obtained */  
module_handle = vpi_handle(vpiModule, net_handle);
```

The PLI application for `$show_all_nets` will need to read the module instance name listed as the argument of the system task. In order to access the argument, the PLI application must first obtain a handle for the instance of the system task which called the application. The object diagram for a system task call shows the following one-to-one relationship:

Figure 3-2: VPI object relationship for an instance of a system task call



Note: the complete object diagram for a system task call is listed in Appendix B, on page 697.

The small circle in an object diagram indicates that a `NULL` is to be used as the reference object. Therefore, the PLI applications can obtain a handle for the system task call using the following C code.

```
vpiHandle systf_handle;  
systf_handle = vpi_handle(vpiSysTfCall, NULL);
```

NOTE → The `NULL` (all upper case letters) that is used by VPI routines is defined in the `stdlib.h` ANSI C library file.

3.5.2 Obtaining a handle for a one-to-many relationship

The routines `vpi_iterate()` and `vpi_scan()` are used to obtain the handles of all objects when there is a one-to-many relationship. `vpi_iterate()` obtains an `iterator` for all target objects in a one-to-many relationship with a reference object. An iterator is essentially a pointer to the next object in a series of objects. The variable in which to store an iterator object is declared as a `vpiHandle` data type. `vpi_scan()` obtains a handle for the next target object that is referenced by an iterator.

vpiHandle vpi_iterate(type, reference)

int tag_constant constant representing an object type.

vpiHandle reference handle for an object.

The `vpi_iterate()` routine returns an *iterator object*. The routine uses the `tag_constant` to determine the type of object for which handles are to be obtained, and the `reference_handle` to determine the source point of the one-to-many relationship.

The *iterator object* returned from `vpi_iterate()` represents the first of the many target objects in the relationship. As each target object is accessed, the iterator is updated automatically to reference the next target object. Conceptually, the many target objects can be thought of as a list of object handles, and the iterator as a pointer to the next object in the list. Note that the usage of a list is purely conceptual—the VPI standard does not require that a simulator create and store lists of target objects. Since the VPI uses the more abstract iterator object to reference each target object, a simulator can maintain its internal storage in any form.

A one-to-many relationship indicates a relationship from a reference object to any number of target objects, including none. For example, a module might not have any nets declared within it, a single net declared, or multiple nets. If there are no target objects in a one-to-many relationship, then `vpi_iterate()` returns a **NULL** as the iterator value.

*The PLI application must check for a **NULL** iterator before attempting to access an object referenced by the iterator.*

vpiHandle vpi_scan(iterator)

vpiHandle iterator handle for an iterator object.

The `vpi_scan()` routine is provided a single input, the iterator object which was returned from `vpi_iterate()`. The `vpi_scan()` routine returns the handle for the target object which the iterator references. After each call to `vpi_scan()`, the iterator object is updated to reference the next target object. When there are no more target objects, the next call to `vpi_scan()` will return **NULL**. In order to access all of the objects in a one-to-many relationship, `vpi_scan()` must be called multiple times, until the return value is **NULL**.

The following C code fragment uses `vpi_iterate()` and `vpi_scan()` to obtain handles for all nets within a module:

```

vpiHandle module_handle, net_iterator, net_handle;
/* assume a module handle has already been obtained */
net_iterator = vpi_iterate(vpiNet, module_handle);
if (net_iterator == NULL)
    vpi_printf("  No nets found in this module\n");
else {
    while ( (net_handle = vpi_scan(net_iterator)) != NULL ) {
        /* access information about the net object */
    }
}

```


For `$show_all_nets`, the PLI application will need to access the first argument of the system task. This first argument will be a module instance name. For example:

```
$show_all_nets(top.i1);
```

Accessing the arguments of a system task/function

A system task can have any number of arguments. As shown in the object diagram in Figure 3-3, there is a one-to-many relationship, indicated by the double arrow, between a system task/function call and the arguments of the system task/function.

Figure 3-3: VPI object relationship for system task arguments



Note: the complete object diagram for a system task call is listed in Appendix B, on page 697.

The `$show_all_nets` application can obtain a handle for the module instance that is named in the first system task argument using `vpi_iterate()` and `vpi_scan()`, as shown in the following C code fragment:

```

vpiHandle systf_handle, module_handle,
           net_iterator, net_handle;

systf_handle = vpi_handle(vpiSysTfCall, NULL);
arg_iterator = vpi_iterate(vpiArgument, systf_handle);
module_handle = vpi_scan(arg_iterator);
vpi_free_object(arg_iterator); /* free iterator memory */

```

This memory required to store the iterator object automatically allocated by the simulator when `vpi_iterate()` is called, and is automatically freed when `vpi_scan()` returns `NULL` after all target objects have been accessed.

In the preceding example, however, `vpi_scan()` was only called one time, which means the memory for the iterator object may not have been automatically freed. Since the `$show_all_nets` application has obtained a handle for the object needed (the first system task argument), the application no longer needs the iterator object and therefore releases the memory for the iterator. The VPI library provides a special routine to release the iterator object memory when the iterator is no longer needed.

int vpi_free_object(object)
vpiHandle **object** handle for an object.

The `vpi_free_object()` routine is used to release the memory which the simulator allocated in order to store the iterator object. Refer to section 4.5.3 on page 108 in Chapter 4 for more details on when to use `vpi_free_object()`.

3.6 Printing messages from VPI applications

The VPI library uses a special routine for printing text messages.

int vpi_printf(format, arg1,...argn)
char * **format** quoted character string of formatted message.
arg1...argn arguments to formatted message string.

`vpi_printf()` is used to print messages from PLI applications. The syntax for this routine is essentially the same as the C `printf()` routine, but there is one very important difference: `vpi_printf()` will print the message to the simulation output window, and to the simulation output log file. By using `vpi_printf()`, a PLI application can generate messages which are part of the simulation, without having to determine where the simulator directs its output. `vpi_printf()` returns the number of characters printed, or `EOF` if an error occurred.

NOTE → The C `printf()` routine will print its message to the operating system's standard output channel. If the simulator uses some other output channel, such as a graphical window, then printing a message using the C `printf()` routine might not be seen by the simulator user. Nor will the message be recorded in any output files generated by the simulator.

3.7 Accessing object properties

Every Verilog object has one or more properties which can be accessed by a PLI application. Some properties are the name of a module or net and the logic value of a net. Most properties will be either an integer value or a string value. The VPI identifies these properties using a property constant.

Two VPI routines are provided to read these types of properties.

int vpi_get (property, object)

int **property** constant representing an object property.
vpiHandle **object** handle for an object.

Returns the value associated with integer and boolean properties of an object. Boolean properties return 1 for true and 0 for false.

char *vpi_get_str (property, object)

int **property** constant representing an object property.
vpiHandle **object** handle for an object.

Returns a pointer to a string containing the value associated with string properties of an object.

3.7.1 Object type properties

Every Verilog object has an integer *type* property, which is accessed using:

```
int obj_type;  
obj_type = vpi_get (vpiType, <object_handle>)
```

The type property identifies what kind of Verilog object is referenced by a VPI handle. This type property is represented by an integer constant, which is defined in the vpi_user.h file. Some example type constants are:

- **vpiModule** — the object handle is referencing a Verilog module instance
- **vpiNet** — the object handle is referencing a Verilog net data type
- **vpiReg** — the object handle is referencing a Verilog reg data type
- **vpiPrimitive** — the object handle is referencing a Verilog primitive instance

The type property can be used many different ways. One common usage is to verify that a handle which was obtained references the type of object expected. For example,

the `$show_all_nets` application requires the first task/function argument be a module instance. The following code fragment uses this type property to verify that the argument provided to `$show_all_nets` is correct:

```
vpiHandle systf_handle, arg_iterator, arg_handle;  
int tfarg_type;  
  
systf_handle = vpi_handle(vpiSysTfCall, NULL);  
arg_iterator = vpi_iterate(vpiArgument, systf_handle);  
arg_handle = vpi_scan(arg_iterator);  
tfarg_type = vpi_get(vpiType, arg_handle);  
if (tfarg_type != vpiModule) {  
    /* report error that argument is not correct */
```

NOTE In the IEEE 1364-1995 standard, the `vpiType` property is defined to be an integer property. Therefore, the call to `vpi_get(vpiType, tfarg_h)` returns the integer value of an object's type constant. These constant values are defined in the `vpi_user.h` header file. In the proposed 1364-1999 standard, `vpiType` will be defined to be both an integer and a string property. The `vpi_get()` routine will return the integer value of the type constant, and `vpi_get_str()` will return the name of the type constant.

3.7.2 Object name properties

Many Verilog objects have one or more *name* properties.

- The property represented by the property constant `vpiName` is the local name of an object. For objects such as nets and variables, the local name is the *declaration name* of the object. For a module or primitive, the local name is the *instance name* within the module in which the module or primitive is used.
- The property represented by the property constant `vpiFullName` is the full hierarchical path name of an object.
- The property represented by the property constant `vpiDefName` is the definition name of a module or primitive.

The following Verilog HDL source code fragment illustrates the difference between the `vpiName`, `vpiFullName` and `vpiDefName` properties.

```

module test;
  wire a, b, ci, sum, co;
  addbit u1 (a, b, ci, sum, co);
endmodule

module addbit (a, b, ci, sum, co);
  input a, b, ci;
  output sum, co;

  wire a, b, ci, sum, co;
  xor g1 (n1, a, b);
  xor #2 g2 (sum, n1, ci);
  and g3 (n2, a, b);
  and g4 (n3, n1, ci);
  or #2 g5 (co, n2, n3);
endmodule

```

local name: "u1"
full name: "test.u1"
definition name: "addbit"

local name: "sum"
full name: "test.u1.sum"

local name: "g1"
full name: "test.u1.g1"
definition name: "xor"

3.8 Reading the logic values of Verilog objects

The VPI library provides a routine to read the logic value of any Verilog object which can contain a value, such as a net or variable.

void vpi_get_value (expr, value)

vpiHandle expr handle for an object.

p_vpi_value value pointer to an application-allocated *s_vpi_value* structure to receive value information.

The Verilog language uses 4-state logic values, comprising logic 0, 1, Z and X. The *vpi_get_value()* routine automatically converts Verilog 4-state logic into various C data types for use in PLI applications. A simple way to represent 4-state logic in C is to use character strings, and this is the method used in the *\$show_all_nets* application. Chapter 5 presents reading and writing Verilog logic values in more detail.



TIP

Using C strings to represent 4-state logic is a simple method for reading and printing a Verilog logic value. However, the automatic conversion from Verilog values to C strings can be very expensive for the run-time performance of a PLI application. If a PLI application will access a large number of values, or if the application will be called many times during a simulation, it is better to use a more efficient format for reading values. Chapter 5 presents all the formats for reading logic values that are available using VPI routines, and discusses performance considerations.

The PLI application must allocate an `s_vpi_value` structure prior to calling `vpi_get_value()`. The definition of the structure is contained in the `vpi_user.h` header file. The PLI application does not define the structure. The application only allocates memory for the structure. The structure definition is:

```
typedef struct t_vpi_value {
    int format;
    union {
        char *str;                      /* string value */
        int scalar;                     /* scalar value */
        double real;                    /* real value */
        struct t_vpi_time   *time;      /* time value */
        struct t_vpi_vecval *vector;    /* vector value */
        struct t_vpi_strengthval *strength; /* strength val */
        char *misc;                     /* reserved */
    } value;
} s_vpi_value, *p_vpi_value;
```

The `s_vpi_value` structure contains two primary fields:

- The `format` field controls how the Verilog logic value should be represented in C. The format is a VPI constant. For example, a format of `vpiBinStrVal` indicates the logic value should be represented using binary numbers with the characters ('0', '1', 'z', and 'x'). A `vpiHexStrVal` format indicates the value should be represented using hexadecimal numbers with the characters ('0' through 'F', 'z', and 'x'). Several other formats are available, which are discussed in Chapter 5.
- The `value` field receives the logic value. This field is a union of C data types, and the format constant determines which field within this union that will be used. For formats which receive the logic value as a string, the `value.str` field will contain a pointer to the string.

The following example retrieves the logic value of a net as a C string:

```
vpiHandle  net_handle;
s_vpi_value current_value;

current_value.format = vpiBinStrVal; /* read as a string */
vpi_get_value(net_handle, &current_value);
vpi_printf(" net %s value is %s (binary)\n",
           vpi_get_str(vpiName, net_handle),
           current_value.value.str);
```

 **NOTE** `vpi_get_value()` automatically allocates storage for the string which contains the logic value. This storage is temporary, and will automatically be freed when another call is made to `vpi_get_value()` or when the PLI application exits.

3.9 Reading the current simulation time

In addition to printing the logic values of nets, the `$show_all_nets` application will print the current simulation time for when the application is called.

The Verilog language uses the ‘timescale compiler directive to establish the time units of a Verilog module, and different modules can represent time in different units. Verilog simulators use an internal simulation time unit, and scale the delays of a module to the internal time unit. The VPI library allows the current simulation time to be represented in either the internal simulation time units or in the time units of any module in the design.

void vpi_get_time (object, time)

vpiHandle object handle for an object.

p_vpi_time time pointer to an application-allocated `s_vpi_time` structure to receive time information.

The `<object_handle>` is a handle for any object in a design. When simulation time is retrieved in the time scale of a module, the module that is used will be the one containing the object.


The `<time_structure_pointer>` is a pointer to an `s_vpi_time` structure to receive the current simulation time.

The PLI application must allocate an `s_vpi_time` structure prior to calling `vpi_get_time()`. The definition of the structure is contained in the `vpi_user.h` header file. The PLI application does not define the structure. The application only allocates memory for the structure. The structure definition is:

```
typedef struct t_vpi_time {
    int type;          /* vpiScaledRealTime, vpiSimTime */
    unsigned int high; /* when using vpiSimTime */
    unsigned int low;  /* when using vpiSimTime */
    double real;       /* when using vpiScaledRealTime */
} s_vpi_time, *p_vpi_time;
```

- The `type` field controls how the Verilog simulation time will be received. The format is set using a constant which is defined in the VPI library.
 - A format of `vpiScaledRealTime` indicates the simulation time will be retrieved as a floating point number and that the time will be scaled to the time units of the module containing the object specified in the call to `vpi_get_time()`. The `real` field receives the simulation time value when the format field is `vpiScaledRealTime`.

- A format of **vpiSimTime** indicates the simulation time should be retrieved as a 64-bit integer and that the time will be in the internal simulation time units. Since the ANSI C language does not define a 64-bit integer data type, the 64-bit value is split into two 32-bit C integers, as shown in the following diagram:



The **\$show_all_nets** application retrieves and prints the current simulation time using **vpi_get_time()** as follows:

```

vpiHandle    module_handle;
s_vpi_time  current_time;

current_time.type = vpiScaledRealTime;
vpi_get_time(module_handle, &current_time);
vpi_printf("\nAt time %2.2f, nets in module %s (%s):\n",
           current_time.real,
           vpi_get_str(vpiFullName, module_handle),
           vpi_get_str(vpiDefName, module_handle));

```

3.10 Controlling simulation from PLI applications

There are occasions when a PLI application needs to control what a Verilog simulator is doing. In the **\$show_all_nets** application, a *compileif routine* will be provided to perform syntax checking. If a serious error is detected, such as the argument provided to **\$show_all_nets** is not a module instance name, then the PLI application needs to abort simulation execution. That is, to treat the error as a fatal error.

The IEEE 1364-1995 standard does not provide the ability for VPI routines to control simulation. These capabilities only exist in the TF library of PLI routines. To control simulation from a PLI application, it is necessary to include the TF routine library, and use the TF simulation control routines such as **tf_error()** and **tf_dofinish()**. Refer to Chapter 10, section 10.4 for a discussion of these routines.

NOTE The proposed IEEE 1364-1999 standard adds new VPI routines to provide simulator control capabilities similar to the TF routines.

3.11 A complete PLI application using VPI routines

Example 3-1 lists a complete PLI application for the `$show_all_nets` application. This application includes three C functions:

- A *registration function* to register the PLI application.
- A *compileif routine* to verify the argument provided to `$show_all_nets`.
- A *callif routine* to access all nets in a module and print the name and current logic value of these nets.

The definition and purpose of the *registration function*, *compileif routine* and *callif routine* of these PLI routines were presented in Chapter 2.

A short Verilog HDL test case for `$show_all_nets` applications and the simulation results follow the listing of the PLI application C code.

3.11.1 PLI application source code for `$show_all_nets`

CD The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.03/show_all_nets_vpi.c
- Verilog test bench: Chapter.03/show_all_nets_test.v
- Verilog-XL results log: Chapter.03/show_all_nets_test.log

Example 3-1: `$show_all_nets` — using VPI routines to access simulation objects

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include "vpi_user.h"     /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"      /* IEEE 1364 PLI TF routine library */
/* using TF routines for simulation control */

/* prototypes of the PLI application routines */
int PLIbook_ShowNets_compileif(), PLIbook_ShowNets_calltf();
```

```
*****
 * $show_all_nets Registration Data
 * (add this function name to the vlog_startup_routines array)
*****
void PLIbook_ShowNets_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname   = "$show_all_nets";
    tf_data.calltf   = PLIbook_ShowNets_calltf;
    tf_data.compiletf = PLIbook_ShowNets_compiletf;
    tf_data.sizetf   = NULL;
    vpi_register_systf(&tf_data);
    return;
}

*****
 * compiletf routine
*****
int PLIbook_ShowNets_compiletf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, arg_handle;
    int        tfarg_type;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handles to system task arguments */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    if (arg_iterator == NULL) {
        vpi_printf("ERROR: $show_all_nets requires 1 argument\n");
        tf_dofinish(); /* abort simulation */
        return(0);
    }

    /* check the type of object in system task arguments */
    arg_handle = vpi_scan(arg_iterator);
    tfarg_type = vpi_get(vpiType, arg_handle);
    if (tfarg_type != vpiModule) {
        vpi_printf("ERROR: $show_all_nets arg1 must be module instance\n");
        vpi_free_object(arg_iterator); /* free iterator memory */
        tf_dofinish(); /* abort simulation */
        return(0);
    }

    /* check that there is only 1 system task argument */
    arg_handle = vpi_scan(arg_iterator);
    if (arg_handle != NULL) {
        vpi_printf("ERROR: $show_all_nets can only have 1 argument\n");
        vpi_free_object(arg_iterator); /* free iterator memory */
        tf_dofinish(); /* abort simulation */
        return(0);
    }
}
```

```
    }
    return(0);
}

/********************* calltf routine *********************/
int PLIbook_ShowNets_calltf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, module_handle,
               net_iterator, net_handle;
    s_vpi_time current_time;
    s_vpi_value current_value;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument */
    /* compiletf has already verified only 1 arg with correct type */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    module_handle = vpi_scan(arg_iterator);
    vpi_free_object(arg_iterator); /* free iterator memory */

    /* read current simulation time */
    current_time.type = vpiScaledRealTime;
    vpi_get_time(systf_handle, &current_time);

    vpi_printf("\nAt time %2.2f, nets in module %s (%s):\n",
              current_time.real,
              vpi_get_str(vpiFullName, module_handle),
              vpi_get_str(vpiDefName, module_handle));

    /* obtain handles to nets in module and read current value */
    net_iterator = vpi_iterate(vpiNet, module_handle);
    if (net_iterator == NULL)
        vpi_printf(" no nets found in this module\n");
    else {
        current_value.format = vpiBinStrVal; /* read values as a string */
        while ( (net_handle = vpi_scan(net_iterator)) != NULL ) {
            vpi_get_value(net_handle, &current_value);
            vpi_printf(" net %-10s value is %s (binary)\n",
                      vpi_get_str(vpiName, net_handle),
                      current_value.value.str);
        }
    }
    return(0);
}
```

3.11.2 A test bench for \$show_all_nets

Example 3-2 lists a simple Verilog HDL design to test *\$show_all_nets*. Figure 3-4, which follows, shows the output of running a simulation with this test bench.

Example 3-2: *\$show_all_nets* — a Verilog HDL test using a PLI application

```
'timescale 1ns / 1ns
module top;
    reg [2:0] test;
    tri [1:0] results;

    addbit i1 (test[0], test[1], test[2], results[0], results[1]);

    initial
        begin
            test = 3'b000;
            #1 test = 3'b011;

            #1 $show_all_nets(top);
            #1 $show_all_nets(i1);

            #1 $stop;
            #1 $finish;
        end
    endmodule

/** A gate level 1 bit adder model ***/
'timescale 1ns / 1ns
module addbit (a, b, ci, sum, co);
    input a, b, ci;
    output sum, co;

    wire a, b, ci, sum, co,
          n1, n2, n3;

    xor (n1, a, b);
    xor #2 (sum, n1, ci);
    and (n2, a, b);
    and (n3, n1, ci);
    or #2 (co, n2, n3);

endmodule
```

3.11.3 Simulation results for \$show_all_nets

Figure 3-4: *\$show_all_nets* — simulation results using a PLI application

The screenshot shows the Cadence Verilog-XL Turbo NT simulation interface. The top menu bar includes File, Control, Set, Show, Select, Tools, Options, and Help. Below the menu is a toolbar with icons for play, stop, and other simulation controls. The main window displays a Verilog code snippet and its simulation results.

```

21 begin
22     test = 3'b000;
23     #1 test = 3'b011;
24
25     #1 $show_all_nets(top);
26     #1 $show_all_nets(ii);
27
28     #1 $stop;
29     #1 $finish;
30 end

```

Scope: top Subscopes:

At time 2.00, nets in module top (top):
net results value is 10 (binary)

At time 3.00, nets in module top.ii (addbit):
net a value is 1 (binary)
net b value is 1 (binary)
net ci value is 0 (binary)
net sum value is 0 (binary)
net co value is 1 (binary)
net n1 value is 0 (binary)
net n2 value is 1 (binary)
net n3 value is 0 (binary)

L28 "show_all_nets_test.v": \$stop at simulation time 4
Type ? for help
C1 >

Ready

3.12 Obtaining handles for reg and variable data types

The Verilog HDL defines two general data type groups, *nets* and *registers*. The register data type group includes the Verilog keywords **reg**, **integer**, **time** and **real**. In the PLI, the term *register* is not used to represent data types. Instead, the PLI treats the **reg** data type as a unique object, and groups the **integer**, **time** and **real** data types into an object class called *variables*.

Only minor changes are needed to enhance the *\$show_all_nets* application so that it can display all the signals of all data types within a module. All that is required is to add additional **vpi_iterate()** and **vpi_scan()** statements to access the other signal data types.

The `vpi_get_value()` is used to read the values of any Verilog data type. The `format` field in the `s_vpi_value` structure establishes the C language data type to be used to represent the value. This gives the PLI application developer complete control over how values are represented in the application.

For the `$show_all_signals` application illustrated in this chapter, the following formats will be used:


- For Verilog net data types, values will be represented as a C string, using a binary format.
- For Verilog `reg` data types, values will be represented as a C string, using a binary format.
- For Verilog `integer` data types, values will be represented as a C `int`.
- For Verilog `real` data types, values will be represented as a C `double`.
- For Verilog `time` data types, values will be represented as a pair of C `unsigned int` variables.

The conversion of the Verilog `time` data type into a pair of C `int` variables is required because, in the Verilog language, the `time` data type stores a 64-bit unsigned value. Within the VPI, a C structure is used to store the upper portion of the 64 bits in a 32-bit unsigned `int`, and the lower 32 bits in another unsigned `int`. These two C integers can then be concatenated together to display the full 64-bit time value.

The structure to store a time value is `s_vpi_time`, and is defined in the `vpi_user.h` VPI library file, as follows:

```
typedef struct t_vpi_time {  
    int type;  
    unsigned int high;  
    unsigned int low;  
    double real;  
} s_vpi_time, *p_vpi_time;
```

Note that this same structure definition is used by several VPI routines. In the context of reading a Verilog time variable, the `type` and `real` fields of this structure are not used and can be ignored. The value of the verilog time variable will be stored in the structure, as shown in the following illustration:



TIP The C language can be used to perform math operations on the high and low order words that make up the 64-bit time value, but it is important that carries and borrows between the words are handled properly. As a convenience, the TF library of the PLI standard contains several routines to perform math operations on the pair of C int variables used to represent a 64-bit Verilog time value. Refer to section 10.10 on page 323 in Chapter 10 for a list of these TF utility routines.

The following C code fragment illustrates the steps required to read the value of a Verilog time variable and print the value in hexadecimal.

```
vpiHandle signal_handle;
s_vpi_value current_value;

current_value.format = vpiTimeVal;
vpi_get_value(signal_handle, &current_value);
vpi_printf(" time %10s value is %x%x\n",
           vpi_get_str(vpiName, signal_handle),
           current_value.value.time->high,
           current_value.value.time->low);
```



NOTE `vpi_get_value()` automatically allocates an `s_vpi_time` structure when it is called using a `vpiTimeVal` format. The memory for this structure is temporary, and will automatically be freed when another call is made to `vpi_get_value()` or when the PLI application exits.

3.12.1 A complete PLI application for \$show_all_signals

Example 3-3 lists the complete C code for the `$show_all_signals` PLI application. This application obtains handles for all signals in a module, including the data types of `net`, `reg`, `integer`, `time` and `real`.



TIP Using structured programming techniques makes it easier to maintain or enhance the functionality of a PLI application. This example uses a more structured programming style, by creating a separate C function, called *PrintSignalValues()*, to print the current logic value. By moving the printing statements into a separate C function, the *calltf routine* is kept shorter and easier to read, plus there is less duplication of code.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.03/show_all_signals_1_vpi.c
- Verilog test bench: Chapter.03/show_all_signals_1_test.v
- Verilog-XL results log: Chapter.03/show_all_signals_1_test.log

Example 3-3: \$show_all_signals, version 1 — printing values of all Verilog types

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include "vpi_user.h"     /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"     /* IEEE 1364 PLI TF routine library */
                       /* using TF routines for simulation control */

/* prototypes of the PLI application routines */
int  PLIbook_ShowSignals_compiletf(), PLIbook_ShowSignals_calltf();
void PLIbook_PrintSignalValues();

/********************* Registration Data
 * (add this function name to the vlog_startup_routines array)
********************/
void PLIbook_ShowSignals_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$show_all_signals";
    tf_data.calltf   = PLIbook_ShowSignals_calltf;
    tf_data.compiletf = PLIbook_ShowSignals_compiletf;
    tf_data.sizetf    = NULL;
    vpi_register_systf(&tf_data);
    return;
}

/********************* compiletf routine
********************/
```

```
int PLIbook_ShowSignals_completf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, arg_handle;
    int        tfarg_type;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handles to system task arguments */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    if (arg_iterator == NULL) {
        vpi_printf("ERROR: $show_all_signals requires 1 argument\n");
        tf_dofinish(); /* abort simulation */
        return(0);
    }

    /* check the type of object in system task arguments */
    arg_handle = vpi_scan(arg_iterator);
    tfarg_type = vpi_get(vpiType, arg_handle);
    if (tfarg_type != vpiModule) {
        vpi_printf("ERROR: $show_all_signals arg 1");
        vpi_printf(" must be a module instance\n");
        vpi_free_object(arg_iterator); /* free iterator memory */
        tf_dofinish(); /* abort simulation */
        return(0);
    }

    /* check that there is only 1 system task argument */
    arg_handle = vpi_scan(arg_iterator);
    if (arg_handle != NULL) {
        vpi_printf("ERROR: $show_all_signals can only have 1 argument\n");
        vpi_free_object(arg_iterator); /* free iterator memory */
        tf_dofinish(); /* abort simulation */
        return(0);
    }
    return(0);
}

/*********************************************
 * calltf routine
 *****/
int PLIbook_ShowSignals_calltf(char *user_data)
{
    vpiHandle    systf_handle, arg_iterator, module_handle,
                 signal_iterator;
    int         format;
    s_vpi_time  current_time;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);
```

```
/* obtain handle to system task argument
 * compiletf has already verified only 1 arg with correct type */
arg_iterator = vpi_iterate(vpiArgument, systf_handle);
module_handle = vpi_scan(arg_iterator);
vpi_free_object(arg_iterator); /* free iterator memory */

/* read current simulation time */
current_time.type = vpiScaledRealTime;
vpi_get_time(systf_handle, &current_time);

vpi_printf("\nAt time %2.2f, signals in module %s (%s):\n",
           current_time.real,
           vpi_get_str(vpiFullName, module_handle),
           vpi_get_str(vpiDefName, module_handle));

/* obtain handles to nets in module and read current value */
signal_iterator = vpi_iterate(vpiNet, module_handle);
if (signal_iterator != NULL)
    PLIbook_PrintSignalValues(signal_iterator);

/* obtain handles to regs in module and read current value */
signal_iterator = vpi_iterate(vpiReg, module_handle);
if (signal_iterator != NULL)
    PLIbook_PrintSignalValues(signal_iterator);

/* obtain handles to variables in module and read current value */
signal_iterator = vpi_iterate(vpiVariables, module_handle);
if (signal_iterator != NULL)
    PLIbook_PrintSignalValues(signal_iterator);

vpi_printf("\n"); /* add some white space to output */
return(0);
}

void PLIbook_PrintSignalValues(vpiHandle signal_iterator)
{
    vpiHandle    signal_handle;
    int         signal_type;
    s_vpi_value current_value;

    while ( (signal_handle = vpi_scan(signal_iterator)) != NULL ) {
        signal_type = vpi_get(vpiType, signal_handle);
        switch (signal_type) {
            case vpiNet:
                current_value.format = vpiBinStrVal;
                vpi_get_value(signal_handle, &current_value);
                vpi_printf(" net      %-10s value is %s (binary)\n",
                           vpi_get_str(vpiName, signal_handle),
                           current_value.value.str);
                break;

            case vpiReg:
                current_value.format = vpiBinStrVal;
```

```
vpi_get_value(signal_handle, &current_value);
vpi_printf(" reg      %-10s value is %s (binary)\n",
           vpi_get_str(vpiName, signal_handle),
           current_value.value.str);
break;

case vpiIntegerVar:
    current_value.format = vpiIntVal;
    vpi_get_value(signal_handle, &current_value);
    vpi_printf(" integer %-10s value is %d (decimal)\n",
               vpi_get_str(vpiName, signal_handle),
               current_value.value.integer);
break;

case vpiRealVar:
    current_value.format = vpiRealVal;
    vpi_get_value(signal_handle, &current_value);
    vpi_printf(" real      %-10s value is %0.2f\n",
               vpi_get_str(vpiName, signal_handle),
               current_value.value.real);
break;

case vpiTimeVar:
    current_value.format = vpiTimeVal;
    vpi_get_value(signal_handle, &current_value);
    vpi_printf(" time      %-10s value is %x%x\n",
               vpi_get_str(vpiName, signal_handle),
               current_value.value.time->high,
               current_value.value.time->low);
break;
}
}
return;
}
```

Example 3-4 lists Verilog source code for testing `$show_all_signals`. This test is similar to the test for `$show_all_nets`, with the difference being the addition of more variables in the test bench level and changing the lower level adder model from a gate level model to an RTL model. Figure 3-5, which follows, shows the simulation results from running a simulation with `$show_all_signals`.

Example 3-4: `$show_all_signals` — aVerilog HDL test for the PLI application

```
'timescale 1ns / 1ns
module top;
  tri [1:0] results;
  integer test;
  real   foo;
  time   bar;
```

```
addbit i1 (test[0], test[1], test[2], results[0], results[1]);  
  
initial  
begin  
    test = 3'b000;  
    foo = 3.14;  
    bar = 0;  
    bar[63:60] = 4'hF;  
    bar[35:32] = 4'hA;  
    bar[31:28] = 4'hC;  
    bar[03:00] = 4'hE;  
  
    #1 test = 3'b011;  
  
    #1 $show_all_signals(top);  
    #1 $show_all_signals(i1);  
  
    #1 $stop;  
    #1 $finish;  
end  
endmodule  
  
/** An RTL level 1 bit adder model **/  
'timescale 1ns / 1ns  
module addbit (a, b, ci, sum, co);  
    input a, b, ci;  
    output sum, co;  
  
    wire a, b, ci;  
    reg sum, co;  
  
    always @ (a or b or ci)  
        {co, sum} = a + b + ci;  
endmodule
```

Figure 3-5: \$show_all_signals — simulation results using the PLI application

The screenshot shows the Cadence Verilog-XL Turbo NT simulation interface. The top menu bar includes File, Control, Set, Show, Select, Tools, Options, and Help. The toolbar contains icons for play, stop, step, zoom, and other simulation controls. The main window displays a Verilog code snippet and its simulation results.

```
31      #1 test = 3'b011;
32
33
34      #1 $show_all_signals(top);
35      #1 $show_all_signals(i1);
36
37      #1 $stop;
38      #1 $finish;
39
40 endmodule
```

Scope: top Subscopes: ▾

At time 2.00, signals in module top (top):

net	results	value is	10 (binary)
integer	test	value is	3 (decimal)
real	foo	value is	3.14
time	bar	value is	f000000ac000000e

At time 3.00, signals in module top.i1 (addbit):

net	a	value is	1 (binary)
net	b	value is	1 (binary)
net	ci	value is	0 (binary)
reg	sum	value is	0 (binary)
reg	co	value is	1 (binary)

L37 "show_all_signals_1_test.v": \$stop at simulation time 4
Type ? for help
C1 > |

Ready

3.13 Obtaining handles to the current hierarchy scope

The Verilog language allows a system task or system function to be invoked from any hierarchy scope. A *scope* in the Verilog HDL is a level of design hierarchy and can be represented by several constructs:

- Module instances
- Named statement groups
- Verilog HDL tasks
- Verilog HDL function

The following example calls the \$show_all_signals from a named statement group:

```
module top;
...
always @(posedge clock)
begin: local
    integer i;
    reg      local_bus;
...
$show_all_signals;
end
endmodule
```

A useful enhancement to the `$show_all_signals` example is to allow either no system task argument or a null system task argument to represent the hierarchy scope which called the `$show_all_signals` system task. The difference between no argument and a null argument is shown in the following two examples.

No system task/function arguments:

```
$show_all_signals;
```

A null system task/function argument:

```
$show_all_signals();
```

The following Verilog source code shows the enhanced usage possibilities for the `$show_all_signals` example:

```
module top;
...
addbit i1 (a, b, ci, sum, co); // instance of an adder
...
always @(sum or co)
$show_all_signals; // list signals in this module
always @(posedge clock)
begin: local
    integer i;
    reg      local_bus;
...
$show_all_signals; // list signals in this block
end
endmodule
```

```
module addbit (a, b, ci, sum, co);
...
always @ (sum or co)
  $show_all_signals(); // list signals in this instance
endmodule
```

In the above example, the `$show_all_signals` applications should search for signal names in the local hierarchy scope, which in one call is a module, and in another call is a named statement group. The name of any type of hierarchy scope can be passed to the system task as a task/function argument, but in this example, `$show_all_signals` is being called with no arguments. To obtain the local hierarchy scope without being passed the scope name requires two steps:

1. `vpi_handle(vpiSysTfCall, NULL)` returns a handle to the system task/function which called the PLI application.
2. `vpi_handle(vpiScope, systf_handle)` returns a handle to the scope containing the system task handle.

Example 3-5 contains the complete listing of the enhanced `$show_all_signals`, with the ability to use either no arguments or null arguments to represent the local design hierarchy scope. The *completf routine* is also enhanced from the previous examples to allow any valid Verilog scope to be used as an argument to `$show_all_signals`.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.03/show_all_signals_2_vpi.c
- Verilog test bench: Chapter.03/show_all_signals_2_test.v
- Verilog-XL results log: Chapter.03/show_all_signals_2_test.log

Example 3-5: `$show_all_signals`, version 2 — obtaining the local scope handle

```
#include <stdlib.h> /* ANSI C standard library */
#include <stdio.h> /* ANSI C standard input/output library */
#include "vpi_user.h" /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
/* using TF routines for simulation control */

/* prototypes of the PLI application routines */
int PLIbook_ShowSignals_completf(), PLIbook_ShowSignals_calltf();
void PLIbook_PrintSignalValues();

***** * $show_all_signals Registration Data
* (add this function name to the vlog_startup_routines array)
*****
```

```
void PLIbook_ShowSignals_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$show_all_signals";
    tf_data.calltf   = PLIbook_ShowSignals_calltf;
    tf_data.compiletf = PLIbook_ShowSignals_compiletf;
    tf_data.sizetf    = NULL;
    vpi_register_systf(&tf_data);
    return;
}

/*********************************************
 * compiletf routine
 *****/
int PLIbook_ShowSignals_compiletf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, arg_handle;
    int        tfarg_type;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handles to system task arguments */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    if (arg_iterator == NULL) {
        return(0); /* no arguments OK; skip remaining checks */
    }

    /* check the type of object in system task arguments */
    arg_handle = vpi_scan(arg_iterator);
    tfarg_type = vpi_get(vpiType, arg_handle);
    switch (tfarg_type) {
        case vpiModule:
        case vpiTask:
        case vpiFunction:
        case vpiNamedBegin:
        case vpiNamedFork:
            break; /* arg is a scope instance; continue to next check */
        case vpiOperation:
            if (vpi_get(vpiOpType, arg_handle) == vpiNullOp)
                break; /* null argument OK; continue to next check */
        default:
            /* wrong type specified for an argument */
            vpi_printf("ERROR: $show_all_signals arg 1");
            vpi_printf(" must be a scope instance or null\n");
            vpi_free_object(arg_iterator); /* free iterator memory */
            tf_dofinish(); /* abort simulation */
            return(0);
    }
}
```

```
/* check that there is only 1 system task argument */
arg_handle = vpi_scan(arg_iterator);
if (arg_handle != NULL) {
    vpi_printf("ERROR: $show_all_signals can only have 1 argument\n");
    vpi_free_object(arg_iterator); /* free iterator memory */
    tf_dofinish(); /* abort simulation */
    return(0);
}
return(0);
}

*****
* calltf routine
*****
int PLIBbook_ShowSignals_calltf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, scope_handle,
               signal_iterator;
    int format;
    s_vpi_time current_time;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    if (arg_iterator == NULL) {
        /* no arguments -- use scope that called this application */
        scope_handle = vpi_handle(vpiScope, systf_handle);
    }
    else {
        /* compiletf has already verified arg is scope instance or null */
        scope_handle = vpi_scan(arg_iterator);
        vpi_free_object(arg_iterator); /* free iterator memory */
        if (vpi_get(vpiType, scope_handle) != vpiModule)
            /* arg isn't a module instance; assume it is null */
        scope_handle = vpi_handle(vpiScope, systf_handle);
    }

    /* read current simulation time */
    current_time.type = vpiScaledRealTime;
    vpi_get_time(systf_handle, &current_time);

    vpi_printf("\nAt time %2.2f, signals in scope %s:\n",
              current_time.real,
              vpi_get_str(vpiFullName, scope_handle));

    /* obtain handles to nets in module and read current value */
    /* nets can only exist if scope is a module */
    if (vpi_get(vpiType, scope_handle) == vpiModule) {
        signal_iterator = vpi_iterate(vpiNet, scope_handle);
        if (signal_iterator != NULL)
```

```
    PLIbook_PrintSignalValues(signal_iterator);
}

/* obtain handles to regs in scope and read current value */
signal_iterator = vpi_iterate(vpiReg, scope_handle);
if (signal_iterator != NULL)
    PLIbook_PrintSignalValues(signal_iterator);

/* obtain handles to variables in scope and read current value */
signal_iterator = vpi_iterate(vpiVariables, scope_handle);
if (signal_iterator != NULL)
    PLIbook_PrintSignalValues(signal_iterator);

vpi_printf("\n"); /* add some white space to output */
return(0);
}

void PLIbook_PrintSignalValues(vpiHandle signal_iterator)
{
    vpiHandle    signal_handle;
    int         signal_type;
    s_vpi_value current_value;

    while ( (signal_handle = vpi_scan(signal_iterator)) != NULL ) {
        signal_type = vpi_get(vpiType, signal_handle);
        switch (signal_type) {
            case vpiNet:
                current_value.format = vpiBinStrVal;
                vpi_get_value(signal_handle, &current_value);
                vpi_printf(" net      %-10s value is %s (binary)\n",
                           vpi_get_str(vpiName, signal_handle),
                           current_value.value.str);
                break;

            case vpiReg:
                current_value.format = vpiBinStrVal;
                vpi_get_value(signal_handle, &current_value);
                vpi_printf(" reg      %-10s value is %s (binary)\n",
                           vpi_get_str(vpiName, signal_handle),
                           current_value.value.str);
                break;

            case vpiIntegerVar:
                current_value.format = vpiIntVal;
                vpi_get_value(signal_handle, &current_value);
                vpi_printf(" integer %-10s value is %d (decimal)\n",
                           vpi_get_str(vpiName, signal_handle),
                           current_value.value.integer);
                break;

            case vpiRealVar:
                current_value.format = vpiRealVal;
```

```
vpi_get_value(signal_handle, &current_value);
vpi_printf(" real %-10s value is %0.2f\n",
           vpi_get_str(vpiName, signal_handle),
           current_value.value.real);
break;

case vpiTimeVar:
    current_value.format = vpiTimeVal;
    vpi_get_value(signal_handle, &current_value);
    vpi_printf(" time %-10s value is %x%x\n",
               vpi_get_str(vpiName, signal_handle),
               current_value.value.time->high,
               current_value.value.time->low);
    break;
}
}
return;
}
```

3.14 Obtaining handles to multiple task/function arguments

Another useful enhancement to the `$show_all_signals` application is to allow multiple hierarchy scopes to be specified at the same time. For example:

```
$show_all_signals(i1, ,top.local);
```

In the above example, there are three system task/function arguments, the second argument being null, to indicate the local hierarchy scope. Only a very minor change is required in the `$show_all_signals` example to support any number of system task arguments. Since the relationship of a task call to its arguments is a one-to-many relationship, the example application has been accessing the first argument by using `vpi_iterate()` to obtain an iterator for all arguments, and then calling `vpi_scan()` just one time in order to access only the first argument. The only change that is needed is to place the call to `vpi_scan()` in loop to access all of the task arguments instead of just the first argument.

Example 3-6 illustrates using a C while loop to access each system task/function argument. In this example another level of structured programming has been added, in order to keep the *calltf routine* short and easy to maintain. In the new structure, the *calltf routine* contains the loop to access each task argument. Within the loop, a *GetAllSignals* function is called. And, within that function, a *PrintSignalValues* function is called. Structured programming is an important technique to use in PLI applications.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.03/show_all_signals_3_vpi.c
- Verilog test bench: Chapter.03/show_all_signals_3_test.v
- Verilog-XL results log: Chapter.03/show_all_signals_3_test.log

Example 3-6: \$show_all_signals, version 3 — obtaining handles for multiple tfargs

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>       /* ANSI C standard input/output library */
#include "vpi_user.h"    /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"    /* IEEE 1364 PLI TF routine library */
                       /* using TF routines for simulation control */

/* prototypes of the PLI application routines */
int  PLIbook_ShowSignals_completf(), PLIbook_ShowSignals_calltf();
void PLIbook_PrintSignalValues(), PLIbook_GetAllSignals();

/********************* $show_all_signals Registration Data
 * (add this function name to the vlog_startup_routines array)
 *****/
void PLIbook_ShowSignals_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$show_all_signals";
    tf_data.calltf   = PLIbook_ShowSignals_calltf;
    tf_data.completf  = PLIbook_ShowSignals_completf;
    tf_data.sizetf    = NULL;
    vpi_register_systf(&tf_data);
    return;
}

/********************* completf routine
 *****/
int PLIbook_ShowSignals_completf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, arg_handle;
    int        tfarg_type, tfarg_num=0;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handles to system task arguments */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    if (arg_iterator == NULL) {
        return(0); /* no arguments OK; skip remaining checks */
    }
}
```

```
}

/* check each argument */
while ( (arg_handle = vpi_scan(arg_iterator)) != NULL ) {
    tfarg_num++;

    /* check the type of object in system task arguments */
    tfarg_type = vpi_get(vpiType, arg_handle);
    switch (tfarg_type) {
        case vpiModule:
        case vpiTask:
        case vpiFunction:
        case vpiNamedBegin:
        case vpiNamedFork:
            break; /* arg is a scope instance; continue to next check */
        case vpiOperation:
            if (vpi_get(vpiOpType, arg_handle) == vpiNullOp) {
                break; /* null argument OK; continue to next check */
            }
        default:
            /* wrong type specified for an argument */
            vpi_printf("ERROR: $show_all_signals arg %d", tfarg_num);
            vpi_printf(" must be a scope instance or null\n");
            vpi_free_object(arg_iterator); /* free iterator memory */
            tf_dofinish(); /* abort simulation */
            return(0);
    }
}
return(0);
}

/*********************************************
 * calltf routine
 *****/
int PLIbook_ShowSignals_calltf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, scope_handle;
    int format;
    s_vpi_time current_time;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* read current simulation time */
    current_time.type = vpiScaledRealTime;
    vpi_get_time(systf_handle, &current_time);

    /* obtain handle to system task argument */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    if (arg_iterator == NULL) {
        /* no arguments -- use scope that called this application */
        scope_handle = vpi_handle(vpiScope, systf_handle);
```

```
    PLIbook_GetAllSignals(scope_handle, &current_time);
}
else {
    /* compiletf has already verified arg is scope instance or null */
    while ( (scope_handle = vpi_scan(arg_iterator)) != NULL ) {
        if (vpi_get(vpiType, scope_handle) != vpiModule) {
            /* arg isn't a module instance; assume it is null */
            scope_handle = vpi_handle(vpiScope, systf_handle);
        }
        PLIbook_GetAllSignals(scope_handle, &current_time);
    }
}
return(0);
}

void PLIbook_GetAllSignals(vpiHandle scope_handle, p_vpi_time
current_time)
{
    vpiHandle signal_iterator;

    vpi_printf("\nAt time %2.2f, signals in scope %s (%s):\n",
               current_time->real,
               vpi_get_str(vpiFullName, scope_handle),
               vpi_get_str(vpiDefName, scope_handle));

    /* obtain handles to nets in module and read current value */
    /* nets can only exist if scope is a module */
    if (vpi_get(vpiType, scope_handle) == vpiModule) {
        signal_iterator = vpi_iterate(vpiNet, scope_handle);
        if (signal_iterator != NULL)
            PLIbook_PrintSignalValues(signal_iterator);
    }

    /* obtain handles to regs in scope and read current value */
    signal_iterator = vpi_iterate(vpiReg, scope_handle);
    if (signal_iterator != NULL)
        PLIbook_PrintSignalValues(signal_iterator);

    /* obtain handles to variables in scope and read current value */
    signal_iterator = vpi_iterate(vpiVariables, scope_handle);
    if (signal_iterator != NULL)
        PLIbook_PrintSignalValues(signal_iterator);

    vpi_printf("\n"); /* add some white space to output */
    return;
}

void PLIbook_PrintSignalValues(vpiHandle signal_iterator)
{
    vpiHandle    signal_handle;
    int         signal_type;
```

```
s_vpi_value current_value;

while ( (signal_handle = vpi_scan(signal_iterator)) != NULL ) {
    signal_type = vpi_get(vpiType, signal_handle);
    switch (signal_type) {
        case vpiNet:
            current_value.format = vpiBinStrVal;
            vpi_get_value(signal_handle, &current_value);
            vpi_printf(" net      %-10s value is %s (binary)\n",
                       vpi_get_str(vpiName, signal_handle),
                       current_value.value.str);
            break;

        case vpiReg:
            current_value.format = vpiBinStrVal;
            vpi_get_value(signal_handle, &current_value);
            vpi_printf(" reg      %-10s value is %s (binary)\n",
                       vpi_get_str(vpiName, signal_handle),
                       current_value.value.str);
            break;

        case vpiIntegerVar:
            current_value.format = vpiIntVal;
            vpi_get_value(signal_handle, &current_value);
            vpi_printf(" integer %-10s value is %d (decimal)\n",
                       vpi_get_str(vpiName, signal_handle),
                       current_value.value.integer);
            break;

        case vpiRealVar:
            current_value.format = vpiRealVal;
            vpi_get_value(signal_handle, &current_value);
            vpi_printf(" real      %-10s value is %0.2f\n",
                       vpi_get_str(vpiName, signal_handle),
                       current_value.value.real);
            break;

        case vpiTimeVar:
            current_value.format = vpiTimeVal;
            vpi_get_value(signal_handle, &current_value);
            vpi_printf(" time      %-10s value is %x%x\n",
                       vpi_get_str(vpiName, signal_handle),
                       current_value.value.time->high,
                       current_value.value.time->low);
            break;
    }
}
return;
}
```

3.15 Summary

The VPI routines in the PLI standard provide complete access to the internal data structures of a Verilog simulation. This access is done using an object oriented method, where Verilog HDL constructs within the simulation data structure are treated as objects. The VPI routines use *handles* to reference these objects.

Each Verilog object has relationships to other objects. Each object also has specific properties, such as a name or a logic value. The VPI library uses a concise set of routines to access all types of objects. Predefined VPI constants represent each object type and the properties associated with the object. Object diagrams show the relationships of one object to other objects and the properties of each object type.

This chapter has focused on how to create PLI applications using the VPI library. The following three chapters will present more detail on the syntax and usage of the 27 VPI routines in the PLI standard. These chapters include several additional examples of PLI applications which use the VPI library.

CHAPTER 4

Details about the VPI Routine Library

The VPI routines are a library of 27 C functions that can interact with Verilog simulators. The previous chapter has introduced how PLI applications are developed using the VPI routines. This chapter presents the VPI library in more detail than the previous chapter, and presents how Verilog objects are documented in the VPI standard. Chapter 5 presents how to read and modify the values of Verilog objects, and Chapter 6 presents how to use VPI routines to synchronize PLI applications with Verilog simulation activity. Appendix B presents the complete syntax of the VPI routine library and the complete set of VPI object diagrams.

The concepts presented in this chapter are:

- PLI application performance considerations
- Verilog objects and object relationships
- Accessing system task and system function arguments
- Traversing Verilog hierarchy to obtain handles for Verilog HDL objects
- Accessing simulation time and time scale factors
- Miscellaneous VPI routines

4.1 PLI application performance considerations

The run-time performance of a simulator can be impacted in either a positive way or a negative way by PLI applications. Often, a complex algorithm can be represented in the C language, using C language data types, much more efficiently than in the hardware-centric Verilog HDL language. The C language can be used for an abstract representation of a design, when 4-state logic, logic transitions, simulation time, and other details are not required, but which a hardware description language must be able to represent. The abstraction that C offers often makes it possible to greatly increase the run-time performance of a simulation algorithm.

However, a poorly thought out PLI application can actually decrease the run-time performance of a simulation. Each call to a routine in the PLI library will take time to be executed. It is important to architect a PLI application to minimize the number of times VPI routines are used.

The following guidelines can help in planning an efficient PLI application:

- Good C programming practices are essential. C programming style and techniques not discussed within the scope of this book.
- Consider every call to a VPI routine as expensive, and try to minimize the number of calls.
- VPI routines which obtain object handles using an object's name are less efficient than VPI routines which obtain object handles based on a relationship to another object.
- Routines which convert logic values from a simulator's internal representation to C strings, and vice-versa, are less efficient than using other C data types. Strings are a convenient means of representing 4-state values in certain types of applications, but strings should be used prudently.
- When the same object must be accessed many times during a simulation, the handle can be obtained once and saved in an application-allocated storage area. Using a pointer to the storage area, a PLI application has immediate access to the object handle, without having to call a VPI routine to obtain the handle each time it is needed. Section 6.2 on page 184 in Chapter 6 discusses allocating storage within a PLI application.
- Use the VPI library to access the unique abilities of hardware description languages, such as representing hardware parallelism and hardware propagation times. Simulator vendors have invested a great deal in optimizing a simulator's algorithms, and that optimization should be utilized, rather than duplicated in a PLI application.

When developing a PLI application, one primary consideration should be how often a PLI application will be called during a simulation. It is well worth the effort to optimizing the performance of an application that is invoked every clock cycle, but may not be as important for an application that is only invoked once during a simulation.

NOTE → The objective of this book is to show how the routines in the VPI library are used. Short examples of using many of these routines are shown in the context of complete PLI applications. In order to meet the book's objectives, the examples presented in this book do not always follow the guidelines of efficient C coding and prudent usage of the VPI routines. It is expected that when parts of these example PLI applications are adapted for other applications, the coding style will also be modified to be more efficient and robust.

4.2 The VPI string buffer

A number of VPI routines return pointers to C character strings, such as the name of a net. These strings are stored in one or more temporary string buffers, and may be overwritten by other calls to VPI routines which return strings. Therefore, a PLI application should use the string pointer returned by a VPI routine immediately. If a string needs to be preserved, the PLI application should copy the string into its own storage space. Following are two examples of using strings in a PLI application.

Read a string and use it immediately:

```
char *string_p;      /* string pointer only, no storage */
string_p = vpi_get_str(vpiName, net_handle);
vpi_printf("string_p points to %s\n", string_p);
```

Read a string and copy it to application-allocated storage for later use:

```
char *string_p;      /* string pointer only, no storage */
char *string_keep;   /* another string pointer only */
string_p = vpi_get_str(vpiName, net_handle);
string_keep = malloc(strlen(string_p)+1);
strcpy(string, string_p); /* save string for later use */
```

NOTE → If application-allocated storage is to remain valid from one call to the application to another call, then the PLI application must maintain a pointer to the memory. Section 6.2 on page 184 in Chapter 6 discusses allocating memory and maintaining pointers.

4.3 VPI error handling

A well written PLI application will perform error checking on the values returned by VPI routines. If a VPI routine failed to return a valid value, passing the invalid return to another VPI routine may lead to unexpected behavior of the Verilog simulator, including program crashes. As an example, if `vpi_handle()` could not locate a target object, it will return a `NULL` instead of a valid handle. If the `NULL` were then passed as an input to `vpi_get_str()` as the reference argument, an error of some type will occur.

A `NULL` return from `vpi_handle()` is that routine's *exception value*, which indicates that an error occurred. A well written PLI application will check for these exception values and act accordingly. For example:

```
task_handle = vpi_handle(vpiSystfCall, NULL)
if (task_handle == NULL) {
    vpi_printf("ERROR: could not obtain task handle\n");
    return(0);
}
```

Most VPI routines have a specific exception return value when the routine cannot perform the operation requested. Routines which return integer or boolean values will return 0 if an error occurs. Routines which return double-precision values will return 0.0. Routines which return handles will return `NULL`. Routines which return string pointer will return `NULL` (note that `NULL` is not the same a null string).

For VPI routines which return integer, boolean or double values, the exception value could be a legitimate value. Therefore, it might not be possible to determine if an error occurred based on the exception return value for which to check.

The VPI library provides a useful routine called `vpi_chk_error()`, which is used to check for errors and to report detailed information about an error. This routine returns 0 if the previous call to a VPI routine was successful, and *non-zero* if the call resulted in an error. This return can be used as a true/false test to determine if an error occurred in the previous VPI routine call when the routine does not have an exception value. `vpi_chk_error()` is also useful for debugging problems in a PLI application.

```
int vpi_chk_error (info)
    p_vpi_error_info info      pointer to an application-allocated s_vpi_error_info
                                structure to receive error information.
```

Every VPI routine except `vpi_chk_error()` will set or clear an internal VPI error status flag, which is common to all VPI routines. When the flag is set by an error, it

will remain set until the next call to a VPI routine changes the flag. `vpi_chk_error()` only reads the error status flag, and does not modify it.

The input to `vpi_chk_error()` is a pointer to an `s_vpi_error_info` structure. If an error in the previous call to VPI routine occurred, `vpi_chk_error()` will fill in the fields of this structure with information about the error. This information includes the simulator product name and version, and the file name and line number containing the system task/function instance which called the PLI application. Using this information, a PLI application can generate meaningful error messages to aid in debugging the cause of an error. If the details about the error are not needed, a NULL can be specified as the input to `vpi_chk_error()`.

The definition of the `s_vpi_error_info` structure is in `vpi_user.h`.

```
typedef struct t_vpi_error_info {
    int      state;          /* vpiCompile, vpiPLI, vpiRun */
    int      level;          /* vpiNotice, vpiWarning, vpiError,
                                vpiSystem, vpiInternal */
    char   *message;
    char   *product;
    char   *code;
    char   *file;
    int     line;
} s_vpi_error_info, *p_vpi_error_info;
```

The Verilog simulator will fill in the fields of the `s_vpi_error_info` structure when `vpi_chk_error()` is called. The `state` field represents when the error occurred, using one of the constants `vpiCompile`, `vpiPLI` or `vpiRun`. The `level` field indicates the severity of the error, using one of the constants `vpiNotice`, `vpiWarning`, `vpiError`, `vpiSystem` or `vpiInternal`. The `message` field, `product` field and `code` field are pointers to strings, but the PLI standard does not define the wording of the strings. A simulator may use these fields for any message, and the message may vary from one simulator product to another. The `file` field is a pointer to a string containing Verilog source code file name which contains the instance of the system task/function that called the PLI application. The `line` field will be filled with the line number of the line which contains the system task/function instance.



NOTE It is possible for a VPI call to not be associated with a Verilog HDL source code line (such as when a VPI simulation callback occurs). In these situations, the file name pointer will be set to NULL. In many C compliers, printing a NULL with the `%s` format will result in an error. It is therefore important that a PLI application verify that the file name pointer is valid before printing the file name.



TIP

The `vpi_chk_error()` routine can provide important information about problems in a PLI application, and is a valuable debug utility. However, excessive use of this routine can degrade the run-time performance of simulation. For best performance, an application should use a VPI routine's exception return value for basic error checking, and only use `vpi_chk_error()` when there is no exception value or when the additional debug information is required.

In the following example, `vpi_chk_error()` is used to check for an error after `vpi_handle()` is called, and, if an error occurred, report a detailed error message. In this example, C conditional compilation is used to only include `vpi_chk_error()` when debugging an application.

```
#define PLIbookDebug 1 /* set to 0 to omit debug messages */
#ifndef PLIbookDebug
    s_vpi_error_info err; /* allocate a VPI error structure */
#endif

primitive_handle = vpi_handle(vpiPrimitive, NULL);
#ifndef PLIbookDebug /* if error, generate verbose debug message */
    if (vpi_chk_error(&err)) {
        vpi_printf("\nERROR:\n");
        vpi_printf(" Product: %s Code: %s\n", err.product, err.code);
        vpi_printf(" Message: %s\n", err.message);
        if (err.file != NULL)
            vpi_printf(" File: %s Line: %d\n\n", err.file, err.line);
    }
#else /* if error, generate basic error message */
    if (primitive_handle == NULL)
        vpi_printf("\nERROR: could not obtain primitive handle\n");
#endif
```


4.4 VPI object diagrams

The IEEE 1364 PLI standard includes an *object diagram* for each object which VPI routines can access. These object diagrams document:

- The *properties* of the object. For example, a net object has *name*, *vector size*, and *logic value* properties (as well as several other properties).
- The *relationships* of the object. Relationships indicate how an object is connected to or contained within other objects within a Verilog data structure. For example, a net is contained within a module, and may also be connected to other objects, such as a module port or primitive terminal.

The object diagrams use enclosures and arrows. The type of object is listed within each enclosure, and the relationships to other objects are shown as arrows between the enclosures. The properties of the object are listed below the diagrams. Figure 4-1 shows part of the object diagram for a Verilog module object.

Figure 4-1: Partial VPI object diagram for Verilog modules



NOTE: Not all objects and properties are listed in this example. Refer to Appendix B.1.2 on page 683 for the complete module diagram.

<i>int</i>	vpiType	returns vpiModule
<i>bool</i>	vpiTopModule	returns true if a module is a top-level module
<i>int</i>	vpiTimeUnit	returns the module time unit as 2 down to -15, where 2==100 seconds, 1==10s, 0==1s, -1==100ms, -2 ==10ms, -3==1ms,... -6==1us,... -9==1ns,... -12==1ps,... -15==1fs
<i>int</i>	vpiTimePrecision	returns module time precision as 2 down to -15
<i>str</i>	vpiName	returns the module instance name
<i>str</i>	vpiFullName	returns the module full hierarchical path name
<i>str</i>	vpiDefName	returns the module definition name
<i>str</i>	vpiFile	returns the file name containing the module instance
<i>str</i>	vpiLineNo	returns the file line number containing the module instance
<i>str</i>	vpiDefFile	returns the file name where the module is defined
<i>str</i>	vpiDefLineNo	returns the file line number where module is defined


NOTE → This book does not use the IEEE 1364 format for object diagrams. The use of enclosures and relationship arrows is similar, but this book adds information which is not shown in the IEEE object diagrams. This book adds:

- The constant names which are used to traverse from one object to another.
- The return values possible for each property constant.

4.4.1 Object diagram symbols

An object diagram contains four primary symbols and four font type faces:

- A ***solid enclosure***, such as **module** or **port**, designates a Verilog object. The name of the object is shown within the enclosure. The font used for the name has significance:
 - A **non-italicized, bold font** designates that this object is being defined in this diagram. In the diagram for module objects, the name, **module**, is in bold.
 - A non-italicized, non-bold font designates that this object is being referenced in this diagram, but is not being defined. The definition will appear in a different diagram. For example, in the diagram for module objects, the name **port**, is not bolded.
- A ***small dotted enclosure***, such as **variables**, designates a reference to a named class of Verilog objects. A class of objects is a group of several objects which have something in common. The name of the object class is shown within the enclosure, using an **italicized, non-bold font**. The specific objects within a class are listed in the diagram for the class definition.
- A ***large dotted enclosure***, which has small enclosures within it, designates the definition of a class of Verilog objects. The large enclosure contains all of the objects which make up the class. A class of objects may or may not have a name. However, only a named class can be referenced in another diagram. The name of the class is shown at the top of the large enclosure, using an **italicized, bold font**. Two examples of object class definitions are:



- A ***solid circle***, such as **O**, designates that a **NULL** is to be used as the reference object in the VPI routine that traverses to the target object.

4.4.2 Traversing object relationships

The relationships of one object to another object are shown as arrows in the object diagrams. Each arrow indicates the relationship from a reference object (the originating end of the line) to a target object (the terminating end of the line). The type of

arrow at the target object indicates the type of relationship. Most objects can be both a reference object or a target object, depending on which direction the Verilog hierarchy is being traversed.

There are three types of object relationships possible in a Verilog design:

- **One-to-one** relationships are represented by a line which terminates with a single arrow in an object diagram. A one-to-one relationship indicates that a given object is related to, at most, one of another type of object. In the module object diagram shown in Figure 4-1 on page 103, there is a single arrow going from **port** back to **module**, which indicates a given port is only contained within one module.
- **One-to-many** relationships are represented by a line which terminates with a double arrow in an object diagram. A one-to-many relationship indicates that a given object is related to any number of another type of object. In the module object diagram shown in Figure 4-1 on page 103, there is a double arrow going from **module** to **port**, which indicates there may be any number of ports within a module.
- **many-to-one** relationships are represented by a line which originates with double arrows and terminates with a single arrow. A many-to-one relationship indicates that any number of a given object are related to a single one of another type of object. There is only one type of object in Verilog which can have a many-to-one relationship: an interconnection path between modules (e.g.: a data bus might have many drivers connected to the same load).

In most object relationships, the connecting line both originates and terminates with a single or double arrow. In certain relationships, however, the connecting line has no arrow on the originating end, at the reference object. This indicates that the relationship is one-way. That is, the target object can be accessed from the reference object, but it is not possible to get back to the reference object from the target object.

The VPI routines used to traverse from one type of object to another are: `vpi_handle()`, `vpi_iterate()`, `vpi_scan()`, and `vpi_handle_multi()`. The usage of these routines is presented in the next section.

4.5 Obtaining handles for objects

Several VPI routines provide a means for a PLI application to locate and retrieve handles for Verilog objects. Two of these routines were introduced in the previous chapter on using VPI routines, and are presented again in this section, along with the remaining VPI routines which obtain handles for Verilog objects.

4.5.1 Traversing one-to-one relationships

vpi_handle() returns a handle for a target object which has a one-to-one relationship to a reference object. The routine is passed in two inputs.

vpiHandle vpi_handle (type, reference)

int **type** constant representing an object type.
vpiHandle **reference** handle for an object.

In the object diagram for modules shown in Figure 4-1 on page 103, there is a single arrow from **port** to **module**, indicating a one-to-one relationship. Assuming a PLI application had already obtained a handle for a port, the module which contains the port can be obtained using **vpi_handle()**.

```
vpiHandle port_handle, module_handle;
/* add code to obtain handle for a port */
module_handle = vpi_handle(vpiModule, port_handle);
```

The constant **vpiModule** used above indicates the type of object to be retrieved by **vpi_handle()**. The constant name is shown in the object diagram, next to the arrow going to that target object.

4.5.2 Traversing one-to-many relationships

Many objects in Verilog have a one-to-many relationship with other objects. A 2-step process is used to traverse these types of relationships.

1. Obtain a handle for an *iterator object* for all the target objects.
2. Process each object referenced by the iterator object, one-by-one.

vpi_iterate() sets up an iterator object which represents the first of the many target objects, and returns a handle for the iterator. If there are no target objects in the Verilog data structure, then **vpi_iterate()** will return **NULL**. This routine is passed in two inputs.

vpiHandle vpi_iterate (type, reference)

int **type** constant representing an object type.
vpiHandle **reference** handle for an object.

vpi_scan() returns a handle for each target object referenced by the iterator object. The routine is passed in a single input—the handle for the iterator:

vpiHandle vpi_scan (iterator)

vpiHandle iterator handle for an iterator object.

vpi_scan() returns the next target handle referenced by the iterator each time it is called. Therefore, the routine must be called repeatedly (typically using a loop) in order to access all of the target objects. When **vpi_scan()** has returned all objects referenced by the iterator, it will return **NULL** the next time it is called. The **NULL** return can be used as a flag to exit the loop.

NOTE

A valid iterator object must be passed as an input to **vpi_scan()**. It is an error to pass **NULL** as an input to **vpi_scan()**. However, **vpi_iterate()** will return a **NULL** if there are none of the requested target objects in the simulation data structure. For example, if **vpi_iterate()** were called to obtain a handle for an iterator for all ports of a module, and the module had no ports, then **vpi_iterate()** would return **NULL**. The PLI application must verify that the handle for the iterator object is not **NULL** before passing the iterator to **vpi_scan()**.

In the example object diagram for modules shown in Figure 4-1 on page 103, there is a double arrow from **module** to **port**, indicating a one-to-many relationship. Assuming a PLI application had already obtained a handle for a module, the ports of the module can be obtained using **vpi_iterate()** and **vpi_scan()**. For example:

```
vpiHandle module_handle, port_handle, port_iterator;  
/* add code to obtain a handle for a module */  
port_iterator = vpi_iterate(vpiPort, module_handle);  
if (port_iterator != NULL)  
    while ( (port_handle = vpi_scan(port_iterator)) != NULL) {  
        /* process the port handle */  
    }
```

Managing memory for iterator object handles returned by **vpi_iterate()**

Generally, the memory required for an iterator object is automatically maintained by the PLI. The **vpi_iterate()** routine will allocate the memory needed for the iterator, and **vpi_scan()** will automatically free the memory when it returns **NULL** (which indicates there are no more target objects referenced by the iterator). This

automatic memory management relieves the PLI application developer of the need to allocate and de-allocate memory for the iterator object.

NOTE

Because the memory for the iterator object is automatically freed, the iterator object set up by `vpi_iterate()` cannot be reused after `vpi_scan()` has returned `NULL`. Attempting to use an iterator handle after the iterator object memory is freed will result in unpredictable behavior in simulation (perhaps a program crash).

4.5.3 When to use `vpi_free_object()`

Normally, the memory management of iterator objects is automatic, but there is one circumstance which requires special attention. Occasionally, a PLI application might call `vpi_iterate()` to set up an iterator object, but then the application does not call `vpi_scan()` at all, or does not call `vpi_scan()` until the routine returns `NULL`. In this circumstance, the memory allocated by `vpi_iterate()` will not be de-allocated automatically. The PLI application must manually free the iterator object memory by calling the routine `vpi_free_object()`. The `vpi_free_object()` routine returns 1 (for true) if successful, and 0 (for false) if unsuccessful. The syntax for this routine is:

```
int vpi_free_object (object)
vpiHandle  object      handle for an iterator object.
```

In the following example, `vpi_iterate()` is used as a simple true/false test to see if a module has any ports. The handles to the ports are not needed, and so `vpi_scan()` is not called. Therefore, the memory allocated by `vpi_iterate()` must be explicitly de-allocated using `vpi_free_object()`.

```
int module_has_ports(vpiHandle module_handle)
{
    vpiHandle port_iterator;
    port_iterator = vpi_iterate(vpiPort, module_handle);
    if (port_iterator == NULL)
        return(0); /* no ports found */
    else
        vpi_free_object(port_iterator); /* free iterator list */
    return(1); /* ports were found */
}
```

NOTE

Do not free an iterator object which has already been de-allocated! Attempting to do so may result in undefined behavior by a simulator. Once `vpi_scan()` returns `NULL`, there is no need to call `vpi_free_object()` on the iterator, because the simulator has already freed it.

4.5.4 Obtaining inter-module path object handles

In the Verilog language, the output port of one module can be connected to the input port of another module, using a net data type. The connection from an output to an input is referred to as an *inter-module path*. In actual hardware, this inter-connection will have a real delay, but, within the Verilog language, there is no construct to accurately represent that delay. The PLI, however, can add, read and modify inter-module path delays.

A module input port can be driven by any number of module output ports. An example of multiple inter-connections might be a shared data bus which can be driven by several components. Because of the possibility of multiple connections, the VPI library refers to inter-module paths as a many-to-one relationship.

vpi_handle_multi() is used to obtain a handle for an inter-module path. Other VPI routines can then annotate delays to the path. This routine requires three inputs:

```
vpiHandle vpi_handle_multi (type, reference1, reference2)
    int          type      constant of vpiInterModPath.
    vpiHandle   reference1 handle for an output or inout port.
    vpiHandle   reference2 handle for an input or inout port.
```

The constant **vpiInterModPath** is the only type constant supported by **vpi_handle_multi()**. The ports specified must be the same vector size, but they do not need to be in the same level of Verilog hierarchy. If no interconnecting net exists between the two ports, then **vpi_handle_multi()** will return **NULL**. An example of using this routine is:

```
inter_mod_path_h = vpi_handle_multi(vpiInterModPath,
                                      in_port_handle,
                                      out_port_handle);
if (inter_mod_path_h != NULL)
    /* inter connection path not found -- process an error */
else
    /* read or modify the inter-connect delay values */
```

4.5.5 Obtaining object handles using an object's name

vpi_handle_by_name() obtains a handle for an object using the name of the object. The handle for any Verilog object with a **vpiFullName** property in the object diagrams can be obtained using this routine. The routine requires two inputs:

vpiHandle vpi_handle_by_name (name, scope)

*char *name* name of an object.

vpiHandle scope handle for an object with a name scope, or *NULL*.

The name provided can be the local name of the object, a relative hierarchical path name, or a full hierarchical path name. The PLI will search for the object, using the name search rules of the Verilog language. The full search rules are outside the scope of this book. Briefly, however, Verilog searches for a name in the local hierarchy scope first, then as a relative path, then as a full path. If the object cannot be found, then *vpi_handle_by_name()* will return *NULL*.

A *NULL* can be specified in place of the scope handle, which indicates that the search should begin at top level of the Verilog design hierarchy.

 **TIP** *vpi_handle_by_name()* is an expensive routine in terms of simulation performance, and should be used judiciously. It is much more efficient to obtain a handle for an object based on its relationship to some other object.

In the following code fragment, the name of a primitive instance and a delay value are read from a file. The handle for the primitive is obtained using *vpi_handle_by_name()*, and the delay is annotated onto the primitive.

```
char      prim_name[64];
vpiHandle prim_handle;
double    new_delay;

fscanf(file_p, "%s %f", prim_name, &new_delay);
prim_handle = vpi_handle_by_name(prim_name, NULL);
if (prim_handle != NULL)
    /* add new delay value to the primitive object */
else
    /* error: primitive not found */
```

4.5.6 Obtaining object handles using an object's index number

vpi_handle_by_index() is used to obtain a handle for an object using the object's index position. This routine requires two inputs:

vpiHandle vpi_handle_by_index (parent, index)

vpiHandle parent handle for an object.

int index index number of an object.

The handle for any Verilog object which has a vpiIndex relationship in the object diagrams can be obtained using this routine. These objects are the bits of a vector net or vector reg, and the words of a memory array or a variable array. Note that `vpi_handle_by_index()` requires the object have an index *relationship*. Some Verilog objects have an index *property*. `vpi_handle_by_index()` does not support index properties.

The bit or word that is represented by an index number is based on the declaration of the vector or array. If an index is out of range, a NULL is returned.

Assuming a handle for a net vector had already been obtained, the following example would obtain a handle for bit number 2 of the vector:


```
vpiHandle net_handle, bit_handle;  
/* add code to obtain handle for a net vector */  
bit_handle = vpi_handle_by_index(net_handle, 2);  
if (bit_handle == NULL)  
    /* process the net bit object */  
else  
    /* error: terminal not found */
```

4.6 System task and system function objects

The Verilog PLI allows users to create user-defined system tasks and user-defined system functions which can be used in Verilog HDL source code. How user-defined system tasks and system functions are defined was presented in Chapter 2.

4.6.1 System task/function arguments

A user-defined system task or system function can have any number of arguments, including none. The VPI standard does not specify an index number for task/function arguments, but the TF and ACC libraries number the arguments from left to right, starting with 1. For convenience in describing PLI applications, this book uses the same numbering scheme in the VPI chapters. In the following example:



- Task/Function argument number 1 is a string, with the value "vectors.pat".
- Task/Function argument number 2 is a signal, with the name input_bus.

4.6.2 Multiple instances of system tasks and system functions

The Verilog HDL source code can reference the same system task/function any number of times. Section 2.3 on page 33 in Chapter 2 discussed the different ways in which multiple instances of a system task/function can occur in Verilog source code. Each instance of a system task/function is unique, and has unique argument values. However, each instance is associated with the same *calltf routine*. For example:

```
always @(posedge clock)
$read_test_vector("A.dat", data_bus);

always @(negedge clock)
$read_test_vector("B.dat", data_bus);
```

In the preceding Verilog source code example, `$read_test_vector` is used two times. The *calltf routine* is a C function, and every instance of `$read_test_vector` will invoke the same C function. Within the simulation data structure, however, each call to the function is unique. In the preceding example, at each positive edge of clock the *calltf routine* will be invoked, and when the routine reads the value of its argument 1, it will retrieve the string "A.dat". At the negative edge of clock, the same *calltf routine* will be executed, but when this call to the routine reads the value of its argument 1, it will retrieve the string "B.dat".

NOTE → *Storage allocated by a PLI application is not unique to each instance!* If the application declares static variables or allocates memory, all instances of the application will share those variables or memory.


If a PLI application requires unique storage, such as storing a different file pointer for each system task instance, then the application must allocate and identify that unique storage. Section 6.2 on page 184 in Chapter 6 presents how instance specific storage can be allocated.

4.6.3 Obtaining a handle for a system task/function instance

An instance of a system task or system function is an object, and the VPI routines can obtain a handle for this object. This handle will be unique for each instance of the system task/function. A partial object diagram for a system task/function is shown in Figure 4-2.

Figure 4-2: Object diagram for a task/function call (partial)

NOTE: Refer to Appendix B.1.17 on page 697 for the complete object diagram for ports.



In this diagram, the objects **sys task call** and **sys func call** are part of an object group, called **tf call**. These objects can be accessed from a small circle, using the constant **vpiSysTfCall**. The small circle indicates that a **NULL** is used as the reference handle for **vpi_handle()**. Therefore, a handle for the system task/function instance which called a PLI application is obtained using:

```

vpiHandle systf_handle;
systf_handle = vpi_handle(vpiSysTfCall, NULL);
  
```

4.6.4 Accessing the arguments of system tasks and system functions

Once a handle for an instance of a system task or system function is obtained, the arguments to the system task/function can be accessed by obtaining handles for the arguments. In the Verilog HDL, many different values and data types can be used as a task/function argument—an integer number, a variable name, a net name, a module instance name, or a literal string (such as a file name), are just a few of the legal arguments. In the object diagram for **tf call**, the argument object is a reference to an

“expression” group, which includes all the object types which can be used as an argument. Since a system task/function can have any number of arguments, the line from `_tfcall__`, to `_expr__`, terminates with a double arrow, indicating a one-to-many relationship.

The following two examples illustrate using VPI routines to access the arguments of a system task/function. These examples are complete C functions which can be called from other PLI applications.

Example 4-1 lists a simple C function called `PLIbook_numargs_vpi()`, which uses the VPI routines to count and return the number of arguments in a system task/function call.

Example 4-2 lists a C function called `PLIbook_getarg_handle_vpi()`, which uses the VPI routines to return a handle for a specific system task/function argument. The input to this routine is an index number of the argument, where the left-most argument is index number 1, and the numbers increases from left to right.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.04/vpi_utilities.c`
- Verilog test bench: `Chapter.04/vpi_utilities_test.v`
- Verilog-XL results log: `Chapter.04/vpi_utilities_test.log`

Example 4-1: `PLIbook_numargs_vpi()` — accessing system task/function arguments

```
int PLIbook_numargs_vpi()
{
    vpiHandle systf_h, arg_itr, arg_h;
    int tfnum = 0;

    #ifdef PLIbookDebug
        s_vpi_error_info err; /* structure for error handling */
    #endif

    systf_h = vpi_handle(vpiSysTfCall, NULL);
    #ifdef PLIbookDebug /* if error, generate verbose debug message */
        if (vpi_chk_error(&err)) {
            vpi_printf("ERROR: PLIbook_numargs_vpi() could not obtain handle
to systf call\n");
            vpi_printf("File %s, Line %d: %s\n",
                err.file, err.line, err.message);
        }
    #else /* if error, generate brief error message */
        if (systf_h == NULL)
            vpi_printf("ERROR: PLIbook_numargs_vpi() could not obtain handle
to systf call\n");
    
```

```
#endif

arg_itr = vpi_iterate(vpiArgument, systf_h);
#ifndef PLIbookDebug /* if error, generate verbose debug message */
    if (vpi_chk_error(&err)) {
        vpi_printf("ERROR: PLIbook_numargs_vpi() could not obtain iterator
to systf args\n");
        vpi_printf("File %s, Line %d: %s\n",
                   err.file, err.line, err.message);
    }
#else /* if error, generate brief error message */
    if (systf_h == NULL)
        vpi_printf("ERROR: PLIbook_numargs_vpi() could not obtain iterator
to systf args\n");
#endif

while (arg_h = vpi_scan(arg_itr) ) {
    tfnum++;
}

return(tfnum);
}
```

Example 4-2: PLIbook_getarg_handle_vpi() — accessing system task/function args

```
vpiHandle PLIbook_getarg_handle_vpi(int argNum)
{
    vpiHandle systf_h, arg_itr, arg_h;
    int i;
#ifndef PLIbookDebug
    s_vpi_error_info err; /* structure for error handling */
#endif

    if (argNum < 1) {
        vpi_printf("ERROR: PLIbook_getarg_handle_vpi() argNum invalid\n");
        return(NULL);
    }

    systf_h = vpi_handle(vpiSysTfCall, NULL);
#ifndef PLIbookDebug /* if error, generate verbose debug message */
    if (vpi_chk_error(&err)) {
        vpi_printf("ERROR: PLIbook_getarg_handle_vpi() could not obtain
handle to systf call\n");
        vpi_printf("File %s, Line %d: %s\n",
                   err.file, err.line, err.message);
    }

```

```
#else /* if error, generate brief error message */
    if (systf_h == NULL) {
        vpi_printf("ERROR: PLIbook_getarg_handle_vpi() could not obtain
handle to systf call\n");
        return(NULL);
    }
#endif

arg_itr = vpi_iterate(vpiArgument, systf_h);
#ifndef PLIbookDebug /* if error, generate verbose debug message */
    if (vpi_chk_error(&err)) {
        vpi_printf("ERROR: PLIbook_getarg_handle_vpi() could not obtain
iterator to systf args\n");
        vpi_printf("File %s, Line %d: %s\n",
                   err.file, err.line, err.message);
    }
#else /* if error, generate brief error message */
    if (systf_h == NULL) {
        vpi_printf("ERROR: PLIbook_getarg_handle_vpi() could not obtain
iterator to systf args\n");
        return(NULL);
    }
#endif

for (i=1; i<=argNum; i++) {
    arg_h = vpi_scan(arg_itr);
#ifndef PLIbookDebug /* if error, generate verbose debug message */
    if (vpi_chk_error(&err)) {
        vpi_printf("ERROR: PLIbook_getarg_handle_vpi() could not obtain
handle to systf arg %d\n", i);
        vpi_printf("File %s, Line %d: %s\n",
                   err.file, err.line, err.message);
    }
#endif
    if (arg_h == NULL) {
        vpi_printf("ERROR: PLIbook_getarg_handle_vpi() systf arg %d out-
of-range\n",
                   argNum);
        return(NULL);
    }
}
vpi_free_object(arg_itr); /* free iterator -- didn't scan all args */

return(arg_h);
}
```

**TIP**

The preceding example 4-2 is not as efficient for simulation run-time performance as it could be. In the preceding example, the routines `vpi_iterate()` and `vpi_scan()` are called every time the `PLibook_getarg_handle_vpi()` application is invoked, in order to obtain the handle for the requested system task/function argument. A more efficient method would be to store all system task/function argument handles in storage that is allocated by the `PLibook_getarg_handle_vpi()` application, and then return the stored handle each time the application is called. The basic steps involved in storing the handles are:

- The first time the `PLibook_getarg_handle_vpi()` application is called, allocate an array of `vpiHandle` variables, with 1 element in the array for each system task/function argument.
- Call `vpi_iterate()` and `vpi_scan()` one time to obtain handles to the system task/function arguments, and store the handles in the array.
- Each time the application is called, use the system task/function argument index number to retrieve the appropriate handle from the array

Note that, in order to store the system task/function arguments in an array, a unique array must be allocated for each instance of a system task/function which uses the `PLibook_getarg_handle_vpi()` application. Section 6.2 on page 184 in Chapter 6 discusses how to allocate storage for each instance of a system task/function.

The `vpi_utilities.c` file that contains the preceding example on the CD accompanying this book also contains an example of the `PLibook_getarg_handle_vpi()`, which uses the more efficient array method to store the system task/function argument handles.

4.7 Traversing Verilog hierarchy using object relationships

Using the object diagrams, the Verilog design hierarchy can be traced from one object to any other object anywhere in a Verilog design. Traversing the design hierarchy often requires following object relationships across several diagrams, until the desired destination is attained.

For example, suppose a PLI application had obtained a handle for a module, and the application needs to locate every module output, where the output is also connected to a module path delay. The following Verilog source code shows the starting and ending objects for which handles are desired in this example.

```

start with a handle
to a module
----->
module dff (clk, d, q, qb);
  input clk, d;
  output q, qb;
  ff_prim g1 (q, d, clk, );
  not g2 (qb, q);
endmodule
----->
obtain handles to all output
ports which are connected to
a module path delay (port q
in this example)


```

The following steps show how the relationships in the object diagrams can be followed to traverse from the starting point in a Verilog design to the desired destination:


1. The **module** diagram (shown in Figure 4-3, which follows) shows a one-to-many connection from **module** to **mod path**.

Using **vpi_iterate(vpiModPath, module_handle)**, an iterator for all module paths in a module can be obtained.

Figure 4-3: Object diagram for Verilog modules (partial)




2. The **mod path** object diagram (shown in Figure 4-4, on the next page) shows a one-to-many relationship from **mod path** to **path term**. This connection can be traversed with specific constants to obtain handles for the path input terminals, output terminals, or the data terminal. Since this example application is looking for the output ports connected to module paths, the output terminal is the relationship which needs to be followed. Using **vpi_iterate(vpiModPathOut, path_handle)**, an iterator for all path outputs can be obtained.

Figure 4-4: Object diagram for module paths (partial)

- The **path term** object in the preceding **mod path** object diagram shows a one-to-one connection to an **expr** object group (expr stands for expression). This object relationship is traversed using **vpi_handle(vpiExpr, path_term_handle)**.

The object diagram for the expressions (shown in Figure 4-5, which follows) lists many types of objects. A knowledge of the Verilog HDL language is required, in order to know which types of objects in the expression diagram can be connected to an output path terminal. In the context of a connection from a module path terminal to a module output port, the only type of object which the Verilog HDL syntax will allow is a net, represented by the **nets** object group. Therefore, in this context, the object diagram for nets will show the next relationship in the connection from a module path to a port.

Figure 4-5: Object diagram for expressions (partial)

- The **nets** object diagram (shown in Figure 4-6, below) shows two types of one-to-many connections from **nets** to **ports**. One connection is accessed using **vpiPorts**, and the other using **vpiPortInst**. The connection accessed by **vpiPorts** is the connection to a port of the module which contains the net. The connection accessed by **vpiPortInst** is the connection to a port of

a module instance within the current module. The following Verilog model illustrates the difference between the two types of objects:

```
module chip (clock, in1, out1, out2);
  ...
  dff u1 (clock, in1, out1, out2);
endmodule


module dff (clk, d, q, qb);
  ...
endmodule
```

module ports are accessed using **vpiPort**

module *instance* ports are accessed using **vpiPortInst**

For the example being illustrated in this section (finding the module ports which are connected to a module path output), the type of port object needed is accessed using **vpi_iterate(vpiPort, net_handle)**, which will return a list of all ports of the module to which the net is connected.

Figure 4-6: Object diagram for nets (partial)



Example summary

To summarize the example described in this section, four object connections were followed to traverse the Verilog hierarchy from a module object to the output ports connected to a path delay. The connection path is from **module** to **mod path** to **path term** to **nets**, (in the expression group) to **ports**.

The following C source code illustrates how these multiple connections through the Verilog hierarchy are traversed.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.04/list_pathout_ports_vpi.c
- Verilog test bench: Chapter.04/list_pathout_ports_test.v
- Verilog-XL results log: Chapter.04/list_pathout_ports_test.log

Example 4-3: \$list_pathout_ports — traversing Verilog hierarchy

```
int PLIbook_ListPort_calltf(char *user_data)
{
    vpiHandle module_handle, systf_h, arg_itr,
               path_itr, path_handle,
               term_itr, term_handle,
               port_itr, port_handle,
               net_handle;

    /* get module handle from first system task argument. Assume the */
    /* compilorf routine has already verified correct argument type. */
    systf_h = vpi_handle(vpiSysTfCall, NULL);
    if (systf_h == NULL) {
        vpi_printf("ERROR: list_pathout_ports could not obtain handle to
systf call\n");
        return(0);
    }
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    if (systf_h == NULL) {
        vpi_printf("ERROR: list_pathout_ports could not obtain iterator to
systf args\n");
        return(0);
    }
    module_handle = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* free itr since did not scan until null */

    vpi_printf("\nModule %s\n", vpi_get_str(vpiDefName, module_handle));

    path_itr = vpi_iterate(vpiModPath, module_handle);
    if (path_itr == NULL) {
        vpi_printf("    No module paths found\n");
        return(0);
    }
    while (path_handle = vpi_scan(path_itr)) {
        term_itr = vpi_iterate(vpiModPathOut, path_handle);
        if (term_itr == NULL) {
            vpi_printf("    No path output terminal found\n");
            break; /* go to next path */
        }
        while (term_handle = vpi_scan(term_itr)) {
            net_handle = vpi_handle(vpiExpr, term_handle);
            port_itr = vpi_iterate(vpiPort, net_handle);
            if (port_itr == NULL) {

```


```

        vpi_printf("  Path output does not connect to a port\n");
        break; /* go to next path output terminal */
    }
    while (port_handle = vpi_scan(port_itr)) {
        vpi_printf("  Port %s is connected to a path delay output\n",
                   vpi_get_str(vpiName, port_handle));
    }
}
return(0);
}

```

4.7.1 Traversing hierarchy across module ports

The Verilog hierarchy connections can be traversed across module boundaries by following the relationships within and without module ports. A module port has two connections, as shown in the following diagram.




The ***lowconn*** (hierarchically lower connection) is the signal *inside* the module that is connected to a port. In the object diagrams, the internal connection is shown as an ***expression***, because a number of different types of objects can be connected to the port. Within a module, the Verilog HDL syntax restricts the expression connected to a port to the Verilog data types of ***nets***, ***regs*** and ***variables*** (except real variables). The expression connected to the port can be a scalar signal (1-bit wide), a vector, a bit or part select of a vector, or a concatenation of any valid data type. Any of these expression types can be connected to an output port, but only expressions using the ***net*** class of data types can be connected to an input or inout port. The Verilog HDL syntax also allows the bits of the ***lowconn*** connections to have different port directions, referred to as a ***mixed I/O*** port.

The ***highconn*** (hierarchically higher connection) is the signal *outside* the module that is connected to a port. External to a module, the expression connected to the port can be a scalar signal, a vector, a bit or part select of a vector, or a concatenation of signals, a constant, a literal value, an operation, or the return of a function call. Any of

these expression types can be connected to an input port, but only expressions using the *net* class of data types can be connected to an output or inout port.

The object diagram for *ports*, (shown in Figure 4-7) shows a one-to-one connection to either the *highconn* and *lowconn* connection to the port. The connecting object is an *expr*, object group, to allow for different types of objects which may be connected to the port. Using the handle for the *expr*, object, the type property of the object can be accessed to determine what is connected to the port. The object diagram for that type of object can then be used to continue traversing the Verilog design hierarchy. The object type property is accessed using `vpi_get(vpiType, <expression_handle>)`.

Figure 4-7: Object diagram for ports (partial)



The following example accesses all ports of a module and lists the port name, port size, port direction, and type of object connected as the *lowconn* and *highconn*. The name, size and direction are all properties of the module port.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.04/port_info_vpi.c
- Verilog test bench: Chapter.04/port_info_test.v
- Verilog-XL results log: Chapter.04/port_info_test.log

Example 4-4: \$port_info — traversing hierarchy across module ports

```

int PLIbook_PortInfo_calltf(char *user_data)
{
    vpiHandle systf_h, arg_itr, mod_h,
               port_itr, port_h, lowconn_h, highconn_h;
    int        lowconn_type, highconn_type;
  
```

```
/* get module handle from first system task argument. Assume the */
/* compilertf routine has already verified correct argument type. */
systf_h = vpi_handle(vpiSysTfCall, NULL);
if (systf_h == NULL) {
    vpi_printf("ERROR: list_pathout_ports could not obtain handle to
systf call\n");
    return(0);
}
arg_itr = vpi_iterate(vpiArgument, systf_h);
if (systf_h == NULL) {
    vpi_printf("ERROR: list_pathout_ports could not obtain iterator to
systf args\n");
    return(0);
}
mod_h = vpi_scan(arg_itr);
vpi_free_object(arg_itr); /* free itr since did not scan until null */

vpi_printf("\nModule %s (instance %s)\n",
          vpi_get_str(vpiDefName, mod_h),
          vpi_get_str(vpiFullName, mod_h));

port_itr = vpi_iterate(vpiPort, mod_h);
if (!port_itr) {
    vpi_printf("    No ports found\n");
    return(0);
}
while (port_h = vpi_scan(port_itr)) {
    vpi_printf("    Port name is %s\n", vpi_get_str(vpiName, port_h));
    vpi_printf("    Size is %d\n", vpi_get(vpiSize, port_h));
    switch (vpi_get(vpiDirection, port_h)) {
        case vpiInput: vpi_printf("        Direction is input\n"); break;
        case vpiOutput: vpi_printf("        Direction is output\n"); break;
        case vpiInout: vpi_printf("        Direction is inout\n"); break;
    }
    lowconn_h = vpi_handle(vpiLowConn, port_h);
    lowconn_type = vpi_get(vpiType, lowconn_h);
    vpi_printf("    Low conn data type is %s\n",
               PLIbook_fetch_type_str_vpi(lowconn_type));

    highconn_h = vpi_handle(vpiHighConn, port_h);
    if (!highconn_h) {
        vpi_printf("    No high conn\n");
        return(0);
    }
    highconn_type = vpi_get(vpiType, highconn_h);
    vpi_printf("    High conn data type is %s\n",
               PLIbook_fetch_type_str_vpi(highconn_type));
}
return(0);
}
```

```
char *PLIbook_fetch_type_str_vpi(int type)
{
    switch (type) {
        case vpiNet           : return("vpiNet");
        case vpiNetBit         : return("vpiNetBit");
        case vpiReg            : return("vpiReg");
        case vpiRegBit         : return("vpiRegBit");
        case vpiPartSelect    : return("vpiPartSelect");
        case vpiIntegerVar    : return("vpiIntegerVar");
        case vpiTimeVar        : return("vpiTimeVar");
        case vpiMemoryWord    : return("vpiMemoryWord");
        case vpiVarSelect      : return("vpiVarSelect");
        case vpiConstant       : return("vpiConstant");
        case vpiParameter      : return("vpiParameter");
        case vpiFuncCall       : return("vpiFuncCall");
        case vpiSysFuncCall   : return("vpiSysFuncCall");
        case vpiOperation      : return("vpiOperation");
        default                : return("UNDEFINED TYPE");
    }
}
```

NOTE The preceding example prints the name of the type constant which represents the type property of the object connected to a port. The IEEE 1364-1995 Verilog standard only provides access to the integer value of the type constant (using `vpi_get(vpitype, object_handle)`). This example passes this integer value to a function which maps the integer value to the name of the constant. The proposed IEEE 1364-1999 standard will be able to access the name of type constants directly, using `vpi_get_str(vpitype, object_handle)`.

4.7.2 Traversing through multiple levels of hierarchy

In some PLI applications, it may be necessary to traverse through multiple levels of design hierarchy. For instance, an application might need to locate every Verilog primitive in an entire design. This requires starting at the very top of the design hierarchy tree, and descending down through every possible branch of the tree, searching for primitive instances. The PLI application will need to traverse an unknown number of hierarchy levels, because the number of branches in the tree, and the depth of the hierarchy are not available.

The most straightforward programming method to traverse an unknown number of hierarchy levels is to use recursive calls to a C function. A C function is passed in a handle for a module, and then searches for module instances within that module. For each module instance found, the function calls itself, passing the module instance handle as an input.

Example 4-5 lists the PLI source code for an application called `$count_all_prims`. The application does not require any inputs. Instead, the application searches for all top level modules, and then traverses all levels from the modules on down and counts the number of primitives found in each module instance as the hierarchy is traversed. A summary report is printed after all hierarchy levels are traversed.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.04/count_all_prims_vpi.c
- Verilog test bench: Chapter.04/count_all_prims_test.v
- Verilog-XL results log: Chapter.04/count_all_prims_test.log

Example 4-5: `$count_all_prims` — traversing multiple levels of hierarchy

```
*****
 * Calltf application
 *****
int PLIbook_CountPrims_calltf(char *user_data)
{
    vpiHandle top_mod_itr, top_mod_h;
    int total_prims = 0;
    top_mod_itr = vpi_iterate(vpiModule, NULL); /*get top modules*/
    while (top_mod_h = vpi_scan(top_mod_itr)) {
        total_prims += PLIbook_find_child_mod(top_mod_h);
    }
    vpi_printf("Total number of primitives is %d\n\n", total_prims);
    return(0);
}

*****
 * Function to look for module instances in local scope.
 * THIS FUNCTION CALLS ITSELF RECURSIVELY.
 *****
int PLIbook_find_child_mod(vpiHandle this_mod_h)
{
    vpiHandle child_mod_itr, child_mod_h;
    int prims_in_child;

    prims_in_child = PLIbook_count_local_prims(this_mod_h);
    child_mod_itr = vpi_iterate(vpiModule, this_mod_h);
    if (child_mod_itr != NULL)
        while (child_mod_h = vpi_scan(child_mod_itr))
            prims_in_child += PLIbook_find_child_mod(child_mod_h);
    return(prims_in_child);
}
```

```
*****
 * Function to count primitives in local scope.
 ****
int PLIBook_count_local_prims(vpiHandle module_h)
{
    vpiHandle prim_itr, prim_h;
    int prims_in_mod = 0;

    prim_itr = vpi_iterate(vpiPrimitive, module_h);
    if (prim_itr != NULL)
        while (prim_h = vpi_scan(prim_itr))
            prims_in_mod++;
    return (prims_in_mod);
}
```

4.8 Writing messages to files

The VPI library provides a routine for writing information to disk files, and special routines for controlling those files.

unsigned int vpi_mcd_open (file_name)

char *file_name name of a file to be opened.

This routine opens a file for writing from a PLI application. The input to the routine is a string containing the file name to be opened. The routine returns an integer, which is a *multi-channel-descriptor (mcd)*. If the file cannot be opened, a value of 0 is returned for the *mcd*. If the file is already open, the value of the file's current *mcd* is returned.

The *mcd* value returned by *vpi_mcd_open()* will have a single bit set to represent an open file. Twenty-nine application-specified files can be opened at a time, plus three output channels which are reserved by the PLI. The reserved channels are represented by the first three bits of the *mcd* integer for the following special meanings:

- **bit 0** (a decimal 1) represents the operating system's standard output channel.
- **bit 1** (a decimal 2) represents the operating system's standard error channel.
- **bit 2** (a decimal 4) represents the simulator's log file. A log file may or may not be maintained by the Verilog simulator. If the simulator does maintain a log file, it generally is used to capture all output which is written to the simulator's display window.

int vpi_mcd_printf (mcd, format, arg1,...argn)

int **mcd** multi-channel descriptor of open files.
char ***format** quoted character string of formatted message.
arg1...argn arguments to formatted message string.

vpi_mcd_printf() writes messages to files which were opened using *vpi_mcd_open()*. The first argument is an *mcd* value. The remaining arguments use the same syntax as the C *fprintf()* function. The routine returns the number of characters written, or EOF if an error occurred.

char*vpi_mcd_name (mcd)

unsigned int mcd multi-channel descriptor representing an open file

This routine returns the name of a file represented by an *mcd*. The argument to the routine is the *mcd* returned when a file was opened. The *mcd* value used must represent a single file (e.g.: only 1 bit may be set in the *mcd*).

unsigned int vpi_mcd_close (mcd)

unsigned int mcd multi-channel descriptor representing open files

vpi_mcd_close() closes files which were opened using *vpi_mcd_open()*. The argument is an *mcd* value. Multiple files can be closed at the same time by logically OR'ing several *mcd* values together. It is illegal to use the *mcd* values which represent simulator's output channel, stderr, and the simulator's log file. The routine returns 0 if successful, and the *mcd* of files not closed if an error occurs.

NOTE → The *mcd* value returned by *vpi_mcd_open()* is not the same as the *mcd* value returned by the built-in *\$fopen* system function in Verilog. The *mcd* values cannot be shared between the PLI applications and the Verilog HDL models.

NOTE → The proposed IEEE 1364-1999 standard will add new file I/O capabilities to the Verilog HDL, which will use *file descriptors* (*fd*'s) instead of multi-channel descriptors. The VPI routines will also be enhanced to support *fd*'s, and will be able to share *fd*'s with the Verilog HDL models.

4.9 Reading and using simulation times

The VPI routines can retrieve the current simulation time and information about the time scaling applied to the model. The VPI routines use two terms when referring to simulation time:

- **Module time units** are the units of time within a specific Verilog module. The `'timescale` directive in Verilog indicates what time units are used in the modules which follow the directive. The time scale also indicates a time precision, which is how many decimal points of accuracy are permitted. Time values with more decimal points than the precision are rounded off. Each module in a design can have a different time scale, so one model can specify delays in nanoseconds, with two decimal points of precision, and another model can specify delays in microseconds, with three decimal points.
- **Simulation time units** are the units of time used internally within the simulator. The simulator scales the times in all modules to the simulation time units. Most Verilog simulators will determine the finest precision of all modules in a simulation, and set that precision as the simulation time units.

Within a Verilog HDL model, time can be specified either as an integer or as a floating point real number. Within a Verilog simulation, however, time is represented as a 64 bit unsigned integer. As a Verilog model is compiled or loaded into simulation, time values in the module are rounded off to the specified precision and then scaled to the simulation time units. As an example:

```
'timescale 1ns/10ps
module A;
    ...
    nand #5.63 n1 (y, a, b);
    ...
endmodule

'timescale 1us/100ns
module B;
    ...
    nand #3.581 n1 (y, a, b);
    ...
endmodule
```

In this example, a simulator will determine that ten-picosecond units of time is the finest precision used by all modules, and set that as the simulation time units. The simulator will then scale the 5.63 nanosecond delay in module A to 563 10-picosecond units, and it will scale the 3.581 microsecond delay in module B to 360,000 10-picosecond units (the rounding to the module's time precision occurs before the time value is scaled to the simulator's time units).

4.9.1 Reading the current simulation time

`vpi_get_time()` retrieves the current simulation time. The routine requires two inputs: an object handle and a pointer to an `s_vpi_time` structure.

void vpi_get_time (object, time)

vpiHandle object handle for an object.

p_vpi_time time pointer to an application-allocated *s_vpi_time* structure to receive time information.

The *vpi_get_time()* routine retrieves simulation time and places in an *s_vpi_time* structure. The time can be represented in the simulation time units or a module's time units. An object handle must be specified when time is to be retrieved in a module's time units. The *vpi_get_time()* routine will scale the time value to the time scale of the module containing the object.


The *s_vpi_time* structure can represent the simulation time as either a pair of C integers (to hold the 64-bit Verilog time value), or as a double-precision value. The *s_vpi_time* structure is defined in the *vpi_user.h* file, as follows:

```
typedef struct t_vpi_time
{
    int type;
    unsigned int high, low;
    double real;
} s_vpi_time, *p_vpi_time;
```

The **type** field in the structure is set by the PLI application to the constant: **vpiSimTime** or **vpiScaledRealTime**. This field determines whether the simulation time should be retrieved as a real number or as integers.

- If the type is set to **vpiScaledRealTime**, then the current simulation time is placed into the **real** field of the structure. The simulation time is returned in the time units of the module containing the reference object specified as an input to *vpi_get_time()*. If a NULL is specified for the reference object, the simulation time is returned in the simulator's time units.
- If the type is set to **vpiSimTime**, then the current simulation time is placed into the pair of C integers in the structure, so that the lower 32 bits of time are in the **low** integer, and the upper 32 bits of time are in the **high** integer. The simulation time is always returned in the simulator's time units. The object handle is not used and can be set to NULL.

When time is retrieved as a pair of C integers, the time will be stored in the structure, as shown in the following illustration:



TIP Using `vpiSimTime` is better for simulation performance. The Verilog language standard specifies that time is a 64-bit unsigned integer. Retrieving the simulation time as integers is faster than requiring the simulator convert time to a real number.

The following code fragment retrieves the current simulation time, scaled to a module's time scale, using `vpi_get_time()`. There are three basic steps required to retrieve the time:

1. Allocate an `s_vpi_time` structure.
2. Set the type field.
3. Call `vpi_get_time()` with an object handle and a pointer to the time structure as inputs.

```

vpiHandle module_handle;
s_vpi_time current_time;

/* get handle for the module (or any object in the module */

current_time.type = vpiScaledRealTime;
vpi_get_time(module_handle, &current_time);

vpi_printf("Current time is %f\n", current_time.real);

```

4.9.2 Reading time scale factors

The PLI application can also read the time scale factors of a module. The time units and time precision are properties of a module, and are accessed using `vpi_get()` for the properties of `vpiTimeUnit` and `vpiTimePrecision`. Both properties are integers which represents a unit of time, as shown in table 4-1.

The simulation time unit can also be accessed using the `vpiTimeUnit` and `vpiTimePrecision` properties, and using a NULL for the reference handle.

The time units and time precision are represented as the magnitude of 1 second, which is the exponent of 1 second times 10^n . For example, 1 nanosecond is 1 second times 10^{-9} , so the integer value used to represent nanoseconds is -9.

Time Unit or Time Precision	Unit of Time Represented
2	100 seconds (1×10^2)
1	10 seconds (1×10^1)
0	1 second (1×10^0)
-1	100 milliseconds (1×10^{-1})
-2	10 milliseconds (1×10^{-2})
-3	1 millisecond (1×10^{-3})
-4	100 microseconds (1×10^{-4})
-5	10 microseconds (1×10^{-5})
-6	1 microsecond (1×10^{-6})
-7	100 nanoseconds (1×10^{-7})
-8	10 nanoseconds (1×10^{-8})
-9	1 nanosecond (1×10^{-9})
-10	100 picoseconds (1×10^{-10})
-11	10 picoseconds (1×10^{-11})
-12	1 picosecond (1×10^{-12})
-13	100 femtoseconds (1×10^{-13})
-14	10 femtoseconds (1×10^{-14})
-15	1 femtosecond (1×10^{-15})

Table 4-1: Time unit values for `vpiTimeUnit` and `vpiTimePrecision`

4.10 User-defined invocation options

The PLI provides a means for application developers to create user-defined invocation options. This capability makes it possible to configure PLI applications or pass data to an application from the invocation command line of a Verilog simulator.

int vpi_get_vlog_info (info)

p_vpi_vlog_info info pointer to an application-allocated *s_vpi_vlog_info* structure to receive invocation information.

vpi_get_vlog_info() returns information about the Verilog simulator that is running the PLI application. The routine returns 1 (for true) if it was successful, and 0 (for false) if an error occurred in retrieving the simulator information. The information retrieved includes:

- The simulator product name.
- The simulator version.
- The number of invocation options with which simulation was invoked.
- The values of the invocation options.

The input to *vpi_get_vlog_info()* is a pointer to an *s_vpi_vlog_info* structure which has been allocated by the PLI application. The routine will fill in the fields of the structure. Note that for string values, such as the product name, a pointer to the string is placed into the structure. The string is stored in a temporary buffer which is allocated and maintained by the simulator. The structure is defined in the *vpi_user.h* file, as follows:

```
typedef struct t_vpi_vlog_info
{
    int argc;
    char **argv;
    char *product;
    char *version;
} s_vpi_vlog_info, *p_vpi_vlog_info;
```

The **product** and **version** strings are defined by the simulator, and will be different for each simulator product.

The **argc** and **argv** values are the same command line values defined in the C language. **argc** is the number of invocation command arguments, and **argv** is a pointer to an array of strings, where each string is one argument from the command line.

The ability to check the simulator command line options makes it possible to create user-defined invocation options. Two applications for user-defined options are:

- To enable debug messages when debugging a PLI application. A verbose mode invocation option could be used to specify that debug messages should be printed.
- To pass file names or other values from the command line to a PLI application.

Parsing the -f command file invocation option

The IEEE 1364-1995 Verilog standard does not specify any invocation options for Verilog simulators. Every Verilog simulator, however, has adopted a small number of de facto standard invocation options. One of these is the **-f** option. This invocation option specifies that the file name which follows the option contains additional command line invocation arguments.

When the `argv` value is **-f**, the next `argv` will be a pointer to a NULL terminated array of pointers to strings. Element 0 in the array will contain the name of the file specified with the **-f** option, and the remaining elements in the array will be invocation commands contained in the file (comments are not included).

For example, assume a command file named **run.f** contained the following:

```
my_chip.v  
my_test.v  
+my_debug
```

If simulation were invoked with the command:

```
verilog -f run.f -s
```

Then `argv` would point to the following array of strings:

```
argv[0] -> "verilog"  
[1] -> "-f"  
[2] -----> [0] -> "run.f"  
[1] -> "my_chip.v"  
[2] -> "my_test"  
[3] -> "+my_debug"  
[5] -> NULL  
[3] -> "-s"
```

Example 4-6 illustrates a PLI application called `$test_invoke_options`. This application uses `vpi_get_vlog_info()` to test to see if simulation was invoked with any

user-specified option on the operating system command line or in a simulation command file, and return a true/false result. An example of using this applications is:

```
if ($test_invoke_options("+verbose"))
    $monitor("data_bus = %b", data_bus);
```

To show how invocation options are retrieved, this example also prints out all invocation options with which a simulation was invoked.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.04/invite_options_vpi.c
- Verilog test bench: Chapter.04/invite_options_test.v
- Verilog-XL results log: Chapter.04/invite_options_vpi.log

Example 4-6: `$test_invoke_options` — testing for user-defined invocation options

```
#define PLIbook_verbose 1 /* uncomment to list all invoke options */

#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include "vpi_user.h"     /* IEEE 1364 PLI VPI routine library */

/* prototypes of routines in this PLI application */
int PLIbook_TestInvokeOptions_calltf(),
    PLIbook_TestInvokeOptions_completf(),
    PLIbook_TestInvokeOptions_sizetf(),
    PLIbook_GetOptions();
void PLIbook_ScanCommandFile();

/********************* sizetf application *****/
int PLIbook_TestInvokeOptions_sizetf(char *user_data)
{
    return(1); /* $test_invoke_options returns 1-bit value */
}

/********************* calltf application *****/
int PLIbook_TestInvokeOptions_calltf(char *user_data)
{
    vpiHandle      systf_h, arg_itr, arg_h;
    char          *option_name;
    s_vpi_value   value_s;
    s_vpi_vlog_info sim_info;
    int           found;
```

```
vpi_get_vlog_info(&sim_info);

/* get system function arg--completetf already verified correctness */
systf_h = vpi_handle(vpiSysTfCall, NULL);
arg_itr = vpi_iterate(vpiArgument, systf_h);
arg_h   = vpi_scan(arg_itr);
vpi_free_object(arg_itr); /* free iterator -- did not scan to null */

/* read target option name from first tfarg */
value_s.format = vpiStringVal;
vpi_get_value(arg_h, &value_s);
option_name = value_s.value.str;

/* test for target option and return true/false to system function */
found = PLIbook_GetOptions(option_name,sim_info.argvc,sim_info.argv);
value_s.format = vpiIntVal;
value_s.value.integer = found;
vpi_put_value(systf_h, &value_s, NULL, vpiNoDelay);
return(0);
}

int PLIbook_optfound = 0; /* global variable for option found flag */
int PLIbook_indent = 0;   /* global variable to format text indenting */

int PLIbook_GetOptions(char *option, int argc, char **argv)
{
    int i;
    PLIbook_optfound = 0;
    PLIbook_indent = 0;
    for (i=0; i<argc; i++) {
        #ifdef PLIbook_verbose
            vpi_printf("%s\n", *argv);
        #endif
        if (strcmp(*argv, option) == 0) {
            PLIbook_optfound = 1;
        }
        if (strcmp(*argv, "-f") == 0) {
            argv++; /* next arg is address to array of strings */
            i++;
            PLIbook_ScanCommandFile(option, (char **)argv);
        }
        argv++; /* increment to next argument */
    }
    return(PLIbook_optfound);
}

void PLIbook_ScanCommandFile(char *option, char **arg)
{
    int i;

    #ifdef PLIbook_verbose
        PLIbook_indent += 4; /* increase text indentation */
    #endif
}
```

```
#endif
while ( *arg != NULL ) { /* loop until null termination */
    #ifdef PLIbook_verbose
        for (i=0; i<=PLIbook_indent; i++)
            vpi_printf(" ");
        vpi_printf("%s\n", *arg);
    #endif
    if (strcmp(*arg, option) == 0) {
        PLIbook_optfound = 1;
    }
    if (strcmp(*arg, "-f") == 0) {
        arg++; /* next arg is address to array of strings */
        PLIbook_ScanCommandFile(option, (char **) *arg);
    }
    arg++;
}
#endif
PLIbook_indent -= 4; /* decrease text indentation */
#endif
return;
}
```

4.11 Comparing VPI handles

Handles for Verilog objects can be obtained several different ways. Occasionally, a PLI application might need to test to see if two handles reference the same object.

int vpi_compare_objects (object1, object2)

vpiHandle object1 handle for an object.

vpiHandle object2 handle for an object.

vpi_compare_objects() returns a 1 (for true) if two handles reference the same object, and a 0 (for false) if they do not. A VPI handle is an abstraction used to reference an object within simulation, and all VPI routines use this abstraction to access information about the object. This layer of abstraction allows PLI applications to be portable to any number of Verilog simulators, because the abstract handle is a layer between the PLI application and the internal data structures of the simulator. Since the handle for an object is an abstraction, it is not possible to determine handle equivalence using the C ‘==’ operator.

4.12 Summary

The VPI routines in the PLI standard provide a simple, yet powerful, way to access what is happening in a Verilog simulation. The VPI routines provide this access by treating Verilog HDL constructs as objects, and using *handles* to reference these objects.

The VPI object diagrams serve as the primary documentation on the relationships of one object to another. The diagrams also document every property that is available for each object, such as names and logic values. Using the object diagrams, it is possible to traverse from any point in a Verilog design hierarchy to any other point in the hierarchy.

The next chapter will present the VPI routines which are used to read and modify logic values and delay values. Chapter 6 presents synchronizing PLI applications with activity occurring within a simulation. These chapters include several complete examples of PLI applications which use the VPI library of routines.

CHAPTER 5

Reading and Modifying Values Using VPI Routines

This chapter applies many of the concepts presented in the previous chapter, and adds the concepts of reading and modifying simulation values using the VPI routines.

The concepts presented in this chapter are:

- Objects which have logic values
- Reading logic values from Verilog simulations
- Writing logic values into Verilog simulations
- Objects which have delay values
- Reading delay values from Verilog simulations
- Writing delay values into Verilog simulations

5.1 Accessing objects which have logic values

The VPI routines can read the values of several different types of objects, and can modify the values of certain object types. In order to read an object's logic value or write a new value into an object, the object must have a value property in the VPI object diagrams. The objects which have this property are:

- Any ***net*** data type: scalar, vector, part-selects and bit-selects of vectors
- The ***reg*** data type: scalar, vector, part-selects and bit-selects of vectors
- The ***integer*** and ***time*** variable data types

- The *real* and *realtime* variable data types
- A *memory* word select
- A *variable array* word select
- A *parameter* constant (read only)
- A *specparam* constant, including a *specparam attribute* (read only)
- A literal *integer* value (read only)
- A literal *real* value (read only)
- A *string* value (read only)
- A *function call*
- A *system function call*
- A *sequential user-defined primitive*
- A *user-defined primitive table entry* (read only)

To access the value of an object, a handle for the object must be obtained. The routines to obtain handles were presented in the previous chapter.

5.2 Reading object logic values

Using the handle of an object, a PLI application can read the logic value property of the object. A single VPI routine is used to read the logic value of any type of object and any data type of logic value.

5.2.1 Working with a 4-logic value, multiple strength level system

The Verilog HDL supports 4 logic values, **0**, **1**, **z** and **x** and multiple levels of signal strength. There are also two ambiguous logic values, represented by **L** (low) and **H** (high), and many ambiguous strength values. The C programming language does not directly represent the same information, and so the VPI routines provide several ways to automatically translate values between Verilog and C. The translation converts Verilog 4-state logic into the following C types:

- **A C integer:** Verilog scalar and vector logic values are converted to a single C integer. The Verilog 4-state logic is converted to 2-state logic values of 0 and 1. Logic values of z and x values are converted to 0, and strength levels are ignored.
- **A C double:** Verilog real number values, scalar values and vector values are converted to a C double precision value. Real numbers in Verilog store 2-state decimal values, which convert directly to C doubles. Scalar and vector values in Verilog use

4-state logic, which is converted to 2-state logic. Strength levels are ignored.

- **A C string:** Verilog scalar and vector logic values are converted to a C character string. The Verilog 4-state logic is converted to the letters “0”, “1”, “z” and “x”. Logic strength levels are ignored.
- **A C constant:** Verilog scalar values are converted to a C integer constants. The Verilog 4-state logic is converted to the constants **vpi0**, **vpi1**, **vpiZ**, **vpiX**, **vpiH**, **vpiL**, or **vpiDontCare**. Logic strength levels are ignored.
- **A C aval/bval structure:** Verilog scalar and vector logic values are converted to a C structure which encodes each bit of a Verilog 4-state value to a pair of bits in C, referred to as an *aval/bval* pair. An array of *aval/bval* pairs is used to encode Verilog vectors of any size. The logic strength levels are ignored.
- **A C strength structure:** Verilog scalar logic values are converted to a C structure. The logic value is represented as a C constant of **vpi0**, **vpi1**, **vpiZ** or **vpiX**. The logic strengths as two C integer values, which represent the strength encoding defined in the Verilog HDL standard.



The format in which a value is read can have an impact on the run-time performance of a PLI application. The fastest run-time performance will be achieved when a value is retrieved in a format closest to the format in which a value is saved in the simulation structure. For Verilog scalar and vector nets and regs, this format is the *aval/bval* pair. For Verilog integers, C integers are most efficient, and for Verilog reals, C doubles are most efficient. The least efficient method for run-time performance is to retrieve a logic value as a C string.

5.2.2 The `vpi_get_value()` routine

The VPI routine which reads logic values is **`vpi_get_value()`**. The routine automatically converts Verilog logic values to C representations. The syntax for this routine is:

`void vpi_get_value (expr, value)`

<code>vpiHandle expr</code>	handle for an object which has a value property.
<code>p_vpi_value value</code>	pointer to an application-allocated <code>s_vpi_value</code> structure to receive value information.

The object’s logic value is retrieved into an **`s_vpi_value`** structure. The PLI application must allocate this structure, and pass a pointer to the structure as an input to **`vpi_get_value()`**.

Following are two ways of allocating an `s_vpi_value` structure:

To allocate local, automatic storage which will be freed when the PLI applications exits:

```
s_vpi_value arg_info;
vpi_get_value(tfarg_handle, &arg_info);
```

To allocate persistent storage which can be used by future calls to the PLI application (a pointer to the storage must be maintained by the application):

```
p_vpi_value arg_info;
arg_info = (p_vpi_value)malloc(sizeof(s_vpi_value));
vpi_get_value(tfarg_handle, arg_info);
```

NOTE

An `s_vpi_value` structure can also be statically allocated, which will create storage which remains in effect throughout simulation. However, caution should be observed when using static variables in PLI applications. All instances of a system task or system function will share the same static variable. When different storage is needed for each instance of a system task, it is better to malloc a separate structure for each instance. Section 6.2 on page 184 in Chapter 6 discusses allocation of instance specific storage.

The `s_vpi_value` structure is defined in `vpi_user.h`, and is listed below.

```
typedef struct t_vpi_value {
    int format;
    union {
        char *str;                                /* string value */
        int scalar;                               /* scalar value */
        double real;                             /* real value */
        struct t_vpi_time *time;                /* time value */
        struct t_vpi_vecval *vector;             /* vector value */
        struct t_vpi_strengthval *strength; /* strength val */
        char *misc;                            /* reserved */
    } value;
} s_vpi_value, *p_vpi_value;
```

This `s_vpi_value` structure has two fields, `format` and `value`. The `format` field controls what C language data type should be used to receive the value, and the `value` field receives the value of the object. The definitions of the `format` values and the `value` union fields are listed in the following table.

Format	Definition
vpiBinStrVal	represents a Verilog logic value as a C string, using binary numbers, with 4-state logic
vpiOctStrVal	represents a Verilog logic value as a C string, using octal numbers, with 4-state logic
vpiDecStrVal	represents a Verilog logic value as a C string, using decimal numbers, with 4-state logic
vpiHexStrVal	represents a Verilog logic value as a C string, using hexadecimal numbers, with 4-state logic
vpiScalarVal	represents a Verilog scalar logic value as C constants, which represents Verilog 4-state logic
vpiStrengthVal	represents a Verilog scalar value as a VPI strength structure containing the logic value and Verilog strength levels
vpiVectorVal	represents a Verilog vector value as an array of VPI aval/bval structures, encoded to represent Verilog 4-state logic
vpiIntVal	represents a Verilog logic value as a C integer, with 2-state logic
vpiTimeVal	represents a Verilog time value as a VPI time structure
vpiRealVal	represents a Verilog logic value as a C double precision number, with 2-state logic
vpiStringVal	represents a Verilog string value as a C string
vpiObjTypeVal	allows simulation to pick the best way to represent a value

Table 5-1: The s_vpi_value structure format constants

The value field of the s_vpi_value structure is a union of C data types. Which C data type is used to receive the Verilog value is controlled by the format field of the structure:

- The **scalar** field is used if the format field is vpiScalarVal. The vpi_get_value() routine will fetch the object's logic value, and return one of the constants: **vpi0**, **vpi1**, **vpiZ**, **vpiX**, **vpiH**, **vpiL**, or **vpiDontCare**. The latter constant is used when reading the values of UDP tables.
- The **integer** field is used if the format field is vpiIntVal. The vpi_get_value() routine will fetch the object's logic value, convert the value into 2-state logic (logic X and Z are converted to 0), and return the value as a C integer. The maximum number of bits which can be retrieved using this format is limited by the size of a C integer. If the number of bits of the Verilog object is

greater than the size of an integer, the most significant bits (the left-most bits) of the object are truncated. If the number of bits of the Verilog object is less than the size of an integer, the value is retrieved into the right-most bits of the integer, and the left-most bits are zero filled. Refer to section 5.2.3 on page 145 for more details on reading Verilog integer values.

- The **real** field is used if the format field is `vpiRealVal`. The `vpi_get_value()` routine will fetch the object's logic value, convert the value into 2-state logic (logic X and Z are converted to 0), and return the value as a C double. Refer to section 5.2.4 on page 146 for more details on reading Verilog real values.
- The **str** field in the value union is used if the format field is `vpiBinStrVal`, `vpiOctStrVal`, `vpiDecStrVal`, `vpiHexStrVal`, or `vpiStringVal`. The `vpi_get_value()` routine will fetch the object's logic value, convert the value into an ASCII string (which can represent all Verilog logic values, including X and Z), and return a pointer to the string. Note that the string will be stored in a temporary buffer, and may be overwritten by other VPI routines which return pointers to strings. Refer to section 5.2.5 on page 148 for more details on reading Verilog string values as C strings, and to section 5.2.6 on page 149 for more details on reading Verilog logic values as C strings.
- The **vector** field is used if the format field is `vpiVectorVal`. The `vpi_get_value()` routine will fetch the object's logic value into an array of `s_vpi_vecval` structures. Refer to section 5.2.7 on page 150 for more details on reading Verilog vector values.
- The **strength** field is used if the format field is `vpiStrengthVal`. The `vpi_get_value()` routine will fetch the object's logic value into an `s_vpi_strengthval` structure. Refer to section 5.2.8 on page 155 for more details on reading Verilog strength values.
- The **time** field is used if the format field is `vpiTimeVal`. The `vpi_get_value()` routine will fetch the object's logic value into an `s_vpi_time` structure. Refer to section 5.2.9 on page 158 for more details on reading Verilog time variable values.



NOTE When reading logic values, the memory for any time, strength or vector structures required to store the value is allocated and maintained by the simulator. The simulator will fill in the value, and place a pointer to the memory in the appropriate field of the value union in the `s_vpi_value` structure.

The storage allocated by the simulator is temporary, and will only remain valid until the next call to `vpi_get_value()` or until the PLI application exits.

If the PLI application needs to preserve the value, the application must allocate memory for its own storage, and copy the values from the simulator's temporary storage into the user-allocated storage.

5.2.3 Reading 2-state logic as C integers

The `vpi_get_value()` can be set to read Verilog values into a C integer by setting the format flag in the `s_vpi_value` structure to `vpiIntVal`. Verilog 4-state logic will be converted into C 2-state logic by converting logic Z and X to 0. As a standard data type in C, the value can be easily manipulated if needed.

The Verilog value which is read can be any Verilog data type, and it will be converted to a C integer. Verilog real values will be rounded to an integer.

NOTE An integer value in Verilog can be a 1-bit scalar value or a vector of any bit size. The maximum value which can be read is constrained by the size of a C integer. If a Verilog vector is wider than the a C integer, then the left-most bits (the most significant bits) of the Verilog vector are truncated.

The steps to read the value of an object as a C integer are:

1. Allocate an `s_vpi_value` structure.
2. Set the `format` field in the structure to `vpiIntVal`.
3. Call `vpi_get_value()`, giving a pointer to the `s_vpi_value` structure as an input, along with a handle for the object from which to read the logic value.
4. Read the logic value of the object from the `value.integer` field of the `s_vpi_value` structure.

Example 5-1, on the following page, lists a C function called `PLibook_getarg_intval_vpi()`. This routine reads the logic value of a Verilog system task/function argument as an integer value. The input provided to the routine is the index number of the task/function argument, where the left-most argument is number 1, and the index number increases from left to right. Note that this example calls another utility presented earlier in this book, `PLibook_getarg_handle_vpi()`, to obtain a handle for a system task/function argument. This utility was listed in example 4-2 on page 115 in Chapter 4.

CD The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.05/vpi_utilities.c`
- Verilog test bench: `Chapter.05/vpi_utilities_test.v`
- Verilog-XL results log: `Chapter.05/vpi_utilities_test.log`

Example 5-1: *PLIbook_getarg_intval_vpi()* — reading Verilog values as C integers

```

int PLIbook_getarg_intval_vpi(int argNum)
{
    vpiHandle arg_h;
    s_vpi_value argVal;
    #ifdef PLIbookDebug
        s_vpi_error_info err; /* structure for error handling */
    #endif

    arg_h = PLIbook_getarg_handle_vpi(argNum);
    if (arg_h == NULL) {
        vpi_printf("ERROR: PLIbook_getarg_intval_vpi() could not obtain arg
handle\n");
        return(0);
    }
    argVal.format = vpiIntVal;
    vpi_get_value(arg_h, &argVal);
    #ifdef PLIbookDebug /* if error, generate verbose debug message */
        if (vpi_chk_error(&err)) {
            vpi_printf("ERROR: PLIbook_getarg_intval_vpi() could not obtain
arg value\n");
            vpi_printf("File %s, Line %d: %s\n",
                       err.file, err.line, err.message);
            return(0);
        }
    #endif
    return(argVal.value.integer);
}

```

5.2.4 Reading 2-state logic as C doubles

The `vpi_get_value()` reads Verilog values into a C double by setting the format flag in the `s_vpi_value` structure to `vpiRealVal`. Verilog 4-state logic will be converted into C 2-state logic by converting logic Z and X to 0. The Verilog value which is read can be any Verilog data type, and it will be converted to a C double.

The steps to read the value of an object as a C double are:

1. Allocate an `s_vpi_value` structure.
2. Set the `format` field in the structure to `vpiRealVal`.
3. Call `vpi_get_value()`, giving a pointer to the `s_vpi_value` structure as an input, along with a handle for the object from which to read the logic value.
4. Read the logic value of the object from the `value.real` field of the `s_vpi_value` structure.

Example 5-2 lists a C function called *PLIbook_getarg_realval_vpi()*, which reads the logic value of a system task/function argument as a C double. The input provided to the routine is the index number of the task/function argument, where the left-most argument is number 1, and the index number increases from left to right. This example calls another utility presented earlier in this book, *PLIbook_getarg_handle_vpi()*, to obtain a handle for a system task/function argument. This utility was listed in example 4-2 on page 115 in Chapter 4.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/vpi_utilities.c
- Verilog test bench: Chapter.05/vpi_utilities_test.v
- Verilog-XL results log: Chapter.05/vpi_utilities_test.log

Example 5-2: *PLIbook_getarg_realval_vpi()* — reading Verilog values as C doubles

```
double PLIbook_getarg_realval_vpi(int argNum)
{
    vpiHandle arg_h;
    s_vpi_value argVal;
    #ifdef PLIbookDebug
        s_vpi_error_info err; /* structure for error handling */
    #endif

    arg_h = PLIbook_getarg_handle_vpi(argNum);
    if (arg_h == NULL) {
        vpi_printf("ERROR: PLIbook_getarg_realval_vpi() could not obtain arg
handle\n");
        return(0);
    }
    argVal.format = vpiRealVal;
    vpi_get_value(arg_h, &argVal);
    #ifdef PLIbookDebug /* if error, generate verbose debug message */
        if (vpi_chk_error(&err)) {
            vpi_printf("ERROR: PLIbook_getarg_realval_vpi() could not obtain
arg value\n");
            vpi_printf("File %s, Line %d: %s\n",
                       err.file, err.line, err.message);
            return(0.0);
        }
    #endif

    return(argVal.value.real);
}
```

5.2.5 Reading Verilog string values into C strings

The `vpi_get_value()` reads Verilog strings into a C string by setting the format flag in the `s_vpi_value` structure to `vpiStringVal`. A pointer to the string will be placed in the `value.str` field of the `s_vpi_value` structure. The steps for reading strings is the same as for reading integer and double values, except that the format specified is `vpiStringVal`. The following example lists another utility application, `PLIbook_getarg_stringval_vpi()`. This utility reads the string value of a system task/function argument as a C string.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/vpi_utilities.c
- Verilog test bench: Chapter.05/vpi_utilities_test.v
- Verilog-XL results log: Chapter.05/vpi_utilities_test.log

Example 5-3: `PLIbook_getarg_stringval_vpi()` — reading Verilog values as C strings

```
char *PLIbook_getarg_stringval_vpi(int argNum)
{
    vpiHandle arg_h;
    s_vpi_value argVal;
    #ifdef PLIbookDebug
        s_vpi_error_info err; /* structure for error handling */
    #endif

    arg_h = PLIbook_getarg_handle_vpi(argNum);
    if (arg_h == NULL) {
        vpi_printf("ERROR: PLIbook_getarg_stringval_vpi() could not obtain
arg handle\n");
        return(0);
    }
    argVal.format = vpiStringVal;
    vpi_get_value(arg_h, &argVal);
    #ifdef PLIbookDebug /* if error, generate verbose debug message */
        if (vpi_chk_error(&err)) {
            vpi_printf("ERROR: PLIbook_getarg_stringval_vpi() could not obtain
arg value\n");
            vpi_printf("File %s, Line %d: %s\n",
                       err.file, err.line, err.message);
            return("");
        }
    #endif

    return(argVal.value.str);
}
```

5.2.6 Reading 4-state logic as C strings

The `vpi_get_value()` can read any Verilog logic value as a C string by setting the format flag in the `s_vpi_value` structure to one of the constants: `vpiBinStrVal`, `vpiOctStrVal`, `vpiDecStrVal` or `vpiHexStrVal`. The easiest method of representing Verilog 4-state logic in the C language is to use C character strings. The logic values Z and X are represented with the characters "Z" and "X". As a character string, Verilog vectors of any width can be represented. The values can be represented in binary, octal, decimal or hexadecimal by setting `format` field to the appropriate constant in the `s_vpi_value` structure. The C string functions contained in standard C libraries, such as `strings.h`, can be used to manipulate the string values.



TIP Reading a Verilog logic value is a simple way to represent 4-state logic in a printable form in the C language, and a simple way to read Verilog vectors of any bit width. However, the conversion from a Verilog logic value to a C string is very expensive for simulation run-time performance. The best simulation performance will be achieved when Verilog values are read into the C format that most closely resembles how the value is stored in a Verilog simulation.

Example 5-4 finds all the nets in a Verilog module and reads each net's logic value as a binary string value. A module instance name is passed to the routine as the first system task/function argument.



CD The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/show_all_nets_vpi.c
- Verilog test bench: Chapter.05/show_all_nets_test.v
- Verilog-XL results log: Chapter.05/show_all_nets_test.log

Example 5-4: using `vpi_get_value()` to read 4-state values as a C string

```
int PLIbook_ShowNets_calltf(char *user_data)
{
    vpiHandle    systf_handle, arg_iterator, module_handle,
                net_iterator, net_handle;
    s_vpi_time   current_time;
    s_vpi_value  current_value;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument */
    /* compiletf has already verified only 1 arg with correct type */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    module_handle = vpi_scan(arg_iterator);
    vpi_free_object(arg_iterator); /* free iterator memory */
```

```

/* read current simulation time */
current_time.type = vpiScaledRealTime;
vpi_get_time(systf_handle, &current_time);

vpi_printf("\nAt time %2.2f, nets in module %s (%s):\n",
           current_time.real,
           vpi_get_str(vpiFullName, module_handle),
           vpi_get_str(vpiDefName, module_handle));

/* obtain handles to nets in module and read current value */
net_iterator = vpi_iterate(vpiNet, module_handle);
if (net_iterator == NULL)
    vpi_printf(" no nets found in this module\n");
else {
    current_value.format = vpiBinStrVal; /* read values as a string */
    while ( (net_handle = vpi_scan(net_iterator)) != NULL ) {
        vpi_get_value(net_handle, &current_value);
        vpi_printf(" net %-10s value is %s (binary)\n",
                   vpi_get_str(vpiName, net_handle),
                   current_value.value.str);
    }
}
return(0);
}

```

5.2.7 Reading Verilog 4-state logic vectors as encoded aval/bval pairs


The `vpi_get_value()` routine with a `vpiVectorVal` format will retrieve an object's 4-state logic value as an encoded pair of C integers. The encoding uses an *aval/bval* pair of C integers to represent the 4 logic values of Verilog, with one bit of each *aval/bval* pair representing one bit of a Verilog logic value. The encoding is:

aval/bval pair	Verilog logic value represented
0/0	0
1/0	1
0/1	Z
1/1	X

Table 5-2: aval/bval logic value encoding

The VPI library assumes C integers are 32-bits wide, and uses the aval/bval pair to encode up to 32 bits of a Verilog vector. By using an array of aval/bval integer pairs, vector lengths of any size may be represented. The representation of a 40-bit vector in Verilog can be visualized as:

For the Verilog declaration: `reg [39:0] data;`



The Verilog language supports any numbering convention for a vector's bit numbers. The least significant bit of the Verilog vector can be the smallest bit number, such as bit 0 (which is referred to as little endian convention). Or, the least significant bit of the Verilog vector can be the largest bit number, such as bit 39 (which is referred to as big endian convention). Verilog does not require there be a bit zero at all. Each of the following examples are valid vector declarations in Verilog:


```

reg [39:0] data;      /* little endian -- LSB is bit 0 */
reg [0:39] data2;    /* big endian     -- LSB is bit 39 */
reg [40:1] data3;    /* little endian -- LSB is bit 1 */

```

The bit numbering used in Verilog does not affect the aval/bval representation of the Verilog vector. In the array of aval/bval pairs, the LSB of the Verilog vector will always be the LSB of the first C integer in the array, and the MSB of the Verilog vector will always be the last bit in the array which is used. The following diagram illustrates the aval/bval array for a Verilog vector declared with a big endian convention.

For the Verilog declaration: `reg [1:40] data;`



lsb = least significant bit

msb = most significant bit

The aval/bval pair is declared in an `s_vpi_vecval` structure, which is defined in the `vpi_user.h` VPI library. The structure definition is:

```
typedef struct t_vpi_vecval
{
    int aval, bval;
} s_vpi_vecval, *p_vpi_vecval;
```

To read a task/function argument's 4-state logic value using `vpi_get_value()` involves four basic steps:

1. Allocate an `s_vpi_value` structure.
2. Set the format field in the structure to `vpiVectorVal`.
3. Call `vpi_get_value()`, giving a pointer to the `s_vpi_value` structure as an input, along with a handle for the object from which to read the logic value.
4. Read the logic value of the object from the `value.vector` field of the `s_vpi_value` structure.

The `value.vector` field is a pointer to an array of `s_vpi_vecval` structures, with each element in the array representing 32 bits of a Verilog vector. The formula for determining the number of elements is:

```
number_of_array_elements = ((vector_size - 1) / 32 + 1);
```

The vector size is a property of the Verilog vector object, and can be accessed using:

```
vector_size = vpi_get(vpiSize, vector_handle);
```

Once the number of elements are known, the value of each 32-bit group of the Verilog vector can be accessed by reading the aval/bval pair of each `s_vpi_vecval` structure in the array. Within each aval/bval pair, an individual bit of the Verilog vector can be accessed by masking out all of the other bits in the aval/bval pair.



NOTE The array of `s_vpi_vecval` structures is allocated by the simulator, and is only guaranteed to be valid until the next call to `vpi_get_value()`, or until the PLI application exits. If the values need to be maintained, the PLI application must allocate its own array of structures and copy the values from the simulator's storage.

Example 5-5 illustrates using `vpi_get_value()` to read the aval/bval encoded 4-state value of a vector. The example assumes the vector is passed to the PLI application as the first system task/function argument. The value of the vector is then printed one bit at a time. For simplicity in this example, it is assumed that the least-significant bit of the Verilog vector is bit 0 (little endian).



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/read_vecval_vpi.c
- Verilog test bench: Chapter.05/read_vecval_test.v
- Verilog-XL results log: Chapter.05/read_vecval_test.log

Example 5-5: \$read_vecval — reading 4-state vector values as aval/bval pairs

```

int PLIbook_ReadVecVal_calltf(char *user_data)
{
    vpiHandle    systf_handle, arg_iterator, arg_handle, vector_h;
    s_vpi_value vector_val;      /* structure to receive vector value */
    int i, vector_size, array_size, avalbit, bvalbit, bit_num;
    char vlogval;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument
       compiletf has already verified only 1 arg with correct type */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    vector_h = vpi_scan(arg_iterator);
    vpi_free_object(arg_iterator); /* free iterator memory */

    vector_size = vpi_get(vpiSize, vector_h); /* determine number of...*/
    array_size  = ((vector_size-1) / 32 + 1); /* ...elements in array */

    vector_val.format = vpiVectorVal;           /* set value format field */

    vpi_get_value(vector_h, &vector_val); /* read vector's logic value */

    vpi_printf("\nVector %s encoded value:\n",
              vpi_get_str(vpiName, vector_h));
    for (i=0; i<array_size; i++) {
        /* the following loop assumes the Verilog LSB is bit 0 */
        for (bit_num=0; bit_num<=31; bit_num++) {
            avalbit=PLIbook_getbit(vector_val.value.vector[i].aval, bit_num);
            bvalbit=PLIbook_getbit(vector_val.value.vector[i].bval, bit_num);
            vlogval=PLIbook_get_4state_val(avalbit, bvalbit);
            vpi_printf(" bit[%2d]  aval/bval = %d/%d  4-state value = %c\n",
                      (i*32+bit_num), avalbit, bvalbit, vlogval);
            /* quit when reach last bit of Verilog vector */
            if ((i*32+bit_num) == vector_size-1) break;
        }
    }
    return(0);
}

*****
* Function to determine if a specific bit is set in a 32-bit word.
* Sets the least-significant bit of a mask value to 1 and shifts the
* mask left to the desired bit number.
*****
int PLIbook_getbit(int word, int bit_num)
{
    int mask;
    mask = 0x00000001 << bit_num;
    return((word & mask)? TRUE: FALSE);
}

```

```
*****
 * Function to convert aval/bval encoding to 4-state logic represented
 * as a C character.
 ****
char PLIbook_get_4state_val(int aval, int bval)
{
    if      (!bval && !aval) return('0');
    else if (!bval &&  aval) return('1');
    else if ( bval && !aval) return('z');
    else
        return('x');
}
```

5.2.8 Representing Verilog strength values in C

The `vpi_get_value()` routine with a `vpiStrengthVal` format can be used to retrieve the 4-state logic value and strength level of a scalar net object or a bit select of a vector net object. The information is retrieved into an `s_vpi_strengthval` structure. The memory for the structure is allocated by the simulator, and a pointer to the memory is returned in the `value.strength` field of the `s_vpi_value` structure. The structure definition is:

```
typedef struct t_vpi_strengthval
{
    int logic;
    int s0, s1;
} s_vpi_strengthval, *p_vpi_strengthval;
```


The logic value of the scalar signal is represented as one of the constants: `vpi0`, `vpi1`, `vpiZ` or `vpix`.

The strength level is represented as a pair of C integers which contain the logic 0 and logic 1 strength components of the net. The Verilog language has 8 named strength levels for logic zero and 8 strength levels for a logic one, plus a number of ambiguous strength levels. Each named strength level is represented by a Verilog keyword, as shown in Table 5-3, which follows:

Strength Level	Strength Name	Specification Keyword
7	Supply Drive	supply0 supply1
6	Strong Drive	strong0 strong1
5	Pull Drive	pull0 pull1
4	Large Capacitance	large
3	Weak Drive	weak0 weak1
2	Medium Capacitance	medium
1	Small Capacitance	small
0	High Impedance	highz0 highz1

Table 5-3: Verilog HDL strength levels and keywords

Within Verilog, the strength of a signal is stored as two 8-bit bytes, as shown in the diagram below:



The full details on how the two strength bytes represent Verilog logic strength are defined in the Verilog language, and are outside the scope of this book. Refer to the list of Verilog HDL books on page 7 for suggestions on where to find more information on Verilog HDL strengths.

In the PLI, the strength value returned by `vpi_get_value()` is represented by a pair of C integers with a value from 00 (hex) to FF (hex). The value indicates which bits of the corresponding Verilog strength byte are set (only 8 bits of each C integer are used). Note that both integers represent bit 0, the highz bit, as the least significant bit of the C integer.



NOTE The `s_vpi_strengthval` structure is allocated by the simulator, and is only guaranteed to be valid until the next call to `vpi_get_value()`, or until the PLI application exits. If the values need to be maintained, the PLI application must allocate its own structure, and copy the strength value from the simulator's storage.

The following example illustrates how to read the 4-state logic value and strength value of scalar nets using `vpi_get_value()`. The example assumes that a scalar Verilog net is passed in as the first system task/function argument.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/read_strengthval_vpi.c
- Verilog test bench: Chapter.05/read_strengthval_test.v
- Verilog-XL results log: Chapter.05/read_strengthval_test.log

Example 5-6: \$read_strengthval — reading logic and strength values

```
int PLIbook_ReadStrengthVal_calltf(char *user_data)
{
    vpiHandle systf_h, arg_itr, arg_h, net_h;
    s_vpi_value net_val; /* structure to receive net value */

    /* obtain a handle to the system task instance */
    systf_h = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument
       compiletf has already verified only 1 arg with correct type */
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    net_h = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* free iterator memory */

    net_val.format = vpiStrengthVal; /* set value format field */
    vpi_get_value(net_h, &net_val); /* read net's strength value */

    vpi_printf("\nNet %s: ", vpi_get_str(vpiName, net_h));
    vpi_printf("value=%c strength0=%2x(hex) strength1=%2x(hex)\n\n",
              net_val.value.strength->logic,
              net_val.value.strength->s0,
              net_val.value.strength->s1);

    return(0);
}
```

```
*****
 * Function to convert VPI logic constant to a character
 ****
char PLIbook_DecodeBitValue(int bit_constant)
{
    switch (bit_constant) {
        case vpi0:      return('0'); break;
        case vpi1:      return('1'); break;
        case vpiZ:      return('Z'); break;
        case vpiX:      return('X'); break;
        case vpiL:      return('L'); break;
        case vpiH:      return('H'); break;
        case vpiDontCare: return('?'); break;
        default:        return('U'); /* undefined value passed in */
    }
}
```

5.2.9 Reading Verilog time variables C

The Verilog language stores a time variable as a 64-bit unsigned integer. The `vpi_get_value()` routine, with a `vpiTimeVal` format, can be used to retrieve the value stored in a time variable. The information is retrieved into an `s_vpi_time` structure. The memory for the structure is allocated by the simulator, and a pointer to the memory is returned in the `value.time` field of the `s_vpi_value` structure.

The `s_vpi_time` structure is defined in the `vpi_user.h` file, as follows:

```
typedef struct t_vpi_time
{
    int type;
    unsigned int high, low;
    double real;
} s_vpi_time, *p_vpi_time;
```

When reading the value of a time variable, the 64-bit time value will be retrieved as two unsigned C integers, with the upper 32 bits of the variable stored in the `high` field, and the lower 32 bits of the variable stored in the `low` field of the structure. The `type` field and `real` field in the structure are not used when time variables are read.

NOTE → The `s_vpi_timeval` structure is allocated by the simulator, and is only guaranteed to be valid until the next call to `vpi_get_value()`, or until the PLI application exits. If the values need to be maintained, the PLI application must allocate its own structure, and copy the time value from the simulator's storage.

The basic steps required to retrieve the value are:

1. Allocate an `s_vpi_value` structure.
2. Set the `format` field to `vpiTimeVal` field.
3. Call `vpi_get_value()`, with an object handle and a pointer to the value structure as inputs.
4. Read the time value from the `value.time->high` and `value.time->low` fields of the time structure.

The calltf application shown in Example 5-7 retrieves the value of a time variable, using `vpi_get_value()`.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/read_timeval_vpi.c
- Verilog test bench: Chapter.05/read_timeval_test.v
- Verilog-XL results log: Chapter.05/read_timeval_test.log

Example 5-7: \$read_timeval — reading time variable values

```
int PLIBook_ReadTimeVal_calltf(char *user_data)
{
    vpiHandle systf_h, arg_itr, arg_h, timevar_h;
    s_vpi_value timevar_val; /* structure to receive variable value */

    /* obtain a handle to the system task instance */
    systf_h = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument
       compiletf has already verified only 1 arg with correct type */
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    timevar_h = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* free iterator memory */

    timevar_val.format = vpiTimeVal; /* set value format field */

    vpi_get_value(timevar_h, &timevar_val); /* read variable's value */

    vpi_printf("\nTime Variable %s: ", vpi_get_str(vpiName, timevar_h));
    vpi_printf(" hi-word/lo-word = %d/%d\n\n",
              timevar_val.value.time->high,
              timevar_val.value.time->low);
    return(0);
}
```

5.2.10 Reading Verilog values without determining the Verilog data type

The `vpi_get_value()` provides a convenient format mode that makes it possible to read any Verilog data type into C data types, without having to first determine the data type of the Verilog object. When the `format` field of the `s_vpi_value` structure is set to `vpiObjTypeVal`, the simulator will retrieve the Verilog data type into the C data type that most closely matches the Verilog type. The simulator automatically allocates whatever storage is needed, and places a pointer to the storage in the appropriate value field of the `s_vpi_value` structure. The simulator also sets the `format` field of the `s_vpi_value` structure with the corresponding format constant to indicate how the retrieved value is stored, overriding the `vpiObjTypeVal` set before `vpi_get_value()` was called.

The Verilog data types are converted to C types as follows:

- Verilog integer variables are converted using the `vpiIntVal` format.
- Verilog real variables are converted using the `vpiRealVal` format.
- Verilog strings are converted using the `vpiStringVal` format.
- Verilog time variables are converted using the `vpiRealVal` format.
- Verilog scalar reg or bit select of a vector reg or integer or time variable are converted using the `vpiScalarVal` format.
- Verilog scalar net or bit select of a vector net are converted using the `vpiStrengthVal` format.
- Verilog vectors are converted using the `vpiVecVal` format.

The automatic Verilog type recognition makes it easy for a PLI application to read any Verilog data type, by simply obtaining a handle to an object with a logic value, setting the format to `vpiObjTypeVal`, calling `vpi_get_value()`, and then examining to `format` field of the `s_vpi_value` structure, to determine how the value was retrieved.

 **NOTE** Any structures to receive the value are allocated by the simulator, and are only guaranteed to be valid until the next call to `vpi_get_value()`, or until the PLI application exits. If the values need to be maintained, the PLI application must allocate its own structures, and copy the value from the simulator's storage.

5.3 Writing values to Verilog objects

Writing a value to a Verilog object is simply a reverse process from reading a value. A single VPI routine is used to write values onto any type of object.

vpi_put_value() converts a value represented as C data type into Verilog 4-state logic, and writes the value to a Verilog object. The value to be written can be represented a variety of ways in the C language. These representations are the same as was described in section 5.2.1 on page 140 for reading values, which include representing Verilog logic values as a character string, as a scalar value represented using VPI constants, as an integer, as a double, as an array of *s_vpi_vecval* structures, as an *s_vpi_time* structure, or as an *s_vpi_strengthval* structure.

The syntax for *vpi_put_value()* is:

vpiHandle vpi_put_value (object, value, time, flag)

<i>vpiHandle object</i>	handle for an object which has a value property.
<i>p_vpi_value value</i>	pointer to an application-allocated <i>s_vpi_value</i> structure containing value information.
<i>p_vpi_time time</i>	pointer to an application-allocated <i>s_vpi_time</i> structure containing the propagation delay information.
<i>int flag</i>	constant representing the delay mode.

The *s_vpi_value* structure that is pointed to in the value field must be allocated by the PLI application, and the fields within the structure set to the value to be written to the object.

NOTE When writing logic values, the PLI application must allocate and maintain all storage elements required to represent the logic value in C. Pointers to this user-allocated storage will be passed to *vpi_put_value()*. This is exactly the opposite of *vpi_get_value()*, where the simulator allocates and maintains any necessary storage.

The third argument to *vpi_put_value()* is pointer to an *s_vpi_time* structure. This structure is allocated by the PLI application and set to a propagation delay value. The *vpi_put_value()* routine can write the value to the object immediately or using the simulator's event scheduling mechanism. By using the event scheduler, a value can be written to an object later in the current simulation time step, or scheduled to occur in any future simulation time step.

The last argument to `vpi_put_value()` is a propagation flag, indicating how the propagation delay should be applied, using the simulator's event scheduler. The flag is one of the following constants:

- **vpiNodelay** indicates no propagation delay is to be used. The object may be a Verilog reg, variable, memory word, sequential UDP or system function. When this flag is used, the propagation delay argument is not used and can be set to `NULL`.
- **vpiInertialDelay** indicates inertial delay propagation is to be used. Any pending events which are scheduled for the object are cancelled. The object may be a Verilog reg, variable, memory word or sequential UDP.
- **vpiPureTransportDelay** indicates transport delay propagation is to be used. Any pending events which are scheduled for the object remain scheduled (no events are cancelled). The object may be a Verilog reg, variable, memory word, variable array word or sequential UDP.
- **vpiTransportDelay** indicates a modified transport delay propagation is to be used. Any pending events are cancelled which are scheduled for the object at a later time than this new event. The object must be a Verilog reg, variable, memory word, variable array word or sequential UDP.
- **vpiForceFlag** indicates the value is to be forced onto the object, overriding any existing values. No other changes can occur on the object except another force until the force is released. No propagation delay is used, and the `time_p` argument may be set to `NULL`. The object may be any Verilog object which has a value property. Only one force value may exist for an object at a time. Setting a force value on an object will replace any existing force value, regardless of whether the force was set within the PLI or within the Verilog HDL.
- **vpiReleaseFlag** indicates any existing force on the object is to be released.
- **vpiReturnEvent** indicates that `vpi_put_value()` should return a handle for the scheduled event. Using this handle, the event can be removed from the simulator's event queue. This flag needs to be logically OR'ed with any of the flags `vpiInertialDelay`, `vpiTransportDelay` or `vpiPureTransportDelay`. If this flag is not set, `vpi_put_value()` will return `NULL`. A `NULL` is also returned if `vpi_put_value()` is called with no delay.
- **vpiCancelEvent** indicates an event should be removed from the simulator's queue. To use this flag, the handle passed to `vpi_put_value()` must be an event object which was returned by a previous call to `vpi_put_value()` which had used the `vpiReturnEvent` flag. It is not an error to attempt to cancel an event which has already transpired—the call to `vpi_put_value()` simply does nothing.

To write a value to an object requires the following steps:

1. Obtain a handle for an object—the object must have a value property, in order to write a value to the object.

2. Allocate an **s_vpi_value** structure.
3. Allocate an **s_vpi_time** structure (if a propagation delay will be used).
4. Allocate variables or structures to store the logic value to be written onto the object, using the appropriate C data type for the format of the value.
5. Load the value memory with the value to be written.
6. Set the appropriate field in the **value** union of the **s_vpi_value** structure to a pointer to the value.
7. Set the **format** field in the **s_vpi_value** structure to indicate how the logic value is represented in C.
8. Set the **type** field in the **s_vpi_time** structure to indicate how the delay value is represented. The delay type is one of the following VPI constants:
 - **vpiScaledRealTime** indicates the delay is represented as a C double. The delay value will be scaled to the time units and precision of the module containing the object onto which the value will be written.
 - **vpiSimTime** indicates the delay is represented as a pair of C integers, which contain the high order 32 bits and the low order 32 bits of the 64-bit simulation time. The delay value is in the simulator's internal time units, and will *not* be scaled to the object's time scale.
9. Set the delay value in the appropriate field of the **s_vpi_time** structure.
10. Call **vpi_put_value()**, with pointers to the **s_vpi_value** and **s_vpi_time** structures, and the constant representing how the value is to be scheduled in simulation time.

The following code fragment takes a value represented as a C character string, and writes the value onto the second argument of a system task, using transport delay event scheduling. The procedure for writing a value represented in any of the other C types is very similar to this example, and therefore separate examples of each value representation are not shown. The code fragment shown in this example is adapted from the **\$read_stimulus_ba** application, which is presented in full in the next chapter as Example 6-4 on page 214 in Chapter 6.

```
int PLIbook_ReadNextStim()
{
    ...
    s_vpi_time        time_s;
    s_vpi_value       value_s;

    /* obtain system task handle and Verilog vector handle */
    ...
}
```

```

/* read next line from the file */
...

time_s.type = vpiScaledRealTime;
time_s.real = delay;
value_s.format = vpiBinStrVal;
value_s.value.str = vector;
vpi_put_value(vector_h, &value_s, &time_s,
              vpiTransportDelay);

return(0);
}

```

No delay versus zero delay

The Verilog PLI standard makes a distinction between putting a value into simulation with no delay and putting a value into simulation with zero delay.

- The **vpiNoDelay** delay flag indicates that the value will be written into Verilog simulation instantly. When the PLI application returns back to simulation, any values written to an object or system function return using these routines will already be in effect for Verilog HDL statements to use.
- The **vpiIntertialDelay**, **vpiTransportDelay** and **vpiPureTransportDelay** flags schedule a value to be written into simulation. If a delay of zero is specified, the value is scheduled to be written to the object later in the current simulation time step. When the system task returns back to simulation, the scheduled value will not yet have taken effect, and other Verilog HDL statements scheduled to execute in the same simulation time step may or may not see the new value of the object (depending on where the value change which was scheduled by the PLI falls in the simulator's event queue in relation to other Verilog HDL events).

The following simple Verilog HDL source code illustrates the potential problem of putting a value into simulation using a delay of zero.

```

module test;
  reg [7:0] reg1, reg2;
  initial
    begin
      reg1 = 0; reg2 = 0;
      $put_value(reg1, reg2);
      $display("reg1=%d    reg2=%d", reg1, reg2);
      $strobe ("reg1=%d    reg2=%d", reg1, reg2);
      #1 $finish;
    end
endmodule

```

If the *calltf routine* for *\$put_value* writes a value to *reg1* using **vpiNoDelay**, then when *\$put_value* returns to the simulation, and the *\$display* statement prints the value of *reg1*, the *new* value will be printed.

If, however, the *calltf routine* for *\$put_value* writes a value to *reg2* using **vpiInertialDelay** or **vpiScaledRealTime** with a delay of zero, then, when *\$put_value* returns to the simulation, the *\$display* statement will print the *old* value of *reg2*. The old value is printed because the value written by the PLI has been scheduled to take place in the current time step, but it will not yet have taken effect. The *\$strobe* statement which follows the *\$display* will print the new value of both *reg1* and *reg2*, because the definition of *\$strobe* is to print its message at the end of the current simulation time step, after all value changes for that moment in time have taken effect.

5.4 Returning logic values of system functions

The **vpi_put_value()** routine is also used to write the return value of a system function into simulation. The object handle for where the value is to be written is the system function which called the PLI application.

Rules for returning values to system functions

There are two important restrictions on returning values to a system function:

- A value can only be written to the return of a system function from a *calltf routine*, which is when the system function is active. The *calltf routine* is invoked when the system function is encountered by the simulator when simulation is running. Return values cannot be written from VPI *simulation callback routines*, because the simulation is not executing the statement containing the function at the times these routines are invoked.
- When returning a value to a system function, the **vpiNoDelay** propagation flag must be used. It is illegal to specify a propagation delay when returning a value to a system function. If a delay is specified, the value will not be written and the VPI error flag will be set. Use **vpi_chk_error()** to determine if the flag is set.

Types of system functions

The VPI allows for several types of system functions. The type of function is established when the system function is registered through a call to **vpi_register_systf()**. Refer to section 2.11 on page 47 in Chapter 2 for a full description of registering system functions. In brief, the process involves allocating

an `s_vpi_systf_data` structure, setting the `type` field to `vpiSysFunction`, and setting the `sysfunctype` field to one of the following VPI constants:

- `vpiSysFuncInt` indicates the system function will return a Verilog integer value. Verilog integers are 32-bit signed values.
- `vpiSysFuncReal` indicates the system function will return a Verilog real value. Verilog reals are double precision floating point values.
- `vpiSysFuncTime` indicates the system function will return a Verilog time value. Verilog time variables are 64-bit unsigned values.
- `vpiSysFuncSized` indicates the system function will return a Verilog scalar or vector value. Verilog scalars are 1-bit wide and vectors can be any width. Vectors are always unsigned values.



NOTE The proposed IEEE 1364-1999 standard will change the names of the function type constants to `vpiIntFunc`, `vpiRealFunc`, `vpiTimeFunc`, and `vpiSizedFunc`, respectively. This change is to allow the constant names to be used to describe both system functions and Verilog HDL functions. For backward compatibility, the old constant names will be aliased to the new names.

Setting the function return size

In order for a Verilog simulator to compile Verilog source code, it must determine the return widths of functions. The system function types `vpiSysFuncInt`, `vpiSysFuncReal` and `vpiSysFuncTime` have predefined return sizes. The return size for `vpiSysFuncSized` system functions, however, is defined by the PLI application. The return size is established using the VPI *sizetf routine*. This routine is called by the simulator's compiler or loader prior to simulation time 0. The routine uses the C return value to inform the simulator of the vector size that the *calltf routine* will write back to the simulation. The *sizetf routine* is only used if the type of system function is `vpiSysFuncSized`. If no *sizetf routine* is registered for a `vpiSysFuncSized` system function, the default return size is 32 bits. The following example illustrates a *sizetf routine* to specify that the *calltf routine* will return a 64-bit vector.

```
int rand64_sizetf(char *user_data)
{
    return(64); /* $rand64 returns 64-bit values */
}
```



NOTE If a system function is used multiple times in the Verilog source code, the *sizetf routine* is only called once. The return value from this call is used to determine the return size for all instances of the system function, including any interactive calls to the function.

Example 5-8, shown below, implements a `$realpow` system function. This application returns a double-precision value representing m^n (m to the power of n). This example is an enhanced version of the `$pow` that was shown as Example 2-1 on page 44 in Chapter 2, which returned a 32-bit sized vector value.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/realpower_vpi.c
- Verilog test bench: Chapter.05/realpower_test.v
- Verilog-XL results log: Chapter.05/realpower_test.log

Example 5-8: `$realpow` — returning values of system functions

```
#include <math.h>      /* ANSI C standard input/output library */
int PLIBook_RealPow_calltf(char *user_data)
{
    s_vpi_value value_s;
    vpiHandle systf_handle, arg_itr, arg_handle;
    double base, exp, result;

    value_s.format = vpiRealVal;

    /* obtain handle to system task arguments;
       compilertf has already verified only 2 args with correct types */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);
    arg_itr = vpi_iterate(vpiArgument, systf_handle);

    /* read base value of system function arg 1 */
    arg_handle = vpi_scan(arg_itr);
    vpi_get_value(arg_handle, &value_s);
    base = value_s.value.real;

    /* read base value of system function arg 2 */
    arg_handle = vpi_scan(arg_itr);
    vpi_get_value(arg_handle, &value_s);
    exp = value_s.value.real;
    vpi_free_object(arg_itr); /* free iterator--did not scan till null */

    /* calculate result of base to power of exponent */
    result = pow(base,exp );

    /* write result to simulation as return value $realpow_func */
    value_s.value.real = result;
    vpi_put_value(systf_handle, &value_s, NULL, vpiNoDelay);
    return(0);
}
```

5.5 Using specparam constants as model attributes

Often a PLI application needs information about a model which is not part of the model functionality. A standard cell delay calculator, for example, might need the rise, slope and load factors of the cells used in a design. Information that is needed by the PLI application can be stored within the Verilog model using Verilog parameter constants or specparam constants. The PLI application can then obtain a handle for the constant and read its value.

The older ACC generation of the PLI standard supports a special usage of the Verilog specparam constant, called an *attribute*. This special usage requires adding a dollar sign (\$) to the end of the specparam constant name. An attribute specparam can be associated with all objects within a module, or with specific objects in a module.

- A general attribute is a specparam constant with a name which ends with a dollar sign.
- An object-specific attribute is a specparam constant with a base name which ends with a dollar sign, followed by the name of some object in the module.

The following example illustrates three specparam attributes:

```
module AN2 (o, a, b); // 2-input AND gate standard cell
  output o;
  input a, b;
  and (o, a, b);

  specify
    specparam BaseDelay$ = 2.2; //general attribute
    specparam InputLoad$a = 0.2; //object-specific attribute
    specparam InputLoad$b = 0.3; //object-specific attribute
  endspecify
endmodule
```

A special set of routines in the ACC portion of the PLI standard can read the value of the specparam constant using the attribute name and a handle for the object associated with that attribute. These ACC routines by-pass the need to first obtain a handle for the specparam constant. The ACC routines will search first for an object-specific attribute, then a general attribute, and finally return a default value.

The VPI library does not have a direct counterpart to the ACC attribute routines. However, it is easy to implement the same functionality. Example 5-9 lists a C function to fetch the value of a specparam attribute constant. The name of the attribute and a handle for an object are passed into the function. The attribute value is returned as a

C double. If no attribute can be found, a default value is returned. The steps performed by this example are:

1. Access the name of the object
2. Concatenate the name of the object to the end of the attribute name.
3. Use `vpi_handle_by_name()` to try to obtain a handle for the specparam constant with the attribute name and object name (the object-specific attribute).
4. If an object-specific attribute is not found, use `vpi_handle_by_name()` to try to obtain a handle for the specparam constant with just the attribute name (the general attribute).
5. If a general attribute is not found, return the default value. The default value is passed into the function as its third input.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/read_attribute_vpi.c
- Verilog test bench: Chapter.05/read_attribute_test.v
- Verilog-XL results log: Chapter.05/read_attribute_test.log

Example 5-9: using `vpi_get_value()` to read specparam attribute values

```
int PLIBbook_ReadAttribute_calltf(char *user_data)
{
    vpiHandle      systf_h, mod_h, arg_itr, arg2_h, port_itr, port_h;
    double         attribute_value;
    char          *attribute_name;
    s_vpi_value   attribute_name_s;

    /* obtain handle to system task arguments;
       compiletf has already verified only 2 args with correct types */
    systf_h = vpi_handle(vpiSysTfCall, NULL);
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    mod_h = vpi_scan(arg_itr);

    /* read base value of system function arg 2 */
    arg2_h = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* free iterator--did not scan till null */
    attribute_name_s.format = vpiStringVal;
    vpi_get_value(arg2_h, &attribute_name_s);
    attribute_name = attribute_name_s.value.str;

    vpi_printf("\nModule %s:\n", vpi_get_str(vpiDefName, mod_h));
    port_itr = vpi_iterate(vpiPort, mod_h);
    while ( (port_h = vpi_scan(port_itr)) != NULL) {
        attribute_value = PLIBbook_GetAttribute(port_h, attribute_name, 9.9);
        vpi_printf("  Port name = %s, attribute %s for this port = %2.2f\n",
                  port_h, attribute_name, attribute_value);
    }
}
```

```
        vpi_get_str(vpiName, port_h),
        attribute_name, attribute_value);
    }
    return(0);
}

//*****************************************************************************
* Function to return a specparam attribute value.
*****
double PLIbook_GetAttribute(vpiHandle obj_h, char *attribute,
                             double default_value)
{
    vpiHandle module_h, param_h;
    char *object_name;
    char param_name[1024]; /* character string to hold attribute name */
    s_vpi_value param_val; /* structure to receive attribute value */

    param_val.format = vpiRealVal;
    module_h = vpi_handle(vpiModule, obj_h); /* get parent module */

    /* build specparam name out of object name and attribute name */
    object_name = vpi_get_str(vpiName, obj_h);
    strcpy(param_name, attribute);
    strcat(param_name, object_name);

    /* try to get a handle to the object specific attribute */
    param_h = vpi_handle_by_name(param_name, module_h);
    if (!vpi_chk_error(NULL)) {
        vpi_get_value(param_h, &param_val);
        if (!vpi_chk_error(NULL))
            return(param_val.value.real); /* found specific attribute */
    }

    /* try to get a handle to a general attribute */
    strcpy(param_name, attribute);
    param_h = vpi_handle_by_name(param_name, module_h);
    if (!vpi_chk_error(NULL)) {
        vpi_get_value(param_h, &param_val);
        if (!vpi_chk_error(NULL))
            return(param_val.value.real); /* found general attribute */
    }

    /* failed to find object-specific or general attribute specparam */
    return(default_value);
}
```

5.6 Accessing Verilog memory arrays and variable arrays

The VPI routines can be used to both read and modify the contents of Verilog memory arrays and variable arrays. Memory arrays are one dimensional arrays of the Verilog `reg` data type. Variable arrays are one dimensional arrays of Verilog `integer` and `time` data types. The Verilog HDL syntax for array declarations is:

```
reg [<msb>:<lsb>] <memory_name> [<first_addr>:<last_addr>];  
integer <memory_name> [<first_address>:<last_address>];  
time <memory_name> [<first_address>:<last_address>];
```

The object diagrams for **memory** and **variables**, show the object relationships and properties for arrays. Refer to Appendix B, sections B.1.9 and B.1.10 for these diagrams.

Memory arrays

A memory (which is an array or `reg` signals) has a one-to-many relationship to **memory word**. Each memory word is an element in the array. The **vpiSize** property of a memory object indicates how many elements are in the array. The expression which can be accessed using **vpiLeftRange** represents the starting address of an array, and the expression which is accessed using **vpiRightRange** represents the ending address of the array.

The **vpiSize** property of a memory word indicates the bit width of each word. The expression which can be accessed using **vpiLeftRange** represents the bit number of the most-significant bit of the word, and the expression which is accessed using **vpiRightRange** represents the bit number of the most-significant bit of the word. The expression which can be accessed using **vpiIndex** represents the address of that word within the array.

Variable arrays

A **variables** object is represented just a little differently than a memory, because a variable object may be a single variable or may be an array of variables (whereas a memory is always an array). The boolean property **vpiArray** is used to determine if the variables object is an array.

- If the **vpiArray** property is false, then the variable object is a single variable. The **vpiSize** property will be the number of bits in the variable, and the expressions accessed by **vpiLeftRange** and **vpiRightRange** will indicate the most-significant and least-significant bit numbers, respectively.

- If the **vpiArray** property is true, then the variable object is an array of variables. The **vpiSize** property will be the number of elements in the array, and the expressions accessed by **vpiLeftRange** and **vpiRightRange** will indicate the first address and last address of the array, respectively.

Each element in a variable array is a **var select**. The expression which can be accessed using **vpiIndex** represents the address of that element within the array, and the **vpiSize** property indicates the bit-width of the element. There is no access to the most-significant bit and least-significant bit expressions of a variable select.

NOTE → The proposed IEEE 1364-1999 standard will add multi-dimensional arrays. The VPI object diagrams for reg and variables will be enhanced in the 1999 standard to support this new capability. The changes will not affect the diagram for memory, which remains a one-dimensional array of reg.

Handles for memory arrays and variable arrays can be obtained from a system task/function argument, or by iterating on `reg` or variable objects within a module. A specific memory word or a variable select is an expression, for which handles may be obtained in several ways, such as an argument of a system task/function or as part of a Verilog statement. From either object, it is easy to traverse to the other, using the one-to-one or one-to-many relationship shown in the object diagram.

The logic value of any element in a memory or variable array can be read or modified by obtaining a handle for the array element and using `vpi_get_value()` or `vpi_put_value()`. The entire contents of a memory array can be accessed by obtaining a handle for the array, and using `vpi_iterate()` to obtain handles for all of the elements in the array.

5.7 Reading and modifying delay values

There are several types of Verilog objects which have delays values. The VPI routines can both read and modify the delays of these objects.

5.7.1 Verilog objects which have delay values

Several types of constructs in the Verilog language can have delay values. How delays are represented for these different types of objects is part of the Verilog HDL standard, and is outside the scope of this book. The objects which can have delays are:

- **Primitive instances** can have delays specified for 1, 2 or 3 output transitions.
- **Module paths** can have delays specified for 1, 2, 3, 6 or 12 output transitions.

- **Module input ports** can have delays specified for 1, 2, or 3 output transitions. There is no construct in the Verilog language to represent Module Input Port Delays (MIPD's); only the PLI can add or modify delays on input ports.
- **Module interconnect paths** can have delays specified for 1, 2, or 3 output transitions. An interconnect path is the connection from the output of one module to the input of another module. There is no construct in the Verilog language to represent interconnect path delays; only the PLI can add or modify delays on input ports.
- **Timing constraint checks** can have delays specified for 1 limit for each constraint in the check. Timing constraints are represented in Verilog using \$setup, \$hold, etc.
- **Continuous assignments** can have delays specified for 1, 2, or 3 output transitions.
- **Procedural time controls (#)** can have 1 delay specified. This delay represents the time before a statement is executed, rather than an output transition delay.

In addition to multiple delay transitions, each delay value can be a set of delays representing a minimum, typical and maximum delay range. For example, a Verilog bufuf1 tri-state buffer gate can represent a propagation delay as:

A tri-state buffer with no delays:

```
bufuf1 g1 (...);
```

A tri-state buffer with delay of 5 for rising, falling, and turn-off transitions:

```
bufuf1 #5 g2 (...);
```

A tri-state buffer with separate delays for rising, falling, and turn-off transitions:

```
bufuf1 #(3, 4, 5) g3(...);
```

A tri-state buffer with separate minimum:typical:maximum delay sets for rising, falling and turn-off transitions:

```
bufuf1 #(2:3:4, 3:4:5, 5:6:7) g4 (...);
```

5.7.2 Reading an object's delay values

vpi_get_delays() reads the delay values of any type of object which has a delay property. The routine has two inputs, a handle for an object, and a pointer to an **s_vpi_delay** structure.

void vpi_get_delays (object, delay)

vpiHandle object	handle for an object which has a delay property.
p_vpi_delay delay	pointer to an application-allocated s_vpi_delay structure to receive delay information.

The **s_vpi_delay** structure is defined in vpi_user.h, and is listed below.

```
typedef struct t_vpi_delay
{
    struct t_vpi_time *da;
    int no_of_delays;
    int time_type;
    int mtm_flag;
    int append_flag;
    int pulsere_flag;
} s_vpi_delay, *p_vpi_delay;
```

There are six fields in the **s_vpi_delay** structure, which must be set by the PLI application before calling **vpi_get_delays()**.

- The **da** field of the structure must be set to an array of **s_vpi_time** structures. The delay values retrieved from the object will be loaded into this array.

NOTE → The PLI application must allocate the array of **s_vpi_time** structures, and the array must have at least as many elements as the number of delay values which are to be read. Refer to table 5-4 for how the number of elements in the array is determined.

- The **no_of_delays** field of the structure is set by the PLI application to the number of delay *transitions* which are to be retrieved from the object.
- The **time_type** field of the structure is set by the PLI application to one of the constants **vpiScaledRealTime** or **vpiSimTime**. The type indicates how the simulator should retrieve and store the delay values of the object. For **vpiScaledRealTime**, the delays are retrieved as double precision values and stored in the **real** field of the **s_vpi_time** structures pointed to by the **da** field. The delay values will be scaled to the time scale of the module containing the object from which the values are read. For **vpiSimTime**, the delays are retrieved as a pair of integer values, and stored in the **high** and **low** fields of the **s_vpi_time** structures pointed to by the **da** field. The delays are represented in the simulator's time units.
- The **mtm_flag** field is set by the PLI application to **0** (false) or **1** (true). False indicates that a single delay (the typical delay value) should be retrieved for each transition. True indicates that a minimum, typical, maximum delay set should be retrieved for each transition. Note that most Verilog simulators only store a single delay value for each transition, regardless of what is represented in the original Verilog source code. By default, these simulators store the typical delay value for each transition, but through invocation options or compiler directives, may store the minimum delay value or the maximum delay value for each transition. If the **mtm_flag** is set to true, and the simulator has only stored a single value for each transition, that one value will be returned for all three delays of the minimum, typical, maximum set.

- The **append_flag** field is not used by `vpi_get_delays()`.
- The **pulsere_flag** field is set by the PLI application to **0** (false) or **1** (true). False indicates that a single delay or delay set (depending on the **mtm_flag**) should be retrieved for each transition. True indicates that a pulse control set of values should be retrieved for each delay. The pulse control set comprises of the delay value, a *pulse reject limit*, and a *pulse error limit*. If the **mtm_flag** is also true, then a pulse control set will be retrieved for each delay in the minimum:typical:maximum set.

Before calling `vpi_get_delays()`, the PLI application must allocate an array of `s_vpi_time` structures to receive the delay values. The total number of elements in the array must be at least as large as the number of delay values to be retrieved. This total is controlled by the settings of **no_of_delays**, **mtm_flag**, and **pulsere_flag**. The number of required elements in the array can vary from 1 element (1 delay transition, typical delay mode, and no pulse control values) to 108 elements (12 delay transitions, minimum:typical:maximum delays, and pulse control values). Table 5-4 shows how to determine the total number of delays which will be retrieved.

Flag Values	Number of Array Elements	Order of retrieved delays			
		array element		object delay	
<code>mtm_flag = 0</code> <code>pulsere_flag = 0</code>	<code>no_of_delay * 1</code>	[0]	receives	1st delay	
		[1]	receives	2nd delay	
		
<code>mtm_flag = 1</code> <code>pulsere_flag = 0</code>	<code>no_of_delay * 3</code>	[0]	receives	1st delay	min value
		[1]			typ value
		[2]			max value
		[3]	receives	2nd delay	min value
		[4]			typ value
		[5]			max value
	
<code>mtm_flag = 0</code> <code>pulsere_flag = 1</code>	<code>no_of_delay * 3</code>	[0]	receives	1st delay	delay value
		[1]			reject limit
		[2]			error limit
		[3]	receives	2nd delay	delay value
		[4]			reject limit
		[5]			error limit
	

Flag Values	Number of Array Elements	Order of retrieved delays			
		array element		object delay	
		[0]	receives	1st delay	min value
		[1]			min reject limit
		[2]			min error limit
		[3]			typ value
		[4]			typ reject limit
		[5]			typ error limit
		[6]			max value
		[7]			max reject limit
		[8]			max error limit
		[9]	receives	2nd delay	min value
		[10]			min reject limit
		[11]			min error limit
		[12]			typ value
		[13]			typ reject limit
		[14]			typ error limit
		[15]			max value
		[16]			max reject limit
		[17]			max error limit
	

Table 5-4: Number of elements and order of the delay array (continued)

To read the delay values of an object requires the following basic steps:

1. Obtain a handle for an object with a delay property.
2. Allocate memory for an `s_vpi_delay` structure.
3. Allocate an array of `s_vpi_time` structures.
4. Set the appropriate fields in the `s_vpi_delay` structure.
5. Call `vpi_get_delays()` with a handle for the object and a pointer to the `s_vpi_delay` structure.

Example 5-10 lists a C function which reads the minimum, typical, maximum rise and fall delays of a module path. Note that module paths do not have a name property in the Verilog object diagrams. In order to print something meaningful, this example creates a module path name by retrieving the names of the nets connected to the first input of the path and the first output of the path, and concatenating these names together with a '\$' in between the two names.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/read_delays_vpi.c
- Verilog test bench: Chapter.05/read_delays_test.v
- Verilog-XL results log: Chapter.05/read_delays_test.log

Example 5-10: \$read_delays — reading module path delays

```
int PLIbook_ReadDelays_calltf(char *user_data)
{
    vpiHandle systf_h, arg_itr, mod_h, path_itr, path_h;

    /* obtain a handle to the system task instance */
    systf_h = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument; */
    /* compiletf has already verified only 1 arg with correct type */
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    mod_h = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* free iterator--did not scan to null */

    vpi_printf("\nModule %s paths:\n", vpi_get_str(vpiDefName, mod_h));
    path_itr = vpi_iterate(vpiModPath, mod_h);
    if (path_itr != NULL)
        while ((path_h = vpi_scan(path_itr)) != NULL) {
            PLIbook_PrintDelays(path_h);
        }
    else
        vpi_printf(" No path delays found.\n");
    return(0);
}

void PLIbook_PrintDelays(vpiHandle modpath_h)
{
    char *path_name;           /* pointer to path name string */
    s_vpi_delay delay_struct; /* structure to setup delays */
    s_vpi_time delay_array[6]; /* structure to receive delays */

    delay_struct.da          = delay_array;
    delay_struct.no_of_delays = 2;
    delay_struct.time_type   = vpiScaledRealTime;
    delay_struct.mtm_flag    = 1;
    delay_struct.append_flag = 0;
    delay_struct.pulsere_flag = 0;

    vpi_get_delays(modpath_h, &delay_struct);
    path_name = PLIbook_BuildPathName(modpath_h);
```

```
vpi_printf("Delays for %s = (%.2f:.2f:.2f, %.2f:.2f:.2f)\n",
    path_name,
    delay_array[0].real, delay_array[1].real, delay_array[2].real,
    delay_array[3].real, delay_array[4].real, delay_array[5].real);
}

char *PLIbook_BuildPathName(vpiHandle modpath_h)
{
    vpiHandle term_itr, term_h, net_h;
    static char path_name[2050]; /* character array to hold path name */
    char *term_name;

    path_name[0] = '\0'; /* clear the path name string */

    term_itr = vpi_iterate(vpiModPathIn, modpath_h);
    if (term_itr == NULL)
        return("UNKNOWN PATH NAME");
    term_h = vpi_scan(term_itr);
    net_h = vpi_handle(vpiExpr, term_h);
    if (net_h == NULL)
        return("UNKNOWN PATH NAME");
    term_name = vpi_get_str(vpiName, net_h);
    strcat(path_name, term_name);
    vpi_free_object(term_itr); /* free iterator--did not scan to null */

    strcat(path_name, "$");

    term_itr = vpi_iterate(vpiModPathOut, modpath_h);
    if (term_itr == NULL)
        return("UNKNOWN PATH NAME");
    term_h = vpi_scan(term_itr);
    net_h = vpi_handle(vpiExpr, term_h);
    if (net_h == NULL)
        return("UNKNOWN PATH NAME");
    term_name = vpi_get_str(vpiName, net_h);
    strcat(path_name, term_name);
    vpi_free_object(term_itr); /* free iterator--did not scan to null */

    return(path_name);
}
```

5.7.3 Writing delay values into an object

vpi_put_delays() writes delays values into an object. This routine requires two inputs, a handle for an object which has a delay property, and a pointer to an s_vpi_delay structure. The syntax of this routine is:

void vpi_put_delays (object, delay)

vpiHandle object handle for an object with a delay property.

p_vpi_delay delay pointer to an application-allocated *s_vpi_delay* structure containing delay information.

The process of writing delays into an object is nearly the reverse of reading delays, with the exception that the PLI application fills in the array of *s_vpi_time* structures with the values to be written into the object. The steps are:

1. Obtain a handle for an object with a delay property
2. Allocate memory for an *s_vpi_delay* structure.
3. Allocate an array of *s_vpi_time* structures.
4. Fill in the appropriate fields of the *s_vpi_time* structures with the delay values to be written onto the object. The delay values can be represented as double precision values or as pairs of C integers. Which way delays are represented is controlled by the setting of the *type* field in the *s_vpi_delay* structure. Note that the *type* also determines whether or not the delay values will be scaled to the object's time scale. A type of *vpiScaledReadTime* will scale the delay values to the time scale of the module which contains the object, and *vpiSimTime* will not scale the delay values (the delays are in the internal simulation time units).
5. Set the appropriate fields in the *s_vpi_delay* structure. The settings for these fields are nearly the same as with *vpi_get_delays()*, but with one additional flag. The **append_flag** is set to a **0** (false) or **1** (true). A false indicates that *vpi_put_delays()* will replace any existing delays for the object with the delays being written. A true indicates *vpi_put_delays()* will add the delays being written to any existing delays for the object.
6. Call *vpi_put_delays()* with a handle for the object and a pointer to the *s_vpi_delay* structure.

Example 5-11 shows a useful PLI application called *\$mipd_delays*. The Verilog HDL does not have a construct to represent module input port delays—these delays may only be represented through the PLI. *\$mipd_delays* provides a means for a Verilog model to add delays to the input port of a module.

The usage of this application is:

```
$mipd_delays(<port_name>, <t_rise>, <t_fall>, <t_toZ>)
```

Example:

```
$mipd_delays(in1, 2.7, 2.2, 1.0);
```



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.05/set_mipd_delays_vpi.c
- Verilog test bench: Chapter.05/set_mipd_delays_test.v
- Verilog-XL results log: Chapter.05/set_mipd_delays_test.log

Example 5-11: \$mipd_delays — using vpi_put_delays()

```

int PLIBook_SetMipd_calltf(char *user_data) {
    vpiHandle tfarg_itr, tfarg_h, port_itr, port_h;
    int i;
    s_vpi_delay delay_struct;      /* structure to delay setup */
    s_vpi_time delay_array[3];    /* structure to hold delays */
    s_vpi_value tfarg_value;      /* structure to hold tfarg values */

    delay_struct.da            = delay_array;
    delay_struct.no_of_delays = 3;
    delay_struct.time_type   = vpiScaledRealTime;
    delay_struct.mtm_flag     = 0;
    delay_struct.append_flag  = 0;
    delay_struct.pulsere_flag = 0;
    tfarg_value.format        = vpiRealVal;

    /* obtain handle to system task arguments;
       compiletf has already verified only 4 args with correct types */
    tfarg_itr = vpi_iterate(vpiArgument, vpi_handle(vpiSysTfCall, NULL));
    tfarg_h = vpi_scan(tfarg_itr); /* read 1st tfarg */
    port_itr = vpi_iterate(vpiPort, tfarg_h);
    port_h = vpi_scan(port_itr);
    vpi_free_object(port_itr); /* free iterator--did not scan to null */
    for (i=0; i<=2; i++) {
        vpi_get_value(vpi_scan(tfarg_itr), &tfarg_value); /* read tfargs */
        delay_array[i].real = tfarg_value.value.real;
    }
    vpi_free_object(tfarg_itr); /* free iterator--did not scan to null */

    vpi_put_delays(port_h, &delay_struct);
}

/*********************************************
 * compiletf routine
 *****/
int PLIBook_SetMipd_compiletf(p_cb_data cb_data)
{
    vpiHandle systf_h, tfarg_itr, tfarg_h, port_itr, port_h;
    int i;

    systf_h = vpi_handle(vpiSysTfCall, NULL);

```

```
tfarg_itr = vpi_iterate(vpiArgument, systf_h);
if (tfarg_itr == NULL) {
    vpi_printf("ERROR: $mipd_delays requires 4 arguments\n");
    tf_dofinish();
    return(0);
}
tfarg_h = vpi_scan(tfarg_itr);
if (vpi_get(vpiType, tfarg_h) != vpiNet) {
    vpi_printf("ERROR: $mipd_delays arg1 must be module input\n");
    tf_dofinish();
    return(0);
}
port_itr = vpi_iterate(vpiPort, tfarg_h);
if (port_itr == NULL) {
    vpi_printf("$mipd_delays arg1 not connected to module port\n");
    tf_dofinish();
    return(0);
}
port_h = vpi_scan(port_itr);
vpi_free_object(port_itr); /* free iterator--did not scan to null */
if (vpi_get(vpiType, port_h) != vpiPort) {
    vpi_printf("$mipd_delays arg1 not connected to an input port\n");
    tf_dofinish();
    return(0);
}
for (i=2; i<=4; i++) {
    tfarg_h = vpi_scan(tfarg_itr);
    if (vpi_get(vpiType, tfarg_h) != vpiConstant) {
        vpi_printf("$mipd_delays arg %d must be a number\n", i);
        tf_dofinish();
        return(0);
    }
}
if (vpi_scan(tfarg_itr) != NULL) {
    vpi_printf("ERROR: $mipd_delays requires only 4 arguments\n");
    vpi_free_object(tfarg_itr); /* free iterator--did not scan to null*/
    tf_dofinish();
    return(0);
}
return(0);
}
```

5.8 Summary

This chapter has presented the VPI routines which read and modify logic values and delay values in Verilog simulations. The VPI routines provide a means for PLI applications to access the logic and delay values of any object within simulation. This interface layer makes it possible for a PLI application to be portable to any IEEE 1364 compliant simulator, without concern for how the each simulator stores this complex information. Several short PLI applications illustrate the concepts presented.

The next chapter presents another powerful aspect of the VPI routines, the ability to synchronize PLI application activity with simulation activity.

CHAPTER 6

Synchronizing to Verilog Simulations Using VPI Callbacks

This chapter presents how to use a powerful feature of the VPI library—the ability to schedule with a Verilog simulator to invoke a PLI application for specific simulation events and simulation times. This ability to use the simulator’s event scheduling mechanism makes it possible to synchronize PLI applications with activities that occurs during a simulation.

The concepts presented in this chapter are:

- Sharing data between the routines of a PLI applications
- Creating a persistent work area unique to each instance of a system task/function
- Registering callbacks for when simulation actions occur
- Registering callbacks for specific simulation times
- Registering callbacks for when logic value changes occur
- Registering callbacks for when procedural statements are executed

6.1 PLI application callbacks

The Programming Language Interface provides a means whereby Verilog simulator users can write C language functions which are called by the Verilog simulator. The VPI standard refers to these calls as *callback routines* to the PLI application. The VPI provides two types callback routines:

- *System task/function callback routines* — these routines are executed when the

simulator encounters a user-defined system task or system function, either during compile time or during simulation run time. The system task/function callback routines are *calltf routines*, *compiletf routines* and *sizetf routines*.

- **Simulation callback routines** — these routines are executed when specific types of events occur, generally during simulation, such as time advancing to a specified time step or a logic value change.

The two types of callback routines need to be *registered* with the simulator. The VPI library provides routines for performing this registration, as well as a routine to remove a requested callback which has not yet transpired.

Chapter 2 showed how system task/function callback routines are registered. This chapter presents how to register and use the *simulation callback routines*.

6.2 Creating an instance specific work area

At times, it may be desirable to share information between different PLI applications that are associated with a system task/function, such as between a *compiletf routine* and a *calltf routine*. At other times, it may be necessary to preserve information from one call to a PLI application to another call of the same application. An example of where both of the above capabilities might be needed is an application to read test vectors from a pattern file. This type of application would need to open a disk file from which to read vectors, and then read a line from the file each time the system task is called by simulation.

An example test vector reader called *\$read_test_vector* is used in this chapter. The arguments for application are a file name and a Verilog reg vector. An example of using *\$read_test_vector* is:

```
always @(posedge clock)
$read_test_vector("A.dat", data_bus);
```

Since the test vector file “A.dat” should only be opened once, rather than at every positive edge of clock, a convenient place to open the file is in a *simulation callback* at the start-of-simulation. The test vector file can be opened at this time, and the file pointer saved for later use.

A PLI application is a C function, which is called by the Verilog simulator. Local variables within a C function are *automatic*, which means any variables (such as a variable to store a file pointer) do not remain allocated from one call of the function to

another call. Also, since the variables are local to a function, other C functions cannot use the value of the variables.

The work of actually reading values from the file will occur in the *calltf routine* for *\$read_test_vector*. The *calltf routine* will be invoked each time simulation encounters the system task, which, in the above example, is at each positive edge of clock. Since the *calltf routine* did not open the disk file, it does not have direct access to the file pointer.

A common solution in the C programming language to preserve data is to declare a static variable, instead of an automatic variable. Sharing common data between C functions can be done using global variables.

These C programming techniques can cause serious problems in PLI applications!

The Verilog HDL allows a system task/function to be used multiple times in the Verilog source code, and each module which uses the task/function can be instantiated multiple times. Each occurrence in each instance of a module becomes a unique *instance* of the system task/function. However, the same C function will be called by the simulator for each instance. A static or global variable cannot hold different values for each instance of the system task/function. Consider the following example:

```
always @(posedge clock)
$read_test_vector("A.dat", data_bus);

always @(negedge clock)
$read_test_vector("B.dat", data_bus);
```

In this example, the *\$read_test_vector* instance that is called at the positive edge of clock needs to save a file pointer to “*A.dat*”, and the *\$read_test_vector* that is called at the negative edge of clock needs to save a file pointer to “*B.dat*”. Since both instances of *\$read_test_vector* will invoke the same C function, a static or global variable cannot be used, because the two instances would share the same single copy of the variable. The variable can only hold a single file pointer, and there is no simple way of knowing which instance of *\$read_test_vector* was the last one to store a value in the variable.

Using static or global variables is a sure way to have problems when there are multiple instances of a system task or system function. Therefore, a PLI application must allocate storage that is unique to each instance of a system task/function. In the older TF portion of the PLI standard, an instance specific storage location is allocated automatically by the simulator, and is available for use by a PLI application whenever needed. This instance-specific storage is referred to as the *TF work area*, and special TF routines are provided to read and write values in the work area. The work area is shared by all routines which are associated with a system task/function (such as the

(*checktf* and *calltf routines*), and a unique work area is created for each instance of a system task/function.

The VPI routines do not include an automatic work area like the TF routines. Instead, a PLI application must allocate its own work area, if one is required.

NOTE The proposed IEEE 1364-1999 standard will add VPI routines to store and retrieve values from a simulator-maintained work area.

6.2.1 Creating instance specific work area storage

The *compiletf routine* will be called one time for each instance of a system task or system function. This makes the *compiletf routine* a convenient place to set up a work area for each instance. However, the *compiletf routine* is invoked prior to simulation, by either the simulator's compiler or as the simulator builds its internal data structures. The IEEE 1364 standard does not guarantee that memory allocated by a *compiletf routine* will be available when simulation starts running.

Instead, the *compiletf routine* can be used to schedule a call to a *simulation callback routine* at the start of actual simulation. The simulation callback routine can then allocate memory for a work area for each instance of a system task or system function.

Each instance of a system task or system function has a unique handle, and this handle can be used to designate the instance of the task/function with which a work area is associated. The unique handle for the system task/function instance must be obtained by the *compiletf routine* and passed to the *simulation callback routine*, so that the handle can be used to tag the work area memory at the time the work area is allocated.

Example 6-1 lists two convenient utilities which store and retrieve values from a work area that is associated with a specific instance of a system task or system function.

- **PLIbook_set_vpiworkarea()** stores a character pointer in a work area. This routine will automatically allocate storage for the work area the first time the routine is called for a specific task/function instance.
- **PLIbook_get_vpiworkarea()** retrieves a character pointer from the work area, using the system task/function instance handle to identify the work area for that instance.

The *PLIbook_set_vpiworkarea()* function will allocate a work area that is specific to an instance of a system task or system function the first time it is called for that instance. The work area is created using a simple linked-list stack, so that a work area can be allocated for any number of system task/function instances. The pointer to the

last address in the stack is maintained as a global variable, so that all PLI applications and C functions called by the simulator can easily access the stack.

The work area contains a single data field, which stores a `char*` pointer. The `char*` C data type is typically a 32-bit storage element, which can store a 32-bit value or an address pointer.



TIP The following example uses a simple linked-list for the work area stack. A simple hashing method is used to locate the work area within the stack for a specific instance of a system task/function. These simple methods are fine when a PLI application is not instantiated a large number of times (which would create a large stack) and when the PLI application is only invoked a few times during a simulation (which would require frequently hashing through the stack). If a PLI application is to be instantiated many times in a design, or invoked frequently during simulation, then other C programming methods for creating and searching stacked information should be used. Refer to advanced C programming books for more efficient methods of allocating and managing stacks and queues.



CD The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.06/vpi_utilities.c
- C test function: Chapter.06/use_workarea_vpi.c
- Verilog test bench: Chapter.06/use_workarea_test.v
- Verilog-XL results log: Chapter.06/use_workarea_test.log

Example 6-1: `PLIbook_set_vpiworkarea()`, `PLIbook_get_vpiworkarea()`—instance storage

```
/* Work area structure definition */
typedef struct PLIbook_vpiworkarea *PLIbook_vpiworkarea_p;
typedef struct PLIbook_vpiworkarea {
    vpiHandle systf_h; /* shows which systf instance owns this space */
    char *data;         /* data to be stored in workarea */
    PLIbook_vpiworkarea_p next_workarea;
} PLIbook_vpiworkarea_s;

/* allocate a global stack pointer */
static PLIbook_vpiworkarea_p PLIbook_vpiworkarea_stack = NULL;

***** * PLIbook_set_vpiworkarea()
***** */
void PLIbook_set_vpiworkarea(vpiHandle systf_h, char *data)
{
    PLIbook_vpiworkarea_p workarea = PLIbook_vpiworkarea_stack;
```

```
/* locate the workarea in the stack for this task instance */
while (workarea && (workarea->systf_h != systf_h))
    workarea = workarea->next_workarea;

/* if no work area found for this task instance, create one */
if (workarea == NULL) {
    workarea =
        (PLIbook_vpiworkarea_p)malloc(sizeof(PLIbook_vpiworkarea_s));
    workarea->systf_h = systf_h;      /* set owner of this workarea */
    if (PLIbook_vpiworkarea_stack == NULL) {
        /* work area stack doesn't exist yet, create first location */
        workarea->next_workarea = NULL;
        PLIbook_vpiworkarea_stack = workarea;
    }
    else {
        workarea->next_workarea = PLIbook_vpiworkarea_stack;
        PLIbook_vpiworkarea_stack = workarea;
    }
}

/* store data in the work area */
workarea->data = data;
return;
}

/*********************************************
 * PLIbook_get_vpiworkarea()
 *****/
char *PLIbook_get_vpiworkarea(vpiHandle systf_h)
{
    PLIbook_vpiworkarea_p workarea;

    /* locate the workarea in the stack for this task instance */
    workarea = PLIbook_vpiworkarea_stack;
    while (workarea && (workarea->systf_h != systf_h))
        workarea = workarea->next_workarea;
    if (workarea == NULL) {
        #ifdef PLIbookDebug /* generate verbose debug message */
            vpi_printf("Warning: workarea not found for this instance\n");
        #endif
        return(NULL);
    }
    return(workarea->data);
}
```

6.2.2 Storing a single value in the PLI application work area

Since the work area data field created in example 6-1 is defined to be a character pointer, it can store a single 32-bit value. Any other data type to be stored should be cast to a character pointer. An example of storing a single integer value in the work area is:

```
int my_data = 5;
PLIbook_set_vpiworkarea(systf_handle, (char *)my_data);
```

The following C code fragment shows an example of how the work area functions might be used. This code fragment opens a file in a C function that is executed at the very start of simulation, and saves the file pointer in a work area. Future invocations of the *calltf routine* can then retrieve the file pointer from the work area, in order to read values from the file. This example uses a *simulation callback routine* that is executed at the start of simulation to allocate the work area storage. Creating and registering *simulation callback routines* for the start-of-simulation is presented later in this chapter, in section 6.4 on page 196.

```
*****
 * StartOfSim callback--opens a test vector file and saves the file
 * pointer in a work area storage for this system task instance.
 ****
int PLIbook_StartOfSim(p_cb_data cb_data)
{
    FILE *vector_file;

    /* add code to open the test vector file */

    /* store file pointer in a work area for this task instance */
    PLIbook_set_vpiworkarea(systf_handle, (char *)vector_file);

    return(0);
}

*****
 * calltf routine
 ****
int ReadVectorcalltf(char *user_data)
{
    FILE *vector_file;

    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* get file pointer from a work area for this task instance */
    vector_file = (FILE *)PLIbook_get_vpiworkarea(systf_handle);

    /* read next line from the file */

    return(0);
}
```

6.2.3 Storing multiple values in the PLI application work area

Since the work area data field created in example 6-1 is a character pointer, the work area can be used to store a pointer to a memory location. Multiple values can be stored in the work area by allocating a block of memory and storing a pointer to the memory in the work area.



TIP Use instance specific work areas to improve the performance of PLI applications by reducing the number of times simulation hierarchy needs to be traversed to obtain object handles.

Often, each call to a PLI application will need to access the same objects. In the `$read_test_vector` example, the vector value read from a file will be applied to the second task argument. The *calltf routine* could go through the steps of obtaining a handle for the system task and traversing to the second argument and then obtaining a handle to for the that argument. But the *calltf routine* will be called every clock cycle in a simulation, possibly millions of times. The additional steps of obtaining the same handle every time the *calltf routine* is invoked is not very efficient for simulation run-time performance.

```
always @(posedge clock)
    $read_test_vector("A.dat", data_bus);
always @(negedge clock)
    $read_test_vector("B.dat", data_bus);
```

An instance-specific work area can be used to improve the performance of a PLI application. All handles which will be used in every call to the application can be retrieved one time and saved in a work area. Future calls to the PLI application can retrieve the handles from the work area much more quickly than repeatedly traversing the Verilog hierarchy to obtain the handles.

The PLI application fragment shown below is taken from the `$read_test_vector` application. This fragment stores both a file pointer and the handle to the second task argument in the work area. In this example, a structure called `PLIbook_MyAppData` is defined to store the two values to be stored. At the start of simulation, a simulation callback routine performs three operations:

- Allocates memory for one `PLIbook_MyAppData` structure.
- Fills the fields in the structure.
- Stores a pointer to the allocated memory in the work area.

Creating and registering *simulation callback routines* for the start-of-simulation is presented later in this section, on page 196. The complete \$read_test_vector is presented in that section , as Example 6-2 on page 200.

```
*****
* define storage structure for file pointer and vector handle.
*****
typedef struct PLIbook_MyAppData {
    FILE *file_ptr;      /* test vector file pointer */
    vpiHandle obj_h;    /* pointer to store an object handle */
} PLIbook_MyAppData_s, *PLIbook_MyAppData_p;

*****
* StartOfSim callback--opens a test vector file and saves the file
* pointer and the system task handle in the work area storage.
*****
int PLIbook_StartOfSim(p_cb_data cb_data)
{
    FILE           *vector_file;
    vpiHandle      systf_h, tfarg1_h, tfarg2_h;

    PLIbook_MyAppData_p vector_data; /* MyAppData structure pointer */

    vector_data =
        (PLIbook_MyAppData_p)malloc(sizeof(PLIbook_MyAppData_s));

    /* get tfarg handles */

    /* read file name from first tfarg and open the file */

    /* store file pointer and tfarg2 handle in the work area */
    /* for this instance */
    vector_data->file_ptr = vector_file;
    vector_data->obj_h = tfarg2_h;
    PLIbook_set_vpiworkarea(systf_h, (char *)vector_data);

    return(0);
}

*****
* calltf routine
*****
int ReadVectorcalltf(char *user_data)
{
    FILE           *vector_file;
    vpiHandle      systf_h, tfarg2_h;
    PLIbook_MyAppData_p vector_data; /* MyAppData structure pointer */

    systf_h = vpi_handle(vpiSysTfCall, NULL);

    /* get MyAppData pointer from work area for this task instance */
    vector_data = (PLIbook_MyData_p)PLIbook_get_vpiworkarea(systf_h);
```

```

vector_file = vector_data->file_ptr;
tfarg2_h    = vector_data->obj_h;

/* read next line from the file and apply vector to simulation */

return(0);
}

```

6.3 Registering simulation callback routines

The `vpi_register_cb()` routine registers a callback to a *simulation callback routine* for various simulation activity that can occur during simulation. These callbacks can occur any time during a simulation, and are not related the execution of a user-defined system task/function. The types of activity which can occur during simulation that can result in a PLI application being called are divided into four major categories:

- **Simulation actions**, such as the start of simulation or the end of simulation.
- **Simulation features**, such as entering or exiting an interactive debug mode.
- **Simulation time activity**, such as the end of the current simulation time step or the advancement to a specific simulation time.
- **Simulation events**, such as a logic value change or procedural statement execution.

The syntax of the `vpi_register_cb()` routine is:

`vpiHandle vpi_register_cb (data)`

`p_cb_data data` pointer to an application allocated `s_cb_data` structure containing callback information.

A callback to a PLI routine for any of these activities is not automatic. Some other PLI routine, such as a *callif routine*, must register a callback request with the simulator for when the desired activity occurs. To register a callback, the following steps are needed:

1. Allocate an `s_cb_data` structure.
2. Fill in the fields of the structure.
3. Call `vpi_register_cb()` with a pointer to the structure as an input.

The `s_cb_data` structure is defined in `vpi_user.h`, as follows:

```
typedef struct t_cb_data
{
    int reason;
    int (*cb_rtn)();
    vpiHandle obj;
    p_vpi_time time;
    p_vpi_value value;
    int index;
    char *user_data;
} s_cb_data, *p_cb_data;
```

The fields within the structure are:

- **reason** — an integer constant which represents what simulation activity will cause the callback to occur. The constants are defined in the `vpi_user.h`. The names and descriptions of these reason constants are presented on the following pages, in sections 6.4 through 6.7.
- **cb_rtn** — the name of the PLI routine which should be called when the specified simulation activity occurs.
- **obj** — a handle for an object. If a callback is requested for logic value changes on a specific net, for example, then a handle for that net must be provided in this field. Not all callback reasons require an object handle. The descriptions of each callback reason on the following pages indicate if an object handle is required for that callback.
- **time** — a pointer to an `s_vpi_time` structure. If a callback is requested at some specific simulation time, the time structure needs to be filled with time value of when the callback should occur. The descriptions of each callback reason on the following pages indicate if a time structure is required for that callback.
- **value** — a pointer to an `s_vpi_value` structure. If a callback is requested for logic value changes on a specific net, then the `format` field in the value structure determines how the simulator will return the value of the object. The descriptions of each callback reason on the following pages indicate if a value structure is required for that callback.
- **index** — not used when a callback is registered.
- **user_data** — a user data value. The value placed in `user_data` will be passed to the routine that is called when the specified simulation activity occurs. If this field is not needed by the routine, it should be set to NULL.

Several examples of registering simulation activity callbacks are presented in the following sections of this chapter.

6.3.1 C function inputs to a simulation callback routine

When the simulator calls the registered routine, a pointer to an **s_cb_data** structure is passed as an input to the routine. Within this structure is information about the callback.

- For *simulation action callbacks* and *simulation feature callbacks*, the structure will contain the `user_data` value specified when the callback was registered.
- For *simulation time callbacks*, the structure will contain the current simulation time and the `user_data` value.
- For *simulation event callbacks*, the structure will contain the current simulation time, the logic value of the object on which the event occurred, the index and object handle, if they pertain to the type of event, and the `user_data` value.



NOTE The structure passed into the *simulation callback routine* is allocated and maintained by the simulator, and will be freed after the *simulation callback routine* exits. This input is the same type of structure as was used to register the callback, but it is not the same structure allocation that was used for registration.

6.3.2 Using the user_data value

The user data value is not required for any simulation activity callback, but it is a very useful way to pass information to a callback. The `user_data` is a character pointer, which can be used to store a single 32-bit value or to point to a block of data. The `user_data` value is passed to the *simulation callback routine* each time the simulator calls the routine.



NOTE If a pointer to a block of data is specified in the `user_data` field, the PLI application must maintain this memory. Local variables used in the function that registers the callback will not be available when the callback occurs. The PLI application should allocate memory to store the data, and free the memory when it is no longer needed.

6.3.3 Using system task/function arguments in a simulation callback

A *simulation callback routine* is not directly associated with an instance of a system task or system function, and therefore cannot directly obtain a handle for a specific task/function instance. In many PLI applications, however, the simulation callback originated from a specific task/function instance, and the simulation callback may need to access information about that instance, such as the system task/function argu-

ments. As an example, assume the Verilog source code contained the following instances of `$read_test_vector`:

```
always @(posedge clock)
$read_test_vector("A.dat", data_bus);

always @(negedge clock)
$read_test_vector("B.dat", data_bus);
```

One way to implement this PLI application is to use the following basic steps:

1. In the *completf routine*, first verify the correctness of the task function arguments, and then register a simulation callback for the start-of-simulation. The *completf routine* is called for each instance of a system task, so, in the above example, two start-of-simulation callbacks will be registered (assuming the module which contains the instances of `$read_test_vector` is only instantiated one time in the Verilog design).
2. In the *simulation callback routine* for start-of-simulation, read the file name listed in the first system task argument, open the file for reading, and save the file pointer in an instance-specific work area.
3. In the *calltf routine* (which is executed twice every clock cycle in this example), retrieve the file pointer from the work area and read the next vector from the file.

In the preceding sequence, step two, the simulation callback for start-of-simulation, will occur two times, once for each instance of `$read_test_vector`. The PLI application which is called will need to obtain a handle to the specific instance of `$read_test_vector` which registered the callback, so that the correct file name is read from argument 1 and the file pointer is stored in the correct work area.

The *simulation callback routine* is not associated with any system task/function, however. Therefore, the `user_data` is used to pass the handle for the `$read_test_vector` task instance to the *simulation callback routine*. When the *completf routine* sets up the simulation callback for start-of-simulation, the *completf routine* can obtain the handle for that specific instance of the system task/function, and then store the handle in the `user_data` field for the callback. The *simulation callback routine* at start-of-simulation can then retrieve the system task/function handle from the `user_data`, and access the arguments of that instance of the system task/function.

Example 6-2 on page 200 presents the `$read_test_vector` PLI application, and illustrates using the `user_data` field to store and pass a system task/function handle.

6.3.4 Removing scheduled callbacks

A callback which has been requested, but which has not yet transpired, is referred to as a *scheduled callback*. A scheduled callback can be removed at any time using `vpi_remove_cb()`. The syntax for this routine is:

```
int vpi_remove_cb (cb_object)
    vpiHandle   cb_object   handle for a callback object.
```

The `vpi_remove_cb()` routine removes callbacks to PLI applications which were registered with `vpi_register_cb()`. The routine returns 1 (for true) if successful, and 0 (for false) if an error occurred. The callback handle is no longer valid after the callback is removed.

When a callback is registered, `vpi_register_cb()` returns a handle for the scheduled callback. If a PLI application might need to remove a callback at a future time, the application should save the scheduled callback handle. If `vpi_remove_cb()` is called to remove a callback which has already transpired, the removal request is ignored.

6.4 Simulation action and simulation feature callbacks

Callbacks to a PLI application can be scheduled for when specific actions occur during a simulation. The IEEE 1364 standard divides action related callbacks into three groups: *required action callbacks*, *optional feature callbacks* and *vendor-specific feature callbacks*. The VPI standard defines constants to represent the reasons callbacks can occur in the first two groups. Reason constants for the third group will be defined by the simulator vendor.

Required action callbacks must be implemented in all IEEE 1364 compliant simulators. These callback reasons are described in Table 6-1.

Constant	Definition
<code>cbEndOfCompile</code>	calls a PLI application at the end of simulation data structure compilation or build (immediately before time 0)
<code>cbStartOfSimulation</code>	calls a PLI application at the start of simulation (beginning of the time 0 simulation cycle)

Table 6-1: Required VPI simulation action callback constants

Constant	Definition
cbEndOfSimulation	calls a PLI application at the end of simulation (e.g.: the <code>\$finish</code> system task was executed)
cbError	calls a PLI application if a run-time error occurred while executing the Verilog HDL portion of the simulation
cbPLIError	calls a PLI application if a run-time error occurred while executing a PLI application
cbTchkViolation	calls a PLI application if a Verilog HDL timing check violation occurred

Table 6-1: Required VPI simulation action callback constants (continued)

Optional feature callbacks might be implemented in an IEEE 1364 compliant simulator, but they are not required. These callback reasons are listed in table 6-2.

Constant	Definition
cbEnterInteractive	calls a PLI application when simulation enters interactive debug mode (e.g.: a <code>\$stop</code> system task executed)
cbExitInteractive	calls a PLI application when simulation exits interactive mode
cbInteractiveScopeChange	calls a PLI application when a simulation command to change interactive scope is executed
cbUnresolvedSystf	calls a PLI application when a system task/function is encountered that is not pre-defined in the simulator or is not registered using <code>vpi_register_systf()</code>
cbStartOfSave	calls a PLI application when a simulation checkpoint is started (e.g.: a <code>\$save</code> command is executed)
cbEndOfSave	calls a PLI application when a simulation checkpoint is completed
cbStartOfRestart	calls a PLI application when a restart from a simulation checkpoint is started (e.g.: a <code>\$restart</code> command is executed)
cbEndOfRestart	calls a PLI application when a restart from a simulation checkpoint is completed

Table 6-2: Optional VPI simulation feature callback constants

6.4.1 Required settings to register simulation action/feature callbacks

To register a callback for a specific simulation action or feature, the following fields in the **s_cb_data** structure must be set:

- **reason** must be set to one of the simulation action or feature constants.
- **cb_rtn** must specify the name of the PLI application which should be called when the action occurs.
- **user_data** (optional) can be set to point to a user data value, if needed. If this field is not used, it should be set to NULL.

The remaining fields are not required for simulator action and feature callbacks, and will be ignored. The unused fields can be set to NULL or left unspecified.

It is not an error to register a callback for an optional feature that is not implemented in a simulator. The constants for these optional reasons are defined in the IEEE 1364 standard, which means the PLI application will compile and link correctly with any IEEE 1364 compliant simulator. If the simulator has not implemented the optional functionality, the callback will simply never occur.

More caution is needed for vendor-specific feature callbacks. If a simulator vendor adds a reason constant to a simulator, that constant may not exist, or might be defined with a different value in another simulator. If a PLI application uses vendor-specific callback reasons, then it is recommended the application be written in such a way that those callbacks can be omitted when the application is compiled for other simulators (perhaps using C conditional compilation directives).

6.4.2 When a simulation action/feature callback occurs

When a simulator action or feature related callback occurs, the simulator will allocate an **s_cb_data** structure and pass a pointer to the structure as an input to the *callback routine* which is called. For action and feature related callbacks, the only field in the **s_cb_data** structure that is used is the **user_data** field. This field will contain whatever **user_data** value was specified when the callback was registered. In the example which follows, the **user_data** field will contain a handle for the specific instance of the system task which requested the simulation action callback.

6.4.3 How to register a simulation callback before time zero

Most simulation callbacks will be registered from another VPI routine, such as a *calltf routine*. These callbacks are typically registered after simulation starts running, and schedule the callback to occur at some future simulation action, time or event.

In order to register a callback for cbEndOfCompile or cbStartOfSimulation, however, the call to vpi_register_cb() must occur prior to simulation time zero. There are two places where a simulation time callback can be registered prior to time zero, each with unique and important capabilities:

- A C function with the call to vpi_register_cb() can be listed in the *vlog_startup_routines* array. The functions listed in this array are executed by the simulator prior to the end of compilation. *Note that the functions listed in the start up array are only executed one time.* This means the end-of-compile or start-of-simulation callback will only occur one time, regardless of how many instances of a system task/function occur in the Verilog hierarchy. Also note that the functions called from the start up array are not associated with any specific instance of a system task or system function, and therefore cannot obtain a handle for a specific system task/function.
- The call to vpi_register_cb() can be listed in a *completf routine*. The *completf routine* is invoked automatically as the simulation data structure is being built, so both end-of-compile and start-of-simulation callbacks can be registered from the *completf routine*. *Note that the completf routine is invoked for each instance of a system task/function.* This means the end-of-compile or start-of-simulation callback may be executed multiple times, once for each instance of the system task/function. Also the *completf routine* is associated with a task/function instance, and can access information about the task/function to pass on to the simulation callback, such as the handle to that instance of a system task/function.

6.4.4 Limitations on callbacks registered before time zero

The cbEndOfCompile and cbStartOfSimulation callbacks have certain limitations which must be observed. These callbacks occur after the simulation data structure has been fully created, but before any simulation values have been propagated. This means all reg and variables will be uninitialized, and nets may not yet reflect the values of their drivers. Reading logic values prior to time zero is not an error, but probably has little meaning.

Other than reading logic values, there are no other restrictions on what can be done at these callback times. Some activity that is common and appropriate for these callback reasons are:

- Traverse the design hierarchy to collect handles or other information that will be needed by the *calltf routine* or other callbacks as simulation time advances.
- Allocate memory for use during simulation.
- Set up an instance specific shared work area.
- Other activity that only needs to be performed one time during simulation.

An example of using simulation action callbacks

Example 6-2 lists the complete `$read_test_vector` PLI application, which reads test vectors from a file and passes the vector value to simulation. This example utilizes a number of concepts that have been presented in this chapter and in previous chapters. Some of the key VPI concepts used are:

- Using multiple C functions to create a complete PLI application, including a *callif routine*, a *compileif routine*, and a *callback routine* for a start-of-simulation action.
- Using of the *compileif routine* to a register simulation callback for each instance of `$read_test_vector`.
- Using *simulation action callbacks*; specifically a callback for the reason `cbStartOfSimulation`.
- Using an instance-specific work area to pass data from one function to another within the `$read_test_vector` application and to preserve data over simulation time.

The usage of `$read_test_vector` is:

```
$read_test_vector("<file_name>", <reg_vector>);
```

The application reads vectors from a file one line at a time. Therefore `$read_test_vector` should be invoked in a loop that will read vectors until the end of simulation or the end of the vector file. For example:

```
always @(posedge clock)
$read_test_vector("A.dat", data_bus);
```



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.06/read_test_vector_vpi.c
- Verilog test bench: Chapter.06/read_test_vector_test.v
- Verilog-XL results log: Chapter.06/read_test_vector_test.log

Example 6-2: `$read_test_vector` — using VPI simulation action callbacks

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include <malloc.h>        /* ANSI C standard memory allocation library */
#include "vpi_user.h"      /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"      /* IEEE 1364 PLI TF routine library
                           (using TF routines for simulation control) */

/* include utility routines to work with tfargs */
#include "vpi_utilities.c"
```

```

/* prototypes of routines in this PLI application */
int PLIbook_ReadVector_calltf(), PLIbook_ReadVector_compiletf(),
    PLIbook_StartOfSim();

/********************* Define storage structure for file pointer and vector handle. *****/
typedef struct PLIbook_Data {
    FILE *file_ptr;      /* test vector file pointer */
    vpiHandle obj_h;    /* pointer to store handle for a Verilog object */
} PLIbook_Data_s, *PLIbook_Data_p;

/********************* VPI Registration Data *****/
void PLIbook_ReadVector_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$read_test_vector";
    tf_data.calltf   = PLIbook_ReadVector_calltf;
    tf_data.compiletf = PLIbook_ReadVector_compiletf;
    tf_data.sizetf   = NULL;
    tf_data.user_data = NULL;
    vpi_register_systf(&tf_data);
}

/********************* compiletf routine *****/
int PLIbook_ReadVector_compiletf(char *user_data)
{
    s_cb_data cb_data_s;
    vpiHandle systf_h, arg_itr, arg_h;
    int tfarg_type, err = 0;
    char *file_name;

    systf_h = vpi_handle(vpiSysTfCall, NULL);
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    if (arg_itr == NULL) {
        vpi_printf("ERROR: $read_test_vector requires 2 arguments\n");
        tf_dofinish();
        return(0);
    }
    arg_h = vpi_scan(arg_itr); /* get handle for first tfarg */
    if (vpi_get(vpiType, arg_h) != vpiConstant) {
        vpi_printf("$read_test_vector arg 1 must be a quoted file name\n");
        err = 1;
    }
    else if (vpi_get(vpiConstType, arg_h) != vpiStringConst) {
        vpi_printf("$read_test_vector arg 1 must be a string\n");
        err = 1;
    }
}

```

```
arg_h = vpi_scan(arg_itr); /* get handle for second tfarg */
tfarg_type = vpi_get(vpiType, arg_h);
if ( (tfarg_type != vpiReg) &&
     (tfarg_type != vpiIntegerVar) &&
     (tfarg_type != vpiTimeVar) ) {
    vpi_printf("$read_test_vector arg 2 must be a register type\n");
    err = 1;
}
if (vpi_scan(arg_itr) != NULL) {
    vpi_printf("read_test_vector requires only 2 arguments\n");
    vpi_free_object(arg_itr);
    err = 1;
}
if (err)
    tf_dofinish();

/* setup a callback for start of simulation */
cb_data_s.reason      = cbStartOfSimulation;
cb_data_s.cb_rtn       = PLIbook_StartOfSim;
cb_data_s.obj          = NULL;
cb_data_s.time         = NULL;
cb_data_s.value        = NULL;
cb_data_s.user_data   = (char *)systf_h; /* pass systf_h to callback */
vpi_register_cb(&cb_data_s);
return(0);
}

*****
* calltf application
*****
int PLIbook_ReadVector_calltf(char *user_data)
{
    s_cb_data      data_s;
    s_vpi_time    time_s;
    s_vpi_value   value_s;
    FILE          *vector_file;
    vpiHandle     systf_h, arg2_h;
    PLIbook_Data_p vector_data; /* pointer to a ReadVecData structure */
    char          vector[1024]; /* fixed vector size, could use malloc*/

    systf_h = vpi_handle(vpiSysTfCall, NULL);

    /* get ReadVecData pointer from a work area for this task instance */
    /* the data in the work area was loaded at the start of simulation */
    vector_data = (PLIbook_Data_p)PLIbook_get_vpiworkarea(systf_h);
    vector_file = vector_data->file_ptr;
    arg2_h      = vector_data->obj_h;

    /* read next line from the file */
    if ( (fscanf(vector_file, "%s\n", vector)) == EOF) {
        vpi_printf("$read_test_vector reached End-Of-File\n");
        fclose(vector_data->file_ptr);
        tf_dofinish();
        return(0);
    }
}
```

```
/* write the vector to the second system task argument */
value_s.format = vpiBinStrVal;
value_s.value.str = vector;
vpi_put_value(arg2_h, &value_s, NULL, vpiNoDelay);

return(0);
}

*****
* StartOfSim callback -- opens the test vector file and saves the
* pointer and the system task handle in the work area storage.
*****
int PLIbook_StartOfSim(p_cb_data cb_data)
{
    s_vpi_value    argVal;
    char           *file_name;
    FILE          *vector_file;
    vpiHandle     systf_h, arg_itr, arg1_h, arg2_h;

    PLIbook_Data_p vector_data; /* pointer to a ReadVecData structure */

    vector_data = (PLIbook_Data_p)malloc(sizeof(PLIbook_Data_s));

    /* retrieve system task handle from user_data */
    systf_h = (vpiHandle)cb_data->user_data;

    /* get argument handles (completetf already verified only 2 args) */
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    arg1_h  = vpi_scan(arg_itr);
    arg2_h  = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* free iterator -- did not scan to null */

    /* read file name from first tfarg */
    argVal.format = vpiStringVal;
    vpi_get_value(arg1_h, &argVal);
    if (vpi_chk_error(NULL)) {
        vpi_printf("ERROR: $read_test_vector could not get file name\n");
        return(0);
    }
    file_name = argVal.value.str;

    if ( !(vector_file = fopen(file_name, "r")) ) {
        vpi_printf("$read_test_vector could not open file %s", file_name);
        tf_dofinish;
        return(0);
    }

    /* store file pointer and tfarg2_h in work area for this instance */
    vector_data->file_ptr = vector_file;
    vector_data->obj_h = arg2_h;
    PLIbook_set_vpiworkarea(systf_h, (char *)vector_data);

    return(0);
}
```

NOTE

In the preceding example, the PLI application will cause simulation to exit when the end of the test vector file is reached. The IEEE 1364-1995 standard does not provide a VPI routine for this functionality, and therefore this application calls the `tf_dofinish()` routine from the TF library to exit simulation. The proposed 1364-1999 standard will add a new VPI routine, which will provide similar functionality as `tf_dofinish()` and `tf_dostop()`.

6.5 Simulation time callbacks

Simulation callbacks can be registered to occur at specific times during a simulation. These types of callbacks utilize the event scheduling mechanism of the simulator to schedule the simulation callback at a specific point in simulation time.

The reason constants for simulation time callbacks are listed in table 6-3.

Constant	Definition
<code>cbReadWriteSynch</code>	calls a PLI application after the execution of all known events in the specified time step. The PLI can schedule additional events for the same simulation time
<code>cbReadOnlySynch</code>	calls a PLI application after the execution of all events in the specified time step. The PLI <i>cannot</i> schedule additional events for the same simulation time. However, another <code>cbReadOnlySynch</code> callback can be scheduled.
<code>cbNextSimTime</code>	calls a PLI application before the execution of any events in the next simulation time step which has simulation events scheduled
<code>cbAtStartOfSimTime</code>	calls a PLI application at a specific simulation time (absolute to time 0), before the execution of any simulation events in the time step
<code>cbAfterDelay</code>	calls a PLI application after a specified amount of time (relative to the current time), before execution of any simulation events in that time step

Table 6-3: VPI simulation time callback constants

Required settings to register simulation time callbacks

To register a simulation time callback, the following fields in the **s_cb_data** structure must be set:

- **reason** must be set to one of the simulation time constants.
- **cb_rtn** must specify the name of the PLI routine which should be called when the specified time activity occurs.
- **time** must specify a pointer to an **s_vpi_time** structure. The **type** field controls how time is specified for when the callback is to occur, and the way that time is passed to the callback routine when the callback occurs.
- **obj** must specify a handle for an object if the **type** field in the time structure is **vpiScaledRealTime**. The simulation time scale for the module which contains the object will be used to scale the time value. If the time type is **vpiSimTime**, the **obj** field is not used, and should still be set to **NULL**.
- **user_data** (optional) can be set to a user data value, if needed. If this field is not used, it should be set to **NULL**.

The **value** and **index** fields are not required for simulator time callbacks, and will be ignored. The unused fields can be set to **NULL** or left unspecified.

When a simulation time callback occurs

When a simulator time callback occurs, the simulator will allocate an **s_cb_data** structure and pass a pointer to the structure as an input to the *simulation callback routine* which is called. The fields in the **s_cb_data** structure which the simulator will fill are:

- The **user_data** field of the structure will contain whatever **user_data** value was specified when the callback was registered.
- The **time** field will contain a pointer to a time structure. Within the time structure, will be the simulation time in which the callback occurred (the current simulation time). The time value will be stored in the format that was specified in the time **type** field when the callback was registered—**vpiScaledRealTime** will store time values in the **real** field, and **vpiSimTime** will store time values in the **high** and **low** fields. The **vpiSuppressTime** constant is not allowed when registering simulation time callbacks.

6.5.1 Simulation callbacks at the end of a time step

A PLI application can schedule simulation time callbacks to a *simulation callback routine* at the end of the current simulation time step or at the end of a future time step. This capability allows PLI applications to synchronize themselves with the activity in simulation.

A **cbReadWriteSynch** callback will call a PLI application after the execution of all known events in the specified time step. The PLI can schedule additional events in that simulation time step.

A **cbReadOnlySynch** callback will call a PLI application after the execution of all events in the specified time step. The PLI *cannot* schedule additional events in that simulation time step. However, a PLI application is permitted to schedule another **cbReadOnlySynch** callback.

For both of these callbacks, the time step in which the callback occurs is set using the **s_vpi_time** structure when the callback is registered. When the **type** field in the time structure is **vpiSimTime** or **vpiScaledRealTime**, the callback will occur at the time specified in the **high** and **low** fields or the **real** field of the time structure, using a relative delay from the time the callback was registered.

There are many reasons a PLI application might need to synchronize activity to the end of a simulation time step. One situation might be to communicate all logic value changes in simulation for that moment in simulation time to a C language model, and pass any C model value changes back to the Verilog simulation. This type of callback would require a read-write synchronization, which would use a simulation time callback scheduled with a **cbReadWriteSynch** reason.

Another example of where synchronization to the end of a simulation time step is shown in the following example:

```
always @(a or b)
$my_strobe(sum);           always @(a or b)
                           sum = a + b;
```

parallel (concurrent) activity

The *callif* routine for **\$my_strobe** will be called every time **a** or **b** changes, at which time the routine might need to read the value of argument 1, which is the **sum** signal in Verilog. At the same moment in simulation time, however, the **sum** variable is also scheduled to change value in the simulation. This is a classic race condition in hardware simulation that is caused by the concurrent activity of reading and writing values

of the same object in the same simulation time step. The Verilog standard states that most types of concurrent activity can be processed in any order by the simulation, which means the *calltf routine* can be executed either before or after the `sum` signal changes at the positive edge of the clock. The outcome of this race condition is not predictable.

The race condition in the above example can be resolved by synchronizing the `$my_strobe` PLI application to the very end of the simulation time step in which the change on `a` or `b` occurs. The `$my_strobe` would require a read-only synchronization, which is a simulation time callback scheduled with the `cbReadOnlySynch` reason. When a *simulation callback routine* is called at the end of the time step, any and all statements for that time step will have been executed, and the values of the arguments to the system task or system function will be at their most current value. For the `$my_strobe` example, when the *calltf routine* is invoked when `a` changes, it can schedule a callback to a *simulation callback routine* at the end of that same time step in read-only mode. The *simulation callback routine* can then read the value of `sum`, and be assured that it has the most current value for that moment in time.

The VPI provides two methods of synchronizing to the end of a simulation time step:

`cbReadWriteSynch` schedules a simulation time callback to a *simulation callback routine* at the end of a specified simulation time step in a *read-write* mode. In this mode, the *simulation callback routine* can read the values of a Verilog objects, and it can also modify the values of objects in the current simulation time (using `vpi_put_value()` with no delay), as well as schedule values to be written in future simulation times (using `vpi_put_value()` with a delay).

NOTE

Changing the value of an object may cause additional events in the current step. The new events will be processed, but the `cbReadWriteSynch` simulation callbacks which have already been executed for that time step will *not* be re-invoked. Only new events will be processed.


`cbReadOnlySynch` schedules a simulation time callback to a *simulation callback routine* at the end of a specified simulation time step in a *read-only* mode. In this mode, the *simulation callback routine* is only permitted to read values. The *simulation callback routine* is not allowed to modify values in the current time. Values can be scheduled to be written at a future time using `vpi_put_value()` with a delay.

Two fundamental differences should be noted between VPI routines and TF routines regarding synchronization to the end of a time step. First, the equivalent TF routines (`tf_synchronize()` and `tf_rosynchronize()`) can only schedule a callback to a *misctf routine* at the end of the current simulation time step—they cannot schedule synchronization callbacks at a future time. Second, in a read-only synchronization

callback using TF routines, it is illegal to write values into simulation at any time, including future times. The VPI *simulation callback routines* provide much more capability and flexibility for synchronizing PLI applications to simulation time.

The following diagram illustrates the activity that occurs within a simulation time step in Verilog.

Figure 6-1: Organization of events in a Verilog simulation time step



In essence, there are four distinct regions of events within a time step, which are referred to as *slots* in the above figure. Within each slot, certain types of simulation events are scheduled to be executed. The Verilog standard allows simulators to optimize the events within a slot in any order, and to intermix the types of events within each slot in any order.

NOTE → *This illustration of a simulation time step is purely conceptual.* While the illustration is accurate, it is an abstraction of the much more detailed description of simulation event scheduling that is specified in the IEEE 1364 Verilog standard.

A simulation will not proceed to slot 2 until it has executed all scheduled events within slot 1. However, the events within slot 2 may cause new events to be added to slot 1. Simulation must then return to slot 1, and execute the new events, which may result in new events in slot 2. Note that when simulation returns to a slot, only new events are processed—events which have already been executed have been removed from that slot's event list, and will not be executed a second time. Assuming there are

no zero-delay infinite loops in the Verilog code, simulation will eventually complete all events in slot 1 and 2, and then proceed to slot 3.

In slot 3, any registered simulation time callbacks with **cbReadWriteSynch** will be invoked. In this mode, the PLI is allowed to schedule new events in the current simulation time step. These new events will be scheduled in slot 1, forcing the simulator to once again return to slot 1 to process the new events. Note again, that only new events will be processed when the simulator returns to slot 1.

Once all events in slots 1, 2 and 3 have been processed, simulation will proceed to slot 4. In this slot, any registered simulation time callbacks with **cbReadOnlySynch** will be invoked. In the read-only synchronize mode, the PLI is allowed to read information from the simulation, but the PLI may *not* schedule new events in the current simulation time step. Slot 4 represents the true end of the current simulation time step.

Example 6-3, below, illustrates using a simulation time callback registered with **cbReadOnlySynch** to implement **\$my_strobe**. An example of using **\$my_strobe** is:

```
always @(a or b)           always @(a or b)
$my_strobe(sum);          sum = a + b;

```

parallel (concurrent) activity

In this example usage, the *calltf routine* for **\$my_strobe** will be executed when a or b changes, at the same moment the signal **sum** will be changed. Since there is no way to determine if the *calltf routine* will be executed before or after the change to **sum**, the *calltf routine* does not read the value of **sum**. Instead, the *calltf routine* schedules a *read-only synchronize* callback to a *simulation callback routine* at the end of the current time step, and the value of **sum** is read and printed from the *simulation callback routine*.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.06/my_strobe_vpi.c
- Verilog test bench: Chapter.06/my_strobe_test.v
- Verilog-XL results log: Chapter.06/my_strobe_test.log

Example 6-3: \$my_strobe — simulation callbacks at the end of the current time

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include "vpi_user.h"     /* IEEE 1364 PLI VPI routine library */
```

```
#include "veriuser.h" /* IEEE 1364 PLI TF routine library
                      (using TF routines for simulation control) */

/* prototypes of routines in this PLI application */
int PLIBook_MyStrobe_calltf(), PLIBook_MyStrobe_compiletf(),
    PLIBook_EndOfTimeStep_callback();

/********************* VPI Registration Data *********************/
void PLIBook_MyStrobe_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname   = "$my_strobe";
    tf_data.calltf   = PLIBook_MyStrobe_calltf;
    tf_data.compiletf = PLIBook_MyStrobe_compiletf;
    tf_data.sizetf   = NULL;
    tf_data.user_data = NULL;
    vpi_register_systf(&tf_data);
}

/********************* calltf application ********************/
PLIBook_MyStrobe_calltf(char *user_data)
{
    vpiHandle systf_h, arg_itr, arg_h;
    s_vpi_time time_s;
    s_cb_data cb_data_s;

    /* obtain a handle to the system task instance */
    systf_h = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument; */
    /* compiletf has already verified only 1 arg with correct type */
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    arg_h = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* free iterator--did not scan to null */

    /* setup end-of-time step callback */
    time_s.low  = 0;
    time_s.high = 0;
    time_s.type = vpiSimTime;
    cb_data_s.reason      = cbReadOnlySynch;
    cb_data_s.user_data  = (char *)arg_h;
    cb_data_s.cb_rtn     = PLIBook_EndOfTimeStep_callback;
    cb_data_s.obj        = arg_h;
    cb_data_s.time       = &time_s;
    cb_data_s.value      = NULL;
    vpi_register_cb(&cb_data_s);
}
```

```
*****
 * Value change callback application
 ****
PLIBbook_EndOfTimeStep_callback(p_cb_data cb_data_p)
{
    s_vpi_time    time_s;
    s_vpi_value   value_s;
    static char   new_value[] = "Z";
    vpiHandle     arg_h = (vpiHandle)cb_data_p->user_data;

    value_s.format = vpiBinStrVal;
    vpi_get_value(arg_h, &value_s);
    vpi_printf("$my_strobe: At %d: \t %s = %s\n",
               cb_data_p->time->low,
               vpi_get_str(vpiFullName, arg_h),
               value_s.value.str);
}

```

6.6 Application-scheduled callbacks at a future simulation time

The VPI provides three ways in which a PLI application can schedule simulation callbacks at a future simulation time.

cbAtStartOfSimTime schedules a simulation time callback at a specific simulation time, using the absolute time in which the callback should occur. Absolute time is always relative to time 0, and is not affected by the current simulation time. The callback will occur at the beginning of the specified time step, before any other simulation events are processed in that time step.

cbAfterDelay schedules a simulation time callback at a future simulation time, using a relative time from the current simulation time. The callback will occur at the beginning of the specified time step, before any other simulation events are processed in that time step.

cbNextSimTime schedules a simulation time callback at a future simulation time, which is determined by whenever the simulator has scheduled its next event to be executed. The callback will occur at the beginning of the specified time step, before any other simulation events are processed in that time step. If the current simulation time is 100, and the next scheduled event is time 125, then the simulation callback will be scheduled before any other simulation events are processed in time 125.

For the cbAtStartOfSimTime and cbAfterDelay callbacks, the time step in which the callback occurs is set using the s_vpi_time structure when the callback is registered. When the type field in the time structure is vpiSimTime or vpiScaledRealTime, the callback will occur at the time specified in the high and low fields or the real field of the time structure. The vpiSuppressTime time type is not allowed with simulation time related callbacks.


The vpi_remove_cb() routine can be used to remove a scheduled callback which has not transpired. In order to remove a callback, the handle for the callback must have been preserved by the PLI application. The callback handle is returned by vpi_register_cb() when the callback is registered.

An example of scheduling simulation callbacks at future times

Example 6-4 illustrates using a simulation callback for a cbAfterDelay reason. The example implements a system task called \$read_stimulus_ba, which uses a *callback routine* to read one line at a time from a stimulus vector file. Each line contains a simulation time and a test vector, as follows:

time	vector
10	11111111xxxxxx
17	00000000zzzzzz
30	1000000011011101

The *callback routine* for \$read_stimulus_ba will need to process the following loop as long as there are more test vectors in the file:

- 
1. Read a time and test vector from the file.
 2. Schedule test vector to be applied to simulation at the desired time.
 3. Schedule the *simulation callback routine* to be called again once the test vector has been applied.

The usage for this example is:

```
initial  
$read_stimulus_<base><delay_type>("file", verilog_reg);
```

where:

```
<base>      is b or h (for binary or hex vectors).  
<delay_type> is a or r for absolute or relative times.  
"file"       is the name of the file to be read, in quotes.  
verilog_reg  is a verilog register data type.
```

For example:

```
initial  
  $read_stimulus_ba("readstim.pat", input_vector);
```

Notice that `$read_stimulus_ba` is called from a Verilog `initial` procedure, which means the *calltf routine* for `$read_stimulus_ba` will only be invoked one time throughout simulation. The *calltf routine* is used to schedule an immediate callback to a *simulation callback routine*. The *simulation callback routine* will then read the first line from the file, and schedule the test vector to be applied at the simulation time specified by the delay (using `vpi_put_value()`). After scheduling the vector to be applied, the *simulation callback routine* will schedule another callback to itself, this time using the same delay with which the test vector was scheduled to be applied. At that future simulation time, the *simulation callback routine* is invoked again, and the cycle is repeated. When the last line of the file is read, the *simulation callback routine* will cause the simulation to exit.

Example 6-4 on page 214 presents the `$read_stimulus_ba` PLI application. This PLI application utilizes several types of VPI callbacks:

- A *completelf system task callback* is used to perform syntax checking and to schedule a simulation callback for the start of simulation.
- A *calltf system task callback* is used to open the stimulus file and invoke the first call to a simulation callback which reads one line at a time from the file.
- A *simulation action callback* scheduled for the start of simulation is used to allocate an instance-specific block of persistent memory to be used by the other *simulation callback routines*. A pointer to the memory is saved in an instance-specific work area.
- A *simulation time callback* scheduled for a future time is used to retrieve a stimulus pattern and a simulation time from the file, and schedule the simulator to apply the test pattern and the designated simulation time.
- A *simulation time callback* scheduled for the end of the current time step is used to terminate simulation when `$read_stimulus_ba` has reached the end of the stimulus pattern file. Instead of terminating simulation immediately on end of file, a read-only synchronous callback is scheduled to allow any other activity in the current simulation time step to complete first.

This example also illustrates other important concepts which have been presented in this and earlier chapters on the VPI library. Some of these concepts are:

- Using invocation option to configure a PLI application. In this example, invoking simulation with `+readstim_debug` will enable additional error checking and output messages as the lines area read from the stimulus pattern file.
- Using instance-specific work areas to store information that needs to be shared by different *simulation callback routines*.
- Using the system task/function user_data to allow different system task names to call the same PLI application *simulation callback routines*. In this example, the stimulus values contained in the stimulus file can be specified as either hex patterns or binary patterns, and the simulation times can be specified as relative delays or as absolute time from time 0. This flexibility is accomplished by registering several versions of the `$read_stimulus_ba` system task name, all of which call the same PLI application, but with different user-data values:
 - `$read_stimulus_ba` has a user-data of “`ba`”
 - `$read_stimulus_br` has a user-data of “`br`”
 - `$read_stimulus_ha` has a user-data of “`ha`”
 - `$read_stimulus_hr` has a user-data of “`hr`”

As with many of the examples in this book, the purpose of the example is to show how specific PLI routines are used. The example does not emphasize efficient C programming techniques. In order to make the code functionality more obvious, the format of the stimulus file has been kept simple, and C variables are used in many places where pointers or nested function calls could have been used.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.06/PLIbook_read_stimulus_vpi.c`
- Verilog test bench: `Chapter.06/PLIbook_read_stimulus_test.v`
- Verilog-XL results log: `Chapter.06/PLIbook_read_stimulus_test.log`

Example 6-4: `$read_stimulus_ba` — scheduling callbacks at a future time

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>       /* ANSI C standard input/output library */
#include <malloc.h>       /* ANSI C standard memory allocation library */
#include "vpi_user.h"     /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"     /* IEEE 1364 PLI TF routine library
                           (using TF routines for simulation control) */

/* include utility routines to work with tfargs */
#include "vpi_utilities.c"
```

```

/* prototypes of routines in this PLI application */
int PLIbook_ReadStim_calltf(), PLIbook_ReadStim_completetf();
int PLIbook_StartOfSim(), PLIbook_ReadNextStim(),
PLIbook_ReadStimEnd();

/*********************************************************/
/* Define storage structure for file pointer and vector handle.
*********************************************************/
typedef struct ReadStimData {
    FILE *file_ptr;      /* test vector file pointer */
    vpiHandle obj_h;    /* pointer to store handle for a Verilog object */
    int mode;           /* 0 & 1 = binary values, 2 & 3 = hex values */
    /* 0 & 2 = absolute time, 1 & 3 = relative time */
    int debug;          /* print debug messages if true */
} s_ReadStimData, *p_ReadStimData;

/*********************************************************/
/* VPI Registration Data
*********************************************************/
void PLIbook_ReadStim_register()
{
    s_vpi_systf_data tf_data;
    tf_data.type      = vpiSysTask;
    tf_data.calltf    = PLIbook_ReadStim_calltf;
    tf_data.completetf = PLIbook_ReadStim_completetf;
    tf_data.sizetf    = NULL;

    tf_data.tfname     = "$read_stimulus_ba"; /* binary, absolute time */
    tf_data.user_data = "ba";
    vpi_register_systf(&tf_data);

    tf_data.tfname     = "$read_stimulus_br"; /* binary, relative time */
    tf_data.user_data = "br";
    vpi_register_systf(&tf_data);

    tf_data.tfname     = "$read_stimulus_ha"; /* hex, absolute time */
    tf_data.user_data = "ha";
    vpi_register_systf(&tf_data);

    tf_data.tfname     = "$read_stimulus_hr"; /* hex, relative time */
    tf_data.user_data = "hr";
    vpi_register_systf(&tf_data);
}

/*********************************************************/
/* completetf routine
*********************************************************/
int PLIbook_ReadStim_completetf(char *user_data)
{
    s_cb_data   cb_data_s;
    vpiHandle   systf_h, arg_itr, arg_h;

```

```
int          tfarg_type, err = 0;
char        *file_name;

systf_h = vpi_handle(vpiSysTfCall, NULL);
arg_itr = vpi_iterate(vpiArgument, systf_h);
if (arg_itr == NULL) {
    vpi_printf("ERROR: $read_stimulus_?? requires 2 arguments\n");
    tf_dofinish();
    return(0);
}
arg_h = vpi_scan(arg_itr); /* get handle for first tfarg */
if (vpi_get(vpiType, arg_h) != vpiConstant) {
    vpi_printf("$read_stimulus_?? arg 1 must be a quoted file name\n");
    err = 1;
}
else if (vpi_get(vpiConstType, arg_h) != vpiStringConst) {
    vpi_printf("$read_stimulus_?? arg 1 must be a string\n");
    err = 1;
}
arg_h = vpi_scan(arg_itr); /* get handle for second tfarg */
tfarg_type = vpi_get(vpiType, arg_h);
if ( (tfarg_type != vpiReg) &&
     (tfarg_type != vpiIntegerVar) &&
     (tfarg_type != vpiTimeVar) ) {
    vpi_printf("$read_stimulus_?? arg 2 must be a register type\n");
    err = 1;
}
if (vpi_scan(arg_itr) != NULL) {
    vpi_printf("read_stimulus_?? requires only 2 arguments\n");
    vpi_free_object(arg_itr);
    err = 1;
}
if (err)
    tf_dofinish();

/* setup a callback for start of simulation */
cb_data_s.reason      = cbStartOfSimulation;
cb_data_s.cb_rtn       = PLIbook_StartOfSim;
cb_data_s.obj          = NULL;
cb_data_s.time         = NULL;
cb_data_s.value        = NULL;
cb_data_s.user_data   = (char *)systf_h; /* pass systf_h to callback */
vpi_register_cb(&cb_data_s);

return(0);
}

*****  

* calltf routine -- registers an immediate callback to the  

* ReadNextStim application.  

*****  

int PLIbook_ReadStim_calltf(char *user_data)
{
```

```

s_cb_data      data_s;
s_vpi_time    time_s;
vpiHandle     systf_h;
p_ReadStimData StimData; /* pointer to a ReadStimData structure */

/* get ReadStimData pointer from work area for this task instance */
systf_h = vpi_handle(vpiSysTfCall, NULL);
StimData = (p_ReadStimData)PLIbook_get_vpiworkarea(systf_h);

/* look at user data to set the stimulus mode flag */
if      (strcmp(user_data, "ba") == 0) StimData->mode = 0;
else if (strcmp(user_data, "br") == 0) StimData->mode = 1;
else if (strcmp(user_data, "ha") == 0) StimData->mode = 2;
else                                     StimData->mode = 3;

/* setup immediate callback to ReadNextStim routine */
systf_h      = vpi_handle(vpiSysTfCall, NULL);
time_s.type   = vpiSimTime;
time_s.low    = 0;
time_s.high   = 0;
data_s.reason = cbReadWriteSynch;
data_s.cb_rtn  = PLIbook_ReadNextStim;
data_s.obj     = NULL;
data_s.time    = &time_s;
data_s.value   = NULL;
data_s.user_data = (char *)systf_h;
vpi_register_cb(&data_s);

return(0);
}

*****
* ReadNextStim callback -- Reads a time and vector from a file.
* Schedules the vector to be applied at the specified time.
* Schedules a callback to self that same time (to read next line).
*****
int PLIbook_ReadNextStim(p_cb_data cb_data)
{
    char vector[1024]; /* fixed max. size, should use malloc instead */
    int          delay;
    vpiHandle    systf_h;
    s_cb_data    data_s;
    s_vpi_time   time_s;
    s_vpi_value  value_s;
    p_ReadStimData StimData; /* pointer to a ReadStimData structure */

    /* retrieve system task handle from user_data */
    systf_h = (vpiHandle)cb_data->user_data;

    /* get ReadStimData pointer from work area for this task instance */
    StimData = (p_ReadStimData)PLIbook_get_vpiworkarea(systf_h);

    /* read next line from the file */

```

```
if ( (fscanf(StimData->file_ptr, "%d %s\n", &delay, vector)) == EOF) {  
    /* At EOF, schedule ReadStimEnd callback at end of this time */  
    time_s.type      = vpiSimTime;  
    time_s.low       = 0;  
    time_s.high      = 0;  
    data_s.reason    = cbReadOnlySynch;  
    data_s.cb_rtn    = PLIbook_ReadStimEnd;  
    data_s.obj       = NULL;  
    data_s.time      = &time_s;  
    data_s.value     = NULL;  
    data_s.user_data = (char *)StimData->file_ptr;  
    vpi_register_cb(&data_s);  
}  
  
if (StimData->debug) {  
    vpi_printf("Values read from file: delay=%d vector=%s\n",  
               delay, vector);  
}  
  
/* convert absolute delay from file to relative delay if needed */  
time_s.type = vpiScaledRealTime;  
if (StimData->mode == 0 || StimData->mode == 2) {  
    vpi_get_time(cb_data->obj, &time_s);  
    time_s.real = ((double)delay - time_s.real);  
}  
else  
    time_s.real = (double)delay;  
  
/* schedule the vector to be applied after the delay period */  
if (StimData->mode == 0 || StimData->mode == 1)  
    value_s.format = vpiBinStrVal;  
else  
    value_s.format = vpiHexStrVal;  
value_s.value.str = vector;  
vpi_put_value(StimData->obj_h, &value_s, &time_s, vpiTransportDelay);  
  
/* schedule callback to this routine when time to read next vector */  
data_s.reason      = cbAfterDelay;  
data_s.cb_rtn      = PLIbook_ReadNextStim;  
data_s.obj         = systf_h; /* object required for scaled delays */  
data_s.time        = &time_s;  
data_s.value       = NULL;  
data_s.user_data   = (char *)systf_h;  
vpi_register_cb(&data_s);  
if (vpi_chk_error(NULL))  
    vpi_printf("An error occurred registering ReadNextStim callback\n");  
return(0);  
}  
  
*****  
* StartOfSim callback -- opens the test vector file and saves the  
* file pointer and other info in an instance-specific work area.  
*****
```

```

int PLIBook_StartOfSim(p_cb_data cb_data)
{
    char            *file_name;
    FILE           *vector_file;
    vpiHandle      systf_h, tfarg1_itr, tfarg2_h;
    s_cb_data      data_s;
    s_vpi_time     time_s;
    s_vpi_value    argVal;
    s_vpi_vlog_info options_s;
    p_ReadStimData StimData; /* pointer to a ReadStimData structure */
    int            i, debug;

    /* retrieve system task handle from user_data */
    systf_h = (vpiHandle)cb_data->user_data;

    /* get tfarg handles (compiletf already verified args are correct) */
    tfarg1_itr = vpi_iterate(vpiArgument, systf_h);
    tfarg2_h = vpi_scan(tfarg1_itr);
    tfarg1_h = vpi_scan(tfarg1_itr);
    vpi_free_object(tfarg1_itr); /* free iterator--did not scan to null */

    /* read file name from first tfarg */
    argVal.format = vpiStringVal;
    vpi_get_value(tfarg1_h, &argVal);
    if (vpi_chk_error(NULL)) {
        vpi_printf("ERROR: $read_stimulus_?? could not get file name\n");
        return(0);
    }
    file_name = argVal.value.str;
    if ( !(vector_file = fopen(file_name, "r")) ) {
        vpi_printf("$read_stimulus_?? could not open file %s\n", file_name);
        tf_dofinish();
        return(0);
    }

    /* check for +readstim_debug invocation option */
    debug = 0; /* assume not invoked with debug flag */
    vpi_get_vlog_info(&options_s);
    for (i=1; i<options_s.argvc; i++) {
        if (strcmp(options_s.argv[i], "+readstim_debug") == 0) {
            debug = 1; /* invocation option found */
            break;
        }
    }

    /* allocate memory to store information about this instance */
    StimData = (p_ReadStimData)malloc(sizeof(s_ReadStimData));
    StimData->file_ptr = vector_file;
    StimData->obj_h = tfarg2_h;
    StimData->debug = debug;
    PLIBook_set_vpiworkarea(systf_h, (char *)StimData);
    return(0);
}

```

```
*****
 * End-Of-Simulation callback -- close file and exit simulation.
 ****
int PLIBook_ReadStimEnd(p_cb_data cb_data_p)
{
    vpi_printf("$read_stimulus_?? reached End-Of-File.\n");
    fclose((FILE *)cb_data_p->user_data);
    tf_dofinish();
    return(0);
}
```

6.7 Simulation event callbacks

Simulation callbacks can be registered for when specific events transpire during simulation.

Constant	Definition
cbValueChange	calls a PLI application after a logic value change on an expression or terminal
cbStmt	calls a PLI application before execution of a procedural statement
cbForce	calls a PLI application after a force has occurred on a simple expression
cbRelease	calls a PLI application after a release has occurred on a simple expression
cbAssign	calls a PLI application after a procedural assign has been executed on a simple expression
cbDeassign	calls a PLI application after a procedural de-assign has been executed on a simple expression
cbDisable	calls a PLI application after a procedural disable has been executed on a named statement group, HDL task or HDL function

Table 6-4: VPI simulation event callback constants

Required settings to register simulation event callbacks

To register a simulation event callback, the following fields in the **s_cb_data** structure must be set:

- **reason** must be set to one of the simulation event constants.
- **cb_rtn** must specify the name of the PLI routine which should be called when the specified event activity occurs.
- **obj** must specify a handle for an object. The object must be appropriate for the callback reason. If the reason is cbForce or cbRelease, and the object field is NULL, the simulation callback will be called for every force or release.
- **time** must specify a pointer to an **s_vpi_time** structure. The time structure **type** field in the time structure can be set to **vpiScaledRealTime** or **vpiSimTime**. The time type determines how the simulation time will be passed to the *simulation callback routine* when a callback occurs. If the simulation time is not needed by the *simulation callback routine*, the type field can be set to **vpiSuppressTime**.
- **value** must specify a pointer to an **s_vpi_value** structure for all event reasons except cbStmt. The **format** field in the value structure must be set to a value format constant to indicate how the logic value of the object will be returned to the *simulation callback routine*. The format constants are listed in table 6-5, which follows. Refer to section 5.2 on page 140 in Chapter 5 for a full description of each value format and how the Verilog logic value is represented. If the new logic value of the object is not needed by the *simulation callback routine*, the format field can be set to **vpiSuppressVal**. For cbStmt simulation event callbacks, the value field is not used, and will be ignored.
- **user_data** (optional) can be set to a user data value, if needed. If this field is not used, it should be set to NULL.

Format	Definition
vpiBinStrVal	represents a Verilog logic value as a C string, using binary numbers, with 4-state logic
vpiOctStrVal	represents a Verilog logic value as a C string, using octal numbers, with 4-state logic
vpiDecStrVal	represents a Verilog logic value as a C string, using decimal numbers, with 4-state logic
vpiHexStrVal	represents a Verilog logic value as a C string, using hexadecimal numbers, with 4-state logic

Table 6-5: The **s_vpi_value** structure format constants

Format	Definition
vpiScalarVal	represents a Verilog scalar logic value as C constants, which represents Verilog 4-state logic
vpiStrengthVal	represents a Verilog scalar value as a VPI strength structure containing the logic value and Verilog strength levels
vpiVectorVal	represents a Verilog vector value as an array of VPI aval/bval structures, encoded to represent Verilog 4-state logic
vpiIntVal	represents a Verilog logic value as a C integer, with 2-state logic
vpiTimeVal	represents a Verilog time value as a VPI time structure
vpiRealVal	represents a Verilog logic value as a C double precision number, with 2-state logic
vpiStringVal	represents a Verilog string value as a C string
vpiObjTypeVal	allows simulation to pick the best way to represent a value
vpiSuppressVal	no logic value is retrieved when a simulation callback occurs

Table 6-5: The `s_vpi_value` structure format constants (continued)

A `cbValueChange` callback reason can be registered for objects of: nets, regs, variables, memories, memory word selects, variable arrays, variable array selects, module ports, and primitive terminals. A value change is any event in Verilog simulation that results in either the logic value or strength value of the object to change.

When a simulation event callback occurs

When a simulator event callback occurs, the simulator will allocate an `s_cb_data` structure, and pass a pointer to the structure as an input to the PLI routine which is called. The fields in the `s_cb_data` structure which are filled in by the simulator are:

- The `user_data` field of the structure will contain the `user_data` value which was specified when the callback was registered.
- The `time` field will contain a pointer to a time structure with the simulation time in which the callback occurred (the current simulation time). The time will be stored in the appropriate time value field, based on the time `type` used when the callback was registered.
- The `value` field will contain a pointer to a value structure, which will contain the new logic value of the specified object, in the format specified when the callback was registered. For `cbForce` and `cbAssign` callbacks, the value in the value structure will show the resultant value of the left-hand side of the statement. For

cbRelease and **cbDeassign** statements, the value will contain the value of the object after the release has occurred. The value field is not used for **cbStmt** callbacks.

- For **cbValueChange** callback reasons, the **obj** field will be a handle for the object which changed (the same object for which the callback was registered). For **cbForce**, **cbRelease**, **cbAssign** and **cbDeassign** callback reasons, the **obj** field will be a handle for the force, release, assign or deassign statement which was executed on the specified object.
- If the object for which the callback was registered is a word of a memory array or variable array, then the **index** field will contain the index into the array.

All memory used by the **s_cb_data** structure that is passed to the *simulation callback routine* is allocated and maintained by the simulator. This is temporary storage, and the pointers to the time and value structures will not remain valid after the *simulation callback routine* exits. If the PLI application needs to preserve the time or logic values, then the application must allocate its own memory and copy the information.

An example of using simulation event callbacks for logic value changes

The following example implements an application called **\$my_monitor**. This application is passed a module instance name as an input. The application locates all nets within the module, and registers for **cbValueChange** callbacks for each net. When the callback occurs, the application registers a simulation time callback to end of the current time step, and prints the new value of the net, along with the net's name and the current simulation time.

The usage of **\$my_monitor** is:

```
initial
$my_monitor(<module_instance_name>);
```

CD The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.06/my_monitor_vpi.c
- Verilog test bench: Chapter.06/my_monitor_test.v
- Verilog-XL results log: Chapter.06/my_monitor_test.log

Example 6-5: **\$my_monitor** — scheduling simulation callbacks at a future time

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include <malloc.h>        /* ANSI C standard memory allocation library */
```

```
#include "vpi_user.h" /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h" /* IEEE 1364 PLI TF routine library
                      (using TF routines for simulation control) */

/* prototypes of routines in this PLI application */
int PLIbook_MyMonitor_calltf(), PLIbook_MyMonitor_completetf(),
    PLIbook_MyMonitor_callback();

/********************* VPI Registration Data ********************/
void PLIbook_MyMonitor_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$my_monitor";
    tf_data.calltf    = PLIbook_MyMonitor_calltf;
    tf_data.completetf = PLIbook_MyMonitor_completetf;
    tf_data.sizetf    = NULL;
    tf_data.user_data = NULL;
    vpi_register_systf(&tf_data);
}

/********************* completf application ********************/
int PLIbook_MyMonitor_completetf(char *user_data)
{
    vpiHandle systf_handle, arg_iterator, arg_handle;
    int        tfarg_type;

    /* obtain a handle to the system task instance */
    systf_handle = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handles to system task arguments */
    arg_iterator = vpi_iterate(vpiArgument, systf_handle);
    if (arg_iterator == NULL) {
        vpi_printf("ERROR: $my_monitor requires 1 argument\n");
        tf_dofinish(); /* abort simulation */
        return(0);
    }

    /* check the type of object in system task arguments */
    arg_handle = vpi_scan(arg_iterator);
    tfarg_type = vpi_get(vpiType, arg_handle);
    if (tfarg_type != vpiModule) {
        vpi_printf("ERROR: $my_monitor arg1 must be module instance\n");
        vpi_free_object(arg_iterator); /* free iterator memory */
        tf_dofinish(); /* abort simulation */
        return(0);
    }
}
```

```

/* check that there is only 1 system task argument */
arg_handle = vpi_scan(arg_iterator);
if (arg_handle != NULL) {
    vpi_printf("ERROR: $my_monitor can only have 1 argument\n");
    vpi_free_object(arg_iterator); /* free iterator memory */
    tf_dofinish(); /* abort simulation */
    return(0);
}
return(0);
}

*****
* calltf application
*****
PLIBook_MyMonitor_calltf(char *user_data)
{
    vpiHandle systf_h, arg_itr, mod_h, net_itr, net_h;
    s_vpi_time time_s;
    s_vpi_value value_s;
    s_cb_data cb_data_s;
    char *net_name_temp, *net_name_keep;

    /* setup value change callback options */
    time_s.type      = vpiScaledRealTime;
    value_s.format   = vpiBinStrVal;

    cb_data_s.reason = cbValueChange;
    cb_data_s.cb_rtn = PLIBook_MyMonitor_callback;
    cb_data_s.time   = &time_s;
    cb_data_s.value   = &value_s;

    /* obtain a handle to the system task instance */
    systf_h = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handle to system task argument */
    /* compiletf has already verified only 1 arg with correct type */
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    mod_h = vpi_scan(arg_itr);
    vpi_free_object(arg_itr); /* free iterator--did not scan to null */

    /* add value change callback for each net in module named in tfarg */
    vpi_printf("\nAdding monitors to all nets in module %s:\n\n",
              vpi_get_str(vpiDefName, mod_h));

    net_itr = vpi_iterate(vpiNet, mod_h);
    while ((net_h = vpi_scan(net_itr)) != NULL) {
        net_name_temp = vpi_get_str(vpiFullName, net_h);
        net_name_keep = malloc(strlen(net_name_temp)+1);
        strcpy(net_name_keep, net_name_temp);
        cb_data_s.obj = net_h;
        cb_data_s.user_data = net_name_keep;
        vpi_register_cb(&cb_data_s);
    }
}

```

```
*****  
* Value change callback application  
*****  
PLIbook_MyMonitor_callback(p_cb_data cb_data_p)  
{  
    vpi_printf("At time %0.2f:\t %s = %s\n",  
              cb_data_p->time->real,  
              cb_data_p->user_data,  
              cb_data_p->value->value.str);  
}
```

6.8 Summary

This chapter has presented one of the more powerful aspects of the VPI routines and the Programming Language Interface—the ability to synchronize PLI application calls with different types of simulation activity. The examples presented in this chapter have shown ways to synchronize PLI applications with the start of simulation, with the end of a simulation time step, with future simulation times, and with logic value changes. Another important topic presented in this chapter was how to allocate and store instance-specific data, and share this data between the different routines that make up a PLI application. A critical consideration in sharing and preserving data is that the data must be *instance specific*. This means unique storage for the data must be allocated for each instance of a system task or system function.

The next chapter presents how the PLI can be used to interface C language models with Verilog simulation. This interface will utilize many of the concepts which were presented in this chapter.

CHAPTER 7

Interfacing to C Models Using VPI Routines

Interfacing C language models to Verilog simulations is a common and powerful application of the Programming Language Interface. The VPI *simulation callback routines* presented in the previous chapter make it easy to create this interface, and to synchronize activity with logic value changes and with simulation time. This chapter shows several ways in which a C model can be interfaced to a Verilog simulation using the VPI routine library (Chapter 13 presents using the TF library for interfacing to C models, and Chapter 18 shows how to use the ACC library to accomplish this same task).

The concepts presented in this chapter are:

- Representing hardware models in C
- Verilog HDL shell modules
- Combinational logic interfaces to C models
- Sequential logic interfaces to C models
- Synchronizing with the end of a simulation time step
- Synchronizing with a future simulation time step
- Multiple instances of a C model
- Creating instance specific storage within C models
- Representing propagation delays in C models

**TIP**

One reason for representing hardware models in the C language is to achieve faster simulation performance. The C programming language allows a very abstract, algorithmic representation of hardware functionality, without representing detailed timing, multi-state logic, hardware concurrency and the many other hardware specific details offered by the Verilog language.

The PLI can be a means to access the efficiency of a highly abstract C model. However, a poorly written PLI application can become a bottleneck that offsets much of the efficiency gains. Care must be taken to write PLI applications that execute as efficiently as possible.

Some guidelines that can help maximize the efficiency and run-time performance of PLI applications are:

- Good C programming practices are essential. General C programming style and techniques are not discussed within the scope of this book.
- Consider every call to a PLI routine as expensive, and try to minimize the number of calls.
- Routines which convert logic values from a simulator's internal representation to C strings, and vice-versa, are very expensive in terms of performance. Best efficiency is attained when the value representation in C is as similar as possible to the value representation in Verilog.
- Use the Verilog language to model the things hardware description languages do well, such as representing hardware parallelism and hardware propagation times. Simulator vendors have invested a great deal in optimizing a simulator's algorithms, and that optimization should be utilized.

**NOTE**

The objective of this book is to show several ways in which the VPI library can be used to interface to C models. Short examples are presented that are written in a relatively easy to follow C coding style. In order to meet the book's objectives, the examples presented in this book do not always follow the guidelines of efficient C coding and prudent usage of the PLI routines. It is expected that when parts of these example PLI applications are adapted for other applications, the coding style will also be modified to be more efficient and robust.

7.1 How to interface C models with Verilog simulations

The power and flexibility of the C programming language and the Verilog PLI provide a wide variety of methods that can be used to interface a Verilog simulation with a C language model. All methods, however, have three essential concepts in common:

- Value changes which occur in the Verilog simulator must be passed to the C model.
- Value changes within the C model must be passed to the Verilog simulation.


- Simulated time in both the Verilog simulation and the C model must remain synchronized.

This chapter will present one of the more common methods of interfacing a Verilog simulation with a C model. The method presented is by no means the only way this interface can be accomplished, and may not always be the most efficient method. However, the method presented does have many advantages, including simplicity to implement, portability to many types of Verilog simulators, and the ability to use the C model any number of times and at any level in the hierarchy of a Verilog design.

The fundamental steps that are presented in this chapter are:

1. Create the C language model as an independent block of pure C code that does not use the PLI routines in any way. The C model will have inputs and outputs, but it will not know where the inputs come from or where the outputs go to. The C code to implement the model might be in the form of a C function with no main function, or it might be a complete C program with its own main function.
2. Create a Verilog HDL *shell module* (also called a *wrapper module*) to represent the inputs and outputs of the C language model. This module will be written completely in the Verilog language, but will not contain any functionality. To represent the functionality of the model, the shell module will call a PLI application.
3. Create a PLI application to serve as an interface between the C model and the Verilog shell module. The PLI application is a communication channel, which:
 - Uses the PLI routines to retrieve data from the Verilog HDL shell module and pass the data to the C model via standard C programming.
 - Uses standard C programming to receive data from the C model, and passes the data to the Verilog shell module via PLI routines.

The following diagram shows how the blocks which are created in these three steps interact with each other.



This chapter presents steps 2 and 3 of this interface method in detail. Step 1 is to model some desired functionality or algorithm in the C language. This step is pure C programming, which does not directly involve the Verilog language or the Verilog PLI. This chapter does not cover how to implement ideas in the C language—the focus is on how to interface that implementation with a Verilog simulation. To maintain this focus, the C model example presented in this chapter will be a practical example, but relatively simple to model in C. The C model example used illustrates all of the important concepts of integrating C models into a Verilog simulation.

7.2 Creating the C language model

A hardware model can be represented in the C programming language in two basic forms, either as a C function or as an independent C program.

7.2.1 Using functions to represent the C model

When the C model is represented as a C function, that function can be linked into the Verilog simulator, together with the PLI application that serves as the interface to the model. The PLI application can then call the function when needed, passing inputs to the function, and receiving outputs from the function. One advantage of representing a C model as a function is the simplicity of passing values to and from the model. Another advantage is ease of porting to different operating systems, since the C model is called directly from the PLI application as a C function. A disadvantage of using a function to represent the C model is that the C model must contain additional code to allow a Verilog design to instantiate the C model multiple times. The model needs to specifically create unique storage for each instance.

7.2.2 Using independent programs to represent the C model

When the C model is represented as an independent program, which means it has its own C *main* function, then the Verilog simulation and the C model can be run as parallel processes on the same or on different computers. The PLI application which serves as an interface between the simulation and the model will need to create and maintain some type of communication channel between the two programs. This communication can be accomplished several ways, such as using the *exec* command in the C standard library. On Unix operating systems, the *fork* or *vfork* commands with either Unix pipes or Unix sockets is an efficient method to communicate with the C model program. On PC systems running a DOS or windows operating system, the *spawn* command can be used to invoke the C model program and establish two-way communications between the PLI application and the C model process.


One of the advantages of representing the C model as an independent model is the ability to have parallel processes running on the same computer or separate computers. Another advantage is that when a Verilog design instantiates multiple instances of the C model, each instance will be a separate process with its own memory storage. The major disadvantage of independent programs when compared to using a C function to represent the C model is that the PLI interface to invoke and communicate with the separate process is more complex, and might be operating system dependent.

David Roberts, of Cadence Design Systems, who reviewed many of the chapters of this book, has provided a full example of representing a C model as a separate C program. This example is included with the CD that accompanies this book.

7.3 A C model example

The C model used for different PLI interfaces shown in this chapter is a scientific Arithmetic Logic Unit, which utilizes the C math library. The C model is represented as a C function, which will be called from the PLI interface mechanism. This model is written entirely with the standard C library routines and C data types, without reference to any PLI routines or PLI data types. This same example is also used in other chapters, to show how a PLI interface to C models can be created using the TF and ACC libraries of the PLI.

The inputs and outputs of the scientific ALU C model are shown below, and Table 7-1 shows the operations which the ALU performs.



exception is set to 1 whenever an operation results in a value which is out of range of the double-precision result.

error is set to 1 whenever an input to an operation is out of range for the operation.

Opcode	C Math Library Operation
0	<code>pow(a, b)</code> — returns a to the power of b
1	<code>sqrt(a)</code> — returns the square root of a
2	<code>exp(a)</code> — returns the natural exponent of a
3	<code>ldexp(a, b)</code> — returns $a * (2 \text{ to the power of } b)$
4	<code>fabs(a)</code> — returns the absolute of a
5	<code>fmod(a, b)</code> — returns the floating remainder of a / b
6	<code>ceil(a)</code> — returns smallest whole number not less than a
7	<code>floor(a)</code> — returns largest whole number not more than a
8	<code>log(a)</code> — returns the natural log of a
9	<code>log10(a)</code> — returns the base 10 log of a
A	<code>sin(a)</code> — returns the sine of a
B	<code>cos(a)</code> — returns the cosine of a
C	<code>tan(a)</code> — returns the tangent of a
D	<code>asin(a)</code> — returns the arcsine of a
E	<code>acos(a)</code> — returns the arccosine of a
F	<code>atan(a)</code> — returns the arctangent of a

Table 7-1: Scientific ALU C model operations

The source code for the scientific ALU is listed in Example 7-1. This version of the ALU uses combinational logic (the outputs change whenever an input changes).



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.07/sci_alu_combinational_vpi.c`
- Verilog shell module: `Chapter.07/sci_alu_combinational_shell.v`
- Verilog test bench: `Chapter.07/sci_alu_combinational_test.v`
- Verilog-XL results log: `Chapter.07/sci_alu_combinational_test.log`

Example 7-1: scientific ALU C model — combinational logic version

```
#include <math.h>
#include <ERRNO.h>
void PLIbook_ScientificALU_C_model(
    double *result,      /* output from ALU */
    int    *excep,       /* output; set if result is out of range */
    int    *err,         /* output; set if input is out of range */
    double a,            /* input */
    double b,            /* input */
    int    opcode)       /* input */

{
    switch (opcode) {
        case 0x0: *result = pow    (a, b);      break;
        case 0x1: *result = sqrt   (a);          break;
        case 0x2: *result = exp    (a);          break;
        case 0x3: *result = ldexp  (a, (int)b);  break;
        case 0x4: *result = fabs   (a);          break;
        case 0x5: *result = fmod   (a, b);       break;
        case 0x6: *result = ceil   (a);          break;
        case 0x7: *result = floor  (a);          break;
        case 0x8: *result = log    (a);          break;
        case 0x9: *result = log10  (a);          break;
        case 0xA: *result = sin    (a);          break;
        case 0xB: *result = cos    (a);          break;
        case 0xC: *result = tan    (a);          break;
        case 0xD: *result = asin   (a);          break;
        case 0xE: *result = acos   (a);          break;
        case 0xF: *result = atan   (a);          break;
    }
    *err   = (errno == EDOM);    /* arg to math func. out of range */
    *excep = (errno == ERANGE); /* result of math func. out of range */
    errno = 0;                  /* clear the error flag */
    if (*err) *result = 0.0;    /* set result to 0 if error occurred */
    return;
}
```

7.4 Creating a Verilog shell module

A **shell module** allows a Verilog design to reference a C model using standard Verilog HDL syntax. The shell module is a Verilog module which has the same input and output ports as the C model, but the module has no functionality modeled within. To represent the module's functionality, the shell module invokes a PLI application, which in turn invokes the C model. A shell module is sometimes referred to as a *wrapper module*, because the module is wrapped around the call to a PLI application.

The shell module for a combinational logic version of the scientific ALU is listed in Example 7-2.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.07/sci_alu_combinational_vpi.c
- Verilog shell module: Chapter.07/sci_alu_combinational_shell.v
- Verilog test bench: Chapter.07/sci_alu_combinational_test.v
- Verilog-XL results log: Chapter.07/sci_alu_combinational_test.log

Example 7-2: Verilog shell module for the scientific ALU C model

```
'timescale 1ns / 1ns
module scientific_alu(a_in, b_in, opcode,
                      result_out, exception, error);
  input  [63:0] a_in, b_in;
  input  [3:0]  opcode;
  output [63:0] result_out;
  output      exception, error;

  reg           exception, error;
  real          a, b, result; // real variables used in this module

  // convert real numbers to/from 64-bit vector port connections
  assign result_out = $realtobits(result);
  always @(a_in) a = $bitstoreal(a_in);
  always @(b_in) b = $bitstoreal(b_in);

  //call the PLI application which interfaces to the C model
  initial
    $scientific_alu(a, b, opcode, result, exception, error);

endmodule
```



In this scientific ALU example, the primary inputs and outputs of the model are double-precision floating point values, represented as Verilog real data types. The Verilog language does not permit real numbers to be connected to module ports. However, the language provides built-in system functions which convert real numbers to 64-bit vectors, and vice-versa, so the real values can be passed through a module port connection. These built-in system functions are `$realtobits()` and `$bitstoreal()`.

The Verilog shell module that represents the C model can be instantiated in a design in the same way as any other Verilog module. For example:

```
module chip (...)
  ...
  scientific_alu u1 (a, b, opcode, result, excep, err);
  ...
endmodule
```

Creating a shell module to represent the C model is not mandatory—the PLI application could be called directly from any place in a Verilog design. However, there are important advantages to using a shell module to represent the C model:

- The shell module provides a simple method to encapsulate the C model.
- The shell module can be instantiated anywhere in a Verilog design hierarchy.
- The shell module can be instantiated any number of times in a Verilog design.
- The shell module can add Verilog HDL delays to the C model, which can accurately represent rise and fall delay delays, state-dependent delays, and timing constraints such as setup times.
- Delays within a shell module can be annotated using delay calculators or SDF files to provide additional delay accuracy for each instance of the shell module.

Section 7.10 later in this chapter discusses how delays can be represented in the Verilog shell module.

7.5 Creating a combinational logic interface to a C model

In a combinational logic model, the outputs of the model continuously reflect the input values of the model. The inputs are asynchronous—when any input changes value, the model outputs are re-evaluated to reflect the input change.

The VPI library can access the arguments of a system task. In the discussion of the Verilog shell module in the previous section of this chapter, it was suggested that a system task be created to represent the C model interface, and that this task list all of the C model inputs and outputs as arguments. This gives a PLI application easy access to these inputs and outputs using the VPI library.

A *simulation callback routine* provides a simple method of creating a combinational logic interface to a C model. A callback can be used to schedule a *simulation callback routine* whenever an argument of a system task changes value. The routine can then read the input values from the system task arguments, and pass the input values to the

C model. The outputs of the C model are then passed back to the Verilog simulation by writing the results onto the system task arguments in the Verilog shell module. Chapter 6 presented how to register callbacks to a *simulation callback routine* for logic value changes.

The basic steps involved with using a *simulation callback routine* to implement a combinational logic interface are:

1. Create a PLI application system task to represent the interface between the Verilog shell module and the C model. The system task is invoked from the shell module, and all of the C model inputs and outputs are listed as arguments to the system task. For example:

```
initial  
  $scientific_alu(a, b, opcode, result, exception, error);
```

Note that, in this example, the system task is called from a Verilog **initial** procedure, which means the system task will only be invoked one time for each instance of the shell module.

2. In the *callif routine* associated with the system task, register logic value change callbacks for a *simulation callback routine*. A callback is registered for each system task argument which represents an input to the C model, and each callback invokes the same *simulation callback routine*.
3. In the *simulation callback routine*, which is called whenever a system task argument changes value, read the values of all inputs and pass the values to the C model. The output values of the C model are returned to the same *simulation callback routine*, which then writes the values to the system task arguments that represent the outputs of the C model.

Obtaining object handles within simulation callback routines

NOTE → The VPI *simulation callback routine* is not directly associated with the system task which represents the C model interface. This means the *callback routine* cannot obtain handles to the system task arguments unless the callback is passed a handle to the system task or is passed the argument handles.

The *simulation callback routine* needs the handles for the system task arguments which represent the C model, in order to read the input values and write the C model output values. Since the *simulation callback routine* is not associated with the system task, access to the task arguments must be passed to the routine. This information can be passed in at least two ways:

- When the *calltf routine* registers the callbacks, the *calltf routine* can obtain a handle for the system task instance, and pass the handle to the *simulation callback routine* through the *user_data* field. The *simulation callback routine* can then use the system task handle to obtain the handles for the task arguments.
- When the *calltf routine* registers callbacks, the *calltf routine* can allocate persistent storage, obtain handles for all task arguments, store the handles in the storage block, and pass a pointer to the storage to the *simulation callback routine* through the *user_data* field. The *simulation callback routine* can then obtain the argument handles from the storage block.

Each of these methods has advantages. Saving just the handle for the system task instance is simpler, and makes the instance handle available for other uses, such as scheduling a callback to the *simulation callback routine* for that instance of the system task. However, each time a value change simulation callback occurs, the *simulation callback routine* must call `vpi_iterate()` and `vpi_scan()` to obtain handles for each of the `$scientific_alu` arguments. Saving all handles in a block of memory means the handles do not need to be obtained each time a value change occurs. Example 7-3, which follows, illustrates passing the system task instance handle. Example 7-4 on page 242 illustrates allocating a block of memory to store the handles for all system task arguments, and passing a pointer to the memory to the consumer routine.

The following example implements a combinational logic interface for the scientific ALU C model.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.07/sci_alu_combinational_vpi.c
- Verilog shell module: Chapter.07/sci_alu_combinational_shell.v
- Verilog test bench: Chapter.07/sci_alu_combinational_test.v
- Verilog-XL results log: Chapter.07/sci_alu_combinational_test.log

Example 7-3: combinational logic C model interface using VPI routines

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include "vpi_user.h"      /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"      /* IEEE 1364 PLI TF routine library
                           (using TF routines for simulation control) */

/* prototypes of routines in this PLI application */
int PLIbook_ScientificALU_calltf(), PLIbook_ScientificALU_compiletf();
int PLIbook_ScientificALU_interface();
```

```
*****
 * VPI Registration Data
 *****
void PLIbook_ScientificALU_register()
{
    s_vpi_systf_data tf_data;
    tf_data.type      = vpiSysTask;
    tf_data.tfname   = "$scientific_alu";
    tf_data.calltf   = PLIbook_ScientificALU_calltf;
    tf_data.compiletf = PLIbook_ScientificALU_compiletf;
    tf_data.sizetf   = NULL;
    tf_data.user_data = NULL;
    vpi_register_systf(&tf_data);
}

*****
 * Definition for a structure to hold the data to be passed from
 * calltf application to the ALU interface.
 *****
typedef struct PLIbook_ScientificALU_data {
    vpiHandle a_h, b_h, opcode_h, result_h, excep_h, err_h;
} PLIbook_ALU_data_s, *PLIbook_ALU_data_p;

*****
 * Value change simulation callback routine: Serves as an interface
 * between Verilog simulation and the C model. Called whenever the
 * C model inputs change value, passes the values to the C model, and
 * puts the C model outputs into simulation.
 *
 * NOTE: The handles for the arguments to $scientific_alu were obtained
 * in the calltf routine and saved in application-allocated memory. A
 * pointer to this memory is passed to this callback via the user_data
 * field.
 *****
int PLIbook_ScientificALU_interface(p_cb_data cb_data)
{
    double      a, b, result;
    int         opcode, excep, err;
    s_vpi_value value_s;

    PLIbook_ALU_data_p ALUdata;

    /* Retrieve pointer to ALU data structure from callback user_data. */
    /* The structure contains the handles for the $scientific_alu args */
    ALUdata = (PLIbook_ALU_data_p)cb_data->user_data;

    /* Read current values of C model inputs from Verilog simulation */
    value_s.format = vpiRealVal;
    vpi_get_value(ALUdata->a_h, &value_s);
    a = value_s.value.real;

    vpi_get_value(ALUdata->b_h, &value_s);
    b = value_s.value.real;
```

```
value_s.format = vpiIntVal;
vpi_get_value(ALUdata->opcode_h, &value_s);
opcode = value_s.value.integer;

***** Call the C model *****/
PLIbook_ScientificALU_C_model(&result, &excep, &err, a, b, opcode);

/* Write the C model outputs onto the Verilog signals */
value_s.format = vpiRealVal;
value_s.value.real = result;
vpi_put_value(ALUdata->result_h, &value_s, NULL, vpiNoDelay);

value_s.format = vpiIntVal;
value_s.value.integer = excep;
vpi_put_value(ALUdata->excep_h, &value_s, NULL, vpiNoDelay);

value_s.value.integer = err;
vpi_put_value(ALUdata->err_h, &value_s, NULL, vpiNoDelay);

return(0);
}

*****
* calltf routine: Registers a callback to the C model interface
* whenever any input to the C model changes value
*****
int PLIbook_ScientificALU_calltf(char *user_data)
{
    vpiHandle instance_h, arg_itr;
    s_vpi_value value_s;
    s_vpi_time time_s;
    s_cb_data cb_data_s;

    PLIbook_ALU_data_p ALUdata;

    /* allocate storage to hold $scientific_alu argument handles */
    ALUdata = (PLIbook_ALU_data_p)malloc(sizeof(PLIbook_ALU_data_s));

    /* obtain a handle to the system task instance */
    instance_h = vpi_handle(vpiSysTfcall, NULL);

    /* obtain handles to system task arguments */
    /* compiletf has already verified arguments are correct */
    arg_itr = vpi_iterate(vpiArgument, instance_h);
    ALUdata->a_h      = vpi_scan(arg_itr); /* 1st arg is a input */
    ALUdata->b_h      = vpi_scan(arg_itr); /* 2nd arg is b input */
    ALUdata->opcode_h = vpi_scan(arg_itr); /* 3rd arg is opcode input */
    ALUdata->result_h = vpi_scan(arg_itr); /* 4th arg is result output */
    ALUdata->excep_h  = vpi_scan(arg_itr); /* 5th arg is excep output */
    ALUdata->err_h    = vpi_scan(arg_itr); /* 6th arg is error output */
    vpi_free_object(arg_itr); /* free iterator--did not scan to null */
}
```

```
/* setup value change callback options */
time_s.type      = vpiSuppressTime;
cb_data_s.reason = cbValueChange;
cb_data_s.cb_rtn = PLIBook_ScientificALU_interface;
cb_data_s.time   = &time_s;
cb_data_s.value  = &value_s;

/* add value change callbacks to all signals which are inputs to */
/* pass pointer to storage for handles as user_data value */
cb_data_s.user_data = (char *)ALUdata;
value_s.format = vpiRealVal;
cb_data_s.obj = ALUdata->a_h;
vpi_register_cb(&cb_data_s);

cb_data_s.obj = ALUdata->b_h;
vpi_register_cb(&cb_data_s);

value_s.format = vpiIntVal;
cb_data_s.obj = ALUdata->opcode_h;
vpi_register_cb(&cb_data_s);

return(0);
}

/*********************************************
 * compilelf routine: Verifies that $scientific_alu() is used correctly
 * Note: For simplicity, only limited data types are allowed for
 * task arguments. Could add checks to allow other data types.
********************************************/
int PLIBook_ScientificALU_compilelf(char *user_data)
{
    vpiHandle systf_h, arg_itr, arg_h;
    int        err = 0;

    systf_h = vpi_handle(vpiSysTfCall, NULL);
    arg_itr = vpi_iterate(vpiArgument, systf_h);
    if (arg_itr == NULL) {
        vpi_printf("ERROR: $scientific_alu requires 6 arguments\n");
        tf_dofinish();
        return(0);
    }

    arg_h = vpi_scan(arg_itr); /* 1st arg is a input */
    if (vpi_get(vpiType, arg_h) != vpiRealVar) {
        vpi_printf("$scientific_alu arg 1 (a) must be a real variable\n");
        err = 1;
    }

    arg_h = vpi_scan(arg_itr); /* 2nd arg is b input */
    if (vpi_get(vpiType, arg_h) != vpiRealVar) {
        vpi_printf("$scientific_alu arg 2 (b) must be a real variable\n");
        err = 1;
    }
}
```

```
arg_h = vpi_scan(arg_itr); /* 3rd arg is opcode input */
if (vpi_get(vpiType, arg_h) != vpiNet) {
    vpi_printf("$scientific_alu arg 3 (opcode) must be a net\n");
    err = 1;
}
else if (vpi_get(vpiSize, arg_h) != 4) {
    vpi_printf("$scientific_alu arg 3 (opcode) must be 4-bit vector\n");
    err = 1;
}

arg_h = vpi_scan(arg_itr); /* 4th arg is result output */
if (vpi_get(vpiType, arg_h) != vpiRealVar) {
    vpi_printf("$scientific_alu arg 4 (result) must be a real var.\n");
    err = 1;
}

arg_h = vpi_scan(arg_itr); /* 5th arg is exception output */
if (vpi_get(vpiType, arg_h) != vpiReg) {
    vpi_printf("$scientific_alu arg 5 (exception) must be a reg\n");
    err = 1;
}
else if (vpi_get(vpiSize, arg_h) != 1) {
    vpi_printf("$scientific_alu arg 5 (exception) must be scalar\n");
    err = 1;
}

arg_h = vpi_scan(arg_itr); /* 6th arg is error output */
if (vpi_get(vpiType, arg_h) != vpiReg) {
    vpi_printf("$scientific_alu arg 6 (error) must be a reg\n");
    err = 1;
}
else if (vpi_get(vpiSize, arg_h) != 1) {
    vpi_printf("$scientific_alu arg 6 (error) must be scalar\n");
    err = 1;
}

if (vpi_scan(arg_itr) != NULL) { /* should not be any more args */
    vpi_printf("ERROR: $scientific_alu requires only 6 arguments\n");
    vpi_free_object(arg_itr);
    err = 1;
}

if (err) {
    tf_dofinish();
    return(0);
}
}
```

7.6 Creating a sequential logic interface to a C model

In a sequential logic model, the outputs of the model change synchronously with an input strobe, such as a positive edge of a clock. There may also be one or more asynchronous inputs, such as a reset signal.

The VPI *simulation callback routine* provides a straightforward way to model a sequential logic interface to a C model. Callbacks are only registered for a logic value change on the clock input and any asynchronous inputs of the C model. Only when those inputs change, are the new input values passed to the C model. The outputs of the C model are then passed back to the Verilog simulation by writing the results onto the system task arguments in the Verilog shell module which represents the outputs of the C model.

The basic steps involved with using the Value Change Link to implement a synchronous sequential logic interface are very similar to implementing a combinational logic interface. The one difference is that simulation callbacks are only registered for specific C model inputs instead of all inputs.

Regardless of whether the interface is combinational or sequential logic, the *simulation callback routine* needs the handles for the system task arguments which represents the C model, in order to read the inputs values and write the C model output values. Access to the task arguments must be passed to the *simulation callback routine* through the callback's `user_data` field.

Example 7-4, which follows, implements a sequential logic interface for the scientific ALU C model, where all inputs are synchronized to value changes of a clock input.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.07/sci_alu_sequential_vpi.c
- Verilog shell module: Chapter.07/sci_alu_sequential_shell.v
- Verilog test bench: Chapter.07/sci_alu_sequential_test.v
- Verilog-XL results log: Chapter.07/sci_alu_sequential_test.log

Example 7-4: sequential logic C model interface using VPI routines

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>        /* ANSI C standard input/output library */
#include "vpi_user.h"     /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"      /* IEEE 1364 PLI TF routine library
                           (using TF routines for simulation control) */
```

```
/* prototypes of routines in this PLI application */
int PLIbook_ScientificALU_calltf(), PLIbook_ScientificALU_compiletf();
int PLIbook_ScientificALU_interface();

/********************* VPI Registration Data *********************/
void PLIbook_ScientificALU_register()
{
    s_vpi_systf_data tf_data;
    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$scientific_alu";
    tf_data.calltf   = PLIbook_ScientificALU_calltf;
    tf_data.compiletf = PLIbook_ScientificALU_compiletf;
    tf_data.sizetf    = NULL;
    tf_data.user_data = NULL;
    vpi_register_systf(&tf_data);
}

/********************* Definition for a structure to hold the data to be passed from
 * calltf application to the ALU interface.
*********************/
typedef struct PLIbook_ScientificALU_data {
    vpiHandle clock_h, a_h, b_h, opcode_h, result_h, excep_h, err_h;
} PLIbook_ALU_data_s, *PLIbook_ALU_data_p;

/********************* Value change simulation callback routine: Serves as an interface
 * between Verilog simulation and the C model. Called whenever the
 * C model inputs change value, passes the values to the C model, and
 * puts the C model outputs into simulation.
 *
 * NOTE: The handles for the arguments to $scientific_alu were obtained
 * in the calltf routine and saved in application-allocated memory. A
 * pointer to this memory is passed to this callback via the user_data
 * field.
*********************/
int PLIbook_ScientificALU_interface(p_cb_data cb_data)
{
    double        a, b, result;
    int          opcode, excep, err;
    s_vpi_value  value_s;

    PLIbook_ALU_data_p ALUdata;

    /* Retrieve pointer to ALU data structure from callback user_data. */
    /* The structure contains the handles for the $scientific_alu args */
    ALUdata = (PLIbook_ALU_data_p)cb_data->user_data;
```

```
/* Read current values of C model inputs from Verilog simulation */
value_s.format = vpiRealVal;
vpi_get_value(ALUdata->a_h, &value_s);
a = value_s.value.real;

vpi_get_value(ALUdata->b_h, &value_s);
b = value_s.value.real;

value_s.format = vpiIntVal;
vpi_get_value(ALUdata->opcode_h, &value_s);
opcode = value_s.value.integer;

***** Call the C model *****/
PLIbook_ScientificALU_C_model(&result, &excep, &err, a, b, opcode);

/* Write the C model outputs onto the Verilog signals */
value_s.format = vpiRealVal;
value_s.value.real = result;
vpi_put_value(ALUdata->result_h, &value_s, NULL, vpiNoDelay);

value_s.format = vpiIntVal;
value_s.value.integer = excep;
vpi_put_value(ALUdata->excep_h, &value_s, NULL, vpiNoDelay);

value_s.value.integer = err;
vpi_put_value(ALUdata->err_h, &value_s, NULL, vpiNoDelay);

return(0);
}

*****
* calltf routine: Registers a callback to the C model interface
* whenever any input to the C model changes value
*****
int PLIbook_ScientificALU_calltf(char *user_data)
{
    vpiHandle    instance_h, arg_itr;
    s_vpi_value  value_s;
    s_vpi_time   time_s;
    s_cb_data    cb_data_s;

    PLIbook_ALU_data_p ALUdata;

    /* allocate storage to hold $scientific_alu argument handles */
    ALUdata = (PLIbook_ALU_data_p)malloc(sizeof(PLIbook_ALU_data_s));

    /* obtain a handle to the system task instance */
    instance_h = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handles to system task arguments */
    /* compiltef has already verified arguments are correct */
    arg_itr = vpi_iterate(vpiArgument, instance_h);
    ALUdata->clock_h = vpi_scan(arg_itr); /* 1st arg is clock input */
```

```
ALUdata->a_h      = vpi_scan(arg_itr); /* 2nd arg is a input */
ALUdata->b_h      = vpi_scan(arg_itr); /* 3rd arg is b input */
ALUdata->opcode_h = vpi_scan(arg_itr); /* 4th arg is opcode input */
ALUdata->result_h = vpi_scan(arg_itr); /* 5th arg is result output */
ALUdata->excep_h  = vpi_scan(arg_itr); /* 6th arg is excep output */
ALUdata->err_h    = vpi_scan(arg_itr); /* 7th arg is error output */
vpi_free_object(arg_itr); /* free iterator--did not scan to null */

/* setup value change callback options */
time_s.type        = vpiSuppressTime;
cb_data_s.reason   = cbValueChange;
cb_data_s.cb_rtn   = PLIbook_ScientificALU_interface;
cb_data_s.time     = &time_s;
cb_data_s.value    = &value_s;

/* add value change callbacks to clock input to the C model, */
/* pass pointer to storage for handles as user_data value */
value_s.format     = vpiSuppressVal;
cb_data_s.user_data = (char *)ALUdata;
cb_data_s.obj      = ALUdata->clock_h;
vpi_register_cb(&cb_data_s);

return(0);
}
```

7.7 Synchronizing with the end of a simulation time step

Within a simulation, it is possible for several signals to change at the same moment of simulation time. In Verilog simulators, the *simulation callback routine* may be called in the middle of a simulation time step for asynchronous value changes, before all input value changes have occurred for that time step.

With a combinational logic interface, the *simulation callback routine* will be called for every input change. In the combinational logic interface presented in Example 7-3, the C model is called each time the *simulation callback routine* is called for an input value change. This is the correct functionality for combinational logic—at the completion of a simulation time step, the outputs from the C model represent the most current input values. However, by synchronizing the call to the C model with the end of the simulation time step in which changes occur, the multiple calls to the C model within a time step could be optimized to a single call.

With a sequential logic interface synchronized to a clock, when the *simulation callback routine* is called at a clock change, other input changes at that moment in simulation time may or may not have occurred. It may be desirable to ensure that the C

model is not called until all inputs have their most current value for the time step in which the clock changes.

By using another *simulation callback routine*, both combinational logic and sequential logic C model interfaces can be synchronized to the end of a current simulation time step. This is done by using the asynchronous value change callback to the *simulation callback routine* to schedule a synchronous callback to a different *simulation callback routine* for end of the current time step. Chapter 6 presented how to register callbacks at the end of the current simulation time step, and this chapter uses those concepts.

Example 7-5 modifies the combinational logic interface presented in Example 7-3. This modified version schedules a simulation callback at the end of a simulation time step in which an input changed value.

This example uses a flag to indicate when a synchronous callback for the end of the current time step has already been scheduled for the current simulation time. This flag is used in the C model interface to prevent more than one synchronous callback to the *simulation callback routine* being requested in the same time step. The flag is stored in the same block of memory that was allocated to store the handles for the `$scientific_alu` arguments. Since this memory block is unique for each instance of a system task, each instance of the C model will have a unique flag.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.07/sci_alu_synchronized_vpi.c`
- Verilog shell module: `Chapter.07/sci_alu_synchronized_shell.v`
- Verilog test bench: `Chapter.07/sci_alu_synchronized_test.v`
- Verilog-XL results log: `Chapter.07/sci_alu_synchronized_test.log`

Example 7-5: C model interface synchronized to the end of a time step

```
#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>       /* ANSI C standard input/output library */
#include "vpi_user.h"    /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"    /* IEEE 1364 PLI TF routine library
                           (using TF routines for simulation control) */

/* prototypes of routines in this PLI application */
int PLIBbook_ScientificALU_calltf(), PLIBbook_ScientificALU_completetf();
int PLIBbook_ScientificALU_interface();
int PLIBbook_EndOfTimeStep_callback();
```

```
*****
 * VPI Registration Data
 *****
void PLIbook_ScientificALU_register()
{
    s_vpi_systf_data tf_data;
    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$scientific_alu";
    tf_data.calltf   = PLIbook_ScientificALU_calltf;
    tf_data.compiletf = PLIbook_ScientificALU_compiletf;
    tf_data.sizetf   = NULL;
    tf_data.user_data = NULL;
    vpi_register_systf(&tf_data);
}

*****
 * Definition for a structure to hold the data to be passed from
 * calltf application to the ALU interface.
 *****
typedef struct PLIbook_ScientificALU_data {
    vpiHandle clock_h, a_h, b_h, opcode_h, result_h, excep_h, err_h;
    short int sync_flag;
} PLIbook_ALU_data_s, *PLIbook_ALU_data_p;

*****
 * Value change simulation callback routine: Schedules a read-write
 * synchronize simulation callback at the end of the current time step.
 * Only schedules one callback for a time step.
 *****
int PLIbook_ScientificALU_interface(p_cb_data cb_data)
{
    s_cb_data          cb_data_s;
    s_vpi_time         time_s;
    PLIbook_ALU_data_p ALUdata;

    /* Retrieve pointer to ALU data structure from VCL user_data field */
    /* The structure contains a flag indicating if a synchronize        */
    /* callback has already been scheduled */
    ALUdata = (PLIbook_ALU_data_p)cb_data->user_data;

    /* If the sync_flag is 0, then no read-write synchronize callback */
    /* has been scheduled for this time step (the sync_flag is set to */
    /* 1 by this routine, and set to 0 by the read-write synchronize */
    /* callback after a callback is processed. */

    if (!ALUdata->sync_flag) {
        /* Schedule a synchronize simulation callback for this instance */
        ALUdata->sync_flag = 1; /* set sync_flag */
        time_s.type        = vpiSuppressTime;
        cb_data_s.reason   = cbReadWriteSynch;
        cb_data_s.user_data = (char *)ALUdata;
        cb_data_s.cb_rtn   = PLIbook_EndOfTimeStep_callback;
        cb_data_s.obj       = NULL;
```

```
    cb_data_s.time      = &time_s;
    cb_data_s.value     = NULL;
    vpi_register_cb(&cb_data_s);
}
return(0);
}

//*****************************************************************************
/* Read-write synchronize simulation callback routine: Serves as an
 * interface between Verilog simulation and the C model. Passes the
 * values to the C model, and puts the C model outputs into simulation.
 *
 * NOTE: The handles for the arguments to $scientific_alu were obtained
 * in the calltf routine and saved in application-allocated memory. A
 * pointer to this memory is passed to this callback via the user_data
 * field.
*/
int PLIBook_EndOfTimeStep_callback(p_cb_data cb_data)
{
    double      a, b, result;
    int         opcode, excep, err;
    s_vpi_value value_s;

    PLIBook_ALU_data_p ALUdata;

    /* Retrieve pointer to ALU data structure from callback user_data. */
    /* The structure contains the handles for the $scientific_alu args */
    ALUdata = (PLIBook_ALU_data_p)cb_data->user_data;

    /* Set the sync_flag to 0 to indicate that this callback has been */
    /* processed */
    ALUdata->sync_flag = 0;

    /* Read current values of C model inputs from Verilog simulation */
    value_s.format = vpiRealVal;
    vpi_get_value(ALUdata->a_h, &value_s);
    a = value_s.value.real;

    vpi_get_value(ALUdata->b_h, &value_s);
    b = value_s.value.real;

    value_s.format = vpiIntVal;
    vpi_get_value(ALUdata->opcode_h, &value_s);
    opcode = value_s.value.integer;

    ***** Call the C model *****
    PLIBook_ScientificALU_C_model(&result, &excep, &err, a, b, opcode);

    /* Write the C model outputs onto the Verilog signals */
    value_s.format = vpiRealVal;
    value_s.value.real = result;
    vpi_put_value(ALUdata->result_h, &value_s, NULL, vpiNoDelay);
```

```
value_s.format = vpiIntVal;
value_s.value.integer = excep;
vpi_put_value(ALUdata->excep_h, &value_s, NULL, vpiNoDelay);

value_s.value.integer = err;
vpi_put_value(ALUdata->err_h, &value_s, NULL, vpiNoDelay);

return(0);
}

/*********************************************************************
 * calltf routine: Registers a callback to the C model interface
 * whenever any input to the C model changes value
 *****/
int PLIbook_ScientificALU_calltf(char *user_data)
{
    vpiHandle instance_h, arg_itr;
    s_vpi_value value_s;
    s_vpi_time time_s;
    s_cb_data cb_data_s;

    PLIbook_ALU_data_p ALUdata;

    /* allocate storage to hold $scientific_alu argument handles */
    ALUdata = (PLIbook_ALU_data_p)malloc(sizeof(PLIbook_ALU_data_s));

    /* obtain a handle to the system task instance */
    instance_h = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handles to system task arguments */
    /* compiletf has already verified arguments are correct */
    arg_itr = vpi_iterate(vpiArgument, instance_h);
    ALUdata->a_h      = vpi_scan(arg_itr); /* 1st arg is a input */
    ALUdata->b_h      = vpi_scan(arg_itr); /* 2nd arg is b input */
    ALUdata->opcode_h = vpi_scan(arg_itr); /* 3rd arg is opcode input */
    ALUdata->result_h = vpi_scan(arg_itr); /* 4th arg is result output */
    ALUdata->excep_h  = vpi_scan(arg_itr); /* 5th arg is excep output */
    ALUdata->err_h    = vpi_scan(arg_itr); /* 6th arg is error output */
    vpi_free_object(arg_itr); /* free iterator--did not scan to null */

    /* setup value change callback options */
    time_s.type      = vpiSuppressTime;
    cb_data_s.reason = cbValueChange;
    cb_data_s.cb_rtn = PLIbook_ScientificALU_interface;
    cb_data_s.time   = &time_s;
    cb_data_s.value  = &value_s;

    /* add value change callbacks to all signals which are inputs to */
    /* pass pointer to storage for handles as user_data value */
    cb_data_s.user_data = (char *)ALUdata;
    value_s.format = vpiRealVal;
    cb_data_s.obj = ALUdata->a_h;
    vpi_register_cb(&cb_data_s);
```

```
cb_data_s.obj = ALUdata->b_h;
vpi_register_cb(&cb_data_s);

value_s.format = vpiIntVal;
cb_data_s.obj = ALUdata->opcode_h;
vpi_register_cb(&cb_data_s);

/* clear the callback sync_flag to indicate that no read-write */
/* synchronize callbacks have been processed */
ALUdata->sync_flag = 0;

return(0);
}
```

7.8 Synchronizing with a future simulation time step

In certain C model applications, it may be necessary to synchronize C model activity with future simulation activity. The VPI routines can also be used to schedule callbacks to a *simulation callback routine* at a specific amount of time in the future, relative to the current simulation time.

The VPI routines can also determine the future simulation time in which the next simulation event is scheduled to occur. This provides a way for a PLI application to synchronize activity for when the Verilog simulator is processing simulation events.

Using the VPI routines to schedule future simulation callbacks was presented in Chapter 6.

7.9 Allocating storage within a C model

Special attention and care must be taken when a C model uses static variables or allocates memory. The Verilog language can instantiate a model any number of times. Each instance of the Verilog shell module creates a unique instance of the system task which invokes the PLI interface to the C model. Therefore, the *calltf routine* which is invoked by a system task instance will be unique. Any memory which is allocated by the *calltf routine* will also be unique for each instance of the system task. However, the *simulation callback routine* used to communicate with the C model is *not* instance specific. Any memory allocated in the *simulation callback routine* will be shared by all instances of the C model.

When a C model is represented as an independent program, multiple instances of the model are not a problem, as each instance will invoke a new process with unique storage for each process. When the C model is represented as a C function, however, multiple instances of the model will share the same function. The model must allow for the possibility of multiple instances, and provide unique storage for each instance.

Example 7-6 presents a latched version of the scientific ALU, which can store the result of a previous operation indefinitely, and Example 7-7 presents a combinational logic interface to the model.

This example allocates storage within the C model, which is unique storage for each version of the C model. The storage for the C model output values is allocated as part of the same instance-specific storage that holds the handles for the `$scientific_alu` arguments. The storage is allocated by the *callif routine* for each system task instance. A pointer to the storage is passed to the *simulation callback routine* as the `user_data` value, which then passes the pointer to the C model as an input to the model function.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.07/sci_alu_latched_vpi.c
- Verilog shell module: Chapter.07/sci_alu_latched_shell.v
- Verilog test bench: Chapter.07/sci_alu_latched_test.v
- Verilog-XL results log: Chapter.07/sci_alu_latched_test.log

Example 7-6: scientific ALU C model with latched outputs

```
*****
 * Structure definition to store output values when the ALU is latched.
 ****
typedef struct PLIbook_ScientificALU_outputs {
    double result;      /* stored result of previous operation */
    int     excep;
    int     err;
} PLIbook_ALU_outputs_s, *PLIbook_ALU_outputs_p;

*****
 * C model with latched outputs. When enable is 1, the ALU returns
 * the currently calculated outputs, and when 0, the ALU returns the
 * latched previous results.
 ****
#include <math.h>
#include <ERRNO.h>
void PLIbook_ScientificALU_C_model(
    double *result,    /* output from ALU */
    int    *excep,     /* output; set if result is out of range */
```

```

int      *err,          /* output; set if input is out of range */
double   a,             /* input */
double   b,             /* input */
int      opcode,        /* input */
int      enable,         /* input; 0 = latched */
PLIbook_ALU_outputs_p LatchedOutputs) /* input */

{
    if (enable) { /* ALU is not latched, calculate outputs and store */
        switch (opcode) {
            case 0x0: LatchedOutputs->result = pow      (a, b);      break;
            case 0x1: LatchedOutputs->result = sqrt     (a);      break;
            case 0x2: LatchedOutputs->result = exp      (a);      break;
            case 0x3: LatchedOutputs->result = ldexp    (a, (int)b); break;
            case 0x4: LatchedOutputs->result = fabs     (a);      break;
            case 0x5: LatchedOutputs->result = fmod     (a, b);      break;
            case 0x6: LatchedOutputs->result = ceil     (a);      break;
            case 0x7: LatchedOutputs->result = floor    (a);      break;
            case 0x8: LatchedOutputs->result = log      (a);      break;
            case 0x9: LatchedOutputs->result = log10   (a);      break;
            case 0xA: LatchedOutputs->result = sin      (a);      break;
            case 0xB: LatchedOutputs->result = cos      (a);      break;
            case 0xC: LatchedOutputs->result = tan      (a);      break;
            case 0xD: LatchedOutputs->result = asin     (a);      break;
            case 0xE: LatchedOutputs->result = acos     (a);      break;
            case 0xF: LatchedOutputs->result = atan     (a);      break;
        }
        LatchedOutputs->err  = (errno == EDOM); /* arg out of range */
        LatchedOutputs->excep = (errno == ERANGE);/* result out of range */
        errno = 0;                                /* clear the error flag */
        if (LatchedOutputs->err) LatchedOutputs->result = 0.0;
    }

    /* return the values stored in the C model */
    *result = LatchedOutputs->result;
    *err    = LatchedOutputs->err;
    *excep  = LatchedOutputs->excep;

    return;
}

```

Example 7-7: combinational logic interface to the latched scientific ALU C model

```

#include <stdlib.h>      /* ANSI C standard library */
#include <stdio.h>       /* ANSI C standard input/output library */
#include "vpi_user.h"    /* IEEE 1364 PLI VPI routine library */
#include "veriuser.h"    /* IEEE 1364 PLI TF routine library
                           (using TF routines for simulation control) */

/* prototypes of routines in this PLI application */
int PLIbook_ScientificALU_calltf(), PLIbook_ScientificALU_compiletf();
int PLIbook_ScientificALU_interface();

```

```
*****
 * VPI Registration Data
 *****
void PLIbook_ScientificALU_register()
{
    s_vpi_systf_data tf_data;
    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$scientific_alu";
    tf_data.calltf   = PLIbook_ScientificALU_calltf;
    tf_data.compiletf = PLIbook_ScientificALU_compiletf;
    tf_data.sizetf   = NULL;
    tf_data.user_data = NULL;
    vpi_register_systf(&tf_data);
}

*****
 * Definition for a structure to hold the data to be passed from
 * calltf application to the ALU interface. Also allocates a structure
 * to store the latched output values of the ALU. This storage is
 * allocated for each instance of the C model.
 *****
typedef struct PLIbook_ScientificALU_data {
    vpiHandle enable_h, a_h, b_h, opcode_h, result_h, excep_h, err_h;
    PLIbook_ALU_outputs_s LatchedOutputs; /* storage for outputs */
} PLIbook_ALU_data_s, *PLIbook_ALU_data_p;

*****
 * Value change simulation callback routine: Serves as an interface
 * between Verilog simulation and the C model. Called whenever the
 * C model inputs change value, passes the values to the C model, and
 * puts the C model outputs into simulation.
 *
 * NOTE: The handles for the arguments to $scientific_alu were obtained
 * in the calltf routine and saved in application-allocated memory. A
 * pointer to this memory is passed to this callback via the user_data
 * field.
 *****
int PLIbook_ScientificALU_interface(p_cb_data cb_data)
{
    double        a, b, result;
    int           opcode, excep, err, enable;
    s_vpi_value  value_s;

    PLIbook_ALU_data_p ALUdata;

    /* Retrieve pointer to ALU data structure from callback user_data. */
    /* The structure contains the handles for the $scientific_alu args */
    ALUdata = (PLIbook_ALU_data_p)cb_data->user_data;

    /* Read current values of C model inputs from Verilog simulation */
    value_s.format = vpiRealVal;
    vpi_get_value(ALUdata->a_h, &value_s);
    a = value_s.value.real;
```

```
vpi_get_value(ALUdata->b_h, &value_s);
b = value_s.value.real;

value_s.format = vpiIntVal;
vpi_get_value(ALUdata->opcode_h, &value_s);
opcode = value_s.value.integer;

vpi_get_value(ALUdata->enable_h, &value_s);
enable = value_s.value.integer;

/* ***** Call the C model *****/
PLIbook_ScientificALU_C_model(&result, &excep, &err, a, b, opcode,
                               enable, &ALUdata->LatchedOutputs);

/* Write the C model outputs onto the Verilog signals */
value_s.format = vpiRealVal;
value_s.value.real = result;
vpi_put_value(ALUdata->result_h, &value_s, NULL, vpiNoDelay);

value_s.format = vpiIntVal;
value_s.value.integer = excep;
vpi_put_value(ALUdata->excep_h, &value_s, NULL, vpiNoDelay);

value_s.value.integer = err;
vpi_put_value(ALUdata->err_h, &value_s, NULL, vpiNoDelay);

return(0);
}

/*********************************************
 * calltf routine: Registers a callback to the C model interface
 * whenever any input to the C model changes value
******************************************/
int PLIbook_ScientificALU_calltf(char *user_data)
{
    vpiHandle     instance_h, arg_itr;
    s_vpi_value   value_s;
    s_vpi_time    time_s;
    s_cb_data    cb_data_s;
    PLIbook_ALU_data_p ALUdata;

    /* allocate storage to hold $scientific_alu argument handles */
    ALUdata = (PLIbook_ALU_data_p)malloc(sizeof(PLIbook_ALU_data_s));

    /* obtain a handle to the system task instance */
    instance_h = vpi_handle(vpiSysTfCall, NULL);

    /* obtain handles to system task arguments */
    /* compiletf has already verified arguments are correct */
    arg_itr = vpi_iterate(vpiArgument, instance_h);
    ALUdata->enable_h = vpi_scan(arg_itr); /* 1st arg is enable input */
    ALUdata->a_h      = vpi_scan(arg_itr); /* 2nd arg is a input */
```

```
ALUdata->b_h      = vpi_scan(arg_itr); /* 3rd arg is b input */
ALUdata->opcode_h = vpi_scan(arg_itr); /* 4th arg is opcode input */
ALUdata->result_h = vpi_scan(arg_itr); /* 5th arg is result output */
ALUdata->excep_h  = vpi_scan(arg_itr); /* 6th arg is excep output */
ALUdata->err_h    = vpi_scan(arg_itr); /* 7th arg is error output */
vpi_free_object(arg_itr); /* free iterator--did not scan to null */

/* setup value change callback options */
time_s.type        = vpiSuppressTime;
cb_data_s.reason  = cbValueChange;
cb_data_s.cb_rtn   = PLIbook_ScientificALU_interface;
cb_data_s.time     = &time_s;
cb_data_s.value    = &value_s;

/* add value change callbacks to all signals which are inputs to */
/* pass pointer to storage for handles as user_data value */
cb_data_s.user_data = (char *)ALUdata;
value_s.format     = vpiRealVal;
cb_data_s.obj      = ALUdata->a_h;
vpi_register_cb(&cb_data_s);

cb_data_s.obj      = ALUdata->b_h;
vpi_register_cb(&cb_data_s);

value_s.format     = vpiIntVal;
cb_data_s.obj      = ALUdata->opcode_h;
vpi_register_cb(&cb_data_s);

cb_data_s.obj      = ALUdata->enable_h;
vpi_register_cb(&cb_data_s);

return(0);
}
```

7.10 Representing propagation delays in a C model

Propagation delays from an input change to an output change in a C model can be represented in two ways:

- Using delays in the PLI interface.
- Using delays in the Verilog shell module.

Delays in the PLI interface are represented by specifying a delay value with the `vpi_put_value()` routine, which writes values onto the system task arguments. Either inertial or transport event propagation can be used, depending on the requirements of the C model. However, using `vpi_put_value()` does not offer a great deal

of flexibility on creating delays which are different for each instance of a model, representing minimum, typical and maximum delays, different delays for rise and fall transitions, or annotating delays using delay calculators or SDF files.

C model propagation delays can also be represented using the pin-to-pin path delays in the Verilog shell module. This method provides the greatest amount of flexibility and accuracy in modeling propagation delays. All path delays constructs can be used, as well and Verilog timing constraints.

Example 7-8 shows adding pin-to-pin path delays to the scientific ALU shell module.

NOTE Some Verilog simulators restrict the use of pin-to-pin path delays and SDF delay back annotation to Verilog models which are represented with Verilog primitives and net data types. To use path delays on a C model with these simulators, buffers must be added to all input and output ports, with net data types connected to the inputs and outputs of these buffers. Example 7-8 illustrates using buffers on all input and output ports of the scientific ALU shell module.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.07/sci_alu_with_delays_vpi.c
- Verilog shell module: Chapter.07/sci_alu_with_delays_shell.v
- Verilog test bench: Chapter.07/sci_alu_with_delays_test.v
- Verilog-XL results log: Chapter.07/sci_alu_with_delays_test.log

Example 7-8: scientific ALU Verilog shell module with pin-to-pin path delays

```
'timescale 1ns / 100ps
module scientific_alu(result_out, exception, error,
                      a_in, b_in, opcode_in);
  output [63:0] result_out;
  output      exception, error;
  input  [63:0] a_in, b_in;
  input  [3:0]  opcode_in;

  wire  [63:0] result_out, result_vector;
  wire  [63:0] a_in, a_vector;
  wire  [63:0] b_in, b_vector;
  wire  [3:0]  opcode_in, opcode_vector;
  wire      exception, error;

  reg      exception_reg, error_reg;
  real    a, b, result; // real variables used in this module
```

```
// convert real numbers to/from 64-bit vector port connections
assign result_vector = $realtobits(result);
always @(a_vector) a = $bitstoreal(a_vector);
always @(b_vector) b = $bitstoreal(b_vector);

//call the PLI application which interfaces to the C model
initial
    $scientific_alu(a, b, opcode_vector,
                    result, exception_reg, error_reg);

specify
    (a_in, b_in *> result_out, exception, error) = (5.6, 4.7);
    (opcode_in  *> result_out, exception, error) = (3.4, 3.8);
endspecify

// add buffers to all ports, with nets connected to each buffer
// (this example uses the array of instance syntax in the
// from the IEEE 1364-1995 Verilog standard
buf result_buf[63:0] (result_out, result_vector);
buf excep_buf        (exception, exception_reg);
buf error_buf         (error,      error_reg);
buf a_buf[63:0]        (a_vector, a_in);
buf b_buf[63:0]        (b_vector, b_in);
buf opcode_buf[3:0]   (opcode_vector, opcode_in);

endmodule
```



NOTE The preceding example uses the array of instances construct from the IEEE 1364-1995 Verilog standard. This construct was not supported by some Verilog simulators at the time this book was written. A work around is to create a separate instance of a buf primitive for each bit or each vector connected to a module port.

7.11 Summary

This chapter has presented several ways in which the VPI library can be used to interface a C language model with Verilog simulations. The VPI *simulation callback routine* provides a means to pass input changes to a C mode. The VPI routines to read and modify logic values allow information to be exchanged with the model. By creating a shell module which contains a system task that invokes the C model interface, the C model can be used in a Verilog design just as any other Verilog module.

This is the end of the discussion of the VPI portion of the Verilog PLI. The VPI library is the newest and most powerful portion of the PLI standard. Part Two of this book presents the TF and ACC portions of the PLI standard.

Part Two:

The TF/ACC Portion

of the Verilog PLI Standard

CHAPTER 8

Creating PLI Applications Using TF and ACC Routines

This chapter uses two short examples to introduce how the Verilog PLI works. The first example will allow a Verilog simulator to execute the ubiquitous *Hello World* C program. The second example uses the PLI to access specific activity within a Verilog simulation, by listing the name and current logic value of a Verilog net. The purpose of these examples is to introduce how to write a PLI application using the TF and ACC routine libraries in the Verilog PLI standard. Subsequent chapters in this part of the book build on the concepts presented in this chapter and show how to write more complex PLI applications.

8.1 The capabilities of the Verilog PLI

The Verilog **Programming Language Interface**, commonly referred to as the Verilog **PLI**, is a user-programmable procedural interface for Verilog digital logic simulators. The PLI gives Verilog users a way to extend the capabilities of a Verilog simulator by providing a means for the simulator to invoke other software programs, and then exchange information and synchronize activity with those programs.

The creators of the Verilog Hardware Description Language (HDL), which was first introduced in 1985, intended that the Verilog language should be focused on modeling digital logic. They felt that design verification tasks not directly related to modeling the design itself, such as reading and writing disk files, were already capably handled in other languages and should not be replicated in the Verilog HDL. Instead, the creators of Verilog added a procedural interface—the PLI—that provides a way for end-users of Verilog to write design verification tasks in the C programming language, and then have Verilog logic simulators execute those programs.

The capabilities of the PLI extend far beyond simple file I/O operations. The Verilog PLI allows Verilog designers to interface virtually any program to Verilog simulators. Some common engineering tasks that are required in the design of hardware, and for which the PLI is aptly fitted, include:

- C-language models

Abstract hardware models are sometimes represented in the C language instead of the Verilog HDL, perhaps to protect intellectual property or to optimize the performance of simulation. The PLI provides a number of ways to pass data from Verilog simulation to a C-language model and from the C-language model back into Verilog simulation.

- Access to programming language libraries

The PLI allows Verilog models and test programs to access the libraries of the C programming language, such as the C math library. Through the PLI, Verilog can pass arguments to a C math function, and pass the return of the math function back to the Verilog simulator.

- Reading test vector files

Test vectors are a common method of representing stimulus for simulation. A file reader program in C is easy to write, and, using the PLI, it is easy to pass values read by the C program to a Verilog simulation.

- Delay calculation

Through the PLI, a Verilog simulator can call an ASIC or FPGA vendor's delay calculator program. The delay calculator can estimate the effects that fanout, voltage, temperature and other variances will have on the delays of ASIC cells or gates. Through the PLI, the delay calculator program can modify the Verilog simulation data structure so that simulation is using the more accurate delays.

- Custom output displays

In order to verify design functionality, a logic simulator must generate outputs. The Verilog HDL provides formatted text output, and most simulators provide waveform displays and other graphical output displays. The PLI can be used to extend these output format capabilities. A video controller, for example, might generate video output for a CRT or other type of display panel. Using the PLI, an engineer can take the outputs from simulation of the video controller, and, while simulation is running, pass the data to a program which can display the data in the same form as the real display output. The PLI allows custom output programs to dynamically read simulation values while simulation is running.

- Co-simulation

Complex designs often require several types of logic simulation, such as a mix of analog and digital simulations, or a mix of Verilog and VHDL simulations. Each simulator type could be run independently on its own regions of the design, but it is far more effective to simulate all regions of the design at the same time, using multiple simulators together. The PLI can be used as a communication channel for Verilog simulators to transfer data to and from other types of simulators.

- Design debug utilities

The Verilog logic simulators on the market vary a great deal in what they provide for debugging design problems. There may be some type of information about a design that is needed, which cannot be accessed by the simulator's debugger. The Verilog PLI can access information deep inside a simulation data structure. It is often a very trivial task to write a short C program that uses the PLI to extract some specific information needed for debugging a design.

- Simulation analysis

Most Verilog logic simulators are intended to simply apply input stimulus to a model and show the resulting model outputs. The PLI can be used to greatly enhance simulation by analyzing what had to happen inside a design to generate the output values. The PLI can be used to generate toggle check reports (how many times each node in a design changed value), code coverage reports (what Verilog HDL statements were exercised by the input tests), power usage (how much energy the chip consumed), etc.

The preceding examples are just a few of the tasks that can be performed using the Verilog PLI. The PLI provides full access to anything that is happening in a Verilog simulation, and allows external programs to modify the simulation. This open access enables truly unlimited possibilities. If an engineer can conceive of an application that requires interaction with Verilog simulations, and can write a program to perform the task, the PLI can be used to interface that application to a Verilog simulator.

8.2 General steps to create a PLI application

A PLI application is a user-defined C language application which can be executed by a Verilog simulator. The PLI application can interact with the simulation by both reading and modifying the simulation logic and delay values.

The general steps to create a PLI application are:

1. Define a **system task** or **system function** name for the application.
2. Write a C language **callif routine** which will be executed by the simulator whenever simulation encounters the system task name or the system function name.
Optionally, additional C language routines can be written which will be executed by the simulator for special conditions, such as when the simulator compiler encounters the system task/function name.
3. **Register** the system task or system function name and the associated C language routines with the Verilog simulator. This registration tells the simulator about the new system task or system function name, and the name of the **callif routine** associated with that task or system function (along with any other routines).
4. **Compile** the C source files which contain the PLI application routines, and link the object files into the Verilog simulator.

8.3 User-defined system tasks and system functions

In the Verilog language, a **system task** or a **system function** is a command which is executed by a Verilog simulator. The name of a system task or a system function begins with a *dollar sign* (\$).

A **system task** is used like a programming statement that is executed from a Verilog *initial procedure* or *always procedure*. For example, to print a message at every positive edge of a clock, the **\$display** system task can be used.

```
always @(posedge clock)
$display("chip_out = %h", chip_out);
```

A **system function** is used like a programming function which returns a value. A function can be called anywhere that a logic value can be used. For example, to assign a random number to a vector at every positive edge of a clock, the **\$random** system function can be used.

```
always @(posedge clock)
vector <= $random();
```

The IEEE 1364 Verilog standard allows system tasks and system functions to be defined in three ways:

- A standard set of built-in system tasks and system functions.

These are defined as part of the IEEE standard, such as `$display`, `$random`, and `$finish`. All IEEE compliant Verilog simulators will have these standard system tasks and system functions.

- Simulator specific system tasks and system functions.

These are proprietary commands which are defined as part of a simulator, and may not exist in other simulators. Examples are `$save` to create a simulation check point file or `$db_settrace` to enable debug tracing.

- User-defined system tasks and system functions.

These are created through the Programming Language Interface. A PLI application developer specifies the name and functionality of the system task/function.

User-defined system task or system function names begin with a dollar sign (\$), just as the built-in system tasks and system functions. The user-defined system task/function name is then associated with a user-defined C application. When a Verilog simulator encounters the system task/function name, it will execute the C application associated with the name. This simple association allows users to extend the capability of a Verilog simulator in any manner desired. Chapter 9 discusses the concept of system tasks and system functions in more detail.

8.4 The `$hello` PLI application example

8.4.1 Step One: defining a `$hello` system task

The well known “Hello world” C program is a quick way to show how PLI applications are created. Though the C program itself is simple, it illustrates one of the most powerful features of the Verilog PLI—the ability for a Verilog designer to extend the Verilog language by having a Verilog simulation dynamically execute a user-defined program while simulation is running.

8.4.2 Step One: defining a `$hello` system task

The first step in creating a PLI application is create a new system task or system function name. A `$hello` user-defined system task will be created for this PLI application. An example of using `$hello` is:

```

module test;
initial
  $hello();
endmodule

```

8.4.3 Step Two: writing a calltf routine for \$hello

The second step in developing a PLI application is to write a C language routine which will be called when a Verilog simulator executes the *\$hello* system task. In this example, the simulator will call a C language routine which prints the message:

Hello World!

The C language routine which will be called by the simulator is referred to as a *calltf routine*. This routine is a user-defined C function and can be any name. Example 8-1 lists the C source code for the *calltf routine* for *\$hello*.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.08/hello_acc.c
- Verilog test bench: Chapter.08/hello_test.v
- Verilog-XL results log: Chapter.08/hello_test.log

Example 8-1: \$hello — a TF/ACC calltf routine for the PLI application

```

#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
#include "acc_user.h" /* IEEE 1364 PLI ACC routine library */

int PLIbook_hello_calltf()
{
  io_printf("\nHello World!\n\n");
  return(0);
}

```

In the above example:

- The header files *veriuser.h* and *acc_user.h* contain the TF and ACC libraries of functions provided by the PLI standard. These libraries are a collection of C functions which can be used by a PLI application to interact with a Verilog simulator. The *veriuser.h* and *acc_user.h* header files are part of the IEEE 1364 Verilog stan-

dard. Note that this example does not use any routines from the ACC library, and so the acc_user.h file does not need to be included in this example.

- The `io_printf()` function used in the preceding example is very similar to the C `printf()` function. The difference is that `io_printf()` will write its message to the simulator's output window, whereas the C `printf()` writes its message to the operating system's standard out channel.

8.4.4 Step Three: Registering the \$hello system task

The third step in creating a new PLI application is to tell the Verilog simulator about the new system task or system function name and the C routines which are associated with the application. This process is referred to as *registering* the PLI application.

The PLI IEEE 1364 standard provides an *interface mechanism* for this process. The interface mechanism defines:

- The type of application, which is a system task or system function.
- The system task or system function name.
- The name of the *callif routine* and other C routines associated with the system task/function.
- Other information about the system task/function required by a simulator.

The TF and ACC interface mechanism for PLI applications is *not* part of the IEEE standard. The standard only defines what the interface mechanism needs to do. The TF and ACC interface mechanism is implemented in different ways in different Verilog simulators. Chapter 9 presents more details on how most Verilog simulators implement the interface mechanism.

8.4.5 Step Four: Compiling and linking the \$hello system task

The final step in creating a new PLI application is to compile the C source code containing the application and linking the compiled files to the Verilog simulator. Once the application has been linked to the simulator, the simulator can invoke the *callif routine* when simulation executes the \$hello system task name. Compiling and linking is not part of the IEEE standard. This process is both specific to the simulator and the operating system on which simulations are run. Appendix A presents the steps for compiling and linking PLI applications for several major Verilog simulators.

8.4.6 Running simulations with the \$hello system task

Once a PLI application has been registered with a Verilog simulator, the new system task or system function can be used in a Verilog model, just as with built-in system tasks and system functions. When the Verilog simulator encounters the user-defined system task, it will call the C routine which has been associated with that application.

The following Verilog model can be used to test the *\$hello* PLI application:

```
module test;
initial
begin
    $hello();
    #10 $stop;
    $finish;
end
endmodule
```

Figure 8-1 shows the results of simulating this Verilog model using the Cadence Verilog-XL simulator with the Cadence SimVision graphical user interface.

Figure 8-1: Simulation results using the *\$hello* PLI application

The screenshot shows the Cadence SimVision interface with the title bar "Cadence Verilog-XL® Turbo NT". The menu bar includes File, Control, Set, Show, Select, Tools, Options, and Help. The toolbar contains icons for play, stop, step, zoom, and other simulation controls. The top status bar displays "Time: 10". The main window has two panes. The left pane is a code editor showing the Verilog source code:

```
08 ****
09 `timescale 1ns / 1ns
10 module test;
11
12     initial
13         begin
14             $hello();
15             #10 $stop;
16             $finish;
17         end
18     endmodule
```

The right pane is a terminal window displaying the simulation output:

```
Compiling source file "hello_test.v"
Highest level modules:
test

Hello World!
.
L15 "hello_test.v": $stop at simulation time 10
Type ? for help
C1 > |
```

The bottom status bar says "Ready".

8.5 The \$show_value PLI application example

The \$hello PLI application, though trivial from a C programming aspect, effectively illustrates the powerful capability provided by the Programming Language Interface—the ability to have a Verilog simulator call a user-defined C routine during simulation. The PLI provides far more capabilities, however. The C routine which is called by the simulator can use the TF and ACC libraries to read information from the simulation data structure, and to dynamically modify logic values and delay values in the data structure.

The system task **\$show_value** illustrates using the PLI to allow a C routine to read current logic values within a Verilog simulation. The **\$show_value** system task requires one argument, which is the name of a net or reg in the Verilog design. The *calltf routine* will then use the library of TF and ACC routines to read the current logic value of that net or reg, and print its name and logic value to the simulator's output screen. An example of using **\$show_value** is:

```
module test;
    reg a, b, ci;
    wire sum, co;
    ...
    initial
        begin
            ...
            $show_value(sum);
        end
    endmodule
```

Two user-defined C routines will be associated with **\$show_value**:

- A C routine to verify that **\$show_value** has the correct type of argument.
- A C routine to print the name and logic value of the signal.

These programs are presented in the following sections.

8.5.1 Writing a checktf routine for \$show_value

The PLI allows users to provide a C routine to verify that a system task or system function is being used correctly and has the correct types of arguments. This C routine is referred to as a *checktf routine*.

The TF and ACC routines are two libraries of C functions which can access the arguments of a system task or system function and determine the Verilog data types of the objects listed in the arguments. This example uses some of these C functions. Only a brief explanation of each function is provided in this chapter. Subsequent chapters will present more details about each function.

The C source code for the *checktf routine* associated with `$show_value` is listed below. The C functions and constants from the PLI library are shown in bold text.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.08/show_value_acc.c
- Verilog test bench: Chapter.08/show_value_test.v
- Verilog-XL results log: Chapter.08/show_value_test.log

Example 8-2: \$show_value — a TF/ACC *checktf routine* for a PLI application

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"           /* IEEE 1364 PLI ACC routine library */

int PLIBbook_ShowVal_Checktf()
{
    int      arg_type;
    handle  arg_handle;

    if (tf_nump() != 1)
        tf_error("$show_value must have 1 argument.");
    else if (tf_typep(1) == TF_NULPARAM)
        tf_error("$show_value arg cannot be null.");
    else {
        arg_handle = acc_handle_tfarg(1);
        arg_type = acc_fetch_type(arg_handle);
        if (!(arg_type == accNet || arg_type == accReg))
            tf_error("$show_value arg must be a net or reg.");
    }
    return(0);
}
```

Observe the following key points in the above example:

- The function `tf_nump()` returns the number of arguments in the system task. This example uses `tf_nump()` to verify `$show_value` has one argument.
- The function `tf_typep()` returns a constant representing the general data type of a task/function argument. `tf_typep()` is used in this example to verify that the argument to `$show_value` is not null.

- The function **acc_handle_tfarg()** returns a pointer to a system task argument. The pointer is stored in a special C data type defined by the PLI, called a **handle**.
- The function **acc_fetch_type()** returns a constant that more closely identifies the type of the argument. In this example, the test is checking that the argument is a Verilog net or reg data type.
- The function **tf_error()** prints an error message and causes simulation to abort.

8.5.2 Writing the calltf routine for \$show_value

The **calltf routine** is the C routine which will be executed when simulation encounters the **\$show_value** system task during simulation.

The following example lists the C source code for the *calltf routine* associated with **\$show_value**. Later chapters define the routines used in this example in full detail. The C functions, constants and data types from the PLI library are shown in bold text.

Example 8-3: \$show_value — a TF/ACC *calltf routine* for a PLI application

```
#include "veriuser.h"
#include "acc_user.h"
int ShowValCall()
{
    handle arg_handle;
    arg_handle = acc_handle_tfarg(1);
    io_printf("Signal %s has the value %s\n",
              acc_fetch_fullname(arg_handle),
              acc_fetch_value(arg_handle, "%b", null));
    return(0);
}
```

In the above example:

- The function **io_printf()** prints a formatted message. It is similar to the C **printf()** function, but prints the message to the simulation output window and the simulation output log file.
- The function **acc_fetch_fullname()** returns the full hierarchical name of a Verilog object.
- The function **acc_fetch_value()** returns the logic value of a Verilog net, reg or variable. The value can be obtained in a variety of formats. In this example, the value is obtained as a string, with a binary representation of the value.

8.5.3 Registering the \$show_value PLI application

The *\$show_value* PLI application needs to be registered with the Verilog simulator, which informs the Verilog simulator about the PLI application and the names of the *calltf routine* and *checktf routine* C functions. The IEEE standard defines that all Verilog simulators should provide an interface mechanism to perform this registration, but the standard does not define how the interface mechanism should be implemented. Chapter 9 shows how most Verilog simulators implement this interface mechanism.

8.5.4 A test case for \$show_value

The following Verilog HDL source code is a small test case for the *\$show_value* application. The example lists the values of a simple 1-bit adder after a few input test values have been applied. Observe how *\$show_value* is used as a Verilog programming statement within the test bench.

Example 8-4: *\$show_value* — a Verilog HDL test case using a PLI application

```
'timescale 1ns / 1ns
module test;
    reg a, b, ci, clk;
    wire sum, co;
    addbit i1 (a, b, ci, sum, co);
    initial
        begin
            clk = 0;
            a = 0;
            b = 0;
            ci = 0;
            #10 a = 1;
            #10 b = 1;

            $show_value(sum);
            $show_value(co);
            $show_value(i1.n3);

            #10 $stop;
            $finish;
        end
    endmodule


/** A gate level 1 bit adder model ***/
'timescale 1ns / 1ns
module addbit (a, b, ci, sum, co);
    input a, b, ci;
    output sum, co;
    wire a, b, ci, sum, co,
```

```
n1, n2, n3;  
xor (n1, a, b);  
xor #2 (sum, n1, ci);  
and (n2, a, b);  
and (n3, n1, ci);  
or #2 (co, n2, n3);  
endmodule
```

8.5.5 Output from running the \$show_value test case

Following is the output from simulating the test case for the `$show_value` application, using the Cadence Verilog-XL simulator with its SimVision user interface.

Figure 8-2: `$show_value` — sample simulation results, using Verilog-XL



8.6 Summary

This chapter has presented two short examples which illustrate how PLI applications are created using the TF and ACC routine libraries in the PLI standard. The *\$hello* PLI application showed how a Verilog simulation can execute another C program. The *\$show_value* application illustrated how arguments can be passed from a Verilog simulation to a PLI application, and how the application can access information about the arguments.

The principles illustrated by the *\$show_value* example can be readily expanded to provide much more powerful and useful capabilities. In Chapter 14, the functionality of *\$show_value* will be extended to create an application called *\$list_nets*, which automatically finds all nets in a module and prints the values of each net. Another example in the same chapter extends *\$show_value* to create *\$list_signals*, which automatically finds all nets, regs and variables in a module and prints the values of each one.

The next chapter presents much greater detail on how PLI application interact with Verilog simulations. Subsequent chapters in Part Two of this book will discuss the complete TF and ACC libraries, and include many examples of how the TF and ACC routines to create PLI applications which to interact with Verilog simulators.

CHAPTER 9

Interfacing TF/ACC based PLI Applications to Verilog Simulators

Don't skip this chapter! This chapter defines the terminology used by the Verilog PLI and how PLI applications which use the TF and ACC libraries are interfaced to Verilog simulators. All remaining chapters in Part Two of this book assume the principles covered in this chapter are understood. The general concepts presented are:

- PLI terms as used in this book
- Generations of the PLI standard
- System tasks and system functions
- How TF and ACC routines work
- A Complete PLI application example
- Interfacing PLI applications to Verilog simulators

9.1 General PLI terms as used in this book

The Verilog PLI has been in use since 1985, and, over the years, many reference manuals, technical papers and training courses have been written about the PLI. Unfortunately, these documents have not used consistent terms for describing the PLI. This inconsistency can make it difficult to cross reference information about the PLI in different documents.

This book strives to clearly define the PLI terms and consistently use this terminology. The general PLI terms used in this book are.

C program (often abbreviated to *program*)

A complete software program written in the C or C++ programming language. A C program must include a C *main* function.

C function

A function written in the C or C++ programming language. A C function does not include a C *main* function.

Verilog function

A Verilog HDL function, written in the Verilog Hardware Description Language. Verilog functions can only be called from Verilog source code.

User-defined system task or system function (abbreviated to *system task/function*)

A user-defined system task or system function is a construct that is used in Verilog HDL source code. The name of a user-defined system task or system function must begin with a dollar sign (\$), and it is used in the Verilog language in the same way as a Verilog HDL standard system task or system function. When simulation encounters the user-defined system task or system function, the simulator will execute a *PLI routine* associated with the system task/function.

PLI routine

A C function which is part of a *PLI application*. The PLI routine is executed by the simulator when simulation encounters a user-defined system task or system function. The TF and ACC portion of the PLI standard define several PLI routine types: *calltf routines*, *checktf routines*, *sizetf routines*, *misctf routines* and *consumer routines*. These terms are defined in more detail later in this chapter.

PLI application (sometimes abbreviated to *application*)

A *user-defined system task or system function* with an associated with a set of one or more *PLI routines*.

PLI library

A library of C functions which are defined in the Verilog PLI standard. PLI library functions are called from *PLI routines*, and enable the *PLI routines* to interact with a Verilog simulation. The IEEE 1364 standard provides three PLI libraries, referred to as the *VPI library*, the *TF library* and the *ACC library*.

VPI routines, TF routines and ACC routines

C functions contained in the *PLI library*. The term *routine* is used for these functions, to avoid confusion with *Verilog functions* and *C functions*.

9.2 System tasks and system functions

The Verilog Hardware Description Language provides constructs called “**system tasks**” and “**system functions**”. A system task/function is not a hardware modeling construct. It is a command that is executed by a Verilog simulator.

The names of system tasks and system functions always begin with a dollar sign (\$). A **system task** is analogous to a subroutine. When the task is called, the simulator branches to a program that executes the functionality of the task. When the task has completed, the simulation returns to the next statement following the task call. A **system function** returns a value. When the function is called, a simulator executes the program associated with the function, which returns a value into the simulation. The return value becomes part of the statement that called the function. The return value of a function may be defined to be an integer, a vector (with a specific bit size) or a floating point number.

9.2.1 Built-in system tasks and system functions

The IEEE 1364 Verilog standard defines a number of system tasks and system functions that are built into all Verilog simulators. Examples of the standard system tasks are \$display, \$stop, and \$finish. Some of the standard system functions are \$time and \$random. These system tasks and system functions are part of the Verilog language, and the syntax and usage of these routines are covered in Verilog language books.

The IEEE 1364 standard also allows simulation vendors to add proprietary system tasks and system functions which are specific to a simulator product. For example, the Cadence Verilog-XL simulator adds the command \$db_settrace, which is a system task that turns on Verilog statement execution tracing to aid in design debugging.

9.2.2 User-defined system tasks and system functions

The Verilog PLI provides a means for Verilog users to add additional system tasks and system functions. When simulation encounters a user-defined system task or system function, the simulator branches to a user-supplied C function, executes that function, and then returns back to executing the Verilog source at the point where the system task/function was called.

User-defined system task and system function names must begin with a dollar sign (\$), and may only use the characters that are legal in Verilog names, which are:

a—z A—Z 1—9 _ \$

Following are examples of legal user-defined system task and system function names:

\$rand64

\$cell_count

\$GetVector

9.2.3 Overloading built-in system tasks and system functions

A user-defined system task or system function can be given the same name as a system task/function that is built into the simulator. This allows a PLI application developer to overload the operation of a built-in system task/function with new functionality. For example, the built-in system function *\$random* will return a 32-bit signed random number, with the random sequence defined by the simulator. If the verification of a design requires a special random number sequence, a random number generator could be written in the C language and then associated with a user-defined system function with the same *\$random* name. The simulator would then call the user PLI application instead of the built-in random number generator.

9.2.4 System tasks are used as procedural programming statements

System tasks are procedural programming statements in the Verilog HDL. This means a system task may only be called from a Verilog **initial** procedure, an **always** procedure, a Verilog HDL **task** or a Verilog HDL **function**. A system task may not be called outside of a procedure, such as from a continuous assignment statement.

The following example calls a user-defined system task from an **always** procedure:

```
always @(posedge clock)
    $read_test_vector("vectors.pat", input_vector);
```

9.2.5 System functions are used as expressions

System functions are an expression in the Verilog HDL. This means a system function may be called anywhere a logic value may be used. The return value of the system function is considered the result of the expression. System functions may be called from a Verilog **initial** procedure, an **always** procedure, a Verilog HDL **task**, a Verilog HDL **function**, an **assign** continuous assignment statement, or as an operand in another expression.

The following examples call a user-defined system function several different ways:

```

always @(posedge clock)
  if (chip_out !== $pow(base,exponent))
    ...
initial
  $monitor("output = %f", $pow(base,exponent));
assign temp = i + $pow(base,exponent);

```

9.2.6 System task/function arguments

A system task or system function may have any number of arguments, including none. The user-supplied PLI application can read the values of the arguments of a system task/function. If the argument is a Verilog `reg` or variable (`integer`, `real` and `time`), then the PLI application can also write values into the task/function argument.

System task/function arguments are typically numbered from 1 to N, with the left-most argument being argument number 1. For example:

```

module top (...);
  ...
  reg [15:0] in1;
  ...
  my_chip u1 (in1, out1);
  ...
  initial
    $cell_count(u1);
  ...
  always @(posedge clock)
    $read_vector_file("vectors.pat", in1);
  ...
endmodule

```

argument 1 is a module instance name

argument 1 is a string

argument 2 is a reg variable

9.3 Instantiated Verilog designs

In a Verilog HDL module, a reference to another module is referred to as a **module instance**. In the following Verilog source code, module `top` contains two *instances* of module `bottom`.

```
module top;
  reg [7:0] in1;
  wire [7:0] out1;
  bottom b1 (in1[7:4], out1[7:4]);
  bottom b2 (in2[3:0], out2[3:0]);
endmodule

module bottom (in, out);
  ...
endmodule
```

The Verilog PLI requires an instantiated Verilog data structure. This means that a Verilog simulator has created a data structure which contains the complete hierarchy tree that has been represented by one module containing instances of other modules. There is no limit to the number of levels of hierarchy in the Verilog language.

The PLI does not work directly with Verilog HDL source code—The PLI is simply a procedural interface provided by a simulator so that PLI applications can access the data structure of a simulation.

Multiple instances of a system task or system function

The system task or system function which invokes a PLI application can be instantiated any number of times in the Verilog source code. Consider the following Verilog HDL source code fragment:

```
module top;
  ...
  middle m1 (...);
  middle m2 (...);

  initial
    $my_app_1(in1, out1);

  always @ (posedge clock)
    $my_app_1(in2, out2);

endmodule

module middle (...);
  ...
  bottom b1 (...);
  bottom b2 (...);

endmodule
```

```

module bottom (...);
  ...
  initial
    $my_app_2(in3, out3);
  always @(posedge clock)
    $my_app_2(in4, out4);
endmodule

```

In the preceding example, four different conditions are shown that can exist in a Verilog simulation.

- A single instance of a system task that is invoked one time

In module **top**, the first *instance* of **\$my_app_1** is in a Verilog **initial** procedure, which will be called once at the beginning of a simulation.

- A single instance of a system task that is invoked many times

In module **top**, the second *instance* of **\$my_app_1** is in a Verilog **always** procedure, which will be called every positive edge of the clock signal.

- Multiple instances of a system task that are each invoked one time

In module **bottom**, the first *instance* of **\$my_app_2** is in a Verilog **initial** procedure, which will be called once at the beginning of a simulation.

- Multiple instances of a system task that are each invoked many times


In module **bottom**, the second *instance* of **\$my_app_2** is in a Verilog **always** procedure, which will be called every positive edge of the clock signal.

Note that in the preceding Verilog source code example, module **bottom** is instantiated two times in module **middle**. The two instances of **bottom**, with two instances of **\$my_app_2** inside **bottom**, creates four unique instances of **\$my_app_2**. Module **middle** is instantiated twice in module **top**, resulting in eight unique instances of **\$my_app_2**.

Thus, in the instantiated simulation data structure for the preceding Verilog source code, there are **two instances** of the **\$my_app_1** system task, and **eight instances** of the **\$my_app_2** system task. Each of these system task instances will have unique logic values for the inputs and outputs. One of the requirements of a PLI application is to allow for multiple unique instances of the application. This is an important consideration when an application allocates static variables or allocates system memory. Chapter 10 discusses these issues in more depth, in section 10.7 on page 313.

9.4 How PLI applications work

The PLI standard allows a Verilog user to create a user-defined system task or system function name and to associate one or more user-defined C routines with the system task/function name. When a Verilog simulator encounters the user-defined system task/function, it will execute the user-defined C routine associated with the name.



9.4.1 The types of PLI routines

The TF and ACC portion of PLI standard defines several types of PLI routines which can be associated with a system task or system function. The type of the routine determines *when* the simulator will execute the routine. Some types of routines are run-time routines, which are invoked during simulation, and some types are compile time routines, which are invoked prior to simulation. The types of PLI routines are:

- *calltf routines*
- *checktf routines*
- *sizetf routines*
- *misctf routines*

The purpose of these routines is defined in the following sections of this chapter.

9.4.2 Associating routine types with system task/functions

A system task/function can have multiple PLI routines associated with it, as long as each routine is a different type. A *\$read_vector_file* application, for example, might have three routines associated with it—one which is called at compile time to perform syntax checking, one which is called when simulation first starts running to open the test vector file, and one which is called during simulation at every positive edge of clock to read vectors from the file.

Up to 4 different C routines may be associated with a system task/function in the TF and ACC portion of the PLI standard: the *calltf routine*, *checktf routine*, *sizetf routine*

and *misctf routine*. The ACC library also provides for another type of PLI routine, called a *Value Change Link consumer routine*. This routine is not directly associated with a system task/function name. Instead, the *consumer routine* is called for logic value changes on specific objects with a simulation data structure.

The following sections of this chapter define in detail the purpose and usage of the *calltf routine*, *checktf routine*, *sizetf routine* and *misctf routine*. The *consumer routine* is discussed in Chapter 17.

NOTE → PLI routines are sometimes referred to as PLI programs or C programs. In proper C terminology, PLI routines are actually C *functions*. A *program* in the C language terminology requires a *main* function. PLI applications do not contain a main function. Instead, the routines which make up the PLI application are linked into a Verilog simulator, and the simulator contains the C main function. This allows the simulator to call the appropriate PLI routine whenever its associated system task/function name is encountered.


9.4.3 C versus C++

Most Verilog simulators expect that PLI applications are written in ANSI C.

The PLI libraries are compliant with the ANSI C standard, but also include the appropriate prototypes and other information required for C++. This allows PLI applications to be developed in either C or C++. However, many Verilog simulators were written in the C language, and may or may not support linking C++ applications to the simulator. PLI application developers who wish to work with C++ should first check the limitations of the simulator product to which the applications will be linked. For maximum portability, PLI applications should be written using ANSI C.


9.5 *calltf routines*

The *calltf routine* is executed when simulation is running. For the \$pow example that follows, at every positive edge of clock the *calltf routine* associated with \$pow will be executed by the simulator. The routine is user-defined, and, for this example, will calculate a 32-bit number representing x to the power of y.



9.6 checktf routines

The **checktf routine** is called by the simulator before simulation starts running—in other words, before simulation time 0. The routine may be called by the simulator’s compiler, or when the simulator loads and prepares its simulation data structure. The purpose of **checktf routine** is to verify that a system task/function is being used correctly (e.g.: to check that the call to the PLI application has the correct number of arguments, and that the arguments are the correct Verilog data types). For example, if the \$pow system function were passed a module name as its second argument instead of a value, then the C math library would probably get an error when trying to use the name as an exponent value. Bad data can result in errors ranging from incorrect results to program crashes. Therefore, it is important to verify that each usage of a system task/function has valid arguments.



TIP Use a *checktf routine* to improve the performance of PLI programs. Since this routine is only called one time prior to simulation time 0, the code to check the correctness of arguments is only executed once. The *calltf routine*, which may be invoked millions of times during a simulation, does not need to repeat the syntax checking, and therefore can execute more efficiently.



NOTE The *checktf routine* will be called one time for each instance of a user-defined system task/function. If a design used the *\$pow* user-defined system function in three different places, the *checktf routine* for *\$pow* would be called three times. If a simulator allows system tasks and system functions to be invoked from the simulator's debug command line, then the *checktf routine* will be called prior to execution of the *calltf routine* for each interactive usage.

Limitations on checktf callbacks

The intent of the *checktf routine* is to verify the correctness of the arguments of an instance of a system task/function. The IEEE 1364 standard does not impose explicit restrictions on what TF and ACC routines will, or will not work in a *checktf routine*.

There are, however, considerations which should be observed. Since the *checktf routine* is called prior to simulation time 0, variables have not yet been assigned any values, and nets which are assigned values as soon as simulation starts may not yet reflect those values.

NOTE → *The checktf routine should only be used for syntax checking!* Do not use this routine to perform run-time duties such as allocating memory or opening files. The PLI standard does ensure that activity performed at compile time will remain in effect at simulation time. A *miscf routine* should be used instead of the *checktf routine* to perform run-time work at the very beginning of a simulation.

To ensure that a PLI application will be portable, only the following activities should be performed in a *checktf routine*:


- Accessing the arguments of an instance of a system task/function. Handles for the arguments can be obtained, and the properties of the arguments can be accessed to verify correctness. The values of literal numbers can be read, but variables and nets will not have been initialized. Note that Verilog parameter constants can be redefined for each instance of a Verilog module. The IEEE 1364 standard does not state whether the *checktf routine* will be called before or after parameter redefinitions have occurred. Reading parameter constant values from a *checktf routine* may not yield the same results on every simulator.
- Using TF and ACC routines such as `io_printf()`, which do not access objects in simulation. Attempting to traverse hierarchy to obtain handles beyond the system task arguments may not work correctly in all simulators. Attempting to write logic or delay values onto objects may result in PLI application errors.

9.7 *sizetf routines*

The *sizetf routine* is only used with system functions which return scalar or vector values (system functions can also return real number values). A system function can return a user-specified number of bits, and the simulator compiler may need to know how many bits to expect from the return, in order to correctly compile the statement from which the system function is called. A *sizetf routine* is called one time, before simulation time 0. The *sizetf routine* returns to the simulator how many bits wide the return value of system function will be.

NOTE → Even if a system function is used multiple times in the Verilog source code, the *sizetf routine* is only called once. The return value from the first call is used to determine the return size for all instances of the system function.

In the `$pow` example, the *calltf routine* will return a 32-bit value. Therefore, the *sizetf routine* associated with `$pow` needs to return a value of 32 to the simulator.



Limitations on sizetf callbacks

The intent of the *sizetf routine* is to notify the simulator compiler of the return size for system functions. The simulator may invoke the *sizetf routine* very early in the compilation phase of a Verilog design, and, at this early stage, the Verilog hierarchy may not have been generated. In addition, the *sizetf routine* is only called one time for a system function name, and the return size applied to all instances of the system function. For these reasons, only standard C language statements and functions should be used in a *sizetf routine*. An error may result if any TF or ACC routines are called from a *sizetf routine*. Any memory or static variables allocated by a *sizetf routine* may not remain in effect for when simulation starts running.

9.8 misctf routines

The *miscrf routine* is called for miscellaneous simulation events while the simulation is running. The IEEE 1364 standard defines the following miscellaneous events:


- End of Verilog source compilation (just before the start of simulation time 0).
- Entering debug mode (such as when the `$stop` built-in system task is executed).
- End of simulation (such as when the `$finish` built-in system task is executed).
- Change of value on a user-defined system task argument.
- End of a simulation time step—the PLI can schedule additional simulation activity in the time step.
- End of a simulation time step—the PLI cannot schedule additional simulation activity in the time step.
- Simulation has reached a specified simulation time step.

In addition to the standard reasons for which the *misctf routine* will be called, the PLI standard allows simulation vendors to add other reasons that are proprietary to features of a simulator. For example, the Cadence Verilog-XL simulator has the ability to save the simulation state to a file and restart simulation from the save file. Since it might be important for a PLI application to know that a save or restart has occurred, Cadence adds save and restart reasons to the Verilog-XL implementation of the PLI.

Each time a *misctf routine* is called, the simulator passes to the routine a C constant which represents the reason for the call. The *misctf routine* can then check the value of the reason constant and act accordingly. The names of the reason constants are defined in the IEEE 1364 standard and are presented in Chapter 12.

A common usage of the *misctf routine* is to perform tasks at the very beginning and the very end of a simulation. A PLI application to read test vectors might need to open a test vector file at the start of simulation, and close the file at the end of simulation.

```
always @(posedge clock)
$read_test_vector("vectors.pat", in1);
```



9.9 PLI routine inputs and outputs

PLI routines are C functions which are called by the Verilog simulator. When the function is called, the simulator will pass specific inputs to the function, and will receive a return value from the function.

The *sizetcf routine*, *checktf routine* and *calltf routine* are defined in the PLI standard to have two C inputs and to return an integer value. The *misctf routine* is defined to have three C inputs and to return an integer value.

PLI routine inputs

The inputs that are passed to the PLI routines are:

- An integer *user_data* value

The *user_data* value is defined along with the system task/function name (refer to section 9.12 on page 297 for more information on the *user_data* value).

- An integer *reason* value

This value is generated by the PLI and is represented by constants which are defined in the *veriuser.h* PLI header file. For *misctf routines*, the reason input is used to specify why the routine was called.

NOTE For the *sizetcf routine*, *checktf routine* and *calltf routine*, the reason input serves no purpose, and can be ignored.

- An integer *paramvc* value (*misctf routines* only)

This integer indicates which argument of a system task changed value. It is only used when a *misctf routine* is called with a reason of **REASON_PARAMVC**, which occurs when an argument of a system task changes value (see Chapter 12, section 12.5 on page 403 for more details on system task argument value change callbacks).

TF/ACC routine outputs

The *sizetcf routine*, *checktf routine*, *calltf routine* and *misctf routine* are all expected to be integer functions in the PLI standard. However, the only type of PLI routine that needs to return a value is the *sizetcf routine*, which returns the bit-width of the system function return value. The return value from a *checktf routine*, *calltf routine* and *misctf routine* are ignored by the simulator.

It is a common practice to declare the *checktf routine*, *calltf routine* and *misctf routine* as void functions, since their return value is ignored, even though all of these routines are expected to be integer functions. However, declaring the functions as a different type from the PLI standard prototype may result in a C compiler warning message when the routine is compiled. To prevent this warning message, the pointer to the void function can be cast to a pointer to an int function. For example:

```
(int(*)())PLIBook_PowCalltf
```

Note, however, that some C compilers might not accept this cast.

9.10 A complete system function example — \$pow

The following example illustrates all the parts of a complete system function, with the user-defined name of *\$pow*.

- The system function will return a 32-bit value.
- The system function requires two arguments, a base value and an exponent value. Both arguments must be numeric integer values.

To implement the *\$pow* functionality, four user-defined PLI routines are used:

- A *sizetf routine* to establish the return size of *\$pow*.
- A *checktf routine* to verify that the *\$pow* arguments are valid values.
- A *calltf routine* to calculate the base to the power of the exponent each time *\$pow* is executed by the simulator.
- A *misctf routine* to print a message when simulation first starts running (immediately prior to simulation time 0).

The *\$pow* system function does not really need the *misctf routine*, but to illustrate using this routine, this example uses a *misctf routine* to print a message at the start of simulation.



The *checktf routine*, *sizetf routine*, *calltf routine* and *misctf routine* are C functions.

These functions may be located in separate files, or they can all be in the same file.

TIP Typically, smaller PLI applications place all the routines in a single file, while larger, more complex applications might break them into multiple files.



TIP The *checktf routine*, *sizetf routine*, *calltf routine*, and *misctf routine* names can be any legal C function name. However, a typical Verilog HDL design may include several PLI applications. Using unique function name conventions can help avoid possible name conflicts when multiple PLI applications are linked together.

For this chapter, the *\$pow* application has been simplified to only work with integer numbers. In practice, this application would probably accept both integer and floating point numbers for inputs, and return a floating point number for the result. Example 11-1 on page 336 of Chapter 11 presents another version of *\$pow* that works with double-precision floating point numbers.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.09/pow_acc.c
- Verilog test bench: Chapter.09/pow_test.v
- Verilog-XL results log: Chapter.09/pow_test.log

Example 9-1: \$pow — a system function using TF/ACC routines

```
#include "veriususer.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"           /* IEEE 1364 PLI ACC routine library */
/*********************************************
 * Sizetf application
 *****/
int PLIBbook_pow_sizetf()
{
    return(32);      /* $pow returns 32-bit values */
}

/*********************************************
 * Checktf application
 *****/
int PLIBbook_pow_checktf()
{
    static int valid_types[4] = {accReg, accIntegerVar, accConstant, 0};
    handle arg_handle;

    if (tf_nump() != 2)
        tf_error("$pow must have 2 arguments.\n");
    else if (tf_typep(1) == tf_nullparam)
        tf_error("$pow arg 1 cannot be null.\n");
    else if (tf_typep(2) == tf_nullparam)
        tf_error("$pow arg 2 cannot be null.\n");
    else {
        arg_handle = acc_handle_tfarg(1);
        if (!(acc_object_in_typelist(arg_handle, valid_types)) )
            tf_error("$pow arg1 must be number, variable or net.\n");
    }
}
```

```

arg_handle = acc_handle_tfarg(2);
if (!(acc_object_in_typelist(arg_handle, valid_types)) )
    tf_error("$pow arg2 must be number, variable or net.\n");
}
return(0);
}

*****
* Calltf application
*****
#include <math.h>
int PLIbook_pow_calltf()
{
    int base, exp, result;

    base    = tf_getp(1);          /* read base value from tfarg 1 */
    exp     = tf_getp(2);          /* read exponent value from tfarg 2 */
    result = (int)pow( (double)base, (double)exp );
    tf_putp(0,result);           /* return result */
    return(0);
}

*****
* Misctf application
*****
int PLIbook_pow_misctf(int user_data, int reason)
{
    if (reason == reason_endofcompile)
        io_printf("\n$pow PLI application is being used.\n\n");
    return(0);
}

```

9.11 Interfacing PLI applications to Verilog simulators

As the previous sections in this chapter have shown, a PLI application may comprise:

- A system task or system function name
- A *calltf routine*
- A *checktf routine*
- A *sizetf routine*
- A *misctf routine*

After these items have been defined, the PLI application developer must carry out two basic steps:

1. Associate the system task/function name to the various application routines.
2. Link the applications into a Verilog simulator, so the simulator can call the appropriate routine when the system task/function name is encountered.

The PLI standard provides an *interface mechanism* to make the associations between the system task/function name and the application routines. There are two generations of the interface mechanism, one that was created for the older TF and ACC libraries, and a newer mechanism created for the VPI library. This chapter presents the TF and ACC interface mechanism. The newer VPI interface mechanism is presented in Part One of this book.

The TF and ACC interface mechanism is derived from the 1990 OVI PLI 1.0 standard. This older interface mechanism defines what all Verilog simulators should provide for interfacing PLI applications to a simulator, but the older interface does not define how the interface should be implemented. The IEEE 1364 standard includes this older interface mechanism to document the method for backward compatibility.

The TF/ACC interface mechanism is used to specify:

1. A system task/function *name*.
2. The application *type*, which is a *task*, *function* or *real function* (the IEEE 1364-1995 standard omitted the description of real functions from the body of the standard—this errata is corrected in the proposed IEEE 1364-1999 standard).
3. *Pointers* to the C functions for the *calltf routine*, *checktf routine*, *sizetf routine* and *misctf routine*, if the routines exist (it is not required—and often not necessary—to provide each type of routine).
4. An integer *user_data* value, which the simulator will pass to the *calltf routine*, *checktf routine*, *sizetf routine* and *misctf routine* each time they are called.

9.11.1 The defacto standard veriusertfs array

NOTE → The IEEE 1364 Verilog standard does not specify *how* a Verilog simulator should implement the TF and ACC interface mechanism—the standard only defines *what* the interface must specify. Every Verilog simulator provides a unique method for users to specify the interface items. However, most simulators use a special C array called a *veriusertfs* array, to specify the interface information. This section presents this defacto standard method. Appendix A presents the steps involved to compile and link PLI applications for several Verilog simulators.

The IEEE 1364 Verilog standard fully documents the VPI interface mechanism as the standard method for interfacing PLI applications. Prior to the IEEE standardization of

Verilog, there was no standard method for how the information about a PLI application was specified. Different simulator companies found different ways to specify the system task/function name and the associated routines. In order to preserve backward compatibility with existing implementations, the IEEE 1364 standards committee chose to not standardize any of the existing methods of specifying the TF/ACC PLI application information. Instead, the standard only specifies what information needs to be provided to interface PLI applications written with the TF and ACC libraries.

Many Verilog simulators have adopted a similar method of specifying the information about a PLI application. This method uses a C array called *veriusertfs* (which stands for *Verilog user tasks and functions*). The veriusertfs array is the original method implemented in the Cadence Verilog-XL product when the PLI was first introduced in the mid 1980's.

NOTE → The *veriusertfs* array is not specified in the IEEE 1364 standard. Though many simulators use this array, it is not required, and some simulators use different methods of specifying the PLI application information for the TF/ACC interface.

The typical *veriusertfs* array definition is:

```
s_tfcell veriusertfs[] =
{
    { type, user_data, checktf_app, sizetf_app,
      calltf_app, misctf_app, "tf_name", 1, 0, 0 },
    { type, user_data, checktf_app, sizetf_app,
      calltf_app, misctf_app, "tf_name", 1, 0, 0 },
    ...
} ; /* first field in final array cell is 0 */
```

Each array cell is an **s_tfcell** structure, which contains several fields. These fields specify the information about a PLI application. There can be any number of cells in the array. A cell with the first field set to 0 is used to denote the last cell in the array. The **s_tfcell** structure is not defined in the IEEE standard. However, nearly all simulators use the same structure definition, which is:

```

typedef struct t_tfcell {
    short type;           /* one of the constants: usertask,
                           userfunction, userrealfunction */
    short data;           /* data passed to user routine */
    int (*checktf)();     /* pointer to the checktf routine */
    int (*sizetf)();      /* pointer to the sizetf routine */
    int (*calltf)();      /* pointer to the calltf routine */
    int (*misctf)();      /* pointer to the misctf routine */
    char *tfname;         /* name of the system task/function */
    int forwref;          /* usually set to 1 */
    char *tfveritool;     /* usually ignored */
    char *tferrmessage;   /* usually ignored */
} s_tfcell, *p_tfcell;

```

The meaning of each field in the `s_tfcell` structure is explained in Table 9-1.

veriusertfs Field	Definition
type	Defines the type of a PLI application as being either a system task or a system function. This field must be set to the C constant usertask , userfunction or userrealfunction (these constants are defined in the TF library of the PLI standard).
user_data	Specifies an integer value—the value will be passed to the PLI application routines each time a routine is called.
checktf_app	Specifies a pointer to the C function that should be called by a simulator for the application's <i>checktf routine</i> .
sizetf_app	Specifies a pointer to the C function that should be called by a simulator for the application's <i>sizetf routine</i> .
calltf_app	Specifies a pointer to the C function that should be called by a simulator for the application's <i>calltf routine</i> .
misctf_app	Specifies a pointer to the C function that should be called by a simulator for the application's <i>misctf routine</i> .
"tf_name"	Specifies the name of the system task/function; must be a quoted literal string beginning with a \$.
forwref	Specifies instance name forward referencing as true (1) or false (0). This is not part of the IEEE standard and is ignored by most Verilog simulators. In the Cadence Verilog-XL simulator, setting to 1 makes Verilog-XL IEEE compliant by allowing module and primitive instance names to be used as system task/function arguments (a 0 makes instance names illegal in a system task/function argument).

Table 9-1: Typical fields in the defacto standard veriusertfs array

veriusertfs Field	Definition
tfveritool	This is not part of the IEEE standard and is ignored by most Verilog simulators.
tferrmessage	This is not part of the IEEE standard and is ignored by most Verilog simulators.

Table 9-1: Typical fields in the defacto standard veriusertfs array (continued)

The *veriusertfs* array can specify any number of PLI applications. A sample *veriusertfs* array is listed in example 9-2, with the entries for the *\$show_value* application presented in Chapter 8 on page 269 and the *\$pow* application presented in this chapter, in section 9.10 on page 290.



Do not specify the veriusertfs array in the same file as the PLI application!

- Not all Verilog simulators use the *veriusertfs* array to specify PLI application information.
 - The array is not standardized, and may be different in different simulators.
 - The C language does not allow multiple global arrays with the same name. If two applications both contained a *veriusertfs* array definition, the applications could not be used together.

Example 9-2: example veriusertfs array for registering TF/ACC PLI applications

```

/* prototypes for the PLI application routines */
extern int PLIbook_ShowVal_checktf(), PLIbook_ShowVal_calltf();
extern int PLIbook_pow_checktf(), PLIbook_pow_sizetf(),
          PLIbook_pow_calltf(), PLIbook_pow_misctf();

/* the veriusertfs array */
s_tfcell veriusertfs[] =
{
    {usertask,
     0,                                /* type of PLI routine */
     /* user_data value */
     PLIbook_ShowVal_checktf,           /* checktf routine */
     0,                                /* sizetf routine */
     PLIbook_ShowVal_calltf,           /* calltf routine */
     0,                                /* misctf routine */
     "$show_value",                   /* system task/function name */
     1                                 /* forward reference = true */
    },
}

```

```

{userfunction,
 0,                                /* type of PLI routine */
PLIbook_pow_checktf,                /* user_data value */
PLIbook_pow_sizetf,                 /* checktf routine */
PLIbook_pow_calltf,                 /* sizetf routine */
PLIbook_pow_misctf,                 /* calltf routine */
"$pow",                            /* misctf routine */
1,                                 /* system task/function name */
                                  /* forward reference = true */
},
{0} /*** final entry must be 0 ***/
};



---



```

9.12 Using the `user_data` field

When a system task or system function is registered, a `user_data` value can be specified. The `user_data` is an integer value. This value is passed to the *calltf routine*, *checktf routine*, *sizetf routine* and *misctf routine* as a C function input each time the simulator calls one of these routines.

In the following example, two different system tasks, `$read_test_vector_bin` and `$read_test_vector_hex`, are associated with the same *calltf routine*. Assuming the defacto standard *veriusertfs* table is used, the registration for these applications is:

```

s_tfcell veriusertfs[] =
{
  {usertask,
   1,                                /* type of PLI routine */
   0,                                /* user_data value */
   0,                                /* checktf routine */
   0,                                /* sizetf routine */
   PLIbook_ReadVector_calltf,          /* calltf routine */
   0,                                /* misctf routine */
   "$read_test_vector_bin",           /* system task/function name */
   1,                                 /* forward reference = true */
 },
 {usertask,
  2,                                /* type of PLI routine */
  0,                                /* user_data value */
  0,                                /* checktf routine */
  0,                                /* sizetf routine */
  PLIbook_ReadVector_calltf,          /* calltf routine */
  0,                                /* misctf routine */
  "$read_test_vector_hex",            /* system task/function name */
  1,                                 /* forward reference = true */
 },
};



---



```

Both system tasks invoke the same *calltf routine*, but the `user_data` value for the two system task names is different. Therefore, the *calltf routine* can check the `user_data` value to determine which system task name was used to call the routine.

The `user_data` value is passed to the PLI routine as a C function input of type `int`. The following example illustrates reading the `user_data` value for the two system tasks registered in the previous example.

```
int PLIBook_ReadVector_calltf(int user_data)
{
    if (user_data == 1)
        /* read test vectors as binary values */
    else if (user_data == 2)
        /* read test vectors as hex values */
}
```



NOTE The `user_data` value listed in the register routine will be common to all instances of the registered system task. If, for example, the system task `$read_test_vector_hex` were used in two places in the Verilog design, then there would be two instances of the system task. Each task will be passed the same `user_data` value. Chapter 10, section 10.7 on page 313 discusses how to allocate data that is unique to each instance of a system task/function.

9.13 Compiling and linking PLI applications

After the C source files for the PLI applications have been defined and the interface information specified, the C source code must be compiled and linked into a Verilog simulator. This allows the simulator to call the appropriate PLI routine when the PLI application system task or system function is encountered by the simulator. The compiling and linking process is not part of the IEEE standard for the PLI. This process is defined by the simulator vendor, and is specific to both the simulator and the operating system on which the simulator is being run. Appendix A presents the instructions for compiling and linking PLI applications for several of the Verilog simulators which were available at the time this book was written.

9.14 Summary

This chapter has presented the major parts of PLI applications which are developed using the TF and ACC libraries. The main parts of an application are the name of a system task or system function, a *checktf routine*, a *calltf routine*, a *sizetf routine*, and a *miscif routine*. A system function called \$pow, which calculates x^y and returns a 32-bit result, was used to illustrate the major parts of a PLI application. After defining the PLI application, the next step is to inform the simulator about the application. This is done through the *PLI interface mechanism*. The IEEE 1364 standard defines what the interface mechanism must specify, but the TF/ACC libraries do not define how the interface should be implemented. Most, but not all, Verilog simulators use a veriusertfs array to specify the interface information.

The remaining chapters in this part of the book examine the TF and ACC libraries in detail, and present a number of examples of using the routines in these libraries. Appendix A presents how PLI applications written using the TF and ACC libraries are linked into different Verilog simulators.

CHAPTER 10

How to Use the TF Routines

The **TF** routines are a library of 104 C functions that can interact with Verilog simulators. The discussion of TF routines is divided into four chapters: This chapter presents the TF routine library and how to use many of the TF routines. Chapter 11 presents how to read the values of the arguments of system tasks and system functions using TF routines. Chapter 12 presents how to use TF routines in conjunction with *miscif routines*, and Chapter 13 shows how to use the TF routines to interface C language models to Verilog simulators. Appendix B presents the complete syntax of the TF routine library.

The concepts presented in this chapter are:

- An overview of the library of TF routines
- Printing messages from PLI applications
- Obtaining information about system task/function arguments
- Using the TF work area
- Miscellaneous TF routines

10.1 The TF Library

The **TF** routines in the IEEE 1364 PLI standard are the oldest of three generations of the PLI standard (followed by the ACC routines and then the VPI routines). TF stands for **T**ask/**F**unction, because most of the TF routines deal with accessing information

about system task/function arguments. There are also a number of utility routines in the TF library for printing messages, controlling simulation, etc.

The library of TF routines is defined in a C header file called *veriuser.h*, which is part of the IEEE 1364 standard. In addition to the TF function definitions, the header file also defines a large number of C constants and C structures used by the TF routines. A complete list of the routines in the TF library, along with any structure definitions and constant names used with each TF function, is provided in Appendix B.

All PLI applications that reference TF routines or constants must include the *veriuser.h* file. For example:

```
#include "veriuser.h"      /* IEEE 1364 PLI TF routine library */
```

10.1.1 Using TF routines in conjunction with the ACC and VPI routines

The TF library can be used entirely by itself to create complete and useful PLI applications. The TF library and the ACC library are also designed to be used together. The ACC routines complement and greatly extend the capabilities of the TF library.

The VPI library was designed to replace both the TF and ACC libraries with a more concise, more robust, and more versatile procedural interface. It is *not* intended to have TF routines and VPI routines used in the same application. However, there is no conflict between the two libraries, which means multiple PLI applications which use the different libraries can be linked into the same Verilog simulation.

NOTE ➤ In the IEEE 1364-1995 PLI standard, there is one set of TF routines which have no counterpart in the VPI library, and so it is very common to include those TF routines in a VPI based PLI applications. Those routines are the ones which control a Verilog simulator, such as `tf_dofinish()`, `tf_dostop()`, `tf_error()` and `tf_writesave()`. The proposed 1364-1999 standard adds these capabilities to the VPI library, so that it will not be necessary to use the TF library at all in PLI applications developed with the VPI library.

The IEEE 1364 standard includes the TF and ACC libraries, in order to provide backward compatibility and portability of older PLI applications with modern Verilog simulators. The official policy of the IEEE 1364 standards committee is that as improvements and enhancements are added to the Verilog HDL, only the VPI library of the PLI will be expanded for those new features. There will be no new additions to the TF library.

10.1.2 Advantages of the TF library

The TF library has a major advantage over the ACC and VPI library, which offers a compelling reason to continue to develop and use PLI applications written solely with the TF library. The TF routines work directly with the arguments of system tasks and system functions. Unlike the VPI and ACC libraries, the TF routines cannot access the internal data structures of a simulation. This restricted access to the simulation data structure allows simulators to more efficiently optimize the simulation data structure, which can dramatically increase the run-time performance of a simulation. When a PLI application includes routines from the VPI and ACC libraries, this level of optimization may not be possible.


NOTE Not all simulators perform optimizations based on the type of PLI libraries which are used. Consult the documentation of specific simulators to determine if restricting PLI applications to only using the TF library can improve simulation performance.

10.2 System tasks and system functions

The Verilog PLI allows users to create user-defined system tasks and system functions which can be used in Verilog HDL source code. Refer back to Chapter 9 for more information creating user-defined system tasks and system functions.

10.2.1 System task/function arguments

A user-defined system task or system function can have any number of arguments, including none. The arguments are numbered from left to right, starting with argument number 1. In the following example:



The diagram shows two callout boxes. The first box, labeled "TF arg #1", points to the string argument "vectors.pat" in the code. The second box, labeled "TF arg #2", points to the signal argument "input_bus".

```
always @(posedge clock)
    $read_test_vector("vectors.pat", input_bus);
```

- Task/Function argument number 1 is a string, with the value “vectors.pat”.
- Task/Function argument number 2 is a signal, with the name `input_bus`.

Most of the TF routines in the PLI library are used for the purpose of accessing information about user-defined system task/function arguments.

NOTE

The IEEE 1364-1995 standard for the TF routines, and the older OVI PLI 1.0 standard from which the IEEE standard was derived, refer to task/function *arguments* as *parameters*. While *parameter* is an alternate programming term for a function argument, it is too easily confused with the Verilog HDL keyword *parameter*. To avoid this confusion, the ACC and VPI generations of the PLI refer to task/function arguments as *arguments*. In the proposed 1364-1999 standard, the word *parameter* will be changed to *argument* for TF routines, so that the standard uses consistent terminology for all routines. This book uses the term *argument* throughout.

10.2.2 Multiple instances of system tasks and system functions

The Verilog HDL source code can reference the same system task or system function any number of times. For example:

```
always @(posedge clock)
  $read_test_vector("A.dat", data_bus);

always @(negedge clock)
  $read_test_vector("B.dat", data_bus);
```

Just as a Verilog module can be used, or “*instantiated*”, many times in a design, every occurrence of a system task/function is a separate and unique *instance*. Each instance of *\$read_test_vector* in the above example has different arguments. The PLI recognizes that each instance is unique, and keeps track of each instance. Therefore, at each positive edge of clock, the *calltf routine* associated with one instance of *\$read_test_vector* will be invoked, and at the negative edge of clock the *calltf routine* associated with a different instance of *\$read_test_vector* will be executed. It is important to understand that the Verilog simulator will call the same C functions for each instance of the system task/function, but the inputs and data associated with each call will be unique for each instance.

As an example, *\$read_test_vector* might have a *calltf routine* named *gvCall()* associated with it. When the instance of *gvCall()* that is invoked at the positive edge of clock reads the value of task/function argument 1, it will see the string “*A.dat*”. When the instance of *gvCall()* that is invoked at the negative edge of clock reads task/function argument number 1, it will see the string “*B.dat*”.

10.2.3 Instance specific system task/function routines

A PLI *calltf routine*, *misctf routine*, *checktf routine* and *sizetf routine* are directly associated with the name of a system task or system function. (This association is part of

the PLI interface mechanism, which was presented in Chapter 9.) Each of these routines can access the arguments of the system task/function which caused the routine to be called. If these routines call another C function, that function can also access the arguments of the system task/function.

The ACC library also provides for another type of PLI routine, called a *consumer routine*. Defining and using *consumer routines* is defined in Chapter 17. A *consumer routine* is *not* associated with the name of a system task or system function. Therefore, a *consumer routine* cannot directly access the arguments of system tasks/functions. The TF library also provides a means for PLI applications to indirectly access system task/function arguments. Indirect access is done by obtaining a pointer for an instance of a system task/function and then using an instance specific version of TF routines to access the arguments of that system task/function instance.

An instance pointer for a system task or system function is obtained using the **tf_getinstance()** routine. The syntax of this routine is:

```
char *tf_getinstance()
```

The **tf_getinstance()** routine returns a pointer to the system task/function which called the PLI application.

There are instance-specific counterparts for many of the routines in the TF library. These instance-specific routines use the instance pointer to access the task/function arguments of a specific task/function instance. For example, the TF routine **tf_getp()** retrieves the value of an argument for the instance of the system task/function that called the PLI application. A counterpart to **tf_getp()** is **tf_igetp()**. This routine reads the value of an argument of a *different instance* of a system task/function, using the instance pointer of the other system task/function.

The typical usage of these routines is for a *calltf routine* or a *misctf routine* to obtain the pointer for the system task/function instance which called the routine. This handle is then saved in the *user_data* field of a *consumer routine*. This gives the consumer routine access to the instance handle, so that the consumer routine can access the arguments of the system task/function.

The instance specific versions of each TF routines are *not* presented in these chapters which discuss using the TF routines. The routines are listed in Appendix C, which presents the full syntax of the TF library. Functionally, the instance-specific versions of each routine work the same as the regular versions of the same routines. Example 13-7 on page 438 of Chapter 13 shows an example of one way in which the **tf_getinstance()** routine might be used. Chapter 18 also contains examples of using other instance specific routines.

NOTE

The TF system task/function instance pointer is not the same as the ACC system task/function instance pointer. The ACC library also has instance specific versions of the routines which access the arguments of a system task or system function. However, the instance pointer returned from `tf_getinstance()` is a different data type than the ACC handle returned from `acc_handle_tfinst()`. The instance pointer from a routine in one library should not be used with routines in the other library.

10.3 The PLI string buffer

A number of TF routines return pointers to C character strings, such as a file name or other name that is an argument to a system task/function. The string is stored in a temporary string buffer that is maintained by the PLI. This temporary buffer is limited in size, and will be overwritten by other calls to TF routines which return strings.

A PLI application should use the string pointer returned by a TF routine immediately. After another call is made to a TF routine which retrieves a string, there is no guarantee that the first string pointer will still be valid. If a string needs to be preserved, the PLI application should copy the string into application-allocated storage space. Following are two examples of using strings returned by TF routines. The TF routines used in these examples are described in more detail later in this chapter.

Read a string and use it immediately:

```
char *string_p;      /* string pointer only, no storage */
string_p = tf_strgetp(1, 'b');
io_printf("string_p points to %s\n", string_p);
```

Read a string and copy it to application-allocated storage for later use:

```
char *string_p;      /* string pointer only, no storage */
char *string_keep;   /* another string pointer only */
string_p = tfstrgetp(1, 'b'); /* get string from 1st arg */
string_keep = malloc(strlen(string_p)+1);
strcpy(string, string_p); /* save string for later use */
```

10.4 Controlling simulation

There are two TF routines which can control the run-time execution of a simulation.

void tf_dostop ()

Executes the same functionality as the Verilog *\$stop* built-in system task, which provides a means for a PLI application to send the simulation to the simulator's interactive debug environment. For example, a PLI application could be watching for a certain error condition to occur in a design (perhaps multiple bus drivers being enabled at the same time). When the error occurs, the PLI application could print a run-time warning message and then halt the simulation to allow an engineer to debug the problem using the simulator's debugger. When the user continues simulation from the interactive debug mode, the PLI application will resume execution at the statement which follows the call to *tf_dostop()*.

void tf_dofinish()

Executes the same functionality as the Verilog *\$finish* built-in system task, which provides a means for a PLI application to cause a simulation to exit, and in the process close any child processes or files that were opened by the simulator. An example of where *tf_dofinish()* can be used is with a PLI application to read test vectors from a file. When the file reader has reached the end of the test vector file, it can exit the simulation.

10.5 Printing messages

The TF routines provide several ways to print messages.

void io_printf (format, arg1,...arg12)

*char * format* quoted character string of formatted message.
arg1...arg12 arguments to formatted message string.

The *io_printf()* routine is used to print formatted text messages. The syntax for this routine is similar to the C *printf()* routine, but with a maximum of 12 arguments. The *io_printf()* routine prints the message to the simulation output window, and to the simulation output log file. By using *io_printf()*, a PLI application can generate messages which are part of the simulation, without having to determine where the simulator directs its output.



The C `printf()` routine will print its message to the operating system's standard output channel. If the simulator uses some other output channel, such as a graphical window, then printing a message using the C `printf()` routine might not be seen by the simulator user. Nor will the message be recorded in any output files generated by the simulator.

Following is an example of using `io_printf()`:

```
io_printf("Module %s has %d nets\n", module_name, num_nets);
```

void io_mcdprintf (mcd, format, arg1,...arg12)

int mcd multi-channel descriptor of open files.
char *format quoted character string of formatted message.
arg1...arg12 arguments to formatted message string.

The `io_mcdprintf()` routine is used to write messages into files that were opened by the Verilog HDL `$fopen` system function. The *multi-channel descriptor* (mcd) returned by `$fopen` must be passed to the PLI application as an argument to a user-defined system task/function or its value obtained using an ACC routine. The syntax of `tf_warning()` is similar to the C `printf()` routine, but the maximum number of arguments that can be printed in a message is 12.

void tf_warning (format, arg1,...arg5)

char *format quoted character string of formatted message.
arg1...arg5 arguments to formatted message string.

The `tf_warning()` routine prints a text message to the output window used by the simulator that is running the PLI application. The formatting of the message is not defined by the IEEE 1364 standard, and different simulators will format the message in different ways. The syntax of `tf_warning()` is similar to the C `printf()` routine, but the maximum number of arguments that can be printed in a message is 5. For example:

```
tf_warning("Reached end-of-file in test vector file %s\n",
           file_name);
```

void tf_error (format, arg1,...arg5)

char *format quoted character string of formatted message.
arg1...arg5 arguments to formatted message string.

The `tf_error()` routine prints a text message to the output window of the simulator that is running the PLI application. The formatting of the message is not defined by the IEEE 1364 standard, and different simulators will format the message in different ways. The syntax of `tf_error()` is similar to the C `printf()` routine, but the maximum number of arguments that can be printed in a message is 5. For example:

```
tf_error("Could not open test vector file %s\n", file_name);
```

In addition to printing a message, `tf_error()` causes the simulation to abort if it is called from a *checktf routine*. This important feature allows PLI application developers to write syntax checking routines which abort a simulation before a *calltf routine* tries to use incorrect information (e.g.: if a task/function argument is expected to have a module instance name, but a user has specified a file name instead).

NOTE Simulation will only abort if `tf_error()` is called from a *checktf routine*. `tf_error()` may also be used in a *calltf routine* or a *misctf routine* to print runtime error messages, but simulation will not be aborted.

TIP Some Verilog simulators do not adhere to the IEEE 1364 standard, in that they do not abort simulation when `tf_error()` is called from a *checktf routine*. For these simulators, the `tf_dofinish()` routine can be called after calling `tf_error()` to force an abort for a fatal error.

`void tf_message (level, facility, code, format, arg1,...arg5)`

<code>int</code>	<code>level</code>	a constant representing the error severity level. One of: <code>ERR_ERROR</code> , <code>ERR_SYSTEM</code> , <code>ERR_INTERNAL</code> , <code>ERR_WARNING</code> , <code>ERR_MESSAGE</code> .
<code>char</code>	<code>* facility</code>	quoted character string appended to the output message. Must be 10 or less characters.
<code>char</code>	<code>* code</code>	quoted character string appended to the output message. Must be 10 or less characters.
<code>char</code>	<code>* format</code>	quoted character string of formatted message.
	<code>arg1...arg5</code>	arguments to formatted message string.

The `tf_message()` routine is essentially a combination of `tf_warning()` and `tf_error()`. The routine prints a text message with a maximum of 5 arguments. A severity level is specified that can be an error, warning or informational level. The severity level is specified as one of the constants: `ERR_ERROR`, `ERR_SYSTEM`, `ERR_INTERNAL`, `ERR_WARNING`, or `ERR_MESSAGE`. If `tf_message()` is called from a *checktf routine*, the first three severity levels will cause simulation to abort (like `tf_error()`). The `tf_message()` routine also allows multiple messages to be queued and then printed together (see the description of `tf_text()`, which follows).

The `tf_message()` routine defines the error message format (whereas `tf_warning()` and `tf_error()` use the format defined by the simulator that is running the PLI application). The format is based on the Cadence Verilog-XL error/warning message format, which consists of a text message and two short text codes. The codes are limited to a maximum of 9 characters. An example usage of `tf_message()` is:

```
tf_message(ERR_ERROR, "User", "TFARG",
           "Arg %d is illegal in $read_test_vector", argnum);
```

The above example will print a message similar to:

```
ERROR: Arg 2 is illegal in $read_test_vector [User-TFARG]
```

void tf_text (format, arg1,...arg5)

`char * format` quoted character string of formatted message.
`arg1...arg5` arguments to formatted message string.

The `tf_text()` routine is used in conjunction with `tf_message()`. `tf_text()` allows multiple messages to be queued, which will all be printed when `tf_message()` is called. The following pseudo-code checks for three different types of errors, and uses `tf_text()` to queue any error messages. Then, if there were any errors, `tf_message()` is called to print the error messages and abort simulation.

```
bool err = FALSE;
if wrong number of arguments
    tf_text("$read_test_vector requires 2 args\n");
    err = TRUE;
if arg 1 is not a string
    tf_text("$read_test_vector arg 1 must be quoted file name\n");
    err = TRUE;
if arg 2 is not a reg data type
    tf_text("$read_test_vector arg 2 must be a reg data type\n");
    err = TRUE;
if (err)
    tf_message(ERR_ERROR, "User", "TFARG", "System task usage error");
```

10.6 Checking system task/function arguments

Three TF routines provide information about system task/function arguments that can be useful for verifying the correctness of arguments or performing other duties that affect the system task/function arguments.

```
int tf_numargs( )
```

The `tf_numargs()` routine returns the number of arguments of the user-defined system task/function that called the PLI application. `tf_numargs()` does not take any inputs. Two common ways to use `tf_numargs()` are:

- In a *checktf routine*, to verify that a system task/function instance has the correct number of arguments. For example:

```
if (tf_numargs() != 2 )
    tf_error("$read_test_vector has %d args, requires 2.\n",
             tf_numargs());
```

- In any PLI application, to loop through the task/function arguments for processing. For example:


```
numargs = tf_numargs();
for (i=1; i<= numargs; i++)
    ...
```



TIP For faster program execution, do not use `tf_numargs()` in a loop control test. In the above example, `tf_numargs()` is only called one time, and its return value is saved in a variable. If the loop had been coded as:

```
for (i=1; i<=tf_numargs(); i++)
then tf_numargs() would be called every pass of the loop. The additional calls to the TF routine are not necessary, yet require time to be executed.
```

In the Verilog HDL, it is legal for system tasks and system functions to have *null* arguments. The count returned by `tf_numargs()` routine includes null arguments. The `tf_typep()` routine can be used to determine if an argument is null. The following examples will return the argument counts shown:



int tf_typep(n)

int n index number of a system task/function argument.

tf_typep() returns a constant that represents the general data type of a specific system task/function argument. The input to **tf_typep()** is the index number of the system task/function argument of interest. The constants returned by **tf_typep()** are defined in the *veriuser.h* TF library. They are:

- **TF_NULLPARAM** — the argument is null.
- **TF_STRING** — the argument is a literal string.
- **TF_READONLY** — the argument is a literal integer number, an integer constant, or a Verilog net data type (wire, wand, wor, tri, triand, trior, trireg, supply0 or supply1).
- **TF_READONLYREAL** — the argument is a literal real number or a real constant.
- **TF_READWRITE** — the argument is a Verilog register variable data type (reg, integer, time).
- **TF_READWRITEREAL** — the argument is a Verilog real variable data type (real, realtime).

Following are two examples of using **tf_typep()**:

```
if (tf_typep(1) != TF_STRING)
    tf_error("$read_test_vector arg 1 must be a string\n");

if (tf_typep(2) != TF_READWRITE)
    tf_error("$read_test_vector arg 2 must be a reg type\n");
```

int tf_sizep(n)

int n index number of a system task/function argument.

The **tf_sizep()** routine returns the number of bits in a system task/function argument. The input to **tf_sizep()** is the index number of the argument. There are many ways the bit size of an argument can be useful. Following are two examples:

```
char *vector;
if (tf_sizep(1) != 1)
    tf_error("$my_app arg 1 must be scalar\n");
/* allocate memory for a string the size of TF arg 2 */
vector = malloc(tf_sizep(2) + 1);
```

10.7 The TF work area

The TF routines provide a *work area* pointer, where a PLI application can store a pointer to application data. Each instance of a system task/function is automatically allocated a work area by the PLI, and this work area is shared by the *checktf routine*, *misctf routine* and *calltf routine* associated with that system task/function.

The work area is defined as a `char * workarea` data type, which is typically a 32-bit storage space. A single value can be stored in the work area, or a pointer to a block of user-allocated memory can be stored.

```
void tf_setworkarea(workarea)
```

`char * workarea` pointer to a string or block of memory.

```
char *tf_getworkarea()
```

`char * tfinst` pointer to an instance of a system task/function.

The `tf_setworkarea()` routine is used to write a character pointer into the TF work area, and the `tf_getworkarea()` routine retrieves the pointer. There are two very powerful ways to use this work area, both of which greatly simplify writing PLI applications. The work area can be used to:

- Pass data from one PLI application routine to another application routine, such as from a *misctf routine* to a *calltf routine*.
- Store data that will be used in a future call to the same PLI application routine.

As an example of where both of the above capabilities might be needed, the `$read_test_vector` application might need to open a disk file from which to read vectors, and then read a line from the file each time the system task is called by simulation. The usage of the `$read_test_vector` example is:

```
always @(posedge clock)
$read_test_vector("A.dat", data_bus);
```

Since the test vector file should only be opened once, rather than at every positive edge of clock, a convenient place to open the file is in the *misctf routine* for `$read_test_vector`. The *misctf routine* is called for miscellaneous reasons, such as `REASON_ENDOFCOMPILE`, which indicates the start of actual simulation. The test vector file can be opened at this time, and the file pointer saved in the TF work area for `$read_test_vector`. More information about *misctf routines* is presented in Chapter 12.

The *calltf routine* for \$read_test_vector will perform the work of reading values from the file and applying them to simulation. The *calltf routine* will be invoked each time simulation encounters the system task, which, in the preceding example of \$read_test_vector, is at each positive edge of clock. Since the *calltf routine* did not open the disk file, it does not have direct access to the file pointer. But, the *calltf routine* shares the work area with the *misctf routine*. Therefore, the *misctf routine* can store the file pointer at the beginning of simulation, and, each time the *calltf routine* is invoked, it can retrieve the file pointer from the work area.

Example 10-1, which follows, contains a partial *misctf routine* and *calltf routine* for the \$read_test_vector example, which uses the TF work area to store a test vector file pointer. The complete code for \$read_test_vector is listed at the end of this chapter, as example 10-3 on page 326.

Example 10-1: using the tf_setworkarea() and tf_getworkarea() routines

```
*****
 * misctf routine
 ****
int PLIBook_ReadVector_misctf(int user_data, int reason)
{
    FILE *in_file;

    ...

    if (reason == REASON_ENDOFCOMPILE) { /* time to open vector file */
        in_file = fopen(tf_getcstringp(1), "r");

        tf_setworkarea((char*)in_file); /* save file pointer in workarea */

    ...

}

*****
 * calltf routine
 ****
int PLIBook_ReadVector_calltf() {
    FILE *in_file;

    ...

    in_file = (FILE*)tf_getworkarea(); /* retrieve file pointer */

    /* read next test vector from file */
    ...
}
```

10.7.1 Using the TF work area instead of a static or global variable

A TF work area is automatically allocated and maintained for each instance of a user-defined system task/function. This means that if there are multiple instances of a system task/function, then each instance has its own unique work area.

For example, consider the Verilog HDL code:

```
always @(posedge clock)
    $read_test_vector("A.dat", data_bus);

always @(negedge clock)
    $read_test_vector("B.dat", data_bus);
```

In this example, the `$read_test_vector` instance that is called at the positive edge of clock needs to save a file pointer to “*A.dat*”, and the `$read_test_vector` that is called at the negative edge of clock needs to save a file pointer to “*B.dat*”.

A common solution in the C programming language to preserve data is to declare a static variable instead of an automatic variable. Sharing common data between C functions can be done using global variables.

These C programming techniques can cause serious problems in PLI applications!

The Verilog HDL allows a system task/function to be used multiple times in the Verilog source code, and each module which uses the task/function can be instantiated multiple times. Each occurrence in each instance of a module becomes a unique *instance* of the system task/function. However, the same C function will be called by the simulator for each instance.

Since both instances of `$read_test_vector` in the preceding example will invoke the same C function, a static or global variable cannot be used, because the two instances would share the same single copy of the variable. The variable can only hold a single file pointer, and there is no simple way of knowing which instance of `$read_test_vector` was the last one to store a value in the variable. Using static or global variables is a sure way to have problems when there are multiple instances of a system task or system function.

The TF work area avoids the problems of static and global variables. The PLI automatically allocates a unique TF work area for each instance of a system task or system function. If a Verilog design includes two instances of the `$read_test_vector` application, each instance will have a unique work area in which to save its file pointer.

10.7.2 Storing multiple values in the TF work area

The work area only stores a single character pointer. This character pointer can be used to store multiple values, by allocating a block of memory in which to store the values and saving the pointer to the memory block in the work area.

The following code fragment, example 10-2, stores a file pointer, a character pointer, and two integers in the TF work area. In this example, a structure called `PLIbook_my_data` is defined to contain the three values to be stored. At the start of simulation, the `misctf` routine performs three operations:

- Allocate memory for one `PLIbook_my_data` structure.
- Fill the fields in the structure (if needed at this time).
- Store a pointer to the allocated memory in the work area.

Example 10-2: storing multiple values in the TF work area

```

typedef struct PLIbook_my_data {
    FILE *file_p;           /* pointer to a file */
    char *tfinst_p;         /* a character pointer */
    int temp1, temp2;       /* other integer data */
} PLIbook_my_data_s, *PLIbook_my_data_p;

void my_misc(int user_data, int reason)
{
    PLIbook_my_data_p data_p;
    switch(reason) {
        case REASON_ENDOFCOMPILE: /* misctf called just before time 0 */
            data_p = (PLIbook_my_data_p)malloc(sizeof(PLIbook_my_data_s));
            data_p->file_p = fopen(tf_getcstringp(1), "r");
            data_p->tfinst_p = tf_getinstance();
            tf_setworkarea((char *)data_p); /* store pointer to data */
            break;
        ...
    }
}


```

10.8 Reading and using simulation times

The Verilog HDL represents simulation time as a 64 bit unsigned integer. Since the ANSI C standard does not provide an integer data type that is guaranteed to be 64 bits on all operating systems, the Verilog PLI uses a pair of C unsigned integers to store the full 64 bits of simulation time. The lower 32 bits of Verilog time are placed in one

integer, and the upper 32 bits of the Verilog time will be stored in a second integer, as shown in the following illustration:

```
unsigned int high, low;
```



Within each Verilog module, time is scaled to a simulation time unit and time precision, using the `'timescale` compiler directive. The TF routines can retrieve the current scaled simulation time and information about a module's time scaling.

The `'timescale` directive in Verilog indicates what time units are used in the modules which follow the directive. The time scale also indicates a time precision, which is how many decimal points of accuracy are permitted in the following modules. Delays with more resolution than the precision are rounded off. Each module in a design can have a different time scale, so one model can specify delays in nanoseconds, with two decimal points of precision, and another model can specify delays in microseconds, with three decimal points. A Verilog simulator will scale the times in each module to the simulator's internal time units.

The `'timescale` directive is part of the Verilog language standard, and is not explained in this book. Refer to any good book on Verilog or the IEEE 1364 Verilog Language Reference Manual for more information on time scaling in Verilog. Page 7 lists a few sources for more information on the Verilog language.

10.8.1 Retrieving the current simulation time

```
int tf_gettime()
```

The `tf_gettime()` routine retrieves the lower 32-bits of current simulation time as a C integer.

int tf_getlongtime(&hightime)

*int * hightime* pointer to the upper 32-bits of the time.

This routine retrieves the current 64-bit simulation time. The lower 32-bits are returned, and the upper 32-bits are loaded into **hightime**.

double tf_getrealtime()

This routine retrieves the current simulation time as a real number.

The **tf_gettime()**, **tf_getlongtime()**, **tf_getrealtime()** and **tf_strgettime()** routines return the current simulation time. For each of these routines, the time is automatically scaled from simulation time units to the time units and precision of the module which called the PLI application. Since **tf_gettime()** and **tf_getlongtime()** return the time as an integer, any precision will be lost.

10.8.2 Retrieving a future simulation time

int tf_getnextlongtime (&lowtime, &hightime)

*int * lowtime* pointer to the lower 32-bits of the time.
*int * hightime* pointer to the upper 32-bits of the time.

The **tf_getnextlongtime()** routine returns the next simulation time in which a simulation event is scheduled, and returns a flag indicating the significance of the time retrieved. The full 64-bits of simulation time is retrieved, with the lower 32-bits loaded into **lowtime**, and the upper 32-bits loaded into **hightime**. The simulation time is returned in simulation time units—it is not scaled to the time scale of the module which called the PLI application.

In order for **tf_getnextlongtime()** to determine when the next event is scheduled, it must be called at the very end of the current simulation time step. The *misctf routine* can be scheduled to be called at the end of a time step, using the routine **tf_rosynchronize()** (refer to section 12.3 on page 395 for a description of using this routine). If **tf_getnextlongtime()** is not called at the very end of a time step, it retrieves the current simulation time, instead of the next simulation time. **tf_getnextlongtime()** returns a flag indicating the meaning of the simulation time retrieved, as follows:

- 0 **tf_getnextlongtime()** was called from a *misctf routine* at the end of a simulation time step (REASON_ROSYNCH). The time retrieved is the time the next simulation event is scheduled to occur.

- 1 `tf_getnextlongtime()` was called when there are no more simulation events scheduled (simulation is over). A 0 is returned for the simulation time.
- 2 `tf_getnextlongtime()` was *not* called from a *misctf routine* at the end of a simulation time step. The time retrieved is the current simulation time.

10.8.3 Retrieving time scale and precision values

`int tf_gettimeunit()`

This routine returns an integer representing the time scale unit of the module containing the instance of the system task/function which called the PLI application.

`int tf_igettimeunit(tfinst)`

`char * tfinst` pointer to an instance of a system task/function, or `null`.

This routine returns an integer representing the time scale unit of the module containing a specific instance of a system task/function. Use the `tfGetInstance()` routine to obtain the instance pointer.

`int tf_gettimeprecision()`

This routine returns an integer representing the time scale precision of the module containing the instance of the system task/function which called the PLI application.

`int tf_igettimeprecision(tfinst)`

`char * tfinst` pointer to an instance of a system task/function, or `null`.

This routine returns an integer representing the time scale precision of the module containing a specific instance of a system task/function. Use the `tfGetInstance()` routine to obtain the instance pointer.

If the instance specific versions are called with a `null` for the instance pointer, then they return the simulator's internal time units. For most Verilog simulators, the internal time unit is the finest precision of all modules in the Verilog simulation.

All of these routines return an integer which represents a unit of time, as shown in Table 10-1, which follows. The time values are represented as the order of magnitude of 1 second, which is the exponent of 1 second times 10^n .

Exponent Value	Time Unit Represented
2	100 seconds (1×10^2)
1	10 seconds (1×10^1)
0	1 second (1×10^0)
-1	100 milliseconds (1×10^{-1})
-2	10 milliseconds (1×10^{-2})
-3	1 millisecond (1×10^{-3})
-4	100 microseconds (1×10^{-4})
-5	10 microseconds (1×10^{-5})
-6	1 nanosecond (1×10^{-6})
-7	10 nanoseconds (1×10^{-7})
-8	1 nanosecond (1×10^{-8})
-9	1 picosecond (1×10^{-9})
-10	100 picoseconds (1×10^{-10})
-11	10 picoseconds (1×10^{-11})
-12	1 picosecond (1×10^{-12})
-13	100 femtoseconds (1×10^{-13})
-14	10 femtoseconds (1×10^{-14})
-15	1 femtosecond (1×10^{-15})

Table 10-1: Time unit values for tf_gettimeunit() and tf_gettimeprecision()

10.8.4 Utility routines to work with 64-bit time values

The TF library provides special routines to aid in manipulating 64-bit time values.

char *tf_longtime_tostr (lowtime, hightime)

int **lowtime** lower (right-most) 32-bits of a 64-bit unsigned integer.
int **hightime** upper (left-most) 32-bits of a 64-bit unsigned integer.

This routine converts a 64-bit simulation time that is represented as a pair of 32-bit integers to a character string. The routine returns a pointer to the string, which is stored in temporary simulation-allocated storage.

```
void tf_scale_longdelay (tfinst, low1, high1, low2, high2)
char   * tfinst      pointer to an instance of a system task/function.
int    low1         lower 32 bits of first operand.
int    high1        upper 32 bits of first operand.
int    * low2        pointer to lower 32 bits of second operand.
int    * high2       pointer to upper 32 bits of second operand.
```

This routine scales a 64-bit time value that is stored as a pair of 32-bit C integers to the time scale of a specific instance of a system task/function. The value stored in arguments `low1` and `high1` are scaled and deposited to `low2` and `high2`.

```
void   tf_scale_realdelay (tfinst, real1, real2)
char   * tfinst      pointer to an instance of a system task/function.
double real1        first operand.
double * real2       pointer to second operand.
```

Scales a double precision real number time value to the time scale of a specific instance of a system task/function. The value stored in `real1` is scaled and deposited to `real2`.

```
void tf_unscale_longdelay (tfinst, low1, high1, low2, high2)
char   * tfinst      pointer to an instance of a system task/function.
int    low1         lower 32 bits of first operand.
int    high1        upper 32 bits of first operand.
int    * low2        pointer to lower 32 bits of second operand.
int    * high2       pointer to upper 32 bits of second operand.
```

This routine converts a 64-bit time value expressed in simulation time units to the time scale of a specific instance of a system task/function. The time value stored in arguments `low1` and `high1` are converted and deposited into `low2` and `high2`.

```
void tf_unscale_realdelay (tfinst, real1, real2)
char   * tfinst      pointer to an instance of a system task/function.
double real1        real number value of first operand.
double * real2       pointer to real number value of second operand.
int    * high2       pointer to upper 32 bits of second operand.
```

Converts a real number value expressed in simulation time units to the time scale of a specific instance of a system task/function. The value stored in `real1` is converted and deposited to `real2`.

10.9 Reading simulation invocation options

The PLI can access the invocation options used to invoke a Verilog simulator. Using the TF routines, a PLI application can access invocation options which begin with a plus sign (+). The ACC and VPI libraries can access any invocation option.

NOTE The PLI can only access the invocation options used to invoke simulation. Some Verilog simulators use a separate command to compile Verilog HDL source code. When a simulator uses a separate compile command, the invocation options used to compile Verilog models cannot be read by the PLI.

char *mc_scan_plusargs (plusarg)
char * plusarg name of the invocation option.

The mc_scan_plusargs() routine scans the simulators invocations options for a specified user-defined invocation plus option. If the invocation option does not exist, the routine returns 0. If the option exists exactly as specified, the routine returns a null string. If the invocation option exists, and has additional characters beyond the specified option, the additional characters are returned in a string.

The following example performs a true/false test to see if simulation was invoked with a +PLIbook_verbose option:

```
if (mc_scan_plusargs("PLIbook_verbose"))
    io_printf("debug: value read from file is %s\n", vector);
```

Notice that the argument passed to mc_scan_plusargs() is a string containing the name of the invocation option without the plus sign.

The next example checks for the invocation option +file+<file_name> and prints the name of the file specified. If the invocation option is +file+vector1.pat, then the string "vector1.pat" will be printed.

```
if (file_name = mc_scan_plusargs("file+"))
    io_printf("Simulation invoked with +file+ name of %s\n",
              file_name);
else
    io_printf("Simulation was not invoked with +file+ \n");
```

NOTE

Care must be taken when using `mc_scan_plusargs()` as a simple true/false test, without checking the return string. The following test,

```
if (mc_scan_plusargs("test1")
    /* do something */
```

will return a true (a non-null return) for both the invocation option of `+test1` and `+test10`. To determine which option was used, it is necessary to examine the return value from `mc_scan_plusargs()`. The former will return a pointer to a null string, and the latter will return a pointer to a string of "0".

10.10 Utility TF routines

There are several utility TF routines to aid in working with Verilog models and simulation.

Performing math operations on 64-bit values stored as a pair of integers

The routines `tf_add_long()`, `tf_subtract_long()`, `tf_multiply_long()`, `tf_divide_long()`, `tf_compare_long()` provide a means of performing math operations on 64-bit values, which are represented as a pair of C integers.

`void tf_add_long (low1, high1, low2, high2)`

<code>int * low1</code>	pointer to lower 32 bits of first operand.
<code>int * high1</code>	pointer to upper 32 bits of first operand.
<code>int low2</code>	lower 32 bits of second operand
<code>int high2</code>	upper 32 bits of second operand

The `tf_add_long()` routine adds two 64-bit values, and deposits the result back into the first operand.

`void tf_subtract_long (low1, high1, low2, high2)`

<code>int * low1</code>	pointer to lower 32 bits of first operand.
<code>int * high1</code>	pointer to upper 32 bits of first operand.
<code>int low2</code>	lower 32 bits of second operand
<code>int high2</code>	upper 32 bits of second operand

The `tf_subtract_long()` routine subtracts two 64-bit values, and deposits the result back into the first operand.

void tf_multiply_long (low1, high1, low2, high2)

<i>int</i>	<i>* low1</i>	pointer to lower 32 bits of first operand.
<i>int</i>	<i>* high1</i>	pointer to upper 32 bits of first operand.
<i>int</i>	low2	lower 32 bits of second operand
<i>int</i>	high2	upper 32 bits of second operand

The `tf_multiply_long()` routine multiplies two 64-bit values, and deposits the result back into the first operand.

void tf_divide_long (low1, high1, low2, high2)

<i>int</i>	<i>* low1</i>	pointer to lower 32 bits of first operand.
<i>int</i>	<i>* high1</i>	pointer to upper 32 bits of first operand.
<i>int</i>	low2	lower 32 bits of second operand
<i>int</i>	high2	upper 32 bits of second operand

The `tf_divide_long()` routine divides two 64-bit values, and deposits the result into the first operand.

Converting 64-bit values stored as a pair of integers to/from doubles

Two TF routines are provided to aid in converting a 64-bit value that is stored in a pair of C integers into a C double, and vice-versa.

void tf_long_to_real (low, high, real)

<i>int</i>	low	lower (right-most) 32-bits of a 64-bit integer.
<i>int</i>	high	upper (left-most) 32-bits of a 64-bit integer.
<i>double</i>	<i>* real</i>	pointer to a double precision variable.

This routine converts a 64-bit value stored as a pair of C integers to a C double precision value.

void tf_real_to_long (real, low, high)

<i>double</i>	real	a double precision variable.
<i>int</i>	<i>* low</i>	pointer to a variable to receive the lower (right-most) 32-bits of a 64-bit integer.
<i>int</i>	<i>* high</i>	pointer to a variable to receive the upper (left-most) 32-bits of a 64-bit integer.

The `tf_real_to_long()` routine converts a C double precision value to a 64-bit value stored in a pair of C integers.

Getting the name of a module or scope containing a system task/function

`char *tf_mipname()`

The `tf_mipname()` routine returns a string containing the full hierarchical path name of the *module instance* containing the system task/function which called the PLI application. If the system task/function call is located in a scope level within a module, such as a Verilog HDL function, the module instance containing the scope is still returned.

`char *tf_spname()`

The `tf_spname()` routine returns a string containing the full hierarchical path name of the Verilog *scope* containing the system task/function which called the PLI application. A scope may be a module instance, a named statement group, a Verilog task or a Verilog function.

10.11 A complete PLI application using TF Routines

This section contains the complete source code for the `$read_test_vector` PLI applications, which uses a *checktf routine*, a *calltf routine* and a *misctf routine*.

- The *checktf routine* verifies that the first system task/function argument is a string and that the second system task/function argument is writable.
- The *misctf routine* opens the test vector file specified in first system task/function argument.
- The *calltf routine* writes a 4-state value, represented as a string, onto the second system task/function argument.

The usage of `$read_test_vector` is as follows:

```
reg [19:0] input_vector;  
always @(posedge clock)  
  $read_test_vector("vector_file", input_vector);
```

NOTE

Most of the TF routines used in this example have been discussed in this chapter, but two routines are presented in the next chapter. The `tf_getcstringp()` routine is used to read the value of a string in Verilog and convert it to a C string, and the `tf_strdelputp()` routine is used to convert a C string into Verilog logic values.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.010/read_test_vector_tf.c
- Verilog test bench: Chapter.010/read_test_vector_test.v
- Verilog-XL results log: Chapter.010/read_test_vector_test.log

Example 10-3: \$read_test_vector — using the TF routines

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include <stdlib.h>           /* ANSI C standard library */
#include <stdio.h>            /* ANSI C standard input/output library */
/*********************************************************************
 * checktf routine
 *****/
int PLIBbook_ReadVector_checktf()
{
    if (tf_nump() != 2) /* check for two system task/function */
        tf_error("Usage: $read_test_vector(\"<file>\",<reg_variable>);");
    if (tf_typep(1) != TF_STRING) /* check that first arg is a string */
        tf_error("$read_test_vector arg 1 must be a quoted file name");
    if (tf_typep(2) != TF_READWRITE) /* check that 2nd arg is reg type */
        tf_error("$read_test_vector arg 2 must be a register data type");
    return(0);
}

/*********************************************************************
 * misctf routine
 *   Use the misctf routine to open test vector file at the beginning
 *   of simulation, and save the file pointer in the work area for the
 *   instance of $read_test_vector that called the misctf routine.
 *****/
int PLIBbook_ReadVector_misctf(int user_data, int reason)
{
    FILE *in_file;
    char *file_name;

    if (reason == REASON_ENDOFCOMPILE) { /* time to open vector file */
        if ( (in_file = fopen(tf_getcstrngp(1), "r")) == NULL)
            tf_error("$read_test_vector cannot open file %s",
                     tf_getcstrngp(1));
        tf_setworkarea((char*)in_file); /* save file pointer in workarea */
    }
    return(0);
}

/*********************************************************************
 * calltf routine
 *****/
int PLIBbook_ReadVector_calltf() {
    FILE *in_file;
```

```
int    vec_size = tf_sizep(2);           /* bit size of tfarg 2 */
char *vector = malloc(vec_size+1);     /* memory for vector string */
bool  VERBOSE = FALSE;                 /* flag for debug output */

if (mc_scan_plusargs("debug")) VERBOSE = TRUE; /* set verbose flag */

in_file = (FILE*)tf_getworkarea();      /* retrieve file pointer */

if ( (fscanf(in_file,"%s\n", vector)) == EOF) { /* read a vector */
    tf_warning("$read_test_vector reached End-Of-File %s",
               tf_getcstringp(1)); /* get file name from task arg 1 */
    fclose(in_file);
    tf_dofinish(); /* exit simulation at end-of-file */
    return(0);
}
if (VERBOSE)
    io_printf("$read_test_vector: Value read from file=%s\n", vector);

/* write test vector value onto system task arg 2 */
if (!(tf_strdelputp(2,vec_size,'b',vector,0,0)))
    if (VERBOSE)
        tf_error("$read_test_vector could not write to arg 2 at time %s",
                  tf_strgettime() );
return(0);
}
```

10.12 Summary

This chapter has presented the routines contained in the PLI TF library. A complete PLI application, called **\$read_test_vector**, was presented, which used many of these routines. This application reads a test vector file from a file, and passes the value to the Verilog simulator.

The next two chapters are a continuation of the discussion of the TF library. They present how to use the TF routines to read and modify the values of system task/function arguments and how to use the TF library with *misctf routines*. The **\$read_stimulus** example presented in the Chapter 12 is a more elaborate test vector reader, which can read values in any radix, and can read both the vector value and the simulation time to apply the vector from the test vector file.

CHAPTER 11

Reading and Writing Values Using TF Routines

This chapter presents the routines in the TF library that read and modify the values of system task and system function arguments. The TF library provides a number of powerful functions that automatically convert Verilog data types and logic values to C data types and back to Verilog data types and values.

The concepts presented in this chapter are:

- Working with 4-state logic values and 8-level signal strengths
- How the PLI writes values into Verilog simulations
- Returning values of system functions
- Reading and writing Verilog 2-state values
- Reading and writing Verilog 4-state values using C strings
- Reading and writing Verilog 2-state values encoded C integer values
- Reading Verilog logic strengths
- Reading and writing to Verilog memory arrays

11.1 Working with a 4-logic value, multiple strength level system

The Verilog HDL supports 4 logic values, **0**, **1**, **z** and **x** and multiple levels of signal strength, but the C programming language does not directly represent the same information. The TF library offers users several ways to handle the difference in data representation between Verilog and C:

- Convert Verilog logic values from 4-state to 2-state (0 and 1); z and x values are converted to 0, and strength levels are ignored.
- Convert Verilog logic values from 4-state values to C character strings ("0", "1", "Z" and "X"), and ignore logic strength levels.
- Convert Verilog logic values from 4-state values to 2-bit encoded C values, referred to as an *aval/bval* pair, and ignore logic strengths.
- Convert Verilog logic values from 4-state values to 2-bit encoded C values, and convert logic strengths to two C integer values, which represent the strength encoding defined in the Verilog HDL standard.

The TF routines which read and write Verilog logic values automatically convert Verilog logic values to and from the various C representations, making it easy to represent hardware logic in the C language.

11.2 How the PLI writes values into Verilog

The TF routines which write values into task/function arguments and return values of system functions are referred to as “*put*” routines.

11.2.1 Putting values into a system task/function argument

To “*put*” a value into a task/function argument requires passing the index number of the system task/function argument as one of the inputs to the TF routine.

System task/function arguments are numbered from 1 to N, with the left-most argument being argument being number 1. For example:

```
$read_vector_file("vectors.pat", in1);
```

The code above is annotated with two boxes and arrows:

- A box labeled "argument 1 is a string" has an arrow pointing to the string literal "vectors.pat".
- A box labeled "argument 2 is a reg variable" has an arrow pointing to the variable "in1".

In the preceding example, the following TF routine will write *c_value* into the second task/function argument, which is *input_vector*:

```
tf_putchar(2, c_value);
```

The code above is annotated with a box and arrow:

- A box labeled "put 32-bit value on TF arg #2" has an arrow pointing to the argument "2".

In order for the PLI to put a value into a system task/function argument, the argument must be *writable*. That is, the argument must be a Verilog HDL register or variable data type (reg, integer, time, real or realtime). The `tf_typep()` routine can be used to verify that a system task/function argument is writable. For example:

```
if (tf_typep(2) != TF_READWRITE)
    tf_error("$read_test_vector arg 2 must be a variable\n");
```

The `tf_putstr()` and other TF put routines are discussed in more detail in subsequent sections of this chapter.

11.2.2 Returning system function values

To have the PLI return the value of a system function, the index number for the system task/function argument is set to **0**. In order to return a value using index number 0, the PLI application must have been declared as a type of *function* or *real-function* (refer back to Chapter 9, section 9.11 on page 292).

The TF library provides routines which return:

- A 2-state value of up to 32-bits wide, using the `tf_putstr()` routine.
- A 2-state value of up to 64-bits wide, using the `tf_putstrlongp()` routine.
- A double precision floating point value, using the `tf_putstrrealp()` routine.

Using the ACC library or the VPI library, PLI applications can return system function values for vectors of any bit size, as well as signed integers.

As an example of returning a system function value, assume the Verilog HDL code contains the statement:

```
reg [31:0] result;
always @(posedge clock)
    result = $pow(a, b);
```

Then the following TF routine will return `c_value` as the system function return value:

`tf_putstr(0, c_value);`

put 32-bit value on system function return

NOTE

The return of a system function can only be written from a PLI *calltf routine*. The *calltf routine* is the C function that is called when the Verilog simulation executes the system function, and that is the only time the simulation expects a return value from the system function. It is an error to attempt to put a system function return value from a *checktf routine*, *misctf routine* or *sizetf routine*.

Example 9-1 on page 291 of Chapter 9.10 presented the *\$pow* PLI application, which returns a 32-bit value. Example 11-1 on page 336 of this chapter illustrates returning a double-precision value for a *\$realpow* system function.

11.3 Reading and writing 2-state values

Several of the TF routines allow PLI applications to read Verilog logic values from system task/function arguments and automatically convert the values into 2-state C integers or C doubles for use in C programs. These routines automatically convert Verilog 4-state logic values into 2-state values in the C language. Logic X and Z Verilog values are converted to 0 in C, and Verilog strength levels are ignored. The routines require as an input the index number of the task/function argument to be read. A complementary set of routines write C values into a Verilog simulation, and automatically convert C data types into Verilog data types.

11.3.1 Reading 2-state values

int tf_getp(n)

int n index number of a system task/function argument.

The *tf_getp()* routine returns the current value of argument *n* of the system task/function. This routine can access a Verilog vector of up to 32 bits, and returns the Verilog 2-state logic value as a single C 2-state integer, which the PLI assumes is 32 bits wide. An example of using *tf_getp()* is:

```
int c_value;
c_value = tf_getp(1); /* read task/function argument 1 */
```

int tf_getlongp(highvalue, n)

*int *highvalue* pointer to the upper 32-bits of the value.

int n index number of a system task/function argument.

tf_getlongp() retrieves the value of argument *n* of the system task/function argument as a 64-bit integer. The routine accesses a Verilog vector of up to 64 bits, and

returns Verilog 2-state logic value as a pair of C integers, to form a 64-bit vector (again, the PLI assumes C integers are 32-bits). One of the C integers is the return of `tf_getlongp()`. The other C integer is placed into a C `int` variable pointed to as the first input to `tf_getlongp()`. The following example reads task/function argument number 4 as a 64-bit 2-state value:

```
int high_order_bits, low_order_bits;
low_order_bits = tf_getlongp(&high_order_bits, 4);
```

double tf_getrealp(n)

int n index number of a system task/function argument.

The `tf_getrealp()` routine returns the current value of argument *n* of the system task/function as a double. This routine returns a Verilog logic value as a C double (double-precision floating point number). The routine is intended for reading Verilog floating point numbers (literal values with a decimal point or Verilog `real` variables). Real numbers in Verilog are stored using the ANSI definition of a C double precision number. `tf_getrealp()` can also read Verilog vector values. Vectors are automatically converted into 2-state values and from integer to floating point. An example of using `tf_getrealp()` is:

```
double c_value;
c_value = tf_getrealp(3); /* read task/function arg #3 */
```

How Verilog vector widths are converted to C vector widths

The Verilog HDL can represent values in any vector width, from 1-bit wide (scalar) to 1 million bits wide. If a Verilog vector is less than the 32 or 64 bits that the `tf_getp()` and `tf_getlongp()` routines return, the upper bits of the C value are simply zero filled. However, if the Verilog vector size is greater than the 32 or 64 bit return value, then the upper bits (the most-significant bits) of the Verilog value are truncated, which means the value read into the PLI application is not accurate. The routine `tf_sizep()` can be used to determine the vector width of a system task/function argument and flag any potential problems. For example:

```
if (tf_sizep(1) > 64)
    tf_warning("$my_app does not support vectors greater than
               64 bits; some data may be lost\n");
```

11.3.2 Writing 2-state values

The TF library provides 3 routines which write 2-state values represented with C data types onto the arguments of a system task/function. These routines can also be used to write the return value of a system function into simulation.

void tf_putchar(n, value)

<i>int</i>	n	index number of a system task/function arg, or 0.
<i>int</i>	value	a 32-bit integer value.

The `tf_putchar()` routine deposits a 32-bit C integer value onto argument **n** of a system task/function. If **n** is 0, then the value is deposited as the return of a system function. An example of using `tf_putchar()` to return a system function value is:

```
int result
tf_putchar(0, result);
```

void tf_putstr(n, value)

<i>int</i>	n	index number of a system task/function argument, or 0.
<i>int</i>	value	lower (right-most) 32-bits of a 64-bit integer.
<i>int</i>	highvalue	upper (left-most) 32-bits of a 64-bit integer.

`tf_putchar()` deposits a 64-bit integer value, stored as a pair of C integers, onto argument **n** of a system task/function. If **n** is 0, then the value is deposited as the return of a system function. An example of using `tf_putstr()` to write a value into the second system task/function argument is:

```
int high_bits, low_bits;
tf_putstr(2, low_bits, high_bits);
```

void tf_putrealp(n, value)

<i>int</i>	n	index number of a system task/function argument, or 0.
<i>double</i>	value	a double precision real number.

The `tf_putchar()` routine deposits a C double value onto argument **n** of a system task/function. If **n** is 0, then the value is deposited as the return of a system function. An example of using `tf_putstr()` to write a value into the third system task/function argument is:

```
double c_value;
tf_putstr(3, c_value);
```

How Verilog vector widths are converted from C vector widths

When using `tf_putstr()` and `tf_putstrlongp()`, the bit width of a value in C might be different than the bit width of the task/function argument into which the value will be written. When a value is put into a task/function argument, the PLI follows the same rules as Verilog HDL assignment statements when there is a mismatch in the value widths, which are:

- If the C value is wider than the system task/function argument, the left-most bits of the C value are truncated.
- If the C value is narrower than the system task/function argument, the left-most bits of the task/function argument are zero filled.

11.3.3 Returning 2-state system function values

To have the PLI return the value of a system function, the index number for the system task/function argument is set to 0. In order to return a value using index number 0, the PLI application must have been declared as a type of *function* or *real-function* (refer back to Chapter 9, section 9.11 on page 292).

Using TF routines, the following types of C values can be put to a system function return value:

- A *32 bit C integer*, which is returned using `tf_putstr()`. Within the Verilog simulation, the value is received as an unsigned vector of up to 32 bits.
- A pair of *32 bit C integers*, which is returned using `tf_putstrlongp()`. Within the Verilog simulation, the value is received as an unsigned vector of up to 64 bits.
- A *C double*, which is returned using `tf_putstrrealp()`. Within the Verilog simulation, the value is received as a double-precision real number.

Using the ACC library or the VPI library, PLI applications can return system function values for vectors of any bit size. The VPI routines can also return a signed integer value to a system function.

NOTE

The return of a system function can only be written from a PLI *calltf routine*. The *calltf routine* is the C function that is called when the Verilog simulation executes the system function, and that is the only time the simulation expects a return value from the system function. It is an error to attempt to put a system function return value from a *checktf routine*, *miscif routine* or *sizeif routine*.

11.3.4 An example of reading and writing 2-state values

Example 9-1 on page 291 of Chapter 9 shows an example of using the `tf_getp()` and `tf_putp()` routines to read the arguments of the `$pow` system function and return a 32-bit value. Using the TF library, the `$pow` function can easily be modified to read double-precision values as inputs, and return a double-precision value into simulation.

Example 11-1 illustrates using `tf_putp()` or `tf_getrealp()` to read the arguments of a `$realpow` system function. The `tf_typep()` routine is used to determine if the arguments are integer or real number values. The `tf_putrealp()` routine is used to write a real number as the `$realpow` system function return value.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.011/realpow_tf.c`
- Verilog test bench: `Chapter.011/realpow_test.v`
- Verilog-XL results log: `Chapter.011/realpow_test.log`

Example 11-1: `$realpow` — using TF routines to read/write 2-state values

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
/*********************************************
 * Checktf application
 *****/
int PLIBook_realpow_checktf()
{
    int arg_type;

    if (tf_nump() != 2) {
        tf_error("$pow must have 2 arguments.\n");
        return(0);
    }
    arg_type = tf_typep(1);
    if (arg_type == TF_NULLPARAM)
        tf_error("$pow arg 1 cannot be null.\n");
    else
        if (    (arg_type != TF_READONLY)
            && (arg_type != TF_READONLYREAL)
            && (arg_type != TF_READWRITE)
            && (arg_type != TF_READWRITEREAL) ) {
            tf_error("$pow arg 1 must be number, variable or net.\n");
        }
    arg_type = tf_typep(2);
    if (arg_type == TF_NULLPARAM)
        tf_error("$pow arg 2 cannot be null.\n");
    else
```

```
if ( (arg_type != TF_READONLY)
    && (arg_type != TF_READONLYREAL)
    && (arg_type != TF_READWRITE)
    && (arg_type != TF_READWRITEREAL) ) {
    tf_error("$pow arg 2 must be number, variable or net.\n");
}
return(0);
}

*****
* Calltf application
*****
#include <math.h>
int PLIbook_realpow_calltf()
{
    double base, exp, result;
    int     arg_type;

    arg_type = tf_typep(1);
    if ( (arg_type == TF_READONLYREAL)
        || (arg_type == TF_READWRITEREAL) )
        base   = tf_getrealp(1);           /* read double value from tfarg 1 */
    else
        base   = (double)tf_getp(1);    /* read int value from tfarg 1 */

    arg_type = tf_typep(2);
    if ( (arg_type == TF_READONLYREAL)
        || (arg_type == TF_READWRITEREAL) )
        exp    = tf_getrealp(2);         /* read double value from tfarg 2 */
    else
        exp    = (double)tf_getp(2);    /* read int value from tfarg 2 */

    result = pow(base, exp);

    tf_putrealp(0,result);           /* return result */

    return(0);
}
```

11.4 Reading and writing 4-state logic values using C strings

Using strings in the C language is an easy way to represent the 4-state logic values of Verilog. There is a TF routine which allows PLI applications to read Verilog logic values from system task/function arguments and automatically convert the values into 4-state values represented as C strings. A set of TF routines will convert C strings back to Verilog 4-state values and write them onto system task function/arguments. These routines work with Verilog values 0, 1, z, x, but ignore the Verilog strength levels.



TIP Using C strings to represent Verilog logic values is a simple way to represent 4-state logic in a printable form in the C language, and a simple way to work with Verilog vectors of any bit width. However, the conversion from a Verilog logic value to a C string, and vice-versa, is expensive for simulation run-time performance. The best simulation performance will be achieved when Verilog values are represented in the C format that most closely resembles how the value is stored in a Verilog simulation.

11.4.1 Reading values as strings

`char *tf_strgetp(n, format_char)`

int n index number of a system task/function argument.
int format_char character in single quotes representing value format: '**b**' or '**B**' for binary, '**o**' or '**O**' for octal, '**d**' or '**D**' for decimal, '**h**' or '**H**' for hexadecimal.

The `tf_strgetp()` routine returns the value of argument **n** of a system task/function. The routine automatically converts Verilog 4-state logic values into C character strings. The return from `tf_strgetp()` is a pointer to the C string. The string itself is stored in a temporary buffer within the PLI.

The string representation of the Verilog value is in an application-specified radix, which can be binary, octal, decimal or hexadecimal. The first input to `tf_strgetp()` is the index number of the system task/function argument to be read. The second input is a C character (in single quotes), called the *format* character, which indicates the radix in which to represent the value:

- '**b**' or '**B**' indicates the return string should be a binary representation.
- '**o**' or '**O**' indicates the return string should be an octal representation.
- '**d**' or '**D**' indicates the return string should be a decimal representation.
- '**h**' or '**H**' indicates the return string should be a hex representation.

The Verilog value which is read can be a literal number, a constant, a vector of any bit width, a real variable, or a string. The `tf_strgetp()` routine will convert the Verilog value into the C string using the following rules:

- A Verilog vector value will be represented in the C string exactly the same as the value in Verilog. Vectors in Verilog are unsigned, so the value will always be a positive value.
- A Verilog integer value will be represented in the C string as a signed value if a decimal radix is used to read the value. If a binary, octal or hexadecimal radix is used to read the value, then the Verilog integer is represented as an unsigned vector. A neg-

ative integer value read using a binary, octal or hexadecimal radix will be represented in the C string as the two's-complementary of the value.

- A Verilog real value will be represented in the C string as a signed floating point value if a decimal radix is used to read the value. If a binary, octal or hexadecimal radix is used to read the value, then the real number will be rounded off and represented as an integer value. A negative real value read using a binary, octal or hexadecimal radix will be rounded off, and represented in the C string as an unsigned vector containing the two's-complementary of the value.
- A Verilog string value will be converted to a C string. The radix is ignored.

Following is an example of using `tf_strgetp()` to read the value of the second argument system task/function, and represent the value in hex.

```
char *i_str;
i_str = tf_strgetp(2,'h'); /* read arg 2 as 4-state string */
io_printf("Value of arg2 in PLI is %s\n", i_str);
```

NOTE → `tf_strgetp()` returns a pointer to a string, which is stored in a temporary string buffer. Because the buffer is temporary, the pointer will not remain valid. The PLI application should either use the string immediately, or copy the string into application-allocated storage.

11.4.2 Writing values represented with strings into Verilog

The TF library includes a set of routines which perform three key functions:

- Automatically convert logic values represented as C character strings into Verilog 4-state values.
- Write the value into a system task argument.
- Schedule a future time in simulation for the written value to take effect, similar to the propagation delay from an input change on a logic gate to the output change.

`int tf_stredputp(n, length, format, value, delay, mode)`

<code>int</code>	<code>n</code>	index number of a system task/function argument.
<code>int</code>	<code>length</code>	number of bits to be deposited.
<code>int</code>	<code>format</code>	character in single quotes, representing value format.
<code>char</code>	<code>* value</code>	string representing the value to be deposited.
<code>int</code>	<code>delay</code>	32-bit time value representing delay before value is deposited.
<code>int</code>	<code>mode</code>	code representing delay mode: 0 for inertial, 1 for modified transport, 2 for pure transport.

int tf_strlongdelputp(n, length, format, value, lowdelay, highdelay, mode)

<i>int</i>	n	index number of a system task/function argument.
<i>int</i>	length	number of bits to be deposited.
<i>int</i>	format	character in single quotes, representing value format.
<i>char</i>	* value	string representing the value to be deposited.
<i>int</i>	lowdelay	lower (right-most) 32-bits of a 64-bit time value representing delay before value is deposited.
<i>int</i>	highdelay	upper (left-most) 32-bits of a 64-bit time value representing delay before value is deposited.
<i>int</i>	mode	code representing delay mode: <i>0</i> for inertial, <i>1</i> for modified transport, <i>2</i> for pure transport.

int tf_strrealdelputp(n, length, format, value, delay, mode)

<i>int</i>	n	index number of a system task/function argument.
<i>int</i>	length	number of bits to be deposited.
<i>int</i>	format	character in single quotes, representing value format.
<i>char</i>	* value	string representing the value to be deposited.
<i>double</i>	delay	real time value representing delay before value is deposited.
<i>int</i>	mode	code representing delay mode: <i>0</i> for inertial, <i>1</i> for modified transport, <i>2</i> for pure transport.

There are several inputs for `tf_strdelputp()`, `tf_strlongdelputp()` and `tf_strrealdelputp()`:

- **index** is the index number of the system task/function argument into which the value is to be written. It is illegal to specify an index of 0 to put the value as the return value of a system function using these routines.
- **length** is the bit-width of the Verilog vector value the string will be converted into, and should be set to the width of the task/function argument (which can be determined using `tf_sizep()`). The value represented as a C string will be written into the Verilog vector using the assignment rules of the Verilog language, which means that if the value is less than the maximum value which the vector stores, the upper bits of the vector are zero filled. If the value is greater than the maximum value which the vector stores, the most significant bits of the value are truncated.
- **format** is the radix format of the string representation of the value. The format is indicated as one of the following C characters:
 - '**b**' or '**B**' indicates the string is a binary representation.
 - '**o**' or '**O**' indicates the string is an octal representation.
 - '**d**' or '**D**' indicates the string is a decimal representation.
 - '**h**' or '**H**' indicates the string is a hexadecimal representation.

- **value** is a pointer to a string, or a literal string, with the logic value. The string may only contain characters which are legal for the format that is selected. An error status will be returned if the value is invalid, and the value may or may not be written into the simulation. The legal values are listed in table 11-1:

Format	Legal Value Characters
binary	0, 1, x, X, z, Z
octal	0—7, x, X, z, Z
decimal	0—9
hexadecimal	0—9, a, A, b, B, c, C, d, D, e, E, f, F, x, X, z, Z

Table 11-1: Legal characters for the TF string put routines

- **delay** is the amount of propagation delay to transpire between when the put routine is executed and when the logic change occurs in Verilog simulation.
- **mode** is the method of event scheduling that the simulation should use, represented by a literal integer of **0**, **1** or **2**, where **0** is *inertial delay*, **1** is *modified transport delay*, and **2** is *transport delay*. The delay modes define what the Verilog simulator should do if there is already a pending change scheduled for the same signal.
 - *Inertial delay*—all pending events on the signal are cancelled.
 - *Modified transport delay*—all pending events on the signal which are at a later time than this event are cancelled, which means the last event to be scheduled will always be the last event that occurs.
 - *Transport delay*—no pending events on the signal are cancelled, which means the last event to be scheduled may not be the last event which occurs.

The following examples illustrate several ways of writing values represented as strings into Verilog.

Write a literal C string with a hex value onto system task/function argument 2, as a 16-bit vector, and with no propagation delay:

```
tf_strdelpup(2, 16, 'H', "F5xZ", 0, 0);
```

Write a binary value stored in the string **value_str** onto system task/function argument 1, a a 4-bit vector, and with a 25 time unit propagation delay, using inertial delay propagation:

```
char value_str[5];
strcpy(value_str, "1x0z");
tf_strdelpup(1, 4, 'b', value_str, 25, 0);
```

Write a hex value stored in the string vector onto task/function argument 2, using the bit-width of the argument, and with a 5.2 time unit propagation delay, using transport delay propagation:


```
double delay;
char *vector;
vector = malloc(tf_sizep(2) + 1);
strcpy(vector, "zzzzzzzz");
delay = 5.2;
tf_strrealdelputp(2, tf_sizep(2), 'h', vector, delay, 2);
```

Specifying propagation delays when writing values:

`tf_strdelputp()` uses a 32-bit integer value to specify the propagation delay time. `tf_strlongdelputp()` uses a 64-bit time value, specified using two 32-bit integers for the high-order and low-order words of the time. `tf_strrealdelputp()` uses a double precision real number to specify the delay time.

To represent a 64 bit time value, the Verilog PLI uses a pair of C unsigned integers to store the full 64 bits of simulation time. The lower 32 bits of Verilog time are placed in one integer, and the upper 32 bits of the Verilog time are stored in a second integer, as shown in the following illustration:

```
unsigned int high, low;
```



Within simulation, delays are scaled to a time unit specified in the Verilog models (using the `'timescale` compiler directive). Time scaling permits each Verilog model to represent time in different units. For example, one model can represent time in nanoseconds with two decimal points of precision, and another model can represent delays in microseconds with no decimal points of precision.

The `tf_strdelputp()`, `tf_strlongdelputp()` and `tf_strrealdelputp()` routines schedule an event at a future simulation time, using the time scale of the Verilog module from which the PLI application was called. Therefore, a time value of 15

in one of these routines could represent 15 nanoseconds, 15 picoseconds, or some other time unit. The following example illustrates this automatic time scaling:

Assume a PLI application contained the following statement:

```
tf_strdelputp(1, 4, 'b', value_str, 15, 0);
```

In the following Verilog module, the delay of 15 in the PLI application would represent 15 nanoseconds.

```
'timescale 1ns/1ns
module test_chip1;
...
always @(posedge clk)
$read_test_vector("chip1.vectors", chip1_in);
...
```

In the next Verilog module, the same delay of 15 in the PLI application would represent 15 picoseconds.

```
'timescale 1ps/1fs
module test_chip2;
...
always @(posedge clk)
$read_test_vector("chip2.vectors", chip2_in);
...
```

Having the TF *put* routines use the calling module's time scale is generally desirable. It means the PLI application developer does not need to be concerned with time scales—the PLI automatically adjusts time to match the model in which the application is used. If needed, however, other TF routines allow the developer to read a module's time scale information, and to manually scale delays within a PLI application. Refer to Chapter 10, section 10.8 on page 316 for details on these routines.

Limitations on using TF string put routines

There are two important restrictions on using the `tf_strdelputp()`, `tf_strlongdelputp()` and `tf_sturrealdelputp()` routines.

- These routines cannot be used to specify the return value of a system function. The Verilog standard requires that system function return values must be returned immediately, with no delay. Only `tf_putp()`, `tf_putlongp()` and `tf_putrealt()` can be used to return the value of a system function (the ACC and VPI libraries also have routines which can write the return value of system functions).

- These routines may not be used from a *misctf routine* that was called for REASON_ROSYNCH. At this read-only synchronization callback, the TF routines are prohibited from scheduling any new events in simulation.

No delay versus zero delay

The Verilog PLI makes a distinction between putting a value into simulation with no delay and putting a value into simulation with zero delay.

- The `tf_putstr()`, `tf_putstrlongp()` and `tf_putstrrealp()` routines write a value into Verilog simulation instantly. When the PLI application returns back to simulation, any values written to a system task argument or system function return using these routines will already be in effect for Verilog HDL statements to use.
- The `tf_strdelputp()`, `tf_strlongdelputp()` and `tf_strrealdelputp()` routines schedule a value to be written into simulation. If a delay of zero is specified, the value is scheduled to be written into the task argument later in the current simulation time step. *When the system task returns back to simulation, the scheduled value will not yet have taken effect*, and other Verilog HDL statements scheduled to execute in the same simulation time step may or may not see the new value of the system task argument (depending on where the value change which was scheduled by the PLI falls in the simulator's event queue in relation to other Verilog HDL events).

The following simple Verilog HDL source code illustrates the potential problem of putting a value into simulation using a delay of zero.

```
module test;
  reg [7:0] reg1, reg2;
  initial
    begin
      reg1 = 0; reg2 = 0;
      $put_value(reg1, reg2);
      $display("reg1=%d    reg2=%d", reg1, reg2);
      $strobe ("reg1=%d    reg2=%d", reg1, reg2);
      #1 $finish;
    end
endmodule
```

If the *calltf routine* for `$put_value` puts a value into `reg1` using `tf_putstr()`, then when `$put_value` returns to the simulation, and the `$display` statement prints the value of `reg1`, the new value will be printed.

If, however, the *calltf routine* for `$put_value` puts a value into `reg2` using `tf_strdelputp()` with a delay of zero, then when `$put_value` returns to the simula-

tion, and the `$display` statement prints the value of `reg2`, the *old* value will be printed. The old value is printed because the value written by the PLI has been scheduled to take place in the current time step, but it will not yet have taken effect. The `$strobe` statement which follows the `$display` will print the new value of both `reg1` and `reg2`, because the definition of `$strobe` is to print its message at the end of the current simulation time step, after all value changes for that moment in time have taken effect.

11.5 Reading Verilog strings

In the Verilog HDL, a system task/function argument can be a string. For example:

```
$read_test_vector("vector_file.pat", input_vector);
```

A string in Verilog is represented as a vector containing the 8-bit ASCII code for each character, with no string termination character. The TF library provides a routine to read a Verilog string and convert it to a C string.

`char *tf_getcstringp(n)`

`int n` index number of a system task/function argument.

`tf_getcstringp()` returns the value of argument `n` of a system task/function argument as a string, and returns `null` if an error occurs. The routine requires the index number of the task/function argument as an input, and expects that the value of the argument be a literal string or a vector with valid ASCII characters. The following example reads a file name stored in system task/function argument 1.

In a *checktf* routine:

```
if (tf_typep(1) != TF_STRING)
    tf_error("$read_vector arg 1 must be literal string\n");
```

In a *calltf* routine:

```
char *file_name;
file_name = tf_getcstringp(1);
```

NOTE → `tf_getcstringp()` returns a string pointer to a string, which is stored in a temporary string buffer. Because the buffer is temporary, the pointer will not remain valid. The PLI application should either use the string immediately, or copy the string into application-allocated storage. Refer to section 10.3 for details on PLI strings.

11.6 Reading and writing 4-state values using aval/bval encoding

There are three TF routines which read and write 4-state values of system task/function arguments, and represent the values as pairs of C integers which encode the Verilog 4-state values. These routines also support Verilog real values and Verilog string values, and represent these values as a C double or C string, respectively (Verilog reals and strings use 2-state logic).

p_tfexprinfo tf_exprinfo(n, info)

<i>int n</i>	index number of a system task/function argument.
<i>p_tfexprinfo *info</i>	pointer to an application-allocated <i>s_tfexprinfo</i> structure to receive the expression information.

The *tf_exprinfo()* routine reads detailed information about a system task/function argument (an *expression*). The information is retrieved into an *s_tfexprinfo* structure, which then contains the general data type of the argument, its current logic value, the bit size of vectors, and the sign property. Integer and literal values in Verilog can be signed or unsigned. The routine returns the value of *info* if successful and 0 if an error occurred.

void tf_evaluatep(n)

<i>int n</i>	index number of a system task/function argument.
--------------	--------------------------------------------------

tf_evaluatep() re-reads the logic value of a system task/function argument. The value of argument *n* must have been previously read using *tf_exprinfo()* before *tf_evaluatep()* may be used. The most current logic value is retrieved into the same C structure which was allocated for *tf_exprinfo()*. Therefore, the PLI application must maintain the *s_tfexprinfo* structure.

void tf_propagatep(n)

<i>int n</i>	index number of a system task/function argument.
--------------	--------------------------------------------------

tf_propagatep() writes a logic value to a system task/function argument. The value of argument *n* must have been previously read using *tf_exprinfo()* before *tf_evaluatep()* may be used. The new logic value to be written is stored in the same C structure which was allocated for *tf_exprinfo()*. Therefore, the PLI application must maintain the *s_tfexprinfo* structure. As soon as the value is written to the argument, the Verilog simulation will propagate the change to any Verilog expressions which read the value of that argument.

The **s_tfexprinfo** structure used by `tf_exprinfo()`, `tf_evaluatep()` and `tf_propagatep()` is defined in the veriuser.h file. The structure definition is listed below, followed by the explanation of the fields in structure.

```
typedef struct t_tfexprinfo
{
    short expr_type;
    short padding;
    struct t_vecval *expr_value_p;
    double real_value;
    char *expr_string;
    int expr_ngroups;
    int expr_vec_size;
    int expr_sign;
    int expr_lhs_select;
    int expr_rhs_select;
} s_tfexprinfo, *p_tfexprinfo;
```



NOTE The PLI application must first allocate memory for an **s_tfexprinfo** structure, and then pass a pointer to the memory to `tf_exprinfo()`. The routine will then fill in the fields of the structure. For example:

To allocate temporary storage that will be freed when the PLI application exits:

```
s_tfexprinfo arg_info;  
tf_exprinfo(n, &arg_info);
```

To allocate persistent storage that can be used each time the PLI application is called:

```
p_tfexprinfo arg_info;  
arg_info = (p_tfexprinfo)malloc(sizeof(s_tfexprinfo));  
tf_exprinfo(n, arg_info);
```



TIP

Using `malloc()`, rather than a local variable, allows the same **s_tfexprinfo** structure to be preserved from one call to the PLI application to another call. The PLI application must maintain a pointer to the structure so that it is available when needed. The TF work area is a good place to store the pointer, because the work area is both persistent and unique to each instance of a system task.

The fields in the **s_tfexprinfo** structure are explained in table 11-2, which follows.

s_tfexprinfo field	Definition
expr_type	An integer constant which represents the general Verilog data type of the system task/function argument. The <code>expr_type</code> determines which of the remaining fields of the <code>s_tfexprinfo</code> structure will be used. The constants are:
	<code>TF_NULLPARAM</code> the argument is null
	<code>TF_STRING</code> the argument is a string
	<code>TF_READONLY</code> the argument is a scalar net, vector net, net bit select, or net part select
	<code>TF_READONLYREAL</code> the arg is a constant real number
	<code>TF_READWRITE</code> the argument is a scalar <code>reg</code> , vector <code>reg</code> , <code>integer</code> , or <code>time</code>
	<code>TF_READWRITEREAL</code> the argument is a <code>real</code> variable
	<code>TF_RWBITLELECT</code> the argument is a bit select of a vector <code>reg</code> , <code>integer</code> , or <code>time</code>
	<code>TF_RWPARTSELECT</code> the argument is a part select of a vector <code>reg</code> , <code>integer</code> , or <code>time</code>
	<code>TF_RWMEMSELECT</code> the argument is a word select of a Verilog memory array
padding	not used
expr_value_p	if <code>expr_type</code> is <code>TF_READONLY</code> , <code>TF_READWRITE</code> , <code>TF_RWBITLELECT</code> , <code>TF_RWPARTSELECT</code> , or <code>TF_RWMEMSELECT</code> , this field contains a pointer to an array of <code>s_vecval</code> structures containing the 4-state logic value
real_value	if <code>expr_type</code> is <code>TF_READONLYREAL</code> , or <code>TF_READWRITEREAL</code> , then this field contains the real value
expr_string	if <code>expr_type</code> is <code>TF_STRING</code> , then this field contains a pointer to the string
expr_ngroups	if <code>expr_type</code> is <code>TF_READONLY</code> , <code>TF_READWRITE</code> , <code>TF_RWBITLELECT</code> , <code>TF_RWPARTSELECT</code> , or <code>TF_RWMEMSELECT</code> , then this field contains the number of elements in the array of <code>s_vecval</code> structures pointed to in the <code>expr_value_p</code> field

Table 11-2: The `tf_exprinfo` structure

s_tfexprinfo field	Definition
expr_vec_size	if <i>expr_type</i> is TF_READONLY, TF_READWRITE, TF_RWBITLELECT, TF_RWPARTSELECT, or TF_RWMEMSELECT, then this field contains the number of bits in task/function argument
expr_sign	contains a flag indicating the sign type of the task/function argument, where 0 represents an unsigned value, and 1 represents a signed value
expr_lhs_select expr_rhs_select	if <i>expr_type</i> is TF_RWBITLELECT, TF_RWPARTSELECT, or TF_RWMEMSELECT, then these fields contain the left-hand and right-hand indices of the bit or part select. Note: some simulators do not use these fields. PLI applications should not depend on these fields, because the application may not be portable to all simulators.

Table 11-2: The `tf_exprinfo` structure (continued)

Representing 4-state logic values as encoded C integers


The expression types `TF_READONLY`, `TF_READWRITE`, `TF_RWBITLELECT`, `TF_RWPARTSELECT` and `TF_RWMEMSELECT` use an *aval/bval* pair of C integers to encode the 4 logic values of Verilog. The encoding uses 1 bit of each aval/bval pair to represent 1 bit of Verilog logic values. The encoding is:

aval/bval pair	Verilog logic value represented
0/0	0
1/0	1
0/1	Z
1/1	X

Table 11-3: aval/bval logic value encoding

The TF library assumes C integers are 32-bits wide, and therefore uses the aval/bval pair to encode up to 32 bits of a Verilog vector. By using an array of aval/bval integer pairs, vector lengths of any size may be represented. The representation of a 40-bit vector in Verilog can be visualized as:

For the Verilog declaration: `reg [39:0] data;`




The Verilog language supports any numbering convention for a vector's bit numbers. The least significant bit of the Verilog vector can be the smallest bit number, such as bit 0 (which is referred to as little endian convention). Or, the least significant bit of the Verilog vector can be the largest bit number, such as bit 39 (which is referred to as big endian convention). Verilog does not require there be a bit zero at all. Each of the following examples are valid vector declarations in Verilog:

```
reg [39:0] data;      /* little endian -- LSB is bit 0 */
reg [0:39] data2;    /* big endian   -- LSB is bit 39 */
reg [40:1] data3;    /* little endian -- LSB is bit 1 */
```

The bit numbering used in Verilog does not affect the aval/bval representation of the Verilog vector. In the array of aval/bval pairs, the LSB of the Verilog vector will

always be the LSB of the first C integer in the array, and the MSB of the Verilog vector will always be the last bit in the array which is used. The following diagram illustrates the aval/bval array for a Verilog vector declared with a big endian convention.

For the Verilog declaration: `reg [1:40] data;`



lsb = least significant bit

msb = most significant bit

The aval/bval pair is declared in an `s_vecval` structure, which is defined in the veri-user.h TF library. The structure definition is:

```
typedef struct t_vecval
{
    int avalbits;
    int bvalbits;
} s_vecval, *p_vecval;
```



The memory for the `s_vecval` structure is allocated and maintained by the simulator. The simulator will fill in the fields of the structure, and place a pointer to the `s_vecval` structure in the `s_tfexprinfo` structure that was allocated by the PLI application. The pointer to the `s_tfexprinfo` structure is passed to the simulator as an argument to `tf_exprinfo()`.

To read a task/function argument's 4-state logic value using `tf_exprinfo()` involves three basic steps:

1. Allocate an `s_tfexprinfo` structure.
2. Call `tf_exprinfo()`, giving a pointer to the `s_tfexprinfo` structure as an argument to `tf_exprinfo()`.
3. Read the logic value of the task/function argument from the `s_tfexprinfo` structure. Use the `expr_type` field in the structure to determine what field in the structure contains the logic value. If the task/function argument is a scalar or vector signal, the logic value will be stored in an array of `s_vecval` structures, which is pointed to in the `expr_value_p` field of the `s_tfexprinfo` structure.

Example 11-2 lists a PLI application which uses `tf_exprinfo()` to read the 4-state value of any type of task/function argument, including any size of vector. Note that example 11-3 on page 354 lists a more efficient version of this example, called `$read_4state_value`.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.011/exprinfo_test_tf.c
- Verilog test bench: Chapter.011/exprinfo_test.v
- Verilog-XL results log: Chapter.011/exprinfo_test.log

Example 11-2: `$exprinfo_test` — using `tf_exprinfo()` to read 4-state values

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
int PLIBook_exprinfoTest_calltf()
{
    s_tfexprinfo info_s;      /* structure for tf_exprinfo() */
    p_vecval     val_array;   /* pointer to value array in info struct */
    int i;

    tf_exprinfo(1, &info_s); /* read expression info for arg 1 */

    io_printf("Expression info:\n");
    switch (info_s.expr_type) {
        case TF_NULLPARAM:    io_printf(" type = TF_NULLPARAM\n"); break;
        case TF_STRING:       io_printf(" type = TF_STRING\n"); break;
        case TF_READONLY:     io_printf(" type = TF_READONLY\n"); break;
        case TF_READONLYREAL: io_printf(" type = TF_READONLYREAL\n"); break;
        case TF_READWRITE:    io_printf(" type = TF_READWRITE\n"); break;
        case TF_READWRITEREAL:io_printf(" type = TF_READWRITEREAL\n");break;
        case TF_RWBITSELECT:  io_printf(" type = TF_RWBITSELECT\n"); break;
        case TF_RWPARTSELECT: io_printf(" type = TF_RWPARTSELECT\n"); break;
        case TF_RWMEMSELECT:  io_printf(" type = TF_RWMEMSELECT\n"); break;
        default: io_printf(" type is unknown (%d)\n", info_s.expr_type);
    }
}
```

```
io_printf(" ngroups = %d\n",      info_s.expr_ngroups);
io_printf(" vector size = %d\n",   info_s.expr_vec_size);
io_printf(" sign = %d\n",          info_s.expr_sign);
io_printf(" LHS select = %d\n",    info_s.expr_lhs_select);
io_printf(" RHS select = %d\n",    info_s.expr_rhs_select);

switch (info_s.expr_type) {
    case TF_STRING:
        io_printf(" string value = %s\n", info_s.expr_string); break;
    case TF_READONLYREAL:
    case TF_READWRITEREAL:
        io_printf(" real value = %f\n", info_s.real_value); break;
    case TF_READONLY:
    case TF_READWRITE:
    case TF_RWBITLETSELECT:
    case TF_RWPARTSELECT:
    case TF_RWMEMSELECT:
        val_array = info_s.expr_value_p;
        io_printf(" vector value (in hex):\n");
        for (i=0; i<info_s.expr_ngroups; i++) {
            io_printf("  avalbits[%d] = %x\n", i, val_array[i].avalbits);
            io_printf("  bvalbits[%d] = %x\n", i, val_array[i].bvalbits);
        }
        break;
}
io_printf("\n\n");
return(0);
}
```

11.6.1 Re-reading values previously read with `tf_exprinfo()`

void tf_evaluatep(n)

int n index number of a system task/function argument.

The `tf_evaluatep()` routine can be used to re-read the logic value of a system task/function argument that was previously read using `tf_exprinfo()`. The `tf_evaluatep()` routine retrieves the current value of the argument, and updates the value in the same structure that was used by `tf_exprinfo()`. Since only the value information in the structure is updated, `tf_evaluatep()` executes more efficiently than calling `tf_exprinfo()` a second time. *Note:* In order to use `tf_evaluatep()`, the PLI application must maintain the `s_tfexprinfo` structure.

The following PLI application example calls `tf_exprinfo()` from a *misctf routine* to read the value of the first task.function argument. The memory for the

`s_tfexprinfo` is allocated using `malloc` so that the memory persists after the *misctf routine* has exited. The pointer to the structure is then stored in the PLI work area. At some later simulation time, a *calltf routine* retrieves the structure pointer from the work area, and calls `tf_evaluatep()` to re-read the task/function argument's value. The PLI work area was presented in Chapter 10, section 10.7 on page 313, and the *misctf routine* is discussed in more detail in Chapter 12.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.011/read_4state_value_tf.c
- Verilog test bench: Chapter.011/read_4state_value_test.v
- Verilog-XL results log: Chapter.011/read_4state_value_test.log

Example 11-3: `$read_4state_value` — using the `tf_evaluatep()` routine

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
/*********************misctf application*************************/
* misctf application
*
* The misctf application is used to call tf_exprinfo() at the
* beginning of simulation, so that the memory allocated by
* tf_exprinfo() is only allocated one time for each instance of
* $read_4state_value.
/*********************misctf application*************************/
int PLIbook_Read4stateValue_misctf(int user_data, int reason)
{
    p_tfexprinfo info_p;      /* pointer to structure for tf_exprinfo() */
    p_vecval     val_array;   /* pointer to value array in info struct */
    int i;
    if (reason != REASON_ENDOFCOMPILE)
        return(0); /* exit now if this is not the start of simulation */

    /* allocate memory for an s_tfexprinfo structure */
    info_p = (p_tfexprinfo)malloc(sizeof(s_tfexprinfo));
    tf_exprinfo(1, info_p); /* read expression info for arg 1 */
    tf_setworkarea((char *)info_p); /* save info pointer in workarea */

    io_printf("Expression info:\n");
    switch (info_p->expr_type) {
        case TF_NULLPARAM:    io_printf(" type = TF_NULLPARAM\n"); break;
        case TF_STRING:       io_printf(" type = TF_STRING\n"); break;
        case TF_READONLY:     io_printf(" type = TF_READONLY\n"); break;
        case TF_READONLYREAL: io_printf(" type = TF_READONLYREAL\n"); break;
        case TF_READWRITE:    io_printf(" type = TF_READWRITE\n"); break;
        case TF_READWRITEREAL:io_printf(" type = TF_READWRITEREAL\n");break;
        case TF_RWBITSELECT:  io_printf(" type = TF_RWBITSELECT\n"); break;
        case TF_RWPARTSELECT: io_printf(" type = TF_RWPARTSELECT\n"); break;
        case TF_RWMEMSELECT:  io_printf(" type = TF_RWMEMSELECT\n"); break;
        default: io_printf(" type is unknown (%d)\n", info_p->expr_type);
    }
}
```

```
io_printf(" ngroups = %d\n", info_p->expr_ngroups);
io_printf(" vector size = %d\n", info_p->expr_vec_size);
io_printf(" sign = %d\n", info_p->expr_sign);
io_printf(" LHS select = %d\n", info_p->expr_lhs_select);
io_printf(" RHS select = %d\n", info_p->expr_rhs_select);
switch (info_p->expr_type) {
    case TF_STRING:
        io_printf(" string value = %s\n", info_p->expr_string); break;
    case TF_READONLYREAL:
    case TF_READWRITEREAL:
        io_printf(" real value = %f\n", info_p->real_value); break;
    case TF_READONLY:
    case TF_READWRITE:
    case TF_RWBITSELECT:
    case TF_RWPARTSELECT:
    case TF_RWMEMSELECT:
        val_array = info_p->expr_value_p;
        io_printf(" vector value (in hex):\n");
        for (i=0; i<info_p->expr_ngroups; i++) {
            io_printf("  avalbits[%d] = %x\n", i, val_array[i].avalbits);
            io_printf("  bvalbits[%d] = %x\n", i, val_array[i].bvalbits);
        }
        break;
    }
    io_printf("\n\n");
    return(0);
}

/*********************************************
 * calltf application
 *****/
int PLIbook_Read4stateValue_calltf()
{
    p_tfexprinfo info_p;      /* pointer to structure for tf_exprinfo() */
    p_vecval     val_array;   /* pointer to value array in info struct */
    int i;

    info_p = (p_tfexprinfo)tf_getworkarea(); /* retrieve info pointer */
    tf_evaluatep(1);                      /* re-read value of arg 1 */

    switch (info_p->expr_type) {
        case TF_STRING:
            io_printf(" string value = %s\n", info_p->expr_string); break;
        case TF_READONLYREAL:
        case TF_READWRITEREAL:
            io_printf(" real value = %f\n", info_p->real_value); break;
        case TF_READONLY:
        case TF_READWRITE:
        case TF_RWBITSELECT:
        case TF_RWPARTSELECT:
        case TF_RWMEMSELECT:
            val_array = info_p->expr_value_p;
            io_printf(" vector value (in hex):\n");
            for (i=0; i<info_p->expr_ngroups; i++) {
```

```

        io_printf("    avalbits[%d] = %x\n", i, val_array[i].avalbits);
        io_printf("    bvalbits[%d] = %x\n", i, val_array[i].bvalbits);
    }
    break;
}
io_printf("\n");
return(0);
}

```

11.6.2 Writing to task/function arguments more than once

void tf_propagatep(n)

int n index number of a system task/function argument.

The `tf_propagatep()` routine is used to write 4-state logic values to a system task/function argument that was previously read using `tf_exprinfo()`. The value to be written must be placed in the same `s_tfexprinfo` structure that was allocated for `tf_exprinfo()`. Scalar and vector logic values are represented using `aval/bval` pairs in an array of `s_vecval` structures, which were allocated by the simulator when `tf_exprinfo()` was called.

`tf_propagatep()` can only modify the value of the *writable* Verilog data types supported by `tf_exprinfo()`. This means the `expr_type` must be `TF_READWRITE`, `TF_READWRITEREAL`, `TF_RWBITSELECT`, `TF_RWPARTSELECT` or `TF_RWMEMSELECT`.

The `tf_propagatep()` routine can also be used to modify and propagate changes to any Verilog memory word. This requires that the entire Verilog memory array be accessed using `tf_nodeinfo()`, instead of `tf_exprinfo()`. The value to be modified is then contained in a character array pointed to in the `s_tfnodeinfo` structure that is used by `tf_nodeinfo()`. Section 11.8 on page 367 discusses using `tf_nodeinfo()` to access and modify the contents of Verilog memories.

NOTE

In order to use `tf_propagatep()` from one call of a PLI application to another, the pointer to the `s_tfexprinfo` or `s_tfnodeinfo` structure must be maintained by the application. Alternatively, the pointer to the field within the structure which contains the value (or pointer to the value) to be modified can be maintained. The PLI work area can be used to save the pointer.

Example 11-4 shows how the logic value of a Verilog signal can be read and then modified using `tf_exprinfo()`, `tf_evaluatep()`, and `tf_propagatep()`.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.011/propagatetp_test_tf.c
- Verilog test bench: Chapter.011/propagatetp_test.v
- Verilog-XL results log: Chapter.011/propagatetp_test.log

Example 11-4: \$propagatetp_test — using tf_propagatetp() to write 4-state values

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
/*********************************************
 * checktf application
 *****/
int PLIbook_propagatetpTest_checktf()
{
    if (tf_numtp() != 1)
        tf_error("Usage error: $propagatetp_test(<signal>);");
    return(0);
}

/*********************************************
 * misctf application
 *
 * The misctf application is used to call tf_exprinfo() at the
 * beginning of simulation, so that the memory allocated by
 * tf_exprinfo() is only allocated one time for each instance of
 * $read_4state_value.
 *****/
int PLIbook_propagatetpTest_misctf(int user_data, int reason)
{
    p_tfexprinfo info_p;      /* pointer to structure for tf_exprinfo() */

    if (reason != REASON_ENDOFCOMPILE)
        return(0); /* exit now if this is not the start of simulation */

    /* allocate memory for an s_tfexprinfo structure */
    info_p = (p_tfexprinfo)malloc(sizeof(s_tfexprinfo));

    tf_exprinfo(1, info_p); /* read expression info for arg 1 */
    if ( (info_p->expr_type != TF_READWRITE)
        && (info_p->expr_type != TF_READWRITEREAL)
        && (info_p->expr_type != TF_RWBITSELECT)
        && (info_p->expr_type != TF_RWPARTSELECT)
        && (info_p->expr_type != TF_RWMEMSELECT) ) {
        io_printf("ERROR: Signal type not supported by $propagatetp_test\n");
        tf_dofinish();
    }
    else
        tf_setworkarea((char *)info_p); /* save info pointer in workarea */
    return(0);
}
```

```
*****
 * calltf application
 ****
/* prototype for subroutine used by calltf routine */
void PLIbook_Print4stateValue();

int PLIbook_propagatepTest_calltf()
{
    p_tfexprinfo info_p;      /* pointer to structure for tf_exprinfo() */
    info_p = (p_tfexprinfo)tf_getworkarea(); /* retrieve info pointer */

    io_printf("$propagatep_test called at time %d to read & modify arg
value\n",
              tf_gettime());
    /* read current value of argument 1 */
    io_printf(" current value:\n");
    tf_evaluatep(1);
    PLIbook_Print4stateValue(info_p);

    /* modify value of argument 1 */
    switch (info_p->expr_type) {
        case TF_READWRITE:
        case TF_RWBITLELECT:
        case TF_RWPARTSELECT:
        case TF_RWMEMSELECT:
            info_p->expr_value_p[0].avalbits++;
            info_p->expr_value_p[0].bvalbits = 0;
            break;
        case TF_READWRITEREAL:
            info_p->real_value++;
            break;
    }
    tf_propagatep(1);

    /* read new value of argument 1 */
    io_printf(" new value:\n");
    tf_evaluatep(1);
    PLIbook_Print4stateValue(info_p);

    return(0);
}

void PLIbook_Print4stateValue(p_tfexprinfo info_p)
{
    int i;

    switch (info_p->expr_type) {
        case TF_READWRITEREAL:
            io_printf(" real value = %0.1f\n", info_p->real_value); break;
        case TF_READWRITE:
        case TF_RWBITLELECT:
        case TF_RWPARTSELECT:
        case TF_RWMEMSELECT:
```

```

case TF_RWMEMSELECT:
    io_printf("  vector value (in hex):\n");
    for (i=0; i<info_p->expr_ngroups; i++) {
        io_printf("    avalbits[%d] = %x\n",
                  i, info_p->expr_value_p[i].avalbits);
        io_printf("    bvalbits[%d] = %x\n",
                  i, info_p->expr_value_p[i].bvalbits);
    }
    break;
}
return;
}

```

11.7 Reading 4-state logic values with strengths

The `tf_nodeinfo()` routine reads information about a system task/function argument (referred to as a *node*), including 4-state logic values and logic strengths.

`p_tfnodeinfo tf_nodeinfo(n, info)`

<code>int n</code>	index number of a system task/function argument.
<code>p_tfnodeinfo *info</code>	pointer to the <code>s_tfnodeinfo</code> structure containing data on a system task/function argument.

The `tf_nodeinfo()` routine retrieves node information about a system task/function argument and places the information into an `s_tfnodeinfo` structure pointed to by `info`. The routine returns the value of `info` if successful, and 0 if an error occurred.

The `tf_nodeinfo()` routine is similar to `tf_exprinfo()`, but differs in these important ways:

- `tf_nodeinfo()` types more closely matches Verilog data types.
- `tf_nodeinfo()` can read the strength values of scalar nets.
- `tf_nodeinfo()` can read and modify the contents of entire Verilog memory arrays (`tf_exprinfo()` can only access word selects of a memory array).
- `tf_nodeinfo()` *cannot* read expressions in system task/function arguments (including literal values, bit and part selects of vectors, and strings).
- `tf_nodeinfo()` *cannot* be used with `tf_evaluatep()` to re-read an argument's value, and `tf_nodeinfo()` cannot be used with `tf_propagatep()` to modify an argument's value.

NOTE

PLI applications which use the `tf_nodeinfo()` routine may not portable to all Verilog simulators. The Verilog language permits arbitrarily complex expressions to be used as system task/function arguments, but the IEEE 1364-1995 Verilog standard does not clearly specify how `tf_nodeinfo()` should handle these expressions. For maximum portability, a PLI application which uses `tf_nodeinfo()` should restrict the system task/function arguments that are accessed with the routine to simple expressions, such as scalar signals, complete vectors, or word selects of arrays. *The `tf_nodeinfo()` routine may not return the same results on all simulators if a system task/function argument is a bit select or part select of a vector.*

`tf_nodeinfo()` uses an **s_tfnodeinfo** structure to receive the value and information of a task/function argument. This structure is defined in the veriuser.h file. The structure is listed below, and table 11-4, which follows, describes the fields of the structure.

```
typedef struct t_tfnodeinfo
{
    short node_type;
    short padding;
    union
    {
        struct t_vecval *vecval_p;
        struct t_strengthval *strengthval_p;
        char *memoryval_p;
        double *real_val_p;
    } node_value;
    char *node_symbol;
    int node_ngroups;
    int node_vec_size;
    int node_sign;
    int node_ms_index;
    int node_ls_index;
    int node_mem_size;
    int node_lhs_element;
    int node_rhs_element;
    int *node_handle;
} s_tfnodeinfo, *p_tfnodeinfo;
```

s_nodeinfo field	Definition	
node_type	A constant which represents the data type of the system task/function argument. The <i>node_type</i> determines which fields of the <i>s_tfexprinfo</i> structure will be used. The constants are:	
	TF_NULL_NODE	arg is null or is not a Verilog data type (e.g.: a number, expression or string)
	TF_REG_NODE	arg is a Verilog scalar or vector <i>reg</i> data type
	TF_INTEGER_NODE	arg is a Verilog <i>integer</i> data type
	TF_TIME_NODE	arg is a Verilog <i>time</i> data type
	TF_REAL_NODE	arg is a Verilog <i>real</i> data type
	TF_NETVECTOR_NODE	arg is a Verilog vector net data type
	TF_NETSCALAR_NODE	arg is a Verilog scalar net data type
padding	arg is a word select of a one-dimensional array of a <i>reg</i> , <i>integer</i> or <i>time</i> types	
	not used	
node_value	A union of C data types that point to the logic value of the system task/function argument. The field within the union which points to the value is controlled by <i>node_type</i> :	
	<p>If <i>node_type</i> is TF_REG_NODE, TF_INTEGER_NODE, TF_TIME_NODE or TF_NETVECTOR_NODE, then <i>node_value.vecval_p</i> contains a pointer to an array of one or more <i>s_vecval</i> structures with the 4-state logic value</p>	
	<p>If <i>node_type</i> is TF_NETSCALAR_NODE, then <i>node_value.strengthval_p</i> contains a pointer to one <i>s_strengthval</i> structure containing the logic strength</p>	
	<p>If <i>node_type</i> is TF_MEMORY_NODE, then <i>node_value.memoryval_p</i> contains a pointer to a character array containing the 4-state logic value of the entire array</p>	
node_symbol	a pointer to a string containing the name of the signal in the system task/function argument	

Table 11-4: The *tf_nodeinfo* structure

s_nodeinfo field	Definition
node_ngroups	If <i>node_type</i> is TF_REG_NODE, TF_INTEGER_NODE, TF_TIME_NODE or TF_NETVECTOR_NODE, this field contains the number of elements in the array of s_vecval structures pointed to by the <i>node_value.vecval_p</i> field If <i>node_type</i> is TF_MEMORY_NODE, then this field contains the number of characters in the character array pointed to by the <i>node_value.memoryval_p</i> field that represent one word in the Verilog memory array
node_vec_size	the number of bits in the system task/function argument
node_sign	a flag indicating the sign type of the task/function argument, 0 indicates an unsigned value, and 1 indicates a signed value
node_ms_index node_ls_index	if <i>expr_type</i> is TF_REG_NODE or TF_NETVECTOR_NODE, then these fields contain the bit number of the most significant bit and the least significant bit of the vector declaration.
node_mem_size	if <i>expr_type</i> is TF_MEMORY_NODE, then this field contains the number of addresses in the Verilog memory array.
node_lhs_element node_rhs_element	if <i>expr_type</i> is TF_REG_NODE or TF_NETVECTOR_NODE, and the argument is a part select of a vector, then these fields contain the bit number of the most significant bit and the least significant bit of the part select. Note: some simulators do not use these fields. PLI applications should not depend on these fields, because the application may not be portable to all simulators.
node_handle	not used. Note: some simulators may store a pointer to the task/function argument signal in this field, but a PLI application should not rely on this data, as it may not be the same in all simulators.

Table 11-4: The tf_nodeinfo structure (continued)

To read a task/function argument's 4-state logic value using `tf_nodeinfo()`, use the following basic steps:

1. Allocate an `s_tfnodeinfo` structure.
2. Call `tf_nodeinfo()` with a pointer to the `s_tfnodeinfo` structure.
3. Read the logic value from the `s_tfnodeinfo` structure. Use the `node_type` to determine what field of the `node_value` union contains the logic value, as described in table 11-4, above.

Reading vector and reg values with tf_nodeinfo()

Reading vector values and scalar `reg` values using `tf_nodeinfo()` is the same as with `tf_exprinfo()`. The logic value is retrieved into an array of `s_vecval` structures containing `aval/bval` pairs to encode the Verilog 4-state logic. The `node_ngroups` field indicates how many elements are in the array. The details on using `s_vecval` structures was presented earlier in this chapter, on page 349.

Reading strength values with tf_nodeinfo()

The logic value and strength value of scalar nets and bit selects of vector nets is represented as a single `s_strengthval` structure. The definition of this structure, as contained in the veriuser.h file, is:


```
typedef struct t_strengthval
{
    int strength0;
    int strength1;
} s_strengthval, *p_strengthval;
```

The Verilog language has 8 strength levels for a logic zero, and 8 strength levels for a logic one. Each strength level is represented by a Verilog keyword, as follows:

Strength Level	Strength Name	Specification Keyword
7	Supply Drive	<code>supply0</code> <code>supply1</code>
6	Strong Drive	<code>strong0</code> <code>strong1</code>
5	Pull Drive	<code>pull0</code> <code>pull1</code>
4	Large Capacitance	<code>large</code>
3	Weak Drive	<code>weak0</code> <code>weak1</code>
2	Medium Capacitance	<code>medium</code>
1	Small Capacitance	<code>small</code>
0	High Impedance	<code>highz0</code> <code>highz1</code>

Table 11-5: Verilog HDL strength levels and keywords

Within Verilog, the strength of a signal is stored as two 8-bit bytes, as shown in the diagram below:



The full details on how the two strength bytes represent Verilog logic strength are defined in the Verilog language, and are outside the scope of this book. Refer to the list of Verilog HDL books on page 7 for suggestions on where to find more information on Verilog HDL strength levels.

In the PLI, the strength value returned by `tf_nodeinfo()` is represented by a pair of C integers with a value from 00 (hex) to FF (hex). The value indicates which bits of the corresponding Verilog strength byte are set. Only 8 bits of each C integer are used. Note that both integers represent bit 0, the highz bit, as the least significant bit of the C integer.

Example 11-5 illustrates how to read the 4-state logic values and strength values of various Verilog HDL data types, using `tf_nodeinfo()`. This example does not read the values of a Verilog memory array. Reading memory values is presented in section 11.8 on page 367.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.011/nodeinfo_test_tf.c
- Verilog test bench: Chapter.011/nodeinfo_test.v
- Verilog-XL results log: Chapter.011/nodeinfo_test.log

Example 11-5: \$nodeinfo_test —using the tf_nodeinfo() routine

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
/*********************************************
 * checktf application
 *****/
int PLIbook_nodeinfoTest_checktf()
{
    if (tf_nump() != 1)
        tf_error("Usage error: $nodeinfo_test(<signal>);");
    return(0);
}

/*********************************************
 * calltf application
 *****/
int PLIbook_nodeinfoTest_calltf()
{
    s_tfnodeinfo    node_info;
    int             i;

    /* Get the nodeinfo structure for tfarg 1 */
    io_printf("Reading Node Information\n");
    if (!tf_nodeinfo(1, &node_info)) {
        tf_error("Error getting tf_nodeinfo for tfarg 1");
        return(0);
    }
    io_printf(" node_symbol is %s\n",
              node_info.node_symbol? node_info.node_symbol: "No Symbol";
    switch (node_info.node_type) {
        case TF_NULL_NODE:
            io_printf(" node_type = TF_NULL_NODE\n"); break;
        case TF_REG_NODE:
            io_printf(" node_type = TF_REG_NODE\n"); break;
        case TF_INTEGER_NODE:
            io_printf(" node_type = TF_INTEGER_NODE\n"); break;
        case TF_TIME_NODE:
            io_printf(" node_type = TF_TIME_NODE\n"); break;
        case TF_REAL_NODE:
            io_printf(" node_type = TF_REAL_NODE\n"); break;
        case TF_NETSCALAR_NODE:
            io_printf(" node_type = TF_NETSCALAR_NODE\n"); break;
        case TF_NETVECTOR_NODE:
            io_printf(" node_type = TF_NETVECTOR_NODE\n"); break;
        case TF_MEMORY_NODE:
            io_printf(" node_type = TF_MEMORY_NODE\n"); break;
        default:
            io_printf(" node_type = unknown (%d)\n\n", node_info.node_type);
    }
    io_printf(" node_ngroups = %d\n",      node_info.node_ngroups);
    io_printf(" node_vec_size = %d\n",     node_info.node_vec_size);
    io_printf(" node_sign = %d\n",         node_info.node_sign);
    io_printf(" node_ms_index = %d\n",     node_info.node_ms_index);
```

```
io_printf(" node_ls_index = %d\n",      node_info.node_ls_index);
io_printf(" node_mem_size = %d\n",       node_info.node_mem_size);
io_printf(" node_lhs_element = %d\n",    node_info.node_lhs_element);
io_printf(" node_rhs_element = %d\n",    node_info.node_rhs_element);

switch (node_info.node_type) {
    case TF_REG_NODE:
        io_printf(" reg value (in hex):\n");
        for (i=0; i<node_info.node_ngroups; i++) {
            io_printf("    avalbits[%d] = %x\n",
                      i, node_info.node_value.vecval_p[i].avalbits);
            io_printf("    bvalbits[%d] = %x\n",
                      i, node_info.node_value.vecval_p[i].bvalbits);
        }
        break;
    case TF_INTEGER_NODE:
        io_printf(" integer value (in hex):\n");
        io_printf("    avalbits[0] = %x\n",
                  node_info.node_value.vecval_p[0].avalbits);
        io_printf("    bvalbits[0] = %x\n",
                  node_info.node_value.vecval_p[0].bvalbits);
        break;
    case TF_TIME_NODE:
        io_printf(" time value (in hex):\n");
        io_printf("    {avalbits[1],avalbits[0]} = %x%x\n",
                  node_info.node_value.vecval_p[1].avalbits,
                  node_info.node_value.vecval_p[0].avalbits);
        io_printf("    {bvalbits[1],bvalbits[0]} = %x%x\n",
                  node_info.node_value.vecval_p[1].bvalbits,
                  node_info.node_value.vecval_p[0].bvalbits);
        break;
    case TF_REAL_NODE:
        io_printf(" real value = %f\n",
                  *node_info.node_value.real_val_p);
        break;
    case TF_NETSCALAR_NODE:
        io_printf(" scalar net value with strength (in hex):\n");
        io_printf("    strength0 = %x\n",
                  node_info.node_value.strengthval_p->strength0);
        io_printf("    strength1 = %x\n",
                  node_info.node_value.strengthval_p->strength1);
        break;
    case TF_NETVECTOR_NODE:
        for (i=0; i<node_info.node_ngroups; i++) {
            io_printf("    avalbits[%d] = %x\n",
                      i, node_info.node_value.vecval_p[i].avalbits);
            io_printf("    bvalbits[%d] = %x\n",
                      i, node_info.node_value.vecval_p[i].bvalbits);
        }
        break;
    case TF_MEMORY_NODE:
        io_printf(" memory arrays are not supported in this example\n");
        break;
```

```
    }
    io_printf("\n");
    return(0);
}
```

11.8 Reading from and writing into Verilog memory arrays

The `tf_nodeinfo()` routine can be used to both read and modify the contents of Verilog memory arrays and variable arrays. In Verilog HDL terminology, a ***memory array*** is a one dimensional array of the Verilog `reg` data type (regardless of the reg vector width). A ***variable array*** is an array of Verilog `integer` or `time` data types. In the IEEE 1364 Verilog standard, the description of the `tf_nodeinfo()` routine does not distinguish between a memory array and a one-dimensional variable array. The routine refers to any one-dimensional array as a memory array.

NOTE → The proposed IEEE 1364-1999 Verilog standard will add multi-dimensional arrays to the Verilog language. The `tf_nodeinfo()` routine cannot access these arrays. Only the VPI routines will support multi-dimensional arrays.

TIP → When the `tf_nodeinfo()` routine is used to access the value of a Verilog array, the routine returns a pointer to the actual storage of the array in the simulation data structure. Once a PLI application has obtained this pointer, the values of the array can be both read and modified any number of times during simulation. It is not necessary to call `tf_nodeinfo()` each time access to the array is required. All that is necessary is to save the pointer to the array which was returned from the first call to `tf_nodeinfo()`.

`tf_nodeinfo()` requires that a word-select of the array be specified in the system task/function argument. For example:

```
reg [23:0] RAM [0:3]; //array with 24-bit words, 4 words deep
initial
$dump_mem_hex(RAM[2]);
```

Note that, although `tf_nodeinfo()` requires that a word select be specified, the routine ignores the value of the word select. The `tf_nodeinfo()` routine will always retrieve a pointer to the entire Verilog array, regardless of the word select which was specified (other routines, such as `tf_exprinfo()`, utilize the word select specified).

The logic values of the Verilog memory array are stored in an array of `char` characters, which is pointed to in the `node_value_p.memoryval_p` field. The C char

data type is used because it is 8 bits wide. This section shows how a Verilog memory array is mapped to a character array in the C language.

Verilog array declaration syntax:

The syntax for declaring a one-dimensional array in Verilog is:

```
reg [<msb>:<lsb>] <memory_name> [<first_addr>:<last_addr>];
integer <array_name> [<first_address>:<last_address>];
time <array_name> [<first_address>:<last_address>];
```

The `integer` and `time` variables in Verilog have predefined word sizes of 32 bits and 64 bits, respectively. The word size of a `reg` data type can be any vector width, from 1-bit wide (scalar) to 1-million bits wide (the upper limit is actually defined by the simulator, but is 1 million bits in most simulators). The bit numbering of a word for a Verilog `reg` data type can use any numbering scheme. The following declarations all declare a 24-bit word size, and a 4-element array size:

```
reg [23:0] RAM1 [0:3]; //least-significant bit is bit 0
reg [0:23] RAM1 [0:3]; //least-significant bit is bit 23
reg [24:1] RAM1 [0:3]; //least-significant bit is bit 1
```

When an `integer` variable is used in a array, it is essentially the same as a `reg` declaration, with the msb as bit 31, and the lsb as bit 0. When a `time` variable is used in a memory array, it is essentially the same as a `reg` declaration, with the msb as bit 63, and the lsb as bit 0.

The array address numbering in Verilog can start and end with any address numbers, and do not need to have an address 0. The following examples all declare an array with 24-bit wide words, and 4 elements in the array:

```
reg [23:0] RAM1 [0:3]; //first word is address 0
reg [23:0] RAM1 [3:0]; //first word is address 3
reg [23:0] RAM1 [1:4]; //first word is address 1
```

11.8.1 Mapping Verilog Memory arrays into a C character array

To represent a Verilog memory array into storage represented in the C language involves three levels of mapping:

- Verilog array addresses need to be mapped to C array addresses.
- Verilog bit numbers within a word need to be mapped to a C representation.
- Verilog 4-state logic needs to be mapped to a C representation.

Mapping Verilog array numbering to C array numbering

The `tf_nodeinfo()` routine makes the bit numbering and address numbering used in Verilog transparent to the PLI application. This is done by mapping C array index numbers to the same position in a Verilog array, rather than mapping a C index number to a Verilog index number. There are two mappings which occur:

- The order of array addresses in the Verilog array are mapped into C array addresses.

In the C language, arrays always begin with address 0. The `tf_nodeinfo()` routine maps a C array address 0 to the lowest word address in the Verilog array address. Note that, since Verilog can define array addressing in either ascending or descending order, the lowest word address number might be either the first address in the Verilog array, or it might be the last address in the Verilog array.

- The bytes within a Verilog array word are mapped to bytes in the C language.


The `tf_nodeinfo()` routine divides a Verilog array word into 8-bit bytes, which are mapped to a C `char` array. The first byte in the C array is always the least-significant byte (the right-most byte) of the Verilog word, regardless of how the bits are numbered in Verilog.

The following diagrams illustrate how a Verilog array is mapped to a C `char` array.

A conceptual view of a Verilog array declaration


A conceptual view of a Verilog `reg` or variable array is:

```
reg [msb:lsb] RAM [first_address:last_address];
```




Two example Verilog memory declarations are:

```
reg [23:0] RAM [0:3]; //lsb is lowest bit, ascending addresses
```




```
reg [0:23] RAM [3:0]; //1sb is highest bit, descending addresses
```




Mapping Verilog array addresses to C array addresses

The `tf_nodeinfo()` routine maps a C array address 0 to the lowest number in the Verilog array address. Since Verilog can define array address in either ascending or descending order, the lowest address number might be either the first address of the Verilog array or it might be the last address in the Verilog array. The following diagrams show how two different Verilog memory declarations are mapped to a C array:

```
reg [23:0] RAM [0:3]; //array with ascending address order
```



```
reg [23:0] RAM [3:0]; //array with descending address order
```




Mapping a Verilog array word into C bytes


The PLI represents each 8-bit byte of a Verilog array word as a *group*. Each group is stored as a C char data type, and the groups which make up a Verilog word are organized into a C char array. The least significant byte of a Verilog array word (which is the right-most byte) is stored in the first group in the C array.

The two diagrams which follow illustrates how two Verilog memory array declarations are represented as groups of 8-bits. *Note the reversed ordering that occurs in these diagrams:* In the Verilog HDL representation, the least-significant byte of a word is always the right most byte, but, in the PLI, the least-significant byte of the Verilog word becomes the first group in the C array. However, the bits within each byte remain in the same order as in the original Verilog word.

```
reg [23:0] RAM [0:3]; //array with LSB the lowest bit number
```



```
reg [0:23] RAM [0:3]; //array with LSB the highest bit number
```



Mapping 4-state values into aval/bval pairs



Each bit in the Verilog array is represented by an aval/bval pair of bits, in order to encode the 4-state logic of Verilog. The encoding is listed in Table 11-6, and is the same encoding as used by `tf_exprinfo()` and other PLI routines which encode 4-state logic.

aval/bval pair	Verilog logic value represented
0/0	0
1/0	1
0/1	Z
1/1	X

Table 11-6: aval/bval logic value encoding


A Verilog array word is stored as a set of two 8-bit groups, with one set containing the aval values and the other set containing the bval values. Conceptually, the two sets of groups for a Verilog memory array can be thought of as follows:

```
reg [23:0] RAM [0:3]; // lsb is lowest bit, ascending addresses
```



The conceptual view of a Verilog memory array shown above is represented as groups of bytes. The PLI stores the Verilog array using a C `char` data type for each group, and uses a character array to store the entire Verilog memory. The aval and bval groups of each memory word are contained in the same array, as follows:

```
reg [23:0] RAM [0:3]; //lsb is lowest bit, ascending addresses
```



11.8.2 Verilog array information that is accessed by `tf_nodeinfo()`

After calling `tf_nodeinfo()` for a system task/function argument that contains a Verilog array word, the `s_tfnodeinfo` structure will contain the following information about the C character array which represents the Verilog array:

- `node_value.memoryval_p` contains a pointer to the start of the character array.
- `node_mem_size` contains the number of words represented by the Verilog array.
- `node_vec_size` contains the number of bits in each Verilog array word.
- `node_ngroups` contains the number of groups in each Verilog array word.

Using the pointer to the C character array, and knowing how many groups make up an array word, a PLI application can read the value of any word or discrete bit within the Verilog array.

Accessing one word of a Verilog array

A full word from a Verilog array can be accessed using the following formula:

Given the following call to `tf_nodeinfo()`:

```
s_tfnodeinfo node_info;
tf_nodeinfo(1, &node_info); /* get info for tfarg 1 */
```

- The beginning of the character array and number of groups in a word are stored in the `s_tfnodeinfo` structure:

```
node_info.node_value.memoryval_p;
node_info.node_ngroups;
```

- Each Verilog array word is represented as an aval/bval pair of 8-bit groups, with two group sets required to represent one Verilog word. The number of groups in the character array which represent one Verilog array word can be calculated as:

```
int word_increment;
word_increment = node_info.node_ngroups * 2;
```


- To access a specific Verilog array word in the C character array, use the formula (assuming the desired word is stored in the variable `memory_address`):

```
char *aval_ptr, *bval_ptr;
aval_ptr = node_info.node_value.memoryval_p
           + (word_increment * memory_address);
bval_ptr = aval_ptr + node_info.node_ngroups;
```

The following example will access the second word (address 1 in this example) of a RAM memory with 24 bit word widths.

```
reg [23:0] RAM [0:3]; //lsb is lowest bit, ascending addresses

s_tfnodeinfo node_info;
tf_node_info(1, &node_info);
```



An example of reading values from a Verilog memory array

Example 11-6 illustrates two PLI applications, `$dump_mem_hex` and `$dump_mem_ascii`. The first application prints the aval/bval pair for each word in a memory array, in hexadecimal. The second application prints the value of just the aval bytes for each word in a memory array, in ASCII. Both applications call the same `checktf routine`, `misctf routine` and `calltf routine`. The `user_data` value associated with the PLI application is used to determine if the application routines were called by `$dump_mem_hex` (a `user_data` of 0) or `$dump_mem_ascii` (a `user_data` of 2).



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.011/dump_memory_tf.c
- Verilog test bench: Chapter.011/dump_and_fill_mem_test.v
- Verilog-XL results log: Chapter.011/dump_and_fill_mem_test.log

Example 11-6: `$dump_mem_??` — using `tf_nodeinfo()` with Verilog memories

```
#include "veriuser.h"          /* IEEE 1364 PLI TF  routine library */

#define HEX    0 /* values of user_data for system task names */
#define ASCII 2

/*********************  
* checktf routine  
*******************/  
int PLIbook_DumpMem_checktf(int user_data)  
{  
    if (tf_numP() != 1)  
        if (user_data == HEX)  
            tf_error("Usage error: $dump_mem_hex(<memory_word_select>);");  
        else  
            tf_error("Usage error: $dump_mem_ascii(<memory_word_select>);");  
    return(0);  
}  
  
/*********************  
* misctf routine  
*  
* The misctf application is used to call tf_nodeinfo() at the  
* beginning of simulation, so that the memory allocated by  
* tf_nodeinfo() is only allocated one time for each instance of  
* $dump_mem_?.  
*******************/  
int PLIbook_DumpMem_misctf(int user_data, int reason)  
{  
    p_tfnodeinfo node_info; /* pointer to structure for tf_nodeinfo() */
```

```
if (reason != REASON_ENDOFCOMPILE)
    return(0); /* exit now if this is not the start of simulation */

/* allocate memory for an s_tfnodeinfo structure */
node_info = (p_tfnodeinfo)malloc(sizeof(s_tfnodeinfo));

/* Get the nodeinfo structure for tfarg 1 */
if (!tf_nodeinfo(1, node_info)) {
    tf_error("Err: $dump_mem_?? could not get tf_nodeinfo for tfarg 1");
    tf_dofinish(); /* about simulation */
    return(0);
}
else if (node_info->node_type != TF_MEMORY_NODE) {
    tf_error("Err: $dump_mem_?? arg is not a memory word -- aborting");
    tf_dofinish(); /* about simulation */
    return(0);
}
else
    tf_setworkarea((char *)node_info); /*put info pointer in workarea*/

return(0);
}

*****
* calltf application
*****
/* prototypes of functions invoked by the calltf routine */
void PLIbook_DumpMemHex();
void PLIbook_DumpMemAscii();

int PLIbook_DumpMem_calltf(int user_data)
{
    p_tfnodeinfo node_info;

    node_info = (p_tfnodeinfo)tf_getworkarea();

    io_printf("\nWithin PLI:\n");
    io_printf(" Memory array width=%d depth=%d ngroups=%d\n",
            node_info->node_vec_size,
            node_info->node_mem_size,
            node_info->node_ngroups);

    if (user_data == HEX)      /* application called by $dump_mem_hex */
        PLIbook_DumpMemHex(node_info);
    else                      /* application called by $dump_mem_ascii */
        PLIbook_DumpMemAscii(node_info);

    return(0);
}
```

```
*****
* Function to dump each word of a Verilog array in hexadecimal
*****
void PLIbook_DumpMemHex(p_tfnodeinfo node_info)
{
    char *aval_ptr, *bval_ptr;
    int    word_increment, mem_address, group_num;

    io_printf(" Current memory contents of aval/bval groups in hex:\n");

    word_increment = node_info->node_ngroups * 2;
    for (mem_address = 0;
         mem_address < node_info->node_mem_size;
         mem_address++) {
        io_printf("    C array address %d:", mem_address);

        /* set pointers to aval and bval words for the address */
        aval_ptr = node_info->node_value.memoryval_p
                   + (mem_address * word_increment);
        bval_ptr = aval_ptr + node_info->node_ngroups;

        /* print groups in word in reverse order so will match Verilog:
           the highest group number represents the left-most byte of a
           Verilog word, the lowest group represents the right-most byte */
        for (group_num = node_info->node_ngroups - 1;
             group_num >= 0;
             group_num--) {
            io_printf("    group %d: %x/%x",
                      group_num, aval_ptr[group_num], bval_ptr[group_num]);
        }
        io_printf("\n");
    }
    io_printf("\n\n");
    return;
}

*****
* Function to dump each word of a Verilog array in ASCII
*****
void PLIbook_DumpMemAscii(p_tfnodeinfo node_info)
{
    char *aval_ptr;
    int    word_increment, mem_address, group_num;

    /* Read current memory values as a string using only aval bits */
    io_printf(" Current memory contents in ASCII are:\n");
    io_printf("  ");

    word_increment = node_info->node_ngroups * 2;
    for (mem_address = 0;
         mem_address < node_info->node_mem_size;
         mem_address++) {
        /* set pointer to aval word for the address */
        aval_ptr = node_info->node_value.memoryval_p
                   + (mem_address * word_increment);
```

```

/* print groups in word in reverse order so will match Verilog:
   the highest group number represents the left-most byte of a
   Verilog word, the lowest group represents the right-most byte */
for (group_num = node_info->node_ngroups - 1;
     group_num >= 0;
     group_num--) {
    io_printf("%c", aval_ptr[group_num]);
}
io_printf("\n\n");
return;
}

```

11.8.3 Modifying values in Verilog memory and variable arrays

The `tf_nodeinfo()` routine retrieves a pointer to the actual storage within simulation of a one-dimensional Verilog `reg` or variable array. A PLI application can use this pointer to modify the values of the `aval/bval` pairs which represent the array.

NOTE *When a PLI application modifies the value of a Verilog array, the new value does not automatically propagate to Verilog statements which read the Verilog array.*

Once a PLI application has modified an array value, the new value is available the next time a Verilog statement reads a word from the array. But, the changes in the Verilog array made by a PLI application do not automatically propagate to Verilog statements which are currently reading the Verilog array. The following two Verilog source code examples illustrate this difference:

Assuming an array declaration of:

```
reg [23:0] RAM [0:3]; //Verilog memory array
```

The following Verilog continuous assignment statement, which reads a word from the Verilog array, will not see a value change caused by a PLI application.

```
assign vector = RAM[1]; //continuously read address 1
```

The Verilog procedure listed next, which reads a word from the Verilog array at a positive edge of clock, will see the value change caused by a PLI application, the next time the statement is executed at a positive edge of clock.

```
always @(posedge clock)
  vector = RAM[1]; //read address 1 each posedge of clock
```

tf_propagatetp() does not work with tf_nodeinfo()

NOTE → The description of `tf_propagatetp()` in the IEEE 1364-1995 Verilog standard contains a serious errata. The proposed IEEE 1364-1999 standard corrects this error.

The description in the IEEE 1364-1995 Verilog standard for the `tf_propagatetp()` routine suggests that the routine can be used to propagate memory array contents to all Verilog HDL constructs which read values from the array. *This is an error*. The `tf_propagatetp()` routine is intended to only propagate the value contained in single word of an array (the word specified in the system task/function argument). The value to be propagated is the value pointed to in the `s_tfexprinfo` structure. The `tf_exprinfo()` routine must be used to setup the `s_tfexprinfo` structure. The value of the memory array word to be propagated is *not* the value pointed to in the `s_tfnodeinfo` structure.

NOTE → In some Verilog simulators, `tf_propagatetp()` may cause some memory word values to propagate within simulation, other than the word contained in the `s_tfexprinfo` structure. This behavior is not part of the Verilog standard, and PLI applications should not depend on changes to memory array contents propagating within the Verilog simulation.

An example of writing values into a Verilog memory

Example 11-7 lists the C source code for a PLI application called `$fill_mem`, which fills every address of a Verilog array with the C language word index of each word in the array (that is, the first address of the Verilog array is loaded with the value of 0, the second address with the value of 1, etc.).

This example uses both `tf_exprinfo()` and `tf_nodeinfo()`, in order to illustrate the difference in using these routines.

- The `tf_exprinfo()` routine modifies the value of just the memory word which is past in as a system task/function argument. The `tf_propagatetp()` routine can be used in conjunction with `tf_exprinfo()` to cause the new value to propagate within simulation.
- The `tf_nodeinfo()` routine modifies the value of any memory location. The new values do not propagate within simulation. However, the new values are in the memory, and will be seen the next time a Verilog statement reads from the array.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.011/fill_memory_tf.c
- Verilog test bench: Chapter.011/dump_and_fill_mem_test.v
- Verilog-XL results log: Chapter.011/dump_and_fill_mem_test.log

Example 11-7: \$fill_mem — using tf_nodeinfo() to modify Verilog memories

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
/********************* checktf application ********************/
int PLIbook_FillMem_checktf()
{
    if (tf_numP() != 2)
        tf_error("Usage: $fill_mem(mem_word_select,word_select_address);");
    return(0);
}

/********************* misctf application ********************/
* The misctf application is used to call tf_nodeinfo() at the
* beginning of simulation, so that the memory allocated by
* tf_nodeinfo() and tf_exprinfo() is only allocated one time for each
* instance of $fill_mem.
/********************* PLIbook_my_data ********************/
typedef struct PLIbook_my_data {
    p_tfnodeinfo node_info; /* pointer to structure for tf_nodeinfo() */
    p_tfexprinfo expr_info; /* pointer to structure for tf_exprinfo() */
} PLIbook_my_data_s, *PLIbook_my_data_p;

int PLIbook_FillMem_misctf(int user_data, int reason)
{
    PLIbook_my_data_p info; /* pointer to info structures */

    if (reason != REASON_ENDOFCOMPILE)
        return(0); /* exit now if this is not the start of simulation */

    /* allocate memory for structure to store info structure */
    info = (PLIbook_my_data_p)malloc(sizeof(PLIbook_my_data_s));

    /* allocate memory for an s_nodeinfo and an s_tfexprinfo structure */
    info->node_info = (p_tfnodeinfo)malloc(sizeof(s_tfnodeinfo));
    info->expr_info = (p_tfexprinfo)malloc(sizeof(s_tfexprinfo));

    /* Get the nodeinfo structure for tfarg 1 */
    if (!tf_nodeinfo(1, info->node_info))
        tf_error("Error: $fill_mem could not get tf_nodeinfo for tfarg 1");
```

```
    tf_dofinish(); /* about simulation */
    return(0);
}
else if (info->node_info->node_type != TF_MEMORY_NODE) {
    tf_error("Error: $fill_mem arg is not a memory word -- aborting");
    tf_dofinish(); /* about simulation */
    return(0);
}

/* Get the exprinfo structure for tfarg 1 */
if (!tf_exprinfo(1, info->expr_info)) {
    tf_error("Error: $fill_mem could not get tf_exprinfo for tfarg 1");
    tf_dofinish(); /* about simulation */
    return(0);
}

tf_setworkarea((char *)info); /* put info pointer in work area */

return(0);
}

/*********************************************
 * calltf application
 *****/
int PLIbook_FillMem_calltf()
{
    int depth, width, ngroups, word_increment, mem_address, i;
    char *mem_ptr, *aval_ptr, *bval_ptr;
    PLIbook_my_data_p info; /* pointer to info structures */

    info = (PLIbook_my_data_p)tf_getworkarea();

    mem_ptr      = info->node_info->node_value.memoryval_p;
    width        = info->node_info->node_vec_size;
    depth        = info->node_info->node_mem_size;
    ngroups      = info->node_info->node_ngroups;

    /* Modify current memory values: set aval bits to memory address,
       set bval bits to 0 (2-state logic) */
    word_increment = ngroups * 2; /* 1 word = aval/bval group set */
    for (mem_address = 0;
         mem_address < 4; /* node_info->node_mem_size; */
         mem_address++) {
        aval_ptr = mem_ptr + (mem_address * word_increment);
        bval_ptr = aval_ptr + ngroups;
        aval_ptr[0] = mem_address;
        bval_ptr[0] = 0x0;
        for (i=1; i<ngroups; i++) {
            aval_ptr[i] = 0x0;
            bval_ptr[i] = 0x0;
        }
    }
}
```

```

mem_address = tf_getp(2); /* get address to propagate changes */
tf_evaluatep(1);          /* update expr_info structure */
info->expr_info->expr_value_p[mem_address].avalbits = mem_address;
info->expr_info->expr_value_p[mem_address].bvalbits = 0x0;

tf_propagatep(1); /* propagate memory changes into simulation */

return(0);
}

```

11.8.4 Accessing one bit within a Verilog array word

A PLI application can access any discrete bit within a Verilog `reg` or variable array. This section shows two C coding methods which can be used to access the individual bits within a Verilog array.

- The first method uses a more obvious C coding style. This example is provided, in order to show how each bit is accessed.
- The second method uses a less intuitive, but more efficient C coding style.

A more obvious, but less efficient method

To access a discrete bit within a Verilog array word requires four basic steps:

1. Select a word from the C `char` array which represents the Verilog array word to be modified, using the same method shown previously for reading and modifying words of Verilog `reg` and variable arrays (see page 373).

```

char *aval_ptr, *bval_ptr;
aval_ptr = node_info.node_value.memoryval_p
           + (word_increment * memory_address);
bval_ptr = aval_ptr + node_info.node_ngroups;

```

2. Determine which 8-bit group in the C `char` array contains the desired bit in the word.

```

int group_num;
group_num = bit_num / 8;

```

The desired bit number is divided by 8, because there are 8 bits in a group.

3. Determine which bit in the group contains the desired bit in the word.

```
int group_bit;
group_bit = bit_num % 8;
```

The desired bit number is divided by 8 using the modulus operator, which returns the remainder of the division. The value of 8 is used because there are 8 bits in a group.

4. Set a mask to block out all unwanted bits in the group.

```
char mask;
mask = 0x01;
mask = mask << group_bit;
```

The mask is set to a hex value of 1, which sets the least-significant bit of the 8-bit byte. This bit is then shifted left to the desired bit within the group.

5. Read the value of the desired bit.

```
char aval_bit_value, bval_bit_value;
aval_bit_value = aval_ptr[group_num] & mask;
bval_bit_value = bval_ptr[group_num] & mask;
```

The value of the desired bit is accessed by logically *anding* the value of the group containing the bit with the mask value. Since only the desired bit in the mask is set, all other bits will be cleared. The resulting value of the aval/bval pair for the desired bit can then be mapped to the 4-state Verilog value.

The following C function shows how discrete bits of a memory word can be accessed using this more obvious, less efficient, C coding style. This C function is an excerpt from a complete example called *\$dump_mem_bin*.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.011/dump_memory_tf.c
- Verilog test bench: Chapter.011/dump_and_fill_mem_test.v
- Verilog-XL results log: Chapter.011/dump_and_fill_mem_test.log

Example 11-8: \$dump_mem_bin — accessing a memory word bit; obvious method

```
void PLIbook_DumpMemBin(p_tfnodeinfo node_info)
{
    char *aval_ptr, *bval_ptr;
    int word_increment, mem_address, word_bit, group_num, group_bit;
    char aval_val, bval_val, bit_mask;

    io_printf(" Current memory contents in binary are:\n");
    word_increment = node_info->node_ngroups * 2;

    for (mem_address = 0;
         mem_address < node_info->node_mem_size;
         mem_address++) {

        /* step 1: set pointers to aval and bval words for the address */
        aval_ptr = node_info->node_value.memoryval_p
                   + (mem_address * word_increment);
        bval_ptr = aval_ptr + node_info->node_ngroups;
        for (word_bit = node_info->node_vec_size - 1;
             word_bit >= 0;
             word_bit--) {

            /* step 2: determine the group which contains the bit number */
            group_num = word_bit / 8;

            /* step 3: determine which bit in the group contains the bit */
            group_bit = word_bit % 8;

            /* step 4: set an 8-bit mask to block all unwanted bits in group */
            bit_mask = 0x01; /* Set mask to most-signif. bit of 8-bit group */
            bit_mask = bit_mask << group_bit; /* Shift to bit to be modified */

            /* step 5: select desired aval and bval bits from the groups */
            aval_val = aval_ptr[group_num] & bit_mask;
            bval_val = bval_ptr[group_num] & bit_mask;

            /* translate aval/bval pair to 4-state logic value */
            if (!bval_val) {
                if (!aval_val) io_printf("0"); /* aval/bval == 0/0 */
                else          io_printf("1"); /* aval/bval == 1/0 */
            }
            else {
                if (!aval_val) io_printf("z"); /* aval/bval == 0/1 */
                else          io_printf("x"); /* aval/bval == 1/1 */
            }
        }
        io_printf("\n");
    }
    io_printf("\n");
    return;
}
```

A more efficient, but less obvious method

Drew Lynch, of Surefire Verification, who reviewed much of this book, suggests the following, more efficient C coding method to access the bits of a Verilog array word. This method uses a combination of the C shift and logical-and operators to select a single bit from a memory word, instead of the mathematical operators and intermediate variables used in the previous example.

```
aval_ptr = node_info->node_value.memoryval_p
           + (mem_address * word_increment);
bval_ptr = aval_ptr + node_info->node_value.ngroups;
aval_bit_value =
    (aval_ptr[word_bit >> 3]) & (1 << (word_bit & 0x7));
bval_bit_value =
    (bval_ptr[word_bit >> 3]) & (1 << (word_bit & 0x7));
 $\brace{}$  select group within word  $\brace{}$  select bit within group
```

The first operation in this expression, `(bval_ptr[word_bit >> 3])`, selects a specific 8-bit group from the aval or bval bits of a Verilog array word. Shifting the bit-select value right three times is equivalent to dividing the bit-select value by 8, using integer division. The word is divided by 8, because there are 8-bits in a group.

The second operation, `(1 << (word_bit & 0x7))`, selects a specific bit from an 8-bit group. By logically-anding the desired bit-select with a value of hex 7, a value of 0 through 7 is derived. This is equivalent to the modulus operation of `(word_bit % 8)`. A literal value of 1 is then shifted left by the result of this modulus operation. This shift will select a specific bit within an 8-bit group.

The final operation in this expression is to logically-and the selected bit within the selected group. This operation will return 1 if that bit is set, and 0 if that bit is not set.

The following C function shows how a discrete bit of a memory word can be accessed using this more efficient C coding style. This C function can replace the less efficient example listed in Example 11-8 on page 384.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.011/dump_memory_tf.c
- Verilog test bench: Chapter.011/dump_and_fill_mem_test.v
- Verilog-XL results log: Chapter.011/dump_and_fill_mem_test.log

Example 11-9: \$dump_mem_bin — accessing a memory word bit; efficient method

```
void PLIbook_DumpMemBin(p_tfnodeinfo node_info)
{
    char *aval_ptr, *bval_ptr;
    int   word_increment, mem_address, word_bit;
    char  aval_val, bval_val;

    io_printf(" Current memory contents in binary are:\n");

    word_increment = node_info->node_ngroups * 2; /* 1 word = aval/bval
pair */
    for (mem_address = 0;
         mem_address < node_info->node_mem_size;
         mem_address++) {
        /* set pointers to aval and bval words for the address */
        aval_ptr = node_info->node_value.memoryval_p
                   + (mem_address * word_increment);
        bval_ptr = aval_ptr + node_info->node_ngroups;

        /* print groups in word in reverse order so will match Verilog:
           the highest group number represents the left-most byte of a
           Verilog word, the lowest group represents the right-most byte */
        for (word_bit = node_info->node_vec_size - 1;
             word_bit >= 0;
             word_bit--) {
            aval_val = (aval_ptr[word_bit >> 3]) & (1 << (word_bit & 0x7));
            bval_val = (bval_ptr[word_bit >> 3]) & (1 << (word_bit & 0x7));

            /* translate aval/bval pair to 4-state logic value */
            if (!bval_val) {
                if (!aval_val)  io_printf("0"); /* aval/bval == 0/0 */
                else          io_printf("1"); /* aval/bval == 1/0 */
            }
            else {
                if (!aval_val)  io_printf("z"); /* aval/bval == 0/1 */
                else          io_printf("x"); /* aval/bval == 1/1 */
            }
        }
        io_printf("\n");
    }
    return;
}
```

Modifying the bits of a Verilog array

A bit of a Verilog `reg` or variable array can be modified by setting the `aval` and `bval` values for the desired bit. The pointer to the desired bit is obtained in the same manner as when reading a bit value. The PLI stores the value of a Verilog array word as 8-bit groups. Therefore, an 8-bit mask value can be logically and'ed or or'ed with the group, to set or clear a specific bit.

The following steps can be used to set a 4-state logic value in the `aval/bval` group pair:

1. Select a word from the C `char` array which represents the Verilog array word to be modified, using the same method shown previously for reading and modifying words of Verilog `reg` and variable arrays.

```
char *aval_ptr, *bval_ptr;
aval_ptr = node_info.node_value.memoryval_p
           + (word_increment * memory_address);
bval_ptr = aval_ptr + node_info.node_ngroups;
```

2. Determine which 8-bit group in the C `char` array contains the desired bit in the Verilog word.

```
int group_num;
group_num = bit_num / 8;
```

3. Determine which bit in the 8-bit group contains the desired bit in the word.

```
int group_bit;
group_bit = bit_num % 8;
```

4. Set a mask to block out all unwanted bits in the group.

```
char mask;
mask = 0x01;
mask = mask << group_bit;
```

The mask is set to a hex value of 1, which sets the least-significant bit of the 8-bit character. The least-significant bit is then shifted left to the desired bit within the group.

5. Modify desired `aval` and `bval` bits from the groups. To clear a bit: logically AND group with the inverse of the mask . To set a bit, logically OR group with the mask. For example:

```
switch (bit_val) {
    case '0': aval_ptr[group_num] = aval_ptr[group_num] & ~bit_mask;
                bval_ptr[group_num] = bval_ptr[group_num] & ~bit_mask;
                break;
    case '1': aval_ptr[group_num] = aval_ptr[group_num] | bit_mask;
                bval_ptr[group_num] = bval_ptr[group_num] & ~bit_mask;
                break;
    case 'z': aval_ptr[group_num] = aval_ptr[group_num] & ~bit_mask;
                bval_ptr[group_num] = bval_ptr[group_num] | bit_mask;
                break;
    case 'x': aval_ptr[group_num] = aval_ptr[group_num] | bit_mask;
                bval_ptr[group_num] = bval_ptr[group_num] | bit_mask;
                break;
}
```



TIP The preceding steps use a less efficient coding style, in order to explicitly show the steps required to modify a specific bit of a Verilog array. A more efficient method of modifying a specific bit would be to use logical operators in a compound expression, instead of using mathematical operators and intermediate variables. The concepts shown in example 11-9 on page 386 can be applied to this example, in order to make the code more efficient.

11.9 Summary

This chapter has presented the routines contained in the PLI TF library, which are used to read and write the logic values of system task and system function arguments, and to specify the return value of a system function. The Verilog HDL provides a number of logic value constructs which do not have a direct counterpart in the C language, such as 4-state logic and vectors with arbitrary bit widths. The TF routines provide several options for converting values from Verilog to C and from C to Verilog.

The next chapter is a continuation of the discussion of the TF library, and presents how to use the TF library with *misctf routines*.

CHAPTER 12

Synchronizing to Verilog Simulations Using Misctf Routines

The *misctf routine* offers several very useful features that can enhance the functionality of PLI applications. There are a number of TF routines which work specifically with the *misctf routine*. The concepts presented in this chapter are:

- The purpose of the *misctf routine*
- Automatic callbacks for simulation events
- Application scheduled callbacks at the end of a simulation time step
- Application scheduled callbacks at future simulation times
- Task/function argument value change callbacks
- Simulation save and restart callbacks

12.1 The purpose of the misctf routine

The PLI *misctf routine* is unique from the *calltf routine*, *checktf routine* and *sizetf routine*. The *misctf routine* is called for a number of miscellaneous reasons during simulation, which are not directly related to the simulation, having encountered or executed the user-defined system task/function name. As an example, if the Verilog HDL source code contained the user-defined system task:

```
always @(posedge clock)
$read_test_vector("vectors.pat", input_bus);
```

The *checktf routine* will be called when the compiler or simulation loader encounters the `$read_test_vector` task, and the *calltf routine* will be called at every positive edge of clock when the simulator executes the `$read_test_vector` task. The *misctf routine* will be called for miscellaneous *simulation events* that might occur during simulation, such as:

- Compilation or loading of the simulation is complete, and simulation is about to start running.
- The simulation executed a break point and entered interactive debug mode.
- The simulation finished and is exiting.
- The state of the simulation is being saved.
- Simulation is being restarted from a saved state.

In addition to the simulation events that will call a *misctf routine*, PLI applications can schedule callbacks to the *misctf routine* for other specific reasons. These reasons include:

- Simulation is finished processing the current simulation time step.
- Simulation has reached a specific future simulation time step.
- An argument of a user-defined system task has changed value.

12.1.1 The misctf reason input

Whenever a *misctf routine* is called, it is passed three C function inputs:

1. ***user_data***: The integer value that was specified when the user-defined task/function name was registered using the PLI interface mechanism. Refer to chapter 9, sections 9.9 and 9.11 for more details on how the *user_data* value.
2. ***reason***: An integer constant that is generated by the Verilog simulator to indicate why the simulator called the *misctf routine*. The reason constants are defined in the veriuser.h file. Example constants are REASON_ENDOFCOMPILE and REASON_FINISH. The complete list of reason constants is listed in tables 12-1 and 12-2, which follow.
3. ***paramvc***: An integer value which indicates the index number of an argument to a system task that changed value. The *paramvc* input is covered in section 12.5 on page 403.

The IEEE 1364 standard lists six reasons for which a *misctf routine* can be called, and which every Verilog simulator should support. The constants for these reasons are listed below. More detailed explanations and examples of using these callback reasons are presented in the remaining sections of this chapter.

Reason Constant	Description
REASON_ENDOFCOMPILE	end of Verilog source compilation/start of execution
REASON_PARAMVC	change of value on a user-defined system task/function argument
REASON_SYNCH	end of a time step flagged by <code>tf_synchronize()</code>
REASON_ROSYNCH	end of a time step flagged by <code>tf_rosynchronize()</code>
REASON.REACTIVATE	a simulation time step has been reached which was flagged using <code>tf_setdelay()</code>
REASON_FINISH	simulation is finished and is preparing to exit

Table 12-1: IEEE 1364 standard *misctf* callback reasons

In addition to the required callback reasons, the IEEE 1364 standard allows simulator vendors to add additional reasons that are features of that simulator. The standard suggests a number of optional misctf callback reasons, along with pre-defined constant names. These reasons may or may not be available in a specific simulator. Be sure to check the features of the simulator, to see if these optional reasons are implemented.

Reason Constant	Description
REASON_PARAMDRC	a strength change occurred on the driver of a user-defined system task/function argument
REASON_FORCE	a procedural <code>force</code> or procedural continuous <code>assign</code> statement was executed
REASON_RELEASE	a procedural <code>release</code> or procedural <code>deassign</code> was executed
REASON_DISABLE	a procedural <code>disable</code> statement was executed
REASON_INTERACTIVE	a simulation breakpoint was executed, such as the <code>\$stop</code> built-in system task
REASON_SCOPE	simulation changed interactive debug scope, such as by executing the <code>\$scope</code> built-in system task
REASON_STARTOFSAVE	simulation has started executing a checkpoint save, such as the <code>\$save</code> built-in system task
REASON_SAVE	simulation has finished executing a checkpoint save
REASON_RESTART	simulation has executed a checkpoint restart, such as the <code>\$restart</code> built-in system task

Table 12-2: IEEE 1364 optional *misctf* callback reasons

Reason Constant	Description
REASON_RESET	simulation has started executing a reset back to time 0, such as executing the \$reset built-in system task
REASON_ENDOFRESET	simulation has finished executing a reset to time 0

Table 12-2: IEEE 1364 optional *misctf* callback reasons (continued)

Example 12-1 illustrates using the reason input. *Notice that the reason the misctf routine is called is passed into the routine as the second C function input.*

Example 12-1: using the *misctf* reason input

```
int PLIbook_ReadVector_misctf(int data, int reason)
{
    FILE *in_file;
    char *file_name;
    if (reason == REASON_ENDOFCOMPILE) { /* time to open vector file */
        if ( (in_file = fopen(tf_getcstringp(1),"r")) == NULL)
            tf_error("$read_test_vector cannot open file %s",
                     tf_getcstringp(1));
        tf_setworkarea((char*)in_file); /* save file pointer in workarea */
    }
    if (reason == REASON_FINISH) { /* time to close vector file */
        in_file = (FILE*)tf_getworkarea(); /* retrieve file pointer */
        fclose(in_file);
    }
    return(0);
}
```

NOTE → The *misctf* routine is called *automatically* for a number of reasons, many of which may not be needed by a particular PLI application. It is imperative to check the reason input, and only perform operations for the desired reasons.

TIP → The *misctf* routine has an inherent performance inefficiency, because it is called automatically for reasons that a PLI application may not need. It is not required for a PLI application to specify a *misctf* routine, and the run-time execution of a simulation can be improved, if no *misctf* routine is specified. However, there are several features of the TF part of the PLI which require using the *misctf* routine. A PLI application developer must decide if the performance impact is worth the value of the extra capabilities.

12.2 Automatic callbacks for simulation events

The *misctf routine* will be called automatically for a number of miscellaneous simulation events. These callbacks to the *misctf routine* allow a PLI application to perform specific operations, or to process data at those times. In the `$read_test_vector` application that was presented in Example 10-3 on page 326 of Chapter 10, the *misctf routine* was used to open a file when the simulator finished compiling/loading (which is just before the simulation starts executing events at time zero).

The automatic simulation event callbacks that will occur in every Verilog simulator are `REASON_ENDOFCOMPILE` and `REASON_FINISH`. Most Verilog simulators also implement `REASON_INTERACTIVE` callbacks, though this callback is optional in the IEEE standard. Some simulators may also implement some or all of the other optional callback reasons listed in Table 12-2 on page 391.

The `REASON_ENDOFCOMPILE` *misctf* callback reason is poorly named. The callback does not occur at the end of compilation, but, instead, the callback occurs just before simulation time zero, when simulation is invoked and the simulation data files are loaded. A better term than end-of-compile would be start-of-simulation, indicating that the Verilog source code has been loaded, and simulation is about to start running.

There are a number of ways the `REASON_ENDOFCOMPILE` can be used. Generally, PLI applications take advantage of the fact that end-of-compile (really start-of-simulation) can only occur once in a simulation. This is a good time to perform any operations that need to be executed just once during the simulation, such as allocating memory for the PLI application, opening disk files, and opening graphics windows.



TIP

Use the TF work area to save memory pointers, file pointers and other information that needs to be preserved during simulation. See section 10.7 on page 313 of Chapter 10 for more information on the TF work area.

The `REASON_FINISH` callback indicates that the simulator is about to exit, which is generally caused by a `$finish` built-in system task having been executed in the Verilog source code, or a `tf_dofinish()` routine having been executed in a PLI application. Many Verilog simulators also provide an exit command from the simulator's interactive debug environment. The `REASON_FINISH` callback allows a PLI application to clean up before simulation exits. For example, if the PLI application had allocated memory, opened disk files, or opened graphics windows, then the application has the opportunity to free the memory and close the files and windows at the time the simulator exits.

The **REASON_INTERACTIVE** callback occurs whenever the simulator encounters a breakpoint and enters the simulator's interactive debug environment. The `$stop` built-in system task and `tf_dostop()` routine are two methods of specifying breakpoints, and most simulators have additional ways to specify breakpoints.

Example 12-2, which is listed below, is an excerpt from a PLI application called `$read_stimulus`, which reads stimulus from a disk file. The stimulus includes both a test vector value and a simulation time to apply the vector. The complete `$read_stimulus` application is listed in example 12-7 on page 411. This example uses the *misctf routine* automatic callbacks to:

- Allocate a block of memory for data storage and open a file at the start of simulation (REASON_ENDOFCOMPILe).
- Free the data structure memory and close the file at the end of simulation (REASON_FINISH).
- Print a status message each time simulation halts at a breakpoint (REASON_INTERACTIVE).
- Save the data structure memory if the simulation state is saved (REASON_SAVE).
- Restore the data structure memory and file read position if the simulation is restarted from a saved state (REASON_RESTART).

Note that some of the automatic callback reasons used in this example are not part of the required callbacks in the IEEE standard. They are, however, part of the recommended optional callbacks specified in the standard. It is not a portability problem to use these optional callbacks, since the constants which represent the optional reasons are defined in the IEEE standard. If a simulator has not implemented the optional callback features, the callback will simply never occur with that simulator.

Example 12-2: pseudo-code for using automatic *misctf routine* callbacks

```
int PLIBbook_ReadVector_misctf(int user_data, int reason)
{
    switch(reason) {
        case REASON_ENDOFCOMPILe: /* misctf called at start of simulation */
            /* allocate a memory block for data storage */
            data_p = (p_stim_data)malloc(sizeof(s_stim_data));
            /* store a pointer to the application data in the work area */
            tf_setworkarea((char *)data_p);
            /* add code to open stimulus file and save pointer to file */
            break;

        case REASON.REACTIVATE: /* misctf called by tf_setdelay */
            /* get the pointer to the data structure from the work area */
            data_p = (p_stim_data)tf_getworkarea();
            /* add code to read next line from file and apply to simulation */
            break;
    }
}
```

```
case REASON_FINISH: /* misctf called at end of simulation */
    /* add code to close files and free any allocated memory */
    break;

case REASON_SAVE: /* misctf called end of simulation */
    /* add code to save PLI application data */
    break;

case REASON_RESTART: /* misctf called end of simulation */
    /* add code to retrieve saved PLI application data */
    break;
}
return(0);
}
```

Checktf routines versus reason_endofcompile callbacks

A *checktf routine* and a *misctf routine* at a REASON_ENDOFCOMPILe callback are both called prior to simulation time zero, and both callbacks occur for each instance of a system task/function. The IEEE standard does not state when in the compile phase the *checktf routine* will be called. Potentially, a *checktf routine* could be called before the simulator has finished building its simulation data structure, which means some compile-time activity—such as parameter redefinitions—may not be completed before the *checktf routine* is called. On the other hand, a *misctf routine* callback for REASON_ENDOFCOMPILe will always occur immediately before simulation starts running, after the simulator has completed building its internal data structures.

To be portable to all Verilog simulators, PLI applications should limit activity in a *checktf routine* to verifying the correctness of system task functions only. Activity such as allocating memory, opening files, and other start-of-simulation actions should not be performed during compile/load time, but instead should be postponed until the *misctf routine* is called at the end-of-compilation.

12.3 Application scheduled callbacks at the end of a time step

A PLI application can schedule callbacks to the *misctf routine* at the end of the current simulation time step. This capability allows PLI applications to synchronize with the activity in simulation. The callback to the *misctf routine* can be scheduled from either a *calltf routine* or a *misctf routine* (but not a *sizetf routine* or *checktf routine*). The TF library allows two types of synchronization callbacks within the current simulation time step:

- **Read-write synchronization** will call the *misctf routine* at the end of all known events in the current simulation time. A PLI application is allowed to write logic values into the simulation in the current time (e.g.: a PLI application can add more events in the current simulation time after all known events have been processed).
- **Read-only synchronization** will call the *misctf routine* at the end of all events in the current simulation time. A PLI application is *not* allowed to write logic values into the simulation in the current or a future simulation time.

There are a number of reasons a PLI application might need to synchronize activity with the end of a simulation time step. One example would be to communicate value changes to a C language model, and have the C model pass value changes for that time step back to simulation. This type of activity would utilize a read-write synchronization.

Another example would be to wait until all activity in a simulation time step is stable before reading logic values from the simulation. This type of activity would utilize a read-write synchronization. In the following Verilog HDL source code:

```
always @(a or b)           always @(a or b)
  $my_strobe(sum);          sum = a + b;
  
                                          ^-----^
                                         parallel (concurrent) activity
```

The *calltf routine* for `$my_strobe` will be called every time `a` or `b` changes value, at which time the PLI application might need to read the value of argument 1 (the `sum` signal). At the same moment in simulation time, however, the `sum` variable is also scheduled to change value in the simulation. This is a classic race condition in Verilog simulation that is caused by the concurrent activity of a value being read and written in the same simulation time step. The Verilog standard states that most concurrent activity can be processed in any order by the simulation, which means that the *calltf routine* can be executed either before or after the `sum` signal changes at the positive edge of the clock. The outcome of this race condition is unpredictable.

The race condition in the above example can be resolved by synchronizing the `$my_strobe` PLI application to the very end of the simulation time step in which the change on `a` or `b` occurs. When a read-only synchronization callback to the *misctf routine* occurs, any and all statements for that time step will have been executed, and the values of the arguments to the system task/function will be at their most current value. For the `$my_strobe` example, when the *calltf routine* is invoked when `a` or `b` changes, it can schedule a read-write synchronize callback to the *misctf routine* at the end of that time step. The *misctf routine* can then read the value of `sum`, and be assured that it has the most current value for that moment in time.

The TF library provides two routines to synchronize to the end of a simulation time step, one for read-write synchronization, and one for read-only synchronization.

`void tf_synchronize()`


The `tf_synchronize()` routine schedules a callback to the *misctf routine* at the end of the current simulation time step in a *read-write* mode. The reason constant passed to the *misctf routine* is `REASON_SYNCH`. In this mode, the *misctf routine* can use TF routines to read the value of a system task/function argument, and can also use TF routines to modify an argument's value (the ACC library can also be used to read and modify the values of objects).

`void tf_rosynchronize()`

`tf_rosynchronize()` schedules a callback to the *misctf routine* at the end of the current simulation time step in a *read-only* mode. The reason constant passed to the *misctf routine* is `REASON_ROSYNCH`. In this mode, the *misctf routine* is only permitted to read values. It is not allowed to modify values in either the current or a future simulation time.

Figure 12-1 illustrates the activity that occurs within a simulation time step in Verilog.

Figure 12-1: Organization of events in a Verilog simulation time step



In essence, there are four distinct regions of events within a time step, which are referred to as *slots* in the preceding figure. Within each slot, certain types of simulation events are scheduled to be executed. The Verilog standard allows simulators to optimize the events within a slot in any order, and to intermix the types of events within each slot in any order.



NOTE *This illustration of a simulation time step is purely conceptual.* While the illustration is accurate, it is an abstraction of the much more detailed description of simulation event scheduling that is specified in the IEEE 1364 Verilog standard.

A simulation will not proceed to slot 2 until it has executed all scheduled events within slot 1. However, the events within slot 2 may cause new events to be added to slot 1. Simulation must then return to slot 1, and execute the new events, which may result in new events in slot 2. Note that when simulation returns to a slot, only new events are processed. Events which have already been executed have been removed from that slot's event list, and will not be executed a second time. Assuming there are no zero-delay infinite loops in the Verilog code, simulation will eventually complete all events in slot 1 and 2, and then proceed to slot 3.

In slot 3, if a callback had been previously scheduled using `tf_synchronize()`, a *misctf routine* will be called with `REASON_SYNCH`. In the synchronize mode, the PLI application is allowed to schedule new events in the current simulation time step. These new events will be scheduled in slot 1, forcing the simulator to once again return to slot 1 to process the new events. Note again, that only new events will be processed when the simulator returns to slot 1.

Once all events in slots 1, 2 and 3 have been processed, simulation will proceed to slot 4. In this slot, a *misctf routine* will be called with `REASON_ROSYNCH`, if a callback had been previously scheduled using `tf_rosynchronize()`. In the read-only synchronize mode, the PLI application is allowed to read information from the simulation, but the PLI may *not* schedule new events in the current simulation time step (the VPI routines are allowed to schedule events in a future simulation time, but the TF and ACC routines are prohibited from scheduling any events from slot 4). Slot 4 represents the true end of the current simulation time step.

When the *misctf routine* is invoked with `REASON_ROSYNCH`, the routine may not modify or schedule the modification of any values in simulation. However, the *misctf routine* is allowed to schedule another read-only synchronize callback to itself, using `tf_rosynchronize()`.

Example 12-3 illustrates using `tf_rosynchronize()` to implement `$my_strobe`.

An example of using `$my_strobe` is:

```
always @(a or b)           always @(a or b)
  $my_strobe(sum);          sum = a + b;

```

parallel (concurrent) activity

In this example usage, the *calltf routine* for `$my_strobe` will be executed when `a` or `b` changes value, at the same moment the signal `sum` will be changed. Since there is no way to determine if the *calltf routine* will be executed before or after the change to `sum`, the *calltf routine* does not read the value of `sum`. Instead, the *calltf routine* schedules a callback to the *misctf routine* at the end of the current time step, and the value of `sum` is read from the *misctf routine*.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.012/my_strobe_tf.c
- Verilog test bench: Chapter.012/my_strobe_test.v
- Verilog-XL results log: Chapter.012/my_strobe_test.log

Example 12-3: `$my_strobe` — using the `tf_rosynchronize()` routine

```
*****
 * calltf routine
 *****
int PLIbook_MyStrobe_calltf() {
    tf_rosynchronize();
    return(0);
}

*****
 * misctf routine
 *****
int PLIbook_MyStrobe_misctf(int user_data, int reason)
{
    if (reason == REASON_ROSYNCH) {
        io_printf("Status of tfarg 1 is %s:\n", tf_strgetp(1, 'b'));
    }
    return(0);
}
```

12.4 Application-scheduled callbacks at a future simulation time

The TF library provides three routines that allow a PLI application to schedule its *misctf routine* to be called at a future simulation time step.

int tf_setdelay(delay)

int delay a 32-bit time value greater than or equal to 0.

int tf_setlongdelay(lowdelay, highdelay)

int lowdelay lower (right-most) 32-bits of a 64-bit time value.

int highdelay upper (left-most) 32-bits of a 64-bit time value.

int tf_setrealdelay (realdelay)

double realdelay a real time value greater than or equal to 0.

The *tf_setdelay()*, *tf_setlongdelay()*, and *tf_setrealdelay()* routines allow a PLI application to schedule its *misctf routine* to be called at a future simulation time step. The three routines allow the simulation time to be specified as a 32-bit integer, a pair of 32-bit integers (for a 64-bit time value) or as a C double. The callback is a scheduled event, which will occur somewhere in the first slot of simulation activity at the future time step. The time value is a relative delay from the current simulation time, and the delay is automatically scaled to the time scale of the Verilog module which contains the system task/function which called the PLI application. The routines will return a 1 (true) if successful, and a 0 (false) if an error occurred.

Any number of callbacks can be scheduled. If multiple callbacks are scheduled in the same time step, the *misctf routine* will be called once for each scheduled callback. Callbacks can also be scheduled in the current time step (using a delay of zero). When the *misctf routine* is invoked by a callback scheduled using any of these routines, the simulator will pass a reason value of **REASON.REACTIVATE**.

The TF library can also remove any callbacks scheduled with *tf_setdelay()*, *tf_setlongdelay()*, and *tf_setrealdelay()* which have not yet transpired.

void tf_clearalldelays()

tf_clearalldelays() will remove all pending (not yet executed) callbacks for a specific instance of a *misctf routine*. Suppose a Verilog model contained two instances of a system task named *\$read_stimulus*, and each instance had scheduled future callbacks. If a PLI application for one of the *\$read_stimulus* instances called *tf_clearalldelays()*, only the pending callbacks for that instance would be removed. The pending callbacks for the other instance would not be affected.

An example of scheduling callbacks at a future time

Example 12-4, which follows, illustrates one usage of `tf_setdelay()`. The example implements a system task called `$read_stimulus`, which uses the *misctf routine* to read one line at a time from a stimulus vector file. Each line contains a simulation time and a test vector, as follows:

time	vector
10	11111111xxxxxxxxxx
17	00000000zzzzzzzz
30	1000000011011101

The *misctf routine* for `$read_stimulus` will need to process the following loop as long as there are more test vectors in the file:

- 
1. Read a time and the next test vector from the file.
 2. Schedule test vector to be applied to simulation at the desired time.
 3. Schedule the *misctf routine* to be called again in the same future time step in which the test vector is applied.

The usage model for `$read_stimulus` is:

```
reg [15:0] input_vector;
initial
  $read_stimulus("read_stimulus.pat", input_vector);
```

Notice that `$read_stimulus` is called from a Verilog `initial` procedure, which means the *calltf routine* for `$read_stimulus` will only be invoked one time throughout simulation. The *calltf routine* calls `tf_setdelay(0)`, which schedules a callback to the *misctf routine* with `REASON.REACTIVATE`. The *misctf routine* will then read the first line from the file, and schedule the test vector to be applied at the simulation time specified by the delay (using `tf_strdelptr()`). After scheduling the vector to be applied, the *misctf routine* calls `tf_setdelay()` with a delay time in order to schedule a future callback to itself, again with `REASON.REACTIVATE`. When the last line of the file is read, the *misctf routine* will cause the simulation to exit. The *misctf routine* is also used to open the test vector file at the beginning of simulation (for `REASON.ENDOFCOMPILE`). Example 12-7 on page 411, listed at the end of this chapter, is a more robust version of the `$read_stimulus` example, with added capabilities for supporting different vector formats and delay specifications.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.012/read_stimulus_short_tf.c
- Verilog test bench: Chapter.012/read_stimulus_short_test.v
- Verilog-XL results log: Chapter.012/read_stimulus_short_test.log

Example 12-4: \$read_stimulus — using the tf_setdelay() routine

```
#include <stdio.h>          /* ANSI C standard I/O library */
#include "veriuser.h"        /* IEEE 1364 PLI TF routine library */
/******************************************************************************/
/* calltf routine                                         */
/******************************************************************************/
int PLIBook_ReadStimulusShort_calltf()
{
    /* call the misctf routine at the end of current      */
    tf_setdelay(0); /* time step, with REASON.REACTIVATE; the misctf      */
    /* routine reads the stimulus file.                      */
    return(0);
}

/******************************************************************************/
/* misctf routine                                         */
/******************************************************************************/
int PLIBook_ReadStimulusShort_misctf(int user_data, int reason)
{
    int delay, foo;
    FILE *file_p;           /* pointer to the test vector file      */
    char vector[1024];     /* stimulus vector -- hard coded limit */

    switch(reason) {
        case REASON_ENDOFCOMPILe: /* misctf called at start of simulation */
            file_p = fopen(tf_getcstringp(1), "r");
            /* store a pointer to the file in the work area */
            tf_setworkarea((char *)file_p);
            break;

        case REASON.REACTIVATE: /* misctf called by tf_setdelay */
            /* get the file pointer from the work area */
            file_p = (FILE *)tf_getworkarea();
            /* read next line from the file */
            if ((fscanf(file_p, "%d %s\n", &delay, vector)) == EOF) {
                io_printf("$read_stimulus reached end-of-file\n");
                return(0);
                break;
            }
            /* schedule the vector to be applied after the delay period */
            tf_strdelpup(2, tf_sizep(2), 'b', vector, delay, 0);
            /* call this routine back after delay time */
            tf_setdelay(delay);
            break;
    }
}
```

```
    return(0);
}
```

12.5 System task/function argument value change callbacks

A PLI application can schedule with the simulator to call the *misctf routine* whenever an argument to a system task changes value.

void tf_asynchon()

The `tf_asynchon()` routine enables asynchronous callbacks to the *misctf routine* whenever an argument of a system task changes value. The reason constant which the simulator will pass to the *misctf routine* is `REASON_PARAMVC`. The simulator also passes a third C function input to the *misctf routine*, called the `paramvc` input, which contains the index of the system task argument which changed value. Using this index number, a PLI application can then read the new value of the task argument, or do any other processing desired at the time of the value change.

NOTE → Asynchronous callbacks may only be used with user-defined system tasks. The routine `tf_asynchon()` is ignored if it is called by a PLI application that is defined as a system function.

void tf_asynchoff()

`tf_asynchoff()` disables the argument change callbacks for the current instance of the system task.

The following example illustrates the basic usage of `tf_asynchon()`. In this example, the asynchronous callbacks are enabled by the *calltf routine*. The PLI application is invoked by `$my_monitor1`, as follows:

```
initial
$my_monitor1(clock, d, q);
```



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.012/my_monitor1_tf.c
- Verilog test bench: Chapter.012/my_monitor1_test.v
- Verilog-XL results log: Chapter.012/my_monitor1_test.log

Example 12-5: \$my_monitor1 — asynchronous argument value change callbacks

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
/*********************************************
 * calltf routine
 *****/
int PLIbook_MyMonitor1_calltf() {
    tf_asynchon(); /* enable asynchronous misctf callbacks */
    return(0);
}

/*********************************************
 * misctf routine
 *****/
int PLIbook_MyMonitor1_misctf(int user_data, int reason, int paramvc)
{
    if (reason == REASON_PARAMVC) {
        io_printf("At %s: tfarg %d changed, new value is %s:\n",
                  tf_strgettime(), paramvc, tf_strgetp(paramvc, 'b'));
    }
    return(0);
}
```

In the preceding example, if two task arguments change at the same moment in simulation time, the *misctf routine* could be called twice, asynchronously. This may be desirable in some applications, but in other applications it may be desirable to only process data at the end of the simulation time step, when all activity for that moment of time is complete. The synchronous callback to the *misctf routine* can be combined with asynchronous callbacks to accomplish this.

Synchronous callbacks at the end of the time step are scheduled using `tf_synchronize()` and `tf_rosynchronize()`, which were discussed in section 12.3 on page 395. In order to pass information about which system task arguments changed value (which result in asynchronous callbacks of the *misctf routine*) to the synchronous callback of the *misctf routine*, it is necessary to temporarily save those task arguments which changed value. The PLI provides a mechanism to automatically save this information.

The Verilog simulator maintains two internal flags for each system task argument. The flags are referred to as the *PVC* flags (for *Parameter Value Change*). One of the flags is the *current PVC flag*, and the other is the *saved PVC flag*. Each time an asynchronous callback occurs, the simulation sets the current PVC flag for the task argument that changed value. By itself, the current PVC flag has little value. However, a *misctf routine* can move or copy the current PVC flag to the saved PVC flag. Then, at the end of the simulation time step, the *misctf routine* can examine all saved PVC flags to determine which task arguments changed value sometime during the current time step.



PVC flags are not enabled by a simulator until `tf_asynchon()` has been called to enable asynchronous callbacks to the *misctf routine*.

`int tf_testpvc_flag(n)`

`int n` index number of a system task/function argument, or -1.

The `tf_testpvc_flag()` routine can be used to see if the *current* PVC flag of a specific system task argument is set. The input provided to `tf_testpvc_flag()` is the index number of a task argument. The `tf_testpvc_flag()` can also be used to test the saved PVC flags. When the input value is -1, then `tf_testpvc_flag()` will return the result of a logical OR of all *saved* PVC flags.

`int tf_copypvc_flag(n)`

`int n` index number of a system task/function argument, or -1.

`tf_copypvc_flag()` is used to copy the *current* PVC flag of a specific system task argument to the *saved* PVC flag for that argument. The current PVC flag is not cleared by the copy. `tf_copypvc_flag()` must be called from a *misctf routine* that was called with `REASON_PARAMVC`, which is when the current PVC flag is active. The input provided to `tf_copypvc_flag()` is the index number of a task argument, typically the argument that just changed value. The return value of `tf_copypvc_flag()` is the value of the PVC flag which was just copied. If the input value to `tf_copypvc_flag()` is -1, then the routine returns the result of a logical OR of all *saved* PVC flags.

`int tf_movepvc_flag(n)`

`int n` index number of a system task/function argument, or -1.

The `tf_movepvc_flag()` routine serves a dual role:

- If `tf_movepvc_flag()` is called from a *misctf routine* that was called with `REASON_PARAMVC`, the routine will move the *current* PVC flag of a specific system task argument to the *saved* PVC flag for that argument. The current PVC flag is cleared by the move. The input provided to `tf_movepvc_flag()` is the index number of a task argument, typically the argument that just changed value. The return value of `tf_movepvc_flag()` is the value of the PVC flag which was just moved. If the input value to `tf_movepvc_flag()` is -1, then the routine returns the result of a logical OR of all *saved* PVC flags.
- If `tf_movepvc_flag(-1)` is called from a *misctf routine* that was called with `REASON_SYNCH` or `REASON_ROSYNCH`, the routine will setup for all *saved* PVC flags to be examined by `tf_getpchange()`.

int tf_getpchange(n)

int n index number of a system task/function argument.

tf_getpchange() is used to scan all saved PVC flags and return the index number of each system task argument with PVC flag that is set. *tf_getpchange()* should be used in a *misctf routine* that was called with REASON_SYNCH or REASON_ROSYNCH. The input to *tf_getpchange()* is the index number of a system task argument. The routine returns the index number of the *next* task argument with a saved PVC flag that is set. If there are no task arguments with a greater index number than the input that has a saved PVC flag set, then *tf_getpchange()* will return 0. There are two mandatory rules for using *tf_getpchange()*:

1. *tf_movepvc_flag(-1)* must be called prior to calling *tf_getpchange()*.
2. The first call to *tf_getpchange()* must have 0 as the input value.

Example 12-6 implements *\$my_monitor2*. This example enables asynchronous argument value change callbacks, and then, in any time step that one or more task arguments change value, prints a summary of all changes. The printing of the summary is synchronized to the end of the simulation time step in which the changes occurred.

```
initial
$my_monitor2(clock, d, q);
```



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.012/my_monitor2_tf.c
- Verilog test bench: Chapter.012/my_monitor2_test.v
- Verilog-XL results log: Chapter.012/my_monitor2_test.log

Example 12-6: *\$my_monitor2* — synchronized analysis of argument value changes

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
/*********************************************calltf routine*****/
int PLIBbook_MyMonitor2_calltf() {
    tf_asynchon(); /* enable asynchronous misctf callbacks */
    return(0);
}
```

```
*****
 * misctf routine
 ****
int PLIbook_MyMonitor2_misctf(int user_data, int reason, int paramvc)
{
    int arg_num;

    if (reason == REASON_PARAMVC) {
        /* io_printf("At %s: change detected on tfarg %d\n",
           tf_strgettime(), paramvc); */
        tf_copypvc_flag(paramvc);
        tf_rosynchronize(); /* schedule a callback at end of time step */
    }

    if (reason == REASON_ROSYNCH) {
        io_printf("Reached end of time step %s:\n", tf_strgettime());
        if (tf_movepvc_flag(-1)) { /* only print if something changed */
            arg_num = 0;
            while (arg_num = tf_getpchange(arg_num) ) {
                io_printf(" tfarg %d changed to %s\n",
                          arg_num, tf_strgetp(arg_num, 'b'));
            }
            io_printf("\n");
        }
    }
    return(0);
}
```

12.6 Simulation save and restart callbacks

Some Verilog simulators provide the means to save the simulation state, and restart simulation from the saved state. Since a PLI application has work area storage, and can allocate storage for its own memory, file pointers, etc., it is important that a simulator save and restore PLI application storage as part of the simulation save and restart feature.

If a simulator supports save and restart capabilities, then the *misctf routine* will be called automatically with **REASON_SAVE** at the start of a save state, and **REASON_RESTART** at the end of a simulator restart from a saved file.

int tf_write_save(blockptr, blocklength)

char * **blockptr** pointer to a block of memory.

int **blocklength** length of the block of memory in bytes.

The `tf_write_save()` routine may only be called from the *miscf routine* of a PLI application. `tf_write_save()` instructs the simulator to add an application-specified block of memory to the save file being created by the simulator. The block of memory can contain any information important to the PLI application. The data stored in the PLI work area is not automatically saved. The application should copy that information into the block of memory to be saved before calling `tf_write_save()`. The inputs to `tf_write_save()` are a pointer to the application-specified block of memory, and the length of the memory block. The routine will return a non-zero value if successful, and 0 if an error occurred.

`int tf_read_restart(blockptr, blocklength)`

`char * blockptr` pointer to a block of memory.

`int blocklength` length of the block of memory in bytes.

The `tf_read_restart()` routine may only be called from the *miscf routine* of a PLI application. `tf_read_restart()` retrieves from the simulator the data that was saved by `tf_write_save()`. Before calling `tf_read_restart()`, the PLI application must first allocate memory to receive the restored data. A pointer to the allocated memory, and the length of the saved data to be restored is passed to `tf_read_restart()`. The routine will return 1 if successful, and 0 if an error occurred.



TIP

`tf_read_restart()` retrieves the saved data as an input stream, and automatically maintains a pointer to how much of the saved block of data has been retrieved. This means the block of saved data can be restored using one call to `tf_read_restart()`, or using multiple calls. Often, the length of the block of data that was saved is not known when the restart occurs. An easy way to pass the length of the saved data from the save operation to the restart operation, is to make the first piece of data saved an integer which contains the save block length. Upon restart, `tf_read_restart()` can be called twice: First, with a block length equal to one integer, in order to retrieve the total saved block length, and then a second time, to retrieve the remainder of the saved block.



NOTE *It is mandatory that all data be retrieved when ever a restart occurs.* Because `tf_read_restart()` retrieves the data as a stream, if one PLI application does not retrieve its data, another PLI application, which might have saved data afterwards, will not be able to retrieve the correct data. The PLI standard guarantees that all calls to the *miscf routine* on a restart will occur in the same order that calls occurred on a save. A PLI application should never assume it will be the only PLI application which saved data, and can therefore ignore retrieving all saved data.

When simulation is restarted, any pointers which were saved may no longer be valid. A file pointer, for example, can no longer be used, because the file may have been closed between the save and restart operations, and any application-allocated memory

may have been freed or relocated, making pointers to the memory invalid. The PLI application must define a scheme to re-open files if needed, and to re-allocate memory if needed.



Save and restart capabilities are not part of the IEEE 1364 Verilog standard, and are not implemented in many simulators. To allow PLI applications to be portable, however, the IEEE 1364 standard includes the `tf_write_save()` and `tf_read_restart()` routines as part of the PLI standard. Simulators which do not have save and restart capability will return an error status when the routines are called. This allows a PLI application to use the routines without concern about specific implementations. However, the application should check the error status to determine if the application data was saved or restored by the simulator.

Example 12-7 on page 411 illustrates how to use save and restart with the `$read_stimulus` stimulus pattern file reader. In this example, the *misctf routine* will append the position indicator of the next line in the stimulus file to be read to the simulation save file. When simulation is restarted, the stimulus file is re-opened, and the C file pointer is reset to the saved position indicator. The example also adds a character string to the simulation save file, in order to illustrate how the data can be retrieved when simulation is restarted. Note that the example also saves the file pointer and stimulus pattern string pointers, but these pointers may no longer be valid when simulation is restarted. Therefore these pointers are recreated on a restart.

12.7 A complete example of using misctf routine callbacks

Example 12-7, which follows, is a more extensive stimulus file reader than was presented in example 12-4 earlier in this chapter. The `$read_stimulus_ba` PLI application listed below utilizes several types of *misctf routine* callbacks:

- `REASON_ENDOFCOMPILE` is used to both allocate memory that the `$read_stimulus_ba` application requires, and to open the file containing the stimulus patterns.
- `REASON.REACTIVATE` is used to retrieve a stimulus pattern and a simulation time from the file, and then schedule the simulator to apply the test pattern and the designated simulation time.
- `REASON_ROSYNCH` is used to terminate simulation when `$read_stimulus_ba` has reached the end of the stimulus pattern file. Instead of terminating simulation immediately on end of file, a read-only synchronous callback is scheduled, to allow any other activity in the current simulation time step to complete first.
- `REASON_FINISH` is used to close the stimulus file and free the memory that was allocated at the start of simulation.

- REASON_SAVE is used to save the current position from which the `$read_stimulus_ba` will retrieve the next stimulus pattern from the file.
- REASON_RESTART is used to retrieve the file position indicator, so that simulation can resume at the next test stimulus pattern in the file.

This example also extends the file handling capability by permitting stimulus values to be specified as either hex patterns or binary patterns, and by allowing the simulation times to be specified as relative delays or as absolute time from when the PLI application was invoked. This is accomplished by having several versions of the system task name, all of which call the same PLI application, but with different user-data values:

- `$read_stimulus_ba` has a user-data of 0
- `$read_stimulus_br` has a user-data of 1
- `$read_stimulus_ha` has a user-data of 2
- `$read_stimulus_hr` has a user-data of 3

The usage for this example is:

```
initial
    $read_stimulus_<base><delay_type>("file_name",
verilog_reg);
```

where:

`<base>` is **b** or **h** (for binary or hex vectors).
`<delay_type>` is **a** or **r** for absolute or relative times.
`"file_name"` is the name of the file to be read, in quotes.
`verilog_reg` is a verilog register data type of the same bit width as the patterns to be read.

For example:

```
initial
    $read_stimulus_ba("read_stimulus.pat", input_vector);
```



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.012/read_stimulus_file_tf.c
- Verilog test bench: Chapter.012/read_stimulus_file_test.v
- Verilog-XL results log: Chapter.012/read_stimulus_file_test.log

Example 12-7: \$read_stimulus_ba — using several misctf routine callbacks

```
#include <stdio.h>           /* ANSI C standard I/O library */
#include "veriuser.h"         /* IEEE 1364 PLI TF routine library */

/* prototypes of sub-functions */
char *PLIbook_reason_name();

/*************************************************************************/
/* structure definition for data used by the misctf routine          */
/*************************************************************************/
typedef struct PLIbook_stim_data {
    FILE *file_ptr;        /* pointer to the test vector file */
    long file_position;   /* position within file of next byte to read */
    char *vector;          /* pointer to stimulus vector */
    char dummy_msg[20];   /* dummy message field to show save/restart */
} PLIbook_stim_data_s, *PLIbook_stim_data_p;

/*************************************************************************/
/* checktf routine                                                       */
/*************************************************************************/
int PLIbook_ReadStimulus_checktf()
{
    bool err = FALSE;

    if (tf_nump() != 2) {
        tf_text("$read_stimulus_?? requires 2 arguments\n");
        err = TRUE;
    }
    if (tf_typep(1) != tf_string) {
        tf_text("$read_stimulus_?? arg 1 must be a string\n");
        err = TRUE;
    }
    if (tf_typep(2) != tf_readwrite) {
        tf_text("$read_stimulus_?? arg 2 must be a register data type");
        err = TRUE;
    }
    if (err)
        tf_message(ERR_ERROR, "", "", "");
    return(0);
}

/*************************************************************************/
/* calltf routine                                                       */
/*************************************************************************/
int PLIbook_ReadStimulus_calltf()
{
    /* call the misctf routine at the end of current */
    tf_setdelay(0); /* time step, with REASON.REACTIVATE; the misctf */
                     /* routine reads the stimulus file. */
    return(0);
}
```

```
/*************************************************************************/
/* misctf routine                                                       */
/*************************************************************************/
int PLIbook_ReadStimulus_misctf(int user_data, int reason)
{
    int    delay, foo;
    char   base;
    PLIbook_stim_data_p data_p;

    bool  debug = FALSE;
    if(mc_scan_plusargs("read_stimulus_debug")) debug = TRUE;

    if (debug)
        io_printf("** read_stimulus misctf called for %s at time %s **\n",
                  PLIbook_reason_name(reason), tf_strgettime());

    switch(reason) {

        case REASON_ENDOFCOMPILE: /* misctf called at start of simulation */

            /* Check work area to see if application data is already
             * allocated and the vector file opened. This check is necessary
             * because REASON_RESTART can occur before REASON_ENDOFCOMPILE */
            data_p = (PLIbook_stim_data_p)tf_getworkarea();
            if (data_p)
                return(0); /* abort if work area already contains data */

            /* allocate a memory block for data storage */
            data_p=(PLIbook_stim_data_p)malloc(sizeof(PLIbook_stim_data_s));

            /* store a pointer to the application data in the work area */
            tf_setworkarea((char *)data_p);

            /* fill in the application data fields */
            data_p->file_ptr = fopen(tf_getcstringp(1), "r");
            if ( !data_p->file_ptr ) {
                tf_error("$read_stimulus_?? could not open file %s",
                         tf_getcstringp(1));
                tf_dofinish(); /* exit simulation if cannot open file */
                break;
            }
            data_p->file_position = ftell(data_p->file_ptr);
            data_p->vector = (char *) malloc((tf_sizep(2) * 8) + 1);
            strcpy(data_p->dummy_msg, "Hello world");
            break;

        case REASON.REACTIVATE: /* misctf called by tf_setdelay */

            /* get the pointer to the data structure from the work area */
            data_p = (PLIbook_stim_data_p)tf_getworkarea();
```

```
/* read next line from the file */
if ( (fscanf(data_p->file_ptr, "%d %s\n",
              &delay, data_p->vector)) == EOF) {
    tf_rosynchronize(); /* if end of file, schedule a callback */
    break;               /* at the end of the current time step */
}
if (debug)
    io_printf("** values read from file: delay = %d  vector = %s\n",
              delay, data_p->vector);

/* set flag for test vector radix */
switch(user_data) {
    case 0:
    case 1: base = 'b'; break; /* vectors are in binary format */
    case 2:
    case 3: base = 'h'; break; /* vectors are in hex format */
}

/* convert absolute delays to relative to current time */
switch(user_data) {
    case 0:
    case 2: /* using absolute delays; convert to relative */
        delay = delay - tf_gettime();
}

/* schedule the vector to be applied after the delay period */
tf_strdelputp(2, tf_sizep(2), base, data_p->vector, delay, 0);

/* schedule reactive callback to this routine after delay time */
tf_setdelay(delay);
break;

case REASON_ROSYNCH: /* misctf called at end of time step */
    io_printf("\n$read_stimulus_?? has encountered end-of-file.\n");
    tf_dofinish;
    break;

case REASON_FINISH: /* misctf called at end of simulation */
    io_printf("\nPDLI is processing finish at simulation time %s\n\n",
              tf_strgettime());
    /* get the pointer to the application data from the work area */
    data_p = (PLIbook_stim_data_p)tf_getworkarea();
    /* close file */
    if (data_p->file_ptr) fclose(data_p->file_ptr);
    /* de-allocate storage */
    free(data_p);
    break;
```

```
case REASON_SAVE: /* miscf called for $save */

    /* get the pointer to the application data from the work area */
    data_p = (PLIbook_stim_data_p)tf_getworkarea();

    /* save current file position in the application data */
    data_p->file_position = ftell(data_p->file_ptr);

    /* add application data to simulation save file */
    tf_write_save((char *)data_p, sizeof(PLIbook_stim_data_s));
    if (debug)
        io_printf("\nPLI data saved (last file position was %ld)\n\n",
                  data_p->file_position);
    break;

case REASON_RESTART: /* miscf called end of simulation */

    /* re-allocate memory for PLI application data */
    data_p=(PLIbook_stim_data_p)malloc(sizeof(PLIbook_stim_data_s));

    /* save new application data pointer in work area */
    tf_setworkarea((char *)data_p);

    /* retrieve old application data from save file */
    if (!tf_read_restart((char *)data_p,
                         sizeof(PLIbook_stim_data_s)) ) {
        tf_error("\nError retrieving PLI data from save file!\n");
        return(0);
    }
    if (debug) {
        io_printf("\nPLI data retrieved from save file.\n");
        /* test to see if old application data was restored */
        io_printf(" dummy message = %s, file position = %ld\n\n",
                  data_p->dummy_msg, data_p->file_position);
    }
    /* re-open test vector file */
    data_p->file_ptr = fopen(tf_getcstringp(1), "r");
    if ( !data_p->file_ptr ) {
        tf_error("$read_stimulus_?? could not re-open file %s",
                 tf_getcstringp(1));
        tf_dofinish(); /* exit simulation if cannot open file */
    }

    /* re-position file to next test vector to be read */
    if ( fseek(data_p->file_ptr, data_p->file_position, SEEK_SET) ) {
        tf_error("$read_stimulus_?? could not reposition file");
        tf_dofinish(); /* exit simulation if reposition open file */
    }
    break;
}
return(0);
}
```

```
*****
/* Function to convert reason integer to reason name */
*****
char *PLIbook_reason_name(int reason)
{
    char str[25];
    switch (reason) {
        case REASON_ENDOFCOMPILE : return("REASON_ENDOFCOMPILE"); break;
        case REASON_FINISH       : return("REASON_FINISH"); break;
        case REASON_INTERACTIVE  : return("REASON_INTERACTIVE"); break;
        case REASON_SYNCH        : return("REASON_SYNCH"); break;
        case REASON_ROSYNCH      : return("REASON_ROSYNCH"); break;
        case REASON.REACTIVATE   : return("REASON.REACTIVATE"); break;
        case REASON_PARAMVC      : return("REASON_PARAMVC"); break;
        case REASON_PARAMDRC     : return("REASON_PARAMDRC"); break;
        case REASON_SAVE          : return("REASON_SAVE"); break;
        case REASON_RESTART       : return("REASON_RESTART"); break;
        case REASON_RESET         : return("REASON_RESET"); break;
        case REASON_ENDOFRESET   : return("REASON_ENDOFRESET"); break;
        case REASON_FORCE         : return("REASON_FORCE"); break;
        case REASON_RELEASE       : return("REASON_RELEASE"); break;
    }
    return("Non-standard or Unknown Reason");
}
```

12.8 Summary

This chapter has presented the TF library routines which allow PLI applications to synchronize with other types of activity that can occur during a simulation. The TF routines use the *misctf routine* to do this synchronization. Some types of callback to the *misctf routine* are automatic. The PLI application does not request the callback. Other types of callbacks are application defined, and only occur if the PLI application requests the callback.

The next chapter applies the concepts presented in this chapter and the previous chapter, to show how the TF routines can be used to interface a hardware model written in the C programming language into a Verilog simulation.

CHAPTER 13

Interfacing to C Models Using TF Routines

One of the ways the TF library can be used is to create an interface to hardware models written in the C programming language. The TF library provides routines to read and modify logic values within a simulation, and to synchronize activity with logic value changes and with simulation time. This chapter shows several ways in which a C model can be interfaced to a Verilog simulation using the TF library (Chapter 7 presents using the VPI library for interfacing to C models, and Chapter 18 shows how to use the ACC library to accomplish this same task).

The concepts presented in this chapter are:

- Representing hardware models in C
- Verilog HDL shell modules
- Combinational logic interfaces to C models
- Sequential logic interfaces to C models
- Synchronizing with the end of a simulation time step
- Synchronizing with a future simulation time step
- Multiple instances of a C model
- Creating instance specific storage within C models
- Representing propagation delays in C models

**TIP**

One reason for representing hardware models in the C language is to achieve faster simulation performance. The C programming language allows a very abstract, algorithmic representation of hardware functionality, without representing detailed timing, multi-state logic, hardware concurrency, and the many other hardware specific details offered by the Verilog language.

The PLI can be a means to access the efficiency of a highly abstract C model. However, a poorly written PLI application can become a bottleneck that offsets much of the efficiency gains. Care must be taken to write PLI applications that execute as efficiently as possible.

Some guidelines that can help maximize the efficiency and run-time performance of PLI applications are:

- Good C programming practices are essential. General C programming style and techniques are not discussed within the scope of this book.
- Consider every call to a PLI routine as expensive, and try to minimize the number of calls.
- Routines which convert logic values from a simulator's internal representation to C strings, and vice-versa, are very expensive in terms of performance. Best efficiency is attained when the value representation in C is as similar as possible to the value representation in Verilog.
- Use the Verilog language to model the things hardware description languages do well, such as representing hardware parallelism and hardware propagation times. Simulator vendors have invested a great deal in optimizing a simulator's algorithms, and that optimization should be utilized.

**NOTE**

The objective of this book is to show several ways in which the TF library can be used to interface to C models. Short examples are presented that are written in a relatively easy to follow C coding style. In order to meet the book's objectives, the examples presented in this book do not always follow the guidelines of efficient C coding and prudent usage of the PLI routines. It is expected that, when parts of these example PLI applications are adapted for other applications, the coding style will also be modified to be more efficient and robust.

13.1 How to interface C models with Verilog simulations

The power and flexibility of the C programming language and the Verilog PLI provide a wide variety of methods that can be used to interface a Verilog simulation with a C language model. All methods have three essential concepts in common:

- Value changes which occur in the Verilog simulator must be passed to the C model.
- Value changes within the C model must be passed to the Verilog simulation.


- Simulated time in both the Verilog simulation and the C model must remain synchronized.

This chapter will present some of the more common methods of interfacing a Verilog simulation with a C model. The methods presented are by no means the only ways this interface can be accomplished, and may not always be the most efficient methods. However, the methods presented have many advantages, including simplicity to implement, portability to many types of Verilog simulators, and the ability to use the C model any number of times and anywhere in the hierarchy of a Verilog design.

The fundamental steps that are presented in this chapter are:

- Create the C language model as an independent block of pure C code that does not use the PLI routines in any way. The C model will have inputs and outputs, but it will not know the source of the inputs or the destination of the outputs. The C code to implement the model might be in the form of a C function with no main function, or it might be a complete C program with its own main function.
- Create a Verilog HDL *shell module* (also called a *wrapper module*) to represent the inputs and outputs of the C language model. This module will be written completely in the Verilog language, but will not contain any functionality. To represent the functionality of the model, the shell module will call a PLI application.
- Create a PLI application to serve as an interface between the C model and the Verilog shell module. The PLI application is a communication channel, which:
 - Uses PLI routines to retrieve data from the Verilog HDL shell module, and pass the data to the C model via standard C programming.
 - Uses standard C programming to receive data from the C model, and pass the data to the Verilog shell module via PLI routines.

The following diagram shows how the blocks which are created in these three steps interact with each other.



This chapter presents steps 2 and 3 of this interface method in detail. Step 1 is to model some desired functionality or algorithm in the C language. This step is pure C programming, which does not directly involve the Verilog language or the Verilog PLI. This chapter does not cover how to implement ideas in the C language. The focus is on how to interface that implementation with a Verilog simulation. To maintain this focus, the C model example presented in this chapter will be a practical example, but relatively simple to model in C. The C model example used will illustrate all of the important concepts of integrating C models into a Verilog simulation.

13.2 Creating the C language model

A hardware model can be represented in the C programming language in two basic forms, either as a C function or as an independent C program.

13.2.1 Using functions to represent the C model

When the C model is represented as a C function, that function can be linked into the Verilog simulator, together with the PLI application that serves as the interface to the model. The PLI application can then call the function when needed, passing inputs to the function, and receiving outputs from the function. One advantage of representing a C model as a function is the simplicity of passing values to and from the model. Another advantage is ease of porting to different operating systems, since the C model is called directly from the PLI application as a C function. A disadvantage of using a function to represent the C model is that the C model must contain additional code to allow a Verilog design to instantiate the C model multiple times. The model needs to specifically create unique storage for each instance.

13.2.2 Using independent programs to represent the C model

When the C model is represented as an independent program, which means it has its own C *main* function, then the Verilog simulation and the C model can be run as parallel processes on the same or on different computers. The PLI application which serves as an interface between the simulation and the model will need to create and maintain some type of communication channel between the two programs. This communication can be accomplished several ways, such as using the `exec` command in the C standard library. On Unix operating systems, the `fork` or `vfork` commands with either Unix pipes or Unix sockets is an efficient method to communicate with the C model program. On PC systems running a DOS or windows operating system, the `spawn` command can be used to invoke the C model program and establish two-way communications between the PLI application and the C model process.


One of the advantages of representing the C model as an independent model is the ability to have parallel processes running on the same computer or separate computers. Another advantage is that when a Verilog design instantiates multiple instances of the C model, each instance will be a separate process with its own memory storage. The major disadvantage of independent programs when compared to using a C function to represent the C model is that the PLI interface to invoke and communicate with the separate process is more complex, and might be operating system dependent.

David Roberts, of Cadence Design Systems, who reviewed many of the chapters of this book, has provided a full example of representing a C model as a separate C program. This example is included with the CD that accompanies this book.

13.3 A C model example

The C model used for the different PLI interfaces shown in this chapter is a scientific Arithmetic Logic Unit, which utilizes the C math library. The C model is represented as a C function which will be called from the PLI interface mechanism. This model is written entirely with standard C library routines and C data types, without reference to any PLI routines or PLI data types. This same example is also used in other chapters, to show how a PLI interface to C models can be created, using the VPI and ACC libraries of the PLI.

The inputs and outputs of the scientific ALU C model are shown below, and Table 13-1 shows the operations which the ALU performs.



exception is set to 1 whenever an operation results in a value which is out of range of the double-precision result.

error is set to 1 whenever an input to an operation is out of range for the operation.

Opcode	C Math Library Operation
0	<code>pow(a, b)</code> — returns a to the power of b
1	<code>sqrt(a)</code> — returns the square root of a
2	<code>exp(a)</code> — returns the natural exponent of a
3	<code>ldexp(a, b)</code> — returns a * (2 to the power of b)
4	<code>fabs(a)</code> — returns the absolute of a
5	<code>fmod(a, b)</code> — returns the floating remainder of a / b
6	<code>ceil(a)</code> — returns smallest whole number not less than a
7	<code>floor(a)</code> — returns largest whole number not more than a
8	<code>log(a)</code> — returns the natural log of a
9	<code>log10(a)</code> — returns the base 10 log of a
A	<code>sin(a)</code> — returns the sine of a
B	<code>cos(a)</code> — returns the cosine of a
C	<code>tan(a)</code> — returns the tangent of a
D	<code>asin(a)</code> — returns the arcsine of a
E	<code>acos(a)</code> — returns the arccosine of a
F	<code>atan(a)</code> — returns the arctangent of a

Table 13-1: Scientific ALU C model operations

The source code for the scientific ALU is listed in Example 13-1. This version of the ALU generates a result, but does not store the result. A latched version of the ALU which stores the operation result is listed later in this chapter.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.13/sci_alu_comb_calltf_tf.c`
- Verilog shell module: `Chapter.13/sci_alu_comb_calltf_shell.v`
- Verilog test bench: `Chapter.13/sci_alu_comb_calltf_test.v`
- Verilog-XL results log: `Chapter.13/sci_alu_comb_calltf_test.log`

Example 13-1: scientific ALU C model — combinational logic version

```
#include <math.h>
#include <ERRNO.h>
void PLIbook_ScientificALU_C_model(
    double a,          /* input */
    double b,          /* input */
    int   opcode,      /* input */
    double *result,    /* output from ALU */
    int   *excep,      /* output; set if result is out of range */
    int   *err)        /* output; set if input is out of range */
{
    switch (opcode) {
        case 0x0: *result = pow    (a, b);      break;
        case 0x1: *result = sqrt   (a);         break;
        case 0x2: *result = exp    (a);         break;
        case 0x3: *result = ldexp  (a, (int)b);  break;
        case 0x4: *result = fabs   (a);         break;
        case 0x5: *result = fmod   (a, b);      break;
        case 0x6: *result = ceil   (a);         break;
        case 0x7: *result = floor  (a);         break;
        case 0x8: *result = log    (a);         break;
        case 0x9: *result = log10  (a);         break;
        case 0xA: *result = sin    (a);         break;
        case 0xB: *result = cos    (a);         break;
        case 0xC: *result = tan    (a);         break;
        case 0xD: *result = asin   (a);         break;
        case 0xE: *result = acos   (a);         break;
        case 0xF: *result = atan   (a);         break;
    }
    *err   = (errno == EDOM); /* arg to math func. out of range */
    *excep = (errno == ERANGE); /* result of math func. out of range */
    errno = 0;                 /* clear the error flag */
    if (*err) *result = 0.0;   /* set result to 0 if error occurred */
    return;
}
```

13.4 Creating a Verilog shell module

A **shell module** allows a Verilog design to reference a C model using standard Verilog HDL syntax. The shell module is a Verilog module which has the same input and output ports as the C model, but the module has no functionality modeled within. To represent the module's functionality, the shell module invokes a PLI application, which in turn invokes the C model. A shell module is sometimes referred to as a *wrapper module*, because the module is wrapped around the call to a PLI application.

The shell module for a combinational logic version of the scientific ALU is listed below.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.13/sci_alu_comb_calltf_tf.c
- Verilog shell module: Chapter.13/sci_alu_comb_calltf_shell.v
- Verilog test bench: Chapter.13/sci_alu_comb_calltf_test.v
- Verilog-XL results log: Chapter.13/sci_alu_comb_calltf_test.log

Example 13-2: Verilog shell module for the scientific ALU C model

```
'timescale 1ns / 1ns
module scientific_alu(a_in, b_in, opcode,
                      result_out, exception, error);
    input [63:0] a_in, b_in;
    input [3:0] opcode;
    output [63:0] result_out;
    output      exception, error;

    real         a, b, result; // real variables used in this module
    reg          exception, error;

    // convert real numbers to/from 64-bit vector port connections
    assign result_out = $realtobits(result);
    always @(a_in) a = $bitstoreal(a_in);
    always @(b_in) b = $bitstoreal(b_in);

    //call the PLI application which interfaces to the C model
    //using combinational logic input sensitivity
    always @(a or b or opcode)
        $scientific_alu(a, b, opcode, result, exception, error);

endmodule
```



In this scientific ALU example, the primary inputs and outputs of the model are double-precision floating point values, represented as Verilog `real` data types. The Verilog language does not permit real numbers to be connected to module ports. However, the language provides built-in system functions which convert real numbers to 64-bit vectors, and vice-versa, so the real values can be passed through a module port connection. These built-in system functions are `$realtobits()` and `$bitstoreal()`.

The Verilog shell module that represents the C model can be instantiated in a design in the same way as any other Verilog module. For example:

```
module chip (...)  
    ...  
    scientific_alu u1 (a, b, opcode, result, excep, err);  
    ...  
endmodule
```

Creating a shell module to represent the C model is not mandatory. The PLI application could be called directly from any place in a Verilog design. However, there are important advantages to using a shell module to represent the C model:

- The shell module provides a simple method to encapsulate the C model.
- The shell module can be instantiated anywhere in a Verilog design hierarchy.
- The shell module can be instantiated any number of times in a Verilog design.
- The shell module can add Verilog HDL delays to the C model, which can accurately represent rise and fall delay delays, state-dependent delays, and timing constraints such as setup times.
- Delays within a shell module can be annotated using delay calculators or SDF files for additional delay accuracy for each instance of the shell module. Section 13.10 later in this chapter discusses how delays can be represented in the Verilog shell module.

13.5 Creating a combinational logic interface to a C model

In a combinational logic model, the outputs of the model continuously reflect the input values of the model. The inputs are asynchronous—when any input changes value, the model outputs are re-evaluated to reflect the input change.

The TF library works with the arguments of a system task. In the discussion of the Verilog shell module, in the previous section of this chapter, it was recommended that a system task be created to represent the C model interface, and that this task list all of the C model inputs and outputs as arguments. The reason for listing all inputs and outputs as task arguments is because this gives a PLI application easy access to these signals using the TF library.

For a combinational logic model, the PLI application which represents the interface to the C model must be called whenever one of the inputs to the C model changes value. There are two ways this can be accomplished:

- The simplest method is to use the Verilog HDL to invoke the *calltf routine* of the PLI application each time an input to the C model changes. This requires using a Verilog always procedure with a combinational logic sensitivity list. For the scientific ALU model example, the Verilog shell module would contain:

```
always @(a or b or opcode)
$scientific_alu(a, b, opcode, result, exception, error);
```

- An alternative method is to use the TF routines to invoke the *miscf routine* of the PLI application each time an argument to a system task changes. This requires using a Verilog initial procedure to instantiate the system task which represents the interface to the C model. For the scientific ALU model example, the Verilog shell module would contain:

```
initial
$scientific_alu(a, b, opcode, result, exception, error);
```

Examples of using each of these methods to represent a combinational logic interface to a C model are shown in sections 13.5.1 and 13.5.2.

13.5.1 Using the Verilog HDL to represent a combinational logic interface

When the Verilog HDL is used to invoke a system task whenever an input changes, the *calltf routine* of the PLI application will be invoked for each input change. In the example:

```
always @(a or b or opcode)
$scientific_alu(a, b, opcode, result, exception, error);
```

The *calltf routine* for `$scientific_alu` will need to:

- Read the values of the C model inputs from the system task arguments
- Call the C model and pass the input values
- Write the outputs of the C model onto the system task arguments

Example 13-3 illustrates a complete PLI application for `$scientific_alu`, using the *calltf routine* to read and modify the values of the system task.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.13/sci_alu_comb_calltf_tf.c
- Verilog shell module: Chapter.13/sci_alu_comb_calltf_shell.v
- Verilog test bench: Chapter.13/sci_alu_comb_calltf_test.v
- Verilog-XL results log: Chapter.13/sci_alu_comb_calltf_test.log

Example 13-3: combinational logic C model interface using a *calltf routine*

```
#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
/********************* Calltf routine: Serves as an interface between Verilog simulation
 * and the C model. Called whenever the C model inputs change value,
 * reads the input values, passes the values to the C model, and
 * writes the C model outputs into simulation.
 *****/
int PLIbook_ScientificALU_calltf()
{
    #define ALU_A      1 /* system task arg 1 is ALU A input */
    #define ALU_B      2 /* system task arg 2 is ALU B input */
    #define ALU_OP     3 /* system task arg 3 is ALU opcode input */
    #define ALU_RESULT 4 /* system task arg 4 is ALU result output */
    #define ALU_EXCEPT 5 /* system task arg 5 is ALU exception output */
    #define ALU_ERROR  6 /* system task arg 6 is ALU error output */

    double a, b, result;
    int    opcode, excep, err;

    /* Read current values of C model inputs from Verilog simulation */
    a      = tf_getrealp(ALU_A);
    b      = tf_getrealp(ALU_B);
    opcode = tf_getp(ALU_OP);

    /***** Call C model *****/
    PLIbook_ScientificALU_C_model(a, b, opcode, &result, &excep, &err);

    /* Write the C model outputs onto the Verilog signals */
    tf_putrealp(ALU_RESULT, result);
    tf_putp    (ALU_EXCEPT, excep);
    tf_putp    (ALU_ERROR,  err);

    return(0);
}

/********************* Checktf routine: Verifies that $scientific_alu() is used correctly.
 * Note: For simplicity, only limited data types are allowed for
 * task arguments. Could add checks to allow other data types.
 *****/
int PLIbook_ScientificALU_checktf()
{
    acc_initialize();

    if (tf_nump() != 6)
        tf_error("$scientific_alu requires 6 arguments");

    else {
```

```

if (tf_typep(ALU_A) != TF_READWRITEREAL)
    tf_error("$scientific_alu arg 1 must be a real variable\n");

if (tf_typep(ALU_B) != TF_READWRITEREAL)
    tf_error("$scientific_alu arg 2 must be a real variable\n");

if (tf_typep(ALU_OP) != TF_READONLY)
    tf_error("$scientific_alu arg 3 must be a net\n");
else if (tf_sizep(ALU_OP) != 4)
    tf_error("$scientific_alu arg 3 must be a 4-bit vector\n");

if (tf_typep(ALU_RESULT) != TF_READWRITEREAL)
    tf_error("$scientific_alu arg 4 must be a real variable\n");

if (tf_typep(ALU_EXCEPT) != TF_READWRITE)
    tf_error("$scientific_alu arg 5 must be a reg\n");
else if (tf_sizep(ALU_EXCEPT) != 1)
    tf_error("$scientific_alu arg 5 must be scalar\n");

if (tf_typep(ALU_ERROR) != TF_READWRITE)
    tf_error("$scientific_alu arg 6 must be a reg\n");
else if (tf_sizep(ALU_ERROR) != 1)
    tf_error("$scientific_alu arg 6 must be scalar\n");
}

return(0);
}

```

13.5.2 Using a *misctf* routine to represent a combinational logic interface

The `tf_asynchon()` routine provides a simple method of creating a combinational logic interface to a C model. This routine schedules asynchronous callbacks to the *misctf routine* whenever an argument of a system task changes value. The *misctf routine* can read the input values from the system task arguments, and pass the input values to the C model. The outputs of the C model are then passed back to the Verilog simulation by writing the results onto the system task arguments in the Verilog shell module. Chapter 12 presented how to use the `tf_asynchon()` routine.

By using *misctf routine* callbacks when a system task argument changes value, the *calltf routine* does not need to be invoked for every value change. In the Verilog shell module, the system task which represents the C model interface can be called from a Verilog initial procedure, so that the *calltf routine* is only invoked one time, instead of repeatedly. For example:

```

initial
    $scientific_alu(a, b, opcode, result, exception, error);

```

Using `tf_asynchronous()` is an easy way to accurately represent a combinational logic interface, but there are two cautions which must be observed when using this routine to represent an interface to a C model: First, the *misctf routine* is called for any task argument change, including changes to the C model outputs, and, second, the *misctf routine* is called automatically for many miscellaneous reasons which are not required by the C model interface. The *misctf routine* must filter out these extra calls.



TIP

The VPI and ACC libraries offer a more efficient method of representing a combinational logic interface to a C model. These libraries provide a means to have a PLI application called only when specific signals change value. By placing value change flags on just the inputs to a C model, the PLI application will only be called when an input changes value. However, the `tf_asynchronous()` routine allows an interface to be created very quickly and simply, which can be advantageous. Also, some simulators execute the TF routines more efficiently than the ACC and VPI routines, which, for these simulators, may make using `tf_asynchronous()` a more efficient method, despite the extra calls to the *misctf routine*.

The basic steps involved with using `tf_asynchronous()` and the *misctf routine* to implement a combinational logic interface are:

1. Create a PLI application system task to represent the interface between the Verilog shell module and the C model. The system task is invoked from the shell module, and all of the C model inputs and outputs are listed as arguments to the system task
2. In the *calltf routine* associated with the system task, invoke the `tf_asynchronous()` routine, to enable asynchronous callbacks to the *misctf routine*.
3. In the *misctf routine*, which is called whenever a system task argument changes value, read the values of all inputs and pass the values to the C model. The output values of the C model are returned to the same *misctf routine*, which then writes the values to the system task arguments that represent the outputs of the C model.

The following example implements a combinational logic interface for the scientific ALU C model. The *misctf routine* has checks to avoid calling the C model when it was called for reasons are not due to value changes on the inputs of the C model.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.13/sci_alu_comb_misctf_tf.c
- Verilog shell module: Chapter.13/sci_alu_comb_misctf_shell.v
- Verilog test bench: Chapter.13/sci_alu_comb_misctf_test.v
- Verilog-XL results log: Chapter.13/sci_alu_comb_misctf_test.log

Example 13-4: combinational logic C model interface using a *misctf routine*

```
#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
/*********************  

 * calltf routine: turns on asynchronous callbacks to the misctf  

 * routine whenever an argument to the system task changes value.  

 *******************/  

int PLIbook_ScientificALU_calltf()  
{  

    tf_asynchon();  

    return(0);  

}  

/*********************  

 * misctf routine: Serves as an interface between Verilog simulation  

 * and the C model. Called whenever the C model inputs change value,  

 * reads the input values, and passes the values to the C model, and  

 * writes the C model outputs into simulation.  

 *******************/  

int PLIbook_ScientificALU_misctf(int user_data, int reason, int  

paramvc)  
{  

#define ALU_A      1 /* system task arg 1 is ALU A input */  

#define ALU_B      2 /* system task arg 2 is ALU B input */  

#define ALU_OP     3 /* system task arg 3 is ALU opcode input */  

#define ALU_RESULT 4 /* system task arg 4 is ALU result output */  

#define ALU_EXCEPT 5 /* system task arg 5 is ALU exception output */  

#define ALU_ERROR   6 /* system task arg 6 is ALU error output */  

double a, b, result;  

int      opcode, excep, err;  

/* abort if misctf was not called for a task argument value change */  

if (reason != REASON_PARAMVC)  

    return(0);  

/* abort if task argument that changed was a model output */  

if (paramvc > ALU_OP) /* model outputs are after model inputs */  

    return(0);  

/* Read current values of C model inputs from Verilog simulation */  

a      = tf_getrealp(ALU_A);  

b      = tf_getrealp(ALU_B);  

opcode = tf_getp(ALU_OP);  

***** Call C model *****/  

PLIbook_ScientificALU_C_model(a, b, opcode, &result, &excep, &err);  

/* Write the C model outputs onto the Verilog signals */  

tf_putrealp(ALU_RESULT, result);  

tf_putp    (ALU_EXCEPT, excep);  

tf_putp    (ALU_ERROR, err);  

return(0);  

}
```

13.6 Creating a sequential logic interface to a C model

In a sequential logic model, the outputs of the model change synchronously with an input strobe, such as a positive edge of a clock. There may also be one or more asynchronous inputs, such as a reset signal. As with representing a combinational logic interface, a sequential logic interface can be represented using either the Verilog HDL or within the PLI application.

13.6.1 Using the Verilog HDL to represent a sequential logic interface

When the Verilog HDL is used to represent the sequential logic of the C model interface, the interface will be represented using a *calltf routine*. The Verilog HDL will invoke the system task at each clock cycle, which will result in the *calltf routine* being called synchronously at that time. The Verilog shell module will contain an *always* procedure with a sequential logic sensitivity list. For example:

```
always @(posedge clock)
$scientific_alu(clock, a, b, op, result, excep, err);
```

An asynchronous reset can be represented in the Verilog procedure as follows (this example models an active low reset):

```
always @(posedge clock or negedge rst)
$scientific_alu(clock, rst, a, b, op, result, excep, err);
```

When the Verilog HDL is used to synchronize the call to the PLI application to a clock, the *calltf routine* does not need to be any different than with combinational logic. When a clock change occurs, the *calltf routine* will be invoked, and all inputs to the C model will be read at that time. An example of this method is not shown here, since it is virtually the same as the example of using a *calltf routine* with the combinational logic that was listed previously in example 13-3.

13.6.2 Using the *misctf* routine to represent a sequential logic interface

The TF routines do not provide a convenient method of synchronizing activity to only one input change, such as the clock line. However, the asynchronous callbacks for any input change using *tf_asynchon()* can be used to create a synchronous, sequential logic interface. The process simply involves adding additional filters in the *misctf routine*, so that value changes on input signals other than the clock line are ignored. The basic steps involved with using *tf_asynchon()* to implement a synchronous sequential logic interface are very similar to implementing a combinational

logic interface. The one difference is that C model input values are only read and passed to the C model when the clock input changes.

Example 13-5 implements a sequential logic interface for the scientific ALU C model using the *misctf routine* and asynchronous callbacks for all task argument changes. This example ignores all task argument changes, except for change which represents the clock input.

Note that Verilog uses 4-state logic, so there are 12 possible transitions on the clock signal. Any transition will cause the *misctf routine* to be called. The following example checks that the value of the clock is a logic 1, with the assumption if the clock just changed value, and it is now a logic 1, it must have been a positive edge of the clock. By checking for a logic one, all negative going transitions on the clock line are ignored, as well as transitions from 0 to Z and 0 to X. The positive going transitions that this example will interpret as a positive edge of clock are transitions from 0 to 1, Z to 1, and X to 1. This filtering of transitions is not quite the same as the *posedge* keyword in the Verilog language, where 0 to Z and 0 to X are also treated as positive transitions.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.13/sci_alu_sequential_tf.c
- Verilog shell module: Chapter.13/sci_alu_sequential_shell.v
- Verilog test bench: Chapter.13/sci_alu_sequential_test.v
- Verilog-XL results log: Chapter.13/sci_alu_sequential_test.log

Example 13-5: sequential logic C model interface using TF routines

```
#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
/*********************  
 * calltf routine: turns on asynchronous callbacks to the misctf  
 * routine whenever an argument to the system task changes value  
 ********************/  
int PLIBook_ScientificALU_calltf()  
{  
    tf_asynchon();  
    return(0);  
}  
/*********************  
 * misctf routine: Serves as an interface between Verilog simulation  
 * and the C model. Called whenever the C model inputs change value,  
 * ignores all changes except a positive edge of clock, on the positive  
 * edge of clock, reads the input values, and passes the values to the  
 * C model, and writes the C model outputs into simulation.  
 ********************/
```

```
int PLIbook_ScientificALU_misctf(int user_data, int reason, int
paramvc)
{
    #define ALU_CLOCK 1 /* system task arg 1 is ALU clock input      */
    #define ALU_A     2 /* system task arg 2 is ALU A input          */
    #define ALU_B     3 /* system task arg 3 is ALU B input          */
    #define ALU_OP    4 /* system task arg 4 is ALU opcode input     */
    #define ALU_RESULT 5 /* system task arg 5 is ALU result output   */
    #define ALU_EXCEPT 6 /* system task arg 6 is ALU exception output */
    #define ALU_ERROR  7 /* system task arg 7 is ALU error output     */

    double a, b, result;
    int     opcode, excep, err, clock;

    /* abort if misctf was not called for a task argument value change */
    if (reason != REASON_PARAMVC)
        return(0);

    /* abort if task argument that changed was not the clock input */
    if (paramvc != ALU_CLOCK)
        return(0);

    /* Read current values of C model inputs from Verilog simulation */
    clock = tf_getp(ALU_CLOCK);
    if (clock != 1) /* abort if not a positive edge of the clock input */
        return(0);
    a      = tf_getrealp(ALU_A);
    b      = tf_getrealp(ALU_B);
    opcode = tf_getp(ALU_OP);

    /***** Call C model *****/
    PLIbook_ScientificALU_C_model(clock, a, b, opcode,
                                   &result, &excep, &err);

    /* Write the C model outputs onto the Verilog signals */
    tf_putrealp(ALU_RESULT, result);
    tf_putp    (ALU_EXCEPT, excep);
    tf_putp    (ALU_ERROR, err);

    return(0);
}
```

13.7 Synchronizing with the end of a simulation time step

Within a simulation, several inputs to the C model might change at the same moment in simulation time. The *calltf routine* or the *misctf routine* will be called for each input change, which means these routines may be called before all input value changes have occurred for that time step.

With a combinational logic interface, the *calltf routine* or *misctf routine* will be called for every input change. This is the correct functionality for combinational logic. At the completion of a simulation time step, the outputs from the C model represent the most current input values. However, by synchronizing the call to the C model with the end of the simulation time step in which changes occur, the multiple calls to the C model within a time step could be optimized to a single call.

With a sequential logic interface synchronized to a clock, when the *calltf routine* or *misctf routine* is called at a clock change, other input changes at that moment in simulation time may or may not have occurred. It may be desirable to ensure that the C model is not called until all inputs have their most current value for the time step in which the clock changes.

By using the `tf_synchronize()` routine, both combinational logic and sequential logic C model interfaces can be synchronized to the end of a current simulation time step. This is done by using the call to the *calltf routine* or *misctf routine* when a value change occurred, to schedule a synchronous callback to the *misctf routine* for the end of the current time step.

Note that `tf_synchronize()` does not guarantee that all changes in the current simulation time step have occurred. The `tf_synchronize()` routine schedules a read/write synchronization callback to the *misctf routine* after all known events in a time step have been processed by the simulator. However, other PLI applications may have also scheduled read/write synchronization callbacks, and these applications can schedule new events in the same simulation time step. The events caused by other PLI applications could affect the inputs of the C model after the `tf_synchronize()` callback has occurred. If the C model interface requires absolute assurance that no other inputs will change in the current simulation time step, the PLI application must use the `tf_rosynchronize()` read-only synchronization callback, instead of a read/write synchronization. In a read-only synchronization, the PLI application is not allowed to write return values into the C model using any of the TF put routines (or the `acc_set_value()` routine).

Example 13-6 modifies the combinational logic interface, which was presented previously in Example 13-4. This modified version schedules a *misctf routine* callback for two different reasons:

- At a value change of a task argument, a synchronize callback is scheduled at the end of the current time step.
- At a synchronize callback at the end of a simulation time step, the values of all C model inputs are read and passed to the C model.

This example uses the TF work area to store a flag to indicate when a synchronous callback to the *misctf routine* has already been scheduled for the current simulation time step. This flag prevents more than one synchronous callback to the *misctf routine* being requested in the same time step. Since the TF work area is unique for each instance of a system task, each instance of the C model will have a unique flag.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.13/sci_alu_synchronized_tf.c
- Verilog shell module: Chapter.13/sci_alu_synchronized_shell.v
- Verilog test bench: Chapter.13/sci_alu_synchronized_test.v
- Verilog-XL results log: Chapter.13/sci_alu_synchronized_test.log

Example 13-6: C model interface synchronized to the end of a time step

```
#include <stdio.h>      /* ANSI C standard I/O library */
#include "veriuser.h"    /* IEEE 1364 PLI TF routine library */
/*********************  
 * calltf routine: turns on asynchronous callbacks to the misctf  
 * routine whenever an argument to the system task changes value  
 *******************/  
  
int PLIBbook_ScientificALU_calltf()  
{  
    tf_asynchon();  
    tf_setworkarea(NULL); /* set work area to null */  
    return(0);  
}  
  
/*********************  
 * misctf routine: Serves as an interface between Verilog simulation  
 * and the C model. The misctf routine performs different operations  
 * depending on the reason it is called:  
 * - For a value change callback: schedules a callback to the misctf  
 *   application synchronized to the end of a time step. Only schedules  
 *   one callback for a time step.  
 * - For a synchronize callback: reads the input values, and passes  
 *   the values to the C model, and writes the C model outputs into  
 *   simulation.  
 *******************/
```

```
int PLIBook_ScientificALU_misctf(int user_data, int reason, int
paramvc)
{
    #define ALU_A      1 /* system task arg 1 is ALU A input          */
    #define ALU_B      2 /* system task arg 2 is ALU B input          */
    #define ALU_OP     3 /* system task arg 3 is ALU opcode input      */
    #define ALU_RESULT 4 /* system task arg 4 is ALU result output      */
    #define ALU_EXCEPT 5 /* system task arg 5 is ALU exception output   */
    #define ALU_ERROR   6 /* system task arg 6 is ALU error output       */

    double a, b, result;
    int     opcode, excep, err;

    /* check if misctf was called for a task argument value change */
    if (reason == REASON_PARAMVC) {
        /* abort if task argument that changed was a model output */
        if (paramvc > ALU_OP) /* model outputs are after model inputs */
            return(0);

        /* If the TF work area is null, then no misctf synchronize      */
        /* callback has been scheduled for this time step (the work area */
        /* is set to non-null by this routine, and is set to null by the */
        /* misctf after a synchronize callback is processed.           */
        if (tf_getworkarea() == NULL) {
            /* Schedule a synchronize callback to misctf for this instance */
            tf_synchronize();
            tf_setworkarea("1"); /* set work area to non-null */
        }
        return(0);
    }

    /* check if misctf was called for end-of-time step synchronize */
    if (reason == REASON_SYNCH) {
        /* Read current values of C model inputs from Verilog simulation */
        a      = tf_getrealp(ALU_A);
        b      = tf_getrealp(ALU_B);
        opcode = tf_getp(ALU_OP);

        /***** Call C model *****/
        PLIBook_ScientificALU_C_model(a, b, opcode, &result, &excep, &err);

        /* Write the C model outputs onto the Verilog signals */
        tf_putrealp(ALU_RESULT, result);
        tf_putp    (ALU_EXCEPT, excep);
        tf_putp    (ALU_ERROR, err);
        tf_setworkarea(NULL); /* set work area to null */
    }
    return(0);
}
```

13.8 Synchronizing with a future simulation time step

In certain C model applications, it may be necessary to synchronize C model activity with future simulation activity. The `tf_setdelay()` routine and variations of this routine can be used to schedule a call to the *misctf routine* for a specific amount of time in the future, relative to the current simulation time.

The `tf_getnextlongtime()` routine returns the future simulation time in which the next simulation event is scheduled to occur. This provides a way for a PLI application to synchronize activity for when the Verilog simulator is processing simulation events.

These TF routines for synchronizing with future simulation times were presented in more detail in Chapter 12.

13.9 Allocating storage within a C model

Special attention and care must be taken when a C model uses static variables or allocates memory.

The Verilog language can instantiate a model any number of times. Each instance of the Verilog shell module creates a unique instance of the system task which invokes the PLI interface to the C model. Therefore, the *calltf routine* and *misctf routine* which are invoked by a system task instance will both be unique to a task instance, and any memory which is allocated by the *calltf routine* and *misctf routine* will also be unique for each instance of the system task.

When a C model is represented as an independent program, multiple instances of the model are not a problem, as each instance will invoke a new process with unique storage for each process.

When the C model is represented as a C function, however, multiple instances of the model will share the same function. The C function must allow for the possibility of multiple instances, and provide unique storage for each instance.

Example 13-7 presents a latched version of the scientific ALU, which can store the result of a previous operation indefinitely. Example 13-8 presents a combinational logic interface to this latched model. This example interface is based on the interface method shown previously in example 13-4, which used the *misctf routine*. The same principles apply to the interface method which uses the *calltf routine* as the interface to the C model. This example allocates unique storage within the C model for each

instance of the C model. The instance pointer of the system task which represents the C model interface is used to identify which storage area belongs to which instance of the model. This instance pointer is obtained using `tf_getinstance()` and is passed to the C model as an input to the model function.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.13/sci_alu_latched_tf.c`
- Verilog shell module: `Chapter.13/sci_alu_latched_shell.v`
- Verilog test bench: `Chapter.13/sci_alu_latched_test.v`
- Verilog-XL results log: `Chapter.13/sci_alu_latched_test.log`

Example 13-7: scientific ALU C model with latched outputs

```
*****
 * Definition for a structure to store output values when the ALU is
 * latched. When enable is 1, the ALU returns the currently calculated
 * outputs, and when 0, the ALU returns the latched previous results.
 ****
#include <stdio.h>
typedef struct PLIbook_SciALUoutputs *PLIbook_SciALUoutputs_p;
typedef struct PLIbook_SciALUoutputs {
    char *instance_p; /* shows which task instance owns this space */
    double result;    /* stored result of previous operation */
    int    excep;
    int    err;
    PLIbook_SciALUoutputs_p next_ALU_outputs; /* next stack location */
} PLIbook_SciALUoutputs_s;

/* declare global stack pointer */
static PLIbook_SciALUoutputs_p ALU_outputs_stack = NULL;

*****
 * C model of a Scientific Arithmetic Logic Unit.
 * Latched outputs version.
 ****
#include <math.h>
#include <ERRNO.h>
void PLIbook_ScientificALU_C_model(
    int    enable,      /* input; 0 = latched */
    double a,           /* input */
    double b,           /* input */
    int    opcode,       /* input */
    double *result,     /* output from ALU */
    int    *excep,       /* output; set if result is out of range */
    int    *err,          /* output; set if input is out of range */
    char   *instance_p) /* input; pointer to system task instance */
```

```
{  
    PLIbook_SciALUoutputs_p ALU_outputs;  
  
    /* Locate the output storage in the stack for this model instance */  
    /* If no storage is found, then allocate a storage block and add */  
    /* the storage to the stack. */  
    ALU_outputs = ALU_outputs_stack; /* top-of-stack is in global var. */  
    while (ALU_outputs && (ALU_outputs->instance_p != instance_p))  
        ALU_outputs = ALU_outputs->next_ALU_outputs;  
  
    /* If no storage area found for this model instance, create one */  
    if (ALU_outputs == NULL) {  
        ALU_outputs =  
            (PLIbook_SciALUoutputs_p)malloc(sizeof(PLIbook_SciALUoutputs_s));  
        ALU_outputs->instance_p = instance_p; /* set owner of this space */  
        ALU_outputs->next_ALU_outputs = NULL;  
        ALU_outputs_stack = ALU_outputs; /* save new top-of-stack */  
    }  
  
    if (enable) { /* ALU is not latched, calculate outputs and store */  
        switch (opcode) {  
            case 0x0: ALU_outputs->result = pow      (a, b);      break;  
            case 0x1: ALU_outputs->result = sqrt     (a);         break;  
            case 0x2: ALU_outputs->result = exp      (a);         break;  
            case 0x3: ALU_outputs->result = ldexp   (a, (int)b); break;  
            case 0x4: ALU_outputs->result = fabs     (a);         break;  
            case 0x5: ALU_outputs->result = fmod     (a, b);      break;  
            case 0x6: ALU_outputs->result = ceil     (a);         break;  
            case 0x7: ALU_outputs->result = floor    (a);         break;  
            case 0x8: ALU_outputs->result = log      (a);         break;  
            case 0x9: ALU_outputs->result = log10   (a);         break;  
            case 0xA: ALU_outputs->result = sin      (a);         break;  
            case 0xB: ALU_outputs->result = cos      (a);         break;  
            case 0xC: ALU_outputs->result = tan      (a);         break;  
            case 0xD: ALU_outputs->result = asin    (a);         break;  
            case 0xE: ALU_outputs->result = acos    (a);         break;  
            case 0xF: ALU_outputs->result = atan    (a);         break;  
        }  
        ALU_outputs->err    = (errno == EDOM); /* arg out of range */  
        ALU_outputs->excep = (errno == ERANGE); /* result out of range */  
        errno = 0;                      /* clear the error flag */  
        if (ALU_outputs->err) ALU_outputs->result = 0.0;  
    }  
  
    /* return the values stored in the C model */  
    *result = ALU_outputs->result;  
    *err     = ALU_outputs->err;  
    *excep   = ALU_outputs->excep;  
  
    return;  
}
```

Example 13-8: combinational logic interface to latched scientific ALU C model

```
#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
/*********************  

 * calltf routine: turns on asynchronous callbacks to the misctf  

 * routine whenever an argument to the system task changes value  

*******************/  

int PLIBook_ScientificALU_calltf()  
{  

    tf_asynchon();  

    return(0);  

}  

/*********************  

 * misctf routine: Serves as an interface between Verilog simulation  

 * and the C model. Called whenever the C model inputs change value,  

 * reads the input values, and passes the values to the C model, and  

 * puts the C model outputs into simulation. Passes the instance  

 * pointer of the Verilog system task which represents the C model  

 * to serve as a unique flag within the C model.  

*******************/  

int PLIBook_ScientificALU_misctf(int user_data, int reason, int  

paramvc)  
{  

#define ALU_ENABLE 1 /* system task arg 1 is ALU enable input */  

#define ALU_A 2 /* system task arg 2 is ALU A input */  

#define ALU_B 3 /* system task arg 3 is ALU B input */  

#define ALU_OP 4 /* system task arg 4 is ALU opcode input */  

#define ALU_RESULT 5 /* system task arg 5 is ALU result output */  

#define ALU_EXCEPT 6 /* system task arg 6 is ALU exception output */  

#define ALU_ERROR 7 /* system task arg 7 is ALU error output */  

    double a, b, result;  

    int opcode, excep, err, enable;  

    char *instance_p;  

/* abort if misctf was not called for a task argument value change */  

if (reason != REASON_PARAMVC)  

    return(0);  

/* abort if task argument that changed was a model output */  

if (paramvc > ALU_OP) /* model outputs are after model inputs */  

    return(0);  

    a = tf_getrealp(ALU_A);  

    b = tf_getrealp(ALU_B);  

    opcode = tf_getp(ALU_OP);  

/* Obtain the instance pointer for this system task instance */  

    instance_p = tf_getinstance();  

***** Call C model *****
```

```
PLIbook_ScientificALU_C_model(enable, a, b, opcode,
                               &result, &excep, &err, instance_p);

/* Write the C model outputs onto the Verilog signals */
tf_putstralp(ALU_RESULT, result);
tf_putp    (ALU_EXCEPT, excep);
tf_putp    (ALU_ERROR, err);

return(0);
}
```

13.10 Representing propagation delays in a C model

Propagation delays from an input change to an output change in a C model can be represented in two ways:

- Using delays in the PLI interface.
- Using delays in the Verilog shell module.

Delays in the PLI interface are represented by specifying a delay value with the `tf_strdelputp()`, `tf_strlongdelputp()` and `tf_sturrealdelputp()` routines, which write values onto the system task arguments at a future simulation time step. Either inertial or transport event propagation can be represented, depending on the requirements of the C model. However, using the `tf_strdelputp()` and related routines has disadvantages. The logic values in the PLI application must be converted to strings, in order to put the value onto an argument. This conversion to a string, which the simulator must then convert into a Verilog logic value, is not efficient for simulation run-time performance. Also, the TF routines do not offer a great deal of flexibility on creating delays which are different for each instance of a model, representing minimum, typical and maximum delays, different delays for rise and fall transitions, or annotating delays using delay calculators or SDF files.

C model propagation delays can also be represented using the pin-to-pin path delays in the Verilog shell module. This method provides the greatest amount of flexibility and accuracy in modeling propagation delays. All path delay constructs can be used, as well and Verilog timing constraints.

Example 13-9 illustrates adding pin-to-pin path delays to the scientific ALU shell module.

NOTE

Some Verilog simulators restrict the use of pin-to-pin path delays and SDF delay back annotation to Verilog models which are represented with Verilog primitives and net data types. To use path delays on a C model with these simulators, buffers must be added to all input and output ports, with net data types connected to the inputs and outputs of these buffers. Example 13-9 illustrates using buffers on all input and output ports of the scientific ALU shell module.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.13/sci_alu_with_delays_tf.c
- Verilog shell module: Chapter.13/sci_alu_with_delays_shell.v
- Verilog test bench: Chapter.13/sci_alu_with_delays_test.v
- Verilog-XL results log: Chapter.13/sci_alu_with_delays_test.log

Example 13-9: scientific ALU Verilog shell module with pin-to-pin path delays

```
'timescale 1ns / 100ps
module scientific_alu(a_in, b_in, opcode_in,
                      result_out, exception, error);
  output [63:0] result_out;
  output      exception, error;
  input  [63:0] a_in, b_in;
  input   [3:0] opcode_in;

  wire  [63:0] result_out, result_vector;
  wire  [63:0] a_in, a_vector;
  wire  [63:0] b_in, b_vector;
  wire   [3:0] opcode_in, opcode_vector;
  wire      exception, error;

  reg      exception_reg, error_reg;
  real    a, b, result; // real variables used in this module

  // convert real numbers to/from 64-bit vector port connections
  assign result_vector = $realtobits(result);
  always @(a_vector) a = $bitstoreal(a_vector);
  always @(b_vector) b = $bitstoreal(b_vector);

  //call the PLI application which interfaces to the C model
  initial
    $scientific_alu(a, b, opcode_vector,
                    result, exception_reg, error_reg);

  specifies
    (a_in, b_in *> result_out, exception, error) = (5.6, 4.7);
    (opcode_in  *> result_out, exception, error) = (3.4, 3.8);
  endspecify
```

```
// add buffers to all ports, with nets connected to each buffer
// (this example uses the array of instance syntax in the
// from the IEEE 1364-1995 Verilog standard
buf result_buf[63:0] (result_out, result_vector);
buf excep_buf        (exception, exception_reg);
buf error_buf         (error,      error_reg);

buf a_buf[63:0]        (a_vector, a_in);
buf b_buf[63:0]        (b_vector, b_in);
buf opcode_buf[3:0]    (opcode_vector, opcode_in);

endmodule
```



NOTE The preceding example uses the array of instances construct from the IEEE 1364-1995 Verilog standard. This construct was not supported by some Verilog simulators at the time this book was written. A work around is to create a separate instance of a buf primitive for each bit or each vector connected to a module port.

13.11 Summary

This chapter has presented a few ways in which the TF library can be used to interface a C language model with Verilog simulations. By creating a shell module which contains the system task that invokes the C model interface, the C model can be used in a Verilog design, just as any other Verilog module. The interface between the shell module and the C model can done through a *calltf routine*, using the Verilog HDL to control when the *calltf routine* is invoked. The interface can also be done through the *miscif routine*, using the TF library to control when the *miscif routine* is invoked. The *tf_asynchon()* routine provides a simple means of using the *miscif routine* to pass input changes to a C model between the C model and the Verilog shell module. The TF routines to read and modify logic values allow information to be exchanged in a variety of formats.

CHAPTER 14

How to Use the ACC Routines

This chapter introduces the ACC portion of the PLI standard, and shows how to use the ACC routines to access information within a simulation data structure. Two complete PLI applications, `$show_all_nets` and `$show_all_signals`, will be created in this chapter, to illustrate how the ACC routines work. The remaining chapters in this part of the book then build on the principles presented in this chapter by explaining the routines within the ACC library in much more detail.

The concepts presented in this chapter are:

- An overview of how ACC routines work
- Advantages of the ACC library
- Creating a complete PLI application using the ACC library
- Obtaining handles to Verilog HDL objects
- Accessing properties of Verilog HDL objects
- Reading values of Verilog HDL objects

14.1 Specification of \$show_all_nets and \$show_all_signals

To show how the ACC routines are used, two PLI applications will be created. These examples will be built up, one step at a time, as this chapter progresses.

The first example presented is an application called `$show_all_nets`. The usage of this application is:

```
$show_all_nets(<module_instance_name>) ;
```

This PLI application will:

1. Access the first argument of the system task, which is the name of a module instance.
2. Print the hierarchical path and name of that module, along with the current simulation time.
3. Search for all net signals in the module, and print the data type and current logic value of each net.

This chapter will first illustrate a *checktf routine* for `$show_all_nets`, which verifies that the argument provided as an input is a valid module instance name. Then a *calltf routine* will be created to perform the functionality of the system task.

The second example is a PLI application called `$show_all_signals`. This application prints the current value of all net, reg and variable data types in a module. The usage of this application is:

```
$show_all_signals(<module_instance_name>) ;
```

To illustrate some additional ways to use the ACC routines, two enhancements to the `$show_all_signals` example will be presented. These are:

- Use no argument or a null argument to `$show_all_signals`, to represent the module instance containing the `$show_all_signals` system task.
- Allow multiple arguments to `$show_all_signals`, so the values of signals in several modules can be printed with one call to `$show_all_signals`.

14.2 The ACC routine library

“ACC” stands for “access”, and the library of ACC routines are often referred to as *access routines*. The ACC routines are the second of three primary generations of the PLI functions (the TF routines were the first generation, and the VPI routines are the newest). The primary purpose of the ACC routines is to provide a PLI application access to the internal data structures of a simulation. The ACC routines provide a consistent layer between a user’s PLI application and the underlying data structures of a simulation. The PLI application does not need to know the specifics about how the simulator stores its data, and the same PLI application will work with many different simulators.

The ACC routines treat Verilog HDL constructs as *objects*, and many of the ACC routines provide ways to locate any specific object or type of objects within a simulation data structure. Other ACC routines can then read and modify information about each object.

The ACC library can be divided into five basic groups of routines:

- *handle* routines obtain a handle for one specific Verilog HDL object.
- *next* routines locate and return handles for a specific type of Verilog object.
- *fetch* routines access information about an object.
- *set* routines modify information about an object.
- *miscellaneous* routines perform a variety of operations.

The library of ACC routines is defined in a C header file called **acc_user.h**, which is part of the IEEE 1364 standard. This header file also defines a number of C constants and C structures used by the ACC routines. All PLI applications that use ACC routines must include the **acc_user.h** file.

The ACC library is designed to work with the TF library. An example of including the header files for these two libraries is:

```
#include "veriuser.h"    /* IEEE 1364 TF PLI library */
#include "acc_user.h"    /* IEEE 1364 ACC PLI library */
```

The ACC routines are intended to complement and expand the capabilities of the TF routines, rather than replace the TF routines. While there is some overlap in the functionality provided by the two libraries, there are a number of capabilities which are unique to the TF library—for example: `tf_numP()`, `io_printf()`, `tf_setworkarea()` and `tf_getworkarea()`.

NOTE → The VPI library was designed to replace both the TF and ACC libraries with a more concise, more robust, and more versatile procedural interface. The IEEE 1364 standard includes the TF and ACC libraries, in order to provide backward compatibility and portability of older PLI applications with modern Verilog simulators. The official policy of the IEEE 1364 standards committee is that, as improvements and enhancements are added to the Verilog language, only the VPI library of the PLI will be expanded to support those new features. The TF and ACC libraries will be maintained, but not enhanced, in future versions of the IEEE 1364 standard.

14.3 Advantages of the ACC library

The ACC routines provide direct access to much of a Verilog simulation data structure. This is in contrast to the TF library, which only provides access to the arguments of system tasks and system functions. This direct access allows a PLI application to more fully analyze and interact with a Verilog simulation.

Advantages of ACC routines, compared to TF routines

An important advantage the direct access which ACC routines provide can be seen in the example PLI applications presented in this chapter. These applications print the current logic values of all signals in a module. To implement this functionality with TF routines, every signal name would need to be listed as an argument to the system task. If a module had dozens of signals, listing each one as an argument would be very awkward. Using ACC routines, only one argument needs to be passed to the system task—the name of a module instance. From this starting point, the ACC routines can find all the signals in the module, and directly access the information of those signals.

 **TIP** Using ACC routines can have a negative impact on simulation run-time performance. Using only TF routines can improve the performance of some simulators.

The ACC routines provide arbitrary, run-time access to information within the simulation data structure. The information that ACC routines will be access during simulation cannot be predicted at compile time. This can prevent a compiler from effectively optimizing the simulation data structure for maximum simulation performance. The TF routines restrict access to the arguments of a system task or system function. This restricted access can be determined at compile time, and can therefore be optimized by the compiler.

 **NOTE** Some simulators, such as the Synopsys® VCST™ simulator, require special configurations or invocation options to enable the ACC routines. Refer to Appendix A for more information about specific simulators.

Advantages of ACC routines, compared to VPI routines

The VPI library is designed to replace the ACC library, and also extend the capabilities of the PLI. In the author's opinion, the VPI library is a better choice for creating new PLI applications. However, there are some advantages which the ACC library offers, compared to the VPI library.

First and foremost, the ACC library is supported by virtually every major Verilog simulator, whereas, at the time this book was written, the VPI standard was only supported by a few simulators. The more widespread support of ACC routines makes a PLI application portable to many more simulators and engineering environments.

Second, it is often much easier and faster to develop PLI applications using the ACC library, compared to the VPI library. The ACC library is a much larger library, and often has pre-defined routines which take care of much of the work that a PLI application needs to accomplish. In the smaller VPI library, the PLI application developer must code much of the corresponding functionality by hand. However, the larger ACC library also tends to be less efficient for simulation run-time performance, and, when using ACC routines, PLI application developers may tend to implement poorly structured C code, which also impacts simulation performance and becomes difficult to maintain.

14.4 Verilog HDL objects

The ACC routines treat Verilog HDL constructs as *objects*, and many of the ACC routines provide ways to locate any specific object or type of object within a simulation data structure. Other ACC routines can then read and modify information about each object. The simple Verilog HDL example which follows has several objects which can be accessed by the library of ACC functions.

```
module test;
    reg [1:0] test_in;
    wire [1:0] test_out;
    buf2 u1 (test_in, test_out);
    initial
        begin
            test_in = 3;
            #50 $display("in=%d, out=%d", test_in, test_out);
        end
    endmodule

module buf2 (in, out);
    input [1:0] in;
    output [1:0] out;
    wire [1:0] in, out;
    buf #5 n0 (out[0], in[0]);
    buf #7 n1 (out[1], in[1]);
endmodule
```

In this Verilog HDL example, the objects that a PLI application can access include:

- A top-level module, with the definition name “*test*”. Within this module are:
 - A reg signal, with a vector size of 2 and the name “*test_in*”. The signal will have a logic value which can be read and modified by the PLI application.
 - A wire net, with a vector size of 2 and the name “*test_out*”. The net reflects a resolved logic value which can be read by the PLI application.
 - A module instance, with the definition name “*buf2*” and the instance name “*u1*”. Within this module are:
 - Two ports, with the names “*in*” and “*out*”. Each port has a vector size and direction.
 - Two wire nets, with vector sizes and names. The nets reflect a resolved logic value which can be read by the PLI application.
 - Two primitive instances, with the definition name “*buf*” and the instance names “*n0*” and “*n1*”. Each primitive has a delay value which can be read and modified by the PLI.
 - Terminals on each primitive instance, with bit-selects of specific nets connected to the terminals, such as **out[0]** and **in[0]**.

14.4.1 The ACC handle data type

The ACC routines use a special data type, called a *handle*, to access Verilog HDL objects. The handle data type is defined in the ACC library (in the acc_user.h header file). The declaration type for variables to store a handle is **handle** (spelled with all lower case letters). An example declaration for two handle variables is:

```
handle module_handle, net_handle;
```

There are more than 45 ACC routines that locate objects within a simulation data structure and return handles for the objects. Other ACC routines are used to access information about an object, using the object’s handle as a reference point. The information that can be accessed depends on the type of the object, but might include the object’s name and current logic value.

The object oriented method of accessing information used by the ACC routines is very similar to the object oriented method used by the VPI portion of the PLI standard, which was presented in Part One of this book. The only real difference is that the ACC routines have a more limited list of what Verilog HDL constructs are considered objects. For example, Verilog procedures and procedural statements are not objects in the ACC library, and therefore the ACC routines cannot access the proce-

dural portions of a Verilog design. Verilog memory arrays are also not objects, and cannot be accessed. In the VPI library, virtually everything that exists in the Verilog language is considered an object.

NOTE

Do not share handles between VPI routines and ACC routines! The VPI routines in the PLI standard also use the concept of a handle for referencing Verilog objects. The IEEE 1364 standard does not guarantee that a handle which is obtained with the VPI library will be the same as a handle which is obtained with the ACC library.

14.5 ACC handle routines

The `$show_all_nets` application shown in this chapter will need to obtain a handle to the first system task/function argument, in order to access all signals within a module.


ACC **handle** routines return a handle for a single object. ***These ACC routines are object-specific.*** There is a different routine for each type of object that can be accessed using handle routines. Some examples of ACC handle routines are:

- `acc_handle_tfarg()` obtains a handle for an object named in a system task/function argument.
- `acc_handle_terminal()` obtains a handle for an object connected to the terminal of a primitive.
- `acc_handle_port()` obtains a handle for a module port.

There are 25 ACC handle routines. This chapter introduces these routines. More details on these routines is presented in the next chapter.

Using ACC handle routines

A system task or system function can have any number of arguments, including none. The arguments are numbered from left to right, starting with argument number 1. In the following example:



- Task/Function argument number 1 is a string, with the value “vectors.pat”.
- Task/Function argument number 2 is a signal, with the name `input_bus`.

For the `$show_all_nets` application, the first system task argument will be a module instance name. For example:

```
TF arg #1
always @(posedge clock)
    $show_all_nets(top.i1);
```

The ACC handle routine that is used to obtain a handle for a system task/function argument is `acc_handle_tfarg()`. The syntax of this routine is:

handle acc_handle_tfarg (n)

<i>int n</i>	position number of a PLI system task/function argument.
<i>handle tfinst</i>	handle for an instance of a PLI system task/function.

This routine returns a handle for the object named as an argument in the system task/function which called the PLI application. Arguments are numbered from left to right, beginning with 1.

A handle for the module instance that is named in the first system task argument can be obtained, using the following C code:

```
handle tfarg_handle;
tfarg_handle = acc_handle_tfarg(1);
```

14.6 ACC next routines

The ACC library provides a set of routines to make it easy to access all occurrences of a specific type of object. These routines are referred to as *ACC next routines*. As with the ACC handle routines, *the ACC next routines are object-type specific*, so there is a different ACC next routine for most types of Verilog objects which the ACC routines can access. A few examples are:

- `acc_next_net()` obtains handles for all nets within a module.
- `acc_next_port()` obtains handles for all ports within a module.
- `acc_next_primitive()` obtains handles for all primitive instances within a module.

There are 22 ACC next routines. How these next routines are used is presented in this chapter, and the full list and syntax of the routines is presented in the next chapter.

There are two important terms used with ACC next routines:

- **target objects** are the type of objects for which the ACC next routine will obtain handles. For `acc_next_net()`, the target objects will be Verilog nets.
- **reference objects** are where the ACC next routine will search for the target objects. For example, to find all nets within a module, the reference object is the module.

Most ACC next routines require two inputs:

1. A handle for a reference object.
2. A handle to the previous target object found.

Using ACC next routines

The `$show_all_nets` application will need to access all nets within a module. The `acc_next_net()` routine will be used to access these nets. The syntax for this routine is:

handle acc_next_net (module, prev_net)

handle module handle for a module.

handle prev_net handle for the previous net found; initially `null`.

The reference object for this routine is a handle for a module instance, and the target handle will be for a net. The basic usage of the routine is:

```
next_net_handle = acc_handle_net(module_handle, previous_net_handle)
```

The ACC next routines return the target object handles one handle at a time. In order to locate all of a specific object, the next routine must be placed in a loop. All ACC next routines follow the same rules for how the target object handles are retrieved:

- To locate the first of the target objects within a reference object, the handle for the previous target found must be set to `null`.
- To locate the next target object within a reference object, the previous target handle found must be set to the previous target found.
- When the ACC next routine cannot find any more of target objects, the routine returns a `null`. This return value can be used to terminate the loop.

The `callif` routine for the `$show_all_nets` application will need to find all net signals in a module instance. This can be done using `acc_next_net()`, as follows:

```

handle mod_h, net_h;
mod_h = acc_handle_tfarg(1);
net_h = null; /* initialize the target handle to null */
while ( (net_h = acc_next_net(mod_h, net_h)) != null) {
    /* perform desired operations on the net handle */
}

```

The **null** used in the above example is defined in the IEEE 1364 standard *acc_user.h* file as a long 0. Another common C coding style is to terminate the loop when the value assigned to the target handle is 0, instead of explicitly comparing the assigned value to **null**. For example:

```

net_h = null; /* initialize the target handle to null */
while ( net_h = acc_next_port(mod_h, net_h) ) {
    /* perform desired operations on the port handle */
}

```

NOTE → The **null** (all lower case letters) that is used by ACC routines is defined in the *acc_user.h* ACC library file. This is not the same null as the standard C language **NULL** (all capital letters) defined in *stdlib.h* library.

14.7 Accessing object types and fulltypes

Every Verilog object which can be accessed by ACC routines has a *type* property and a *fulltype* property. These properties identify what Verilog object is referenced by a Verilog handle.

acc_fetch_type() retrieves the *type* property of an object, which identifies the general type of an object. The syntax of this routine is:

int acc_fetch_type (object)
handle object handle for an object.

The type property is an integer constant, such as *accModule*, *accPort*, *accNet*, *accPrimitive*, etc. This property can be used many different ways—one common usage is to verify that a handle which was obtained references the type of object expected. For example, the *\$show_all_nets* application requires that the first task/function argument be a module instance. The PLI application could verify that the argument is correct, using the following code fragment:

```
handle tfarg_handle;
tfarg_handle = acc_handle_tfarg(1);
if (acc_fetch_type(tfarg_handle) != accModule)
    /* report error that argument is not correct */
```

The `acc_fetch_fulltype()` routine retrieves the *fulltype* property of an object.

int acc_fetch_fulltype (object)

handle object handle for an object.

The fulltype property provides more detailed information about an object. For example, if a PLI application has obtained a handle for a Verilog net, the constant returned for `acc_fetch_type(net_handle)` is **accNet**, while the constants returned for `acc_fetch_fulltype(net_handle)` include **accWire**, **accWor**, **accWand**, etc. For a Verilog module, `acc_fetch_type(module_handle)` returns the constant **accModule**, and `acc_fetch_fulltype(module_handle)` returns one of the constants: **accTopModule**, **accModuleInstance**, **accCellInstance**.

The type and fulltype properties are represented by constants with integer values. Printing the type of an object directly would print the integer value of the constant, not the name of the constant. A useful routine for debugging problems in a PLI application is **acc_fetch_type_str()**. This routine takes a type or fulltype constant value as its input, and returns a pointer to a string which contains the actual name of the constant. The syntax of this routine is:

char *acc_fetch_type_str (type)

int type type or fulltype constant.

An example of using the `acc_fetch_type_str()` routine to print the name of a type constant is:

```
int object_type;
object_type = acc_fetch_type(tfarg_handle);
if (object_type != accModule) {
    tf_error("Tfarg type of %s is illegal.\n",
            acc_fetch_type_str(object_type));
}
```

14.8 Accessing the names of objects

Many Verilog objects have one or more ***name*** properties, which can be accessed using ACC routines.

The name properties which an object can have are:

- A ***local name***. For objects such as nets, the local name is the *declaration name* of the object within a module. Note that in Verilog, a net can be implicitly declared by simply referencing the name. For a module or primitive, the local name is the *instance name* within the module that the module or primitive is instantiated.
- A ***hierarchical path name***, which is the Verilog HDL design hierarchy path to an object, starting with the top of the design hierarchy.
- A ***definition name***, which is the *definition name* of a Verilog module or primitive.

The following Verilog HDL source code fragment illustrates the difference between *name*, *full name* and *definition name*.

```
module test;
    wire a, b, ci, sum, co;
    addbit u1 (a, b, ci, sum, co);
endmodule
```

local name: "u1"
 full name: "test.u1"
 definition name: "addbit"


```
module addbit (a, b, ci, sum, co);
    input a, b, ci;
    output sum, co;
    wire a, b, ci, sum, co;
    xor g1 (n1, a, b);
    xor #2 g2 (sum, n1, ci);
    and g3 (n2, a, b);
    and g4 (n3, n1, ci);
    or #2 g5 (co, n2, n3);
endmodule
```

local name: "sum"
 full name: "test.u1.sum"

local name: "g1"
 full name: "test.u1.g1"
 definition name: "xor"

The routines which retrieve an object's name are **acc_fetch_name()**, **acc_fetch_fullname()** and **acc_fetch_defname()**. The syntax of these three routines is:

char *acc_fetch_name (object)

handle object handle for an object.

char *acc_fetch_fullname (object)

handle object handle for an object.

char *acc_fetch_defname (object)

handle object handle for a module or primitive instance.

Each of these routines retrieves the object's name into a temporary string buffer, and returns a pointer to the string. Other ACC routines which return pointers to strings also share this string buffer. This temporary buffer is limited in size, and when the buffer is full, it wraps around to the beginning of the buffer. The buffer can hold multiple strings, but, when it is full, previous strings will be overwritten. A PLI application should use the string pointer returned by an ACC routine immediately. After another call is made to an ACC routine which retrieves a string, there is no guarantee that the first string pointer will still be valid. If a string needs to be preserved, the PLI application should copy the string into application-allocated storage space. Following are two examples of using strings in the PLI.

Read a string and use it immediately:

```
char *string_p;      /* string pointer only, no storage */
string_p = acc_fetch_name(net_handle);
io_printf("string_p points to %s\n", string_p);
```

Read a string and copy it to application-allocated storage for later use:

```
char *string_p;      /* string pointer only, no storage */
char *string_keep;   /* another string pointer */
string_p = acc_fetch_name(net_handle);
string_keep = malloc(strlen(string_p)+1);
strcpy(string, string_p); /* save string for later use */
```

14.9 Reading the logic values of Verilog objects

The ACC routine **acc_fetch_value()** retrieves the value of Verilog objects which contain a logic value. The Verilog language uses 4-state logic values, comprising logic 0, 1, Z and X. The **acc_fetch_value()** routine automatically converts Verilog 4-state logic into various C data types for representation in PLI applications. The simplest way to represent 4-state logic in C is to use character strings, and this is the

method that is used in the `$show_all_nets` application. Chapter 16 presents reading and writing Verilog logic values in more detail.

`char *acc_fetch_value (object, format_str, value)`

<code>handle object</code>	handle for a net or register.
<code>char *format_str</code>	character string controlling the radix of the retrieved value; must be “%b”, “%o”, “%d”, “%h”, “%v” or “%%”.
<code>p_acc_value value</code>	pointer to an application-allocated <code>s_acc_value</code> structure to receive the value as an aval/bval pair. Only used if <code>format_str</code> is “%%”.

This routine has three inputs, but only the first two are used when retrieving values as a string. The third input can be set to `null`, or left unspecified if it is not used.

The format string controls how the Verilog logic value should be represented in the C string. A format of “%b” indicates the logic value should be represented using binary numbers using the characters (‘0’, ‘1’, ‘z’, and ‘x’). A “%h” format indicates the value should be represented using hexadecimal numbers, using the characters (‘0’ through ‘F’, ‘z’, and ‘x’). A “%o” format indicates an octal representation, and a “%d” format represents a decimal representation. Other formats are available, which are discussed in Chapter 16.

Once a handle for a net has been obtained, the value for the net can be retrieved and printed as a string, using `acc_fetch_value()`, as follows:

```
io_printf(" net %s value is %s\n",
          acc_fetch_name(net_h),
          acc_fetch_value(net_h, "%b", null));
```



TIP Using C strings to represent 4-state logic is a simple method for reading and printing a Verilog logic value. However, the automatic conversion from Verilog values to C strings can be very expensive for the run-time performance of a PLI application. If a PLI application will access a large number of values, or if the application will be called many times during a simulation, it is better to use a more efficient format for reading values. Chapter 16 presents all the formats for reading logic values that are available using ACC routines, and discusses performance considerations.

14.10 A complete PLI application using ACC routines

Example 14-1 lists a complete *checktf routine* and *calltf routine* for the *\$show_all_nets* PLI application. Note the mixture of ACC routines and TF routines. These two libraries are designed to complement each other. Retrieving the current simulation time, for example, is done using *tf_strgettime()*. There is no routine in the ACC library to retrieve the current simulation time.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.014/show_all_nets_acc.c
- Verilog test bench: Chapter.014/show_all_nets_test.v
- Verilog-XL results log: Chapter.014/show_all_nets_test.log

Example 14-1: \$show_all_nets — using ACC routines in a PLI application

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/*********************  
 * checktf application  
*****  
int PLIbook_ShowNets_checktf()  
{  
    acc_initialize();  
    if (tf_nump() != 1)  
        tf_error("$show_all_nets must have 1 argument.");  
    else if (tf_typep(1) == TF_NULLPARAM)  
        tf_error("$show_all_nets arg cannot be null.");  
    else if (acc_fetch_type(acc_handle_tfarg(1)) != accModule)  
        tf_error("$show_all_nets arg must be a module instance.");  
    acc_close();  
    return(0);  
}  
*****  
 * calltf application  
*****  
int PLIbook_ShowNets_calltf()  
{  
    handle module_handle, net_handle;  
    acc_initialize();  
    module_handle = acc_handle_tfarg(1);  
    io_printf("\nAt time %s, nets in module %s (%s):\n",  
             tf_strgettime(),  
             acc_fetch_fullname(module_handle),  
             acc_fetch_defname(module_handle));  
    net_handle = null; /* start with known value for target handle */  
    while (net_handle=acc_next_net(module_handle,net_handle)) {
```

```

    io_printf(" %-13s %-13s value is %s (hex)\n",
              acc_fetch_type_str(acc_fetch_fulltype(net_handle)),
              acc_fetch_name(net_handle),
              acc_fetch_value(net_handle, "%h", null));
}
acc_close();
return(0);
}

```

Example 14-2, which follows, lists a simple Verilog HDL design to test `$show_all_nets`. Figure 14-1 which follows shows the output of running simulation with this test design.

Example 14-2: `$show_all_nets` — Verilog HDL test case for the PLI application

```

`timescale 1ns / 1ns
module top;
  reg [2:0] test;
  tri [1:0] results;

  addbit i1 (test[0], test[1], test[2], results[0], results[1]);

  initial
    begin
      test = 3'b000;
      #10 test = 3'b001;

      #10 $show_all_nets(top);
      #10 $show_all_nets(i1);

      #10 $stop;
      #10 $finish;
    end
endmodule

/** A gate level 1 bit adder model ***/
`timescale 1ns / 1ns
module addbit (a, b, ci, sum, co);
  input a, b, ci;
  output sum, co;

  wire a, b, ci, sum, co, n1, n2, n3;

  xor (n1, a, b);
  xor #2 (sum, n1, ci);
  and (n2, a, b);
  and (n3, n1, ci);
  or #2 (co, n2, n3);
endmodule

```

Figure 14-1: `$show_all_nets` — simulation results, using Verilog-XL

The screenshot shows the Cadence Verilog-XL Turbo NT simulation interface. The menu bar includes File, Control, Set, Show, Select, Tools, Options, and Help. The toolbar has icons for play, stop, and other simulation controls. The scope panel shows 'Scope: top' and 'Objects:'. The tools panel includes icons for waveform, logic, and memory. The status bar shows 'Time: 40'. The main window displays Verilog code and its simulation results.

```

22      #10 test = 3'b001;
23
24      #10 $show_all_nets(top);
25      #10 $show_all_nets(i1);
26
27      #10 $stop;
28      #10 $finish;
29  end
30 endmodule

```

Scope top Subscopes

```

At time 20, nets in module top (top):
accTri      results      value is 1 (hex)

At time 30, nets in module top.i1 (addbit):
accWire     a           value is 1 (hex)
accWire     b           value is 0 (hex)
accWire     ci          value is 0 (hex)
accWire     sum         value is 1 (hex)
accWire     co           value is 0 (hex)
accWire     n1           value is 1 (hex)
accWire     n2           value is 0 (hex)
accWire     n3           value is 0 (hex)
L27 "show_all_nets_test.v": $stop at simulation time 40
Type ? for help
C1 >
1
Ready

```

14.11 Accessing handles for reg and variable data types

The Verilog HDL defines two general data type groups, *nets* and *registers*. The register data type group includes the Verilog keywords **reg**, **integer**, **time** and **real**.

In the PLI, the term *register* is not used to represent data types. Instead, the PLI treats the **reg** data type as a unique object, and groups the **integer**, **time** and **real** data types into an object class called *variables*.

Most Verilog HDL objects which can be accessed by the ACC routines have either an object-specific *handle* routine, or an object-specific *next* routine to access a specific type of object. However, there are no object-specific routines for accessing the **reg** and **variable** data types. Instead, there is a generic ACC *next* routine which can be used to obtain handles for a number of different types of objects.

acc_next() is used to obtain handles for multiple types of objects. The syntax of this routine is:

handle acc_next(type_list, scope, prev_object)

int *type_list static array with a list of type or fulltype constants.

handle scope handle for the scope in which to scan for objects.

handle prev_object handle for the previous object found; initially *null*.

This routine requires three inputs:

1. A pointer to a list of *type* or *fulltype* constants. This list must be a static integer array, with **0** in the last element of the array.
2. A handle for the reference object, which determines where the routine will search for the destination objects.
3. A handle to the last target object found.

By providing a list of constants, **acc_next()** can be used in three general contexts:

- **acc_next()** can obtain handles for a specific type of object that does not have an object-specific ACC next routine by providing the *type* constant of that object.
- **acc_next()** can obtain handles for several different types of objects at the same time by providing a list of multiple object *type* constants.
- **acc_next()** can obtain handles for only certain objects within a larger class of objects by providing a list of one or more *fulltype* constants. For example, **acc_next_net()** will access all nets of any net type within a reference module. If only wired-logic net types were desired, **acc_next()** can be used to access only those specific net fulltypes.

The following example illustrates using **acc_next()** to obtain handles for only wired logic nets, and exclude other types of nets:

```
handle mod_h, net_h;

static int wired_nets[5] = {accWand, accWor, accTriand,
                           accTrior, 0};

/* add code to get a module handle */

net_h = null; /* initialize the target handle to null */
while (net_h = acc_next(wired_nets, mod_h, port_h) ) {
    /* perform desired operations on the net handle */
}
```

The `acc_next()` routine only supports a subset of the object type and fulltype constants. Object types which are not supported by `acc_next()` can only be accessed using an object-specific next routine. The supported constants are:

- Verilog data type constants: `accNet`, `accReg`, `accIntegerVar`, `accTimeVar`, `accRealVar`, and `accNamedEvent`.
- Verilog net fulltypes: `accWire`, `accTri`, `accWand`, `accTriand`, `accWor`, `accTrior`, `accTri0`, `accTri1`, `accTrireg`, `accSupply0`, and `accSupply1`.
- Verilog module type and fulltype constants: `accModule`, `accTopModule`, `accModuleInstance`, and `accCellInstance`.
- Verilog primitive fulltypes: `accCombPrim`, `accSeqPrim`, `accAndGate`, `accNandGate`, `accNorGate`, `accOrGate`, `accXorGate`, `accXnorGate`, `accBufGate`, `accNotGate`, `accBufif0Gate`, `accBufif1Gate`, `accNotif0Gate`, `accNotif1Gate`, `accNmosGate`, `accPmosGate`, `accCmosGate`, `accRnmosGate`, `accRpmosGate`, `accRcmosGate`, `accRtranGate`, `accRtranif0Gate`, `accRtranif1Gate`, `accTranGate`, `accTranif0Gate`, `accTranif1Gate`, `accPullupGate`, `accPulldownGate`.

NOTE → The `reg`, `integer`, `time`, `real` and `event` data types do not have a corresponding object-specific ACC next routine. The generic `acc_next()` routine must be used to obtain handles for these objects.

14.11.1 A complete PLI application for `$show_all_signals`

Example 14-3 lists the complete C code for the `$show_all_signals` PLI application. This application uses the generic `acc_next()` routine to obtain handles for all signals in a module, including the data types of `net`, `reg`, `integer`, `time`, and `real`.

CD The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.014/show_all_signals1_acc.c`
- Verilog test bench: `Chapter.014/show_all_signals1_test.v`
- Verilog-XL results log: `Chapter.014/show_all_signals1_test.log`

Example 14-3: \$show_all_signals, version 1 — using the acc_next() routine

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/*********************************************
 * checktf application
*****************************************/
int PLIbook_ShowSignals1_checktf()
{
    acc_initialize();
    if (tf_nump() != 1)
        tf_error("$show_all_signals must have 1 argument.");
    else if (tf_typep(1) == TF_NULLPARAM)
        tf_error("$show_all_signals arg cannot be null.");
    else if (acc_fetch_type(acc_handle_tfarg(1)) != accModule)
        tf_error("$show_all_signals arg must be a module instance.");
    acc_close();
    return(0);
}

/*********************************************
 * calltf application
*****************************************/
int PLIbook_ShowSignals1_calltf()
{
    handle module_h, signal_h;
    static int signal_types[6] = {accNet, accReg, accIntegerVar,
                                 accTimeVar, accRealVar, 0 };
    acc_initialize();
    module_h = acc_handle_tfarg(1);
    io_printf("\nAt time %s, signals in module %s (%s):\n",
              tf_strgettime(),
              acc_fetch_fullname(module_h),
              acc_fetch_defname(module_h));
    signal_h = null;      /* start with known value for target handle */
    while (signal_h = acc_next(signal_types, module_h, signal_h)) {
        io_printf(" %-13s %-13s value is %s (hex)\n",
                  acc_fetch_type_str(acc_fetch_fulltype(signal_h)),
                  acc_fetch_name(signal_h),
                  acc_fetch_value(signal_h, "%h", null));
    }
    acc_close();
    return(0);
}
```

Example 14-4 lists Verilog source code for testing `$show_all_signals`. This test is similar to the test for `$show_all_nets`, but the lower level adder model has been changed from a gate level model to an RTL model in order to use more data types in the Verilog source code. Figure 14-2, which follows, shows the simulation results from running a simulation with `$show_all_signals`.

Example 14-4: `$show_all_signals1` — Verilog HDL test case for the PLI application

```
'timescale 1ns / 1ns
module top;
    integer      test;
    tri [1:0] results;

    addbit i1 (test[0], test[1], test[2], results[0], results[1]);

initial
begin
    test = 3'b000;
    #10 test = 3'b001;

    #10 $show_all_signals1(top);
    #10 $show_all_signals1(i1);

    #10 $stop;
    #10 $finish;
end
endmodule

/** An RTL level 1 bit adder model ***/
'timescale 1ns / 1ns
module addbit (a, b, ci, sum, co);
    input  a, b, ci;
    output sum, co;

    wire  a, b, ci;
    reg   sum, co;

    always @ (a or b or ci)
        {co, sum} = a + b + ci;

endmodule
```

Figure 14-2: `$show_all_signals1` — simulation results, using Verilog-XL

The screenshot shows the Cadence Verilog-XL Turbo NT simulation environment. The top menu bar includes File, Control, Set, Show, Select, Tools, Options, and Help. The toolbar contains icons for Scope, Objects, and Tools. The status bar shows Time: 40. The main window displays the Verilog code for `$show_all_signals1`, which includes a testbench with stimulus and calls to `$show_all_signals1` at time 10. The output pane shows simulation results at time 20 and 30, listing signal names and their hex values. The bottom status bar says Ready.

```

23      #10 test = 3'b001;
24
25      #10 $show_all_signals1(top);
26      #10 $show_all_signals1(i1);
27
28      #10 $stop;
29      #10 $finish;
30
31 endmodule .

```

Scope top Subscopes

Highest level modules:
top

At time 20, signals in module top (top):
accIntegerVar test value is 00000001 (hex)
accTri results value is 1 (hex)

At time 30, signals in module top.i1 (addbit):
accWire a value is 1 (hex)
accWire b value is 0 (hex)
accWire ci value is 0 (hex)
accRegister sum value is 1 (hex)
accRegister co value is 0 (hex)

L28 "show_all_signals1_test.v": \$stop at simulation time 40
Type ? for help
C1 >

14.12 Obtaining handles to the current hierarchy scope

The Verilog language allows a system task or system function to be invoked from any hierarchy scope. A *scope* in the Verilog HDL is a level of design hierarchy, and can be represented by several constructs:

- Module instances
- Named statement groups
- Verilog HDL tasks
- Verilog HDL function

The following example calls the `$show_all_signals` from a named statement group:

```
module top;
  ...
  always @(posedge clock)
    begin: local
      integer i;
      reg      local_bus;
      ...
      $show_all_signals;
    end
endmodule
```

A useful enhancement to the `$show_all_signals` example is to allow either no system task argument or a null system task argument to represent the module instance which called the `$show_all_signals` system task. The difference between no argument and a null argument is shown in the following two examples.

No system task/function arguments:

```
$show_all_signals;
```

A null system task/function argument:

```
$show_all_signals();
```

The following Verilog source code shows the enhanced usage possibilities for the `$show_all_signals` example:

```
module top;
  ...
  addbit i1 (a, b, ci, sum, co); // instance of an adder
  ...
  always @(posedge clock)
    $show_all_signals;           // list signals in this module
    $show_all_signals(i1);      // list signals in instance i1
endmodule

module addbit (a, b, ci, sum, co);
  ...
  always @(sum or co)
    $show_all_signals();        // list signals in this module
endmodule
```

In order to access all signals in a module, the `$show_all_signals` application will need a handle for a module instance, but now the name of the module instance is not passed

to the PLI application as an argument to `$show_all_signals`. The ACC library provides two ways to obtain the handle for the module which called a PLI application.

`acc_handle_calling_mod_m` returns the handle for the module from which a PLI application was called. The syntax for this routine is:

handle acc_handle_calling_mod_m

Note that in the ACC library, this routine is defined as a macro, not a function. Therefore, it should not be called with parentheses at the end of the name.

The `$show_all_signals` PLI application can be easily enhanced to work with no argument or a null argument using the code:

```
if (tf_nump() == 0) /* no task/function arguments */
    mod_h = acc_handle_calling_mod_m;
else if (tf_typep(1) == tf_nullparam) /* null argument */
    mod_h = acc_handle_calling_mod_m;
else /* a task/function argument exists */
    mod_h = acc_handle_tfarg(1);
```

In the Verilog HDL, a level of design hierarchy can be an object other than a module. The following example calls the `$show_all_signals` from a named statement group:

```
module top;
    ...
    always @(posedge clock)
        begin: local
            integer i;
            reg      local_bus;
            ...
            $show_all_signals;      // list signals in this scope
        end
    endmodule
```

In the above example, the `$show_all_signals` applications should search for signal names in the local hierarchy scope, which is a named statement group, instead of a module instance. The name of any type of hierarchy scope can be passed to the system task as a task/function argument, but, in this example, `$show_all_signals` is being called with no arguments. To obtain the local hierarchy scope without being passed, the scope name requires two steps:

1. Call the `acc_handle_tfinst()` routine to obtain a handle for the system task/ function which called the PLI application. This routine does not require any inputs.
 2. Call the `acc_handle_scope()` routine to obtain a handle for the scope containing a reference handle. For this example, the reference handle will be the system task handle.

Example 14-5 contains the complete listing of the enhanced `$show_all_signals`, with the ability to use either no arguments or null arguments to represent the local design hierarchy scope.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.014/show_all_signals2_acc.c
 - Verilog test bench: Chapter.014/show_all_signals2_test.v
 - Verilog-XL results log: Chapter.014/show_all_signals2_test.log

Example 14-5: `$show_all_signals`, version 2—obtaining a handle for the local scope

```

acc_initialize();
if (tf_nump() == 0)
    module_h = acc_handle_scope(acc_handle_tfinst());
else if (tf_typep(1) == tf_nullparam)
    module_h = acc_handle_scope(acc_handle_tfinst());
else
    module_h = acc_handle_tfarg(1);
io_printf("\nAt time %s, signals in module %s (%s):\n",
          tf_strgettime(),
          acc_fetch_fullname(module_h),
          acc_fetch_defname(module_h));
signal_h = null; /* start with known value for target handle */
while (signal_h = acc_next(signal_types, module_h, signal_h)) {
    io_printf(" %-13s %-13s value is %s (hex)\n",
              acc_fetch_type_str(acc_fetch_fulltype(signal_h)),
              acc_fetch_name(signal_h),
              acc_fetch_value(signal_h, "%h", null));
}
acc_close();
return(0);
}

```

14.13 Obtaining handles to multiple task/function arguments

Another useful modification to the `$show_all_signals` application is to allow multiple hierarchy scopes to be specified at the same time. For example:

```
$show_all_signals(i1, ,top.local);
```

In the above example, there are three system task/function arguments, the second argument being null, to indicate the local hierarchy scope.

Example 14-6 illustrates using a C `for` loop to access each system task/function argument.



TIP

This example uses a more structured programming style by creating a separate function called `get_all_signals()` to do the work of searching for all signals and printing the current logic value. By moving this application logic into a separate C function, the *calltf routine* is kept shorter and easier to read, plus there is less duplication of code. Using structured programming techniques makes it easier to maintain or enhance the functionality of the PLI application.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.014/show_all_signals3_acc.c
- Verilog test bench: Chapter.014/show_all_signals3_test.v
- Verilog-XL results log: Chapter.014/show_all_signals3_test.log

Example 14-6: \$show_all_signals, version 3—obtaining handles for multiple tfargs

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/*********************************************
 * checktf application
 *****/
int PLIbook_ShowSignals3_checktf()
{
    int i, numargs;

    acc_initialize();
    numargs = tf_numargs();
    if (numargs == 0)
        return(0); /* no arguments is OK, skip remaining checks */
    for (i = 1; i <= numargs; i++) {
        if (tf_typepep(i) == TF_NULLPARAM)
            break; /* null argument is OK, skip other checks for this arg */
        else if (acc_fetch_type(acc_handle_tfarg(i)) != accModule)
            tf_error("$show_all_signals arg must be a module instance.");
    }
    acc_close();
    return(0);
}

/*********************************************
 * calltf application
 *****/
void PLIbook_GetAllSignals(); /* prototype function used by calltf */

int PLIbook_ShowSignals3_calltf()
{
    handle module_h;
    int i, numargs;
    acc_initialize();
    numargs = tf_numargs();
    if (numargs == 0) {
        module_h = acc_handle_scope(acc_handle_tfinst());
        PLIbook_GetAllSignals(module_h);
    }
    else
        for (i = 1; i <= numargs; i++) {
            if (tf_typepep(i) == tf_nullparam)
                module_h = acc_handle_scope(acc_handle_tfinst());
```

```
    else
        module_h = acc_handle_tfarg(i);
        PLIbook_GetAllSignals(module_h);
    }
    acc_close();
    return(0);
}

void PLIbook_GetAllSignals(handle module_h)
{
    handle signal_h;
    static int signal_types[6] = {accNet, accReg, accIntegerVar,
                                accTimeVar, accRealVar, 0 };
    acc_initialize();
    io_printf("\nAt time %s, signals in module %s (%s):\n",
              tf_strgettime(),
              acc_fetch_fullname(module_h),
              acc_fetch_defname(module_h));
    signal_h = null; /* start with known value for target handle */
    while (signal_h = acc_next(signal_types, module_h, signal_h)) {
        io_printf(" %-13s %-13s value is %s (hex)\n",
                  acc_fetch_type_str(acc_fetch_fulltype(signal_h)),
                  acc_fetch_name(signal_h),
                  acc_fetch_value(signal_h, "%h", null));
    }
    acc_close();
    return;
}
```

14.14 Summary

The ACC routines in the PLI standard provide direct access to what is happening within a Verilog simulation. This access is done using an object oriented method, where most Verilog HDL constructs that can exist in a simulation data structure are treated as objects. The ACC routines use *handles* to reference these objects. A large number of ACC routines are provided to obtain handles for the various types of Verilog HDL objects, and other ACC routines retrieve data about the objects, such as the name of an object or the vector size of a net.

This chapter has focused on how to create PLI applications using the ACC library. The following three chapters will present more detail on the syntax and usage of the 103 ACC routines in the PLI standard. These chapters include several additional examples of PLI applications which use the ACC routines.

CHAPTER 15 *Details on the ACC Routine Library*

The ACC library provides 103 C functions that can interact with Verilog simulators. The previous chapter provided an overview of the ACC library, and how ACC routines are used in PLI applications. This chapter presents a more detailed description of the ACC library, and how ACC routines access information from Verilog simulations.

The concepts presented in this chapter are:

- Initializing and configuring ACC routines
- ACC error handling
- ACC object diagrams and object relationships
- Using ACC *handle* routines
- Using ACC *next* routines
- Miscellaneous ACC routines

15.1 PLI application performance considerations

The run-time performance of a simulator can be impacted in either a positive way or a negative way by PLI applications. Often, a complex algorithm can be represented in the C language, using C language data types, much more efficiently than in the hardware-centric Verilog HDL language. The C language can be used for an abstract representation of a design, when 4-state logic, logic transitions, simulation time, and other details are not required, but which a hardware description language must be able

to represent. The abstraction that C offers often makes it possible to greatly increase the run-time performance of a simulation algorithm.

However, a poorly thought out PLI application can actually decrease the run-time performance of a simulation. Each call to a routine in the PLI library will take time to be executed. It is important to architect a PLI application to minimize the number of times ACC routines are used.

The following guidelines can help in planning an efficient PLI application:

- Good C programming practices are essential. C programming style and techniques not discussed within the scope of this book.
- Consider every call to an ACC routine as expensive, and try to minimize the number of calls.
- ACC routines which obtain object handles using an object's name are less efficient than routines which obtain object handles based on a relationship to another object.
- Routines which convert logic values from a simulator's internal representation to C strings, and vice-versa, are less efficient than using other C data types. Strings are a convenient means of representing 4-state values in certain types of applications, but strings should be used prudently.
- When the same object must be accessed many times during a simulation, the object handle can be obtained once, and saved in application-allocated storage. Using a pointer to the storage, a PLI application has immediate access to the object handle, without having to call an ACC routine to obtain the handle each time it is needed.
- Use the ACC library to access the unique abilities of hardware description languages, such as representing hardware parallelism and hardware propagation times. Simulator vendors have invested a great deal in optimizing a simulator's algorithms, and that optimization should be utilized in a PLI application.

When developing a PLI application, one primary consideration should be how often a PLI application will be called during a simulation. It is well worth the effort to optimizing the performance of an application that is invoked every clock cycle, but may not be as important for an application that is only invoked once during a simulation.

NOTE

The objective of this book is to show how the routines in the ACC library are used. Short examples of using many of these routines are shown in the context of complete PLI applications. In order to meet the book's objectives, the examples presented in this book do not always follow the guidelines of efficient C coding and prudent usage of the ACC routines. It is expected that when parts of these example PLI applications are adapted for other applications, the coding style will also be modified to be more efficient and robust.

15.2 Initializing and configuring ACC routines

Several ACC routines can be configured for how the routines access information in the Verilog simulation data structure. In addition, some ACC routines need to allocate and initialize memory for their operation. Three ACC routines are used to initialize and configure ACC routines.

15.2.1 Initializing and closing the ACC environment

Two special ACC routines, `acc_initialize()` and `acc_close()`, are used to initialize and maintain the environment used by the ACC library.

`bool acc_initialize()`

The `acc_initialize()` routine performs two primary operations: allocate and initialize any memory that is needed by the ACC routines, and set all ACC configurations to their default values. The routine returns true if it was successful, and false if an error occurred.

`void acc_close()`

`acc_close()` frees any memory that was allocated by `acc_initialize()`, and resets all ACC configurations back to their default values.

An example of using `acc_initialize()` and `acc_close()` is:

```
my_calltf_app()
{
    /* declarations */
    acc_initialize();
    /* use routines from ACC library */
    acc_close();
}
```

When to use `acc_initialize()` and `acc_close()`

 **TIP** Simulation run-time performance might be improved through judicious usage of `acc_initialize()` and `acc_close()`. PLI applications which do not use any ACC routines do not need to call `acc_initialize()` and `acc_close()`.

The IEEE standard recommends that `acc_initialize()` be called at the beginning of every PLI application which uses the ACC library, and `acc_close()` be called at the end of the application. Following the IEEE recommendation will ensure that a PLI application is well behaved and portable to all Verilog simulators.

Some—but not all—Verilog simulators do not need `acc_initialize()` and `acc_close()`. In many simulators, the configuration of ACC routines is automatically set to default values when the PLI application is entered. Any memory required by ACC routines is automatically allocated when needed, and automatically freed when not needed. Refer to the documentation of a specific simulator to see if `acc_initialize()` and `acc_close()` can be omitted for that simulator.

Calling `acc_initialize()` and `acc_close()` when the routines are not required may slightly slow down the run-time performance of a simulator. If a PLI application is invoked many times during a simulation, such as at every clock cycle, this performance cost can become expensive. For those simulators where the routines are not needed, initializing and closing the ACC routines is an unnecessary cost to simulation performance.

Three possible ways to utilize `acc_initialize()` and `acc_close()` are:

- Follow the IEEE recommendation, and use `acc_initialize()` and `acc_close()` in all ACC based PLI applications. This provides maximum portability of the PLI application, with minimum optimization.
- Use conditional compilation to include or exclude `acc_initialize()` and `acc_close()`, depending on the simulator for which the PLI application is being compiled. This provides better optimization, but complicates developing and compiling PLI applications.
- Only use `acc_initialize()` and `acc_close()` in PLI applications which use ACC routines that are affected by `acc_configure()` (excluding warning and error configuration, which affect all ACC routines). The ACC routines which can be configured are the routines which are most likely to need initialization. This approach provides a reasonable compromise between portability and performance.

15.2.2 Configuring the ACC environment

Several ACC routines can be configured for how they operate. A special ACC routine is used to configure all of these routines.

bool acc_configure(config_param, value)

int config_param One of the constants listed in table 15-1.

*char *value* configuration value as a character string.

The `acc_configure()` routine configures specific ACC routines. The routine returns 1 if successful, 0 if an error occurred. The name of each configuration is represented by a constant, and the setting of the configuration is represented as a string. Table 15-1 lists the configuration constants and the ACC routines affected.

Configuration Constant	Description and Routines Affected
<code>accDevelopmentVersion</code>	documents the version of the PLI standard for which the PLI application was developed no ACC routines affected
<code>accDisplayErrors</code>	enables or disables printing error messages for run-time errors caused by ACC routines all ACC routines
<code>accDisplayWarnings</code>	enables or disables printing warning messages for run-time warnings caused by ACC routines all ACC routines
<code>accDefaultAttr0</code>	controls the default return value if an attribute is not found <code>acc_fetch_attribute()</code>
<code>accEnableArgs</code>	controls which input arguments must be specified <code>acc_handle_modpath()</code> <code>acc_handle_tchck()</code> <code>acc_handle_scope()</code>
<code>accPathDelimStr</code>	controls the delimiter used in path names <code>acc_fetch_name()</code> <code>acc_fetch_fullname()</code> <code>acc_fetch_attribute()</code>
<code>accMinTypMaxDelays</code>	controls whether min and max delays are used <code>acc_fetch_delays()</code> <code>acc_append_delays()</code> <code>acc_replace_delays()</code>
<code>accPathDelayCount</code>	controls the number of path delays read or modified <code>acc_fetch_delays()</code> <code>acc_append_delays()</code> <code>acc_replace_delays()</code>

Table 15-1: ACC configuration constants

Configuration Constant	Description and Routines Affected
accToHizDelay	controls how turn-off delays are specified acc_fetch_delays() acc_append_delays() acc_replace_delays()
accMapToMipd	controls how interconnect delays are mapped to input ports acc_replace_delays()

Table 15-1: ACC configuration constants (continued)

15.3 ACC routine error handling

The ACC routines have built-in error handling for when a call to a routine cannot perform its operation. For example, an ACC routine to obtain a handle for an object using the object's name will fail if an object of that name does not exist in the design.

The ACC library provides three actions which should occur when an ACC routine cannot perform its operation:

- Set a global error flag, called **acc_error_flag**.
- Display an error message or warning message to the simulator's output window.
- Return an exception value, if the ACC routine returns values.

15.3.1 The global ACC error flag

The global **acc_error_flag** is set to 0 if a call to an ACC routine is successful, and is set to a non-zero value if the call was unsuccessful. The status of the error flag is updated each time an ACC routine is called. Every ACC routine shares the global **acc_error_flag**. The following C code fragment illustrates using the error flag.

```
handle net_handle, module_handle;
char *net_name;

/* add code to read the name of a net from a file */
net_handle = acc_handle_by_name(net_name, module_handle);
if (acc_error_flag) {
    io_printf("Net %s was could not be found\n", net_name);
}
... /* use the net handle obtained */
```

15.3.2 Exception return values

ACC routines which return values have a defined exception value. Most, but not all ACC routines follow the same convention for the exception value. However, there are a few routines which do not adhere to the convention. The exception values used by most ACC routines are listed in Table 15-2. If an ACC routine follows a different convention, its exception value is noted in the description of the routine.

Type of ACC Routine	Exception Value
ACC routines which return int values	0
ACC routines which return double values	0.0
ACC routines which return boolean values	false
ACC routines which return handle values	null
ACC routines which return pointers to character strings	pointer to a null string

Table 15-2: Exception return values for ACC routines

15.3.3 Enabling and disabling ACC error and warning messages

By default, whenever an ACC routine is not successful, an ACC error message is printed in the output channel of the simulator. This is the same output channel used by routines such as `io_printf()`— it is not necessarily the operating system's `stdout` or `stderr` message channels. The ACC library provides a means to disable this automatic error message generation, using `acc_configure()`, as follows:

```
acc_configure(accDisplayErrors, "false");
```

Setting the configuration of `accDisplayErrors` to "true" will re-enable automatic error message generation, as will re-initializing the ACC environment.

Some simulators may also generate ACC warning messages, such as when an invalid input is provided to the ACC routine. Typically, a warning message indicates that a less severe error occurred with the ACC routine. The `acc_error_flag` is not set when a warning message occurs.

The end-user of a PLI application may not be interested in these non-fatal warnings. But, these warnings can be of interest to a PLI application developer. By default, the

generation of ACC warning messages is *disabled*. The `acc_configure()` routine is used to enable warning messages, as follows:

```
acc_configure(accDisplayWarnings, "true");
```



TIP

Warning messages can provide valuable information when developing or debugging a PLI application. Often a warning message can indicate a potential problem with an application that might not be obvious as an application is being tested. It is a good technique to enable warning messages until it is certain that a PLI application is working correctly, and then disable warning messages for better run-time performance of the PLI application. The TF and ACC routines allow user-defined invocation options to be created, and many PLI applications use invocation options as a way to easily enable or disable the ACC warning messages.

15.4 Using ACC object diagrams

The IEEE 1364 Verilog standard defines what Verilog HDL constructs are considered objects in the ACC library. However, in the ACC portion of the IEEE 1364 standard, the relationships between objects are not defined. Object relationships are only defined in the VPI portion of the standard.

This book includes all of the information about ACC objects which is defined in the IEEE 1364 standard, and adds *object diagrams* for each object which the ACC routines can access. These object diagrams document:


- The *properties* of an object. For example, a net object has *name*, *vector size*, and *logic value* properties (as well as several other properties).
- The *relationships* of the object. Relationships indicate how an object is connected to, or contained within, other objects within a Verilog data structure. For example, a net is contained within a module, and may also be connected to other objects, such as a module port or primitive terminal.

The object diagrams for an object are based on enclosures with arrows. The type of object is listed within each enclosure, and the relationships to other objects are shown as arrows between the enclosures. The properties of the object are listed below the diagrams. The specific ACC routines required to access the properties of an object or to traverse to a related object are shown in the diagrams.

The complete set of ACC object diagrams are in Appendix C. This section shows how to read the ACC object diagrams, Section 15.7, later in this chapter, shows how to utilize the diagrams to traverse from any point in a Verilog design hierarchy to any other point.

Following is the object diagram for a Verilog module object:

Figure 15-1: ACC object diagram for Verilog modules




Related Routines:

<code>acc_fetch_type()</code>	returns accModule
<code>acc_fetch_fulltype()</code>	returns accTopModule , accModuleInstance , accCellInstance
<code>acc_fetch_name()</code>	returns the instance name of a module
<code>acc_fetch_fullname()</code>	returns the hierarchical path name of a module
<code>acc_fetch_defname()</code>	returns the definition name of a module
<code>acc_fetch_delay_mode()</code>	returns accDelayModeNone , accDelayModeZero , accDelayModeUnit , accDelayModePath , accDelayModeDistrib , accDelayModeMTM
<code>acc_fetch_timescale_info()</code>	returns the timescale of a module
<code>acc_fetch_location()</code>	returns source file name & line no. containing module instance

15.4.1 Object diagram symbols

An object diagram contains four primary symbols and four font type faces:

- A *solid circle*, such as \circ , designates either the top level of a Verilog hierarchy tree, or the hierarchy scope from which the PLI application was called.
- A *solid enclosure*, such as **module** or **port**, designates a Verilog object. The name of the object is shown within the enclosure. The font used for the name has significance:
 - A **non-italicized, bold font** designates that this object is being defined in this diagram. In the diagram for module objects, the name, **module**, is in bold.
 - A **non-italicized, non-bold font** designates that this object is being referenced in this diagram, but is not being defined. The definition will appear in a different diagram. For example, in the diagram for module objects, the name **port**, is not bolded.
- A *small dotted enclosure*, such as $\langle _variables _ \rangle$, designates a reference to a named class of Verilog objects. A class of objects is a group of several objects which have something in common. The name of the object class is shown within the enclosure, using an *italicized, non-bold font*. The specific objects within a class are listed in the diagram for the class definition.
- A *large dotted enclosure*, which has smaller enclosures within it, designates the definition of a class of Verilog objects. The large enclosure contains all of the objects which make up the class. A class of objects may or may not have a name, but only a named class can be referenced in another diagram. The name of the class is shown at the top of the large enclosure, using an *italicized, bold font*. Two examples of object class definitions are:



15.4.2 Traversing object relationships

The relationship of one object to another object is shown as arrows in the object diagrams. Each arrow indicates the relationship from a reference object (the originating end of the line) to a target object (the terminating end of the line). The type of arrow

at the target object indicates the type of relationship. Most objects can be both a reference object and a target object, depending on which direction the Verilog hierarchy is being traversed.

There are two types of object relationships possible in a Verilog design for which ACC routines can obtain handles:

- **One-to-one** relationships are represented by a line which terminates with a single arrow in an object diagram. A one-to-one relationship indicates that a given object is related to, at most, one of another type of object. In the module object diagram shown in Figure 15-1 on page 481, there is a single arrow going from **port** back to **module**, which indicates a given port is only contained within one module.
- **One-to-many** relationships are represented by a line which terminates with a double arrow in an object diagram. A one-to-many relationship indicates that a given object is related to any number of another type of object. In the module object diagram shown in Figure 15-1 on page 481, there is a double arrow going from **module** to **port**, which indicates there may be any number of ports within a module.

In most object relationships, the connecting line both originates and terminates with a single or double arrow. In certain relationships, however, the connecting line has no arrow on the originating end, at the reference object. This indicates that the relationship is one-way. That is, the target object can be accessed from the reference object, but it is not possible to get back to the reference object from the target object.

Several ACC routines provide a means for a PLI application to locate and retrieve handles for Verilog objects. This chapter will introduce these routines, and subsequent chapters in this part of the book will present several examples of using these routines in practical PLI applications.

There are two general groups of ACC routines used to obtain handles for objects, which are referred to as **ACC handle routines** and **ACC next routines**.

- The **ACC handle routines** obtain handles for one object of a specific object type. For example, the routine `acc_handle_port()` will obtain a handle for one specific port within a module.
- The **ACC next routines** obtain handles for all objects of a specific object type. For example, the routine `acc_next_port()` will obtain handles for all of the ports within a module.

The description and usage of these routines are presented in sections 15.5 and 15.6, which follow.

15.5 Using ACC handle routines

An ACC **handle** is an abstraction used to reference an object within the simulation data structure. ACC routines use this abstraction to access information about the object. This layer of abstraction allows PLI applications to be portable to any Verilog simulator, because the abstract handle is a layer between the PLI application and the internal data structures of the simulator.

ACC **handle routines** return a handle for a single object. Generally, these routines are used when there is a one-to-one relationship shown in the object diagram, which is represented by a single arrow at the terminating end of a relationship line. **ACC handle routines are object-specific**, so there is a different routine for each type of object that can be accessed using handle routines. There are 25 handle routines in the ACC library:

acc_handle_by_name()	acc_handle_scope()
acc_handle_calling_mod_m	acc_handle_path()
acc_handle_condition()	acc_handle_pathin()
acc_handle_conn()	acc_handle_pathout()
acc_handle_datapath()	acc_handle_port()
acc_handle_hiconn()	acc_handle_simulated_net()
acc_handle_interactive_scope()	acc_handle_tchk()
acc_handle_itfarg()	acc_handle_tchkarg1()
acc_handle_loconn()	acc_handle_tchkarg2()
acc_handle_modpath()	acc_handle_terminal()
acc_handle_notifier()	acc_handle_tfarg()
acc_handle_object()	acc_handle_tfinst()
acc_handle_parent()	

The ACC handle routines are straightforward and simple to use. The complete syntax for each ACC handle routine is shown in Appendix D, and is not duplicated in this section. This chapter includes examples of using many of these ACC handle routines, and provides additional description on a few of the routines.

Most ACC handle routines require as an input a handle for a reference object. The routines return a handle for a specific target object. The ACC object diagrams show which ACC handle routine is used to obtain a handle for a specific object. As an example, assume that a PLI application had already obtained a handle for a port, and needs to locate the module which contains that port. The **reference** point will be the

port, and the *target* will be a Verilog module. In the object diagram for modules, which was shown in Figure 15-1 on page 481, there is a single arrow from **port** to **module**, indicating a one-to-one relationship. Next to the target object (the module), the ACC routine to obtain a handle for the module is shown as being **acc_handle_parent()**. Using this information, the following C code fragment can be used to obtain the module handle from the port handle:

```
handle port_handle, module_handle;  
/* add code to obtain handle for a port */  
module_handle = acc_handle_parent(port_handle);
```

15.6 Using ACC next routines

Many objects in Verilog have a one-to-many relationship with other objects. These relationships are represented by a double arrow terminating at the target object in the object diagrams. The *ACC next routines* are used to obtain handles for all of objects in this type of relationship. As with the ACC handle routines, the ACC next routines are object-type specific, so there is a different ACC routines for most types of Verilog objects for which next routines can obtain handles. There are 22 ACC next routines:

acc_next()	acc_next_net()
acc_next_bit()	acc_next_output()
acc_next_cell()	acc_next_parameter()
acc_next_cell_load()	acc_next_port()
acc_next_child()	acc_next_portout()
acc_next_driver()	acc_next_primitive()
acc_next_hiconn()	acc_next_scope()
acc_next_input()	acc_next_specparam()
acc_next_load()	acc_next_tchk()
acc_next_loconn()	acc_next_terminal()
acc_next_modpath()	acc_next_topmod()

The complete syntax for each ACC next routine is shown in Appendix D, and is not duplicated in this section. Chapter 14 discussed how ACC next routines are used (refer back to section 14.6 on page 452). This chapter shows how the ACC object diagrams relate to the ACC next routines.

Most of the ACC next routines require two inputs:

1. A handle for the *reference object*, which determines where the routine will search for the destination objects.
2. A handle for the *previous target object* found.

The ACC object diagrams show which ACC next routine should be used to traverse from one object to another. As an example, suppose a PLI application needed to obtain handles for all ports in a module. In the object diagram for modules shown in Figure 15-1 on page 481, there is a double arrow from **module** to **port**, indicating a one-to-many relationship. The ACC routines listed next to the target objects (the ports) show that **acc_next_port()** can be used to access all ports within that module. For example:

```
port_h = null; /* initialize the target handle for null */
while ( port_h = acc_next_port(mod_h, port_h) ) {
    /* perform desired operations on the port handle */
}
```

Obtaining handles for just one object, using ACC next routines

On occasion, a PLI application might need to obtain a handle for just the first target object in a one-to-many relationship. Or, an application may only need to determine if any target objects exist, without obtaining any target object handles.

For example, an application might need to determine if a Verilog module contains instances of other modules, but the handles for all of the module instances are not needed. The object diagram for modules (refer back to Figure 15-1 on page 481) shows a double arrow from **module** to **module**, indicating a one-to-many relationship from a module to all module instances within that module. The ACC routines listed next to the target object show that **acc_next_child()** can be used to access all module instances within a module. Since this example PLI application does not need the child module handles which would be returned from **acc_next_child()**, the application can simplify the way the next routine is used in two ways:

- The ACC next routine does not need to be called in a loop to access all handles.
- The previous target handle can be replaced with **null**, since the routine will not be called to find a second target object.

The following code fragment illustrates using `acc_next_child()` to perform a true/false test, where a true is returned if a module represents the bottom of the Verilog hierarchy (the module does not contain any module instances).

```
int is_bottom_module(handle this_module_h)
{
    if (acc_next_child(this_module_h, null))
        return(0); /* a module instance was found */
    else
        return(1); /* no child module instances */
}
```

A similar situation that occasionally arises in PLI application is when only the first target object of a one-to-many relationship is desired. Once again, since the second or subsequent target objects will not be accessed, the ACC next routine does not need to be called in a loop, and the previous target handle does not need to be provided. The following code fragment illustrates using an ACC next routine to obtain only the first port of a module.

```
handle first_port_handle;
first_port_handle = acc_next_port(module_handle, null);
```

15.7 Traversing Verilog hierarchy using object relationships

By following the object diagrams and using the appropriate ACC handle and ACC next routines, the Verilog design hierarchy can be traced from one object to any other object anywhere in the Verilog design. Traversing the design hierarchy often requires traversing from one object, to another object, to another object, until the desired destination is attained.

For example, suppose a PLI application had obtained a handle for a module, and the application needs to locate every module output, where the output is also connected to a module path delay. The following Verilog source code shows the starting and ending objects for which handles are desired in this example.

```


start with a handle
to a module
module dff (clk, d, q, qb);
    input clk, d;
    output q, qb;
    ff_prim g1 (q, d, clk, );
    not g2 (qb, q);
specify
    (d *> q) = 2.5;
endspecify
endmodule
obtain handles for all output
ports which are also
connected to a module path
delay (port q in this example)

```

By following the connections in the object diagrams, the following object connections can be traversed:


1. The partial **module** diagram shown in Figure 15-2, which follows, shows a one-to-many connection from **module** to **mod path**. The diagram also shows that the ACC next routine **acc_next_modpath()** is used to traverse from the module to the module path target. Since this is a next routine, a handle for each module path in the module can be obtained.

Figure 15-2: ACC object diagram for Verilog modules (partial)



2. The *mod path* object diagram in Figure 15-3 shows a one-to-many connection from **mod path** to **path term**. This connection can be traversed with specific ACC routines to obtain handles for the path input terminals or output terminals. Since this example application is looking for the output ports connected to module paths, the output terminal will be the connection which needs to be followed. In addition, the diagram shows two ACC routines to obtain handles for the path output terminal. The ACC handle routine, `acc_handle_pathout()`, will return only a single handle, which is the first output of the path. The ACC next routine, `acc_next_output()`, will return handles for all outputs in the path (if the routine is called in a loop).

Figure 15-3: ACC object diagram for module paths (partial)



NOTE: Refer to Appendix D.1.10 on page 745 for the complete object diagram for module paths.

Since the objective for this example PLI application is to find all ports connected to path outputs, the `acc_next_output()` routine will be the right choice for this example. Using this routine, a handle for each path output in the module path can be obtained. Note that this call to `acc_next_output()` must be done for each module path in the module, because the `acc_next_modpath()` routine will return the module path handles one at a time. Nested loops will be used, in order to obtain handles for each path output for each module path. If C while loops are used, the nested loops can be coded as:


```

modpath_h = pathterm_h = null;
while (modpath_h=acc_next_modpath(module_h, modpath_h)) {
    while (pathterm_h=acc_next_output(modpath_h, pathterm_h)) {
        /* continue traversing to port driven by this output */
    }
}
  
```

3. The **path term** object is part of the **mod path** object diagram, shown in Figure 15-3 on the previous page. The **path term** object shows a one-to-one connection to an object group called `'expr'`, (`expr` stands for expression). The diagram indicates that this object relationship is traversed using `acc_handle_conn()`.

The object diagram for the expression group, shown in Figure 15-4, below, lists many types of Verilog objects, including nets, variables, and several other objects. These are the types of objects which can be used in connections and statements in many places in the Verilog language. For example, the connection to a module port could be any of the objects listed in the expression class. The objects referenced within the expression group are defined in other diagrams.


Figure 15-4: ACC object diagram of expressions (partial)



A knowledge of the Verilog HDL language is needed, in order to know which types of objects in the expression diagram can be connected to an output path terminal. In this context, the only type of object which the Verilog HDL syntax will allow is a net, represented by the `'nets'` object group. Therefore, the object diagram for nets will show the next relationship in the connection from a module path to a port.

4. The **nets** object diagram, shown in Figure 15-5 on the following page, shows two types of connections from `'nets'`, to `'ports'`.

Figure 15-5: ACC object diagram for nets (partial)



NOTE: Refer to Appendix D.1.6 on page 741 for the complete object diagram for nets.

One connection to ports is accessed using `acc_next_port()`. The other connection has no terminating arrow, which indicates that the ACC routines cannot traverse hierarchy in that direction. The connection accessed by `acc_next_port()` is the connection from a net within a module to a port of that module. The other connection is from a port of a *module instance* within the current module to a net. The following Verilog model illustrates the difference between the two types of port objects:

```

module chip (clock, in1, out1, out2);
  wire out1, out2;
  dff u1 (clock, in1, out1, out2);
endmodule

module dff (clk, d, q, qb);
  ...
endmodule

```

Annotations explain the differences:

- A callout box points to the `out1` port of the `dff` module with the text: "module ports are accessed from an internal net using `acc_next_port()`".
- A callout box points to the `out1` port of the `chip` module with the text: "module instance ports cannot be accessed from an internal net using ACC routines".

Example summary

In the example described in this section, four object connections were followed to traverse the Verilog hierarchy from a module object to the output ports connected to a path delay. The connection path is from **module** to **mod path** to **path term** to **nets**, (in the expression group) to **ports**.

The C source code listed in Example 15-1 illustrates how these multiple connections through the Verilog hierarchy are traversed. The PLI application is *\$list_pathout_ports*. The usage is:

```
$list_pathout_ports(<module_instance_name>);
```



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.015/list_pathout_ports_acc.c
- Verilog test bench: Chapter.015/list_pathout_ports_test.v
- Verilog-XL results log: Chapter.015/list_pathout_ports_test.log

Example 15-1: *\$list_pathout_ports* — traversing Verilog hierarchy

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/*********************************************
 * checktf application
 *****/
int PLIbook_ListPorts_checktf()
{
    acc_initialize();
    if (tf_nump() != 1)
        tf_error("$list_pathout_ports must have 1 argument.");
    else if (tf_typep(1) == TF_NULLPARAM)
        tf_error("$list_pathout_ports arg cannot be null.");
    else if (acc_fetch_type(acc_handle_tfarg(1)) != accModule)
        tf_error("$list_pathout_ports arg must be a module instance.");
    acc_close();
    return(0);
}

/*********************************************
 * calltf application
 *****/
int PLIbook_ListPorts_calltf()
{
    handle module_h, modpath_h, pathterm_h, net_h, port_h;
    acc_initialize();
```

```


module_h = acc_handle_tfarg(1);
io_printf("\nOutput ports with path delays in module %s:\n",
          acc_fetch_defname(module_h));

modpath_h = pathterm_h = net_h = port_h = null;
while (modpath_h = acc_next_modpath(module_h, modpath_h)) {
    while (pathterm_h = acc_next_output(modpath_h, pathterm_h)) {
        net_h = acc_handle_conn(pathterm_h);
        while (port_h = acc_next_port(net_h, port_h)) {
            io_printf("  Port %s\n", acc_fetch_name(port_h));
        }
    }
}
acc_close();
return(0);
}

```

15.8 Traversing hierarchy across module ports

Verilog hierarchy connections can be traversed across module boundaries by following the connections within and without module ports. A module port has two connections, as shown in the following diagram.




The *loconn* (hierarchically lower connection) is the signal *inside* the module that is connected to a port. In the object diagrams, the internal connection is shown as an *expression*, because a number of different types of objects can be connected to the port. Within a module, the Verilog HDL syntax restricts the expression connected to a port to the Verilog data types of *nets*, *regs* and *variables* (except real variables). The expression connected to the port can be a scalar signal (1-bit wide), a vector, a bit or part select of a vector, or a concatenation of any these signals. For the loconn connection, any of these expression types can be connected to an output port, but only expressions using the *net* class of data types can be connected to an input or inout port.

The ***hiconn*** (hierarchically higher connection) is the signal *outside* the module that is connected to a port. External to a module, the expression connected to the port can be a scalar signal, a vector, a bit or part select of a vector, or a concatenation of these signals. Any of these expression types can be connected to an input port, but only expressions using the *net* class of data types can be connected to an output or inout port. (The Verilog language also allows the *hiconn* connection of an input port to be a constant, a literal value, an operation, or the return of a function call, but the ACC routines cannot access these types of objects).

The Verilog HDL syntax allows concatenation multiple internal signals to be connected to *loconn* of a module port. Each signal can have a different data type, and the bits of the *loconn* connections can have different port directions. The ACC routines refer to a port with a mix of signal directions connected to it as a *mixed I/O* port.

The object diagram for `ports`, shown in Figure 15-6, shows a one-to-many connection to either the *hiconn* or *loconn* connection to the port. The connecting object is an `expr` object group, to allow for the different types of objects which may be connected to the port. Using the handle for the `expr` object, the type property of the object can be accessed using either `acc_fetch_type()` or `acc_fetch_fulltype()` to determine what is connected to the port. The object diagram for that type of object can then be used to continue traversing the Verilog design hierarchy. Accessing the object type property is presented previously, in Chapter 14, section 14.7 on page 454.

Figure 15-6: ACC object diagram for ports (partial)



Both ACC handle routines and ACC next routines can be used to obtain the handles for the hiconn and loconn connections of a port. This provides flexibility for different PLI applications. Table 15-3, on the following page, shows what type of object will be returned, based on different types of port connections and the ACC routine used. The terms used in this table are defined after the table.

Routine	Port Property	Connection Handle
acc_handle_loconn() acc_handle_hiconn()	scalar port	scalar connection
	expanded vector port	vector connection
	unexpanded vector port	vector connection
	bit select of a port	bit select of connection
acc_next_loconn()	scalar port	scalar connection
	expanded vector port	bit select of connection
	unexpanded vector port	vector connection
	bit select of a port	illegal
acc_next_hiconn()	scalar port	scalar connection
	expanded vector port	vector connection
	unexpanded vector port	vector connection
	bit select of a port	illegal

Table 15-3: Port connection handles

Bit-select, part-select and concatenation ports

When a bit-select or part-select of a vector signal is connected to a module port, the port type inherits this property. If a concatenation of several signals is connected to a port, this property is also reflected in the port type. The type property of the port is accessed, using `acc_fetch_fulltype()`, which was presented previously, in Chapter 14, section 14.7 on page 454.

The ACC constants which represent ports are:

- **accScalarPort** — The port has a scalar signal attached to it.

```
module m1 (a);
  input a;
  wire a;
  ...

```

- **accVectorPort** — The port has a vector signal attached to it.

```
module m1 (a);
  input [7:0] a;
  wire [7:0] a;
  ...

```

- **accBitSelectPort** — The port has a bit-select of a vector signal attached to it.

```
module m1 (a[0]);
  input [7:0] a;
  wire [7:0] a;
  ...

```

- **accPartSelectPort** — The port has a part-select of a vector signal attached to it.

```
module m1 (a[7:4], a[3:0]);
  input [7:0] a;
  wire [7:0] a;
  ...

```

- **accConcatPort** — The port has a concatenation of signals attached to it.

```
module m1 (.a{n,m});
  input [3:0] n;
  output [3:0] m;
  wire [3:0] n, m;
  ...

```

Expanded and unexpanded vectors

A module port may be any vector size. If a vector port has a special *expanded* property, then the ACC routines are allowed to access discrete bits of the vector. If the port is *unexpanded*, then certain ACC routines, such as `acc_next_loconn()`, may be restricted to only accessing the complete vector. The expanded or unexpanded property is inherited from the type of loconn signal connected to the port. The routine `acc_object_of_type()` is used to determine if a port is expanded or unexpanded. This routine is presented in Chapter 16, in section 16.2 on page 526.

Expanded vectors are vectors for which the simulator allows access to individual bits within the vector. Access to the bits of a vector is the default in the Verilog language, and can also be explicitly declared in the Verilog source code, using the `scalared` keyword. For example:

```
wire scalared [63:0] data_bus;
```

Unexpanded vectors are vectors for which the simulator may prohibit access to individual bits within the vector. This is a feature which some simulators use to optimize run-time performance. Vector only access must be explicitly declared in the Verilog source code, using the `vectored` keyword. For example:

```
wire vectored [31:0] address_bus;
```

An example of traversing hierarchy across module ports

The following example accesses all ports of a module and lists the port name, port size, port direction, and type of object connected as the *loconn* and *hiconn*. The name, size and direction are all properties of the module port.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.015/port_info_acc.c
- Verilog test bench: Chapter.015/port_info_test.v
- Verilog-XL results log: Chapter.015/port_info_test.log

Example 15-2: \$port_info — traversing Verilog hierarchy across module ports

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/*********************************************************************
 * calltf application
 *****/
void PLIbook_PortInfo_calltf()
{
    handle mod_h, port_h, loconn_h, hiconn_h;

    acc_initialize();
    acc_configure(accDisplayWarnings, "true");

    mod_h = acc_handle_tfarg(1);

    io_printf("\nModule %s:\n", acc_fetch_defname(mod_h));
    io_printf("  Instance name: %s\n", acc_fetch_fullname(mod_h));

    switch ( acc_fetch_fulltype(mod_h) ) {
        case accTopModule:
            io_printf("  Module type: top-level\n");
            break;
        case accModuleInstance:
            io_printf("  Module type: module instance\n");
            break;
        case accCellInstance:
            io_printf("  Module type: cell module\n");
            break;
        default:
            io_printf("  Module type: unknown\n");
    }

    io_printf("  Ports:\n");
    port_h = null;
    while (port_h = acc_next_port(mod_h, port_h)) {
        io_printf("%-8s", acc_fetch_name(port_h) );
        io_printf("%2d-bit ", acc_fetch_size(port_h) );
    }
}
```

```
switch ( acc_fetch_direction(port_h) ) {
    case accInput:
        io_printf("input    ");
        break;
    case accOutput:
        io_printf("output   ");
        break;
    case accInout:
        io_printf("inout   ");
        break;
    case accMixedIo:
        io_printf("mixed input/output   ");
        break;
    default:
        io_printf("unknown direction   ");
}

if (acc_object_of_type(port_h, accExpandedVector))
    io_printf(" Expanded=true   ");
else
    io_printf(" Expanded=false   ");
if (acc_object_of_type(port_h, accUnExpandedVector))
    io_printf(" Unexpanded=true\n");
else
    io_printf(" Unexpanded=false\n");

loconn_h = acc_handle_loconn(port_h);
io_printf("      Loconn type = %s\n",
         acc_fetch_type_str(acc_fetch_fulltype(loconn_h)));


hiconn_h = acc_handle_hiconn(port_h);
if (hiconn_h)
    io_printf("      Hiconn type = %s\n\n",
              acc_fetch_type_str(acc_fetch_fulltype(hiconn_h)));
else
    io_printf("      Hiconn type = none\n\n");
} /* end of next port_h loop */
acc_close();
return;
}
```

15.9 Identifying modules and library cells

In the Verilog language, modules are building blocks which make up a larger design. A design hierarchy tree is formed when one module instantiates another module.

A PLI application often needs to distinguish one level of a design hierarchy from others. An ASIC library vendor, for example, may need to locate the cells from the ASIC library within the rest of the Verilog hierarchy, in order to calculate cell delays or power usage. Figure 15-7 shows a simple hierarchy tree. A PLI application might wish to traverse from the top of the design, and locate the module instances of *cellA*, *cellB*, and *cellC*. To access *cellA* and *cellB*, the PLI application will need to traverse through *modA*. The same PLI application might wish to consider *cellC* a leaf in the hierarchy tree, and ignore the hierarchy below that leaf (the instance of *modB* in this example).

Figure 15-7: Example Verilog design hierarchy tree



For convenience in developing PLI applications, the Verilog language provides a special flag which can be placed in the Verilog source code. This flag designates that the modules which follow the flag are to be treated as cell modules. The flag has no effect on simulation. It is strictly a flag which is used by certain ACC routines to distinguish one type of module from another. The '**'celldefine**' compiler directive is used to set the cell module flag, and the '**'endcelldefine**' directive is used to terminate the flag. All modules which are compiled between the two directives will be flagged as cell modules.

NOTE

The IEEE 1364 standard states that any Verilog module which is loaded into simulation using a library scan option should automatically be flagged as a cell module. In most Verilog simulators, the library scan options are `-v`, `-y`, and the `'uselib` compiler directive. Some simulators provide special invocation options to enable or disable this automatic flagging of modules loaded from a library as being a cell modules.

The ACC library provides two different ACC next routines to find module instances within a module:

handle acc_next_child (module, prev_child)

handle module handle for a module, or *null*.

handle prev_child handle for the previous child found; initially *null*.

The `acc_next_child()` routine returns handles for all instances of any module type within a reference module. If the module handle is *null*, then the next top-level module is returned. In the diagram shown above, if `acc_next_child()` were passed a reference handle for `topmod`, it would retrieve handles for `modA` and `cellC`. Those are the modules which are instantiated within `topmod`.

handle acc_next_cell (module, prev_cell)

handle module handle for a module.

handle prev_cell handle for the previous cell found; initially *null*.

The `acc_next_cell()` routine returns handles for all instances of modules which are flagged as cells. This routine automatically traverses all levels of hierarchy, starting with the reference module. In the hierarchy tree shown in Figure 15-7 on page 499, if `acc_next_cell()` were given a reference handle for `topmod`, it would retrieve handles for `cellA`, `cellB` and `cellC`. Those are the cell modules which are instantiated within and below `topmod`.

Example 15-3 lists a short PLI application which looks for all top-level modules in a design, and then searches all levels of hierarchy below that top level and prints the name of every cell module used in the design.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.015/list_cells_acc.c`
- Verilog test bench: `Chapter.015/list_cells_test.v`
- Verilog-XL results log: `Chapter.015/list_cells_test.log`

Example 15-3: \$list_cells — working with module cells

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"           /* IEEE 1364 PLI ACC routine library */
/*********************************************************************
 * Checktf application
 *****/
int PLIBbook_ListCells_checktf()
{
    acc_initialize();
    if (tf_nump() != 1)
        tf_error("$list_cells must have 1 argument.");
    else if (tf_typep(1) == TF_NULLPARAM)
        tf_error("$list_cells arg cannot be null.");
    else if (acc_fetch_type(acc_handle_tfarg(1)) != accModule)
        tf_error("$list_cells arg must be a module instance.");
    acc_close();
    return(0);
}

/*********************************************************************
 * Calltf application
 *****/
int PLIBbook_ListCells_calltf()
{
    handle mod_h, cell_h;
    int     cell_cnt = 0;

    acc_initialize();
    acc_configure(accDisplayWarnings, "true");

    mod_h = acc_handle_tfarg(1);
    io_printf("\nCells in Module %s, instance %s:\n",
              acc_fetch_defname(mod_h), acc_fetch_fullname(mod_h));
    cell_h = null; /* start with null (no cells found yet) */
    while (cell_h = acc_next_cell(mod_h,cell_h))
        io_printf("  %s (%s)\n",
                  acc_fetch_fullname(cell_h), acc_fetch_defname(cell_h));
        cell_cnt++;
    io_printf("Total cells in this hierarchy tree = %d\n\n", cell_cnt);
    acc_close();
    return(0);
}
```

15.10 Accessing loads and drivers

The ACC library provides three routines to find the drivers and loads of a Verilog net. These routines also distinguish between standard modules and cell modules.

handle acc_next_driver (net, prev_driver)

handle net handle for a scalar net or bit-select of a vector net.
handle prev_driver handle for the previous driver found; initially *null*.

The `acc_next_driver()` returns the handle for the next primitive terminal driver on a net.

handle acc_next_load (net, prev_load)

handle net handle for a scalar net or bit-select of a vector net.
handle prev_load handle for the previous load found; initially *null*.

The `acc_next_driver()` returns the handle for the next primitive terminal load on a net. If there are multiple loads within a module, all loads will be returned .


handle acc_next_cell_load (net, prev_load)

handle net handle for a scalar net or bit-select of a vector net.
handle prev_load handle for the previous load found; initially *null*.

This routine returns the handle for the next primitive terminal of a cell load on a net. Loads that are not in cell modules are not returned. Only the first load within a cell will be returned if there are multiple loads within the cell.

Both `acc_next_load()` and `acc_next_cell_load()` return primitive terminals as the loads of a net. The difference in the routines is which primitive terminals are considered to be loads. If a net fans out to multiple loads within a module, `acc_next_load()` returns a handle for each load. `acc_next_cell_load()` is more discriminatory. This routine only returns the primitive terminals which load a net, if the primitive is within a module that was flagged as a cell, and only one load per cell module will be returned. The following diagram illustrates the difference between the loads returned from each of these routines.

NOTE → `acc_next_driver()`, `acc_next_load()` and `acc_next_cell_load()` only locate primitive input terminals as the drivers or loads of a net. The VPI routines will locate other types of drivers and loads, such as continuous assignment statements.

Figure 15-8: Difference of acc_next_load() and acc_next_cell_load()

15.11 Accessing model timing

The ACC library provides several routines designed specifically to provide access to the timing information within Verilog simulation. This timing information is specific to each *instance* of a module, rather than to each definition of a module. By reading and/or modifying the timing of each instance, PLI applications can greatly increase the timing accuracy of a simulation.

The PLI can access the following timing constructs within a Verilog simulation:

- Primitive delays
- Module path delays
- Timing constraint checks
- Module inter-connect delays

The ACC library provides several routines to obtain handles for these timing objects, which are presented on the following pages. There are also special ACC routines for reading and modifying the values of the delays on each object. Reading and modifying delays is presented in the next chapter.

15.11.1 Obtaining handles for primitives

The Verilog language allows a unique set of delays to be specified on each instance of a primitive. The syntax for specifying primitive delays is part of the Verilog language, and is not presented in this book. From the perspectives of the PLI, there are only two major considerations for accessing primitive delays:

- Verilog primitives can be two-state or three-state devices. A two-state primitive has timing for output transitions from 0 to 1 and from 1 to 0. A three-state primitive has timing for transitions from 0 to 1, 1 to 0, and from any value to hi-impedance.
- Primitive delays can be represented as a single value for each possible transition, or as a minimum:typical:maximum set of delays for each transition.

To access the delays of a primitive, a handle for the primitive must be obtained. The routines which can obtain handles for primitives are:

- `acc_next_primitive()` returns handles for all primitive instances within a module instance.
- `acc_handle_object()` returns a handle for a primitive instance, using the instance name of the primitive, searching in the design hierarchy scope in which the PLI is currently operating.
- `acc_handle_by_name()` returns a handle for a primitive instance, using the instance name of the primitive, searching in the design hierarchy scope specified as an input to the routine.

The syntax and usage of the ACC next routines was presented in Chapter 14, in section 14.6 on page 452. The syntax and usage of `acc_handle_object()` and `acc_handle_by_name()` is discussed in more detail later in this chapter, in section 15.14 on page 516.

15.11.2 Obtaining handles for module paths


The primary Verilog construct for representing the timing of a module is a *module path delay*. Module path delays are also referred to as *pin-to-pin delays*, or simply as *path delays*. These delays are specified within the specify block of a module. The full syntax for specifying module path delays is part of the Verilog language, and is outside the scope of this book. The PLI standard has special terms for the various components of a module path delay, which are presented in the following paragraphs. The components of a module path delay which are identified by the PLI are:

- The *module path*
- The path *input terminals*
- The path *output terminals*


- The ***data path***
- The data path ***source terminals***
- The data path ***destination terminals***
- The path ***conditional expression***

The following diagrams illustrate these components of a module path delay.


Module path delay with single inputs and outputs:




Module path delay with multiple inputs and outputs:



Edge sensitive module path delay:



Conditional module path delay:



There are several ACC handle and ACC next routines used to obtain handles for the different components of a module path. The ACC handle and ACC next routines which obtain handles for module paths and the components of a module path are:

- **acc_next_modpath()** returns handles for all module paths within a module.
- **acc_handle_modpath()** returns a handle for a specific module path within a module instance, using a description of the input and output terminals of the path.
- **acc_handle_pathin()** returns a handle for the first input terminal of a module path.
- **acc_handle_pathout()** returns a handle for the first output terminal of a module path.
- **acc_handle_condition()** returns a handle for the condition expression of a module path or data path.
- **acc_next_input()** returns handles for all input terminals of a module path, or all source terminals of a data path.
- **acc_next_output()** returns handles for all output terminals of a module path, or all destination terminals of a data path.
- **acc_handle_conn()** returns a handle for the net connected to a module path terminal or data path terminal.

The general syntax and usage of the ACC handle routines and ACC next routines was presented in Chapter 14, in sections 14.5 and 14.6.

The ACC routines listed above allow a PLI application to obtain the module path timing information for any module instance in a design. Example 15-1 on page 492, presented earlier in this chapter, illustrates accessing all module paths in a module, and all output terminals of each module path. After a handle is obtained, other ACC routines can then fetch information about the path, such as the polarity and condition edges. Additional ACC routines can read and modify the delay values of each module path, for each instance of the module containing the path. The routines to read and modify module path delays are presented in the next chapter.

Configuring acc_handle_modpath()

The routine `acc_handle_modpath()` uses handles for the path input and output terminals to obtain a handle for the module. This routine can be configured to describe the terminals in different ways. The syntax of this routine is:

handle acc_handle_modpath(object, src_name, dest_name, src_handle, dest_handle)

handle object handle for a module.
char * src_name name of net connected to path source, or *null*.
char * dest_name name of net connected to path destination, or *null*.
handle src_handle (optional) handle for net connected to path source.
handle dest_handle (optional) handle for net connected to path destination.

The `acc_handle_modpath()` routine uses the `acc_configure()` routine to control which arguments are used to describe the path terminals. The syntax for `acc_configure()` was presented earlier in this chapter, in section 15.2 on page 475.

- `acc_configure(accEnableArgs, "no_acc_handle_modpath")` configures `acc_handle_modpath()` to use only the `src_name` and `dest_name` arguments. A literal string or pointer to a string is provided for the arguments. The `src_handle` and `dest_handle` arguments are ignored, and can be set to *null* or dropped from the argument list. This is the default configuration.
- `acc_configure(accEnableArgs, "acc_handle_modpath")` configures `acc_handle_modpath()` to use the `src_name` and `dest_name` arguments, if a literal string or pointer to a string is provided for the arguments. If the name arguments are set to *null*, then the `src_handle` and `dest_handle` arguments are used.

Memory allocation for acc_next_input() and acc_next_output()


The routines `acc_next_input()` and `acc_next_output()` require special attention in a PLI application. Some simulators may need to allocate memory in order to process these routines. This memory will automatically be freed once the PLI application has retrieved the last terminal, which is indicated when the ACC next routine returns a *null*. However, if a PLI application does not call these routines in a loop until the routine returns a *null*, then the memory is not released. The PLI application must then manually free the memory by calling `acc_release_object()`. The syntax of this routine is:

void acc_release_object(object)

handle object handle for a module path or data path terminal.

15.11.3 Obtaining handles for module inter-connect paths

A module inter-connect path is a connection from the output port of one module to the input port of another module.



In this illustration, there is a single net **n1**, but two inter-connect paths, **out** to **in1**, and **out** to **d**.

The `acc_handle_path()` routine obtains a handle for an inter-connect path, using handles for an output port and the input port. The syntax for this routine is:

handle acc_handle_path (source, destination)

handle source handle for a scalar output or inout port, or bit-select of a vector output or inout port.

handle destination handle for a scalar input or inout port, or bit-select of a vector input or inout port.

There is no construct in the Verilog language to represent inter-connect delays. However, using an inter-connect path handle, a PLI application can place delays on the inter-connection and read the value of those delays. Reading and modifying delays is presented in Chapter 14.

15.11.4 Obtaining handles for timing constraint checks

Verilog modules can specify timing constraints between one or more input ports of the module, such as setup and hold times. These constraints are specified in a Verilog `specify` block, and use a special form of built-in system tasks. The names of these tasks are: `$setup`, `$hold`, `$setuphold`, `$skew`, `$recovery`, `$period`, and `$width`. The PLI uses specific terminology to identify the arguments of a Verilog timing constraint.

Example timing check:

```
$setup(data, posedge clock && (reset), 10, err_flag);
```

The diagram shows six arrows pointing from labels to specific parts of the code:

- 1st tchk arg terminal points to the first argument, "data".
- edge specifier points to the second argument, "posedge clock".
- 2nd tchk arg terminal points to the third argument, "(reset)".
- condition (optional) points to the fourth argument, "10".
- tchk limit points to the fifth argument, "err_flag".
- notifier reg (optional) points to the sixth argument, which is missing.

The PLI can access the timing constraint checks for each instance of a module, and can modify the constraint values of each instance. Handles for the optional condition and notifier fields of a timing check can also be accessed, as well as the type of edge specifier.

The Verilog HDL syntax of timing constraint checks is not consistent. Some checks have two terminals, like the `$setup` timing check shown above; and some checks have just one terminal. Most timing checks have a single tchk limit, as shown above, but some timing checks specify two limits. To allow for these variations of syntax in the Verilog language, the ACC library provide specific routines to access the individual fields of each timing check:

- `acc_next_tchk()` returns handles for all timing checks within a module instance.
- `acc_handle_tchk()` returns a handle for a specific timing check, using a description of the timing check arguments.
- `acc_handle_tchkarg1()` returns a handle for the first terminal of a timing check.
- `acc_handle_tchkarg2()` returns a handle for the second terminal of a timing check.
- `acc_handle_conn()` returns a handle for the net connected to a timing check argument.

The general syntax and usage of the ACC handle routines and ACC next routines was presented in Chapter 14, in sections 14.5 and 14.6. After a handle is obtained, other ACC routines can then fetch information about the constraint, such as the polarity and condition edges. Additional ACC routines can read and modify the values of each timing constraint, for each instance of the module containing the constraint. The routines to read and modify module constraint values are presented in the next chapter.

NOTE

The proposed IEEE 1364-1995 standard will add additional timing constraint checks to the Verilog language, and modify the syntax of some existing timing constraint checks. These enhancements will increase the accuracy for modeling deep-submicron technologies. The ACC library will not be enhanced for these new timing constraint checks. Only the VPI library will support these changes.

Configuring acc_handle_tchk()

The routine `acc_handle_tchk()` uses a description of the timing check arguments to obtain a handle for the timing check. This routine can be configured to describe the timing check arguments in different ways. The syntax for `acc_handle_tchk()` is:

`handle acc_handle_tchk(object, type, name1, edge1, name2, edge2,
conn1, conn2)`

<code>handle</code>	<code>object</code>	handle for a module.
<code>int</code>	<code>type</code>	constant representing timing check type. One of: <code>accHold</code> , <code>accNochange</code> , <code>accPeriod</code> , <code>accRecovery</code> , <code>accSetup</code> , <code>accSkew</code> or <code>accWidth</code> .
<code>char</code>	<code>* name1</code>	name of net connected to the 1st tchk argument, or <code>null</code> .
<code>int</code>	<code>edge1</code>	constant representing the edge of the 1st timing check argument (constant names are listed below).
<code>char</code>	<code>* name2</code>	name of net connected to the 2nd tchk argument, or <code>null</code> .
<code>int</code>	<code>edge2</code>	edge of the 2nd tchk arg (constant names are listed below).
<code>handle</code>	<code>conn1</code>	(optional) handle for the net connected to the 1st tchk argument.
<code>handle</code>	<code>conn2</code>	(optional) handle for the net connected to the 2nd tchk argument.

`edge1` and `edge2` identifiers are one of the following:

- One of the constants: `accNoedge`, `accPosedge`, `accNegedge`
- List of constants separated by +: `accEdge01`, `accEdge0x`, `accEdgex1`
- List of constants separated by +: `accEdge10`, `accEdge1x`, `accEdgex0`.

`acc_handle_tchk()` uses the `acc_configure()` routine to control which arguments are used to describe the path terminals. The syntax for `acc_configure()` was presented earlier in this chapter, in section 15.2 on page 475.

- `acc_configure(accEnableArgs, "no_acc_handle_tchk")` configures the `acc_handle_tchk()` routine to use only the `name1` and `name2` arguments. A literal string or pointer to a string is provided for the arguments. The `conn1` and `conn2` arguments are ignored, and can be set to `null` or dropped from the argument list. This is the default configuration.
- `acc_configure(accEnableArgs, "acc_handle_tchk")` configures the `acc_handle_tchk()` routine to use the `name1` and `name2` arguments, if a literal string or pointer to a string is provided for the arguments. If the name arguments are set to `null`, then the `conn1` and `conn2` arguments are used.

15.12 Counting the number of objects

The ACC library provides a convenient routine for counting how many of a specific type of object exist within a reference object.

int acc_count(next_routine, object)

handle *next_routine name of any ACC *next* routine except acc_next_topmod()
 handle object reference object for the ACC *next* routine.

The `acc_count()` routine returns the number of objects which were located by an ACC *next* routine. The inputs to `acc_count()` are the name of a *next* routine, and a reference point in which that routine should search for its target objects. The handles for the objects are not returned.

An example of using `acc_count()` is shown in example 15-4. This example implements a `$count_loads` system function, which counts how many cell loads are driven by an output port of a module, and then returns that count back to the simulation as a system function return. `$count_loads` is passed a module port name as its input.

An example of using `$count_loads` might be:

```
if ($count_loads(chip_out) > 5)
    $display("Warning: output port has too many loads!");
```



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.015/count_loads_acc.c
- Verilog test bench: Chapter.015/count_loads_test.v
- Verilog-XL results log: Chapter.015/count_loads_test.log

Example 15-4: `$count_loads` — using the `acc_count()` routine

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"           /* IEEE 1364 PLI ACC routine library */
/*********************************************
 * Sizetf application
 *****/
int PLIbook_CountLoads_sizetf()
{
    return(0);
}
```

```
*****
 * Checktf application
 ****
int PLIbook_CountLoads_checktf()
{
    static int valid_args[4] = {accNet, accReg, accRegBit, 0};
    int     direction;
    handle tfarg_h, port_h;

    acc_initialize();
    if (tf_numargs() != 1)
        tf_error("$count_loads must have 1 argument.");
    else if (tf_typep(1) == TF_NULLPARAM)
        tf_error("$count_loads arg cannot be null.");
    /* acc_handle_tfarg() returns a loconn handle, not a port handle */
    else if (tf_sizep(1) != 1)
        tf_error("$count_loads arg must be scalar or a bit-select.");
    else {
        tfarg_h = acc_handle_tfarg(1);
        if (!acc_object_in_typelist(tfarg_h, valid_args)) {
            tf_error("$count_loads arg must be a net or reg signal");
            return(0);
        }
        port_h = acc_next_port(tfarg_h, null);
        if (port_h == null) {
            tf_error("$count_loads arg is not connected to a module port.");
            return(0);
        }
        direction = acc_fetch_direction(port_h);
        if (direction != accOutput
            && direction != accInout)
            tf_error("$count_loads arg must be an output or inout port.");
    }
    acc_close();
    return(0);
}

*****
 * Calltf application
 ****
int PLIbook_CountLoads_calltf()
{
    handle loconn_h, port_h, hiconn_h;
    int     load_count;

    acc_initialize();
    acc_configure(accDisplayWarnings, "true");

    /* acc_handle_tfarg() returns a loconn handle, not a port handle */
    loconn_h = acc_handle_tfarg(1);
    port_h = acc_next_port(loconn_h, null);
    hiconn_h = acc_handle_hiconn(port_h);
    load_count = acc_count(acc_next_cell_load, hiconn_h);
```

```
    tf_putstr(0, load_count);
    acc_close();
    return(0);
}
```

15.13 Collecting and maintaining lists of object handles

Sometimes a PLI application may need to use the same ACC next routine to find the same target objects multiple times in the same call to a PLI application. For example, an application might need to find all of some type of object in a module (such as all the primitives) in order perform some type of operation. Then later, in the same call to the PLI application, it might be necessary to again find all of the same objects in the same module, in order to perform some other type of operation.

Sometimes an instance of a PLI application might be called many times during a simulation. For example, an application that prints the logic value of all drivers of a net might be called every clock cycle. In this case, the same ACC next routine might be called millions of times during a simulation. Each call will need to locate and return the same set of objects.

Repeatedly calling the same ACC next routine to locate the same objects can have a negative impact on simulation run-time performance. The ACC library provides a pair of routines to help make a PLI application more efficient when the same ACC next routine needs to be called multiple times.

handle *acc_collect(next_routine, object, count)

*handle *next_routine* name of any ACC *next* routine except acc_next_topmod().

handle object reference object for the ACC *next* routine.

*int *count* pointer to variable to receive number of objects collected.

The **acc_collect()** routine is provided the name of an ACC next routine, a handle for a reference object, and a pointer to an integer variable. The routine will execute the specified ACC next routine and collect a list of all of the target handles. The list of handles is stored in an array of handles, and a pointer to the array is returned. The total number of objects found is placed the integer variable pointed to as the third argument to **acc_collect()**.

Once the array of object handles has been collected, a PLI application can reference the array as often as needed, without having to call the ACC next routine again and again. The array is stored in persistent memory, which is allocated by

`acc_collect()`. Therefore, as long as the pointer for the array is saved by the PLI application, the array can be used as often as needed. If the PLI application saves the pointer to the array, then the array can be used in future calls to the PLI application as well. The TF work area is the proper place to preserve the pointer to the array of object handles. The work area is specific to each instance of a system task/function, and is shared by the *checktf routine*, *calltf routine* and *misctf routine*. Refer to section 10.7 on page 313 of Chapter 10, for details on using the TF work area.

Because the memory for the array of object handles allocated by `acc_collect()` is persistent, the PLI application must explicitly release the memory when the array is no longer required. The `acc_free()` routine is used to release the memory which was allocated by `acc_collect()`. If `acc_free()` is not called, the memory allocated for the array of handles will not be released until the simulation exits. The syntax for `acc_free()` is:

```
void acc_free(array_ptr)
handle *array_ptr    pointer to an array of handles.
```

The following C code shows how a list of handles for all signals in a module can be created, using `acc_collect()`. The array of signal handles is generated at the start of simulation, when the *misctf routine* is called for REASON_ENDOFCOMPILE, and the array memory is freed at the end of simulation, when the *misctf routine* is called for REASON_FINISH. The pointer to the array is stored in the TF work area, so that the *calltf routine* can access the array whenever needed.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.015/display_all_acc.c
- Verilog test bench: Chapter.015/display_all_test.v
- Verilog-XL results log: Chapter.015/display_all_test.log

Example 15-5: \$display_all_nets — using the `acc_collect()` routine

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/*********************  

 * structure definition for data to be passed from misctf to calltf  

 ****/  

typedef struct PLIbook_NetData {
    handle module_h;
    handle *net_array;
    int net_count;
} PLIbook_NetData_s, *PLIbook_NetData_p;
```

```
*****
 * checktf application
 *****/
int PLIbook_DisplayNets_checktf()
{
    acc_initialize();
    if (tf_nump() != 1)
        tf_error("$display_all_nets must have 1 argument.");
    else if (tf_typep(1) == TF_NULLPARAM)
        tf_error("$display_all_nets arg cannot be null.");
    else if (acc_fetch_type(acc_handle_tfarg(1)) != accModule)
        tf_error("$display_all_nets arg must be a module instance.");
    acc_close();
    return(0);
}

*****
 * misctf application
 *****/
int PLIbook_DisplayNets_misctf(int user_data, int reason)
{
    PLIbook_NetData_p net_data;
    handle mod_h;

    acc_initialize();
    acc_configure(accDisplayWarnings, "true");

    switch (reason) {
        case REASON_ENDOFCOMPILE:
            acc_initialize();
            net_data = (PLIbook_NetData_p)malloc(sizeof(PLIbook_NetData_s));
            net_data->module_h = acc_handle_tfarg(1);
            net_data->net_array = acc_collect(acc_next_net,
                                              net_data->module_h,
                                              &net_data->net_count);
            io_printf("Total nets collected in misctf routine = %d\n",
                      net_data->net_count);
            tf_setworkarea((char *)net_data);
            acc_close();
            break;
        case REASON_FINISH:
            net_data = (PLIbook_NetData_p)tf_getworkarea();
            acc_free(net_data->net_array);
            break;
    }
    acc_close();
    return(0);
}

*****
 * calltf application
 *****/
int PLIbook_DisplayNets_calltf()
{
```

```

PLIbook_NetData_p net_data;
int i;

acc_initialize();
acc_configure(accDisplayWarnings, "true");

net_data = (PLIbook_NetData_p)tf_getworkarea();

io_printf("\nAt time %s, nets in module %s (%s):\n",
          tf_strgettime(),
          acc_fetch_fullname(net_data->module_h),
          acc_fetch_defname(net_data->module_h));
for (i=0; i<net_data->net_count; i++) {
    io_printf(" %-13s value is %s (hex)\n",
              acc_fetch_name(net_data->net_array[i]),
              acc_fetch_value(net_data->net_array[i], "%h", null));
}
acc_close();
return(0);
}

```

15.14 Obtaining object handles using an object's name

The ACC library contains two routines which can obtain a handle for an object, using the name of the object as the reference point.



TIP Obtaining an object handle from the name of the object is an expensive process for simulation performance, and should be used judiciously. It is much more efficient to obtain a handle for an object based on its relationship to some other object.

handle acc_handle_by_name(obj_name, scope)

*char * obj_name* name of an object as a string.

handle scope handle for a scope, or *null*.

The *acc_handle_by_name()* routine is used to obtain a handle for an object, using the name of the object. The routine requires two inputs, a string containing the object's name, and a handle for a Verilog hierarchy scope. A scope in the Verilog hierarchy can be a top-level module, a module instance, a named statement group, a Verilog task or a Verilog function. The scope argument can also be set to *null*, which represents the top level of Verilog hierarchy.

The name provided can be the local name of the object, a relative hierarchical path name, or a full hierarchical path name. The PLI will search for the object, using the same search rules as the Verilog language. The full search rules are outside the scope

of this book, but briefly, Verilog searches for a name in the hierarchy scope specified by the scope handle first, then as a relative path, then as a full path. If the object cannot be found, then `acc_handle_by_name()` will return `null`.

The objects for which `acc_handle_by_name()` routine can obtain a handle are listed in Table 15-4.

Modules	Parameters
Primitives	Specparams
Nets	Named blocks
Registers	Verilog HDL tasks
Integer, time and real variables	Verilog HDL functions
Named events	

Table 15-4: Objects supported by `acc_handle_by_name()`

The ACC library provides a second routine which obtains the handle for an object using the object's name.

handle acc_handle_object(obj_name)
*char * obj_name* name of an object.

`acc_handle_object()` is similar to `acc_handle_by_name()`, but differs in where it will search for the object. The `acc_handle_object()` routine will begin searching for the object in the design hierarchy scope in which the PLI is executing. By default, the PLI scope is the Verilog design hierarchy scope from which the PLI application was called. The `acc_set_scope()` routine can be used to change the PLI scope to another point in the Verilog design hierarchy. The `acc_handle_object()` routine will search for the object, using the search rules of the Verilog language. If the object cannot be found by searching from the PLI scope, then a `null` is returned.

The `acc_handle_object()` routine can obtain the handle for the objects listed in Table 15-5. Note that there are three objects for which handles can be obtained using this routine which cannot be obtained with the `acc_handle_by_name()` routine. These are module ports, module paths and data paths.

Modules	Integer, time and real variables
Module ports	Named events
Module paths	Parameters
Data paths	Specparams
Primitives	Named blocks
Nets	Verilog HDL tasks
Registers	Verilog HDL functions

Table 15-5: Objects supported by acc_handle_object()

The C code fragment listed below illustrates obtaining an object handle using the object's name. The name of a primitive instance and a delay value are read from a file, and then the handle for the primitive is obtained, using `acc_handle_by_name()`.

```
char      prim_name[64];
handle   prim_handle;
double   new_delay;

fscanf(file_p, "%s %f", prim_name, &new_delay);
prim_handle = acc_handle_by_name(prim_name, null);
if (prim_handle)
    /* add new delay value to the primitive object */
else
    /* error: primitive not found */
```

15.14.1 Obtaining handles for ports using a port name

The `acc_handle_object()` routine can obtain handles for module ports. However, there is a potential risk when using port names for which a PLI application must make provisions. In most Verilog designs, a module port will be connected to an internal signal of the same name. For example:

```
module my_chip (in1, in2, out);
    input in1, in2;
    output out;
    wire in1, in2, out;
    ...
endmodule
```

In this example, there are two objects with the same name of `in1` within the module `my_chip`, the port and the net. If `acc_handle_object()` were called to obtain a

handle for the name “**in1**”, the IEEE 1364-1995 standard is ambiguous about which object handle should be returned. Most Verilog simulators will return the handle for the internal signal (the loconn), but this is not guaranteed behavior. Therefore, when obtaining handles for either a signal or a module port using the object’s name, the PLI application should always check the type of object for which a handle was returned. If the handle references the wrong object, the appropriate ACC handle or next routine should be used to traverse from the object which was returned to the object desired.

15.14.2 Obtaining handles for module paths and data paths using a name

The `acc_handle_object()` routine can obtain handles for two objects which do not have names. This is done by creating a derived name for the objects. These objects are module paths and data paths. To create the derived name, the name of the signal connected to the input and the name signal connected to the output of the path are concatenated together, with a dollar sign (\$) character between the two names.

In the following example, the derived name of the module path is “*a\$b*”.

```
(a *-> b) = 2.5;
```

15.14.3 Changing the PLI hierarchy scope

The `acc_handle_object()` routine will search for an object within a specific hierarchy scope of the Verilog design hierarchy. This routine will use the *PLI scope* as the reference point for where in the Verilog design hierarchy to search for objects.

When a PLI application is called, the PLI scope is the Verilog HDL hierarchy scope from which the PLI application was called. A Verilog hierarchy scope may be: a top-level module, a module instance, a Verilog HDL task, a Verilog HDL function, a named begin—end statement group, or a named fork—join statement group.

The `acc_set_scope()` routine is used to change the PLI scope to another point in the Verilog design hierarchy. The syntax of `acc_set_scope()` is:

`char *acc_set_scope(module, name)`

`handle module` handle for a module.

`char * name` (optional) name of a module.

`acc_set_scope()` uses the `acc_configure()` routine to control which arguments are used to describe the target hierarchy scope. The syntax for `acc_configure()` was presented earlier in this chapter, in section 15.2 on page 475.

- `acc_configure(accEnableArgs, "no_acc_set_scope")` configures the `acc_set_scope()` routine to only use the `module` argument. The `name` argument is ignored. This is the default configuration.
 - If a handle is specified in the `module` argument, the PLI scope will be set to the hierarchy scope of the module handle.
 - If the `module` argument is set to `null`, the PLI scope will be set to the first top-level module found in the Verilog design hierarchy.
- `acc_configure(accEnableArgs, "acc_set_scope")` configures the `acc_set_scope()` routine to use the `module` argument as the first choice, and the `name` argument as the second choice. This is the default configuration.
 - If a handle is specified in the `module` argument, the PLI scope will be set to the hierarchy scope of the module handle.
 - If the `module` argument is set to `null`, and a literal string or pointer to a string is specified in the `name` argument, the PLI scope will be set to the hierarchy scope specified in the `name` argument.
 - If the `module` argument is set to `null`, and the `name` argument is `null`, the PLI scope will be set to the first top-level module found in the Verilog design hierarchy.

The return from `acc_set_scope()` is a pointer to a string containing the full hierarchical path name of the new PLI scope. The return will be a null if an error occurred.

15.15 Comparing ACC handles

Handles for Verilog objects can be obtained several different ways. Occasionally, a PLI application might need to test to see if two handles reference the same object.

`acc_compare_handles()` returns a boolean `TRUE` if two handles reference the same object, and a `FALSE` if they do not. An ACC handle is an abstraction used to reference an object within simulation, and all ACC routines use this abstraction to access information about the object. This layer of abstraction allows PLI applications to be portable to any number of Verilog simulators, because the abstract handle is a layer between the PLI application and the internal data structures of the simulator. Since the handle for an object is an abstraction, it is not possible to determine handle equivalence using the C '`==`' operator.

15.16 Summary

The ACC routines in the PLI standard provide dynamic access to what is happening within a Verilog simulation. This access is achieved using an object oriented method, where most Verilog HDL constructs that can exist in a simulation data structure are treated as objects. The ACC routines use *handles* to reference these objects. A large number of ACC routines are provided to obtain handles for the various types of Verilog HDL objects, and other ACC routines retrieve information about the objects, such as the name of an object or the vector size of a net. This chapter has presented many of the 103 ACC routines in the PLI standard, and has shown how these routines are used. The following chapters include several complete examples of PLI applications which use many of the ACC routines that have been introduced in this chapter.

CHAPTER 16

Reading and Modifying Values Using ACC Routines

The ACC routines provide access to the simulation values of Verilog objects. These values include information about the object, logic values and delay values. This access allows a PLI application to both read and modify what is happening during a simulation. This chapter presents how to use the ACC routines which read and modify values.

The concepts presented in this chapter are:


- Reading the type, full type and special-type properties of objects
- Using ACC *fetch* routines
- Reading an object's source code file location
- Reading the simulation invocation commands
- Reading the values of system task and system function arguments
- Reading net, reg, variable and UDP logic values
- Modifying logic values
- Reading delay values
- Modifying delay values
- Reading parameter and specparam values
- Reading specparam attribute values
- Reading and modifying pulse control attribute values

16.1 Using ACC fetch routines

Once a handle for an object has been obtained, a number of different properties for the object can be retrieved. The ACC routines which retrieve properties of an object are referred to as ACC *fetch* routines. These routines are property-specific, so there is a different fetch routine for each type of property which can be accessed.

The object diagrams show what properties can be accessed for each object, and which fetch routine is used to access a property. This information is listed in a table below the diagram of the object. For example, the object diagram for modules contains the following table:

Figure 16-1: Partial ACC object diagram for Verilog modules



Related Routines:

acc_fetch_type()	returns accModule
acc_fetch_fulltype()	returns accTopModule , accModuleInstance , accCellInstance
acc_fetch_name()	returns the instance name of a module
acc_fetch_fullname()	returns the hierarchical path name of a module
acc_fetch_defname()	returns the definition name of a module
acc_fetch_delay_mode()	returns accDelayModeNone , accDelayModeZero , accDelayModeUnit , accDelayModePath , accDelayModeDistrib , accDelayModeMTM
acc_fetch_timescale_info()	returns the timescale of a module
acc_fetch_location()	returns source file name & line no. containing module instance

The ACC fetch routines are property-type specific, so there is a different ACC routines for most types of Verilog properties. There are 33 ACC fetch routines:

acc_fetch_argc()	acc_fetch_location()
acc_fetch_argv()	acc_fetch_name()
acc_fetch_attribute()	acc_fetch_paramtype()
acc_fetch_attribute_int()	acc_fetch_paramval()
acc_fetch_attribute_str()	acc_fetch_polarity()
acc_fetch_defname()	acc_fetch_precision()
acc_fetch_delay_mode()	acc_fetch_pulsere()
acc_fetch_delays()	acc_fetch_range()
acc_fetch_delays()	acc_fetch_size()
acc_fetch_direction()	acc_fetch_tfarg()
acc_fetch_edge()	acc_fetch_tfarg_int()
acc_fetch_fullname()	acc_fetch_tfarg_str()
acc_fetch_fulltype()	acc_fetch_timescale_info()
acc_fetch_index()	acc_fetch_type()
acc_fetch_itfarg()	acc_fetch_type_str()
acc_fetch_itfarg_int()	acc_fetch_value()
acc_fetch_itfarg_str()	

Most ACC fetch routines require as an input, a handle for the object for which the property is to be accessed. For example, to retrieve the name of a module, a handle for the module must be provided as an input to `acc_fetch_name()`.

```
handle module_handle;
char *module_name;

/* obtain a handle for a module */

module_name = acc_fetch_name(module_handle);
```

The full syntax for each ACC fetch routine is listed in Appendix D, and this information is not duplicated in this chapter. The objective of this chapter is to show how several of these ACC fetch routines can be used in PLI applications.

16.2 Reading object type properties

Every Verilog object which can be accessed by ACC routines has a *type* property and a *fulltype* property. In addition, certain Verilog objects have a *special-type* property, such as a flag to indicate if a net is scalar (1-bit wide) or vector (multiple bits wide).

These properties identify what Verilog object is referenced by a Verilog handle.

- The *type* property identifies the general type of an object. For example, all net data types have the same type property of accNet.
- The *fulltype* property identifies the specific type of an object. For example, a net data type will have one of the fulltype properties: accWire, accWand, accWor, accTri, accTriand, accTrrior, accTrireg, accTri0, accTri1, accSupply0 or accSupply1.
- The *special-type* property identifies special attributes an object might have. A net, for example, can have the special-type properties of accScalar, accVector, accExpandedVector and accCollapsedNet.

The ACC object diagrams for each Verilog object list the name of the type, fulltype and special-type constants for that object. Examples of these diagrams are shown in this chapter, and the complete set of diagrams are contained in Appendix D.

acc_fetch_type() retrieves the *type* property of an object, which identifies the general type of an object. The syntax of this routine is:

```
int acc_fetch_type(object)
handle    object      handle for an object.
```

The type property is an integer constant, such as accModule, accPort, accNet, accPrimitive, etc. This property can be used many different ways. One common usage is to verify that a handle which was obtained references the type of object expected. For example, the \$show_all_nets application requires that the first task/function argument be a module instance. The PLI application could verify that the argument is correct, using the following code fragment:

```
handle tfarg_handle;
tfarg_handle = acc_handle_tfarg(1);
if (acc_fetch_type(tfarg_handle) != accModule)
/* report error that argument is not correct */
```

acc_fetch_fulltype() retrieves the *fulltype* property of an object. The syntax is:

```
int acc_fetch_fulltype (object)
handle    object      handle for an object.
```

The fulltype property provides more detailed information about an object. For example, if a PLI application has obtained a handle for a Verilog net, the constant returned for `acc_fetch_type(net_handle)` is **accNet**, while the constants returned for `acc_fetch_fulltype(net_handle)` include **accWire**, **accWor**, **accWand**, etc. For a Verilog module, `acc_fetch_type(module_handle)` returns the constant **accModule**, and `acc_fetch_fulltype(module_handle)` returns one of the constants: **accTopModule**, **accModuleInstance**, **accCellInstance**.

The type and fulltype properties are represented by constants with integer values. Printing the type of an object directly would print the integer value of the constant, rather than the name of the constant. A useful routine for debugging problems in a PLI application is **acc_fetch_type_str()**. This routine takes a type or fulltype constant value as its input and returns a pointer to a string which contains the actual name of the constant. The syntax of this routine is:

```
char *acc_fetch_type_str(type)
int      type      type or fulltype constant.
```

An example of using the `acc_fetch_type_str()` routine to print the name of a type constant is:

```
int object_type;
object_type = acc_fetch_type(tfarg_handle);
if (object_type != accModule) {
    tf_error("Tfarg type of %s is illegal.\n",
            acc_fetch_type_str(object_type));
}
```

Testing for special-type properties

There are no ACC routines to fetch the special-type properties of objects. Instead, two ACC routines test to see if an object has a special-type property, and return true if the property exists and false if it does not. These true/false tests can be used with any of the type, fulltype or special-type constants.

bool acc_object_of_type(object, type)

handle object handle for an object.

int type type, fulltype or special-type property constant.

The `acc_object_of_type()` routine returns true if an object has a specific type, fulltype or special-type property. For example, the following test could be used to perform different operations if an object handle is referencing a module that is at top-level of the design hierarchy or a cell module at the bottom of the design hierarchy (the properties `accTopModule` and `accCellInstance` are fulltypes).

```
handle mod_handle;
/* add code to get a module handle */
if (acc_object_of_type(mod_handle, accTopModule))
    /* process top-level modules */
else if (acc_object_of_type(mod_handle, accCellInstance))
    /* process cell-level modules */
```

bool acc_object_in_typelist(object, type_list)

handle object handle for an object.

int *type_list static array of type, fulltype and special-type property constants.

The `acc_object_in_typelist()` routine returns true if an object has any of a list of type, fulltype or special-type properties. The list of types must declared as a static integer array which contains the type, fulltype or special-type constants. The last element in the array must be 0. The following example tests to see if a net is any type of wired logic, using the fulltype property constants which represent the Verilog net data types:

```
handle net_handle;
static int valid_types[5] = {accWand, accWor, accTriand,
                            accTrior, 0};
/* add code to get a net handle */
if (acc_object_in_typelist(net_handle, valid_types))
    /* process wired-logic nets */
```

The special-type properties

Certain Verilog objects have special-type properties which can be useful in a PLI application. For example, Verilog nets have the special-type properties **accScalar** and **accVector**, which indicate if the net is 1-bit wide or multiple bits wide. The special-type properties are represented with the following constants:

- **accScope** indicates that an object has its own hierarchy scope. Objects which can have hierarchy scope are modules, named begin—end statement groups, named fork—join statement groups, Verilog HDL tasks and Verilog HDL functions. Objects which have scope can have local reg, variables and parameters declared.
- **accModPathHasIfnone** indicates that a module path object has an *ifnone* condition. This Verilog language keyword is used to indicate a default delay for modules with conditional path delays.
- **accScalar** indicates that a net or reg object is scalar (1-bit wide)
- **accVector** indicates that a net or reg object is 2 or more bits wide
- **accExpandedVector** indicates a vector net for which the simulator allows access to individual bits within the vector. Access to all bits of a vector is the default in the Verilog language, and can also be explicitly declared in the Verilog source code, using the *scalared* keyword. For example:


```
wire scalared [63:0] data_bus;
```

- **accUnExpandedVector** indicates a vector net for which the simulator prohibits access to individual bits within the vector. This is a feature which some simulators use to optimize run-time performance. If a simulator has flagged a vector as unexpanded, then some ACC routines will only be able to obtain a handle for the complete vector. Vector only access must be explicitly declared in the Verilog source code, using the *vectored* keyword. For example:

```
wire vectored [31:0] address_bus;
```

 **NOTE** A net will always test true for at least one of the **accExpandedVector** and **accUnexpandedVector** properties. These properties are not mutually exclusive, and a simulator can return true for both properties. When both properties are true, the rules for expanded vectors will take precedence.

- **accCollapsedNet** indicates that a net object has been removed from the simulation by collapsing the net onto an equivalent net. The equivalent net is referred to as the simulated net. Net collapsing may or may not occur, depending on the optimizations performed by a simulator. An example of where net collapsing might occur is when two modules are connected together using an intermediate net.



In this example, the nets **y**, **n3** and **in1** are all tied together, effectively becoming the same net. A Verilog simulator might maintain all three nets in its simulation data structure, or it might collapse two of the nets into the third net in order to optimize the simulation performance. When net collapsing occurs, handles for the collapsed nets can still be obtained (using `acc_next_net()`, for example). However, the collapsed nets may not reflect the actual logic value of the design if the simulator is only propagating value changes on the equivalent simulated net. If a net has been collapsed onto another net, the routine `acc_handle_simulated_net()` will retrieve the handle for the equivalent net.

The following example uses the `acc_object_of_type()` routine to determine if a net is a scalar net or a vector net, so it can process scalar nets differently than vector nets:

```

handle net_handle;

/* add code to get a net handle */

if (acc_object_of_type(net_handle, accScalar))
    /* process scalar nets */
else if (acc_object_of_type(net_handle, accVector))
    /* process vector nets */

```

16.3 Accessing an object's source code location

The ACC library can access the Verilog HDL source code location for an object.

void acc_fetch_location(location, object)

p_location location	pointer to an application-allocated s_location structure to receive the location.
handle object	handle for an object

`acc_fetch_location()` is passed an object handle and a pointer to an **s_location** structure. The Verilog HDL source code location for that object is then

retrieved into the structure fields by the simulator. The `s_location` structure is defined in `acc_user.h`, as follows:

```
typedef struct t_location
{
    int line_no;
    char *filename;
} s_location, *p_location;
```

The `acc_fetch_location()` routine retrieves the file name into the PLI string buffer, and places a pointer to the string in the `s_location` structure. The string is stored in the temporary ACC string buffer, and should be used immediately, or copied into application-allocated storage.

Not all Verilog objects are supported with `acc_fetch_location()`. Table 16-1 lists the types of objects for which the source code location can be accessed, and what location will be returned.

Object Type	Source Location Returned
Modules	Module instantiation line
Module ports	Module definition
Module paths	Module path line
Data paths	Module path line
Primitives	Instantiation line
Explicit nets	Definition line
Implicit nets	Line where first used
Registers	Definition line
Integer, time and real variables	Definition line
Named events	Definition line
Parameters	Definition line
Specparams	Definition line
Named blocks	Definition line
Verilog HDL tasks	Definition line
Verilog HDL functions	Definition line

Table 16-1: Objects supported with `acc_fetch_location()`

The following code fragment illustrates using `acc_fetch_location()` to print the file name and location of an object.

```
handle obj_h;
s_location source_location;
/* get handle for some object */
acc_fetch_location(&source_location, obj_h);
io_printf ("%s is defined in file %s at line %d.\n",
           acc_fetch_fullname(obj_h),
           source_location->filename,
           source_location->line_no);
}
```

16.4 Reading the simulation invocation commands

The ACC library provides a means for PLI application developers to create user-defined invocation options. This capability makes it possible to configure PLI applications or to pass data to an application from the invocation command line of a Verilog simulator. Two ACC routines are used to read the simulation invocation commands:

int acc_fetch_argc()

The `acc_fetch_argc()` routine returns the number of command line arguments given on the command line used to invoke a Verilog simulator.

char **acc_fetch_argv()

The `acc_fetch_argv()` routine returns a pointer to an array of character string pointers containing the command line arguments used to invoke a Verilog simulator.

The **argc** and **argv** values are the same command line values defined in the C language. **argc** is the number of invocation command arguments, and **argv** is a pointer to an array of strings, where each string is one argument from the command line.

The ability to check the simulator command line options makes it possible to create user-defined invocation options. Two applications for user-defined options are:

- To enable debug messages when debugging a PLI application. A verbose mode invocation option could be used to specify that debug messages should be printed.
- To pass file names or other values from the command line to a PLI application.

Parsing the -f command file invocation option

The IEEE 1364-1995 Verilog standard does not specify any invocation options for Verilog simulators. Every Verilog simulator, however, has adopted a small number of de facto standard invocation options. One of these is the **-f** option. This invocation option specifies that the file name which follows the option contains additional command line invocation arguments.

When the `argv` value is **-f**, the next `argv` will be a pointer to a NULL terminated array of pointers to strings. Element 0 in the array will contain the name of the file specified with the **-f** option, and the remaining elements in the array will be invocation commands contained in the file. Comments are not included.

For example, assume that a command file named **run.f** contained the following:

```
my_chip.v  
my_test.v  
+my_debug
```

If simulation were invoked with the command:

```
verilog -f run.f -s
```

Then `argv` would point to the following array of strings:

```
argv[0] -> "verilog"  
[1] -> "-f"  
[2] -----> [0] -> "run.f"  
              [1] -> "my_chip.v"  
              [2] -> "my_test"  
              [3] -> "+my_debug"  
              [5] -> NULL  
[3] -> "-s"
```

Example 16-1 illustrates a PLI application called `$print_invoke_commands`. This application uses `acc_fetch_argc()` and `acc_fetch_argv()` routines to print all command line arguments, including commands from within **-f** command files.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.16/Invoke_Commands_acc.c
- Verilog test bench: Chapter.16/Invoke_Commands_Test.v
- Verilog-XL results log: Chapter.16/Invoke_Commands_Test.log

Example 16-1: \$print_Invoke_Commands — printing invocation commands

```
#include <stdio.h>           /* ANSI C standard I/O library */
#include "veriuser.h"         /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"         /* IEEE 1364 PLI ACC routine library */
/*********************************************************************
 * calltf application
 *****/
/* prototypes of subroutines used by calltf routine */
void PLIbook_ScanCommandFile();

int PLIbook_InvokeCommands_calltf()
{
    int      argc, i;
    char **argv;

    acc_initialize();
    acc_configure(accDisplayWarnings, "true");

    argc = acc_fetch_argc();
    argv = acc_fetch_argv();

    io_printf("\nSimulation invocation commands:\n");
    for (i=0; i<argc; i++) {
        io_printf(" %s\n", *argv);
        if (strcmp(*argv, "-f") == 0) {
            argv++; /* next arg is address to array of strings */
            i++;
            PLIbook_ScanCommandFile((char **)argv);
        }
        argv++; /* increment to next argument */
    }
    io_printf("\n\n");
    acc_close();
    return(0);
}

int PLIbook_indent = 0; /* global variable to format text indenting */

void PLIbook_ScanCommandFile(char **arg)
{
    int i;
    PLIbook_indent += 4; /* increase text indentation */
    while ( *arg != NULL ) { /* loop until null termination */
```

```
for (i=0; i<=PLIbook_indent; i++)
    io_printf(" ");
io_printf("%s\n", *arg);
if (strcmp(*arg, "-f") == 0) {
    arg++; /* next arg is address to array of strings */
    PLIbook_ScanCommandFile((char **) *arg);
}
arg++;
}
PLIbook_indent -= 4; /* decrease text indentation */
return;
}
```

16.5 Accessing objects in simulation which have logic values

In order to read an object's logic value or write a new value into an object, a handle for the object must first be obtained. The ACC routines can read the values of several different types of objects, but can only modify the values of certain object types. The Verilog objects for which the ACC routines can access logic values are:

- A **parameter** constant (read only)
- A **specparam** constant (read only)
- A **specparam attribute** constant (read only)
- Any **net** data type: scalar, vector, part-selects and bit-selects of vectors
- The **reg** data type: scalar, vector, part-selects and bit-selects of vectors
- The **integer**, **time** and **real** variable data types
- A **memory** word select
- A literal **integer** value (read only)
- A literal **real** value (read only)
- A **string** value (read only)
- A **function call** (read only)
- A **system function call**
- A **sequential user-defined primitive**

The routines to obtain handles for the various types of Verilog objects were presented in the previous chapter.

16.6 Reading the values of system task/function arguments


Many of the objects which have logic values can be used as arguments to a system task or system function. The ACC routines offer two methods to access the values of these arguments:

- Read the value directly, using specific ACC routines.
- Obtain a handle for the argument, and read or modify the value of the object, using the appropriate ACC routine for the type of object.

This section shows how to use the specific ACC routines which can read the value of a task/function argument directly, and other sections in this chapter show how to use the ACC routines which read the values of different types of objects, using a handle for the object.

16.6.1 System task/function arguments

A user-defined system task or system function can have any number of arguments, including none. The arguments are numbered from left to right, starting with argument number 1. In the following example:



- Task/Function argument number 1 is a string, with the value “vectors.pat”.
- Task/Function argument number 2 is a signal, with the name input_bus.

16.6.2 Multiple instances of system tasks and system functions

The Verilog HDL source code can reference the same system task or system function any number of times. For example:

```

always @(posedge clock)
    $read_test_vector("A.dat", data_bus);

always @(negedge clock)
    $read_test_vector("B.dat", data_bus);
  
```

Just as a Verilog module can be used, or “*instantiated*”, many times in a design, every occurrence of a system task/function is a separate and unique *instance*. Each instance of `$read_test_vector` in the above example has different arguments. The PLI recognizes that each instance is unique, and keeps track of each instance. Therefore, at each positive edge of clock, the *calltf routine* associated with one instance of `$read_test_vector` will be invoked, and at the negative edge of clock, the *calltf routine* associated with a different instance of `$read_test_vector` will be executed. It is important to understand that the Verilog simulator will call the same C functions for each instance of the system task/function, but the inputs and data associated with each call will be unique for each instance.

As an example, `$read_test_vector` might have a *calltf routine* named `gvCall()` associated with it. When the instance of `gvCall()` that is invoked at the positive edge of clock reads the value of task/function argument 1, it will see the string “`A.dat`”. When the instance of `gvCall()` that is invoked at the negative edge of clock read task/function argument number 1, it will see the string “`B.dat`”.

Reading the values of system task/function arguments

The ACC library provides three specific routines to directly read the values of system task/function arguments:

`double acc_fetch_tfarg(n)`

`int n` position number of a PLI system task/function arg.

`int acc_fetch_tfarg_int(n)`

`int n` position number of a PLI system task/function arg.

`char *acc_fetch_tfarg_str(n)`

`int n` position number of a PLI system task/function arg.

The `acc_fetch_tfarg()` routine returns the value of a task/function argument as a C double. The `acc_fetch_tfarg_int()` routine returns the value of a task/function argument as a C integer.

The `acc_fetch_tfarg_str()` routine converts the value of the argument to a string, which is stored in the temporary ACC string buffer. The routine returns a pointer to the string. The primary purpose of `acc_fetch_tfarg_str()` is to read strings from the Verilog language into strings in the C language. Strings are stored differently in Verilog than in C, so this routine converts the string from one format to the other. If the value of the task/function argument is not a string, each 8 bits of the value are converted into an ASCII character, and the value is retrieved as a string.

The `$read_test_vector` system task can be used to illustrate how these ACC routines are used. An example of this system task is:

```
$read_test_vector("vectors.pat", input_bus);
```

The following C code fragment reads the values of the two system task/function arguments. The first argument is read as a string, and the second as an integer.

```
char *file_name;
int bus_value;

file_name = acc_fetch_tfarg_str(1);
bus_value = acc_fetch_tfarg_int(2);
```

16.6.3 Testing for errors when reading system task/function arguments

An error can occur when reading the value of a system task/function argument if the index number specified is out of range, or if the argument does not have a logic value. The `acc_fetch_tfarg()` routine will return an exception value of `0.0` if an error occurs. The `acc_fetch_tfarg_int()` routine will return an exception value of `0`. The `acc_fetch_tfarg_str()` routine will return an exception value of `null`.

Since the exception values of `0.0` and `0` could be legitimate values, it may not be possible to determine that an error occurred, based on the return value. The `acc_error_flag` will be set if an error occurs, and this flag can be used to test for errors. The `acc_error_flag` was presented in Chapter 15, section 15.2.2 on page 476.

16.6.4 Instance specific system task/function routines

A PLI `calltf routine`, `miscif routine`, `checktf routine` or `sizetf routine` is directly associated with the name of a system task or system function. This association is part of the PLI interface mechanism, which was presented in Chapter 9. Each of these routines can access the arguments of the system task/function which caused the routine to be called, using `acc_handle_tfarg()`, `acc_handle_tfarg_int()` or `acc_handle_tfarg_str()`. If these routines call another C function, that function can also access the arguments of the system task/function.

The ACC library also provides for another type of PLI routine, called a *consumer routine*. Defining and using *consumer routines* is defined in the next chapter of this book. A *consumer routine* is *not* associated with the name of a system task or system function. Therefore, a *consumer routine* cannot directly access the arguments of system tasks/functions. The PLI also provides a means for PLI applications to indirectly

access system task/function arguments. Indirect access is done by obtaining a handle for an instance of a system task/function, and then using a specific set of ACC routines to access the arguments of that system task/function instance.

An instance of a system task/function is an object, and the ACC routines can obtain a handle for that object using, the **acc_handle_tfinst()** routine. The syntax of this routine is:

handle acc_handle_tfinst()

The **acc_handle_tfinst()** routine returns a handle for the system task/function which called the PLI application.

There are instance-specific counterparts to the routines which read the values of task/function arguments, or which obtain a handle for an argument:

double acc_fetch_itfarg(n, tfinst)

int n position number of a PLI system task/function arg.
handle tfinst handle for an instance of a PLI system task/function

int acc_fetch_itfarg_int(n, tfinst)

int n position number of a PLI system task/function arg.
handle tfinst handle for an instance of a PLI system task/function

char *acc_fetch_itfarg_str(n, tfinst)

int n position number of a PLI system task/function arg.
handle tfinst handle for an instance of a PLI system task/function

handle acc_handle_itfarg(n, tfinst)

int n position number of a PLI system task/function argument.
handle tfinst handle for an instance of a PLI system task/function.

Each of these routines require two inputs, the index number of a task/function argument, and a handle for a task/function instance.

The typical usage of these routines is for a *calltf routine* or a *misctf routine* to obtain the handle for the system task/function instance which called the routine. This handle is then saved in the *user_data* field of a *consumer routine*. This gives the consumer routine access to the instance handle, so that the consumer routine can access the arguments of the system task/function. Example 18-6 on page 622 of chapter 18 shows examples of using these instance specific ACC routines.

NOTE

The TF system task/function instance pointer is not the same as the ACC system task/function instance pointer. The TF library also has instance specific versions of the routines which access the arguments of a system task or system function. However, the instance pointer returned from `tf_getinstance()` is a different data type than the ACC handle returned from `acc_handle_tfinst()`. The instance pointer from a routine in one library should not be used with routines in the other library.

16.6.5 Modifying the values of system task/function arguments

The values of system task/function arguments can be modified by obtaining a handle for the argument. The `acc_handle_tfarg()` routine is used to obtain the argument handle. Note that not all legal arguments to a system task/function have logic values which can be modified. For example, a literal string or a module instance name are valid arguments, but have fixed values which cannot be modified by the PLI. The routines `acc_fetch_type()` and `acc_fetch_fulltype()` can be used to determine the type of object which is being used as a task/function argument. The subsequent sections of this chapter present how to read and modify values of different object types.

16.7 Reading object logic values

Using the handle for an object, a PLI application can read the logic value property of the object. The same ACC routine is used to read the logic value of any type of object and any data type of logic value.

16.7.1 Working with a 4-logic value, multiple strength level system

The Verilog HDL supports 4 logic values, **0**, **1**, **z** and **x** and multiple levels of signal strength. There are also two ambiguous logic values, represented by **L** (low) and **H** (high), and many ambiguous strength values. The C programming language does not directly represent the same information. The ACC routine to read logic values provides several ways to automatically translate values between Verilog and C. The translation converts Verilog 4-state logic into the following C types:

- **A C integer:** Verilog scalar and vector logic values are converted to a single C integer. The Verilog 4-state logic is converted to 2-state logic values of 0 and 1. Logic values of z and x values are converted to 0, and strength levels are ignored.

- **A C double:** Verilog real number values, scalar values and vector values are converted to a C double precision value. Real numbers in Verilog are 2-state decimal values, which convert directly to C doubles. Scalar and vector 4-state logic is converted to 2-state logic, and strength levels are ignored.
- **A C string:** Verilog scalar and vector logic values are converted to a C character string. The Verilog 4-state logic is converted to the letters “0”, “1”, “z” and “x”. Logic strength levels for scalar nets are represented using the 3 character mnemonics defined in the Verilog language. Strength levels are ignored for Verilog vectors and other data types.
- **A C constant:** Verilog scalar values are converted to a C integer constants. The Verilog 4-state logic is converted to the constants **acc0**, **acc1**, **accZ** and **accX**. Logic strength levels are ignored.
- **A C aval/bval structure:** Verilog scalar and vector logic values are converted to a C structure which encodes each bit of a Verilog 4-state value to a pair of bits in C, referred to as an *aval/bval* pair. An array of *aval/bval* pairs is used to encode Verilog vectors of any size. The logic strength levels are ignored.

16.7.2 The acc_fetch_value() routine

A single ACC routine is used to read the logic value of most Verilog objects.

char *acc_fetch_value(object, format_str, value)

handle object	handle for a net or register.
char *format_str	character string controlling the radix of the retrieved value; must be “%b”, “%o”, “%d”, “%h”, “%v” or “%%”.
p_acc_value value	pointer to an application-allocated s_acc_value structure to receive the value as an <i>aval/bval</i> pair. Only used if format_str is “%%”.

The **acc_fetch_value()** routine converts Verilog logic values into C language representations. The logic value that is retrieved can be passed to the PLI application in one of two ways:

- The retrieved value can be saved in the temporary ACC string buffer, and a pointer to the string is returned by **acc_fetch_value()**.
- The retrieved value can be stored in an **s_acc_value** structure allocated by the PLI application. The structure can receive the Verilog logic value in a variety of C data types.

Retrieving values as character strings

Retrieving values as a C string is the easiest way to represent any size of Verilog vector and any Verilog logic value. This format is also an easy way to retrieve values using `acc_fetch_value()`. All that is needed is to obtain an object handle and call `acc_fetch_value()`, saving the string pointer that is returned in a `char *` variable, if desired.



TIP *The format in which a value is read can have an impact on the run-time performance of a PLI application.* The fastest run-time performance will be achieved when a value is retrieved in a format closest to the format in which a value is saved in the simulation structure. For Verilog scalar and vector nets and regs, this format is the `aval/bval` pair. For Verilog integers, C integers are most efficient, and for Verilog reals, C doubles are most efficient. The least efficient method for run-time performance is to retrieve a logic value as a C string.

The `acc_fetch_value()` routine has three inputs, but only the first two are used when retrieving values as a string. The third input should be set to `null`. It is not used when retrieving values as C strings.

The first input to `acc_fetch_value()` is a handle for the object from which the logic value is to be read. The second input to a format string, which controls how the Verilog value will be represented in the C string. Table 16-2 shows the formats which are supported for returning string values:

Format String	Return Value Description
<code>"%b"</code>	value is retrieved as a C string with a binary representation ('0', '1', 'z', 'x')
<code>"%o"</code>	value is retrieved as a C string with an octal representation ('0' through '7', 'z', 'Z', 'x', 'X')
<code>"%d"</code>	value is retrieved as a C string with a decimal representation ('0' through '9', 'z', 'Z', 'x', 'X')
<code>"%h"</code>	value is retrieved as a C string with a hexadecimal representation ('0' through 'F', 'z', 'Z', 'x', 'X')
<code>"%v"</code>	value is retrieved as a C string with a 3-character strength representation, such as <code>st0, we1</code> or <code>65X</code>
<code>"%%"</code>	value is retrieved into an <code>s_acc_value</code> structure, instead of a C string (see section 16.7.7 on page 548)

Table 16-2: `acc_set_value()` format strings

The formats which return the logic value as a string ("%b", "%o", "%d" and "%h") use the same representation as the Verilog language built-in \$display system task.

The following example illustrates using acc_set_value() to return the logic values of all nets in a module, represented in a binary notation.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.16/list_nets_acc.c
- Verilog test bench: Chapter.16/list_nets_test.v
- Verilog-XL results log: Chapter.16/list_nets_test.log

Example 16-2: \$list_nets — using acc_fetch_value() to read values as strings

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/*********************************************************************
 * calltf application
 *****/
int PLIbook_ListNets_calltf()
{
    handle module_h, net_h;
    acc_initialize();
    module_h = acc_handle_tfarg(1);
    io_printf("\nNet values in module %s:\n",
              acc_fetch_fullname(module_h));
    net_h = null;      /* start with known value for target handle */
    while (net_h = acc_next_net(module_h, net_h)) {
        io_printf(" %-10s: %s\n",
                  acc_fetch_name(net_h),
                  acc_fetch_value(net_h, "%b", null));
    }
    acc_close();
    return(0);
}
```



NOTE The value retrieved from acc_fetch_value() for any of the C string representations will be stored in the ACC string buffer. A pointer to the string is returned by acc_fetch_value(). The ACC string buffer storage is temporary, so the value should be used immediately, or copied into application-allocated memory.

Retrieving values into an s_acc_value structure

`acc_fetch_value()` can convert Verilog logic values into several C data types. This is done by setting the format string to “%%”. The value is then retrieved into an **s_acc_value** structure. The PLI application must first allocate this structure, and pass a pointer to the structure as the third input to `acc_fetch_value()`. The return value from `acc_fetch_value()` is not used, and should be ignored.

The **s_acc_value** structure is defined in `acc_user.h`, and is listed below.

```
typedef struct t_setval_value
{
    int format;
    union
    {
        char *str;
        int scalar;
        int integer;
        double real;
        p_acc_vecval vector;
    } value;
} s_setval_value, *p_setval_value, s_acc_value, *p_acc_value;
```

This **s_acc_value** structure has two fields, **format** and **value**. The **format** field controls what C language data type should be used to receive the value. The **value** field is a union of C data types which receive the value of the object. The definitions of the **format** values and the **value** union fields are listed in the following table.

Format	Definition
accBinStrVal	retrieves the Verilog logic value as a C string, using binary numbers, with 4-state logic
accOctStrVal	retrieves the Verilog logic value as a C string, using octal numbers, with 4-state logic
accDecStrVal	retrieves the Verilog logic value as a C string, using decimal numbers, with 4-state logic
accHexStrVal	retrieves the Verilog logic value as a C string, using hexadecimal numbers, with 4-state logic

Table 16-3: The **s_acc_value** structure format constants

Format	Definition
accScalarVal	retrieves the Verilog scalar logic value as one of the C constants acc0 , acc1 , accZ or accX
accIntVal	retrieves the Verilog logic value as a C integer, with 2-state logic
accRealVal	retrieves the Verilog logic value as a C double precision number, with 2-state logic
accStringVal	retrieves the Verilog string value as a C string
accVectorVal	retrieves the Verilog vector value as an array of C integers, encoded to represent Verilog 4-state logic

Table 16-3: The **s_acc_value** structure format constants (continued)

The **value** field of the **s_acc_value** structure is a union of C data types. Which C data type is used to receive the Verilog value is controlled by the **format** field of the structure:

- The **value.str** field in the **value** union is used if the format is **accBinStrVal**, **accOctStrVal**, **accDecStrVal**, **accHexStrVal**, or **accStringVal**. The **acc_fetch_value()** routine will fetch the object's logic value, and convert the value into an ASCII string. The routine will place a pointer to the string in the **value.str** field. The string is stored in the temporary ACC string buffer, and may be overwritten by other ACC routines which return pointers to strings.
- The **value.scalar** field is used if the format is **accScalarVal**. The **acc_fetch_value()** routine will fetch the object's logic value, and place one of the constants: **acc0**, **acc1**, **accZ** or **accX** in the **value.scalar** field.
- The **value.integer** field is used if the format is **accIntVal**. The **acc_fetch_value()** routine will fetch the object's logic value, and convert the value into 2-state logic. The routine will place the value in the **value.integer** field.
- The **value.real** field is used if the format is **accRealVal**. The **acc_fetch_value()** routine will fetch the object's logic value, convert the value into 2-state logic, and place the value in the **value.real** field.
- The **value.vector** field is used if the format is **accVectorVal**. The **acc_fetch_value()** routine will fetch the object's logic value into an array of **s_acc_vecval** structures, and place a pointer to the memory in the **value.vector** field of the **s_acc_value** structure. Refer to section 16.7.7 on page 548, for more details on reading vector values.

Before `acc_fetch_value()` can be called using the “`%%`” format, the PLI application must first allocate an `s_acc_value` structure. Following are three ways in which the structure might be allocated:

- Allocate an automatic variable of the `s_acc_value` type. The storage allocated will automatically be freed when the PLI application exits:

```
s_acc_value obj_value;  
acc_fetch_value(obj_handle, "%%", &obj_value);
```

- Allocate persistent storage, which can be maintained from one call of the PLI application to another. The pointer to the storage can be preserved in the TF work area.

```
p_acc_value obj_value;  
obj_value = (p_acc_value)malloc(sizeof(s_acc_value));  
acc_fetch_value(obj_handle, "%%", obj_value);
```

- Allocate static storage, which can be initialized at the time of allocation. The initialization is not repeated for each call to the PLI application, which can improve run-time efficiency.

```
static s_acc_value obj_value = {accBinStringVal};
```

NOTE If static storage is used, the initial values of the structure should not be changed within the PLI application code. All instances of the PLI application will share the same static structure, so if one call to the application changes the structure values, it will affect instances of the PLI application.

16.7.3 Reading 4-state logic as C strings, using `acc_fetch_value()`

The `accBinStrVal`, `accOctStrVal`, `accDecStrVal` and `accHexStrVal` formats of the `s_acc_value` structure are nearly identical to the respective “`%b`”, “`%o`”, “`%d`”, and “`%h`” string formats that can be specified as an argument to `acc_fetch_value()`. The only difference is that, when using a “`%%`” format, the pointer to the string will be placed in the `value.str` field of the `s_acc_value` structure, instead of being returned by the `acc_fetch_value()` function. The string will be stored in the temporary ACC string buffer.

16.7.4 Reading 2-state logic as a C integer, using `acc_fetch_value()`

The `accIntVal` format will retrieve Verilog logic values into a C integer. The `acc_fetch_value()` routine will convert Verilog 4-state logic into C 2-state logic. The Verilog value which is read can be any Verilog data type.

NOTE

The maximum value which can be read is constrained by the size of a C integer. An integer value in Verilog can be a vector of any bit width. If the value stored in a Verilog vector is greater than the maximum value which can be stored as a C integer, then the left-most bits (the most significant bits) of the Verilog vector are truncated.

The steps to read the value of an object as a C integer are:

1. Allocate an `s_acc_value` structure.
2. Set the `format` field in the structure to `accIntVal`.
3. Call `acc_fetch_value()`, giving a pointer to the `s_acc_value` structure as an input, along with a handle for the object from which to read the logic value.
4. Read the logic value of the object from the `value.integer` field of the `s_acc_value` structure.

The following C code fragment illustrates reading the logic value of a Verilog integer into a C integer using `acc_fetch_value()`.

```
handle tfarg_h;
s_acc_value val_s;

tfarg_h = acc_handle_tfarg(1);

val_s.format = accIntVal;

acc_fetch_value(tfarg_h, "%d", &val_s);

io_printf("%s = %d\n",
          acc_fetch_name(tfarg_h),
          val_s.value.integer);
```

16.7.5 Reading 2-state logic as a C double, using `acc_fetch_value()`

The `accRealVal` format will retrieve Verilog real values into a C double. The `acc_fetch_value()` routine will convert any other Verilog value into a 2-state integer value, and then to a C double.

The steps to read the value of an object as a C double are:

1. Allocate an `s_acc_value` structure.
2. Set the `format` field in the structure to `accRealVal`.
3. Call `acc_fetch_value()`, giving a pointer to the `s_acc_value` structure as an input, along with a handle for the object from which to read the logic value.

4. Read the logic value of the object from the **value.real** field of the **s_acc_value** structure.

16.7.6 Reading Verilog string values into C strings

The **acc_fetch_value()** routine will automatically convert a Verilog string value to a C string, and place a pointer to the string in the **value.str** field of the **s_acc_value** structure. The intended use of the **accStringVal** format is to read Verilog string values. If the value of the object is not a Verilog string, then each 8 bits of the Verilog value will be converted to an ASCII character.

The steps to read the value of an object as a C double are:

1. Allocate an **s_acc_value** structure.
2. Set the **format** field in the structure to **accStringVal**.
3. Call **acc_fetch_value()**, giving a pointer to the **s_acc_value** structure as an input, along with a handle for the object from which to read the logic value.
4. Read the string from the **value.str** field of the **s_acc_value** structure.

16.7.7 Reading Verilog 4-state logic vectors as encoded aval/bval pairs

The **acc_fetch_value()** routine with an **accVectorVal** format will retrieve an object's 4-state logic value as an encoded pair of C integers. The encoding uses an **aval/bval** pair of C integers to represent the 4-state logic values of Verilog. One bit of each **aval/bval** pair represents a corresponding bit of the Verilog logic value. The encoding is shown in Table 16-4. The same encoding is used by the TF library and the VPI library:


aval/bval pair	Verilog logic value represented
0/0	0
1/0	1
0/1	Z
1/1	X

Table 16-4: **aval/bval** logic value encoding

The ACC library assumes C integers are 32-bits wide, and therefore uses the **aval/bval** pair to encode up to 32 bits of a Verilog vector. By using an array of **aval/bval**

integer pairs, vector lengths of any size may be represented. The representation of a 40-bit vector in Verilog can be visualized as:

For the Verilog declaration: `reg [39:0] data;`



lsb = least significant bit

msb = most significant bit


The Verilog language supports any numbering convention for a vector's bit numbers. The least significant bit of the Verilog vector can be the smallest bit number, such as bit 0 (which is referred to as little endian convention). Or, the least significant bit of the Verilog vector can be the largest bit number, such as bit 39 (which is referred to as big endian convention). Verilog does not require that there be a bit zero at all. Each of the following examples are valid vector declarations in Verilog:

```
reg [39:0] data;      /* little endian -- LSB is bit 0 */
reg [0:39] data2;    /* big endian   -- LSB is bit 39 */
reg [40:1] data3;    /* little endian -- LSB is bit 1 */
```

The bit numbering used in Verilog does not affect the `aval/bval` representation of the Verilog vector. In the array of `aval/bval` pairs, the LSB of the Verilog vector will always be the LSB of the first C integer in the array, and the MSB of the Verilog vec-

tor will always be the last bit in the array which is used. The following diagram illustrates the aval/bval array for a Verilog vector declared with a big endian convention.

For the Verilog declaration: `reg [1:40] data;`



The aval/bval pair is declared in an `s_acc_vecval` structure, which is defined in the `acc_user.h` file. The structure definition is:

```
typedef struct t_acc_vecval
{
    int aval;
    int bval;
} s_acc_vecval, *p_acc_vecval;
```

To read a task/function argument's 4-state logic value using `acc_fetch_value()` involves four basic steps:

1. Allocate an `s_acc_value` structure.
2. Allocate memory for an array of `s_acc_vecval` structures. The array must have 1 element for each 32 bits of the Verilog vector.
3. Set the `format` field in the structure to `accVectorVal`.
4. Set the `value.vector` field in the structure to the pointer to the `s_acc_vecval` array.
5. Call `acc_fetch_value()`, passing a pointer to the `s_acc_value` structure as an input, along with a handle for the object from which to read the logic value.
6. Read the logic value of the object from the `value.vector` field of the `s_acc_value` structure.

To allocate memory for the `s_acc_vecval` array requires first determining the number of elements that will be needed in the array. Since there will be one element for each 32 bits of a Verilog vector, the number of array elements can be calculated using:

```
number_of_array_elements = ((vector_size - 1) / 32 + 1);
```

The `vector_size` of the Verilog vector object can be retrieved using the `acc_fetch_size()` routine. For example:

```
vector_size = acc_fetch_size(vector_handle);
```

Once the number of elements are known, the value of each 32-bit group of the Verilog vector can be accessed by reading the `aval/bval` pair of each `s_acc_vecval` structure in the array. Within each `aval/bval` pair, an individual bit of the Verilog vector can be accessed by masking out the other bits in the `aval/bval` pair.

NOTE When reading values using the `accVectorVal` format, the memory to store the array of `s_acc_vecval` structures must be allocated and maintained by the PLI application. If the PLI application needs to preserve the array for future calls to the application, a pointer to the array must be saved. One way to preserve the pointer is to allocate persistent memory for the `s_acc_value` structure, and save a pointer to this structure in the TF work area. The pointer to the `s_acc_vecval` array can then be saved in the `value.vector` field of the `s_acc_value` structure.

Example 16-3 illustrates using `acc_fetch_value()` to read the `aval/bval` encoded 4-state value of a vector. This example assumes the vector is passed to the PLI application as the first system task/function argument. The value of the vector is then printed one bit at a time. For simplicity in this example, it is assumed that the least-significant bit of the Verilog vector is bit 0.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.16/read_vecval_acc.c
- Verilog test bench: Chapter.16/read_vecval_test.v
- Verilog-XL results log: Chapter.16/read_vecval_test.log

Example 16-3: *\$read_vector_value* — reading vector values as aval/bval pairs

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/********************* calltf application *********************/
int PLIbook_ReadVecVal_calltf()
{
    handle      vector_h;
    s_acc_value vector_val;      /* structure to receive vector value */
    int         i, vector_size, array_size, avalbit, bvalbit, bit_num;
    char        vlogval;

    vector_h = acc_handle_tfarg(1);

    vector_size = acc_fetch_size(vector_h); /* determine number of...*/
    array_size = ((vector_size-1) / 32 + 1); /* ...elements in array */

    vector_val.value.vector = (p_acc_vecval)malloc(array_size*
sizeof(p_acc_vecval));

    vector_val.format = accVectorVal;           /* set value format field */

    acc_fetch_value(vector_h,"%",&vector_val); /* read vector's value */

    io_printf("\nVector %s encoded value:\n",
              acc_fetch_name(vector_h));
    for (i=0; i<array_size; i++) {
        /* the following loop assumes the Verilog LSB is bit 0 */
        for (bit_num=0; bit_num<=31; bit_num++) {
            avalbit=PLIbook_getbit(vector_val.value.vector[i].aval, bit_num);
            bvalbit=PLIbook_getbit(vector_val.value.vector[i].bval, bit_num);
            vlogval=PLIbook_get_4state_val(avalbit, bvalbit);
            io_printf(" bit[%2d]  aval/bval = %d/%d  4-state value = %c\n",
                      (i*32+bit_num), avalbit, bvalbit, vlogval);
            /* quit when reach last bit of Verilog vector */
            if ((i*32+bit_num) == vector_size-1) break;
        }
    }
    return(0);
}
```

```

*****
 * Function to determine if a specific bit is set in a 32-bit word.
 * Sets the least-significant bit of a mask value to 1 and shifts the
 * mask left to the desired bit number.
 *****
int PLIbook_getbit(int word, int bit_num)
{
    int mask;
    mask = 0x00000001 << bit_num;
    return((word & mask)? TRUE: FALSE);
}

*****
 * Function to convert aval/bval encoding to 4-state logic represented
 * as a C character.
 *****
char PLIbook_get_4state_val(int aval, int bval)
{
    if      (!bval && !aval) return('0');
    else if (!bval &&  aval) return('1');
    else if ( bval && !aval) return('z');
    else                           return('x');
}

```

16.8 Writing values into Verilog objects

A single ACC routine is used to write values onto any type of object which the ACC supports modifying values.

int acc_set_value(object, value, delay)

<i>handle object</i>	handle for a net, register or sequential UDP.
<i>p_setval_value value</i>	pointer to an application-allocated <i>s_setval_value</i> structure containing the value to be set.
<i>p_setval_delay delay</i>	pointer to an application-allocated <i>s_setval_delay</i> structure containing a propagation delay value.

The *acc_set_value()* routine converts a value represented as a C data type into Verilog 4-state logic, and writes the value into a Verilog object. The value to be written can be represented a variety of ways in the C language. These representations are the same as was described in section 16.7.1 on page 540 for reading values.

There are three inputs to `acc_set_value()`:

1. A handle for the object into which the value is to be written.
2. A pointer to an `s_setval_value` structure. This structure must be allocated by the PLI application, and the appropriate field within the structure set to the value to be written into the object.
3. A pointer to an `s_setval_delay` structure. This structure is allocated by the PLI application, and set to a propagation delay value.

The `s_setval_delay` structure is used to specify when the simulator should apply the value which is being written by the PLI application. The structure is defined in `acc_user.h`, as follows:

```
typedef struct t_setval_delay
{
    s_acc_time time;
    int model;
} s_setval_delay, *p_setval_delay;
```

The `model` field in the `s_setval_delay` structure controls how the value to be written will propagate within the simulation. The `acc_set_value()` routine can either write the value into the object immediately, or use the simulator's event scheduling mechanism. By using the event scheduler, a value can be scheduled to occur in a future simulation time step.

The `model` field is set to one of the following constants:

- `accNoDelay` indicates no propagation delay is to be used. The object may be a Verilog reg, variable, memory word, sequential UDP or system function. When this flag is used, the time field in the `s_setval_delay` structure is not used, and can be set to `null`.
- `accInertialDelay` indicates that inertial delay propagation is to be used. Any pending events which are scheduled for the object are cancelled. The object may be a Verilog reg, variable, or memory word.
- `accPureTransportDelay` indicates transport delay propagation is to be used. Any pending events which are scheduled for the object remain scheduled (no events are cancelled). The object may be a Verilog reg, variable, memory word or variable array word.
- `accTransportDelay` indicates a modified transport delay propagation is to be used. Any pending events for the object which are scheduled at a later time than this new event are cancelled. The object must be a Verilog reg, variable, memory word or variable array word.

- **accAssignFlag** indicates that the value is to be continuously assigned into the object, overriding any existing values. The value is written into the object in the same manner as the Verilog HDL procedural `assign` statement. No propagation delay is used, and the `time` field of the `s_setval_delay` structure can be set to `null`. The object may be a Verilog `reg` or variable. Only one procedural continuous assign value may exist for an object at a time. Setting another assign value on an object will replace any existing procedural continuous assign value, regardless of whether the assign was set within the PLI or within the Verilog HDL.
- **accDeassignFlag** indicates that any existing procedural continuous assign on the object is to be deassigned. This is the same functionality as the Verilog HDL procedural `deassign` statement.
- **accForceFlag** indicates that the value is to be forced into the object, overriding any existing values. The value is written into the object in the same manner as the Verilog HDL procedural `force` statement. No propagation delay is used, and the `time` field of the `s_setval_delay` structure may be set to `null`. The object may be a Verilog `reg`, variable, memory word, variable array word, or net. Only one force value may exist for an object at a time. Setting a force value on an object will replace any existing force value, regardless of whether the force was set within the PLI or within the Verilog HDL.
- **accReleaseFlag** indicates that any existing force on the object is to be released. This is the same functionality as the Verilog HDL procedural `release` statement.

The `time` field of the `s_setval_delay` structure is a pointer to an `s_acc_time` structure. This structure is used to hold the propagation delay which is to be used by `acc_set_value()`.

The `s_acc_time` structure is defined in `acc_user.h`, and is shown below:

```
typedef struct t_acc_time
{
    int type;
    int low,
        high;
    double real;
} s_acc_time, *p_acc_time;
```


The `type` field in the `s_acc_time` structure controls how the time value will be specified. The type must be set to one of the following constants:

- **accRealTime** indicates that the delay is represented as a C double. The delay value will be scaled to the time units and precision of the module containing the object into which the value will be written.

- **accTime** indicates that the delay is represented as a pair of C integers, which contain the high order 32 bits and the low order 32 bits of the 64-bit simulation time. The delay value will be scaled to the time units and precision of the module containing the object into which the value will be written.
- **accSimTime** indicates that the delay is represented as a pair of C integers, which contain the high order 32 bits and the low order 32 bits of the 64-bit simulation time. The delay value is in the simulator's internal time units, and will *not* be scaled to the object's time scale.

To represent a 64 bit time value, the Verilog PLI uses a pair of C unsigned integers to store the full 64 bits of simulation time. The lower 32 bits of Verilog time are placed in one integer, and the upper 32 bits of the Verilog time are stored in a second integer, as shown in the following illustration:

```
int high, low;
```



An example of using acc_set_value()

To write a value into an object requires the following steps:

1. Obtain a handle for an object—the object must have a value property, in order to write a value into the object.
2. Allocate an **s_acc_value** structure.
3. Allocate an **s_setval_delay** structure.
4. Allocate an **s_acc_time** structure (if a propagation delay is to be specified).
5. Allocate an array of **s_acc_vecval** structures, if needed (for the **accVectorVal** format).
6. Load the value variables or structures with the value to be written.
7. Set the **format** field in the **s_acc_value** structure to indicate how the logic value is represented in C.

8. Set the appropriate field in the **value** union of the **s_acc_value** structure to a pointer to the value.
9. Set the **model** field in the **s_setval_delay** structure to indicate the type of propagation delay which should be used.
10. Set the **time** field of the **s_setval_delay** structure to a pointer to the **s_acc_time** structure that was allocated, or **null** if no delay is to be used.
11. Set the **type** field in the **s_acc_time** structure to indicate how the delay value is represented.
12. Set the delay value in the appropriate field of the **s_acc_time** structure.
13. Call **acc_set_value()**, with pointers to the **s_setval_value** and **s_setval_delay** structures.

The following code fragment writes a value represented as a C character string onto the second argument of a system task, using transport delay. The procedure for writing a value represented in other C data types is very similar to this example. The code fragment example assumes a **\$read_test_vector** system task, where the second argument is a Verilog reg data type. An example of this system task is:

```
reg [22:0] input_vector;  
always @(posedge clock)  
  $read_test_vector("vector_file.pat", input_vector);
```

The C code to write a value into the second system task argument is.

```
handle tfarg_h;  
char test_vector[1024];  
double delay;  
s_acc_value value_s;  
s_setval_delay delay_s;  
/* read test value and delay from file *.  
tfarg_h = acc_handle_tfarg(2);  
value_s.format = accStrBinVal;  
value_s.value.str = test_vector;  
delay_s.model = accTransportDelay;  
delay_s.time.type = accRealTime;  
delay_s.time.real = delay;  
acc_set_value(tfarg_h, &value_s, &delay_s);
```

No delay versus zero delay

The Verilog PLI standard makes a distinction between putting a value into simulation with no delay and putting a value into simulation with zero delay.

- The **accNoDelay** delay flag indicates that the value will be written into Verilog simulation instantly. When the PLI application returns back to simulation, any values written to an object or system function return using these routines will already be in effect for Verilog HDL statements to use.
- The **accInertialDelay**, **accTransportDelay** and **accPureTransportDelay** flags schedule a value to be written into simulation. If a delay of zero is specified, the value is scheduled to be written into the object later in the current simulation time step. When the system task returns back to simulation, the scheduled value will not yet have taken effect. Other Verilog HDL statements scheduled to be executed in the same simulation time step may or may not see the new value of the object (depending on where the value change which was scheduled by the PLI falls in the simulator's event queue, in relation to other Verilog HDL events).

The following simple Verilog HDL source code illustrates the potential problem of putting a value into simulation using a delay of zero.

```
module test;
    reg [7:0] reg1, reg2;
    initial
        begin
            reg1 = 0; reg2 = 0;
            $put_value(reg1, reg2);
            $display("reg1=%d    reg2=%d", reg1, reg2);
            $strobe ("reg1=%d    reg2=%d", reg1, reg2);
            #1 $finish;
        end
    endmodule
```

If the *calltf routine* for **\$put_value** puts a value into **reg1** using **accNoDelay**, then when **\$put_value** returns to the simulation, and the **\$display** statement prints the value of **reg1**, the *new* value will be printed.

If, however, the *calltf routine* for **\$put_value** writes a value into **reg2** using **accInertialDelay** with a delay of zero, then, when **\$put_value** returns to the simulation, the **\$display** statement will print the *old* value of **reg2**. The old value is printed because the value written by the PLI has been scheduled to take place in the current time step, but will not yet have taken effect. The **\$strobe** statement which follows the **\$display** will print the new value of both **reg1** and **reg2**, because the definition of **\$strobe** is to print its message at the end of the current simulation time step, after all value changes for that moment in time have taken effect.

16.9 Returning logic values of system functions

When a handle for a system function is passed to the `acc_set_value()` routine, the value will be written as the system function's return value. The `acc_handle_tfinst()` routine is used to obtain the handle for the system function.

Rules for returning values to system functions

There are two important restrictions on returning values to a system function:

- A value can only be written to the return of a system function from a *calltf routine*, which is when the system function is active. The *calltf routine* is invoked when the system function is encountered by the simulator while simulation is running. System function return values cannot be written from a *miscif routine*, *sizeif routine* or *checktf routine* because the simulation is not executing the statement containing the function at the times these routines are invoked.
- It is illegal to specify a propagation delay when returning a value to a system function. If a delay is specified, the value will not be written. To return a value with zero delay, specify **accNoDelay** for the delay flag.

Types of system functions

The ACC standard allows for two types of system functions:

- Real functions, which will return a Verilog real value. Verilog real variables are double precision floating point values.
- Sized integer functions, which return a Verilog scalar or vector value. Verilog scalars are 1-bit wide, and vectors can be any width. The `acc_set_value()` routine can return a system function value of any vector width.

The type of function is established when the system function is registered through the TF/ACC PLI interface mechanism. Refer to Chapter 9, section 9.11 on page 292, for a full description of registering system functions.

Example 16-4, shown below, implements a system function that returns a 72-bit value. The value is represented using aval/bval pairs to encode 4-state logic in the PLI application. To keep this example simple, the 72-bit value that is returned is hard-coded into the PLI application.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.16/func_72bit_acc.c
- Verilog test bench: Chapter.16/func_72bit_test.v
- Verilog-XL results log: Chapter.16/func_72bit_test.log

Example 16-4: \$func_72bit — returning 4-state vector values with system functions

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"           /* IEEE 1364 PLI ACC routine library */
/*********************************************
 * Sizetf application
 *****/
int PLIBook_Func72bit_sizetf()
{
    return(72);    /* $pow returns 32-bit values */
}

/*********************************************
 * Calltf application
 *****/
int PLIBook_Func72bit_calltf()
{
    handle systf_handle;
    s_setval_value value;
    s_setval_delay delay;
    p_acc_vecval val_array;
    int array_size, i;

    /* declare an array of aval/bval pairs for the vector size */
#define VEC_SIZE 72 /* hard coded 72-bit vector for this example */
    array_size = ((VEC_SIZE-1)/32)+1;
    val_array = (p_acc_vecval)malloc(sizeof(s_acc_vecval) * array_size);

    /* set value of vector aval/bval pairs */
    for (i=0; i<array_size; i++) {
        val_array[i].aval = 0xAAAAAAAA; /* aval bits encode logic 0 & 1 */
        val_array[i].bval = 0x00000000; /* bval bits encode logic Z & X */
    }

    systf_handle = acc_handle_tfinst();
    value.format = accVectorVal;
    value.value.vector = val_array;
    delay.model = accNoDelay;

    acc_set_value(systf_handle, &value, &delay); /* set sysfunc return */
    return(0);
}
```

16.10 Verilog objects which can have delay values

Several types of constructs in the Verilog language can have delay values, and each type of object can store a different number of delay values, representing different delays for different output transitions. In addition, each output transition can be represented by a single delay value, referred to as a typical delay, or each output transition can be represented by a set of three delay values for a minimum, typical and maximum delay range.

Table 16-5 lists the types of objects for which ACC routines can access delay values. The table shows the number of output transitions for each type of object that can be represented in the Verilog HDL source code, and the number of transitions which will be accessed by the PLI for that object.

Verilog Object	Verilog HDL Source Code		Delays Accessed by PLI	
	Number of Delay Transitions ¹ (and order specified)		Number of Transitions ¹ (and order accessed)	
2-state primitive	0	zero delay	2	rise, fall
	1	all transitions have same delay		
	2	separate delays for rise, fall		
3-state primitive	0	zero delay	3	rise, fall, toZ
	1	all transitions have same delay		
	2	separate delays for rise, fall		
	3	separate delays for rise, fall, toZ		
module path	(number of values specified is mapped to 12 transitions within simulation data structure)		(number of transitions accessed is set by accPathDelayCount)	
	1	all transitions have same delay	1	all transitions
	2	separate delays for rise, fall	2	rise, fall
	3	separate delays for rise, fall, toZ	3	rise, fall, toZ
	6	separate delays for 0->1, 1->0, 0->Z, Z->1, 1->Z, Z->0	6	0->1, 1->0, 0->Z, Z->1, 1->Z, Z->0
	12	separate delays for 0->1, 1->0, 0->Z, Z->1, 1->Z, Z->0, 0->X, X->1, 1->X, X->0, X->Z, Z->X	12	0->1, 1->0, 0->Z, Z->1, 1->Z, Z->0, 0->X, X->1, 1->X, X->0, X->Z, Z->X

Table 16-5: Verilog objects which can have delays

Verilog Object	Verilog HDL Source Code		Delays Accessed by PLI	
	Number of Delay Transitions ¹ (and order specified)		Number of Transitions ¹ (and order accessed)	
module port	(cannot be represented in Verilog HDL)		3	rise, fall, toZ
module port bit	(cannot be represented in Verilog HDL)		3	rise, fall, toZ
module interconnect	(cannot be represented in Verilog HDL)		3	rise, fall, toZ
timing constraint	1	timing limit	1	timing limit

¹A transition can be a typical delay value, or a min:typ:max set of delay values

Table 16-5: Verilog objects which can have delays (continued)

Definitions of terms used in table 16-5.

- **2-state primitive:** an instance of a Verilog primitive which does not use high impedance. These are: and, nand, or, nor, xor, xnor, buf, not, pullup, pulldown and *user-defined primitives*.
- **3-state primitive:** an instance of a Verilog primitive which uses high impedance. These are: bufif0, bufif1, notif0, notif1, cmos, rcmos, pmos, rpmos, tran, rtran, tranif0, rtranif0, tranif1, rtranif1.
- **Module paths:** a path delay specified in a Verilog specify block from an input or inout port of a module to an output or inout port of the same module.
- **Module input ports (MIPD):** a delay on all signals which fan-in to an input port of a module instance. There is no construct in the Verilog language to represent module input port delays; but the PLI can add or modify delays on input ports.
- **Module input port bits:** a delay on all signals which fan-in to a bit of a vector port.
- **Module interconnect paths:** a connection from an output port of one module to the input port of another module. There is no construct in the Verilog language to represent interconnect path delays; but the PLI can add or modify delays on interconnect paths.
- **Timing constraint checks:** a timing constraint represented in a Verilog specify block, using \$setup, \$hold, etc.

Minimum, typical and maximum delays

Each Verilog delay transition can have:

- A single delay value, representing the typical delay.
- Three delay values, representing the minimum, typical, maximum delay range for that transition.

The Verilog language can represent minimum, typical and maximum delay values. Many Verilog simulators, however, only store a single value in the simulation data structure for each transition. The value that is stored is usually controlled by an invocation option (such as `+mindelays`). By default, the typical delay value is stored for each transition. For these types of simulators, the PLI does not know whether the single value stored for each transition was originally the minimum, typical, or maximum delay. The PLI considers the single value stored to always be the typical delay.

16.11 Reading module time scale information

Within a Verilog HDL model, delay values can be specified either as an integer or as a floating point real number. Within a Verilog simulation, however, time is represented as a 64 bit unsigned integer. As a Verilog model is compiled or loaded into simulation, delay values in the module are scaled to a simulation time unit. The scaling factor is specified with the `'timescale` compiler directive.

The `'timescale` directive in Verilog indicates what time units are used in the modules which follow the directive. The directive also indicates a time precision, which is how many decimal points of accuracy are permitted. Time values with more decimal points than the precision are rounded off. Each module in a design can have a different time scale, so one model can specify delays in nanoseconds with two decimal points of precision and another model can specify delays in microseconds with three decimal points. Most Verilog simulators will determine the finest precision of all modules in a simulation, and set that precision as the simulation time units. The simulator then scales the times in all modules to the simulation time units. As an example:

```
'timescale 1ns/10ps //nanosecond units, 2 decimal points
module A;
  ...
  nand #5.63 n1 (y, a, b);
  ...
endmodule
```

```
'timescale 1us/100ns //microsecond units, 3 decimal points
module B;
...
nand #3.581 n1 (y, a, b);
...
endmodule
```

In this example, a simulator will determine that ten-picosecond units of time is the finest precision used by all modules, and set that as the simulation time units. The simulator will then scale the 5.63 nanosecond delay in module **A** to 563 10-picosecond units, and it will scale the 3.581 microsecond delay in module **B** to 360,000 10-picosecond units (the rounding to the module's time precision occurs before the time value is scaled to the simulator's time units).

16.11.1 Reading time scale factors

The ACC routines can access both the time units and the time precision of a module.

void acc_fetch_timescale_info(object, timescale)

handle object handle for a module instance, module definition, PLI system task/function, or *null*.

p_timescale_info timescale pointer to an application-allocated **s_timescale_info** structure to receive the timescale information.

The **acc_fetch_timescale_info()** routine retrieves the time scale factors of a module, or a **\$timeformat** built-in system task. The information is retrieved into an application-allocated **s_timescale_info** structure, which is defined in **acc_user.h**, as follows:

```
typedef struct t_timescale_info
{
    short unit;
    short precision;
} s_timescale_info,
*p_timescale_info;
```

The object handle passed to **acc_fetch_timescale_info()** can be one of three types of objects:

- *A handle for a module instance or top-level module* — the time scale factors of that module will be retrieved and stored in the **s_timescale_info** structure.
- *A handle for a system task or system function* — the time scale factors of the module containing the instance of the task/function will be retrieved and stored in the

`s_timescale_info` structure. Handles for system tasks and system functions are obtained with the routine `acc_handle_tfinst()`.

- null — the time scale factors of the active `$timeformat` built-in Verilog system task will be retrieved and stored in the `s_timescale_info` structure.

The time units and time precision are retrieved as integers, which represent the exponent of the scale factor. For example, 1 nanosecond is 1 second times 10^{-9} , so the integer value used to represent nanoseconds is -9. Table 16-6 shows the representations of all time units and precisions supported in the Verilog language.

Exponent Value	Time Unit Represented
2	100 seconds (1×10^2)
1	10 seconds (1×10^1)
0	1 second (1×10^0)
-1	100 milliseconds (1×10^{-1})
-2	10 milliseconds (1×10^{-2})
-3	1 millisecond (1×10^{-3})
-4	100 microseconds (1×10^{-4})
-5	10 microseconds (1×10^{-5})
-6	1 microsecond (1×10^{-6})
-7	100 nanoseconds (1×10^{-7})
-8	10 nanoseconds (1×10^{-8})
-9	1 nanosecond (1×10^{-9})
-10	100 picoseconds (1×10^{-10})
-11	10 picoseconds (1×10^{-11})
-12	1 picosecond (1×10^{-12})
-13	100 femtoseconds (1×10^{-13})
-14	10 femtoseconds (1×10^{-14})
-15	1 femtosecond (1×10^{-15})

Table 16-6: Time unit exponents for `acc_fetch_timescale_info()`

int acc_fetch_precision()

The `acc_fetch_precision()` routine retrieves the internal simulation time units. The internal time units represent the finest increment of time for which the simulator can schedule events. In most Verilog simulators, the internal time unit will be the finest time precision of a Verilog modules which were compiled into the simulation. `acc_fetch_precision()` does not have any inputs, and the return value is an integer which represents the magnitude of 1 second in the same way as with `acc_fetch_timescale_info()`.

The following code fragment retrieves the time scale information of a module instance, the module containing the system task, and the current `$timeformat` system task.

```
s_timescale_info ts_info;
handle module_h, tfinst_h;

/* obtain handle for a module instance */
/* obtain handle for a system task instance */

acc_fetch_timescale_info(module_h, &ts_info);
io_printf("\nModule %s: units = %d precision = %d\n",
          acc_fetch_fullname(module_h),
          ts_info.unit, ts_info.precision);

acc_fetch_timescale_info(tfinst_h, &ts_info);
io_printf("\nSystem task: units = %d precision = %d\n",
          ts_info.unit, ts_info.precision);

acc_fetch_timescale_info(null, &ts_info);
io_printf("\n$timeformat units = %d precision = %d\n\n",
          ts_info.unit, ts_info.precision);
```

16.11.2 Reading the current simulation time

The ACC library does not have a routine to retrieve the current simulation time. The TF routines should be used to read the current simulation time in ACC applications. These routines are: `tf_gettime()`, `tf_getlongtime()`, `tf_getrealtime()` and `tf_str_gettime()`.

16.12 Reading delay values

There are several types of Verilog objects which have delays values. The ACC routines can both read and modify the delays of these objects.

16.12.1 Verilog objects which have delay values

Several types of constructs in the Verilog language can have delay values. How delays are represented for these different types of objects is part of the Verilog HDL standard, and is outside the scope of this book. The objects which can have delays are:

- **Primitive instances** can have delays specified for 1, 2 or 3 output transitions.
- **Module paths** can have delays specified for 1, 2, 3, 6 or 12 output transitions.
- **Module input ports** can have delays specified for 1, 2, or 3 output transitions. There is no construct in the Verilog language to represent Module Input Port Delays (MIPD's); only the PLI can add or modify delays on input ports.
- **Module interconnect paths** can have delays specified for 1, 2, or 3 output transitions. An interconnect path is the connection from the output of one module to the input of another module. There is no construct in the Verilog language to represent interconnect path delays; only the PLI can add or modify delays on input ports.
- **Timing constraint checks** can have delays specified for 1 limit for each constraint in the check. Timing constraints are represented in Verilog using \$setup, \$hold, etc.
- **Continuous assignments** can have delays specified for 1, 2, or 3 output transitions. The ACC routines cannot access the delays of continuous assignments. Only the VPI routines can access these delays.
- **Procedural time controls (#)** can have 1 delay specified. This delay represents the time before a statement is executed, rather than an output transition delay. The ACC routines cannot access the procedural delays. Only the VPI routines can access these delays.

In addition to multiple delay transitions, each delay value can be a set of delays representing a minimum, typical and maximum delay range. For example, a Verilog bufif1 tri-state buffer gate can represent a propagation delay as:

A tri-state buffer with no delays:

```
bufif1 g1 (...);
```

A tri-state buffer with delay of 5 for rising, falling, and turn-off transitions:

```
bufif1 #5 g2 (...);
```

A tri-state buffer with separate delays for rising, falling, and turn-off transitions:

```
bufuf1 #(3, 4, 5) g3(...);
```

A tri-state buffer with separate minimum:typical:maximum delay sets for rising, falling and turn-off transitions:

```
bufuf1 #(2:3:4, 3:4:5, 5:6:7) g4 (...);
```

16.12.2 Reading an object's delay values

A single ACC routine, `acc_fetch_delays()`, is used to read the delay values of any Verilog object. This single routine can be configured to read either typical delay values, or to read minimum, typical and maximum delay values. The arguments which are passed to the routine will change according to the configuration.

The syntax for `acc_fetch_delays()` when configured to read just typical delays is:

`bool acc_fetch_delays(object, d1, d2, ... d12)`

`handle object` handle for a primitive, module path, timing check, module input port, module input port bit select, or inter-module path.

`double *d1...*d12` pointer to variables to receive delay values.

The syntax for `acc_fetch_delays()` when configured to read minimum, typical and maximum delays is:

`bool acc_fetch_delays(object, dset_array)`

`handle object` handle for a primitive, module path, timing check, module input port, module input port bit select, or inter-module path.

`double *dset_array` pointer to an array of delay values.

The `acc_fetch_delays()` routine retrieves the current delay values of a *primitive instance*, a *module path*, a *module input port*, a *bit-select of a module input port* or a *module interconnect path*. The routine returns TRUE if it was successful, and FALSE if an error occurred.



The delay values retrieved into the PLI application will automatically be scaled to the time scale of the module containing the object from which the delays are read.

`acc_fetch_delays()` is overloaded in such a way that the number of arguments and the C data types of the arguments required by `acc_fetch_delays()` change, depending on:

- The type of object from which delays are being read.
- The setting of the ACC configuration for `accPathDelayCount`.
- The setting of the ACC configuration for `accMinTypMaxDelays`.

16.12.3 Reading typical delays

`acc_fetch_delays()` can be configured to read either the typical delays or the minimum, typical, maximum delays of an object. The default configuration is to read the typical delays of an object. This configuration can also be explicitly specified, using `acc_configure()`, as follows:

```
acc_configure(accMinTypMaxDelays, "false");
```

In the typical delay mode, `acc_fetch_delays()` requires as its arguments:

- A handle for the object from which to read the delay values.
- A pointer to a C double variable for the first delay value to be read.
- A pointer to a C double variable for the second delay value to be read.
- A pointer to a C double variable for the third delay value to be read.
- ...
- A pointer to a C double variable for the last delay value to be read.

The number of delay values which will be retrieved depends on the type of object. The `acc_fetch_type()` or `acc_object_of_type()` routines can be used to determine the object type.

- **2** delay values if the object is a *2-state primitive instance*.
- **3** delay values if the object is a *3-state primitive instance*.
- **3** delay values if the object is a *module input port*.
- **3** delay values if the object is a *module interconnect path*.
- **1** delay value if the object is a *timing constraint* check.
- **1, 2, 3, 6, or 12** delay values if the object is a *module path*. The number of delays is controlled by the configuration of `accPathDelayCount`. The default is 6.

The number of *module path* delay transitions to be retrieved by `acc_fetch_delays()` can be configured by the PLI application. A module path will always have 12 output transitions stored in the Verilog simulation data structure. The `acc_fetch_delays()` routine will automatically map the 12 delay values of the module path to the number of transitions requested by the PLI application. The configuration is set, using `acc_configure()`, as one of the following:

```
acc_configure(accPathDelayCount, "1");
acc_configure(accPathDelayCount, "2");
acc_configure(accPathDelayCount, "3");
acc_configure(accPathDelayCount, "6");
acc_configure(accPathDelayCount, "12");
```

The default configuration is "6".

It is an error to call `acc_fetch_delays()` with too few arguments for the number of delays of an object. If too many arguments are specified, the extra arguments are ignored. The maximum number of arguments which can be specified when reading only typical delay values is 13 (the object handle plus pointers to 12 delay variables).

Example 16-5 illustrates using `acc_fetch_delays()` to read the typical delays of all primitives in a module. Three delay values are retrieved, which represent the rise, fall and turn-off delays of the primitive.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.16/list_prim_delays_acc.c`
- Verilog test bench: `Chapter.16/list_prim_delays_test.v`
- Verilog-XL results log: `Chapter.16/list_prim_delays_test.log`

Example 16-5: `$list_prim_delays` — reading typical delays

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/*********************calltf application*****calltf application*****/
* calltf application
*********************calltf application*****calltf application*****/
int PLIBbook_PrimDelays_calltf()
{
    handle module_h, prim_h;
    double rise, fall, toZ;
```

```
static int three_state[7] = {accBufif0Gate, accBufif1Gate,
                            accNotif0Gate, accNotif1Gate,
                            accTranif0Gate, accTranif1Gate, 0};

acc_initialize();
acc_configure(accDisplayWarnings, "true");
acc_configure(accMinTypMaxDelays, "false");

module_h = acc_handle_tfarg(1);
io_printf("\nPrimitives in module %s:\n",
         acc_fetch_fullname(module_h));
prim_h = null; /* start with known value for target handle */
while (prim_h = acc_next_primitive(module_h, prim_h)) {
    io_printf(" %-8s instance %-4s: ",
              acc_fetch_defname(prim_h),
              acc_fetch_name(prim_h));

    if (acc_object_in_typeplist(prim_h, three_state)) {
        acc_fetch_delays(prim_h, &rise, &fall, &toZ);
        io_printf("rise=%2.2f, fall=%2.2f, toZ=%2.2f\n", rise, fall, toZ);
    }
    else {
        acc_fetch_delays(prim_h, &rise, &fall);
        io_printf("rise=%2.2f, fall=%2.2f\n", rise, fall);
    }
}

acc_close();
return(0);
}
```

16.12.4 Reading minimum, typical and maximum delays

`acc_fetch_delays()` can be configured to read the minimum, typical, maximum delays for each output transition of an object. This configuration must be explicitly specified, using `acc_configure()`, as follows:

```
acc_configure(accMinTypMaxDelays, "true");
```

In the min:typ:max delay mode, `acc_fetch_delays()` requires as its inputs:

- A handle for the object from which to read the delay values.
- A pointer to an array of C double variables, with one element in the array for each delay value to be read.

`acc_fetch_delays()` will retrieve a minimum, typical and maximum delay value for each delay transition that is read from the object. The total number of delay values which will be retrieved depends on the type of object, as follows:

- **6** delay values if the object is a *2-state primitive instance*.
- **9** delay values if the object is a *3-state primitive instance*.
- **9** delay values if the object is a *module input port*.
- **9** delay values if the object is a *module interconnect path*.
- **3** delay value if the object is a *timing constraint check*.
- **3, 6, 9, 18 or 36** delay values if the object is a module path. The number of delays is controlled by the configuration of `accPathDelayCount`.

Table 16-7 shows the order in which `acc_fetch_delays()` retrieves delay values.

Number of Delay Transitions	Order of retrieved delays			
	array element		object delay	
1 transition	[0]	receives	all transitions	min value
	[1]			typ value
	[2]			max value
2 transitions	[0]	receives	rise transition	min value
	[1]			typ value
	[2]			max value
	[3]	receives	fall transition	min value
	[4]			typ value
	[5]			max value
3 transitions	[0]	receives	rise transition	min value
	[1]			typ value
	[2]			max value
	[3]	receives	fall transition	min value
	[4]			typ value
	[5]			max value
	[6]	receives	turn-off transition	min value
	[7]			typ value
	[8]			max value

Table 16-7: Number elements and order of delay array

Number of Delay Transitions	Order of retrieved delays			
	array element		object delay	
6 transitions	[0]	receives	1st transition ¹	min value
	[1]			typ value
	[2]			max value

	[15]	receives	6th transition	min value
	[16]			typ value
	[17]			max value
12 transitions	[0]	receives	1st transition ¹	min value
	[1]			typ value
	[2]			max value

	[33]	receives	12th transition	min value
	[34]			typ value
	[35]			max value

Table 16-7: Number elements and order of delay array (continued)

¹In the preceding table, the order of the transitions for the 6 and 12 transition sets is:
 $0 \rightarrow 1, 1 \rightarrow 0, 0 \rightarrow Z, Z \rightarrow 1, 1 \rightarrow Z, Z \rightarrow 0, 0 \rightarrow X, X \rightarrow 1, 1 \rightarrow X, X \rightarrow 0, X \rightarrow Z, Z \rightarrow X$


NOTE

The Verilog language can represent minimum, typical and maximum delay values, but many Verilog simulators only store a single value in the simulation data structure for each transition. The value that is stored is usually controlled by an invocation option, and, by default, will be the typical delay value. If the simulator has not stored the full set of values, then the value that was stored for each transition will be used for all three minimum, typical and maximum fields in the delay array.

The array of C doubles is declared by the PLI application, and must be large enough to receive all the delay values which will be retrieved. Declaring the array with too few elements may not result in a PLI error, but could potentially crash a PLI application. It is not an error to declare the array with more elements than are needed. The additional elements in the array will be ignored by `acc_fetch_delays()`.

Example 16-6 lists a C function which reads the minimum, typical, maximum rise and fall delays of all module paths in a module. Note that module paths do not have an actual name in the Verilog language. The ACC routines create a derived module path name by concatenating together the names of the nets connected to the first input of the path and the first output of the path, with a '\$' in between the two names.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.16/list_path_delays_acc.c
- Verilog test bench: Chapter.16/list_path_delays_test.v
- Verilog-XL results log: Chapter.16/list_path_delays_test.log

Example 16-6: *\$list_path_delays* — reading min:typ:max delay values

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"           /* IEEE 1364 PLI ACC routine library */
/*********************************************
 * calltf application
 *****/
int PLIBook_PathDelays_calltf()
{
    handle module_h, path_h;
    double delay_set[18];
    int i;

    acc_initialize();
    acc_configure(accDisplayWarnings, "true");

    acc_configure(accMinTypMaxDelays, "true");
    acc_configure(accPathDelayCount, "6");

    module_h = acc_handle_tfarg(1);
    io_printf("\nPath delays in module %s:\n",
              acc_fetch_fullname(module_h));
    path_h = null;      /* start with known value for target handle */
    while (path_h = acc_next_modpath(module_h, path_h)) {
        io_printf(" %-12s : ",
                  acc_fetch_name(path_h));
        acc_fetch_delays(path_h, delay_set);
        for (i=0; i<18; i++) {
            if ( i == 0 ) /* format output like Verilog syntax */
                io_printf("(");
            else if ( (i % 3) )
                io_printf(":");
            else
                io_printf(",");
            io_printf("%1.1f", delay_set[i]);
        }
        io_printf("\n\n");
    }
    acc_close();
    return(0);
}
```

16.13 Writing delay values into an object

A pair of ACC routines, `acc_append_delays()` and `acc_replace_delays()`, are used to modify the delay values of Verilog objects. These routines work the same as when reading delays values. The routines can be configured to modify either typical delay values, or to modify minimum, typical and maximum delay values. The arguments which are passed to the routines change according to the configuration.

The `acc_append_delays()` routine adds delays to the existing delays of an object. Delays can be appended to a *primitive instance*, a *module path*, a *module input port* and a *bit-select of a module input port*. Delays cannot be appended to a module inter-connect path.

The `acc_replace_delays()` routine replaces any existing delays of an object. Delays can be replaced on a *primitive instance*, a *module path*, a *module input port*, a *bit-select of a module input port* and a *module interconnect path*.

The syntax for these routines, when configured to modify just typical delays, is:

`bool acc_append_delays(object, d1, d2,... d12)`
handle **object** handle for a primitive, module path, timing check, module
 input port or module input port bit select.
double d1...d12 delay values.

`bool acc_replace_delays(object, d1, d2,...d12)`
handle **object** handle for a primitive, module path, timing check, module
 input port, module input port bit select, or inter-module
 path.
double d1...d12 delay values.

The syntax for these routines, when configured to modify minimum, typical and maximum delays, is:

`bool acc_append_delays(object, dset_array)`
handle **object** handle for a primitive, module path, timing check, module
 input port or module input port bit select.
double * dset_array pointer to an array of delay values.

bool acc_replace_delays(object, dset_array)

handle object handle for a primitive, module path, timing check, module input port, module input port bit select, or inter-module path.

double *dset_array pointer to an array of delay values.



The delay values specified in the PLI application will automatically be scaled to the time scale of the module containing the object which is being modified.

The number of arguments and the types of arguments provided to `acc_append_delays()` and `acc_replace_delays()` will vary, based on the type of the object and settings of the ACC configurations. The factors which affect the arguments provided to these routines are:

- The type of object for which delays are being modified.
- The setting of the ACC configuration for `accPathDelayCount`.
- The setting of the ACC configuration for `accMinTypMaxDelays`.
- The setting of the ACC configuration for `accToHizDelay`.

16.13.1 Setting the number of module path delays to be modified

When the PLI modifies the delays of a module path, the number of transitions to be modified can be configured by the PLI application. The configuration is set by `acc_configure()`, as one of the following configurations:

```
acc_configure(accPathDelayCount, "1");
acc_configure(accPathDelayCount, "2");
acc_configure(accPathDelayCount, "3");
acc_configure(accPathDelayCount, "6");
acc_configure(accPathDelayCount, "12");
```

The default configuration is "6".

16.13.2 Setting the calculation of turn-off delays

Verilog three-state primitives, module input ports, and module interconnect paths can have three transitions: the *rise time*, *fall time* and *turn-off time*. The turn-off time is the amount of time it takes to transition from any logic value to high impedance.

The **accToHiZDelay** configuration determines whether a turn-off transition delay should be calculated by the PLI or specified as an input to `acc_append_delays()` or `acc_replace_delays()`. The configuration is specified as follows:

```
acc_configure(accToHiZDelay, "min");
acc_configure(accToHiZDelay, "max");
acc_configure(accToHiZDelay, "average");
acc_configure(accToHiZDelay, "from_user");
```

The first three configurations enable the PLI to automatically calculate the turn-off delay time, based on the rise and fall transition times. The calculation can use the shortest, the longest or the average of the rise and fall times. The “`from_user`” configuration indicates that the turn-off time for objects with three transitions will be specified as an input to `acc_append_delays()` or `acc_replace_delays()`.

The default setting of `accToHiZDelay` is “`from_user`”.

 **NOTE** The setting of `accToHiZDelay` is ignored if the object is a module path.

16.13.3 Modifying typical delays

`acc_append_delays()` and `acc_replace_delays()` can be configured to modify either the typical delays or the minimum, typical, maximum delays of an object. The default configuration is to modify the typical delays of an object. This configuration can also be explicitly specified, using `acc_configure()`, as follows:

```
acc_configure(accMinTypMaxDelays, "false");
```

In the typical delay mode, `acc_append_delays()` and `acc_replace_delays()` require as inputs:

- A handle for the object from which to modify the delay values.
- A C double variable or literal value for the first delay value to be modified.
- A C double variable or literal value for the second delay value to be modified.
- A C double variable or literal value for the third delay value to be modified.
- ...
- A C double variable or literal value for the last delay value to be modified.

The number of delay inputs will be based on the type of object and the ACC configurations:

- **2** delay values if the object is a *two-state primitive instance*.
- **2** delay values if the object is a *three-state primitive instance, module input port* or *module interconnect path* and accToHIZDelay is configured for “min”, “max” or “average”.
- **3** delay values if the object is a *three-state primitive instance, module input port* or *module interconnect path* and accToHIZDelay is configured for “from_user”.
- **1** delay value if the object is a timing constraint check.
- **1, 2, 3, 6, or 12** delay values if the object is a module path. The number of delays is controlled by the configuration of accPathDelayCount (refer back to section 16.13.1 on page 576).

It is an error to call acc_append_delays() and acc_replace_delays() with too few arguments for the type of object. If too many arguments are specified, the unnecessary arguments are ignored. The maximum number of arguments which can be specified when modifying typical delay values is 13 (the object handle plus 12 delay values).

The following C code fragment illustrates how to append typical rise and fall delays onto a Verilog primitive.

```
handle prim_h;
double add_rise, add_fall;
s_acc_value val_s;

acc_configure(accMinTypMaxDelays, "false");
acc_configure(accToHIZDelay, "average");
...
acc_append_delays(prim_h, add_rise, add_fall);
```

16.13.4 Modifying minimum, typical and maximum delays

acc_append_delays() and acc_replace_delays() can be configured to modify the minimum, typical, maximum delays for each output transition of an object. This configuration must be explicitly specified using acc_configure(), as follows:

```
acc_configure(accMinTypMaxDelays, "true");
```

In the min:typ:max delay mode, acc_append_delays() and acc_replace_delays() require as inputs:

- A handle for the object from which to read the delay values.
- A pointer to an array of C double variables, with one element in the array for each delay value to be modified.

The total number of delay values which will be modified depends on the type of object and the ACC configurations, as follows:

- **2** delay values if the object is a *two-state primitive instance*.
- **2** delay values if the object is a *three-state primitive instance, module input port* or *module interconnect path* and accToHIZDelay is configured for “min”, “max” or “average”.
- **3** delay values if the object is a *three-state primitive instance, module input port* or *module interconnect path* and accToHIZDelay is configured for “from_user”.
- **1** delay value if the object is a timing constraint check.
- **1, 2, 3, 6, or 12** delay values if the object is a module path. The number of delays is controlled by the configuration of accPathDelayCount.

The delay values to be written onto the object are loaded into the array by the PLI application prior to calling acc_append_delays() or acc_replace_delays(). The order of the delay values in the array is the same as with acc_fetch_delays(), which was shown earlier, in Table 16-7 on page 572.

NOTE → The Verilog language can represent minimum, typical and maximum delay values, but many Verilog simulators only store a single value in the simulation data structure for each transition. The value that is stored is usually controlled by an invocation option, and, by default, will be the typical delay value. If the simulator does not store the full set of values, then only the appropriate minimum, typical or maximum fields in the delay array will be used by the simulator.

The array of C doubles is declared by the PLI application, and must be large enough to hold all delay values which will be modified. Declaring the array without enough elements may result in an error, but could potentially crash a PLI application. It is not an error to declare the array with more elements than are needed.

Example 16-7 shows a useful PLI application called \$mipd_delays. The Verilog HDL does not have a construct to represent module input port delays, but these delays can

be added through the PLI. `$mipd_delays` provides a means for a Verilog model to add delays to the input port of a module. The usage of this application is:

```
$mipd_delays(<port_name>, <d1>, <d2>, ... <d9>)
```

Example:

```
$mipd_delays(in1, 1.4, 1.6, 1.9, 1.1, 1.3, 1.5, 0.6, 0.8, 0.9);
```



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.16/mipd_delays_acc.c`
- Verilog test bench: `Chapter.16/mipd_delays_test.v`
- Verilog-XL results log: `Chapter.16/mipd_delays_test.log`

Example 16-7: `$mipd_delays` — modifying min:typ:max delays

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"           /* IEEE 1364 PLI ACC routine library */
/*********************************************
 * calltf application
 *****/
int PLIBook_MipdDelays_calltf()
{
    double delay_array[9];
    double rise, fall, toZ;
    handle port_h;
    int i;

    acc_initialize();
    acc_configure(accDisplayWarnings, "true");

    acc_configure(accMinTypMaxDelays, "true");
    port_h = acc_handle_tfarg(1);

    /* most simulators return loconn handle, not port handle */
    if ( (acc_fetch_type(port_h) != accPort)
        && (acc_fetch_type(port_h) != accPortBit) )
        port_h = acc_next_port(port_h, null);
    if ( (acc_fetch_type(port_h) != accPort)
        && (acc_fetch_type(port_h) != accPortBit) ) {
        io_printf("ERR: $mipd_delays could not obtain port handle\n");
        return(0);
    }

    for (i = 0; i < 9; i++)
        delay_array[i] = acc_fetch_tfarg(i+2);

    acc_replace_delays(port_h, delay_array);
```

```
/* verify new delays took affect */
acc_configure(accMinTypMaxDelays, "false");
rise = fall = toZ = 0.0;
acc_fetch_delays(port_h, &rise, &fall, &toZ);
io_printf("Port %s new delays: (%1.2f, %1.2f, %1.2f)\n\n",
          acc_fetch_name(port_h),
          rise, fall, toZ);

return(0);
acc_close();
}
```

16.14 Reading parameter constant values

The Verilog HDL has two types of constants, `parameter` and `specparam`. Each type of constant can store integer values, real values or string values. There are specific ACC routines provided to determine what type of value is stored in a constant, and to read the value of the constant.

int acc_fetch_paramtype(object)

`handle object` handle for a parameter or specparam.

The `acc_fetch_paramtype()` routine returns an integer constant which represents what type of data that is stored in a Verilog `parameter` or `specparam` constant. The constant will be: `accIntegerParam` for integer values, `accRealParam` for floating point values, or `accStringParam` for ASCII string values

double acc_fetch_paramval(object)

`handle object` handle for a parameter or specparam.

`acc_fetch_paramval()` returns the value of what is stored in a `parameter` or `specparam` constant. *The value is always returned as a C double.* If the value stored in the parameter is an integer, the return from `acc_fetch_paramval()` can be cast to an `int`. A Verilog parameter can store 4-state logic. The `acc_fetch_paramval()` routine returns 2-state logic by converting logic X and Z to 0.

If the value stored in a constant is an ASCII string, `acc_fetch_paramval()` will retrieve the string into the ACC string buffer, and return a pointer to the string as a C double. Strings stored in the string buffer are temporary, and should be used immediately, or copied to application-allocated storage.

The following example illustrates accessing and printing the value of all parameter constants in a Verilog model.

NOTE

Some C compilers do not allow casting a C double to a string pointer. Therefore, this example first casts the double to an int, and then casts the int to a char *.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.16/list_parameters_acc.c
- Verilog test bench: Chapter.16/list_parameters_test.v
- Verilog-XL results log: Chapter.16/list_parameters_test.log

Example 16-8: *\$list_params* — reading parameter and specparam values

```
#include "veriuser.h"          /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"          /* IEEE 1364 PLI ACC routine library */
/********************* calltf application *****/
int PLIBook_ListParams_calltf()
{
    handle module_h, param_h;
    acc_initialize();
    acc_configure(accDisplayWarnings, "true");

    module_h = acc_handle_tfarg(1);
    io_printf("\nConstants in module %s:\n",
              acc_fetch_fullname(module_h));
    param_h = null;
    while (param_h = acc_next_parameter(module_h, param_h)) {
        io_printf(" Parameter %s is: ", acc_fetch_fullname(param_h));
        switch(acc_fetch_paramtype(param_h) ) {
            case accRealParam:
                io_printf("%f\n", acc_fetch_paramval(param_h));
                break;
            case accIntegerParam:
                io_printf("%d\n", (int)acc_fetch_paramval(param_h));
                break;
            case accStringParam:
                io_printf("%s\n", (char*)(int)acc_fetch_paramval(param_h));
        }
    }
    acc_close();
    return(0);
}
```

16.15 Using specparam constants as model attributes

Often a PLI application needs information about a model which is not part of the model functionality. A standard cell delay calculator, for example, might need the rise, slope and load factors of the cells used in a design. Information that is needed by the PLI application can be stored within the Verilog model, using Verilog parameter constants or specparam constants. The PLI application can then obtain a handle for the constant and read its value.

The ACC library supports a special usage of the Verilog specparam constant, called an **attribute**. This special usage requires adding a dollar sign (\$) to the end of the specparam constant name. An attribute specparam can be associated with all objects within a module, or with specific objects in a module.

- A **general attribute** is a specparam constant with a name which ends with a dollar sign. Every object within a Verilog module will be associated with a general attribute.
- An **object-specific attribute** is a specparam constant with a base name which ends with a dollar sign, followed by the name of some object in the module. Only the object which is named will be associated with that attribute.

The following example illustrates three specparam attributes:

```
module AN2 (o, a, b); // 2-input AND gate standard cell
    output o;
    input a, b;
    and (o, a, b);

    specify
        specparam BaseDelay$ = 2.2; //general attribute
        specparam InputLoad$a = 0.2; //object-specific attribute
        specparam InputLoad$b = 0.3; //object-specific attribute
    endspecify
endmodule
```

A specific set of ACC routines read the value of the specparam constants. These routines use the attribute name and a handle for the object associated with that attribute. These ACC routines by-pass the need to first obtain a handle for the specparam constant. The routines will search first for an object-specific attribute, then for a general attribute, and finally return a default value. Using these routines, it is not necessary to first obtain a handle for the specparam constant.

double acc_fetch_attribute(object, attribute, default)

handle object handle for a named object.
*char * attribute* name of the attribute associated with the object.
double default default value to be returned if the attribute does not exist for the object.

int acc_fetch_attribute_int(object, attribute, default)

handle object handle for a named object.
*char * attribute* name of the attribute associated with the object.
int default default value to be returned if the attribute does not exist for the object.

char *acc_fetch_attribute_str(object, attribute, default)

handle object handle for a named object.
*char * attribute* name of the attribute associated with the object.
double default default value to be returned if the attribute does not exist for the object.

The `acc_fetch_attribute()` routine returns the value of a specparam attribute as a C double. `acc_fetch_attribute_int()` returns the value as a C integer. `acc_fetch_attribute_str()` returns the value as a pointer to a C string. The string is stored in the temporary ACC string buffer.

These ACC routines access a specparam constant, using the name of the attribute (including the dollar sign) and a handle for the object with which the attribute is associated. For example, to read the value of one of the following attributes:

```
specparam InputLoad$a = 0.2; //object-specific attribute
specparam InputLoad$b = 0.3; //object-specific attribute
```

The following C code can be used to read the value of the attribute `InputLoad$b`:

```
double load_in;
handle port_h;

/* add code to obtain a handle for module input port b */

load_in = acc_fetch_attribute(port_h, "InputLoad$", 1.0);
```

The ACC routines will search first for an object-specific attribute, then for a general attribute, and finally return a default value.

The default value can be specified in two different ways:

- By setting the ACC configuration to return a default of zero.
- By providing a default value as a third argument to the fetch attribute routine.

The three fetch attribute routines use the `acc_configure()` routine to determine the source of the default return value. The syntax for `acc_configure()` was presented in Chapter 15, in section 15.2 on page 475.

- `acc_configure(accDefaultAttr0, "true")`
configures the fetch attribute routines to return a default value of zero (represented as 0.0, 0 or "0"). With this configuration, the third argument to the fetch attribute routines is not needed, and will be ignored. This is the default configuration.
- `acc_configure(accDefaultAttr0, "false")`
configures the fetch attribute routines to use the third argument of the routine as the default value.

16.16 Reading and modifying path pulse controls

The Verilog HDL defines a special `specparam` attribute to control how glitches propagate across a module path. The attribute name is **PATHPULSE\$**. The ACC library provides specific routines to read and modify the pulse control values.

A *pulse* is two transitions on the same module path that occur in a shorter period of time than the path delay. Pulse control values determine whether a pulse of a certain width can pass through a module path. The pulse control values consist of a *reject limit* and an *error limit* pair of values, where:

- The *reject limit* sets the threshold for when a pulse is rejected. Any pulse less than the *reject limit* will not propagate to the output of the path.
- The *error limit* sets the threshold for when a pulse generates an error. Any pulse less than the *error limit* and greater than or equal to the *reject limit* will propagate a logic X to the path output.
- A pulse that is greater than or equal to the *error limit* will propagate to the path output.

As with other attributes, the **PATHPULSE\$** attribute can be an object-specific attribute, which is associated with one specific module path, or it can be a general attribute which is associated with all module paths. The following Verilog source code fragment shows both types of pulse control attributes.

```

module and3 (out, in1, in2, in3);
    output out;
    input  in1, in2, in3;
    ...
    specify
        (in1 => out) = (4, 6);
        (in2 => out) = (4, 5);
        (in3 => out) = (3, 4);

        specparam PATHPULSE$in1$out = 2, 3;
        specparam PATHPULSE$ = 0, 4;

    endspecify
endmodule

```

In this example, the module path from `in1` to `out` has a *reject limit* of 2 and an *error limit* of 3. All other paths in the module have a *reject limit* of 0 and an *error limit* of 4.

There are two important facts to note about the syntax of `PATHPULSE$` attributes:

1. The attribute is assigned two values, the *reject limit* and the *error limit*. This usage does not conform to the normal syntax for `specparam` definitions. Normally, a `specparam` is assigned a single value. Because of this exception, the ACC routines to read `specparam` values (`acc_fetch_paramval()` and `acc_fetch_attribute()`) cannot be used to read the `PATHPULSE$` attribute values. These routines would only return the first value of the pulse control limits.
2. The object specific attribute appends the name of a module path to the end of the `PATHPULSE$` attribute name. However, module paths do not have a name in the Verilog language. Therefore, the `PATHPULSE$` attribute creates a path name by concatenating the path input name and the path output name together with a \$ (dollar sign) between the names. If the module path specifies multiple inputs or outputs, the names of the first input and output in the path are used to represent the path name.

The ACC library provides a specific set of routines to read and modify pulse control values.

`bool acc_fetch_pulsere(object, r1, e1, ... r12, e12)`

`handle object` handle for a module path.

`double *r1... *r12` pointer to variables to receive reject limit values.

`double *e1... *e12` pointer to variables to receive error limit values.

The `acc_fetch_pulsere()` routine retrieves the current pulse control limits of a module path. The inputs to this routine are a handle for a module path, and pointers to two C double variables for each pulse control value to be retrieved.

`bool acc_append_pulsere(object, r1, e1... r12, e12)`

`handle object` handle for a module path.

`double r1... r12` pulse reject limit values.

`double e1... e12` pulse error limit values.

The `acc_append_pulsere()` routine adds to the current pulse control limits of a module path.

`bool acc_replace_pulsere(object, r1, e1... r12, e12)`

`handle object` handle for a module path.

`double r1... r12` pulse reject limit values.

`double e1... e12` pulse error limit values.

The `acc_replace_pulsere()` routine replaces the current pulse control limits of a module path. The values specified are relative to the delays of the module path.

`void acc_set_pulsere(object, r_percent, e_percent)`

`handle object` handle for a module path.

`double r_percent` pulse reject limit.

`double e_percent` pulse error limit.

The `acc_set_pulsere()` routine replaces the current pulse control limits of a module path. The values specified are a percentage of the current delays of the module path.

The inputs to these routines are a handle for a module path, and a pair of C double variables for each pulse control value to be read or modified. The number of reject/error pairs required is controlled by the configuration of `accPathDelayCount`.

```
acc_configure(accPathDelayCount, "1");
acc_configure(accPathDelayCount, "2");
acc_configure(accPathDelayCount, "3");
acc_configure(accPathDelayCount, "6");
acc_configure(accPathDelayCount, "12");
```

The default configuration is "6".

It is an error to specify too few input arguments. If too many arguments are specified, the unnecessary arguments are ignored. The maximum number of arguments is 25 (a module path handle plus 12 pairs of reject limit and error limit variables).

NOTE The path control values specified in the PLI application will automatically be scaled to the time scale of the module containing the path which is being modified.

Within the Verilog language, a single *reject limit* and an *error limit* is specified for each module path. This single pulse control set is applied to all output transitions of the path. The PLI provides a much greater level of control over pulse control values, by allowing different reject and error limits to be specified for each output transition that can occur for a path delay. Table 16-8 shows how the number of paths specified are mapped to the pulse control values for path output transitions.

accPathDelayCount configuration	path delay output transitions represented
“1”	One reject/error pair of values for all transitions
“2”	One reject/error set of values for rising transitions One reject/error set of values for falling transitions
“3”	One reject/error set of values for rising transitions One reject/error set of values for falling transitions One reject/error set of values for toZ transitions
“6” (the default)	One reject/error set of values for 0->1 transitions One reject/error set of values for 1->0 transitions One reject/error set of values for 0->Z transitions One reject/error set of values for Z->1 transitions One reject/error set of values for 1->Z transitions One reject/error set of values for Z->0 transitions
“12”	One reject/error set of values for 0->1 transitions One reject/error set of values for 1->0 transitions One reject/error set of values for 0->Z transitions One reject/error set of values for Z->1 transitions One reject/error set of values for 1->Z transitions One reject/error set of values for Z->0 transitions One reject/error set of values for 0->X transitions One reject/error set of values for X->1 transitions One reject/error set of values for 1->X transitions One reject/error set of values for X->0 transitions One reject/error set of values for X->Z transitions One reject/error set of values for Z->X transitions

Table 16-8: Configuring the number of pulse control values

16.17 Summary

This chapter has presented the ACC routines which read property values, read/modify logic values and read/modify delay values in Verilog simulations. The ability to access objects anywhere in a Verilog design hierarchy, and to read and modify the values of those objects, provides a great deal of power for PLI applications.

The next chapter presents another powerful aspect of the ACC library, the ability to synchronize PLI application activity with simulation logic value changes. When the concepts presented in this chapter are combined with those presented in the next chapter, another useful capability of the PLI is made possible—interfacing C language models to Verilog simulations. Examples of how this is done are presented in Chapter 18.

CHAPTER 17

Synchronizing to Simulations Using the Value Change Link

One of the important capabilities provided by the ACC library is the ability to have a Verilog simulator call a PLI application whenever specific objects in the simulation change logic value. This chapter presents how to use the *Value Change Link* (*VCL*) routines of the ACC library.

The concepts presented in this chapter are:

- An overview of the VCL routines
- Adding VCL flags
- Removing VCL flags
- Using the VCL *consumer routine*
- Synchronizing PLI applications to simulation activity
- Interfacing C language models to Verilog simulations

17.1 An overview of the VCL routines

The *Value Change Link* (or *VCL*) routines in the ACC library schedule with a Verilog simulator to have a PLI application called whenever a specific object changes logic value or strength value. The application that is called is a VCL *consumer routine*.

There are many ways a PLI application can utilize the Value Change Link capability. Just a few possibilities are:

- A graphical display, such as a waveform display. Selected signals in a design can be monitored and each logic change on the signal recorded in a data base which can then be displayed graphically.
- To monitor test vector coverage of a design. All critical signals in a design can be monitored and a count maintained for the number of times each signal changes value. A summary report can be generated to show how often different parts of a design changed value during a simulation.
- To create an asynchronous interface to a C language model. When an input to the model changes during a Verilog simulation, the PLI *consumer routine* is called and can communicate the value change to the C model.

The typical steps of using the Value Change Link routines in a PLI application are:

1. A *calltf routine* or *misctf routine* obtains a handle for an object that the routine wishes to monitor for value changes.
2. The *calltf routine* or *misctf routine* adds a VCL flag to the object. The flag indicates the name of the *consumer routine* that should be called for logic value changes.
3. When the object changes logic value during simulation, the simulator calls the *consumer routine*.
4. The *consumer routine* processes whatever the PLI application needs to do, such as reading the new logic value of the object that changed.

17.2 Adding and removing VCL flags on Verilog objects

Value Change Link flags are added to an object using **acc_vcl_add()**, and are removed using **acc_vcl_delete()**.

17.2.1 Adding VCL flags

void acc_vcl_add(object, consumer, user_data, vcl_flag)

handle	object	handle for a net, register, event, port, primitive output terminal or primitive inout terminal.
int	* consumer	unquoted name of a C consumer routine.
char	* user_data	user-defined data value.
int	vcl_flag	constant: <i>vcl_verilog_logic</i> , <i>vcl_verilog_strength</i> .

The `acc_vcl_add()` routine adds a Value Change Link flag to a Verilog object. From that point on, until the flag is removed, whenever a value change occurs on that object, the specified *consumer routine* will be called.

Most objects which have logic values can have a Value Change Link flag added to the object. These objects are:

- Any *net* data type, including scalar nets, vector nets, part-selects of vector nets and bit-selects of vector nets
- The *reg* data type, including scalar regs, vector regs, part-selects of vector regs and bit-selects of vector regs
- The *integer* and *time* variable data types
- The *real* and *realtime* variable data types
- The *event* data type
- A *primitive output or inout terminal*
- A *scalar module port* or a *bit-select of a module port*

NOTE When a handle for a module port or port bit is specified, the VCL flag is placed on the loconn of the port. The loconn is the signal inside the module which is connected to that port. When the *consumer routine* is called by the simulator, the information passed to the *consumer routine* about the change will be based on the loconn signal.

The `consumer` argument passed to `acc_vcl_add()` is a pointer to the C function which the simulator should call when a change occurs on the object. A pointer to the function is the unquoted name of the function, without the parentheses after the function name. The PLI standard expects the *consumer routine* to be a C `int` function.

The `reason_flag` indicates what types of value changes should cause the *consumer routine* to be called. The `reason_flag` is one of the following two constants:

- `vcl_verilog_logic` indicates the *consumer routine* should be called for logic value changes on the object.
- `vcl_verilog_strength` indicates the *consumer routine* should be called for both logic value changes and strength level changes on the object.

The `user_data` value specified as an argument to `acc_vcl_add()` will be passed to the *consumer routine* whenever a callback occurs. Examples of using the `user_data` value are presented later in this chapter. The `user_data` is a character pointer, which can store a 32-bit value, or a pointer to a block of data. If no `user_data` value is needed, the argument should be set to null.

The `user_data` value can be used to pass the handle for an object to the *consumer routine*. Because the `user_data` is a pointer, several values can be stored in an application-allocated block of memory, and a pointer to the memory block placed in the `user_data` field.

When a pointer to a memory block is placed in the `user_data` field, the PLI application must ensure that the memory is persistent, and will be available when the *consumer routine* is called. Since the VCL flag will be added from a C function that was called by the simulator, such as a *calltf routine*, local automatic storage cannot be used to store data that is to be accessed from the *consumer routine*. Automatic storage will automatically be freed when the C function exits.

The following code fragment places a VCL flag on to a net, and stores a pointer to the name of the net in the `user_data` field. The net name is stored in memory that was allocated using `malloc`. Note that, in this example, the casting of the `user_data` is not necessary, since the value is already a character pointer. The casting is shown here to illustrate the need to cast values that are not the correct type (such as a handle for an object).

```
handle  net_handle;
char    *net_name;
char    *net_name_keep;

/* obtain a handle for a net from the first task arg */
net_handle = acc_handle_tfarg(1);

/* allocate memory for the name of the net */
net_name = acc_fetch_name(net_handle);
net_name_keep = malloc(strlen(net_name)+1);
strcpy(net_name_keep, net_name); /* save net_name */

/* add a VCL flag to net--user_data is pointer to net name */
acc_vcl_add(net_handle,
            my_consumer_routine,
            (char*)net_name_keep,
            vcl_verilog_logic);
```

Multiple VCL flags on an object

An object can have any number of VCL flags added to it, as long as each flag is unique. Either the name of the *consumer routine* or the value of the `user_data` must be different in order to make the flag unique. Using a different reason constant does not make the VCL flag unique.

17.2.2 Removing VCL flags from Verilog objects

```
void acc_vcl_delete(object, consumer, user_data, vcl_flag)
```

handle object handle for an object with a VCL flag.
*int *consumer* unquoted name of a C consumer function.
*char *user_data* user-defined data value.
int vcl_flag constant: *vcl_verilog*.

`acc_vcl_delete()` removes a Value Change Link flag from a Verilog object. The inputs to `acc_vcl_add()` are:

- A *handle* for an object which has a VCL flag.
- The *consumer routine name* which was specified when the VCL flag was added.
- The *user_data* value which was specified when the VCL flag was added.
- A *reason_flag* which must be constant **vcl_verilog**.

An object can have several VCL flags attached to it. Therefore, the arguments to `acc_vcl_delete()` must specify the same *consumer routine name* and the same *user_data* value as in the `acc_vcl_add()` which created the VCL flag.

17.3 Using the VCL consumer routine

When the simulator calls the *consumer routine*, the simulator allocates an **s_vc_record** structure and passes a pointer to the structure as an input to the *consumer routine*. The structure will contain information about the value change that occurred, including the simulation time and some form of representing the new logic value of the object.

The **s_vc_record** structure is defined in `acc_user.h`, and is listed on the next page.

```

typedef struct t_vc_record
{
    int vc_reason;
    int vc_hightime;
    int vc_lowtime;
    char *user_data;
    union
    {
        unsigned char logic_value;
        double real_value;
        handle vector_handle;
        s_strengths strengths_s;
    } out_value;
} s_vc_record, *p_vc_record;

```


The fields of the `s_vc_record` structure are filled in by the simulator. The `vc_reason` is a constant which represents the type of Verilog object which changed value, as listed in table 17-1.

vc_reason Constant	Description
<code>logic_value_change</code>	a scalar net or bit-select of a vector net changed logic value
<code>strength_value_change</code>	a scalar net or bit-select of a vector net changed strength
<code>vector_value_change</code>	a vector net or part-select of a vector net changed logic value
<code>sregister_value_change</code>	a scalar register or bit-select of a vector register changed logic value.
<code>vregister_value_change</code>	a vector register or part-select of a vector register changed logic value
<code>integer_value_change</code>	an integer variable changed value
<code>real_value_change</code>	a real variable changed value
<code>time_value_change</code>	a time variable changed value
<code>event_value_change</code>	an event data type had an event

Table 17-1: vc_reason constants

`vc_hightime` and `vc_lowtime` contain the simulation in which the value change occurred (which is the current simulation time). The 64-bit simulation time is stored as two 32-bit C integers, as shown in the following diagram.

```
unsigned int vc_hightime, vc_lowtime;
```



The simulation time is represented in the internal simulation time units. It is *not* scaled to the time units of the module containing the object with the VCL flag.

The **out_value** field contains information about the new logic value of the object which has the VCL flag. Because there are many types of objects which can have VCL flags, the **out_value** is a union of C data types. The field within the union that contains the value information is based on the **vc_reason** constant. Table 17-2 shows which **out_value** union field is used for the different reason constants.

vc_reason	out_value Field With New Value	The New Logic Value Representation
<code>logic_value_change</code>	<code>logic_value</code>	one of the constants: <code>vc0</code> <code>vc1</code> <code>vcX</code> <code>vcZ</code>
<code>strength_value_change</code>	<code>strengths_s</code>	a structure with the logic and strength
<code>vector_value_change</code>	<code>vector_handle</code>	a handle for a vector net or part-select of a vector net
<code>sregister_value_change</code>	<code>logic_value</code>	one of the constants: <code>vc0</code> <code>vc1</code> <code>vcX</code> <code>vcZ</code>
<code>vregister_value_change</code>	<code>vector_handle</code>	a handle for a vector reg or part-select of a vector reg
<code>integer_value_change</code>	<code>vector_handle</code>	a handle for an integer variable
<code>real_value_change</code>	<code>real_value</code>	the value of a real variable
<code>time_value_change</code>	<code>vector_handle</code>	a handle for a time variable
<code>event_value_change</code>	<code>none</code>	(event types have no logic value)

Table 17-2: s_vc_record out_value member fields

When the object which changed logic value is a vector signal, the simulator passes a handle for the object to the *consumer routine*. Using the handle, the routine can

retrieve the logic value using `acc_fetch_value()`. This allows the routine to retrieve the vector value in a wide variety of formats.

When the object which changed value is a scalar net or bit select of a vector net, and the VCL flag was for `vcl_verilog_logic`, the new value is represented as a constant, which encodes the Verilog 4-state logic. Note that the constants used to represent the logic values are defined as `char` data type. This is unique from most other PLI constants, which are of `int` data types.

When the object which changed is a scalar net or bit select of a vector net, and the VCL flag was for `vcl_verilog_strength`, the new value is represented as an `s_strengths` structure. Within the strength structure, the new logic value is represented with a constant which encodes the Verilog 4-state logic. The constants are the same as with scalar logic value changes: `vcl0`, `vcl1`, `vclX` and `vclZ`. The strength is represented as a pair of constants, which represent the 8 strength levels available for a logic 0 and logic 1. The strength constants are: `vclSupply`, `vclStrong`, `vclPull`, `vclLarge`, `vclWeak`, `vclMedium`, `vclSmall` and `vclHighz`. The `s_strengths` structure is defined in `acc_user.h`.

```
typedef struct t_strengths
{
    unsigned char logic_value;
    unsigned char strength1;
    unsigned char strength2;
} s_strengths, *p_strengths;
```

17.4 An example of using Value Change Link routines

This section contains a complete PLI application which uses the ACC Value Change Link capability. The example implements an application called `$my_monitor`, which monitors value changes of either scalar or vector nets, regs or variables. When a signal changes, the application prints the current simulation time, the old logic value and the new logic value of the signal. (Chapter 18 contains additional examples of using the Value Change Link routines.)



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.17/my_monitor_acc.c
- Verilog test bench: Chapter.17/my_monitor_test.v
- Verilog-XL results log: Chapter.17/my_monitor_test.log

Example 17-1: \$my_monitor — using the acc_vcl_add() routine

```
#include "veriuser.h"           /* IEEE 1364 PLI TF routine library */
#include "acc_user.h"           /* IEEE 1364 PLI ACC routine library */

/* prototype of consumer routine */
int PLIbook_MyMonitor_consumer();

/********************* checktf routine *********************/
int PLIbook_MyMonitor_checktf()
{
    int      i;
    static int validTypes[53] = {accNet, accNetBit, accReg, accRegBit, 0};

    acc_initialize();
    for (i=1; i<=tf_nump(); i++) {
        if ( !(acc_object_in_typeplist(acc_handle_tfarg(i), validTypes)) )
            tf_error("$my_monitor argument %d must be a net or reg", i);
        else if ( tf_sizep(i) != 1)
            tf_error("$my_monitor argument %d must be scalar", i);
    }
    acc_close();
    return(0);
}

/********************* calltf routine *********************/
typedef struct PLIbook_MyMon_t {
    char    signalName[256]; /* signal names--up to 255 characters */
    char    lastValue[2];   /* scalar logic value stored as a string */
} PLIbook_MyMon_s, *PLIbook_MyMon_p;

PLIbook_MyMonitor_calltf() {
    handle signal;
    int      i, numargs = tf_nump();

    /* allocate memory for an array of p_monitor structures */
    PLIbook_MyMon_p monArray; /* starting address for the array */
    monArray=(PLIbook_MyMon_p)malloc(numargs*(sizeof(PLIbook_MyMon_s)));

    acc_initialize();
    /* save name and current logic value of each signal */
    for (i=0; i<numargs; i++) {
        signal = acc_handle_tfarg(i+1);
        strcpy(monArray[i].signalName, acc_fetch_fullname(signal));
        strcpy(monArray[i].lastValue, acc_fetch_value(signal, "%b", null));
    /* add a VCL flag to each net--user_data is a pointer to saved info */
        acc_vcl_add(signal,
                    PLIbook_MyMonitor_consumer,
                    (char*)&monArray[i],

```

```

        vcl_verilog_logic);
}
acc_close();
return(0);
}
*****
* consumer routine
*****
int PLIbook_MyMonitor_consumer(p_vc_record vc_record)
{
    char newValue[2];
    /* retrieve pointer to data structure array from user_data field */
    PLIbook_MyMon_p ArrayElem_p = (PLIbook_MyMon_p)vc_record->user_data;
    switch (vc_record->vc_reason) { /* check reason call-back occurred */
        case logic_value_change: /* scalar net changed */
        case sregister_value_change : { /* scalar register changed */
            switch (vc_record->out_value.logic_value) { /* convert value */
                case vc10: strcpy(newValue, "0"); break; /* to string */
                case vc11: strcpy(newValue, "1"); break;
                case vc1X: strcpy(newValue, "x"); break;
                case vc1Z: strcpy(newValue, "z"); break;
            }
            io_printf("At time %4d: %-20s last value=%s new value=%s\n",
                      vc_record->vc_lowtime, ArrayElem_p->signalName,
                      ArrayElem_p->lastValue, newValue);
            strcpy(ArrayElem_p->lastValue, newValue); /* save the new value */
        }
    }
    return(0);
}

```

17.5 Summary

This chapter has presented one of the more powerful capabilities of the ACC library—the ability to have a Verilog simulator call a PLI application whenever any specific objects change logic value. The ACC routines used for Value Change Links are simple to use.

The next chapter shows additional ways to use the ACC Value Change Link capability. The VCL callbacks will be used to create both combinational and sequential interfaces between a Verilog simulation and hardware designs modeled in the C programming language.

CHAPTER 18

Interfacing to C Models Using ACC Routines

The Value Change Link in the ACC routines can be used to create a simple and efficient interface to C language models. The ACC routines can also be used in conjunction with the TF routines for synchronizing C models with changes in simulation time. The power of the ACC and TF libraries together with the flexibility of C programming provide countless ways to represent a C model interface. This chapter shows several of these ways.

The concepts presented in this chapter are:

- Representing hardware models in C
- Verilog HDL shell modules
- Combinational logic interfaces to C models
- Sequential logic interfaces to C models
- Synchronizing with the end of a simulation time step
- Synchronizing with a future simulation time step
- Multiple instances of a C model
- Creating instance specific storage within C models
- Representing propagation delays in C models

**TIP**

One reason for representing hardware models in the C language is to achieve faster simulation performance. The C programming language allows a very abstract, algorithmic representation of hardware functionality, without representing detailed timing, multi-state logic, hardware concurrency and the many other hardware specific details offered by the Verilog language.

The PLI can be a means to access the efficiency of a highly abstract C model. However, a poorly written PLI application can become a bottleneck that offsets much of the efficiency gains. Care must be taken to write PLI applications that execute as efficiently as possible.

Some guidelines that can help maximize the efficiency and run-time performance of PLI applications are:

- Good C programming practices are essential. General C programming style and technique are not discussed within the scope of this book.
- Consider every call to a PLI routine as expensive, and try to minimize the number of calls.
- Routines which convert logic values from a simulator's internal representation to C strings, and vice-versa, are very expensive in terms of performance. Best efficiency is attained when the value representation in C is as similar as possible to the value representation in Verilog.
- Use the Verilog language to model the things hardware description languages do well, such as representing hardware parallelism and hardware propagation times. Simulator vendors have invested a great deal in optimizing a simulator's algorithms, and that optimization should be utilized.

**NOTE**

The objective of this book is to show several ways in which the ACC library can be used to interface to C models. Short examples are presented that are written in a relatively easy to follow C coding style. In order to meet the book's objectives, the examples presented in this book do not always follow the guidelines of efficient C coding and prudent usage of the PLI routines. It is expected that when parts of these example PLI applications are adapted for other applications, the coding style will also be modified to be more efficient and robust.

18.1 How to interface C models with Verilog simulations

The power and flexibility of the C programming language and the Verilog PLI provide a wide variety of methods that can be used to interface a Verilog simulation with a C language model. All methods have three essential concepts in common:

- Value changes which occur in the Verilog simulator must be passed to the C model.
- Value changes within the C model must be passed to the Verilog simulation.


- Simulated time in both the Verilog simulation and the C model must remain synchronized.

This chapter will present some of the more common methods of interfacing a Verilog simulation with a C model. The methods presented are by no means the only ways this interface can be accomplished, and may not always be the most efficient method. However, the methods presented have many advantages, including simplicity to implement, portability to many types of Verilog simulators, and the ability to use the C model any number of times and anywhere in the hierarchy of a Verilog design.

The fundamental steps that are presented in this chapter are:

1. Create the C language model as an independent block of pure C code that does not use the PLI routines in any way. The C model will have inputs and outputs, but it will not know the source of the inputs or the destination of the outputs. The C code to implement the model might be in the form of a C function with no main function, or it might be a complete C program with its own main function.
2. Create a Verilog HDL *shell module* (also called a *wrapper module*) to represent the inputs and outputs of the C language model. This module will be written completely in the Verilog language, but will not contain any functionality. To represent the functionality of the model, the shell module will call a PLI application.
3. Create a PLI application to serve as an interface between the C model and the Verilog shell module. The PLI application is a communication channel, which:
 - Uses PLI routines to retrieve data from the Verilog HDL shell module and pass the data to the C model via standard C programming.
 - Uses standard C programming to receive data from the C model, and pass the data to the Verilog shell module via PLI routines.

The following diagram shows how the blocks which are created in these three steps interact with each other.



This chapter presents steps 2 and 3 of this interface method in detail. Step 1 is to model some desired functionally or algorithm in the C language. This step is pure C programming, which does not directly involve the Verilog language or the Verilog PLI. This chapter does not cover how to implement ideas in the C language—the focus is on how to interface that implementation with a Verilog simulation. To maintain this focus, the C model example presented in this chapter will be a practical example, but relatively simple to model in C. The C model example used illustrates all of the important concepts of integrating C models into a Verilog simulation.

18.2 Creating the C language model

A hardware model can be represented in the C programming language in two basic forms, either as a C function or as an independent C program.

18.2.1 Using functions to represent the C model

When the C model is represented as a C function, that function can be linked into the Verilog simulator, together with the PLI application that serves as the interface to the model. The PLI application can then call the function when needed, passing inputs to the function, and receiving outputs from the function. One advantage of representing a C model as a function is the simplicity of passing values to and from the model. Another advantage is ease of porting to different operating systems, since the C model is called directly from the PLI application as a C function. A disadvantage of using a function to represent the C model is that the C model must contain additional code to allow a Verilog design to instantiate the C model multiple times. The model needs to specifically create unique storage for each instance.

18.2.2 Using independent programs to represent the C model

When the C model is represented as an independent program, which means it has its own C *main* function, then the Verilog simulation and the C model can be run as parallel processes on the same or on different computers. The PLI application which serves as an interface between the simulation and the model will need to create and maintain some type of communication channel between the two programs. This communication can be accomplished several ways, such as using the *exec* command in the C standard library. On Unix operating systems, the *fork* or *vfork* commands with either Unix pipes or Unix sockets is an efficient method to communicate with the C model program. On PC systems running a DOS or windows operating system, the *spawn* command can be used to invoke the C model program and establish two-way communications between the PLI application and the C model process.


One of the advantages of representing the C model as an independent model is the ability to have parallel processes running on the same computer or separate computers. Another advantage is that when a Verilog design instantiates multiple instances of the C model, each instance will be a separate process with its own memory storage. The major disadvantage of independent programs when compared to using a C function to represent the C model is that the PLI interface to invoke and communicate with the separate process is more complex, and might be operating system dependent.

David Roberts, of Cadence Design Systems, who reviewed many of the chapters of this book, has provided a full example of representing a C model as a separate C program. This example is included with the CD that accompanies this book.

18.3 A C model example

The C model used for the different PLI interfaces shown in this chapter is a scientific Arithmetic Logic Unit, which utilizes the C math library. The C model is represented as a C function, which will be called from the PLI interface mechanism. This model is written entirely with standard C library routines and C data types, without reference to any PLI routines or PLI data types. This same example is also used in other chapters, to show how a PLI interface to C models can be created using the VPI and TF libraries of the PLI.

The inputs and outputs of the scientific ALU C model are shown below, and Table 18-1 shows the operations which the ALU performs.



exception is set to 1 whenever an operation results in a value which is out of range of the double-precision result.

error is set to 1 whenever an input to an operation is out of range for the operation.

Opcode	C Math Library Operation
0	<code>pow(a, b)</code> — returns a to the power of b
1	<code>sqrt(a)</code> — returns the square root of a
2	<code>exp(a)</code> — returns the natural exponent of a
3	<code>ldexp(a, b)</code> — returns a * (2 to the power of b)
4	<code>fabs(a)</code> — returns the absolute of a
5	<code>fmod(a, b)</code> — returns the floating remainder of a / b
6	<code>ceil(a)</code> — returns smallest whole number not less than a
7	<code>floor(a)</code> — returns largest whole number not more than a
8	<code>log(a)</code> — returns the natural log of a
9	<code>log10(a)</code> — returns the base 10 log of a
A	<code>sin(a)</code> — returns the sine of a
B	<code>cos(a)</code> — returns the cosine of a
C	<code>tan(a)</code> — returns the tangent of a
D	<code>asin(a)</code> — returns the arcsine of a
E	<code>acos(a)</code> — returns the arccosine of a
F	<code>atan(a)</code> — returns the arctangent of a

Table 18-1: Scientific ALU C model operations

The source code for the scientific ALU is listed in Example 18-1. This version of the ALU generates a result, but does not store the result. A latched version of the ALU which stores the operation result is listed later in this chapter.



The source code for this example is on the CD accompanying this book.

- Application source file: `Chapter.18/sci_alu_combinational_acc.c`
- Verilog shell module: `Chapter.18/sci_alu_combinational_shell.v`
- Verilog test bench: `Chapter.18/sci_alu_combinational_test.v`
- Verilog-XL results log: `Chapter.18/sci_alu_combinational_test.log`

Example 18-1: scientific ALU C model — combinational logic version

```
#include <math.h>
#include <ERRNO.h>
void PLIbook_ScientificALU_C_model(
    double a,          /* input */
    double b,          /* input */
    int   opcode,      /* input */
    double *result,    /* output from ALU */
    int   *excep,      /* output; set if result is out of range */
    int   *err)        /* output; set if input is out of range */
{
    switch (opcode) {
        case 0x0: *result = pow    (a, b);      break;
        case 0x1: *result = sqrt   (a);         break;
        case 0x2: *result = exp    (a);         break;
        case 0x3: *result = ldexp  (a, (int)b);  break;
        case 0x4: *result = fabs   (a);         break;
        case 0x5: *result = fmod   (a, b);      break;
        case 0x6: *result = ceil   (a);         break;
        case 0x7: *result = floor  (a);         break;
        case 0x8: *result = log    (a);         break;
        case 0x9: *result = log10  (a);         break;
        case 0xA: *result = sin    (a);         break;
        case 0xB: *result = cos    (a);         break;
        case 0xC: *result = tan    (a);         break;
        case 0xD: *result = asin   (a);         break;
        case 0xE: *result = acos   (a);         break;
        case 0xF: *result = atan   (a);         break;
    }
    *err   = (errno == EDOM); /* arg to math func. out of range */
    *excep = (errno == ERANGE); /* result of math func. out of range */
    errno = 0;                /* clear the error flag */
    if (*err) *result = 0.0;   /* set result to 0 if error occurred */
    return;
}
```

18.4 Creating a Verilog shell module

A **shell module** allows a Verilog design to reference a C model using standard Verilog HDL syntax. The shell module is a Verilog module which has the same input and output ports as the C model, but the module has no functionality modeled within. To represent the module's functionality, the shell module invokes a PLI application, which in turn invokes the C model. A shell module is sometimes referred to as a *wrapper module*, because the module is wrapped around the call to a PLI application.

The shell module for a combinational logic version of the scientific ALU is listed below.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.18/sci_alu_combinational_acc.c
- Verilog shell module: Chapter.18/sci_alu_combinational_shell.v
- Verilog test bench: Chapter.18/sci_alu_combinational_test.v
- Verilog-XL results log: Chapter.18/sci_alu_combinational_test.log

Example 18-2: Verilog shell module for the Scientific ALU C model

```

`timescale 1ns / 1ns
module scientific_alu(a_in, b_in, opcode,
                      result_out, exception, error);
  input  [63:0] a_in, b_in;
  input  [3:0]  opcode;
  output [63:0] result_out;
  output      exception, error;

  real          a, b, result; // real variables used in this module
  reg           exception, error;

  // convert real numbers to/from 64-bit vector port connections
  assign result_out = $realtobits(result);
  always @(a_in) a  = $bitstoreal(a_in);
  always @(b_in) b  = $bitstoreal(b_in);

  //call the PLI application which interfaces to the C model
  initial
    $scientific_alu(a, b, opcode, result, exception, error);

endmodule

```



NOTE In this scientific ALU example, the primary inputs and outputs of the model are double-precision floating point values, represented as Verilog `real` data types. The Verilog language does not permit real numbers to be connected to module ports. However, the language provides built-in system functions which convert real numbers to 64-bit vectors, and vice-versa, so the real values can be passed through a module port connection. These built-in system functions are `$realtobits()` and `$bitstoreal()`.

The Verilog shell module that represents the C model can be instantiated in a design in the same way as any other Verilog module. For example:

```
module chip (...)  
    ...  
    scientific_alu u1 (a, b, opcode, result, excep, err);  
    ...  
endmodule
```

Creating a shell module to represent the C model is not mandatory—the PLI application could be called directly from any place in a Verilog design. However, there are important advantages to using a shell module to represent the C model:

- The shell module provides a simple method to encapsulate the C model.
- The shell module can be instantiated anywhere in a Verilog design hierarchy.
- The shell module can be instantiated any number of times in a Verilog design.
- The shell module can add Verilog HDL delays to the C model, which can accurately represent rise and fall delay delays, state-dependent delays, and timing constraints such as setup times.
- Delays within a shell module can be annotated using delay calculators or SDF files for additional delay accuracy for each instance of the shell module.

Section 18.10 later in this chapter discusses how delays can be represented in the Verilog shell module.

18.5 Creating a combinational logic interface to a C model

In a combinational logic model, the outputs of the model continuously reflect the input values of the model. The inputs are asynchronous—when any input changes value, the model outputs are re-evaluated to reflect the input change.

The ACC Value Change Link provides an efficient method of creating a combinational logic interface to a C model. VCL flags can be attached to each input of the C model. Whenever an input changes, the new input values can be passed to the C model. The outputs of the C model are then passed back to the Verilog simulation by writing the results onto the system task arguments in the Verilog shell module. How the ACC Value Change Link is used was presented in Chapter 17.

The basic steps involved with using the Value Change Link to implement a combinational logic interface are:

1. Create a PLI application system task to represent the interface between the Verilog shell module and the C model. The system task is invoked from the shell module, and all of the C model inputs and outputs are listed as arguments to the system task. For example:

```
initial  
    $scientific_alu(a, b, opcode, result, exception, error);
```

Note that, in this example, the system task is called from a Verilog **initial** procedure, which means the task will only be invoked one time for each instance of the shell module.

2. In the *calltf routine* associated with the system task, add VCL flags to each system task argument which represents an input to the C model. The same VCL *consumer routine* is listed for each model input.
3. In the VCL *consumer routine*, which is called whenever an input changes value, read the values of all inputs and pass the values to the C model. The output values of the C model are returned to the same *consumer routine*, which then writes the values to the system task arguments that represent the outputs of the C model.

Obtaining object handles within VCL consumer routines

NOTE → The VCL *consumer routine* is not directly associated with the system task which represents the C model interface. This means the *consumer routine* cannot obtain handles to the system task arguments using `acc_handle_tfarg()`, unless a handle to the system task is passed to the *consumer routine*.

The VCL *consumer routine* needs the handles for the system task arguments which represent the C model, in order to read the input values and write the C model output values. Since the *consumer routine* is not associated with the system task, access to the task arguments must be passed to the *consumer routine* through the VCL `user_data` field. This information can be passed in either of two ways:

- When the *calltf routine* sets up the VCL callbacks, the *calltf routine* can obtain a handle for the system task instance, and pass the handle to the *consumer routine* through the VCL `user_data` field. The *consumer routine* can then obtain the handle for any of the system task arguments using `acc_handle_itfarg()`, which is the instance specific version of `acc_handle_tfarg()`.
- When the *calltf routine* sets up the VCL callbacks, the *calltf routine* can allocate persistent storage. The handles for the system task arguments can be stored in the memory which was allocated. A pointer to the storage can be passed to the *consumer routine* through the VCL `user_data` field. The *consumer routine* can then obtain the handles from the storage block.

Each of these methods has advantages. Saving just the handle for the system task instance is simpler, and makes the instance handle available for other uses, such as scheduling a callback to the *miscif routine* for that instance of the system task. Saving all handles in a block of memory means the handles do not need to be obtained each time a value change occurs. Example 18-3, which follows, illustrates passing the system task instance handle, and example 18-4 on page 615 illustrates allocating a block of memory to store the handles for all system task arguments, and passing a pointer to the memory to the consumer routine.

The following example implements a combinational logic interface for the scientific ALU C model.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.18/sci_alu_combinational_acc.c
- Verilog shell module: Chapter.18/sci_alu_combinational_shell.v
- Verilog test bench: Chapter.18/sci_alu_combinational_test.v
- Verilog-XL results log: Chapter.18/sci_alu_combinational_test.log

Example 18-3: combinational logic C model interface using ACC routines

```
#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
#include "acc_user.h" /* IEEE 1364 PLI ACC routine library */

#define ALU_A      1 /* system task arg 1 is ALU A input */
#define ALU_B      2 /* system task arg 2 is ALU B input */
#define ALU_OP     3 /* system task arg 3 is ALU opcode input */
#define ALU_RESULT 4 /* system task arg 4 is ALU result output */
#define ALU_EXCEPT 5 /* system task arg 5 is ALU exception output */
#define ALU_ERROR   6 /* system task arg 6 is ALU error output */

/*********************  
 * VCL simulation callback routine: Serves as an interface between  
 * Verilog simulation and the C model. Called whenever the C model  
 * inputs change value, passes the values to the C model, and puts  
 * the C model outputs into simulation.  
 *****/  
int PLIBook_ScientificALU_interface(p_vc_record vc_record)  
{  
    double          a, b, result;  
    int            opcode, excep, err;  
    handle         instance_h, result_h, excep_h, err_h,  
                  a_h, b_h, opcode_h;  
    s_setval_value value_s;  
    s_setval_delay delay_s;  
    s_acc_time     time_s;
```

```
acc_initialize();

/* Retrieve instance handle from VCL user_data field */
instance_h = (handle)vc_record->user_data;

/* Obtain handles to all task arguments */
result_h = acc_handle_itfarg(ALU_RESULT, instance_h);
excep_h = acc_handle_itfarg(ALU_EXCEPT, instance_h);
err_h = acc_handle_itfarg(ALU_ERROR, instance_h);
a_h = acc_handle_itfarg(ALU_A, instance_h);
b_h = acc_handle_itfarg(ALU_B, instance_h);
opcode_h = acc_handle_itfarg(ALU_OP, instance_h);

/* Read current values of C model inputs from Verilog simulation */
value_s.format = accRealVal;
acc_fetch_value(a_h, "%", &value_s);
a = value_s.value.real;

acc_fetch_value(b_h, "%", &value_s);
b = value_s.value.real;

value_s.format = accIntVal;
acc_fetch_value(opcode_h, "%", &value_s);
opcode = value_s.value.integer;

***** Call C model *****/
PLIbook_ScientificALU_C_model(a, b, opcode, &result, &excep, &err);

/* Write the C model outputs onto the Verilog signals */
delay_s.model = accNoDelay;
delay_s.time = time_s;
delay_s.time.type = accRealTime;
delay_s.time.real = 0.0;

value_s.format = accRealVal;
value_s.value.real = result;
acc_set_value(result_h, &value_s, &delay_s);

value_s.format = accIntVal;
value_s.value.integer = excep;
acc_set_value(excep_h, &value_s, &delay_s);

value_s.value.integer = err;
acc_set_value(err_h, &value_s, &delay_s);

acc_close();
return(0);
}
```

```
*****
 * calltf routine: Registers a callback to the C model interface
 * whenever any input to the C model changes value
 ****
int PLIbook_ScientificALU_calltf()
{
    handle instance_h, a_h, b_h, opcode_h;

    acc_initialize();

    /* get handles for signals in task args which are C model inputs */
    a_h      = acc_handle_tfarg(ALU_A);
    b_h      = acc_handle_tfarg(ALU_B);
    opcode_h = acc_handle_tfarg(ALU_OP);

    /* get handles for this system task instance to pass to VCL app. */
    instance_h = acc_handle_tfinst();

    /* add VCL flags to all signals which are inputs to the C model */
    /* pass handle for task instance as the user_data value */
    acc_vcl_add(a_h, PLIbook_ScientificALU_interface,
                (char*)instance_h, vcl_verilog_logic);
    acc_vcl_add(b_h, PLIbook_ScientificALU_interface,
                (char*)instance_h, vcl_verilog_logic);
    acc_vcl_add(opcode_h, PLIbook_ScientificALU_interface,
                (char*)instance_h, vcl_verilog_logic);

    acc_close();
    return(0);
}

*****
 * checktf routine: Verifies that $scientific_alu() is used correctly.
 * Note: For simplicity, only limited data types are allowed for
 * task arguments. Could add checks to allow other data types.
 ****
int PLIbook_ScientificALU_checktf()
{
    acc_initialize();

    if (tf_nump() != 6)
        tf_error("$scientific_alu requires 6 arguments");

    else {
        if (!(acc_object_of_type(acc_handle_tfarg(ALU_A), accRealVar)))
            tf_error("$scientific_alu arg 4 must be a real variable\n");

        if (!(acc_object_of_type(acc_handle_tfarg(ALU_B), accRealVar)))
            tf_error("$scientific_alu arg 5 must be a real variable\n");

        if (!(acc_object_of_type(acc_handle_tfarg(ALU_OP), accWire)))
            tf_error("$scientific_alu arg 6 must be a net\n");
        else if (acc_fetch_size(acc_handle_tfarg(ALU_OP)) != 4)
            tf_error("$scientific_alu arg 6 must be a 4-bit vector\n");
    }
}
```

```

if (!(acc_object_of_type(acc_handle_tfarg(ALU_RESULT), accRealVar)))
    tf_error("$scientific_alu arg 1 must be a real variable\n");

if (!(acc_object_of_type(acc_handle_tfarg(ALU_EXCEPT), accReg)))
    tf_error("$scientific_alu arg 2 must be a reg\n");
else if (acc_fetch_size(acc_handle_tfarg(ALU_EXCEPT)) != 1)
    tf_error("$scientific_alu arg 2 must be scalar\n");

if (!(acc_object_of_type(acc_handle_tfarg(ALU_ERROR), accReg)))
    tf_error("$scientific_alu arg 3 must be a reg\n");
else if (acc_fetch_size(acc_handle_tfarg(ALU_ERROR)) != 1)
    tf_error("$scientific_alu arg 3 must be scalar\n");
}

acc_close();
return(0);
}

```

18.6 Creating a sequential logic interface to a C model

In a sequential logic model, the outputs of the model change synchronously with an input strobe, such as a positive edge of a clock. There may also be one or more asynchronous inputs, such as a reset signal.

The ACC Value Change Link is a straightforward way to model a sequential logic interface to a C model. VCL flags are attached only to the clock input and any asynchronous inputs of the C model. Only when those inputs change, are the new input values passed to the C model.

The basic steps involved with using the Value Change Link to implement a synchronous sequential logic interface are very similar to implementing a combinational logic interface. The one difference is that VCL flags are only attached to specific C model inputs instead of all inputs.

Regardless of whether the interface is combinational or sequential logic, the VCL *consumer routine* needs the handles for the system task arguments which represents the C model, in order to read the inputs values and write the C model output values. Access to the task arguments must be passed to the *consumer routine*.

Example 18-4, which follows, illustrates allocating a block of memory to store the handles for all system task arguments. This example implements a sequential logic interface for the scientific ALU C model, where all inputs are synchronized to value changes of a clock input.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.18/sci_alu_sequential_acc.c
- Verilog shell module: Chapter.18/sci_alu_sequential_shell.v
- Verilog test bench: Chapter.18/sci_alu_sequential_test.v
- Verilog-XL results log: Chapter.18/sci_alu_sequential_test.log

Example 18-4: sequential logic C model interface using ACC routines

```
#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
#include "acc_user.h" /* IEEE 1364 PLI ACC routine library */

#define ALU_CLOCK 1 /* system task arg 1 is ALU clock input */
#define ALU_A 2 /* system task arg 2 is ALU A input */
#define ALU_B 3 /* system task arg 3 is ALU B input */
#define ALU_OP 4 /* system task arg 4 is ALU opcode input */
#define ALU_RESULT 5 /* system task arg 5 is ALU result output */
#define ALU_EXCEPT 6 /* system task arg 6 is ALU exception output */
#define ALU_ERROR 7 /* system task arg 7 is ALU error output */

/*********************  
 * Definition for a structure to hold the data to be passed from  
 * calltf routine to the ALU interface (a VCL consumer routine).  
 *******************/  
  
typedef struct PLIbook_SciALU_data {  
    handle clock_h, a_h, b_h, opcode_h, result_h, excep_h, err_h;  
} PLIbook_SciALU_data_s, *PLIbook_SciALU_data_p;  
  
/*********************  
 * VCL simulation callback routine: Serves as an interface between  
 * Verilog simulation and the C model. Called whenever the C model  
 * clock input changes value, reads the values of all model inputs,  
 * passes the values to the C model, and writes the C model outputs  
 * into simulation.  
 *******************/  
  
int PLIbook_ScientificALU_interface(p_vc_record vc_record)  
{  
    double a, b, result;  
    int opcode, excep, err;  
    s_setval_value value_s;  
    s_setval_delay delay_s;  
    s_acc_time time_s;  
  
    PLIbook_SciALU_data_p ALUdata;  
  
    acc_initialize();  
  
    /* Retrieve pointer to ALU data structure from VCL user_data field */  
    ALUdata = (PLIbook_SciALU_data_p)vc_record->user_data;
```

```
/* Read current values of C model inputs from Verilog simulation */
value_s.format = accRealVal;
acc_fetch_value(ALUdata->a_h, "%%", &value_s);
a = value_s.value.real;

acc_fetch_value(ALUdata->b_h, "%%", &value_s);
b = value_s.value.real;

value_s.format = accIntVal;
acc_fetch_value(ALUdata->opcode_h, "%%", &value_s);
opcode = value_s.value.integer;

***** Call C model *****/
PLIBook_ScientificALU_C_model(0, a, b, opcode, &result, &excep, &err);

/* Write the C model outputs onto the Verilog signals */
delay_s.model      = accNoDelay;
delay_s.time        = time_s;
delay_s.time.type   = accRealTime;
delay_s.time.real   = 0.0;

value_s.format      = accRealVal;
value_s.value.real  = result;
acc_set_value(ALUdata->result_h, &value_s, &delay_s);

value_s.format      = accIntVal;
value_s.value.integer = excep;
acc_set_value(ALUdata->excep_h, &value_s, &delay_s);

value_s.value.integer = err;
acc_set_value(ALUdata->err_h, &value_s, &delay_s);

acc_close();
return(0);
}

*****
* calltf routine: Registers a callback to the C model interface
* whenever the clock input to the C model changes value
*****
int PLIBook_ScientificALU_calltf()
{
    PLIBook_SciALU_data_p ALUdata;

    acc_initialize();

    ALUdata=(PLIBook_SciALU_data_p)malloc(sizeof(PLIBook_SciALU_data_s));

    /* get handles for all signals in Verilog which connect to C model */
    ALUdata->clock_h  = acc_handle_tfarg(ALU_CLOCK);
    ALUdata->a_h       = acc_handle_tfarg(ALU_A);
    ALUdata->b_h       = acc_handle_tfarg(ALU_B);
```

```
ALUdata->opcode_h = acc_handle_tfarg(ALU_OP);
ALUdata->result_h = acc_handle_tfarg(ALU_RESULT);
ALUdata->excep_h  = acc_handle_tfarg(ALU_EXCEPT);
ALUdata->err_h    = acc_handle_tfarg(ALU_ERROR);

/* add VCL flag to the clock input to the C model */
/* pass pointer to storage for handles as user_data value */
acc_vcl_add(ALUdata->clock_h, PLIbook_ScientificALU_interface,
            (char*)ALUdata, vcl_verilog_logic);

acc_close();
return(0);
}
```

18.7 Synchronizing with the end of a simulation time step

Within a simulation, several inputs to the C model might change at the same moment in simulation time. In Verilog simulators, the Value Change Link *consumer routine* will be called for each input change, which means these routine may be called before all input value changes have occurred for that time step.

With a combinational logic interface, the *consumer routine* will be called for every input change. This is the correct functionality for combinational logic—at the completion of a simulation time step, the outputs from the C model represent the most current input values. However, by synchronizing the call to the C model with the end of the simulation time step in which changes occur, the multiple calls to the C model within a time step could be optimized to a single call.

With a sequential logic interface synchronized to a clock, when the *consumer routine* is called at a clock change, other input changes at that moment in simulation time may or may not have occurred. It may be desirable to ensure that the C model is not called until all inputs have their most current value for the time step in which the clock changes.

By using both the ACC and TF libraries, both combinational logic and sequential logic C model interfaces can be synchronized to the end of a current simulation time step. This is done by using the VCL *consumer routine* to schedule a callback to the *misctf routine* for that instance of the system task which represents the C model.

Since the VCL *consumer routine* is not directly associated with a system task, a handle to the system task must be passed to the *consumer routine*. The *consumer routine* can then schedule a callback to the *misctf routine* using `tf_isynchronize()`, which is the instance specific version of `tf_synchronize()`.

Example 18-5 modifies the combinational logic interface presented in Example 18-3. This modified version schedules a *misctf* synchronize callback at the end of a simulation time step in which an input changed value.

This example uses the TF work area to store a flag to indicate when a callback to the *misctf routine* has already been scheduled for the current simulation time step. This flag is used in the C model interface to prevent more than one callback to the *misctf routine* being requested in the same time step. Since the TF work area is unique for each instance of a system task, each instance of the C model will have a unique flag.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.18/sci_alu_synchronized_acc.c
- Verilog shell module: Chapter.18/sci_alu_synchronized_shell.v
- Verilog test bench: Chapter.18/sci_alu_synchronized_test.v
- Verilog-XL results log: Chapter.18/sci_alu_synchronized_test.log

Example 18-5: C model interface synchronized to the end of a time step

```
#include <stdio.h>      /* ANSI C standard I/O library */
#include "veriuser.h"    /* IEEE 1364 PLI TF  routine library */
#include "acc_user.h"    /* IEEE 1364 PLI ACC routine library */

#define ALU_A      1 /* system task arg 1 is ALU A input      */
#define ALU_B      2 /* system task arg 2 is ALU B input      */
#define ALU_OP     3 /* system task arg 3 is ALU opcode input */
#define ALU_RESULT 4 /* system task arg 4 is ALU result output */
#define ALU_EXCEPT 5 /* system task arg 5 is ALU exception output */
#define ALU_ERROR   6 /* system task arg 6 is ALU error output */

/*********************  
 * VCL simulation callback routine: Schedules a callback to the misctf  
 * application synchronized to the end of a time step. Only schedules  
 * one callback for a time step  
 *******************/  
  
int PLIbook_ScientificALU_interface(p_vc_record vc_record)  
{  
    char *instance_p;  
  
    /* Retrieve instance handle from VCL user_data field */  
    instance_p = vc_record->user_data;  
  
    /* If the TF work area for this instance is NULL, then no misctf */  
    /* synchronize callback has been scheduled for this time step (the */  
    /* work area is set to non-null by this routine, and is set to */  
    /* NULL by the misctf after a callback is processed. */
```

```
if (tf_igetworkarea(instance_p) == NULL) {
    /* Schedule a synchronize callback to misctf for this instance */
    tf_isynchronize(instance_p);
    tf_isetworkarea("1", instance_p); /* set work area to non-null */
}

return(0);
}

/*********************misctf routine: Serves as an interface between Verilog simulation
* and the C model. Called by the VCL consumer application whenever
* the C model inputs change value, reads the values of all inputs,
* passes the values to the C model, and writes the C model outputs
* into simulation.
******************/int PLIBook_ScientificALU_misctf(int user_data, int reason)
{
    double          a, b, result;
    int             opcode, excep, err;
    handle          instance_h, result_h, excep_h, err_h,
                    a_h, b_h, opcode_h;
    s_setval_value value_s;
    s_setval_delay delay_s;
    s_acc_time     time_s;

    if (reason != REASON_SYNCH)
        return(0); /* abort misctf if not called for synchronize reason */

    acc_initialize();

    /* Set the TF work area for this instance to null (a flag that
     * this callback has been processed)
     */
    tf_setworkarea(null);

    /* Obtain handles to all task arguments */
    a_h      = acc_handle_tfarg(ALU_A);
    b_h      = acc_handle_tfarg(ALU_B);
    opcode_h = acc_handle_tfarg(ALU_OP);
    result_h = acc_handle_tfarg(ALU_RESULT);
    excep_h = acc_handle_tfarg(ALU_EXCEPT);
    err_h    = acc_handle_tfarg(ALU_ERROR);

    /* Read current values of C model inputs from Verilog simulation */
    value_s.format = accRealVal;
    acc_fetch_value(a_h, "%", &value_s);
    a = value_s.value.real;

    acc_fetch_value(b_h, "%", &value_s);
    b = value_s.value.real;

    value_s.format = accIntVal;
    acc_fetch_value(opcode_h, "%", &value_s);
    opcode = value_s.value.integer;
```

```
***** Call C model *****/
PLIbook_ScientificALU_C_model(a, b, opcode, &result, &excep, &err);

/* Write the C model outputs onto the Verilog signals */
delay_s.model      = accNoDelay;
delay_s.time       = time_s;
delay_s.time.type = accRealTime;
delay_s.time.real  = 0.0;

value_s.format     = accRealVal;
value_s.value.real = result;
acc_set_value(result_h, &value_s, &delay_s);

value_s.format     = accIntVal;
value_s.value.integer = excep;
acc_set_value(excep_h, &value_s, &delay_s);

value_s.value.integer = err;
acc_set_value(err_h, &value_s, &delay_s);

acc_close();
return(0);
}

*****
* calltf routine: Registers a callback to the C model interface
* whenever any input to the C model changes value
*****
int PLIbook_ScientificALU_calltf()
{
    handle a_h, b_h, opcode_h;
    char *instance_p;

    acc_initialize();

    /* get handles for signals in task args which are C model inputs */
    a_h      = acc_handle_tfarg(ALU_A);
    b_h      = acc_handle_tfarg(ALU_B);
    opcode_h = acc_handle_tfarg(ALU_OP);

    /* get pointer for this system task instance to pass to VCL app. */
    instance_p = tf_getinstance();

    /* set the TF work area for this instance to null */
    tf_setworkarea(NULL);

    /* add VCL flags to all signals which are inputs to the C model */
    /* pass handle for task instance as the user_data value */
    acc_vcl_add(a_h, PLIbook_ScientificALU_interface,
                instance_p, vcl_verilog_logic);
    acc_vcl_add(b_h, PLIbook_ScientificALU_interface,
                instance_p, vcl_verilog_logic);
```

```
    acc_vcl_add(opcode_h, PLIbook_ScientificALU_interface,  
                instance_p, vcl_verilog_logic);  
  
    acc_close();  
    return(0);  
}
```

18.8 Synchronizing with a future simulation time step

In certain C model applications, it may be necessary to synchronize C model activity with future simulation activity. The `tf_setdelay()` routine and variations of this routine from the TF library can be used to schedule a call to the *miscf routine* for a specific amount of time in the future, relative to the current simulation time.

The `tf_getnextlongtime()` routine returns the future simulation time in which the next simulation event is scheduled to occur. This provides a way for a PLI application to synchronize activity for when the Verilog simulator is processing simulation events.

These TF routines for synchronizing with future simulation times were presented in more detail in Chapter 12.

18.9 Allocating storage within a C model

Special attention and care must be taken when a C model uses static variables or allocates memory.

The Verilog language can instantiate a model any number of times. Each instance of the Verilog shell module creates a unique instance of the system task which invokes the PLI interface to the C model. Therefore, the *callif routine* which is invoked by a system task instance, the handle for the task instance, and the handles for the arguments for the task instance will all be unique for each system task instance. Any memory which is allocated by the *callif routine* will also be unique for each instance of the system task.

When a C model is represented as an independent program, multiple instances of the model are not a problem, as each instance will invoke a new process with unique storage for each process. When the C model is represented as a C function, however, mul-

multiple instances of the model will share the same function. The model must allow for the possibility of multiple instances, and provide unique storage for each instance.

In addition, the VCL *consumer routine* is not unique to each instance of a system task. Any static variables or memory allocated by a *consumer routine* will be common to all instances of the C model. The *consumer routine* must therefore devise a scheme to allocate unique memory storage for each instance of a system task which can call the routine.

One simple way to create unique storage for each instance is to pass the handle for the system task instance to the *consumer routine* using the *user_data* field of the VCL callback. The *consumer routine* can then allocate memory for each task instance which calls it, and use the instance handle to associate the memory block with a specific task instance.

Example 18-6 presents a latched version of the scientific ALU, which can store the result of a previous operation indefinitely, and Example 18-7 presents a combinational logic interface to the model.

This example allocates unique storage within the C model for each instance of the C model. The instance handle of the system task which represents the C model interface is used to identify which storage area belongs to which instance of the model. This instance handle is obtained by the *calltf routine* for a system task instance, passed to the Value Change Link *consumer routine* through the *user_data* field, and then passed to the C model as an input to the model function.



The source code for this example is on the CD accompanying this book.

- Application source file: Chapter.18/sci_alu_latched_acc.c
- Verilog shell module: Chapter.18/sci_alu_latched_shell.v
- Verilog test bench: Chapter.18/sci_alu_latched_test.v
- Verilog-XL results log: Chapter.18/sci_alu_latched_test.log

Example 18-6: scientific ALU C model with latched outputs

```
*****
 * Definition for a structure to store output values when the ALU is
 * latched. When enable is 1, the ALU returns the currently calculated
 * outputs, and when 0, the ALU returns the latched previous results.
*****
#include <stdio.h>
typedef struct PLIbook_SciALUoutputs *PLIbook_SciALUoutputs_p;
typedef struct PLIbook_SciALUoutputs {
    char *instance_p; /* shows which task instance owns this space */
    double result;    /* stored result of previous operation */
}
```

```
int      excep;
int      err;
PLIbook_SciALUoutputs_p next_ALU_outputs; /* next stack location */
} PLIbook_SciALUoutputs_s;

/* declare global stack pointer */
static PLIbook_SciALUoutputs_p ALU_outputs_stack = NULL;

*****  
* C model of a Scientific Arithmetic Logic Unit.  
* Latched outputs version.  
*****  
#include <math.h>
#include <ERRNO.h>
void PLIbook_ScientificALU_C_model(
    int      enable,      /* input; 0 = latched */
    double   a,           /* input */
    double   b,           /* input */
    int      opcode,      /* input */
    double *result,      /* output from ALU */
    int     *excep,       /* output; set if result is out of range */
    int     *err,          /* output; set if input is out of range */
    char    *instance_p) /* input; pointer to system task instance */
{
    PLIbook_SciALUoutputs_p ALU_outputs;

    /* Locate the output storage in the stack for this model instance */
    /* If no storage is found, then allocate a storage block and add */
    /* the storage to the stack. */
    ALU_outputs = ALU_outputs_stack; /* top-of-stack is in global var. */
    while (ALU_outputs && (ALU_outputs->instance_p != instance_p))
        ALU_outputs = ALU_outputs->next_ALU_outputs;

    /* If no storage area found for this model instance, create one */
    if (ALU_outputs == NULL) {
        ALU_outputs =
            (PLIbook_SciALUoutputs_p)malloc(sizeof(PLIbook_SciALUoutputs_s));
        ALU_outputs->instance_p = instance_p; /* set owner of this space */
        ALU_outputs->next_ALU_outputs = NULL;
        ALU_outputs_stack = ALU_outputs; /* save new top-of-stack */
    }

    if (enable) { /* ALU is not latched, calculate outputs and store */
        switch (opcode) {
            case 0x0: ALU_outputs->result = pow      (a, b);      break;
            case 0x1: ALU_outputs->result = sqrt     (a);         break;
            case 0x2: ALU_outputs->result = exp      (a);         break;
            case 0x3: ALU_outputs->result = ldexp   (a, (int)b); break;
            case 0x4: ALU_outputs->result = fabs     (a);         break;
            case 0x5: ALU_outputs->result = fmod     (a, b);      break;
            case 0x6: ALU_outputs->result = ceil     (a);         break;
            case 0x7: ALU_outputs->result = floor     (a);        break;
            case 0x8: ALU_outputs->result = log      (a);         break;
        }
    }
}
```

```

    case 0x9: ALU_outputs->result = log10 (a);      break;
    case 0xA: ALU_outputs->result = sin   (a);      break;
    case 0xB: ALU_outputs->result = cos   (a);      break;
    case 0xC: ALU_outputs->result = tan   (a);      break;
    case 0xD: ALU_outputs->result = asin  (a);      break;
    case 0xE: ALU_outputs->result = acos  (a);      break;
    case 0xF: ALU_outputs->result = atan  (a);      break;
}
ALU_outputs->err  = (errno == EDOM); /* arg out of range */
ALU_outputs->excep = (errno == ERANGE); /* result out of range */
errno = 0;                                /* clear the error flag */
if (ALU_outputs->err) ALU_outputs->result = 0.0;
}

/* return the values stored in the C model */
*result = ALU_outputs->result;
*err     = ALU_outputs->err;
*excep   = ALU_outputs->excep;

return;
}

```

The C source code listing in Example 18-7 illustrates a combinational logic interface for the latched scientific ALU model.

Example 18-7: combinational logic interface to latched scientific ALU C model

```

#include "veriuser.h" /* IEEE 1364 PLI TF routine library */
#include "acc_user.h" /* IEEE 1364 PLI ACC routine library */

#define ALU_ENABLE 1 /* system task arg 1 is ALU enable input */
#define ALU_A      2 /* system task arg 2 is ALU A input */
#define ALU_B      3 /* system task arg 3 is ALU B input */
#define ALU_OP     4 /* system task arg 4 is ALU opcode input */
#define ALU_RESULT 5 /* system task arg 5 is ALU result output */
#define ALU_EXCEPT 6 /* system task arg 6 is ALU exception output */
#define ALU_ERROR  7 /* system task arg 7 is ALU error output */

*****  

* VCL simulation callback routine: Serves as an interface between  

* Verilog simulation and the C model. Called whenever the C model  

* inputs change value, passes the values to the C model, and puts  

* the C model outputs into simulation.  

*****  

int PLIbook_ScientificALU_interface(p_vc_record vc_record)
{
    double          a, b, result;
    int            opcode, excep, err, enable;
    handle         instance_h, result_h, excep_h, err_h,
    a_h, b_h, opcode_h, enable_h;

```

```
s_setval_value  value_s;
s_setval_delay  delay_s;
s_acc_time      time_s;

acc_initialize();

/* Retrieve instance handle from VCL user_data field */
instance_h = (handle)vc_record->user_data;

/* Obtain handles to all task arguments */
enable_h = acc_handle_itfarg(ALU_ENABLE, instance_h);
a_h      = acc_handle_itfarg(ALU_A,           instance_h);
b_h      = acc_handle_itfarg(ALU_B,           instance_h);
opcode_h = acc_handle_itfarg(ALU_OP,          instance_h);
result_h = acc_handle_itfarg(ALU_RESULT,       instance_h);
excep_h  = acc_handle_itfarg(ALU_EXCEPT,      instance_h);
err_h    = acc_handle_itfarg(ALU_ERROR,        instance_h);

/* Read current values of C model inputs from Verilog simulation */
value_s.format = accRealVal;
acc_fetch_value(a_h, "%%", &value_s);
a = value_s.value.real;

acc_fetch_value(b_h, "%%", &value_s);
b = value_s.value.real;

value_s.format = accIntVal;
acc_fetch_value(opcode_h, "%%", &value_s);
opcode = value_s.value.integer;

acc_fetch_value(enable_h, "%%", &value_s);
enable = value_s.value.integer;

***** Call C model *****/
PLIbook_ScientificALU_C_model(enable, a, b, opcode,
                               &result, &excep, &err,
                               (char *)instance_h);

/* Write the C model outputs onto the Verilog signals */
delay_s.model     = accNoDelay;
delay_s.time      = time_s;
delay_s.time.type = accRealTime;
delay_s.time.real = 0.0;

value_s.format     = accRealVal;
value_s.value.real = result;
acc_set_value(result_h, &value_s, &delay_s);

value_s.format     = accIntVal;
value_s.value.integer = excep;
acc_set_value(excep_h, &value_s, &delay_s);

value_s.value.integer = err;
acc_set_value(err_h, &value_s, &delay_s);
```

```

    acc_close();
    return(0);
}

/*****
 * calltf routine: Registers a callback to the C model interface
 * whenever any input to the C model changes value
 *****/
int PLIbook_ScientificALU_calltf()
{
    handle instance_h, enable_h, a_h, b_h, opcode_h;

    acc_initialize();

    /* get handles for signals in task args which are C model inputs */
    enable_h = acc_handle_tfarg(ALU_ENABLE);
    a_h      = acc_handle_tfarg(ALU_A);
    b_h      = acc_handle_tfarg(ALU_B);
    opcode_h = acc_handle_tfarg(ALU_OP);

    /* get handles for this system task instance to pass to VCL app. */
    instance_h = acc_handle_tfinst();

    /* add VCL flags to all signals which are inputs to the C model */
    /* pass handle for task instance as the user_data value */
    acc_vcl_add(enable_h, PLIbook_ScientificALU_interface,
                (char*)instance_h, vcl_verilog_logic);
    acc_vcl_add(a_h, PLIbook_ScientificALU_interface,
                (char*)instance_h, vcl_verilog_logic);
    acc_vcl_add(b_h, PLIbook_ScientificALU_interface,
                (char*)instance_h, vcl_verilog_logic);
    acc_vcl_add(opcode_h, PLIbook_ScientificALU_interface,
                (char*)instance_h, vcl_verilog_logic);

    acc_close();
    return(0);
}

```

18.10 Representing propagation delays in a C model

Propagation delays from an input change to an output change in a C model can be represented in two ways:

- Using delays in the PLI interface.
- Using delays in the Verilog shell module.

Delays in the PLI interface are represented by specifying a delay value with the `acc_set_value()` routine which writes values onto the system task arguments. Either inertial or transport event propagation can be used, depending on the requirements of the C model. However, using `acc_set_value()` does not offer a great deal of flexibility on creating delays which are different for each instance of a model. Nor can the `acc_set_value()` routine represent minimum, typical and maximum delays, different delays for rise and fall transitions, or annotation of delays using delay calculators or SDF files.

C model propagation delays can also be represented using the pin-to-pin path delays in the Verilog shell module. This method provides the greatest amount of flexibility and accuracy in modeling propagation delays. All path delays constructs can be used, as well and Verilog timing constraints.



Some Verilog simulators restrict the use of pin-to-pin path delays and SDF delay back annotation to Verilog models which are represented with Verilog primitives and net data types. To use path delays on a C model with these simulators, buffers must be added to all input and output ports, with net data types connected to the inputs and outputs of these buffers.

In Chapter 13, example 13-9 on page 442 illustrates adding pin-to-pin path delays to the scientific ALU shell module.

18.11 Summary

This chapter has presented just a few ways in which the ACC library can be used to interface a C language model with Verilog simulations. The Value Change Link routines in the ACC library provide an efficient means to pass input changes to a C model. The ACC routines to read and modify logic values allow information to be exchanged in a variety of formats. By creating a shell module which contains the system task that invokes the C model interface, the C model can be used in a Verilog design just as any other Verilog module.

Appendices

These appendices presents the complete IEEE 1364-1995 PLI libraries for the VPI, TF and ACC libraries. Each library routine is listed with its return value type, the input value types, and a brief description of the routines purpose. Object data diagrams are provided for the VPI and ACC libraries to show each Verilog object which can be accessed by those routines. The diagrams also show the properties of those objects and the relationships of each object to other objects. The names of all C constants and structure definitions used with the library are listed with the routine descriptions or in the data diagrams. Also included, is a description of how PLI applications are linked into several major Verilog simulators.

APPENDIX A

Linking PLI Applications to Verilog Simulators

PLI applications must be linked into Verilog simulators. There are many simulators available, and each simulator has a unique method of linking PLI applications. This appendix presents how PLI applications are linked into a few of the simulators at the time this book was written. The topics presented include:

- The IEEE 1364 interface mechanism for VPI based applications
- The IEEE 1364 interface mechanism for TF and ACC based applications
- Linking – PLI applications to the following simulators:
 - **Verilog-XL™** from Cadence Design Systems, Inc.
 - **NC-Verilog™** from Cadence Design Systems, Inc.
 - **ModelSim™** from Model Technology, Inc.
 - **Polaris™** from Avant! Corporation
 - **Silos III™** from Simucad®, Inc.
 - **VCSTM** from Synopsys®, Inc.
 - **VeriBest™** from VeriBest, Inc.

With two exceptions, the simulators are listed alphabetically by product name. The exceptions are that the Cadence Verilog-XL and NC-Verilog simulators are listed first, simply because Verilog-XL is the original Verilog simulator and many companies have copied some of all of Verilog-XL's method of interfacing PLI applications.

Company names and product names listed in this appendix are trademarks or registered trademarks of the respective company with which they are associated. Other names may also be trademarks or registered trademarks of their respective companies.

A.1 The PLI Interface Mechanism

A Verilog PLI application comprises of:

- A system task or system function name
- A set of C functions for:
 - A *calltf routine*
 - A *compiletf routine* or *checktf routine*
 - A *sizetf routine*
 - A *misctf routine* (for PLI applications written with the TF/ACC routines).

After the system task/function name and PLI application C functions have been defined, two actions are required:

1. Associate the system task/function name to the various application routines.
The PLI standard provides an *interface mechanism* to make the associations between the system task/function name and the application routines. There are two generations of the interface mechanism, one that was created for the older TF and ACC libraries, and a newer mechanism created for the VPI library.
2. Link the applications into a Verilog simulator, so the simulator can call the appropriate routine when the system task/function name is encountered.

The PLI standard does not provide any guidelines on how PLI applications should be linked into a Verilog simulator. There are many different C compilers and operating systems available to Verilog users, and each compiler and operating system has unique methods for compiling and linking programs.

A.1.1 Interfacing PLI applications using the VPI library

The IEEE 1364 standard defines a *VPI Interface Mechanism* for associating PLI applications that use any of the TF, ACC or VPI routines from the PLI libraries. The VPI interface mechanism involves creating a register function, which associates the system task/function name with the application routines. After the register function is defined, the Verilog simulator must be notified about the registration function.

The VPI Interface Mechanism is defined as part of the IEEE 1364 standard, providing a consistent method for all Verilog simulators to use. The VPI interface specifies:

- A system task/function *name*.
- The application *type*, which is a *task*, *sized function*, *integer function*, *time function* or *real function*.

- Pointers to the C functions for a *calltf routine*, *compiletf routine* and *sizetf routine*, if the routines exist. It is not required—and often not necessary—to provide each class of routine.
- A character pointer *user_data* value, which the simulator will pass to the *calltf routine*, *compiletf routine* and *sizetf routine* each time they are called.

The process of specifying the PLI application information is referred to as *registering* the application. To register a PLI application, the information about the application is specified in an **s_vpi_systf_data** structure. This structure is defined as part of the VPI standard, in the PLI *vpi_user.h* file. The definition is:

```
typedef struct t_vpi_systf_data {
    int type;
    int subtype;
    char *tfname;
    int (*calltf)();
    int (*compiletf)();
    int (*sizetf)();
    char *user_data;
} s_vpi_systf_data, *p_vpi_systf_data;
```

Table A-1 explains the fields of the **s_vpi_systf_data** structure:

s_vpi_systf_data Field	Definition
type	Defines the type of a PLI application as being either a system task or a system function. This field must be set to the C constant: vpiSysTask or vpiSysFunc .
subtype	Defines the return type of a system function. This field is only used if the type field is vpiSysFunc , in which case the subtype must be set to the C constant: vpiSysFuncInt , vpiSysFuncReal , vpiSysFuncTime or vpiSysFuncSized .
tfname	Specifies the name of the system task/function; must be a quoted literal string.
calltf	Specifies a pointer to the C function that will be called by a simulator for the application's <i>calltf routine</i> .
compiletf	Specifies a pointer to the C function that will be called by a simulator for the application's <i>compiletf routine</i> .

Table A-1: VPI interface mechanism **s_vpi_systf_data** structure fields

s_vpi_systf_data Field	Definition
sizetf	Specifies a pointer to the C function that will be called by a simulator for the application's <i>sizetf routine</i> .
user_data	Specifies a character pointer—the value of the pointer will be passed to the PLI application routines each time a routine is called.

Table A-1: VPI interface mechanism `s_vpi_systf_data` structure fields

NOTE → The proposed IEEE 1364-1999 standard will change the names of the constants involved with system task/functions:

- `vpiSysFuncInt` changes to `vpiIntFunc`
- `vpiSysFuncTime` changes to `vpiTimeFunc`
- `vpiSysFuncReal` changes to `vpiRealFunc`
- `vpiSysFuncSized` changes to `vpiSizedFunc`

The constant names from the 1364-1995 standard will be aliased to the new constant names, to provide backward compatibility.

The steps required to register a PLI application using the VPI interface

The following steps are used to register a system task or system function using the VPI interface mechanism:

1. Create a C function to register the system task/function. The C function name is application-defined and can be any legal C name.
2. Allocate an `s_vpi_systf_data` C structure.
3. Fill in the fields of the structure with the information about the system task or system function.
4. Register the system task/function by calling the VPI routine `vpi_register_systf()`.
5. Add the name of the C function created in step 1 to a C language array called `vlog_startup_routines`. This array is typically contained in a C source file provided with the simulator, called `vpi_user.c`, though the IEEE 1364 standard does not require that the file be called that name.

The following example registers a PLI application called `$hello`.

Example A-1: VPI register function for the \$hello system function

```
/* prototypes of PLI application routine names */
int PLIbook_hello_calltf(), PLIbook_hello_compiletf();

void PLIbook_hello_register()
{
    s_vpi_systf_data tf_data;

    tf_data.type      = vpiSysTask;
    tf_data.tfname    = "$hello";
    tf_data.calltf   = PLIbook_hello_calltf;
    tf_data.compiletf = PLIbook_hello_compiletf;
    tf_data.sizetf    = NULL;
    vpi_register_systf(&tf_data);
}
```

Notifying Verilog simulators about the VPI register functions

Once the register function has been defined, a Verilog simulator must be notified of the name of the register function, so that the simulator can call the functions and register the PLI applications. The VPI standard requires that all Verilog simulators provide a special array, called *vlog_startup_routines*, in order to notify the simulators about the register functions. All PLI applications which will be called by the simulator should have a register function listed in this array.

An example *vlog_startup_routines* array is listed below, with an entry for the \$hello PLI application register function.

Example A-2: sample *vlog_startup_routines* array

```
/* prototypes of the PLI application register routines */
extern void PLIbook_hello_register();

void (*vlog_startup_routines[])() =
{
    /*** add user entries here ***/
    PLIbook_hello_register,
    0 /*** final entry must be 0 ***/
};
```



NOTE The IEEE 1364 standard does not define where the *vlog_startup_routines* array should be located. Consult the reference manual of the simulator for the location of the start-up array used by that simulator.

**TIP**

Do not place the `vlog_startup_routines` array in the same file as the PLI application! The `vlog_startup_routines` is a global array, and the C language does not permit multiple global arrays with the same name. In a typical design environment, PLI applications will come from several sources, such as internally developed applications and 3rd party applications. If the `vlog_startup_routines` array and a PLI application are in the same file, then the source code for the application must be available whenever another PLI application needs to be added to the start-up array. If two PLI applications include a `vlog_startup_routines` array in the application, then the object files for both applications could not be used together.

A.1.2 Interfacing PLI applications using the TF/ACC libraries

The TF/ACC interface mechanism is derived from the 1990 OVI PLI 1.0 standard. This older interface mechanism defines what all Verilog simulators should provide for interfacing PLI applications to a simulator, but does not define how the interface should be implemented.

The TF/ACC interface mechanism is used to specify:

- A system task/function ***name***.
- The application ***type***, which is a *task*, *function* or *real function* (the IEEE 1364-1995 standard omitted the description of real functions from the body of the standard. This errata is corrected in the proposed IEEE 1364-1999 standard).
- ***Pointers*** to the C functions for the *calltf routine*, *checktf routine*, *sizetf routine* and *misctf routine*, if the routines exist. It is not required—and often not necessary—to provide each type of routine).
- An integer ***user_data*** value, which the simulator will pass to the *calltf routine*, *checktf routine*, *sizetf routine* and *misctf routine* each time they are called.

The defacto standard TF/ACC veriusertfs array

Many Verilog simulators have adopted a similar method of specifying the information about a PLI application. This method uses a C array called ***veriusertfs*** (which stands for *Verilog user tasks and functions*). The *veriusertfs* array is the original method implemented in the Cadence Verilog-XL product when the PLI was first introduced in the mid 1980's.

NOTE

The *veriusertfs* array is not specified in the IEEE 1364 standard. Though many simulators use this array, it is not required, and some simulators use different methods of specifying the PLI application information for the TF/ACC interface.

The typical *veriusertfs* array definition is:

```
s_tfcell veriusertfs[] =
{
    { type, user_data, checktf_app, sizetf_app,
      calltf_app, misctf_app, "tf_name", 1, 0, 0 },
    { type, user_data, checktf_app, sizetf_app,
      calltf_app, misctf_app, "tf_name", 1, 0, 0 },
    ...
    {0}, /* first field in final array cell is 0 */
};
```

Each array cell is an **s_tfcell** structure, which contains several fields. These fields specify the information about a PLI application. There can be any number of cells in the array. A cell with the first field set to 0 is used to denote the last cell in the array. The **s_tfcell** structure is not defined in the IEEE standard. However, nearly all simulators use the same structure definition, which is:

```
typedef struct t_tfcell {
    short type;           /* one of the constants: usertask,
                           userfunction, userrealfunction */
    short data;           /* data passed to user routine */
    int (*checktf)();     /* pointer to the checktf routine */
    int (*sizetf)();      /* pointer to the sizetf routine */
    int (*calltf)();      /* pointer to the calltf routine */
    int (*misctf)();      /* pointer to the misctf routine */
    char *tfname;         /* name of the system task/function */
    int forwref;          /* usually set to 1 */
    char *tfveritool;     /* usually ignored */
    char *tferrrmassage; /* usually ignored */
} s_tfcell, *p_tfcell;
```

The meaning of each field in the **s_tfcell** structure is explained in Table A-2.

veriusertfs Field	Definition
type	Defines the type of a PLI application as being either a system task or a system function. This field must be set to the C constant usertask , userfunction or userrealfunction (these constants are defined in the veriuser.c file).
user_data	Specifies an integer value—the value will be passed to the PLI application routines each time a routine is called.
checktf_app	Specifies a pointer to the C function that should be called by a simulator for the application's <i>checktf routine</i> .
sizetcf_app	Specifies a pointer to the C function that should be called by a simulator for the application's <i>sizetcf routine</i> .
calltf_app	Specifies a pointer to the C function that should be called by a simulator for the application's <i>calltf routine</i> .
misctf_app	Specifies a pointer to the C function that should be called by a simulator for the application's <i>misctf routine</i> .
"tf_name"	Specifies the name of the system task/function; must be a quoted literal string beginning with a \$.
forwref	Specifies instance name forward referencing as true (1) or false (0). This field is ignored by most Verilog simulators. In the Cadence Verilog-XL simulator, setting to 1 makes Verilog-XL IEEE compliant by allowing module and primitive instance names to be used as system task/function arguments (a 0 makes instance names illegal in a system task/function argument).
tfveritool	Ignored by most Verilog simulators.
tferrmessage	Ignored by most Verilog simulators.

Table A-2: Typical fields in the defacto standard veriusertfs array

The *veriusertfs array* can specify any number of PLI applications. A sample *veriusertfs array* with information about a PLI application is listed in example A-3.

NOTE → *Do not specify the veriusertfs array in the same file as the PLI application!*

- Not all Verilog simulators use the *veriusertfs array* to specify PLI application information.
- The array is not standardized, and may be different in different simulators.
- The C language does not allow multiple global arrays with the same name. If two applications both contained a *veriusertfs array* definition, the applications could not be used together.

Example A-3 illustrates an example veriusertfs array with the information required for the \$hello PLI application.

Example A-3: sample veriusertfs array, as used by many Verilog simulators

```
/* prototypes of the PLI application routines */
extern int PLIbook_hello_calltf(), PLIbook_hello_checktf();

/* the veriusertfs table */
s_tfcell veriusertfs[] =
{
    {usertask,                      /* type of PLI routine */
     0,                            /* user_data value */
     PLIbook_hello_checktf,        /* checktf routine */
     0,                            /* sizetf routine */
     PLIbook_hello_calltf,         /* calltf routine */
     0,                            /* miscf routine */
     "$hello",                     /* system task/function name */
     1                            /* forward reference = true */
},
{0} /*** final entry must be 0 ***/
};
```

A.1.3 Should PLI routines be **int** or **void** functions?

The IEEE 1364 standard defines that all PLI routines, such as *calltf routines*, should be of type **int**. However, the simulator only uses the return value from the *sizetf routine*. The return value from the *calltf routine*, *checktf routine* and *miscf routine* is ignored. Since the return value from these routines is ignored, it is a common practice to declare these routines as **void** functions. Declaring the functions as a different type from the PLI standard prototype may result in a C compiler warning message when the routine is compiled. To prevent this warning message, the pointer to the **void** function can be cast to a pointer to an **int** function. For example:

```
(int(*)())PLIbook_Pow_calltf
```

Note, however, that some C compilers might not permit casting a function pointer from one type to another.

A.1.4 C versus C++

Most Verilog simulators expect PLI applications to be written in ANSI C. The PLI standard includes a library of C functions as well as definitions of C constants, special data types and structures. The PLI library is compliant with the ANSI C standard.

Many Verilog simulators were written in the C language, and may or may not support linking C++ applications to the simulator. PLI application developers who wish to work with C++ should first check the limitations of the simulator products to which the applications will be linked. For maximum portability, it is recommended that PLI applications be written in ANSI C.

A.1.5 Static linking versus dynamic linking

There are two general methods in which a PLI application can be linked into a Verilog simulator: static linking and dynamic linking.

Many Verilog simulators statically link PLI applications into the simulator. Static linking requires:

1. The simulator vendor must provide all of the object files for the simulator.
2. The end-user must create a new simulator executable by relinking the simulator's object files with the compiled PLI application object files.

Static linking is simple, but does have some drawbacks. The relinking process often depends on operating system graphical libraries being present (such as the X-11 libraries on some Unix systems). If a required operating system object file is missing, is an incorrect version, or is in a different location than expected, the Verilog simulator will not link properly.

Some Verilog simulators dynamically link PLI applications into the simulator. With this method:

1. Object files are compiled separately as position independent code, to form object libraries (or dll's in Microsoft Windows).
2. The standard simulator executable dynamically loads the separately compiled PLI applications when needed.

With dynamic linking, the simulator vendor does not need to provide the object files for the Verilog simulator. A major advantage of this is that the user does not need to create new, customized simulator executables, and does not need to worry about having all the correct operating system object files. A disadvantage of dynamic linking is a dependency on operating system and compiler versions.

A.2 Linking PLI app's to the Cadence Verilog-XL simulator

The Cadence Verilog-XL™ simulator is the original Verilog simulator, and was first introduced in 1985. At the time this book was written, Verilog-XL ran on all major Unix workstations and on PC's with the Windows NT operating system.

A.2.1 IEEE 1364-1995 compliance

The release of Verilog-XL used in this book is version 2.6 (also called release 97B). This release supports the full IEEE 1364-1995 PLI standard (the TF, ACC and VPI routines).

Verilog-XL uses the standard IEEE 1364 *vpi_user.h*, *veriuser.h* and *acc_user.h* header files, which define the VPI, TF and ACC PLI libraries. These header files are included in the installation directory of Verilog-XL.

NOTE Verilog-XL also provides a number of extensions to the PLI standard, which support proprietary features of the Verilog-XL simulator. The extensions to the PLI standard are specified in three proprietary header files called *vpi_user_cds.h*, *vxl_veriuser.h* and *vxl_acc_user.h*. It is not necessary to include these header files in a PLI application, unless the proprietary extensions to the PLI are being used.

A.2.2 Specifying PLI application information

Verilog-XL has two PLI interface mechanisms:

- A VPI interface mechanism
- A TF/ACC interface mechanism

Sections A.2.3 and A.2.4, which follow, present first the VPI interface mechanism and then the older TF/ACC interface mechanism.

The Verilog-XL simulator also offers the user a choice of either statically or dynamically linking PLI applications into the simulator. Section A.2.5 presents how to dynamically link PLI applications on Unix systems and section A.2.6 presents the dynamic linking instructions for Windows NT operating systems. Statically linking applications in Verilog-XL uses a similar procedure as dynamically linking, but is only available with Unix operating systems. This book does not present the static linking method.

A.2.3 Interfacing VPI applications with Verilog-XL

The VPI interface mechanism defines a standard method that all simulators should use. This method involves creating a application-defined register function, as discussed earlier in section A.1.1 on page 632. After the register function is defined, it must be listed in a *vlog_startup_routines* array.

For Verilog-XL, the *vlog_startup_routines* array is contained in a file called *vpi_user.c*. The startup array can contain any number of register functions.

To specify a PLI application using the VPI interface mechanism in Verilog-XL:

1. Copy a C source file called *vpi_user.c*, that is provided with Verilog-XL.
2. In the copy of the *vpi_user.c* file, edit the *vlog_startup_routines* array to add the information for each VPI register function to the array.

The *vpi_user.c* file for Verilog-XL is located in the installation directory of the simulator software. On both Unix and NT systems, the typical location is:

<cadence_install_directory>/tools/verilog/src/vpi_user.c

NOTE → *Do not modify the original vpi_user.c file!* Make a copy of the file, and edit the copy. If the original file were to become corrupted, it could require re-installing the Cadence simulator software.

An example *vpi_user.c* file is listed below, with entries for a register function for a *\$hello* PLI application. The lines that were added for this example are listed in bold.

Example A-4: Sample Verilog-XL *vpi_user.c* file (partial listing)

```
#include "vpi_user.h"
#include "vpi_user_cds.h"

/* prototypes of the PLI application routines */
extern void PLIbook_ShowValue_register();
extern void PLIbook_pow_register();

void (*vlog_startup_routines[])() =
{
    /*** add user entries here ***/
    PLIbook_ShowValue_register,
    PLIbook_pow_register,
    0 /* *** final entry must be 0 ***/
};
```

After the copy of vpi_user.c has been modified, the PLI applications and the vpi_user.c file need to be compiled and linked. Refer to section A.2.5 on page 644 for details on compiling and linking with Verilog-XL.

A.2.4 Interfacing TF/ACC applications with Verilog-XL

Verilog-XL includes an older interface mechanism that does not use any VPI routines. This older interface mechanism uses the same *veriusertfs* array as most other Verilog simulators to specify PLI application information. Refer back to Section A.1.2 on page 636 for a description of the *veriusertfs* array. The *veriusertfs* array, along with other information used by Verilog-XL, is specified in a file called *veriuser.c*. To use this older Verilog-XL PLI interface:

1. Copy a C source file called *veriuser.c*, that is provided with Verilog-XL.
2. In the copy of the *veriuser.c* file, edit the *veriusertfs* array to add the information for each PLI application to the array.

The *veriuser.c* file for Verilog-XL is located in the installation directory of the simulator software. On Unix and NT systems, the typical location is:

```
<cadence_install_directory>/tools/verilog/src/veriuser.c
```

NOTE *Do not modify the original veriuser.c file!* Make a copy of the file, and edit the copy. If the original file were to become corrupted, it could require re-installing the Cadence simulator software.

A sample Verilog-XL *veriuser.c* file is listed in example A-5, below, for a *\$show_value* and a *\$pow* PLI application. The information added for this application is shown in bold type.

Example A-5: Sample Verilog-XL *veriuser.c* file (partial listing)

```
#include "veriuser.h"
#include "vxl_veriuser.h"

/* prototypes for the PLI application routines */
extern int PLIbook_ShowVal_checktf(), PLIbook_ShowVal_calltf();
extern int PLIbook_pow_sizetf(), PLIbook_pow_checktf(),
          PLIbook_pow_calltf(), PLIbook_pow_misctf();

/* the veriusertfs table */
s_tfcell veriusertfs[] =
{
    {usertask,                      /* type of PLI routine */
     0,                            /* user_data value */
     PLIbook_ShowVal_checktf,      /* checktf routine */
```

```

        0,                                /* sizetf routine */
        PLIbook_ShowVal_calltf,           /* calltf routine */
        0,                                /* misctf routine */
        "$show_value",                  /* system task/function name */
        1                                 /* forward reference = true */
    },

    {userfunction,
        0,                                /* type of PLI routine */
        /* user_data value */
        PLIbook_pow_checktf,             /* checktf routine */
        PLIbook_pow_sizetf,              /* sizetf routine */
        PLIbook_pow_calltf,              /* calltf routine */
        PLIbook_pow_misctf,              /* misctf routine */
        "$pow",                          /* system task/function name */
        1                                 /* forward reference = true */
    },

    {0} /*** final entry must be 0 ***/
};


```



TIP ***Do not place PLI applications in the veriuser.c file!*** The C programming language does not permit multiple global arrays with the same name, such as the *veriusertfs array*. In a typical design environment, PLI applications will come from several sources, such as internally developed applications, 3rd party applications, and ASIC vendor applications. If the *veriusertfs array* and a PLI application are in the same file, then the source code for the application must be available whenever another PLI application needs to be added to the table. If two PLI applications were to both include the *veriusertfs array* in the application, then the object files for both applications could never be used together.

A.2.5 Compiling and linking PLI applications in Verilog-XL using Unix operating systems

The Unix version of Verilog-XL includes a utility program called *vconfig*. This program is used to configure the options to be linked into the Verilog-XL simulator, such as which waveform display to use.

The *vconfig* program prompts the user with a series of questions, including the choice to dynamically or statically link PLI applications. The steps for dynamic linking and static linking are the same—the only difference is in the answers provided to the *vconfig* program. Only the process of dynamic linking is illustrated in this book.

An example output from running the *vconfig* program for dynamic linking is listed below. The answers required for the VPI example of the *\$show_value* and *\$pow* PLI applications are shown in bold.

Example A-6: Sample Verilog-XL *vconfig* program output

```
*****
VCONFIG VERSION 3.7.3, THE CONFIGURATION PROGRAM
FOR THE Verilog-XL FAMILY OF PRODUCTS
*****
o This program will prompt you for the Verilog-XL configuration
options with which to build your new Verilog-XL executable.

o The default answers to all questions appear inside of '[]'.

o Once VCONFIG completes, the script produced must be run to
build your new Verilog-XL executable.
*****
Please enter the name of the output script.
[cr_vlog] : cr_vlog
*****
Please choose a target.
valid choices:1) Stand Alone,
              2) Backplane,
              3) Verilog Export,
              4) VHDL Import,
              5) Dynamic PLI libraries only,
<Enter 1-5> [1] :5
*****
Pls choose a target.
Valid choices:1) build libpli,
              2) build libvpi,
              3) both
<Enter 1-3> [1] :2
*****
The user template file 'vpi_user.c' must always be
included in the link statement. What is the path
name of this file ?
[/net/cds/tools/verilog/src/vpi_user.c] :PLIBbook_vpi_user.c
List the files one at a time, terminating
the list with a single '.'
-----> show_value_vpi.c
-----> pow_vpi.c
-----> .

*** SUCCESSFUL COMPLETION OF VCONFIG ***
*** EXECUTE THE SCRIPT:cr_vlog TO BUILD:Dynamic VPI library
```

When the *vconfig* program completes, it generates a Unix C-shell script, which by default is called *cr_vlog*. This script contains the C compile and link commands for the PLI application. If dynamic linking was selected, the *cr_vlog* script will generate a C object library file. If static linking was selected, then the *cr_vlog* script will generate a new Verilog-XL executable with the PLI applications linked into the executable.

The *cr_vlog* script that is created by *vconfig* program is executed at a Unix prompt. This script does not require additional input from the user.

 **TIP** The *vconfig* program does not need to be run each time the Verilog-XL simulator is relinked. The primary purpose of *vconfig* is to configure the options for Verilog-XL and generate the *cr_vlog* shell script. Once the *cr_vlog* script has been created, it can be re-run as often as needed, and can be edited to make changes (such as adding or changing the names of PLI application source files).

 **TIP** The *cr_vlog* shell script does not take advantage of the C Make utility. Many PLI applications developers use the information in a *cr_vlog* script to create their own make files.

A.2.6 Compiling and linking PLI applications in Verilog-XL using the Windows NT operating system

Compiling and linking PLI applications on the Windows NT operating system requires four basic steps:

1. Set a system environment variable, **CDS_INST_DIR**. This variable is set to the full path of the installation directory of the Cadence software.
2. Copy a C make file. Cadence provides two make file examples, *libvpi.mak* for compiling applications that use the VPI interface mechanism, and *libpli.mak* for compiling applications that use the TF/ACC interface mechanism. The generic make files provided by Cadence are located in the directory:

```
<cadence_install_directory>/tools/verilog/examples/
```

3. Edit the copy of the make file to:
 - Add the name of the modified *vpi_user.c* file (from section A.2.3) if using the *libvpi.mak* file, or add the name of the modified *veriuser.c* file (from section A.2.4) if using the *libpli.mak* file.
 - Add the names of the C source files containing the PLI applications.

Note: the example make files provided by Cadence are for the Microsoft Visual C++ compiler. If a different C compiler is used, the compile and link options in the make file will need to be modified accordingly.

4. Run the C compiler's make utility. For example, the Microsoft Visual C++ make utility command is:

```
nmake -f libvpi.mak
```

The *libvpi.mak* make file will generate a dynamically linked library file called *libvpi.dll* and the *libpli.mak* make file will generate a dynamically linked library file called *libpli.dll*.

Example A-7 lists a partial Cadence *libvpi.mak* file for compiling with Visual C++. The additions for compiling and linking a *\$show_value* and a *\$pow* PLI application are shown in bold.

Example A-7: Sample Cadence *libvpi.mak* file (partial listing)

```
#  
# SRCS set to the list of sources that comprise your VPI application  
#  
SRCS = show_value_vpi.c \  
      pow_vpi.c \  
      PLibook_vpi_user.c  
#  
  
OBJS = $(SRCS:.c=.obj)  
CFLAGS = -DMSC -DWIN32 -I$(CDS_INST_DIR)/tools/verilog/include \  
        -I$(USER_INCLUDE) -MD -O2  
.c.obj:  
$(CC) $(CFLAGS) -c $<  
  
libvpi.dll: $(OBJS)  
link -dll /out:$@ $(OBJS) $(CDS_INST_DIR)/tools/verilog/lib/verilog.lib
```

A.2.7 Running Verilog-XL with PLI applications

Once the PLI applications have been compiled and linked, whether on Unix or Windows NT, they can be executed by the Verilog simulator.


If the PLI applications are dynamically linked, the simulation is run using the standard Verilog-XL simulator. The simulator will dynamically load the PLI object libraries when simulation is invoked. The standard Verilog-XL executable name is *verilog*. In order for Verilog-XL to load a dynamically linked PLI library file, it must be in the operating system's dynamic linker search path. This search path is operating system dependent. For example on the Sun Solaris operating system, the linker uses the **LD_LIBRARY_PATH** environment variable, and on Windows NT, the linker uses the **PATH** environment variable. Verilog-XL will also search for the dynamically loaded library files in the directory specified with a **+loadvpi** invoca-

tion option for VPI applications and a `+loadpli11` invocation option for TF/ACC applications.

If the PLI applications are statically linked into Verilog-XL, then simulations are run using the newly created simulator executable.

Figure A-1 shows the results of running the Verilog-XL simulator with the `$pow` and `$show_value` PLI applications.

Figure A-1: Verilog-XL simulator using the `$pow` and `$show_value` PLI applications



A.3 *Linking to NC-Verilog from Cadence Design Systems, Inc.*

The Cadence NC-Verilog™ simulator is Cadence's second generation Verilog simulator. NC-Verilog is a "native compiled" code simulator, which compiles Verilog source code into machine-dependent object files. As part of the native compilation, NC-Verilog can highly optimize the object files for a specific operating system. At the time this book was written, NC-Verilog ran on several Unix workstations. Support for PC's with the Windows NT operating system was planned.

A.3.1 IEEE 1364-1995 compliance

The release of NC-Verilog used in this book is version 2.1. This release supports the full IEEE 1364-1995 PLI standard (the TF, ACC and VPI routines).

NC-Verilog uses the standard IEEE 1364 *vpi_user.h*, *veriuser.h* and *acc_user.h* header files, which define the VPI, TF and ACC PLI libraries. These header files are included in the installation directory of NC-Verilog.

NOTE → NC-Verilog also provides a number of extensions to the PLI standard, which support proprietary features of the NC-Verilog simulator. The extensions to the PLI standard are specified in three proprietary header files called *vpi_user_cds.h*, *vxl_veriuser.h* and *vxl_acc_user.h*. It should not necessary to include these header files in a PLI application, unless the proprietary extensions to the PLI are being used.

A.3.2 Specifying PLI application information

NC-Verilog supports two PLI interface mechanisms:

- The IEEE 1364 standard VPI interface mechanism
- The older TF/ACC interface mechanism, from the Cadence Verilog-XL simulator.

Sections A.2.3 and A.2.4, which follow, present first the VPI interface mechanism and then the older TF/ACC interface mechanism.

The NC-Verilog simulator also offers the user a choice of either statically or dynamically linking PLI applications into the simulator. Section A.2.5 presents how to dynamically link PLI applications. Statically linking applications in NC-Verilog uses a similar procedure as dynamically linking. This book does not present the static linking method.

A.3.3 Interfacing VPI applications with NC-Verilog

The VPI interface mechanism defines a standard method that all simulators should use. This method involves creating a application-defined register function, as discussed earlier in section A.1.1 on page 632. After the register function is defined, it must be listed in a *vlog_startup_routines* array.

In NC-Verilog, the *vlog_startup_routines* array is contained in a file called *vpi_user.c*. The startup array can contain any number of register functions.

To specify a PLI application using the VPI interface mechanism in NC-Verilog :

1. Copy a C source file called *vpi_user.c*, that is provided with NC-Verilog.
2. In the copy of the *vpi_user.c* file, edit the *vlog_startup_routines* array to add the information for each VPI register function to the array.

The *vpi_user.c* file for NC-Verilog is located in the installation directory of the simulator software. The typical location is:

```
<cadence_install_directory>/tools/inca/src/vpi_user.c
```

 **NOTE** *Do not modify the original vpi_user.c file!* Make a copy of the file, and edit the copy. If the original file were to somehow become corrupted, it could require re-installing the Cadence simulator software.

NC-Verilog uses the same *vpi_user.c* file as the Cadence Verilog-XL simulator. Refer back to Example A-4 on page 642 for a sample *vpi_user.c* file, with entries for the register functions for a *\$show_value* and a *\$pow* PLI application.

After the copy of *vpi_user.c* has been modified, the PLI applications and the *vpi_user.c* file need to be compiled and linked. Refer to section A.2.5 on page 644 for details on compiling and linking with NC-Verilog.

A.3.4 Interfacing TF/ACC applications with NC-Verilog

NC-Verilog also supports the older interface mechanism from Verilog-XL, which does not use any of the VPI routines from the PLI standard. This older TF/ACC interface mechanism uses the same *veriusertfs* array which most other Verilog simulators use to specify PLI application information. Refer back to Section A.1.2 on page 636 for a description of the *veriusertfs* array. The older NC-Verilog interface mechanism uses a file called *veriuser.c* to specify the interface information.

To specify a PLI application using this older NC-Verilog interface mechanism:

1. Copy a C source file called *veriuser.c*, that is provided with NC-Verilog.
2. In the copy of the *veriuser.c* file, edit the *veriuserifs array* to add the information for each PLI application to the array.

The *veriuser.c* file for NC-Verilog is located in the installation directory of the simulator software. On Unix and NT systems, the typical location is:

```
<cadence_install_directory>/tools/inca/src/veriuser.c
```

 **Do not modify the original veriuser.c file!** Make a copy of the file, and edit the copy. If the original file were to somehow become corrupted, it could require re-installing the Cadence simulator software.

NC-Verilog uses the same *vpi_user.c* file as the Cadence Verilog-XL simulator. Refer back to Example A-4 on page 642 for a sample *vpi_user.c* file, with entries for a *\$show_value* and a *\$pow* PLI application.

After the copy of *vpi_user.c* has been modified, the PLI applications and the *vpi_user.c* file need to be compiled and linked. Refer to section A.2.5 on page 644 for details on compiling and linking with NC-Verilog.

A.3.5 Compiling and linking PLI applications in NC-Verilog using Unix operating systems

NC-Verilog uses the C make utility to compile and link PLI applications. Cadence provides an example make file, which specifies:

- The names of the PLI application object file(s) to be compiled and linked.
- The names of the modified *vpi_user.c* file and/or the modified *veriuser.c* file.
- Whether to create shared object libraries (.so or .dll files), dynamically linked object files, or statically linked files.

There are four basic steps involved in compiling and linking a PLI application with NC-Verilog:

1. Set a system environment variable, **CDS_INST_DIR**. This variable is set to the full path of the installation directory of the Cadence software.
2. Copy the Cadence make file example, called **makefile.nc**. This make file is used for compiling applications that use either the VPI interface mechanism or the TF/ACC interface mechanism. The example make file provided by Cadence is located in the directory:

```
<cadence_install_directory>/tools/inca/examples/
```

3. Edit the copy of the make file to:

- Add the names of the C source files containing the PLI applications.
- Add the name of the modified `vpi_user.c` file (from section A.2.3), or add the name of the modified `veriuser.c` file, or both.

Example A-8, which follows, shows where this information is specified in the make file.

Note: the example make file provided by Cadence is for a specific C compiler, based on the operating system being used. If a different C compiler is used, the compile and link options in the make file will need to be modified accordingly.

4. Run the C compiler's make utility.

The Cadence example make file can create a shared object library, a dynamically linked object file, or a statically linked executable file. This is controlled by a command line argument to the make utility, using the argument `shared_lib`, `dynamic` or `static`.

For example:

```
make -f makefile.nc shared_lib
```

When `shared_lib` is used, the make file will generate dynamically linked object library files, such as `libvpi.so` and `libpli.so`.

Example A-7 lists a partial `makefile.nc` file, with the additions for compiling and linking a PLI application. The additions to compile and link the VPI version of a `$show_value` and a `$pow` PLI application are shown in bold.

Example A-8: Sample Cadence `makfile.nc` file (partial)

```
#  
# Makefile.nc  
#  
...  
#  
# The VPI_USER_C macro should be set to the vpi_user.c file containing  
# any user defined 'C' routines for use with VPI (PLI 2.0).  
#  
VPI_USER_C=./PLIBbook_vpi_user.c  
VPI_USER_O=$(TARGETDIR)/vpi_user.o  
...
```

```
#  
# Add any VPI objects which need to be compiled, and can be deleted  
# using the "clean" target here:  
#  
VPI_OBJECTS= show_value_vpi.o pow_vpi.o  
  
...  
  
#  
# The VERIUSER_C macro should be set to the veriuser.c file containing  
# any user define 'C' routines for use with PLI 1.0.  
#  
VERIUSER_C=$(INCA_DIR)/src/veriuser.c  
VERIUSER_O=$(TARGETDIR)/veriuser.o  
  
...  
  
#  
# Add any PLI objects which need to be compiled, and can be deleted  
# using the "clean" target here:  
#  
PLI_OBJECTS=  
  
...
```

A.3.6 Running NC-Verilog with PLI applications

Once the PLI applications have been compiled and linked, they can be executed by the Verilog simulator.

If the PLI applications were linked as shared object libraries, the simulation is compiled and run using the standard NC-Verilog simulator commands. The simulator will dynamically load the PLI object libraries when needed. In order for NC-Verilog to dynamically load a shared object library file, it must be in the operating system's dynamic linker search path. This search path is operating system dependent. For example on the Sun Solaris operating system, the linker uses the **LD_LIBRARY_PATH** environment variable. NC-Verilog will also search for the dynamically loaded library files in the directory specified with a **+loadvpi** invocation option for VPI applications and a **+loadpli1** invocation option for TF/ACC applications.

If the PLI applications are statically linked into NC-Verilog, then the *makefile.nc* make file will create new executable files for the NC-Verilog elaborator NC-Verilog simulator. The new elaborator is then used in place of the Cadence ncelab elaborator when the Verilog models are compiled. The new simulator is used in place of the Cadence ncsim simulator when the Verilog simulation is run.

A.4 Linking PLI applications to the Model Technology *ModelSim* simulator

The **ModelSim™** simulator from Model Technology, Inc. (a wholly owned subsidiary of Mentor Graphics, Inc.) is an easy to use Verilog simulator with a well integrated user interface. ModelSim runs on most Unix operating systems, on the Windows NT operating system and on the Windows 95 operating system.

ModelSim PE version 4.7h was used to test many of the examples in this book, running on a Pentium II laptop with the Windows-95 operating system.

 **NOTE** The procedure for linking PLI applications to the ModelSim simulator is identical for all operating systems. The only operating system specific issues are the C compiler commands for compiling and linking the PLI application source code.

A.4.1 IEEE 1364-1995 compliance

Modelsim PE version 4.7h supports the TF and ACC libraries of the IEEE 1364 Verilog PLI standard, but does *not* support the VPI library of the PLI standard.

ModelSim uses the standard IEEE 1364 *veriuser.h* and *acc_user.h* library files, which define the TF and ACC PLI libraries. These header files are included in the installation directory of ModelSim.

 **NOTE** ModelSim also provides a number of extensions to the PLI standard, primarily to support access to VHDL models, which ModelSim allows to be instantiated within a Verilog design structure. The extensions to the PLI standard are specified in a proprietary header file called *acc_vhdl.h*.

A.4.2 Specifying PLI application information

ModelSim uses the same *veriusertfs array* as most other Verilog simulators to specify PLI application information. Refer back to section A.1.2 on page 636 for a description of the *veriusertfs array*. Unlike other simulators, however, ModelSim does not provide an example file with an empty array pre-declared.

To specify the PLI application for ModelSim:

1. Create a new C source file, which can be any name
2. In the file, define a *veriusertfs array*.

3. Add the information for each PLI application to the array.
4. Add the following C function after the veriusertfs array. This function will be called by the ModelSim simulator to read the entries in the *veriusertfs* array.

```
void init_usertfs()
{
    p_tfcell usertf;

    for (usertf = veriusertfs; usertf; usertf++) {
        if (usertf->type == 0)
            return;
        mti_RegisterUserTF(usertf);
    }
}
```

Example A-9 lists a sample file for specifying the *\$pow* and *\$show_value* PLI applications.

Example A-9: Sample file to specify PLI applications for the ModelSim simulator

```
#include "veriuser.h"

/* prototypes of the PLI application routines */
extern int PLIbook_ShowVal_checktf(), PLIbook_ShowVal_calltf();
extern int PLIbook_pow_sizetf(), PLIbook_pow_checktf(),
          PLIbook_pow_calltf(), PLIbook_pow_misctf();

/* the veriusertfs table */
s_tfcell veriusertfs[] =
{
    {usertask,                      /* type of PLI routine */
     0,                            /* user_data value */
     PLIbook_ShowVal_checktf,      /* checktf routine */
     0,                            /* sizetf routine */
     PLIbook_ShowVal_calltf,       /* calltf routine */
     0,                            /* misctf routine */
     "$show_value",               /* system task/function name */
     1                            /* forward reference = true */
    },
    {userfunction,                  /* type of PLI routine */
     0,                            /* user_data value */
     PLIbook_pow_checktf,         /* checktf routine */
     PLIbook_pow_sizetf,          /* sizetf routine */
     PLIbook_pow_calltf,          /* calltf routine */
     PLIbook_pow_misctf,          /* misctf routine */
     "$pow",                     /* system task/function name */
     1                            /* forward reference = true */
    },
    {0} /*** final entry must be 0 ***/
};
```

```

void init_usertfs()
{
    p_tfcell usertf;
    for (usertf = veriusertfs; usertf; usertf++) {
        if (usertf->type == 0)
            return;
        mti_RegisterUserTF(usertf);
    }
}

```

A.4.3 Compiling and running PLI applications with ModelSim

With ModelSim, running a simulation that has PLI applications requires three steps:

1. Compile and link the PLI applications source files and the file containing the *veriusertfs* array as shared object libraries. All PLI applications can be compiled and linked into a single shared object library, or each PLI application can be compiled and linked into separate shared object libraries. The compile and link commands are specific to the operating system and C compiler. Three examples are:

On a Sun system with the Solaris operating system and the Sun C compiler:

```

cc -c -pic -I<MTI_install_dir>/shared/include <app_source_files>
ld -G -o mti_pli_apps.so <pli_app_object_files>

```

On an HP system with the HPUX operating system and the HP C compiler:

```

cc -c +z -I<MTI_install_dir>/shared/include <app_source_files>
ld -b -o mti_pli_apps.sl <pli_app_object_files> -lc

```

On a PC system with the Windows 95 operating system and the Visual C++ compiler:

```

cl -I<MTI_install_dir>\include -Fmti_pli_apps.dll
<app_source_files> -LD -link -dll -EXPORT:init_usertfs
<MTI_install_dir>\mtipli.lib

```

2. Modify the ModelSim .ini initialization file and specify the names of the shared object libraries in the *Veriuser* variable. This variable is typically found towards the middle of the initialization file, and by default it is commented out. Any number of shared object files may be listed, separated by a white space. As an example, if the shared library created in step 1 was called *mti_pli_apps.so*, then the *Veriuser* variable would be set to:

```


; List of dynamically loaded objects for Verilog PLI applications
Veriuser=mti_pli_apps.so

```

- Invoke the ModelSim simulator in the usual way. The simulator will automatically load and dynamically link the shared object files referenced in the .ini file.

Figure A-2 shows the results of running the ModelSim simulator on a small Verilog design which utilizes the `$show_value` and `$pow` PLI applications.

Figure A-2: *ModelSim* simulator using the `$show_value` and `$pow` PLI app's



A.5 Linking PLI applications to the Avant! Polaris simulator

The **Polaris**™ simulator from Avant!, Inc. is versatile Verilog simulator which runs as either a Verilog HDL interpreter based simulator or a compiled-code based simulator. Polaris runs on most Unix operating systems, on the Linux operating system, on the Windows NT operating system and on the Windows 95 operating system.

Polaris version 1998.2 was used to test some of the examples in this book, running on a Sun Sparc 10 system with the Solaris operating system.

A.5.1 IEEE 1364-1995 compliance

Polaris version 1998.2 supports the TF and ACC libraries of the IEEE 1364 Verilog PLI standard, but does *not* support the VPI library of the PLI standard.

Polaris does *not* use the IEEE 1364 *veriuser.h* and *acc_user.h* library files, which define the TF and ACC PLI libraries. Instead, Polaris uses proprietary files with the same names, which are derived from older Open Verilog International versions of these header files. These proprietary *veriuser.h* and *acc_user.h* header files are included in the installation directory of Polaris.

A.5.2 Specifying PLI application information

Polaris uses the same *veriusertfs array* as most other Verilog simulators to specify PLI application information. Refer back to section A.1.2 on page 636 for a description of the *veriusertfs array*. Polaris provides an empty *veriusertfs array* in a file called either *base_user.h* or *base_sdf.h*.

To specify a PLI application with the Polaris PLI interface mechanism:

1. Copy the Avant! C source files called *baseuser.c* and *baseuser.h*, which are provided with Polaris. Note that in some installations, these files may be named *user_sdf.c* and *base_sdf.h*.
2. In the copy of the *baseuser.h* file, edit the *veriusertfs array* to add the information for each PLI application to the array.

A sample Polaris *baseuser.c* file is listed in example A-10, below, for a *\$show_value* and a *\$pow* PLI application. The *baseuser.c* file uses the #include C preprocessor command to include this *baseuser.h* file.

Example A-10: Sample Polaris *baseuser.c* file

```
s_tfcell veriusertfs[]=
{
    {usertask,
        0,                                /* type of PLI routine */
        PLIbook_ShowVal_checktf,           /* user_data value */
        0,                                /* checktf routine */
        PLIbook_ShowVal_sizetf,            /* sizetf routine */
        0,                                /* calltf routine */
        "$show_value",                   /* misctf routine */
        1                                 /* system task/function name */
        /* forward reference = true */
    },
    {userfunction,
        0,                                /* type of PLI routine */
        PLIbook_pow_checktf,              /* user_data value */
        PLIbook_pow_sizetf,               /* checktf routine */
        PLIbook_pow_calltf,               /* sizetf routine */
        PLIbook_pow_misctf,               /* calltf routine */
        "$pow",                           /* misctf routine */
        1                                 /* system task/function name */
        /* forward reference = true */
    },
{0}
};
```

A.5.3 Compiling and linking PLI applications for Polaris

Polaris statically links PLI applications into the Polaris simulator. The process of compiling and linking PLI applications is nearly the same for the Unix and Windows operating systems. The basic steps involved are:

1. Run a utility program provided with Polaris, called either *fdagenmak* (for the Unix version of Polaris) or *genmake* (for the Windows version of Polaris). This program is used to configure the options to be linked into the Polaris simulator, such as which waveform display to use. The program generates a *make file* for the C compiler that corresponds to the operating system on which the fdagenmak or genmake program was run.
2. Use the C compiler's make utility to generate a new Polaris simulator executable.
3. Invoke simulations in the usual method, but with one additional invocation option: **-pli**.

An example output from running the *fdagenmak* program on a Unix system is listed in Example A-11. The answers required for the examples of the *\$show_value* and *\$pow* PLI applications are shown in bold.

Example A-11: Sample Polaris *fdagenmak* program output

```
Name of Linker (default = fdalink): <CR>
Name of c-compiler (default = cc ): <CR>
Name of C++ compiler (default = CC): <CR>
Do you need to run with SignalScan ?
Enter Yes / No (default No): Yes <CR>
SignalScan library dir (default: ${POLARIS}/designacc/lib/signalscan):
<CR>
Do you need to run with Virsim ?
Enter Yes / No (default No): <CR>
Are you using LMC models?
Enter Yes / No (default No): <CR>
Specify the name of PLI object(s) that need to be linked in.
The generated makefile will contains rules to automatically
compile the corresponding .c, .cxx, .cpp, or .C file(s)to generate the
specified .o file(s).
Enter one object/lib per line. An empty line ends the list
show_value_acc.o <CR>
pow_acc.o <CR>
veriuser.o <CR>
<CR>
Specify additional user libraries that need to be linked in.
Enter one object/lib per line. An empty line ends the list
<CR>
Generating makefile: makefile.pli ...
Generating pli list file: pli.lst ...
Run Polaris-COM with the option -pli=pli.lst
Done !!!
```

When the *fdagenmak* program completes, it generates two files: a C make file called *makefile.pli* and a text file called *pli.lst*.

After running *fdagenmak*, the next step is to run the make utility program with the *makefile.pli* file. This will compile and link the PLI application, and generate a new Polaris simulator executable.

Simulations are then run by using the new Polaris simulator executable. When simulation is invoked, the *pli.lst* file is passed to the Polaris simulator using the **-pli** invocation option. This file informs the simulator of the location of the object files which were created by the make utility. The Polaris simulator will link in the PLI application object files as the Verilog models are compiled and linked.

A.6 Linking to *Silos III* from Simucad, Inc.

The *Silos III™* simulator from Simucad, Inc. is both a digital logic simulator and a digital fault simulator. Silos III runs on several Unix operating systems, on the Windows NT operating system, and on the Windows 95 operating system.

Silos III version 98.100 was used to test several of the examples in this book, running on a Pentium II desktop system with the Windows NT operating system.

A.6.1 IEEE 1364-1995 compliance

Silos III version 98.100 supports a subset of the TF and ACC libraries of the IEEE 1364 Verilog PLI standard, but does *not* support the VPI library of the PLI standard.

Silos III does *not* use the IEEE 1364 *veriuser.h* and *acc_user.h* library files, which define the TF and ACC PLI libraries. Instead, Silos III uses proprietary files with the same names, which are derived from older Open Verilog International versions of these header files. These proprietary *veriuser.h* and *acc_user.h* header files are included in the installation directory of Silos III.

NOTE → Silos III also provides a number of extensions to the PLI standard, primarily to support the fault simulation capabilities of Silos III. The extensions to the PLI standard are specified in a proprietary header file called *ext_user.h*.

There are two deviations which Silos III makes from the IEEE standard that are important to note.

- The IEEE standard defines that the TF routines *tf_error()* and *tf_message()* should print an error message and abort compilation, when called from a *checkif routine*. In Silos III, however these routines only print an error message.
- The IEEE standard specifies that the *tf_dofinish()* routine should exit simulation from any routine. In Silos III, this *tf_dofinish()* does not cause simulation to exit.

A.6.2 Specifying PLI application information

Silos III uses the same *veriusertfs array* as most other Verilog simulators to specify PLI application information. Refer back to section A.1.2 on page 636 for a description of the *veriusertfs array*. Silos III does not provide an example *veriusertfs array*. A file containing this array must be created by the user.

To specify the PLI application for Silos III:

1. Create a new C source file, which can be any name
2. In the file, define a *veriusertfs array*.
3. Add the information for each PLI application to the array. Example A-9 lists a sample file for specifying the \$pow and \$show_value PLI applications.

Example A-12: Sample file to specify PLI applications for the Silos III simulator

```
#include "veriuser.h"

/* prototypes of the PLI application routines */
extern int PLIbook_ShowVal_checktf(), PLIbook_ShowVal_calltf();
extern int PLIbook_pow_sizetf(), PLIbook_pow_checktf(),
          PLIbook_pow_calltf(), PLIbook_pow_misctf();

/* the veriusertfs table */
s_tfcell veriusertfs[] =
{
    {usertask,
     0,                                /* type of PLI routine */
     PLIbook_ShowVal_checktf,           /* user_data value */
     0,                                /* checktf routine */
     PLIbook_ShowVal_calltf,           /* sizetf routine */
     0,                                /* calltf routine */
     "$show_value",                   /* misctf routine */
     1,                                /* system task/function name */
     1}                                /* forward reference = true */

    ,{userfunction,
     0,                                /* type of PLI routine */
     PLIbook_pow_checktf,              /* user_data value */
     PLIbook_pow_sizetf,               /* checktf routine */
     PLIbook_pow_calltf,               /* sizetf routine */
     PLIbook_pow_misctf,               /* calltf routine */
     "$pow",                           /* misctf routine */
     1,                                /* system task/function name */
     1}                                /* forward reference = true */

    ,{0} /* final entry must be 0 */
};

}
```



TIP *Do not place PLI applications in the file with the veriusertfs array!* The C programming language does not permit multiple global arrays with the same name. In a typical design environment, PLI applications will come from several sources. If the *veriusertfs array* and a PLI application are in the same file, then the source code for the application must be available whenever another PLI application needs to be added to the array. If two PLI applications were to both include the *veriusertfs array* in the application, then the object files for both applications could not be used together.

A.6.3 Compiling and linking PLI applications for Silos III using Unix operating systems

The Silos III simulator dynamically loads PLI applications during simulation. The PLI applications source files, including the file containing the *veriusertfs* array, are compiled and linked to into a shared object library file. This shared library is then loaded by the simulator. The steps required to compile and link an application are:

1. Compile each PLI application source file, including the file containing the *veriusertfs* array. The Silos III simulator does not impose any restrictions on the C compiler options and optimization levels used to compile the C source files.
2. Link the object files into a shared library. The linker commands are both operating system and dependent. As an example, the command for the Sun Solaris operating system is:

```
ld -o silos_pli_apps.so -dy -G "*.*"
```

Silos III does not provide an example C make file for compiling and linking PLI applications. PLI application developers may wish to create their own make files.

A.6.4 Compiling and linking PLI applications for Silos III using the Windows NT or Windows 95 operating system

For the Windows NT and Windows 95 operating systems, compiling and linking a PLI application for Silos III involves the following steps:

1. Create a text ***definitions file***, which contains the names of the names of the C functions that make up the PLI application. This file also defines the name of the C array of *s_tfcell* structures. This array is the same as the *veriusertfs array* used by most other Verilog simulators, and by the Unix versions of Silos III. Example A-13, on the next page, shows an example definitions

Note: Silos III, running on Windows NT or Windows 95, does not require that the array of *s_tfcell* structures be called *veriusertfs*. This means a separate array can be created for different PLI applications. It is not necessary to merge multiple PLI applications into a single array. Multiple ***definition files*** can be used to specify multiple *s_tfcell* arrays.

2. Compile and link the PLI application source files to create a dynamically linked library file (a ***dll*** file). The link command should include the names of the definition file(s). With the Microsoft Visual C++ compiler, the **-defs** option is used to specify the definition files.

Silos III does not provide an example C make file for compiling and linking PLI applications. PLI application developers may wish to create their own make files. An example make file is listed in Example XXX, which is shown below.

Example A-13 lists a sample definitions file for the *\$show_value* and *\$pow* PLI applications.

Example A-13: Sample PLI definitions file for the Silos III simulator

```
LIBRARY
EXPORTS
    veriusertfs
    PLIbook_ShowVal_checktf
    PLIbook_ShowVal_calltf
    PLIbook_pow_sizetf
    PLIbook_pow_checktf
    PLIbook_pow_calltf
    PLIbook_pow_misctf
```

Example A-14 lists a simple make file for Visual C++.

Example A-14: Sample Visual C++ make file for the Silos III simulator

```
SILOS_INST_DIR=c:/progra~1/silos3

#
# SRCS set this to the list of source files for the PLI application
#
SRCS = pow_acc.c \
        show_values_acc.c \
        PLIbook_veriuser.c

#
# DEFS set this to the list of definition files for the PLI application
#
DEFS = silos_pli.def

OBJS = $(SRCS:.c=.obj)
CFLAGS = -DMSC -DWIN32 -I$(SILOS_INST_DIR)/pli -MD -O2

.c.obj:
$(CC) $(CFLAGS) -c $<

silos_pli.dll: $(OBJS)
link -dll /out:$@ -def:$(DEFS) $(OBJS) $(SILOS_INST_DIR)/sse.lib
```

A.6.5 Notifying Silos III of the PLI application object files

Once a PLI application has been compiled and linked, the Silos III simulator must be notified of the location of the shared object files. Silos III simulations can be run from either a text-based command line or from a graphical user interface. For the former, the location of the PLI application shared object file(s) are specified as special “*!/pli-load*” command in the Verilog netlist. With the graphical user interface, the location of the PLI shared object file(s) are specified through the project menus and forms.

1. Select the **Project->Files** menu in Silos III.
2. In the **Project Files** dialog box, select the drop down arrow in the **File Group**.
3. Select **PLI Library Files**.
4. Add the PLI shared object library file(s) that contain the PLI applications used by the Verilog design.

Figure A-3 shows the Project Files form for specifying a PLI application shared object library file called *silos_pli.dll*.

Figure A-3: The Silos III Project Files form





Figure A-4 shows the results of running the Silos III simulator on a small Verilog design which utilizes the `$show_value` and `$pow` PLI applications.

Figure A-4: Simucad *Silos III* simulator using the `$show_value` and `$pow` PLI app's



A.7 *Linking to VCS from Synopsys, Inc.*

The **VCS™** simulator from Synopsys is a high performance Verilog simulator that is well known for its RTL simulation speed. VCS was originally developed by Chronologic, Inc. in 1991. Chronologic was acquired by Viewlogic, Inc., which was later acquired by Synopsys. VCS runs on most Unix operating systems, on the Windows NT operating system and on the Windows 95 operating system.

VCS is a command-line based text-only simulator. VCSI, is an enhanced version of VCS with additional interactive debug capabilities, and XVCS adds a complete graphical user interface to the VCSI version. VCSI version 4.1.1 was used to test many of the examples in this book, running on a Pentium II workstation with the Windows NT operating system.

A.7.1 IEEE 1364-1995 compliance

VCSI version 4.1.1, which was used to test examples in this book, supports the TF and ACC libraries of the IEEE 1364 Verilog PLI standard, but does not support the VPI library of the PLI standard. This version of VCS implements the 1990 OVI PLI 1.0 version of the Verilog PLI standard, *not* the IEEE 1364-1995 PLI standard. Since the TF and ACC routines in the IEEE 1364 standard are derived from the OVI PLI 1.0 standard, VCSI version 4.1.1 is nearly compliant with the TF and ACC portion of the IEEE 1364 standard, but deviates from the standard in several ways.

The most critical deviation which VCS has from the IEEE 1364 standard for the TF and ACC libraries is that VCS does not use the IEEE 1364 *veriuser.h* and *acc_user.h* library files, which define the TF and ACC PLI libraries. Instead, VCS uses proprietary *vcsuser.h* and *acc_user.h* files to define these libraries. There are several minor differences in the VCS library definitions and the IEEE standard libraries. For example, the boolean *true* and *false* constants are defined in the IEEE 1364 *veriuser.h* file and not in the *acc_user.h* file, but in VCS they are defined in the *acc_user.h* file and not in *vcsuser.h*.

Before a PLI application can be compiled with VCS, it is necessary to modify the PLI application source code and change the `#include` statements so that they reference the VCS proprietary libraries. Conveniently, VCS defines a C macro name called **vcs**, which can be used in PLI application source code for conditional compilation. This allows the same application can be compiled with VCS as well as any other Verilog simulator that uses the IEEE library header files. For example:

```
#ifdef VCS
#include vcsuser.h /* VCS proprietary TF routine library */
#include acc_user.h /* VCS proprietary ACC routine lib. */
#else
#include veriuser.h /* IEEE 1364 TF routine library */
#include acc_user.h /* IEEE 1364 ACC routine library */
#endif
```

Another deviation VCS makes from the IEEE standard, is that the standard defines that the TF routines `tf_error()` and `tf_message()` should print an error message and abort compilation, when called from a *checktf routine*. In the VCS, however these routines only print an error message. In order to force a fatal error in VCS, the `tf_dofinish()` routine must be called after the message is printed.

A.7.2 Enabling PLI support with VCS

The VCS simulator optimizes the run-time performance of the PLI applications by providing PLI application developers a means to configure the level of access the PLI has into the Verilog simulation data structure.

By default, VCS only supports the TF library of the PLI standard, which allows VCS to optimize simulation for maximum performance. The TF routines can only access the simulation data that is passed to the PLI application through system task or system function arguments. This limited access means that at compile time, VCS can determine exactly what data will be accessed by the simulator, and can optimize the simulation data structure for faster performance. The ACC routines, however, can arbitrarily access information anywhere and at any time in the simulation data structure. This arbitrary access cannot be predicted at compile time, and therefore cannot be optimized as efficiently. VCS enables support for the ACC library in increments, and can specify the added support for the entire simulation or for just certain regions of a design. VCS allows the support for ACC to be enabled:

- Per Verilog module definition
- Per ACC capability (for example: read only access, read and write access, etc.).

The fastest simulation performance is achieved with VCS when only the TF library is used in PLI applications. When ACC routines are used, simulation performance may be impacted, but by allowing a PLI application developer to control how much access into the simulation data structure that the ACC routines have, the impact on simulation performance can be controlled.

18.11.1 Specifying PLI application information—the VCS *PLI table* file

VCS does not use the defacto standard *veriuserif*s array that most other Verilog simulators use to specify PLI application information. Instead, VCS uses a proprietary *PLI table* file. This table specifies:

- The PLI application information.
- The configuration of ACC library access into the simulation data structure.

The PLI table file can contain the specification for any number of PLI applications. Each specification must be on a single line. Comments may be included in the file using a // token to begin the comment and a carriage return to terminate the comment. VCS allows any number of PLI tables to be used. This means a separate PLI table file can be created for each PLI application, and for any given simulation, only the PLI tables which are required need to be specified to the simulator. This is much more flexible than using the C language *veriuserif*s array, which, in order to compile, requires that all PLI applications be defined in a single array.

The general format of the VCS PLI table is:

```
$<system_task_function_name> <PLI_spec> <ACC_spec>
```

The <PLI_spec> is zero, one or more of the specifications listed in Table A-3. The specifications can be listed in any order, but must be listed on the same line.

PLI Specification	Description
call=<routine_name>	Name of the <i>callif routine</i> for the PLI application
check=<routine_name>	Name of the <i>checkif routine</i> for the PLI application
misc=<routine_name>	Name of the <i>misctf routine</i> for the PLI application
size=<number_or_r>	Size of a system function return, or r if the return is a real number
data=<number>	Integer user data value to be passed to the <i>callif routine</i> , <i>checkif routine</i> and <i>misctf routine</i>

Table A-3: The Synopsys VCS PLI table PLI routine specifications

NOTE → VCS infers that the PLI application is a system task if there is no size specification for the system task name. A system function is inferred if there is a size specification.



NOTE VCS does not support the *sizetf routine*, which is part of the IEEE 1364 PLI standard. Instead, VSC defines the function return size explicitly in the VCS PLI table. This is a minor deviation from the standard, since the VCS implementation provides the same functionality.

The <ACC_spec> is zero, one or more of the specifications listed in Table A-4. The specifications can be listed in any order, but must follow the <PLI_spec>, and be listed on the same line as the <PLI_spec>. The format of the <ACC_spec> is:

```
acc=<operation><capability>:<design_scopes>+
```

The <operation> is one of the four tokens listed in Table A-4.

ACC Specification Operation	Description
<code>+ =</code>	Add the specified capabilities to the specified scope
<code>=</code>	The same as <code>+ =</code>
<code>- =</code>	Remove the specified capabilities from the specified scope
<code>: =</code>	Set the specified capabilities to the specified scope

Table A-4: The Synopsys VCS PLI table ACC specification operations

The <capability> of the ACC specification must be one of commands listed in Table A-5.

ACC Capability
read (abbreviation: <code>r</code>) Allow reading values of nets, regs and variables
read_write (abbreviation: <code>rw</code>) Allow reading values of nets, regs and variables, and writing values to regs and variables
callback (abbreviation: <code>cbk</code>) Allow VCL (Value Change Link) callbacks on named objects

Table A-5: The Synopsys VCS PLI table ACC capability specification

ACC Capability	
callback_all (abbreviation: cbka)	Allow VCL callbacks on named and unnamed objects
force (abbreviation: frc)	Allow forcing values onto nets and regs
timning_check_backannotation (abbreviation: tchk)	Allow delay back annotation of timing check limits
gate_backannotation (abbreviation: gate)	Allow delay back annotation of primitives
module_path_backannotation (abbreviation: mp)	Allow delay back annotation of module paths
module_input_port_backannotation (abbreviation: mip)	Allow delay back annotation of module input ports
module_input_port_bit_backannotation (abbreviation: mipb)	Allow delay back annotation of bits of a module input port

Table A-5: The Synopsys VCS PLI table ACC capability specification (continued)

The <scope> of the ACC specification is one of names listed in Table A-6.

ACC Specification Scope	Description
<module_name>	The definition name of any Verilog module. The specified operation and capability will apply to all instances of that module.
%TASK	Any module which contains the system task/function which is being specified. The specified operation and capability will apply to all instances of those modules.
%CELL	Any module flagged as a cell (using the `celldefine compiler directive or the -y or -v invocation options). The specified operation and capability will apply to all instances of those modules.
*	A wild card which indicates all Verilog modules in the design.

Table A-6: The Synopsys VCS PLI table ACC specification scope

If a plus sign (+) is specified after the <scope>, then the ACC capability is applied to both the specified module, and all levels of hierarchy below that module. If the plus sign is not specified, then the capability is only applied to the specified module.

Example A-15 lists a VCS PLI table file for a *\$show_value* application and a *\$pow* application. *Note:* Each PLI application must be specified on a single line. The following example wraps each table entry to a second line to fit the page format of this book, but in the actual PLI table, a table entry cannot be split into multiple lines.

Example A-15: PLI table file to specify PLI applications for the VCS simulator

```
// Example Synopsys VCS PLI table to register PLI applications
// For the book, "Using The Verilog PLI", Chapter 4

$show_value check=PLIbook_ShowVal_checktf call=PLIbook_ShowVal_calltf
  data=0 acc+=read:*
$pow check=PLIbook_pow_checktf call=PLIbook_pow_calltf
  misc=PLIbook_pow_misctf size=32 data=0
```

A.7.3 Compiling and running PLI applications with VCS

NOTE → The procedure for compiling PLI applications to the VCS simulator is identical for all operating systems. Therefore the discussion of VCS presented in this section does not refer to any particular operating system.

Running a simulation with VCS requires two steps:

1. Invoking the VCS compiler.
2. Invoking the simulation executable created by the compiler

To invoke the VCS compiler with a PLI application involves:

- Specifying the names of the PLI table files using the **-P** invocation option.
- Specifying the names of the PLI application source or object files on the invocation command line


For example, a simulation using the *\$show_value* and *\$pow* PLI applications might be invoked using the command:

```
vcsi -Mupdate test.v -P my_pli_apps.tab pow.c showval.c
```

VCS does not require that PLI applications be pre-compiled. The source files for the PLI applications can be specified as inputs to the VCS compiler, and VCS will automatically compile and link the applications into the VCS simulation executable. Alternatively, the PLI application can be pre-compiled into object files or shared library files, and the appropriate file names specified on the VCS compiler command line. If the PLI applications are pre-compiled, the VCS compiler will just link the applications into the VCS simulation executable.

Figure A-5 shows the results of running the VCS simulator on a small Verilog design which utilizes the `$show_value` and the `$pow` PLI applications. The example simulation results is shown using the XVCS graphical debug environment for VCSI.

Figure A-5: Synopsys VCS simulator using for the `$show_value` and `$pow` PLI app's



A.8 Linking to VeriBest from VeriBest, Inc.

The **VeriBest™** Verilog simulator from VeriBest, Inc. is an easy-to-use Verilog simulator which runs as either a Verilog HDL interpreter based simulator or as a compiled-code based simulator. VeriBest runs on major Unix operating systems, and on the Windows NT operating system. VeriBest uses graphical forms and menus to compile and link PLI applications, making it one of the simplest of all Verilog simulators for using the Verilog PLI.

Note: The VeriBest Verilog simulator is based on the **Fintronic FinSim™** simulator. VeriBest adds the graphical user interface, waveform displays and other enhancements to the FinSim simulator. The information regarding Verilog PLI support presented here also applies to the Fintronic FinSim simulator, but, with FinSim, the compile, link, and invocation commands are specified from text-based command lines rather than from graphical forms.

VeriBest version VB98.0B (FinSim version 4.5.26) was used in this book, running on a Pentium II desktop system with the Windows NT operating system.

A.8.1 IEEE 1364-1995 compliance

VeriBest version VB98.0B supports the TF and ACC libraries of the IEEE 1364 Verilog PLI standard, but does *not* support the VPI library of the PLI standard.

VeriBest does *not* use the IEEE 1364 *veriuser.h* and *acc_user.h* library files, which define the TF and ACC PLI libraries. Instead, VeriBest uses proprietary files with the same names, which are derived from older Open Verilog International versions of these header files. These proprietary *veriuser.h* and *acc_user.h* header files are included in the installation directory of VeriBest.

A.8.2 Specifying PLI application information

VeriBest uses the same *veriusertfs array* as most other Verilog simulators to specify PLI application information. Refer back to section A.1.2 on page 636 for a description of the *veriusertfs array*. VeriBest provides an example *veriusertfs array* in a file called *veriuser.c*.



TIP

VeriBest can utilize the same *veriuser.c* file that is provided with the Cadence Verilog-XL simulator, making it easier to use the same *veriusertfs array* with both products. However, the *veriuser.c* file provided with VeriBest includes predefined PLI applications for the VeriBest waveform viewer and state diagram editor.

To specify a PLI application with the VeriBest PLI interface mechanism:

1. Copy a C source file called *veriuser.c*, that is provided with VeriBest.
2. In the copy of the *veriuser.c* file, edit the *veriusertfs* array to add the information for each PLI application to the array.

The *veriuser.c* file for VeriBest is located in the installation directory of the simulator software. On Unix and NT systems, the typical location is:

```
<veribest_install_directory>/vbvlg/include/veriuser.c
```

 **NOTE** *Do not modify the original veriuser.c file!* Make a copy of the file, and edit the copy. If the original file were to somehow become corrupted, it could require re-installing the VeriBest simulator software.

A sample VeriBest *veriuser.c* file is listed in example A-10, below, for a *\$show_value* and a *\$pow* PLI application. The information added for this application is shown in bold type.

Example A-16: Sample VeriBest *veriuser.c* file (partial listing)

```
#include "veriuser.h"
#include "acc_user.h"

...
int sde_debug_check(int data, int reason);
int sde_debug_call(int data, int reason);
int sde_debug_misc(int data, int reason);
int WavesReadNextVector( int data, int reason );
int WavesOpenVectorFile( int data, int reason );

/* prototypes for the PLI application routines */
extern int PLIbook_ShowVal_checktf(), PLIbook_ShowVal_calltf();
extern int PLIbook_pow_sizetf(), PLIbook_pow_checktf(),
          PLIbook_pow_calltf(), PLIbook_pow_misctf();

...
s_tfcell veriusertfs[] = {

...
/* $sde_debug is used for state diagram debugging */
{usertask, 0, sde_debug_check, 0, sde_debug_call, sde_debug_misc,
"$sde_debug", 0},
/* $Waves is for the vector file support in Waves Stimulus editor */
{userfunction, 0, WavesOpenVectorFile , WavesOpenVectorFile ,
WavesOpenVectorFile , 0, "$WavesOpenVectorFile", 0 },
{usertask, 0, WavesReadNextVector , WavesReadNextVector ,
WavesReadNextVector, 0, "$WavesReadNextVector", 0 },
```

```

{usertask,                      /* type of PLI routine */
  0,                            /* user_data value */
  PLIbook_ShowVal_checktf,     /* checktf routine */
  0,                            /* sizetf routine */
  PLIbook_ShowVal_calltf,      /* calltf routine */
  0,                            /* misctf routine */
  "$show_value",                /* system task/function name */
  1                             /* forward reference = true */

},
{userfunction,                  /* type of PLI routine */
  0,                            /* user_data value */
  PLIbook_pow_checktf,         /* checktf routine */
  PLIbook_pow_sizetf,          /* sizetf routine */
  PLIbook_pow_calltf,          /* calltf routine */
  PLIbook_pow_misctf,          /* misc tf routine */
  "$pow",                       /* system task/function name */
  1                             /* forward reference = true */

},
/* all entry must be entered before this line */
{0, 0, 0, 0, 0, 0, 0, 0} /* this must be the last entry */
};

...

```



TIP *Do not place PLI applications in the veriuser.c file!* The C programming language does not permit multiple global arrays with the same name, such as the *veriuserfts* array. In a typical design environment, PLI applications will come from several sources, such as internally developed applications, 3rd party applications, and ASIC vendor applications. If the *veriuserfts* array and a PLI application are in the same file, then the source code for the application must be available whenever another PLI application needs to be added to the table. If two PLI applications were to both include the *veriuserfts* array in the application, then the object files for both applications could never be used together.

A.8.3 Compiling and linking PLI applications for VeriBest

VeriBest uses a graphical user interface to specify all information for compiling and linking PLI applications. This makes it easy to specify the information. More importantly, the user interface makes the operating system specific details transparent to the PLI application user.

To specify a PLI application involves the following steps:

1. Invoke the VeriBest Verilog simulator.

2. In the **VeriBest Verilog Builder** form, select the **PLI Options** button. Figure A-6, on the following page, shows an example of the Builder form.
3. In the **Custom PLI Analyzer** form, complete the fields which specify the PLI application files. Figure A-7 shows an example of the Custom Analyzer form.
 - Enable the **custom** option.
 - Enable the **Add PLI's** option.
 - Enter a name for the custom VeriBest analyzer that will be created.
 - Select the **User PLI C Files** button.
 - Select the **Add** button. Use the browser form that comes up to select the C source files containing the PLI application, including the file that contains the copy of the veriuser.c file.
 - Enter the path for the build directory. This directory contains the C source code files for the PLI application. It is also the directory where the custom VeriBest analyzer will be created.
4. In the **VeriBest Verilog Builder** form, select the **PLI Build** button. This will compile the PLI applications and build a new VeriBest analyzer that has the PLI applications linked into the analyzer.

Once the custom VeriBest analyzer is built, simulations are run using the standard VeriBest commands and forms.

Figure A-8 shows the results of running the VeriBest simulator on a small Verilog design which utilizes the `$show_value` and `$pow` PLI applications.




Figure A-6: The *VeriBest Verilog Builder* form**Figure A-7:** The VeriBest *Custom PLI Analyzer* form

Figure A-8: VeriBest *VeriBest* simulator using the \$show_value and \$pow PLI app's

A.9 Summary

This appendix has presented how Verilog PLI applications are linked into several Verilog simulators. The primary objective has been to show enough information to aid in getting started with using the PLI with these simulators. It was not an objective to present all the features available with these simulators. Readers should always refer to a simulator vendor's documentation for the most up-to-date information on using the Verilog PLI with a specific simulator.

APPENDIX B

The IEEE 1364-1995 VPI Routine Library

Appendix B lists the definitions of all VPI routines contained in the IEEE 1364-1995 library. VPI object diagrams are shown for all Verilog objects which can be accessed by VPI routines. These object diagrams are similar to the object diagrams contained in the IEEE 1364 standard, but are drawn in a different format and with additional information that is not contained in the IEEE standard. For example, the IEEE object diagrams do not list most of the constant names used to traverse hierarchy or access information about an object, preferring to let the reader derive the constant names from the object names.

Note: This appendix is an excerpt from the booklet “Verilog PLI Quick Reference Guide”, published by Sutherland HDL, Inc., Portland, Oregon. Copyright 1995, 1998. Used with permission.

B.1 VPI Objects Relationships

B.1.1 Object diagram legend




```
expression_handle = vpi_handle(vpiExpr, prim_tern_handle);
```



```
net_iterator = vpi_iterate(vpiNet, module_handle);
```

module


A **solid enclosure with bold font** represents an **object definition**. The diagram shows all relationships which can be traversed to and from the object.



A **dotted enclosure with bold-italic font** represents an **object class definition**. The name of the object class is listed at the top of the enclosure (a class may be unnamed). The diagram shows all relationships which can be traversed to and from the object class.

port

A **solid enclosure with standard font** represents an **object reference**. Another diagram defines the referenced object.




A **dotted enclosure with italic font** represents an **object class reference**. Another diagram defines the referenced object class.



A **small circle** indicates a **NULL** is used for the reference handle supplied to **vpi_handle()** or **vpiIterate()**.

B.1.2 VPI module objects




<i>int</i>	vpiType	returns vpiModule
<i>bool</i>	vpiTopModule	returns true if a module is a top-level module
<i>bool</i>	vpiCellInstance	returns true if a module is flagged as a cell (using ‘celldefine’)
<i>bool</i>	vpiProtected	returns true if the module source is protected
<i>int</i>	vpiTimeUnit	returns the module time unit as 2 down to -15, where 2==100 seconds, 1==10s, 0==1s, -1==100ms, -2 ==10ms, -3==1ms,... -6==1us,... -9==1ns,... -12==1ps,... -15==1fs
<i>int</i>	vpiTimePrecision	returns module time precision as 2 down to -15
<i>int</i>	vpiDefNetType	returns the default net type as vpiWire , vpiWand , vpiWor , vpiTri , vpiTri0 , vpiTri1 , vpiTriReg , vpiTriAnd , vpiTriOr , vpiSupply1 , vpiSupply0
<i>int</i>	vpiUnconnDrive	returns the unconnected port drive strength as vpiHighZ , vpiPull1 , vpiPull0

continued on next page

<i>int</i>	vpiDefDelayMode	returns the delay mode of the module as: vpiDelayModeNone , vpiDelayModePath , vpiDelayModeDistrib , vpiDelayModeUnit , vpiDelayModeZero , vpiDelayModeMTM
<i>int</i>	vpiDefDecayTime	returns the default decay time for trireg nets
<i>str</i>	vpiName	returns the module instance name
<i>str</i>	vpiFullName	returns the module full hierarchical path name
<i>str</i>	vpiDefName	returns the module definition name
<i>str</i>	vpiFile	returns the file name containing the module instance
<i>int</i>	vpiLineNo	returns the file line number containing the module instance
<i>str</i>	vpiDefFile	returns the file name where the module is defined
<i>int</i>	vpiDefLineNo	returns the file line number where module is defined


1. **vpi_iterate(vpiModule, NULL)** returns an iterator for all top-level modules.

B.1.3 VPI scope objects:



<i>int</i>	vpiType	returns vpiModule , vpiTaskFunc , vpiTask , vpiFunction , vpiNamedBegin or vpiNamedFork
<i>str</i>	vpiName	returns the name of the scope
<i>str</i>	vpiFullName	returns the full hierarchical path name of the scope

B.1.4 VPI expression objects




<i>int</i>	vpiType	returns vpiPartSelect , vpiOperation , vpiConstant , vpiFuncCall , vpiSysFuncCall , vpiNet , vpiReg , vpiIntegerVar , vpiTimeVar , vpiRealVar , vpiVarSelect , vpiParameter , vpiMemory , vpiMemoryWord
<i>int</i>	vpiOpType	if vpiType is vpiOperation , returns one of the operator constants shown in note 4, below
<i>int</i>	vpiConstType	if vpiType is vpiConstant , returns vpiDecConst , vpiRealConst , vpiBinaryConst , vpiOctConst , vpiHexConst , vpiStringConst
<i>int</i>	vpiSize	returns the bit size of the expression result
<i>str</i>	vpiFile	returns the file name containing the expression
<i>int</i>	vpiLineNo	returns the file line number containing the expression

1. The value of an expression is accessed using **vpi_get_value()**. The value of a memory cannot be accessed; instead, iterate on memory words and access the value of each word.
2. If an operator is **vpiMultiConcat**, the first operand is the multiplier expression.
continued on next page

3. For vectors, **vpiUse** accesses any use of the vector, including part and bit selects. For bit-selects, **vpiUse** accesses any use of that bit and any use of the.parent vector.
4. The constants returned from `vpi_get(vpiOpType, <operation_handle>)` are:

vpiMinusOp	unary minus operator
vpiPlusOp	unary plus operator
vpiNotOp	unary not operator
vpiBitNegOp	bitwise negation operator
vpiUnaryAndOp	bitwise reduction and operator
vpiUnaryNandOp	bitwise reduction nand operator
vpiUnaryOrOp	bitwise reduction or operator
vpiUnaryNorOp	bitwise reduction nor operator
vpiUnaryXorOp	bitwise reduction xor operator
vpiUnaryXNorOp	bitwise reduction xnor operator
vpiSubOp	binary subtraction operator
vpiDivOp	binary division operator
vpiModOp	binary modulus operator
vpiEqOp	binary equality operator
vpiNeqOp	binary inequality operator
vpiCaseEqOp	case equality operator (<code>==</code>)
vpiCaseNeqOp	case inequality operator (<code>!=</code>)
vpiGtOp	binary greater-than operator
vpiGeOp	binary greater-than-or-equal operator
vpiLtOp	binary less-than operator
vpiLeOp	binary less-than-or-equal operator
vpiLShiftOp	binary left shift operator
vpiRShiftOp	binary right shift operator
vpiAddOp	binary addition operator
vpiMultOp	binary multiplication operator
vpiLogAndOp	binary logical and operator
vpiLogOrOp	binary logical or operator
vpiBitAndOp	binary bitwise and operator
vpiBitOrOp	binary bitwise or operator
vpiBitXorOp	binary bitwise xor operator
vpiBitXNorOp	binary bitwise xnor operator
vpiConditionOp	ternary conditional operator
vpiConcatOp	concatenation operator
vpiMultiConcatOp	replication operator (repeated concatenation)
vpiEventOrOp	event or operator
vpiNullOp	null operation
vpiListOp	list of expressions
vpiMinTypMaxOp	min:typ:max: delay expression
vpiPosedgeOp	posedge operator
vpiNegedgeOp	negedge operator


B.1.5 VPI port objects



<i>int</i>	vpiType	returns vpiPort or vpiPortBit
<i>bool</i>	vpiScalar	returns true if the port is scalar
<i>bool</i>	vpiVector	returns true if the port is a vector
<i>bool</i>	vpiExplicitName	returns true if the port is explicitly named
<i>bool</i>	vpiConnByName	returns true if the port instance is connected by name
<i>int</i>	vpiSize	returns the size of the port
<i>int</i>	vpiPortIndex	returns the index position of the port
<i>int</i>	vpiDirection	returns vpiInput , vpiOutput , vpiInout , vpiMixedIO , vpiNoDirection
<i>str</i>	vpiName	returns the name of the port
<i>str</i>	vpiFile	returns the file name containing the port
<i>int</i>	vpiLineNo	returns the file line number of the port


- Properties **vpiPortIndex** and **vpiName** do not apply for port bits.
- vpiPortIndex** can be used to determine the port order.
- Module input port delays (MIPD's) are accessed using **vpi_get_delays()** and **vpi_put_delays()**.
- To access the port connected to an expression, determine the type of the expression and refer to the object diagram for that object type.

B.1.6 VPI I/O declaration objects



<i>int</i>	vpiType	returns vpiIODecl
<i>bool</i>	vpiScalar	returns true if the declaration is scalar
<i>bool</i>	vpiVector	returns true if the declaration is a vector
<i>int</i>	vpiSize	returns the bit size of the declaration
<i>int</i>	vpiDirection	returns vpiInput , vpiOutput , vpiInout , vpiMixedIO , vpiNoDirection
<i>str</i>	vpiName	returns the name of the declaration
<i>str</i>	vpiFile	returns the file name containing the declaration
<i>int</i>	vpiLineNo	returns the file line number of the declaration


B.1.7 VPI net objects



<i>int</i>	vpiType	returns vpiNet
<i>int</i>	vpiNetType	returns vpiWire , vpiWand , vpiWor , vpiTri , vpiTri0 , vpiTri1 , vpiTriReg , vpiTriAnd , vpiTriOr , vpiSupply1 , vpiSupply0
<i>bool</i>	vpiScalar	returns true if the net is scalar
<i>bool</i>	vpiVector	returns true if the net is a vector
<i>bool</i>	vpiExplicitScalared	returns true if the net is explicitly declared scalared
<i>bool</i>	vpiExplicitVectored	returns true if the net is explicitly declared vectored
<i>bool</i>	vpiExpanded	returns true if the net is expanded
<i>bool</i>	vpiNetDeclAssign	returns true if the net decl. is combined with a cont assign
<i>bool</i>	vpiImplicitDecl	returns true if the net is implicitly declared
<i>int</i>	vpiSize	returns the bit size of the net
<i>int</i>	vpiStrength0	returns the logic 0 strength of the net
<i>int</i>	vpiStrength1	returns the logic 1 strength of the net
<i>int</i>	vpiChargeStrength	returns the capacitance strength of the net
<i>str</i>	vpiName	returns the declaration name of the net
<i>str</i>	vpiFullName	returns full hierarchical path name of the net
<i>str</i>	vpiFile	returns the file name containing the net declaration
<i>int</i>	vpiLineNo	returns the file line number of the net declaration

1. Cont assigns and prim terms can only be accessed from scalar nets or bits of vector nets.


B.1.8 VPI reg objects



<i>int</i>	vpiType	returns vpiReg or vpiRegBit
<i>bool</i>	vpiScalar	returns true if the reg is scalar
<i>bool</i>	vpiVector	returns true if the reg is a vector
<i>int</i>	vpiSize	returns the bit size of the reg
<i>str</i>	vpiName	returns the declaration name of the reg
<i>str</i>	vpiFullName	returns the full hierarchical path name of the reg
<i>str</i>	vpiFile	returns the file name containing the reg declaration
<i>int</i>	vpiLineNo	returns the file line number of the reg declaration

1. Continuous assignments and primitive terminals can only be accessed from scalar or bit selects of regs, and are accessed across hierarchical boundaries.
2. For **vpiPortInst** and **vpiPort**, if reference handle is a bit, a handle for a port bit is returned; if reference handle is a vector, a handle for the port is returned.
3. For **vpiDriver** and **vpiLoad**, if the reference object is **vpiReg**, scalar and vector drivers and loads are returned. Use a **vpiRegBit** reference to access bit selects of drivers and loads.
4. For **vpiDriver** and **vpiLoad**, only active assigns and forces are returned.
5. The value of a reg is accessed using **vpi_get_value()** and **vpi_put_value()**.


B.1.9 VPI memory objects



<i>int</i>	vpiType	returns vpiMemory , vpiMemoryWord
<i>int</i>	vpiSize	for a memory, returns the number of words in the memory; for a memory word, returns the number of bits in a word
<i>str</i>	vpiName	returns the name of a memory or memory word
<i>str</i>	vpiFullName	returns the full hierarchical path name of a memory or memory word
<i>str</i>	vpiFile	returns the file name containing the memory declaration
<i>int</i>	vpiLineNo	returns file line number of the memory declaration

1. A memory is a 1-dimensional array of **reg** (e.g.: **reg [7:0] RAM [0:1023];**)
2. For a memory object, **vpiLeftRange** and **vpiRightRange** refer to the starting and ending address of a memory declaration.
3. For a memory word object, **vpiLeftRange** and **vpiRightRange** refer to the MSB and LSB of the word.
4. The value of a memory word is accessed using **vpi_get_value()** and **vpi_put_value()**.


B.1.10 VPI variable objects



<i>int</i>	vpiType	returns vpiIntegerVar , vpiTimeVar , vpiRealVar
<i>bool</i>	vpiArray	the variable is an array (only applies to variables)
<i>int</i>	vpiSize	returns: – the bit size of a variable – the number of words in a variable array – the word size of a var array select
<i>str</i>	vpiName	returns the name of a variable or var array select
<i>str</i>	vpiFullName	returns the full hierarchical path name of a variable or var array select
<i>str</i>	vpiFile	returns file name containing the variable definition or var array select
<i>int</i>	vpiLineNo	returns file line number of the variable definition or var array select


1. For a variable array object, **vpiLeftRange** and **vpiRightRange** refer to the starting and ending address of the variable array declaration.
2. For a var array select object, **vpiLeftRange** and **vpiRightRange** refer to the MSB and LSB of the word.
3. The value of a variable or var array select is accessed using **vpi_get_value()** and **vpi_put_value()**.

B.1.11 VPI named event objects



<i>int</i>	vpiType	returns the type as vpiNamedEvent
<i>str</i>	vpiName	returns the declaration name of the named event
<i>str</i>	vpiFullName	returns the full hierarchical path name of the named event
<i>str</i>	vpiFile	returns the file name containing the named event declaration
<i>int</i>	vpiLineNo	returns the file line number of the named event declaration




B.1.12 VPI continuous assignment objects



<i>int</i>	vpiType	returns vpiContAssign
<i>bool</i>	vpiNetDeclAssign	returns true if the cont assign is combined with a net declaration
<i>int</i>	vpiStrength0	returns the logic 0 strength of the net declaration
<i>int</i>	vpiStrength1	returns the logic 1 strength of the net declaration
<i>str</i>	vpiFile	returns the file name containing the cont assign
<i>int</i>	vpiLineNo	returns the file line number containing the cont assign

1. The delay value of a continuous assignment delay expression is accessed using **vpi_get_delays()**.


B.1.13 VPI parameter, param assign and defparam objects



<i>int</i>	vpiType	returns vpiParameter , vpiParamAssign , vpiDefParam
<i>int</i>	vpiConstType	returns the type of Verilog constant as vpiDecConst , vpiRealConst , vpiBinaryConst , vpiOctConst , vpiHexConst , vpiStringConst
<i>str</i>	vpiName	returns the name of a parameter
<i>str</i>	vpiFullName	returns the full hierarchical path name of a parameter
<i>str</i>	vpiFile	returns the file name containing the parameter, parameter assignment or defparam
<i>int</i>	vpiLineNo	returns the file line number containing the parameter, parameter assignment or defparam

1. The value of a parameter is accessed using **vpi_get_value()**. The value will be the resolved value after parameter redefinitions have been applied.
2. **vpiLhs** from a parameter assignment object returns a handle for the parameter.

B.1.14 VPI primitive, prim term and UDP objects:




<i>int</i>	vpiType	returns vpiGate , vpiSwitch , vpiUDP , vpiPrimTerm , vpiUDPDefn or vpiTableEntry
<i>int</i>	vpiPrimType	for a primitive or UDP definition, returns vpiAndPrim , vpiNandPrim , vpiNorPrim , vpiOrPrim , vpiXorPrim , vpiXnorPrim , vpiBufPrim , vpiNotPrim , vpiBufif0Prim , vpiBufif1Prim , vpiNotif0Prim , vpiNotif1Prim , vpiNmosPrim , vpiPmosPrim , vpiCmosPrim , vpiRnmosPrim , vpiRpmosPrim , vpiRcmosPrim , vpiRtranPrim , vpiRtranif0Prim , vpiRtranif1Prim , vpiTranPrim , vpiTranif0Prim , vpiTranif1Prim , vpiPullupPrim , vpiPulldownPrim , vpiSeqPrim (UDP), vpiCombPrim (UDP)
<i>bool</i>	vpiProtected	for UDP definitions, returns true if the source code is protected
<i>int</i>	vpiSize	returns the number of inputs for a primitive or UDP definition; returns the number of symbol entries for a table entry
<i>int</i>	vpiStrength0	returns the logic 0 strength of the primitive
<i>int</i>	vpiStrength1	returns the logic 1 strength of the primitive
<i>int</i>	vpiIndex	returns the index number or a primitive terminal
<i>int</i>	vpiDirection	for a primitive terminal returns the direction as vpiInput , vpiOutput , vpiInout , vpiMixedIO , vpiNoDirection
<i>str</i>	vpiName	returns the primitive instance name
<i>str</i>	vpiFullName	returns the primitive full hierarchical path name
<i>str</i>	vpiDefName	returns the primitive or UDP definition name

continued on next page

<i>str</i>	vpiFile	returns the file name containing the object
<i>int</i>	vpiLineNo	returns the file line number containing the object

1. Primitive delays are accessed using **vpi_get_delays()** and **vpi_put_delays()**.
2. The logic value of a primitive terminal is accessed using **vpi_get_value()**.
3. The logic value of a sequential UDP primitive may be set using **vpi_put_value()**.
4. Table entry values can be accessed using **vpi_get_value()**; must read the values as a string or vector (ASCII values).


B.1.15 VPI user-defined system task/function instance objects



<i>int</i>	vpiType	returns vpiUserSystf
------------	----------------	-----------------------------

1. **vpi_iterate(vpiUserSystf, NULL)** will access all user-defined system task/function instances in a simulation.
2. The user-defined system task/function registration information can be obtained with **vpi_get_systf_info()**.


B.1.16 VPI task/function declaration objects



<i>int</i>	vpiType	returns vpiTask or vpiFunction
<i>str</i>	vpiFile	returns the file name containing the task or function declaration
<i>int</i>	vpiLineNo	returns the file line number containing the task or function declaration

1. Refer to B.1.17 on page 697 for the object diagram for calls to a task or function.
2. A Verilog function contains an object with the same name, size and type as the function (this is the object to which the output of the function is assigned).
3. The constant **vpiTaskFunc** is missing in the IEEE 1364-1995 vpi_user.h file.
4. The proposed IEEE 1364-1999 standard adds a new property: *int vpiFuncType* returns **vpiIntFunc**, **vpiTimeFunc**, **vpiRealFunc** or **vpiSizedFunc**.


B.1.17 VPI task/function call objects



<i>int</i>	vpiType	returns vpiTaskCall , vpiFuncCall , vpiSysTaskCall , vpiSysFuncCall
<i>int</i>	vpiSysFuncType	if vpiType is vpiSysFuncCall , returns vpiSysFuncInt , vpiSysFuncTime , vpiSysFuncReal , vpiSysFuncSized
<i>bool</i>	vpiUserDefn	returns true if the system task/function is user-defined (a PLI system task or function)
<i>str</i>	vpiName	returns the task/function name
<i>str</i>	vpiFile	returns the file name containing the task/function call
<i>int</i>	vpiLineNo	returns the file line number containing the task/function call


1. **vpi_get_value()** returns the current value of a function call or system function call.
2. **vpi_put_value()** writes the return value of a system function call into simulation.
3. **vpi_handle(vpiSysTfCall, NULL)** accesses the system task/function call which invoked an application.
4. The proposed IEEE 1364-1999 standard changes the names of the constants involved with task/functions:
 - **vpiSysFuncType** changes to **vpiFuncType** — the new property will apply to both functions and system functions
 - **vpiSysFuncInt** changes to **vpiIntFunc**
 - **vpiSysFuncTime** changes to **vpiTimeFunc**
 - **vpiSysFuncReal** changes to **vpiRealFunc**
 - **vpiSysFuncSized** changes to **vpiSizedFunc**.

B.1.18 VPI procedure, procedural block, procedural statement objects




	vpiType	for procedure, returns <code>vpiInitial</code> or <code>vpiAlways</code> for procedural block, returns <code>vpiBegin</code> , <code>vpiNamedBegin</code> , <code>vpiFork</code> , <code>vpiNamedFork</code>
<i>int</i>		for procedural statement, returns <code>vpiAssignment</code> , <code>vpiDelayControl</code> , <code>vpiEventControl</code> , <code>vpiWhile</code> , <code>vpiRepeat</code> , <code>vpiWait</code> , <code>vpiFor</code> , <code>vpiForever</code> , <code>vpiIf</code> , <code>vpiIfElse</code> , <code>vpiCase</code> , <code>vpiAssignStmt</code> , <code>vpiDeassign</code> , <code>vpiForce</code> , <code>vpiRelease</code> , <code>vpiDisable</code> , <code>vpiEventStmt</code> , <code>vpiFuncCall</code> , <code>vpiTaskCall</code> , <code>vpiNullStmt</code>
<i>str</i>	vpiFile	returns the file name containing the object
<i>int</i>	vpiLineNo	returns the file line number containing the object

B.1.19 VPI assignment statement objects



<i>int</i>	vpiType	returns vpiAssignment
<i>bool</i>	vpiBlocking	returns true if the assignment is blocking (=)
<i>str</i>	vpiFile	returns the file name containing the assignment
<i>int</i>	vpiLineNo	returns the file line number containing the assignment


B.1.20 VPI delay, event, repeat controls and event trigger objects



<i>int</i>	vpiType	returns vpiDelayControl , vpiEventControl , vpiRepeatControl , vpiEventStmt
<i>str</i>	vpiFile	returns the file name containing the control
<i>int</i>	vpiLineNo	return the file line number containing the control




1. The delay value of a delay control expression is accessed using **vpi_get_delays()**.

B.1.21 VPI for, while, repeat, wait, forever statement objects



<i>int</i>	vpiType	returns vpiFor , vpiWhile , vpiRepeat , vpiWait , vpiForever
<i>str</i>	vpiFile	file name containing the statement
<i>int</i>	vpiLineNo	file line number containing the statement




B.1.22 VPI if, if-else, case statement objects



<i>int</i>	vpiType	returns vpiIf , vpiIfElse , vpiCase , vpiCaseItem
<i>int</i>	vpiCaseType	returns vpiCaseExact , vpiCaseX , vpiCaseZ
<i>str</i>	vpiFile	returns the file name containing the statement
<i>int</i>	vpiLineNo	returns the file line number containing the statement


1. The case item groups all case conditions which branch to the same statement.
2. **vpi_iterate(vpiExpr, <case_item_handle>)** returns NULL for the default case item since there is no expression for the default case.

B.1.23 VPI assign, deassign, force, release, disable, null statement objects



<i>int</i>	vpiType	returns vpiAssignStmt , vpiDeassign , vpiForce , vpiRelease , vpiDisable , vpiNullStmt
<i>str</i>	vpiFile	returns the file name containing the statement
<i>int</i>	vpiLineNo	returns the file line number containing the statement


B.1.24 VPI specify block specparam objects



<i>int</i>	vpiType	returns vpiSpecParam
<i>int</i>	vpiConstType	returns the type of Verilog constant as vpiDecConst , vpiRealConst , vpiBinaryConst , vpiOctConst , vpiHexConst , vpiStringConst
<i>str</i>	vpiName	returns the name of a specparam
<i>str</i>	vpiFullName	returns the full hierarchical path name of a specparam
<i>str</i>	vpiFile	returns the file name containing the specparam
<i>int</i>	vpiLineNo	returns the file line number containing the specparam

1. The value of a specparam is accessed using **vpi_get_value()**.


B.1.25 VPI specify block module path objects



<i>int</i>	vpiType	returns vpiModPath , vpiPathTerm
<i>str</i>	vpiFile	returns the file name containing the mod path or path term
<i>int</i>	vpiLineNo	returns the file line number containing the mod path or path term
<i>int</i>	vpiPathType	for mod paths, returns vpiPathFull or vpiPathParallel
<i>int</i>	vpiPolarity	for mod paths, returns vpiPositive , vpiNegative or vpiUnknown
<i>int</i>	vpiDataPolarity	for mod paths, returns vpiPositive , vpiNegative or vpiUnknown
<i>int</i>	vpiDirection	for path terms, returns vpiInput , vpiOutput , vpiInout , vpiMixedIO , vpiNoDirection

1. Module paths delays are accessed using **vpi_get_delays()** and **vpi_put_delays()**.


B.1.26 VPI specify block timing check objects



<i>int</i>	vpiType	returns vpiTchk or vpiTchkTerm
<i>str</i>	vpiFile	returns the file name containing the timing check or timing terminal
<i>int</i>	vpiLineNo	returns the file line number containing the timing check or timing terminal
<i>int</i>	vpiTchkType	for timing checks, returns vpiSetup , vpiHold , vpiPeriod , vpiWidth , vpiSkew , vpiRecovery , vpiNoChange , vpiSetupHold
<i>int</i>	vpiEdge	returns vpiNoEdge , vpiEdge01 , vpiEdge10 , vpiEdge0x , vpiEdgex1 , vpiEdge1x , vpiEdgex0 , vpiPosedge , vpiNegedge , vpiAnyEdge

1. The **vpiTchkRefTerm** is the first terminal for all timing checks except \$setup, where **vpiTchkDataTerm** is the first terminal.
2. Timing check delay limit values are accessed using **vpi_get_delays()** and **vpi_put_delays()**.


B.1.27 VPI inter-module path objects



<i>int</i>	vpiType	returns vpiInterModPath
------------	----------------	--------------------------------

1. **vpi_multi(vpiInterModPath, port1, port2)** obtains inter mod path handles.
2. Inter-module path delays are accessed using **vpi_get_delays()** and **vpi_put_delays()**.

B.1.28 VPI simulation callback objects



<i>int</i>	vpiType	returns vpiCallback
------------	----------------	----------------------------

1. Information about the callback object is accessed with **vpi_get_cb_info()**.
2. **vpi_iterate(vpiCallback, NULL)** returns an iterator for all active callbacks which have been registered using **vpi_register_cb()**.


B.1.29 VPI scheduled event objects



<i>int</i>	vpiType	returns vpiSchedEvent
<i>bool</i>	vpiScheduled	returns true if the event is still scheduled (it has not transpired)

1. A scheduled event handle is returned from **vpi_register_cb()** when the event is registered.
2. **vpi_register_cb()** can cancel a scheduled event if **vpiScheduled** is true.


B.1.30 VPI simulation time queue objects



<i>int</i>	vpiType	returns vpiTimeQueue
------------	----------------	-----------------------------

1. The time value of the queue is accessed using **vpi_get_time()**.
2. The time queue objects are returned in increasing order of simulation time.
3. **vpi_iterate()** returns NULL if there is nothing left in the simulation queue.

B.1.31 VPI iterator objects



<i>int</i>	vpiType	returns vpiIterator
------------	----------------	----------------------------

1. **vpi_handle(vpiUse, <iterator_handle>)** returns a handle to the reference object which created the iterator object when **vpi_iterate(<type_constant>, <reference_handle>)** was called.
2. It is possible to have an iterator that does not have a reference object.

B.2 VPI routine definitions (listed alphabetically)

int vpi_chk_error(info)

p_vpi_error_info info pointer to an application-allocated ***s_vpi_error_info*** structure to receive error information.

Returns *false* if the previous call to a VPI routine was successful, and *true* if the call resulted in an error. If an error occurred, information about the error is retrieved into an ***s_vpi_error*** structure pointed to by ***info***. If the information is not required, a ***NULL*** can be passed to the routine.

```
typedef struct t_vpi_error_info {
    int state;          /* vpiCompile, vpiPLI, vpiRun */
    int level;          /* vpiNotice, vpiWarning, vpiError,
                           vpiSystem, vpiInternal */
    char *message;
    char *product;
    char *code;
    char *file;
    int line;
} s_vpi_error_info, *p_vpi_error_info;
```

int vpi_compare_objects(object1, object2)

vpiHandle object1 handle for an object.

vpiHandle object2 handle for an object.

Returns 1 (true) if ***object1*** and ***object2*** handles reference the same object.

int vpi_free_object(object)

vpiHandle object handle for an object.

Frees memory allocated by the PLI for the specified object. Should be called whenever ***vpi_scan()*** is terminated before the last object is found. Returns 1 (true) if successful, and 0 (false) if an error occurred.

int vpi_get(property, object)

int property constant representing an object property.

vpiHandle object handle for an object.

Returns the value associated with integer and boolean properties of an object. Boolean properties return 1 for true and 0 for false. For property ***vpiTimeUnit***, if object is ***NULL***, then the simulation time unit is returned. If an error occurs, the property ***vpiUndefined*** is returned.

void vpi_get_cb_info(object, data)

vpiHandle object handle for an object.
p_cb_data data pointer to an application-allocated *s_cb_data* structure to receive callback information.

Retrieves information about a simulation callback and places in an *s_cb_data* structure.

```
typedef struct t_cb_data {
    int reason;           /* callback reason */
    int (*cb_rtn)();     /* call routine name */
    vpiHandle obj;        /* trigger object */
    p_vpi_time *time;    /* callback time */
    p_vpi_value *value;  /* trigger object value */
    int index;            /* index of memory word or var
                           select that changed value */
    char *user_data;
} s_cb_data, *p_cb_data;
```

void vpi_get_delays(object, delay)

vpiHandle object handle for an object.
p_vpi_delay delay pointer to an application-allocated *s_vpi_delay* structure to receive delay information.

Retrieves an object's delays or pulse control values into an *s_vpi_delay* structure. The *append_flag* and *pulsere_flag* are ignored.

```
typedef struct t_vpi_delay {
    struct t_vpi_time *da; /* ptr to user allocated
                           array of delay values */
    int no_of_delays;     /* number of delay transitions */
    int time_type;        /* vpiScaledRealTime,vpiSimTime,
                           or vpiSuppressTime */
    int mtm_flag;         /* true if using min:typ:max
                           delay value sets */
    int append_flag;       /* true to append delays,
                           false to replace delays */
    int pulsere_flag;     /* true to set pulsere values*/
} s_vpi_delay, *p_vpi_delay;

typedef struct t_vpi_time {
    int type;             /* not used by vpi_get_delays() */
    unsigned int high;    /* when using vpiSimTime */
    unsigned int low;     /* when using vpiSimTime */
    double real;          /* when using vpiScaledRealTime */
} s_vpi_time, *p_vpi_time;
```

char *vpi_get_str(property, object)

int property constant representing an object property.

vpiHandle object handle for an object.

Returns a pointer to a string containing the value of a string properties of an object.

void vpi_get_systf_info(object, data)

vpiHandle object handle for an object.

p_vpi_systf_data data pointer to an application-allocated *s_vpi_systf_data* structure to receive data information.

Retrieves information about a PLI system task/function callback and places in an *s_vpi_systf_data* structure.

```
typedef struct t_vpi_systf_data {
    int type;           /* vpiSysTask, vpiSysFunc */
    int subtype;        /* vpiSysFuncInt, vpiSysFuncReal,
                           vpiSysFuncTime, vpiSysFuncSized */
    char *tfname;       /* quoted task/function name,
                           first character must be $ */
    int (*calltf)();    /* unquoted name of C routine */
    int (*compiletf)(); /* unquoted name of C routine */
    int (*sizetf)();    /* unquoted name of C routine,
                           only used with vpiSysFuncSized
                           callbacks */
    char *user_data;    /* returned with callback */
} s_vpi_systf_data, *p_vpi_systf_data;
```

Note: The proposed IEEE 1364-1999 standard changes the names of the constants involved with task/functions: **vpiSysFuncInt** changes to **vpiIntFunc**, **vpiSysFuncTime** changes to **vpiTimeFunc**, **vpiSysFuncReal** changes to **vpiRealFunc**, and **vpiSysFuncSized** changes to **vpiSizedFunc**.

void vpi_get_time(object, time)

vpiHandle object handle for an object.

p_vpi_time time pointer to an application-allocated *s_vpi_time* structure to receive time information.

Retrieves simulation time and places in an *s_vpi_time* structure. Time is scaled to the timescale of module containing the object. If object is *NULL*, time is returned in simulation time units.

```
typedef struct t_vpi_time {
    int type;           /* vpiScaledRealTime, vpiSimTime,
                           vpiSuppressTime */
    unsigned int high;   /* when using vpiSimTime */
    unsigned int low;    /* when using vpiSimTime */
    double real;        /* when using vpiScaledRealTime */
} s_vpi_time, *p_vpi_time;
```

void vpi_get_value(expr, value)

vpiHandle expr handle for an object.

p_vpi_value value pointer to an application-allocated *s_vpi_value* structure to receive value information.

Retrieves the simulation value of an expression into an *s_vpi_value* structure.

```

typedef struct t_vpi_value {
    int format; /* vpiBinStrVal, vpiOctStrVal, vpiDecStrVal,
                  vpiHexStrVal, vpiScalarVal, vpiIntVal,
                  vpiRealVal, vpiStringVal, vpiVectorVal,
                  vpiStrengthVal, vpiSuppressVal, vpiTimeVal,
                  vpiObjTypeVal */

    union {
        char *str; /* string value */
        int scalar; /* vpi0, vpi1, vpiX,
                      vpiZ, vpiH, vpiL,
                      vpiDontCare */

        int integer; /* integer value */
        double real; /* real value */
        struct t_vpi_time *time; /* time value */
        struct t_vpi_vecval *vector; /* vector value */
        struct t_vpi_strengthval *strength; /* strength value */
        char *misc; /* reserved */
    } value;
} s_vpi_value, *p_vpi_value;

typedef struct t_vpi_time {
    int type; /* vpiScaledRealTime, vpiSimTime,
                vpiSuppressTime */
    unsigned int high; /* when using vpiSimTime */
    unsigned int low; /* when using vpiSimTime */
    double real; /* when using vpiScaledRealTime */
} s_vpi_time, *p_vpi_time;

typedef struct t_vpi_vecval {
    /* one aval/bval pair for each bit in vector */
    int aval, bval; /* bit encoding: a/b: 0/0==0,
                     1/0==1, 1/1==X, 0/1==Z */
} s_vpi_vecval, *p_vpi_vecval;

typedef struct t_vpi_strengthval {
    int logic; /* vpi0, vpi1, vpiX, vpiZ */
    int s0, s1; /* Logical-OR of the constants:
                  vpiSupplyDrive, vpiStrongDrive,
                  vpiPullDrive, vpiWeakDrive,
                  vpiLargeCharge, vpiMediumCharge,
                  vpiSmallCharge, vpiHiZ */
} s_vpi_strengthval, *p_vpi_strengthval;

```

int vpi_get_vlog_info(info)

p_vpi_vlog_info info pointer to an application-allocated *s_vpi_vlog_info* structure to receive invocation information.

Retrieves the simulator's invocation option information into an *s_vpi_vlog_info* structure. Returns 1 (true) if successful and 0 (false) if an error occurred.

```
typedef struct t_vpi_vlog_info {
    int      argc;
    char **argv;
    char   *product;
    char   *version;
} s_vpi_vlog_info, *p_vpi_vlog_info;
```

vpiHandle vpi_handle(type, reference)

int type constant representing an object type.

vpiHandle reference handle for an object.

Returns a handle for an object with a one-to-one relationship to a reference object (single arrows in the VPI relationship diagrams).

vpiHandle vpi_handle_by_index(parent, index)

vpiHandle parent handle for an object.

int index index number of an object.

Returns a handle for an object based on its index number within a parent object. The parent object must have a *vpiIndex* relationship.

vpiHandle vpi_handle_by_name(name, scope)

*char * name* name of an object.

vpiHandle scope handle for an object with a name scope, or *NULL*.

Returns a handle for an object using the name of the object. Objects which can be searched for are those with a *vpiFullName* property. The object is searched for starting at the specified scope, using Verilog name search rules. If scope is *NULL*, the search begins at the top level module.

vpiHandle vpi_handle_multi(type, reference1, reference2)

int type constant of *vpiInterModPath*.

vpiHandle reference1 handle for an output or inout port.

vpiHandle reference2 handle for an input or inout port.

Returns a handle for a module inter-connect path. The ports must be of the same size, and may be at different levels of hierarchy.

vpiHandle vpi_iterate(type, reference)

int type constant representing an object type.

vpiHandle reference handle for an object.

Returns a handle for a *vpiIterator* for objects with a one-to-many relationship to a reference object (double arrows in the VPI relationship diagrams). The *vpiIterator* handle can be passed to *vpi_scan()* to traverse all objects found.

unsigned int vpi_mcd_close(mcd)

unsigned int mcd multi-channel descriptor representing open files

Closes one or more files that were opened with *vpi_mcd_open()*. Returns 0 if successful, and the mcd of files not closed if an error occurs.

char *vpi_mcd_name(mcd)

unsigned int mcd multi-channel descriptor representing an open file

Returns a pointer to a string containing the name of an open file. The **mcd** must point to a single open file.

unsigned int vpi_mcd_open(file_name)

*char *file_name* name of a file to be opened.

Opens a file for writing and returns a multi-channel descriptor number (mcd). Returns 0 if an error occurred.

int vpi_mcd_printf(mcd, format, arg1,...argn)

int mcd multi-channel descriptor of open files.

*char *format* quoted character string of formatted message.

arg1...argn arguments to formatted message string.

Prints a formatted message to one or more open files. Uses a format string with formatting controls similar to C printf. The **mcd** is generated by *vpi_mcd_open()*. The following **mcd** values are pre-defined: 1 is stdout, 2 is stderr, 3 is simulator's current output log file. Returns the number of characters written, or *EOF* if an error occurred.

int vpi_printf(format, arg1,...argn)

*char *format* quoted character string of formatted message.

arg1...argn arguments to formatted message string.

Prints a formatted message to the simulation output channel and the current simulation output log file. Uses one or more format strings with formatting controls similar to C printf. Returns the number of characters written, or *EOF* if an error occurred.

void vpi_put_delays(object, delay)

vpiHandle object handle for an object
p_vpi_delay delay pointer to an application-allocated *s_vpi_delay* structure containing delay information.

Deposits delays or pulse control values stored in an *s_vpi_delay* structure onto an object.

```
typedef struct t_vpi_delay {
    struct t_vpi_time *da; /* ptr to user allocated
                           array of delay values */
    int no_of_delays;     /* number of delays */
    int time_type;        /* vpiScaledRealTime, vpiSimTime,
                           or vpiSuppressTime */
    int mtm_flag;         /* true if using min:typ:max
                           delay value sets */
    int append_flag;      /* true to append delays,
                           false to replace delays */
    int pulsere_flag;     /* true to set pulsere values */
} s_vpi_delay, *p_vpi_delay;

typedef struct t_vpi_time {
    int type;             /* vpiScaledRealTime, vpiSimTime,
                           vpiSuppressTime */
    unsigned int high;    /* when using vpiSimTime */
    unsigned int low;     /* when using vpiSimTime */
    double real;          /* when using vpiScaledRealTime */
} s_vpi_time, *p_vpi_time;
```

vpiHandle vpi_put_value(object, value, time, flag)

vpiHandle object handle for an object
p_vpi_value value pointer to an application-allocated *s_vpi_value* structure containing value information.
p_vpi_time time pointer to an application-allocated *s_vpi_time* structure containing the propagation delay information.
int flag constant representing the delay mode.

Deposits logic values stored in an *s_vpi_value* structure onto an object. The deposit is scheduled after a simulation time stored in an *s_vpi_value* structure, using event scheduling indicated by *flag*:

- A *flag* of *vpiNoDelay*, *vpiInertialDelay*, *vpiTransportDelay* or *vpiPureTransportDelay* is used to deposit values onto registers, memory words, system functions and sequential UDPs.
- A *flag* of *vpiForceFlag* is used to force values onto nets registers, memory words, system functions and sequential UDPs.
- A *flag* of *vpiReleaseFlag* is used to release forced values.

Returns a handle for type *vpiSchedEvent* if *flag* is ORed with *vpiReturnEvent*, otherwise *NULL* is returned. The scheduled event can be canceled by calling *vpi_put_value()* with the *vpiReturnEvent* handle and a *flag* of *vpiCancelFlag*.

```

typedef struct t_vpi_value {
    int format; /* vpiBinStrVal, vpiOctStrVal, vpiDecStrVal,
                  vpiHexStrVal, vpiScalarVal, vpiIntVal,
                  vpiRealVal, vpiStringVal, vpiVectorVal,
                  vpiStrengthVal, vpiSuppressVal, vpiTimeVal,
                  vpiObjTypeVal */

    union {
        char *str;                                /* string value */
        int scalar;                               /* vpi0, vpi1, vpiX,
                                                   vpiZ, vpiH, vpiL,
                                                   vpiDontCare */

        int integer;                            /* integer value */
        double real;                             /* real value */
        struct t_vpi_time *time;                /* time value */
        struct t_vpi_vecval *vector;             /* vector value*/
        struct t_vpi_strengthval *strength;    /* strength value */
        char *misc;                             /* reserved */
    } value;
} s_vpi_value, *p_vpi_value;

typedef struct t_vpi_time {
    int type;           /* vpiScaledRealTime, vpiSimTime,
                           vpiSuppressTime */
    unsigned int high; /* when using vpiSimTime */
    unsigned int low;  /* when using vpiSimTime */
    double real;       /* when using vpiScaledRealTime */
} s_vpi_time, *p_vpi_time;

typedef struct t_vpi_vecval {
    /* one aval/bval pair for each bit in vector */
    int aval, bval; /* bit encoding: a/b: 0/0==0,
                      1/0==1, 1/1==X, 0/1==Z */
} s_vpi_vecval, *p_vpi_vecval;

typedef struct t_vpi_strengthval {
    int logic;          /* vpi0, vpi1, vpiX, vpiZ */
    int s0, s1;         /* Logical-OR of the constants:
                           vpiSupplyDrive, vpiStrongDrive,
                           vpiPullDrive, vpiWeakDrive,
                           vpiLargeCharge, vpiMediumCharge,
                           vpiSmallCharge, vpiHiZ */
} s_vpi_strengthval, *p_vpi_strengthval;

```

vpiHandle vpi_register_cb(data)

p_cb_data data pointer to the *s_cb_data* structure containing callback information.

Registers a simulation callback to a PLI application for specific reasons during simulation. Returns a handle for the callback object. Callbacks can be scheduled for:

- Event related callbacks: *cbValueChange*, *cbStmt*, *cbForce*, *cbRelease*.
- Time related callbacks: *cbAtStartOfSimTime*, *cbReadWriteSynch*, *cbReadOnlySynch*, *cbNextSimTime*, *cbAfterDelay*.
- Action related callbacks: *cbEndOfCompile*, *cbStartOfSimulation*, *cbEndOfSimulation*, *cbError*, *cbTchkViolation*.
- Simulator feature related callbacks—a simulator may have special callbacks, such as: *cbStartOfSave*, *cbEndOfSave*, *cbStartOfRestart*, *cbEndOfRestart*, *cbEnterInteractive*, *cbExitInteractive*, *cbInteractiveScopeChange*, *cbUnresolvedSystf*, etc.

```

typedef struct t_cb_data {
    int reason;           /* callback reason */
    int (*cb_rtn)();      /* call routine */
    vpiHandle obj;        /* trigger object */
    p_vpi_time *time;    /* callback time */
    p_vpi_value *value;  /* trigger object value */
    int index;            /* index of the memory word
                           or var select which
                           changed value */
    char *user_data;      /* passed to callback */
} s_cb_data, *p_cb_data;

typedef struct t_vpi_time {
    int type;             /* vpiScaledRealTime, vpiSimTime,
                           vpiSuppressTime */
    unsigned int high;    /* when using vpiSimTime */
    unsigned int low;     /* when using vpiSimTime */
    double real;          /* when using vpiScaledRealTime */
} s_vpi_time, *p_vpi_time;

typedef struct t_vpi_vecval {
    /* one aval/bval pair for each bit in vector */
    int aval, bval;        /* bit encoding: a/b: 0/0==0,
                           1/0==1, 1/1==X, 0/1==Z */
} s_vpi_vecval, *p_vpi_vecval;

```

vpiHandle vpi_register_systf(data)

p_vpi_systf_data data pointer to the *s_vpi_systf_data* structure containing system task/function information.

Registers a PLI system task/function \$ name and associates the name with PLI C routines. Returns a handle for the callback object. Callbacks can be scheduled for:

- Applications associated as *compiletf* are called when the system task or function name is encountered during compilation or load.
- Applications associated as *sizetf* are called when a system function name is encountered during compilation or load and the function subtype is *vpiFuncSized*.
- Applications associated as *calltf* are called when the system function name is encountered during simulation.

```
typedef struct t_vpi_systf_data {
    int type;          /* vpiSysTask, vpiSysFunc */
    int subtype;       /* vpiSysFuncInt, vpiSysFuncReal,
                        vpiSysFuncTime, vpiSysFuncSized */
    char *tfname;      /* quoted task/function name,
                        first character must be $ */
    int (*calltf)();   /* unquoted name of C routine */
    int (*compiletf)(); /* unquoted name of C routine */
    int (*sizetf)();   /* unquoted name of C routine,
                        only used with vpiSysFuncSized
                        callbacks */
    char *user_data;   /* returned with callback */
} s_vpi_systf_data, *p_vpi_systf_data;
```

Note: The proposed IEEE 1364-1999 standard changes the names of the constants involved with task/functions: **vpiSysFuncInt** changes to **vpiIntFunc**, **vpiSysFuncTime** changes to **vpiTimeFunc**, **vpiSysFuncReal** changes to **vpiRealFunc**, and **vpiSysFuncSized** changes to **vpiSizedFunc**.

int vpi_remove_cb(cb_object)

vpiHandle cb_object handle for a callback object.

Removes callbacks to PLI applications which were registered with *vpi_register_cb()*. Returns 1 (true) if successful, 0 (false) if an error occurred. The callback handle is no longer valid after the callback is removed.

vpiHandle vpi_scan(iterator)

vpiHandle iterator handle for an iterator object.

Traverses all objects pointed to by an iterator object. The iterator object is obtained by calling *vpi_iterate()* for a specific object type. Returns NULL after the last object pointed to by the iterator has been returned. The iterator handle is no longer valid after *vpi_scan()* returns NULL. Note: It is an error to call *vpi_scan()* if the iterator is invalid or NULL

APPENDIX C

The IEEE 1364-1995 TF Routine Library

Appendix C lists the definitions of all TF routines contained in the IEEE 1364-1995 library. The TF library provides access to the arguments of user-defined system tasks and system functions, as well as a number of utility routines such as the ability to print messages and to control Verilog simulations. Note that the TF library refers to the arguments of system tasks and system functions as *parameters*. This should not be confused with the Verilog parameter keyword.

Many of the TF routines have two versions:

- A version which defaults to working with the instance of the system task or system function which called the PLI application.
- A version which requires as one of its inputs, a pointer to an instance of a system task or system function. The instance pointer is obtained using the *tf_getinstance()* routine.

In this appendix, the two versions of a TF routine are listed together, followed by a single description of the routines.

Note: This appendix is an excerpt from the booklet “Verilog PLI Quick Reference Guide”, published by Sutherland HDL, Inc., Portland, Oregon. Copyright 1995, 1998. Used with permission.

C.1 TF routine definitions (listed alphabetically)

void io_mcdprintf(mcd, format, arg1,...arg12)

int mcd multi-channel descriptor of open files.
*char *format* quoted character string of formatted message.
arg1...arg12 arguments to formatted message string.

Prints a formatted message to one or more open files. Uses formatting controls similar to C printf, with a maximum of 12 arguments. The **mcd** values are generated in simulation by the Verilog \$fopen system function, and must be passed to the PLI application as an argument to a system task or function.

void io_printf(format, arg1,...arg12)

*char *format* quoted character string of formatted message.
arg1...arg12 arguments to formatted message string.

Prints a formatted message to the simulator's output channel and output log file. Uses formatting controls similar to C printf, with a maximum of 12 arguments.

char *mc_scan_plusargs(plusarg)

*char *plusarg* name of the invocation option.

Tests to see if a string was included as a + option on the simulator product invocation commands. Returns *null* if the string does not exist. Returns a pointer to a *null string ("0")* if the string exists exactly as tested. Returns a pointer to a string with any suffix characters if the string exists with suffix characters (e.g.: if the invocation option is +size64 and **plusarg** is "size", then a pointer to the string "64" is returned. The plus sign used on the command line is not included as part of the string that is tested.

void tf_add_long(low1, high1, low2, high2)

*int *low1* pointer to lower 32 bits of first operand.
*int *high1* pointer to upper 32 bits of first operand.
int low2 lower 32 bits of second operand
int high2 upper 32 bits of second operand

Adds two 64-bit values and deposits the result back into the first operand.

void tf_asynchoff()

void tf_iasyncoff(tfinst)

*char *tfinst* pointer to an instance of a system task/function.

Disables calling of the misctf application with *reason_paramvc* with *reason_paramdrc* for the calling or specific instance of a system task/function.

```
void tf_asynchon()
void tf_iasynchon(tfinst)
```

*char * tfinst* pointer to an instance of a system task/function.

Enables asynchronous calling of the misctf application for the calling or specific instance of a system task/function with *reason_paramvc* whenever any argument to the system task/function changes value, and with *reason_paramdrc* whenever any argument changes strength.

```
void tf_clearalldelays()
void tf_iclearalldelays(tfinst)
```

*char * tfinst* pointer to an instance of a system task/function.

Clears callbacks to the misctf routine scheduled by *tf_setdelay()* or *tf_isetdelay()* for the calling or specific instance of a system task/function.

```
int tf_compare_long(low1, high1, low2, high2)
```

int low1 lower 32 bits of the first operand.

int high1 upper 32 bits of the first operand.

int low2 lower 32 bits of the second operand.

int high2 upper 32 bits of the second operand.

Compares two 64-bit integers. Returns **0** if equal, **1** if operand 1 is greater than operand 2, **-1** if operand 1 is less than operand 2.

```
int tf_copypvc_flag(n)
```

```
int tf_icopypvc_flag(n, tfinst)
```

int n index number of a system task/function argument, or **-1**.

*char * tfinst* pointer to an instance of a system task/function.

Copies the current PVC flag of argument **n** of the calling or specific instance of a system task/function to the saved PVC flag. Returns the value of the copied flag.

```
void tf_divide_long(low1, high1, low2, high2)
```

*int * low1* pointer to lower 32 bits of first operand.

*int * high1* pointer to upper 32 bits of first operand.

int low2 lower 32 bits of second operand

int high2 upper 32 bits of second operand

Divides two 64-bit values and deposits the result into the first operand.

```
void tf_dofinish()
```

Executes the same functionality as the Verilog *\$finish* task.

void tf_dostop()

Executes the same functionality as the Verilog \$stop task.

void tf_error(format, arg1,...arg5)

*char *format* quoted character string of formatted message.
arg1...arg5 arguments to formatted message string.

Prints a formatted message to the simulator's output channel and output log file. Uses formatting controls similar to C printf, with a maximum of 5 arguments.

void tf_evaluatep(n)**void tf_ievaluatep(n, tfinst)**

int n index number of a system task/function argument.
*char *tfinst* pointer to an instance of a system task/function.

Evaluates the value of argument *n* of the calling or specific instance of a system task/function and places it into an application-allocated *s_tfexprinfo* structure. The *tf_exprinfo()* routine must have been previously called to setup the structure.

p_tfexprinfo tf_exprinfo(n, info)**p_tfexprinfo tf_iexprinfo(n, info, tfinst)**

int n index number of a system task/function argument.
*p_tfexprinfo *info* pointer to application-allocated *s_tfexprinfo* structure to receive the expression information.
*char *tfinst* pointer to an instance of a system task/function.

Retrieves expression information about argument *n* of the calling or specific instance of a system task/function argument and places the information into an *s_tfexprinfo* structure pointed to by *info*. Returns the value of *info* if successful and 0 if an error occurred.

```
typedef struct t_tfexprinfo {
    short expr_type; /*tf_nullparam, tf_string, tf_READONLY,
                      tf_READONLYREAL, tf_READWRITE, tf_READWRITEREAL,
                      tf_RWBITSELECT, tf_RWPARTSELECT, tf_RWMEMSELECT */
    short padding;
    struct t_vecval *expr_value_p;
    double real_value;
    char *expr_string;
    int expr_ngroups;
    int expr_vec_size;
    int expr_sign;
    int expr_lhs_select;
    int expr_rhs_select;
} s_tfexprinfo, *p_tfexprinfo;
```

```
typedef struct t_vecval {
    int avalbits; /* aval/bval encoding: 0/0 == 0, */
    int bvalbits; /* 1/0 == 1, 0/1 == Z, 1/1 == X */
} s_vecval, *p_vecval;
```

char *tf_getcstringp(n)

char *tf_igetcstringp(n, tfinst)

int n index number of a system task/function argument.

*char *tfinst* pointer to an instance of a system task/function.

Returns the value argument **n** of the calling or specific instance of a system task/function argument as a string. Returns *null* if an error occurs.

char *tfGetInstance()

Returns a pointer to the instance of a system task/function that called the PLI application.

int tf_getlongp(highvalue, n)

int tf_igetlongp(highvalue, n, tfinst)

*int *highvalue* pointer to the upper 32-bits of the value.

int n index number of a system task/function argument.

*char *tfinst* pointer to an instance of a system task/function.

Retrieves the value of argument **n** of the calling or specific instance of a system task/function argument as a 64-bit integer. The lower 32-bits are returned, and the upper 32-bits are loaded into **highvalue**.

int tf_getlongtime(hightime)

int tf_igetlongtime(hightime, tfinst)

*int *hightime* pointer to the upper 32-bits of the time.

*char *tfinst* pointer to an instance of a system task/function.

Retrieves the current 64-bit simulation time. The lower 32-bits are returned, the upper 32-bits are loaded into **hightime**. Time is scaled to the time scale of the module containing the calling or specific instance of the system task/function.

int tf_getnextlongtime(lowtime, hightime)

*int * lowtime* pointer to the lower 32-bits of the time.
*int * hightime* pointer to the upper 32-bits of the time.

Retrieves the 64-bit simulation time of the next scheduled simulation event. The lower 32-bits are loaded into **lowtime**, the upper 32-bits are loaded into **hightime**. Time is scaled to the time scale of the module containing the instance of the system task/function. The time loaded and return value are determined by when the routine is called:

- Returns **0** if called from a misctf application that was called with *reason_rosynch*; The time of the next simulation event is retrieved.
- Returns **1** if there are no more simulation events; A time of 0 is retrieved.
- Returns **2** if not called from a misctf application that was called with *reason_rosynch*; The current simulation time is retrieved.

int tf_getp(n)***int tf_igetp(n, tfinst)***

int n index number of a system task/function argument.
*char * tfinst* pointer to an instance of a system task/function.

Returns the current value of argument **n** of the calling or specific instance of a system task/function. If the argument is an integer or a real, the value is returned as an integer. If the argument is a string, a pointer to a string with the value is returned.

int tf_getpchange(n)***int tf_igetpchange(n, tfinst)***

int n index number of a system task/function argument.
*char * tfinst* pointer to an instance of a system task/function.

Returns the number of the next system task/function argument greater than **n** that changed value at the current time step — **n** must be 0 the first time the routine is called in the current call to the PLI application. A 0 is returned if no other arguments changed value or an error occurred. At the current call to the PLI application, *tf_movepvc_flag(-1)* must first be called to store the PVC flags. *tf_asynchron()* or *tf_iasynchron()* must be active to enable PVC flags.

double tf_getrealp(n)***double tf_igetrealp(n, tfinst)***

int n index number of a system task/function argument.
*char * tfinst* pointer to an instance of a system task/function.

Returns the current value of argument **n** of the calling or specific instance of a system task/function as a double. The argument value must be an integer or a real.

`double tf_getrealtime()`

`double tf_igetrealtime(tfinst)`

`char * tfinst` pointer to an instance of a system task/function.

Retrieves the current simulation time as a real number. Time is scaled to the time scale of the module containing the calling or specific instance of a system task/function.

`int tf_gettime()`

`int tf_igettime(tfinst)`

`char * tfinst` pointer to an instance of a system task/function.

Retrieves the lower 32-bits of current simulation time as an integer. Time is scaled to the time scale of the module containing the calling or specific instance of a system task/function.

`int tf_gettimeprecision()`

`int tf_igettimeprecision(tfinst)`

`char * tfinst` pointer to an instance of a system task/function, or `null`.

Returns an integer representing the time scale precision of the module containing the calling or specific instance of a system task/function. If `tfinst` is `null`, the smallest simulation time precision of an instantiated design is returned. Time is represented as: 2==100 seconds, 1==1 s, 0==1 s, -1==100 ms, -2==10 ms, -3==1 ms,... -6==1 us,... -9==1 ns,... -12==1 ps,... -15==1 fs.

`int tf_gettimeunit()`

`int tf_igettimeunit(tfinst)`

`char * tfinst` pointer to an instance of a system task/function, or `null`.

Returns an integer representing the time scale unit of the module containing the calling or specific instance of a system task/function. If `tfinst` is `null`, the smallest simulation time precision of an instantiated design is returned. Time is represented as: 2==100 seconds, 1==10 s, 0==1 s, -1==100 ms, -2==10 ms, -3==1 ms,... -6==1 us,... -9==1 ns,... -12==1 ps,... -15==1 fs.

`char *tf_getworkarea()`

`char *tf_igetworkarea(tfinst)`

`char * tfinst` pointer to an instance of a system task/function.

Returns a pointer to the built-in work area of the calling or specific instance of a system task/function.

`void tf_long_to_real(low, high, real)`

int **low** lower (right-most) 32-bits of a 64-bit integer.

int **high** upper (left-most) 32-bits of a 64-bit integer.

double * **real** pointer to a double precision variable.

Converts a 64-bit long integer to a real number.

`char *tf_longtime_tosstr(lowtime, hightime)`

int **lowtime** lower (right-most) 32-bits of a 64-bit unsigned integer.

int **hightime** upper (left-most) 32-bits of a 64-bit unsigned integer.

Converts simulation time (64-bit unsigned integer) to a character string. Returns a pointer to the string.

`void tf_message(level, facility, code, format, arg1...arg5)`

int **level** a constant representing the error severity level. One of:
ERR_ERROR, *ERR_SYSTEM*,
ERR_INTERNAL, *ERR_WARNING*,
ERR_MESSAGE.

char * **facility** quoted character string appended to the output message.
 Must be 10 or less characters.

char * **code** quoted character string appended to the output message.
 Must be 10 or less characters.

char * **format** quoted character string of formatted message.
arg1...arg5 arguments to formatted message string.

Prints a formatted message to the simulator's output channel and output log file. Uses formatting controls similar to C printf, with a maximum of 5 arguments. If called by a checktf application during compiling or loading the Verilog code, and if severity is *ERR_ERROR*, *ERR_SYSTEM*, *ERR_INTERNAL*, compilation or load is aborted.

`char *tf_mipname()`**`char *tf_imipname(tfinst)`**

char * **tfinst** pointer to an instance of a system task/function.

Returns a pointer to a string containing the full hierarchical path name of the module containing the calling or specific instance of a system task/function.

int tf_movepvc_flag(n)

int tf_imovepvc_flag(n, tfinst)

int n index number of a system task/function argument, or -1.

*char * tfinst* pointer to an instance of a system task/function.

Moves the current PVC flag of argument **n** of the calling or specific instance of a system task/function to the saved PVC flag and clears the current flag. If **n** is -1, then all PVC flags are copied. Returns the value of the moved flag.

void tf_multiply_long(low1, high1, low2, high2)

*int * low1* pointer to lower 32 bits of first operand.

*int * high1* pointer to upper 32 bits of first operand.

int low2 lower 32 bits of second operand

int high2 upper 32 bits of second operand

Multiplies two 64-bit values and deposits the result back into the first operand.

p_tfnodeinfo tf_nodeinfo(n, info)

p_tfnodeinfo tf_inodeinfo(n, info, tfinst)

int n index number of a system task/function argument.

*p_tfnodeinfo *info* pointer to the *s_tfnodeinfo* structure containing data on a system task/function argument.

*char * tfinst* pointer to an instance of a system task/function.

Retrieves node information about argument **n** of the calling or specific instance of a system task/function argument. Places the information in an *s_tfnodeinfo* structure pointed to by **info**. Returns the value of **info** if successful and 0 if an error occurred.

(Structures are listed on the next page).

```

typedef struct t_tfnodeinfo {
    short node_type; /* tf_null_node, tf_reg_node, tf_integer_node,
                      tf_real_node, tf_time_node, tf_netvector_node,
                      tf_netscalar_node, tf_memory_node */
    short padding;
    union {
        struct t_vecval *vecval_p;
        struct t_strengthval *strengthval_p;
        char *memoryval_p;
        double *real_val_p;
    } node_value;
    char *node_symbol;
    int node_ngroups;
    int node_vec_size;
    int node_sign;
    int node_ms_index;
    int node_ls_index;
    int node_mem_size;
    int node_lhs_element;
    int node_rhs_element;
    int *node_handle;
} s_tfnodeinfo, *p_tfnodeinfo;

```

```

typedef struct t_vecval {
    int avalbits; /* aval/bval encoding: 0/0 == 0, */
    int bvalbits; /* 1/0 == 1, 0/1 == Z, 1/1 == X */
} s_vecval, *p_vecval;

```

```

typedef struct t_strengthval {
    int strength0;
    int strength1;
} s_strengthval, *p_strengthval;

```

int tf_numP()

int tf_inumP(tfinst)

*char *tfinst* pointer to an instance of a system task/function.

Returns the number of arguments to the calling or a specific system task/function.

```
void tf_propagatp(n)
void tf_ipropagatp(n, tfinst)
```

int n index number of a system task/function argument.
*char *tfinst* pointer to an instance of a system task/function.

Deposits a value onto argument **n** of the calling or specific instance of a system task/function, and propagates the value to any continuous assignments that read the argument's value. The value to be deposited is stored in a *s_tfexprinfo* structure. The *tf_exprinfo()* routine must have been previously called to allocate memory for the structure.

```
typedef struct t_tfexprinfo {
    short expr_type; /*tf_nullparam, tf_string, tf_READONLY,
                      tf_READONLYREAL, tf_READWRITE, tf_READWRITEREAL,
                      tf_RWBITSELECT, tf_RWPARTSELECT, tf_RWMEMSELECT */
    short padding;
    struct t_vecval *expr_value_p;
    double real_value;
    char *expr_string;
    int expr_ngroups;
    int expr_vec_size;
    int expr_sign;
    int expr_lhs_select;
    int expr_rhs_select;
} s_tfexprinfo, *p_tfexprinfo;
```

```
typedef struct t_vecval {
    int avalbits; /* aval/bval encoding: 0/0 == 0, */
    int bvalbits; /* 1/0 == 1, 0/1 == Z, 1/1 == X */
} s_vecval, *p_vecval;
```

```
void tf_putlongp(n, lowvalue, highvalue)
void tf_iputlongp(n, lowvalue, highvalue, tfinst)
```

int n index number of a system task/function argument, or 0.
int lowvalue lower (right-most) 32-bits of a 64-bit integer.
int highvalue upper (left-most) 32-bits of a 64-bit integer.
*char *tfinst* pointer to an instance of a system task/function.

Deposits a 64-bit integer value onto argument **n** of the calling or specific instance of a system task/function. If **n** is 0, then the value is deposited as the return of a system function.

void tf_putstr(n, value)

void tf_iputstr(n, value, tfinst)

int n index number of a system task/function argument, or 0.

int value a 32-bit integer value.

*char * tfinst* pointer to an instance of a system task/function.

Deposits a 32-bit integer value onto argument *n* of the calling or specific instance of a system task/function. If *n* is 0, then the value is deposited as the return of a system function.

void tf_putstrl(n, value)

void tf_iputstrl(n, value, tfinst)

int n index number of a system task/function argument, or 0.

double value a double precision real number.

*char * tfinst* pointer to an instance of a system task/function.

Deposits a real value onto argument *n* of the calling or specific instance of a system task/function. If *n* is 0, then the value is deposited as the return of a system function.

int tf_read_restart(blockptr, blocklength)

*char * blockptr* pointer to a block of memory.

int blocklength length of the block of memory in bytes.

Reads a block of memory that was saved with *tf_write_save()*. This routine must be called from a misctf application that was called with *reason_restart*. Returns 1 if successful and 0 if an error occurred.

void tf_real_to_long(real, low, high)

double real a double precision variable.

*int * low* pointer to a variable to receive the lower (right-most) 32-bits of a 64-bit integer.

*int * high* pointer to a variable to receive the upper (left-most) 32-bits of a 64-bit integer.

Converts a real number to a 64-bit long integer.

void tf_rosynchronize()

void tf_irosynchronize(tfinst)

*char * tfinst* pointer to an instance of a system task/function.

Schedules a callback to the misctf application for the calling or specific instance of a system task/function. The callback occurs at the end of the current simulation time step with *reason_rosynch*. The PLI is not allowed to schedule any additional events.

```
void tf_scale_longdelay(tfinst, low1, high1, low2, high2)
```

char * tfinst pointer to an instance of a system task/function.
int low1 lower 32 bits of first operand.
int high1 upper 32 bits of first operand.
int * low2 pointer to lower 32 bits of second operand.
int * high2 pointer to upper 32 bits of second operand.

Scales a 64-bit time value to the time scale of a specific instance of a system task/function. The value stored in arguments **low1** and **high1** are scaled and deposited to **low2** and **high2**.

```
void tf_scale_realdelay(tfinst, real1, real2)
```

char * tfinst pointer to an instance of a system task/function.
double real1 first operand.
double * real2 pointer to second operand.

Scales a double precision real number time value to the time scale of a specific instance of a system task/function. The value stored in **real1** is scaled and deposited to **real2**.

```
int tf_setdelay(delay)
```

```
int tf_isetdelay(delay, tfinst)
```

int delay a 32-bit time value greater than or equal to 0.
char * tfinst pointer to an instance of a system task/function.

Schedules a callback to the miscif application for the calling or specific instance of a system task/function. The callback occurs with *reason_reactivate* after the amount of time specified by **delay**. The delay value assumes the time scale of module containing the calling or specific instance of a system task/function. Returns 1 if successful and 0 if an error occurred.

```
int tf_setlongdelay(lowdelay, highdelay)
```

```
int tf_isetlongdelay(lowdelay, highdelay, tfinst)
```

int lowdelay lower (right-most) 32-bits of a 64-bit time value.
int highdelay upper (left-most) 32-bits of a 64-bit time value.
char * tfinst pointer to an instance of a system task/function.

Schedules a callback to the miscif application for the calling or specific instance of a system task/function. The callback occurs with *reason_reactivate* after the amount of time specified by **delay**. The delay value assumes the time scale of module containing the calling or specific instance of a system task/function. Returns 1 if successful and 0 if an error occurred.

```
int tf_setrealdelay(realdelay)
int tf_isetrealdelay(realdelay, tfinst)
```

double realdelay a real time value greater than or equal to 0.

*char *tfinst* pointer to an instance of a system task/function.

Schedules a callback to the miscf application for the calling or specific instance of a system task/function. The callback occurs with *reason_reactivate* after the amount of time specified by **delay**. The delay value assumes the time scale of module containing the calling or specific instance of a system task/function. Returns 1 if successful and 0 if an error occurred.

```
void tf_setworkarea(workarea)
void tf_isetworkarea(workarea, tfinst)
```

*char *workarea* pointer to a string or block of memory.

*char *tfinst* pointer to an instance of a system task/function.

Deposits a pointer value into the built-in work area of the calling or specific instance of a system task/function.

```
int tf_sizep(n)
int tf_isizep(n, tfinst)
```

int n index number of a system task/function argument.

*char *tfinst* pointer to an instance of a system task/function.

Returns the size in number of bits of argument **n** of the calling or specific instance of a system task/function.

```
char *tf_spname()
char *tf_ispname(tfcell)
```

*char *tfinst* pointer to an instance of a system task/function.

Returns a pointer to a string containing the full hierarchical path name of the scope containing the calling or specific instance of a system task/function.

int tf_strdelputp(n, length, format, value, delay, mode)

int tf_istrdelputp(n, length, format, value, delay, mode, tfinst)

<i>int</i>	n	index number of a system task/function argument.
<i>int</i>	length	number of bits to be deposited.
<i>int</i>	format	character in single quotes representing value format: ' <i>b</i> ' or ' <i>B</i> ' for binary, ' <i>o</i> ' or ' <i>O</i> ' for octal, ' <i>d</i> ' or ' <i>D</i> ' for decimal, ' <i>h</i> ' or ' <i>H</i> ' for hexadecimal.
<i>char</i>	* value	string representing the value to be deposited.
<i>int</i>	delay	32-bit time value representing delay before value is deposited.
<i>int</i>	mode	code representing delay mode: <i>0</i> for inertial, <i>1</i> for modified transport, <i>2</i> for pure transport.
<i>char</i>	* tfinst	pointer to an instance of a system task/function.

Deposits a value to argument **n** of the calling or specific instance of a system task. The value to be deposited is represented as a string, in binary, octal, decimal or hexadecimal format. The deposit is scheduled after a delay which must be greater than or equal to 0, using inertial, modified transport or pure transport event scheduling. Time is scaled to the time scale of the module containing the calling or specific instance of a system task. Returns 1 if successful and 0 if an error occurred.

char *tf_strgetp(n, format_char)

char *tf_istrgetp(n, format_char, tfinst)

<i>int</i>	n	index number of a system task/function argument.
<i>int</i>	format_char	character in single quotes representing value format: ' <i>b</i> ' or ' <i>B</i> ' for binary, ' <i>o</i> ' or ' <i>O</i> ' for octal, ' <i>d</i> ' or ' <i>D</i> ' for decimal, ' <i>h</i> ' or ' <i>H</i> ' for hexadecimal.
<i>char</i>	* tfinst	pointer to an instance of a system task/function.

Returns the value of argument **n** of the calling or specific instance of a system task/function. The value is returned as pointer to a string, in binary, octal, decimal or hexadecimal format.

char *tf_strgettime()

Retrieves the current simulation time as a string. Time is scaled to the time scale of the module containing the calling or specific instance of a system task/function.

int tf_strlongdelput(n, length, format, value, lowdelay, highdelay, mode)
int tf_istrlongdelput(n, length, format, value, lowdelay, highdelay,

mode, tfinst)

int	n	index number of a system task/function argument.
int	length	number of bits to be deposited.
int	format	character in single quotes representing value format: ' <i>b</i> ' or ' <i>B</i> ' for binary, ' <i>o</i> ' or ' <i>O</i> ' for octal, ' <i>d</i> ' or ' <i>D</i> ' for decimal, ' <i>h</i> ' or ' <i>H</i> ' for hexadecimal.
char	* value	string representing the value to be deposited.
int	lowdelay	lower (right-most) 32-bits of a 64-bit time value representing delay before value is deposited.
int	highdelay	upper (left-most) 32-bits of a 64-bit time value representing delay before value is deposited.
int	mode	code representing delay mode: <i>0</i> for inertial, <i>1</i> for modified transport, <i>2</i> for pure transport.
char	* tfinst	pointer to an instance of a system task/function.

Deposits a value to argument **n** of the calling or specific instance of a system task. The value to be deposited is represented as a string, in binary, octal, decimal or hexadecimal format. The deposit is scheduled after a delay which must be greater than or equal to 0, using inertial, modified transport or pure transport event scheduling. Time is scaled to the time scale of the module containing the calling or specific instance of a system task. Returns 1 if successful and 0 if an error occurred.

int tf_strealdelput(n, length, format, value, delay, mode)

int tf_istrrealdelput(n, length, format, value, delay, mode, tfinst)

int	n	index number of a system task/function argument.
int	length	number of bits to be deposited.
int	format	character in single quotes representing value format: ' <i>b</i> ' or ' <i>B</i> ' for binary, ' <i>o</i> ' or ' <i>O</i> ' for octal, ' <i>d</i> ' or ' <i>D</i> ' for decimal, ' <i>h</i> ' or ' <i>H</i> ' for hexadecimal.
char	* value	string representing the value to be deposited.
double	delay	real time value representing delay before value is deposited.
int	mode	code representing delay mode: <i>0</i> for inertial, <i>1</i> for modified transport, <i>2</i> for pure transport.
char	* tfinst	pointer to an instance of a system task/function.

Deposits a value to argument **n** of the calling or specific instance of a system task. The value to be deposited is represented as a string, in binary, octal, decimal or hexadecimal format. The deposit is scheduled after a delay which must be greater than or equal to 0, using inertial, modified transport or pure transport event scheduling. Time is scaled to the time scale of the module containing the calling or specific instance of a system task. Returns 1 if successful and 0 if an error occurred.

void tf_subtract_long(low1, high1, low2, high2)

*int * low1* pointer to lower 32 bits of first operand.
*int * high1* pointer to upper 32 bits of first operand.
int low2 lower 32 bits of second operand
int high2 upper 32 bits of second operand

Subtracts two 64-bit values and deposits the result back into the first operand.

void tf_synchronize()

void tf_isynchronize(tfinst)

*char * tfinst* pointer to an instance of a system task/function.

Schedules a callback to the miscf application for the calling or specific instance of a system task/function. The callback occurs at the end of the current simulation time step with *reason_synch*. The PLI is allowed to schedule any additional events at the same or a later time step.

int tf_testpvc_flag(n)

int tf_itestpvc_flag(n, tfinst)

int n index number of a system task/function argument, or -1.
*char * tfinst* pointer to an instance of a system task/function.

Returns the value of the saved PVC flag of argument *n* of the calling or specific instance of a system task/function to the saved PVC flag. If *n* is -1, a logical OR of all saved PVC flags is returned.

void tf_text(format, arg1,...arg5)

*char * format* quoted character string of formatted message.
arg1...arg5 arguments to formatted message string.

Queues a formatted message into an error buffer to be printed when *tf_message()* is called.

int tf_typep(n)

int tf_itypep(n, tfinst)

int n index number of a system task/function argument.
*char * tfinst* pointer to an instance of a system task/function.

Returns a constant representing the type of argument *n* of the calling or specific instance of a system task/function. The type is one of: *tf_nullparam*, *tf_string*, *tf_READONLY*, *tf_READWRITE*, *tf_READONLYREAL*, *tf_READWRITEREAL*. A read/write argument is a Verilog data type that is legal on the left-hand side of a procedural assignment (a register data type).

void tf_unscale_longdelay(tfinst, low1, high1, low2, high2)

*char * tfinst* pointer to an instance of a system task/function.
int low1 lower 32 bits of first operand.
int high1 upper 32 bits of first operand.
*int * low2* pointer to lower 32 bits of second operand.
*int * high2* pointer to upper 32 bits of second operand.

Converts a 64-bit time value expressed in simulation time units to the time scale of a specific instance of a system task/function. The value stored in arguments **low1** and **high1** are converted and deposited to **low2** and **high2**.

void tf_unscale_realdelay(tfinst, real1, real2)

*char * tfinst* pointer to an instance of a system task/function.
double real1 real number value of first operand.
*double * real2* pointer to real number value of second operand.
*int * high2* pointer to upper 32 bits of second operand.

Converts a real number value expressed in simulation time units to the time scale of a specific instance of a system task/function. The value stored in **real1** is converted and deposited to **real2**.

void tf_warning(format, arg1,...arg5)

*char * format* quoted character string of formatted message.
arg1...arg5 arguments to formatted message string.

Prints a formatted message to the simulator's output channel and output log file. Uses formatting controls similar to C printf, with a maximum of 5 arguments. Does not abort compilation or loading.

int tf_write_save(blockptr, blocklength)

*char * blockptr* pointer to a block of memory.
int blocklength length of the block of memory in bytes.

Causes a block of memory to be included in a data file created by the \$save() system task. The routine *tf_write_save()* may only be called by a miscf application that was called with *reason_save*. Returns non-zero if successful and 0 if an error occurred.

APPENDIX D


The IEEE 1364-1995 ACC Routine Library

This appendix lists the definitions of all ACC routines contained in the IEEE 1364-1995 library. In addition, ACC object diagrams are provided for all Verilog objects which can be accessed by ACC routines. These object diagrams do not exist in the IEEE 1364 standard, or in any previous versions of the Verilog PLI standard.


Note: This appendix is an excerpt from the booklet “Verilog PLI Quick Reference Guide”, published by Sutherland HDL, Inc., Portland, Oregon. Copyright 1995, 1998. Used with permission.

D.1 ACC Objects Relationships

D.1.1 Object diagram legend




A **single arrow** represents a **one-to-one relationship** which is traversed using the routine shown, usually an `acc_handle_*` routine. Refer to the description of the routine for the arguments required.



A **double arrow** represents a **one-to-many relationship** which is traversed using the routine shown, usually an `acc_next_*` routine. Refer to the description of the routine for the arguments required.

module

A **solid enclosure with bold font** represents an **object definition**. The diagram shows all relationships which can be traversed to and from the object.




port

A **solid enclosure with standard font** represents an **object reference**. Another diagram defines the referenced object.


variables

A **dotted enclosure with italic font** represents an **object class reference**. Another diagram defines the referenced object class.



A **small circle** represents either the **top of hierarchy** or the **current PLI hierarchy scope**. Objects are accessed using the `acc_handle_*` or `acc_next_*` routine shown.

D.1.2 ACC object data diagram for module




Related Routines:

acc_fetch_type()	returns accModule
acc_fetch_fulltype()	returns accTopModule , accModuleInstance , accCellInstance
acc_fetch_name()	returns the instance name of a module
acc_fetch_fullname()	returns the hierarchical path name of a module
acc_fetch_defname()	returns the definition name of a module
acc_fetch_delay_mode()	returns accDelayModeNone , accDelayModeZero , accDelayModeUnit , accDelayModePath , accDelayModeDistrib , accDelayModeMTM
acc_fetch_timescale_info()	returns the timescale of a module
acc_fetch_location()	returns source file name & line no. containing module instance

1. **acc_handle_parent()** does not work with scope types of **accTask**, **accFunction**, **accNamedBeginStat**, and **accNamedForkStat**.
2. In some simulators, **acc_next_modpath()**, **acc_next_tchk()** and **acc_handle_tchk()** only

D.1.3 ACC object data diagram for scope




Related Routines:

acc_fetch_type()	returns accModule , accFunction , accTask , accStatement
acc_fetch_fulltype()	returns accTopModule , accModuleInstance , accCellInstance , accFunction , accNamedBeginStat , accNamedForkStat
acc_object_of_type()	returns true for special-type of accScope
acc_fetch_name()	returns the instance name of a scope
acc_fetch_fullname()	returns the hierarchical path name of a scope
acc_fetch_location()	returns the source file name and line number containing a scope

1. Some simulators cannot access parameters declared within scope types of **accNamedBeginStat** and **accNamedForkStat** using **acc_next_parameter()**.


D.1.4 ACC object data diagram for expression



Related Routines:

acc_fetch_type()	returns accNet , accReg (or accRegister), accIntegerVar , accTimeVar , accRealVar , accConstant , accParameter , accSpecparam , accPartSelect
acc_fetch_fulltype()	returns accWire , accTri , accWand , accTriand , accWor , accTrior , accTireg , accTri0 , accTri1 , accSupply0 , accSupply1 , accReg (or accRegister), accIntegerVar , accTimeVar , accRealVar , accConstant , accRegBit , accIntegerParam , accRealParam , accStringParam

D.1.5 ACC object data diagram for ports



Related Routines:

acc_fetch_type()	returns accPort or accPortBit
acc_fetch_fulltype()	returns accScalarPort , accVectorPort , accBitSelectPort ¹ , accPartSelectPort ¹ , accConcatPort ¹
acc_fetch_name()	returns the instance name of a port or port bit
acc_fetch_fullname()	returns the hierarchical path name of a port or bit
acc_fetch_direction()	returns accInput , accOutput , accInout , accMixedIO
acc_fetch_index()	returns the port's index position in the module port list (beginning with 0 for the left-most port)
acc_fetch_size()	returns the number of bits in a port
acc_fetch_location()	returns the source file name and line number containing a port
acc_object_of_type()	returns true or false; use to verify a port is accExpandedVector before calling acc_next_bit()
acc_fetch_delays()	returns delays of a <i>scalar</i> input port or input port bit
acc_append_delays()	appends delays to a <i>scalar</i> input port or input port bit
acc_replace_delays()	replaces delays of a <i>scalar</i> input port or input port bit
acc_vcl_add()	adds a VCL flag to a <i>scalar</i> port or port bit
acc_vcl_delete()	removes a VCL flag on a <i>scalar</i> port or port bit


1. **accBitSelectPort** represents “module m1 (a[0])”,
accPartSelectPort represents “module m1 (a[15:8])”,
accConcatPort represents “module m1 (.a({n,m}))”.
2. **acc_next_bit()** requires that a true be returned from the test
acc_object_of_type(<port_handle>, accExpandedVector).
3. **acc_next_bit()** when the port fulltype is **accConcatPort** will return handles for each vector in the concatenation, not each bit of each vector.
4. **acc_handle_loconn()** does not work with ports of **accConcatPort**.
5. **acc_next_loconn()** does not work with ports of **accPartSelectPort**, **accConcatPort**.
(continued on next page)

ACC object data diagram for ports, continued

6. `acc_fetch_name()` and `acc_fetch_fullname()` do not work with ports of `accPartSelectPort` or `accConcatPort`.
7. If a port name is specified as a system task/function argument, `acc_handle_tfarg()` will return a handle for the loconn signal connected to the port.
8. The loconn and hiconn handle returned is determined by the port properties:

routine	port property	connection handle
<code>acc_handle_loconn()</code>	<i>scalar port</i>	<i>scalar connection</i>
	<i>expanded vector port</i>	<i>vector connection</i>
	<i>unexpanded vector port</i>	<i>vector connection</i>
	<i>bit select of a port</i>	<i>bit select of connection</i>
<code>acc_next_loconn()</code>	<i>scalar port</i>	<i>scalar connection</i>
	<i>expanded vector port</i>	<i>bit select of connection</i>
	<i>unexpanded vector port</i>	<i>vector connection</i>
	<i>bit select of a port</i>	<i>illegal</i>
<code>acc_next_hiconn()</code>	<i>scalar port</i>	<i>scalar connection</i>
	<i>expanded vector port</i>	<i>vector connection</i>
	<i>unexpanded vector port</i>	<i>vector connection</i>
	<i>bit select of a port</i>	<i>illegal</i>

D.1.6 ACC object data diagram for nets




Related Routines:

<code>acc_fetch_type()</code>	returns accNet
<code>acc_fetch_fulltype()</code>	returns accSupply0 , accSupply1 , accTri , accTriand , accTrior , accTrireg , accTri0 , accTri1 , accWand , accWire , accWor
<code>acc_object_of_type()</code>	returns true if an object has a special-type property; use to test for: accScalar , accVector , accCollapsedNet , accExpandedVector
<code>acc_fetch_name()</code>	returns the instance name of the net or bit
<code>acc_fetch_fullname()</code>	returns the hierarchical path name of the net or bit
<code>acc_fetch_value()</code>	returns the logic value of a net or net bit
<code>acc_set_value()</code>	sets the logic value of a net or net bit (force only)
<code>acc_fetch_size()</code>	returns the number of bits in a net
<code>acc_fetch_range()³</code>	returns the most significant bit and least significant bit declaration values of a vector net
<code>acc_fetch_location()</code>	returns the source file name and line number containing a net
<code>acc_handle_object()</code>	obtains a handle for a net using the net's name
<code>acc_handle_by_name()</code>	
<code>acc_vcl_add()⁴</code>	adds a value change link flag to a net or net bit
<code>acc_vcl_delete()</code>	removes a value change link flag from a net or bit

1. `acc_next_port()` requires `acc_object_of_type(<net_handle>, accScalar)` return true.
2. `acc_next_port()` with a net handle should return a handle for module **port**, but some simulators return a handle for a **port bit**.
3. `acc_next_bit()` and `acc_fetch_range()` require that a true be returned from `acc_object_of_type(<net_handle>, accExpandedVector)`.
4. Some simulators restrict `acc_vcl_add()` to scalar nets, unexpanded vector nets and bit selects of expanded vector nets.

D.1.7 ACC object data diagram for `regs`




Related Routines:

<code>acc_fetch_type()</code>	returns <code>accReg</code> (or <code>accRegister</code>), <code>accRegBit</code>
<code>acc_fetch_fulltype()</code>	returns <code>accReg</code> (or <code>accRegister</code>), <code>accRegBit</code>
<code>acc_object_of_type()</code>	returns true or false; use to test if a reg has a special-type property of <code>accScalar</code> or <code>accVector</code>
<code>acc_fetch_name()</code>	returns the instance name of reg or reg bit
<code>acc_fetch_fullname()</code>	returns the hierarchical path name of reg or bit
<code>acc_fetch_value()</code>	returns the logic value of a reg or reg bit
<code>acc_set_value()</code>	sets the logic value of a reg or reg bit
<code>acc_fetch_range²()</code>	returns the most significant bit and least significant bit declaration values
<code>acc_fetch_size()</code>	returns the number of bits in a reg
<code>acc_fetch_location()</code>	returns the source file name and line number containing a reg
<code>acc_handle_object()</code> <code>acc_handle_by_name()</code>	obtains a handle for a reg using the name of a reg
<code>acc_vcl_add()</code>	adds a value change link flag to a reg or reg bit
<code>acc_vcl_delete()</code>	removes a value change link flag from an object

1. `acc_next_port()` requires that
`acc_object_of_type(<reg_handle>, accScalar)` return true.

2. `acc_fetch_range()` requires
`acc_object_of_type(<reg_handle>, accVector)` return true.

D.1.8 ACC object diagrams for variables and named event




Related Routines:

acc_fetch_type()	returns accIntegerVar , accTimeVar , accRealVar , accNamedEvent
acc_fetch_fulltype()	returns accIntegerVar , accTimeVar , accRealVar , accNamedEvent
acc_object_of_type()	returns true or false; use to test if a reg has a special-type property of accScalar or accVector
acc_fetch_name()	returns the instance name of the object
acc_fetch_fullname()	returns the hierarchical path name of the object
acc_fetch_value()	returns the logic value of a variable
acc_set_value()	sets the logic value of a variable
acc_fetch_size()	returns the number of bits in a reg or variable
acc_fetch_location()	returns the source file name and line number containing a variable or named event
acc_handle_object() acc_handle_by_name()	obtains a handle for a variable or named event using the name of the object
acc_vcl_add()	adds a value change link flag to an object
acc_vcl_delete()	removes a value change link flag from an object


1. **acc_next_port()** requires that
`acc_object_of_type(<variable_handle>, accScalar)` return true.
2. **acc_fetch_size()** does not apply to real variables or named events.

D.1.9 ACC object diagrams for parameter and specparam




Related Routines:

<code>acc_fetch_type()</code>	returns <code>accParameter</code> or <code>accStatement</code> ¹
<code>acc_fetch_fulltype()</code>	returns <code>accIntegerParam</code> , <code>accRealParam</code> , <code>accStringParam</code>
<code>acc_fetch_paramtype()</code>	returns <code>accIntegerParam</code> , <code>accRealParam</code> , <code>accStringParam</code>
<code>acc_fetch_name()</code>	returns the instance name of a parameter
<code>acc_fetch_fullname()</code>	returns the hierarchical path name of a parameter
<code>acc_fetch_paramval()</code>	returns the logic value of a parameter
<code>acc_fetch_location()</code>	returns source file name and line number containing a parameter
<code>acc_handle_object()</code> <code>acc_handle_by_name()</code>	obtains a handle for a parameter using the name of the object



Related Routines:


<code>acc_fetch_type()</code>	returns <code>accSpecparam</code>
<code>acc_fetch_fulltype()</code>	returns <code>accIntegerParam</code> , <code>accRealParam</code> , <code>accStringParam</code>
<code>acc_fetch_paramtype()</code>	returns <code>accIntegerParam</code> , <code>accRealParam</code> , <code>accStringParam</code>
<code>acc_fetch_name()</code>	returns the instance name of the specparam
<code>acc_fetch_fullname()</code>	returns the hierarchical path name of the specparam
<code>acc_fetch_paramval()</code>	returns the logic value of a specparam
<code>acc_fetch_attribute()</code>	returns the logic value of a specparam attribute associated with specific objects or all objects within a module
<code>acc_fetch_attribute_int()</code>	
<code>acc_fetch_attribute_str()</code>	
<code>acc_configure()</code>	<code>accDefaultAttr0</code> configures the return value of <code>acc_fetch_attribute*()</code> routines. <code>accPathDelimStr</code> configures the delimiter string between path inputs and outputs used by the <code>acc_fetch_attribute*()</code> routines.

D.1.10 ACC object data diagram for **module path**

Related Routines:

acc_fetch_type()	returns accPath , accDataPath , accPathTerminal
acc_fetch_fulltype()	returns accModPath , accDataPath , accPathInput , accPathOutput
acc_fetch_polarity()	for a module path or data path, returns accPositive , accNegative , accUnknown
acc_fetch_edge()	for a path term, returns accNoedge , accPosedge , accNegedge , accEdge01 , accEdge10 , accEdge0x , accEdgex1 , accEdge1x , accEdgex0
acc_fetch_name()	returns a derived name of a module path
acc_fetch_fullname()	returns a derived hierarchical name of a module path
acc_fetch_delays()	returns the delay values of a module path
acc_append_delays()	adds to the delay values of a module path
acc_replace_delays()	replaces the delay values of a module path
acc_fetch_pulsere()	returns the pulse control values of a module path
acc_append_pulsere()	adds to the pulse control values of a module path
acc_replace_pulser e()	replaces the pulse control values of a module path
acc_set_pulser e()	sets the pulse control values of a module path
acc_configure()	accEnableArgs configures the arguments required by acc_handle_modpath() . accPathDelimStr sets delimiter string between path inputs and outputs used by acc_fetch_name() and acc_fetch_fullname() . accPathDelayCount sets number of delays accessed by delay & pulsere routines. accMinTypeMaxDelays sets type of delays accessed by delay routines.
acc_object_of_time	returns true or false. use with accModPathHasNone to test if a


D.1.11 ACC object data diagram for inter-module path



Related Routines:

<code>acc_fetch_type()</code>	returns accWirePath
<code>acc_fetch_fulltype()</code>	returns accIntermodPath
<code>acc_fetch_delays()</code>	returns the delay values of an inter-module path
<code>acc_replace_delays()</code>	sets the delay values of an inter-module path
<code>acc_configure()</code>	accMapToMipd sets how inter-mod path delays are mapped to module input port delays if simulator does not store inter-mod paths.


D.1.12 ACC object data diagram for timing check



Related Routines:

<code>acc_fetch_type()</code>	returns accTchk or accTchkTerminal
<code>acc_fetch_fulltype()</code>	returns accHold , accNochange , accPeriod , accRecovery , accSetup , accSetuphold , accSkew , accWidth , accTchkTerminal
<code>acc_fetch_edge()</code>	for a tchk term, returns accNoedge , accPosedge , accNegedge , accEdge01 , accEdge10 , accEdge0x , accEdgex1 , accEdge1x , accEdgex0
<code>acc_fetch_delays()</code>	returns the timing limit values of a timing check
<code>acc_append_delays()</code>	adds to the timing limit values of a timing check
<code>acc_replace_delays()</code>	sets the timing limit values of a timing check
<code>acc_configure()</code>	accEnableArgs sets arguments required by <code>acc_handle_tchk()</code>
<code>acc_fetch_location()</code>	returns source file name and line number containing timing check


D.1.13 ACC object data diagram for primitives



Related Routines:

<code>acc_fetch_type()</code>	returns accPrimitive , accTerminal
<code>acc_fetch_fulltype()</code>	returns accAndGate , accBufGate , accBufif0Gate , accBufif1Gate , accCmosGate , accNandGate , accNmosGate , accNorGate , accNotGate , accNotif0Gate , accNotif1Gate , accOrGate , accPmosGate , accPulldownGate , accPullupGate , accRcmosGate , accRnmosGate , accRpmosGate , accRtranGate , accRtranif0Gate , accRtranif1Gate , accTranGate , accTranif0Gate , accTranif1Gate , accXnorGate , accXorGate , accCombPrim (UDP), accSeqPrim (UDP), accInputTerminal , accOutputTerminal , accInoutTerminal
<code>acc_fetch_name()</code>	returns the instance name of a primitive
<code>acc_fetch_fullname()</code>	returns the hierarchical path name of the primitive
<code>acc_fetch_defname()</code>	returns the definition name of the primitive
<code>acc_fetch_direction()</code>	for a primitive terminal, returns accInput , accOutput , accInout , accMixedIO
<code>acc_fetch_index()</code>	returns the primitive terminal's index position in a primitive (beginning with 0 as the left-most terminal)
<code>acc_fetch_location()</code>	returns source file name and line number containing a primitive
<code>acc_handle_object()</code> <code>acc_handle_by_name()</code>	obtains a handle for a primitive using the primitive's instance or hierarchical name
<code>acc_fetch_delays()</code>	returns the delay values of a primitive
<code>acc_append_delays()</code>	adds to the delay values of a primitive
<code>acc_replace_delays()</code>	sets the delay values of a primitive
<code>acc_set_value()</code>	sets the logic value of a sequential UDP
<code>acc_vcl_add()</code>	adds Value Change Link flag to primitive input or inout terminal
<code>acc_vcl_delete()</code>	removes VCL flag from a primitive input or inout terminal


D.1.14 ACC object data diagram for system task/function



Related Routines:

<code>acc_fetch_type()</code>	returns accUserTask , accUserFunction , accUserRealFunction
<code>acc_fetch_fulltype()</code>	returns accUserTask , accUserFunction , accUserRealFunction
<code>acc_fetch_tfarg()</code> <code>acc_fetch_itfarg()</code>	returns the value of a system task/function argument as a double-precision value
<code>acc_fetch_tfarg_int()</code> <code>acc_fetch_itfarg_int()</code>	returns the value of a system task/function argument as an integer value
<code>acc_fetch_tfarg_str()</code> <code>acc_fetch_itfarg_str()</code>	returns the value of a system task/function argument as a pointer to a string
<code>acc_set_value()</code>	writes a value onto a system task/function argument or to the return of a system function
<code>acc_fetch_value()</code>	retrieves the current value of a system function
<code>acc_fetch_location()</code>	returns source file name and line number containing the current system task/function instance

D.1.15 ACC object data diagram for PLI scope



Related Routines:

<code>acc_set_scope()</code>	sets the PLI scope, which affects the search path for routines which search for objects by name
<code>acc_configure()</code>	accEnableArgs configures the arguments required by the <code>acc_set_scope()</code> routine

D.1.16 ACC object data diagram for interactive scope



Related Routines:

<code>acc_set_interactive_scope()</code>	sets the simulator's interactive debug scope
------------------------------------------	----------------------------------------------

D.2 ACC routine definitions (*listed alphabetically*)

bool acc_append_delays(object, d1, d2,... d12)

bool acc_append_delays(object, dset_array)

handle object handle for a primitive, module path, timing check, module input port or module input port bit select.

double d1...d12 delay values.

double *dset_array pointer to an array of delay values.

Appends delays onto an object. How delay values are passed is set by *acc_configure(accMinTypmaxDelays,...)*:

- If “*false*” (default), delay values are listed as arguments. A single delay value is used for each transition.
- If “*true*”, delay values are passed as a pointer to an array. Three delay values are used for each transition, representing a min:typ:max delay set.

The number of delays depends on the type of object:

- 2-state primitives have 2 delays or delay sets.
- 3-state primitives and module ports have 2 delays or delay sets if *acc_configure(accToHiZDelay,...)* is “*average*”, “*min*” or “*max*”, and have 3 delays or delay sets if it is “*from_user*”.
- Module paths may have 1, 2, 3, 6 or 12 delays or delay sets, depending on the setting of *acc_configure(accPathDelayCount,...)*.
- Timing checks have 1 delay or delay set; the timing check limit.

Returns 1 if successful, 0 if an error occurred.

bool acc_append_pulsere(object, r1, e1,... r12, e12)

handle object handle for a module path.

double r1... r12 pulse reject limit values.

double e1... e12 pulse error limit values.

Appends pulse control values to a module path. May have 1, 2, 3, 6 or 12 reject and error pairs, depending on the setting of *acc_configure(accPathDelayCount,...)*. Returns 1 if successful, 0 if an error occurred.

void acc_close()

Resets all ACC configuration parameters to default values; frees any memory allocated by *acc_initialize()*.

handle *acc_collect(next_routine, object, count)

*handle *next_routine* name of any ACC *next* routine except acc_next_topmod().

handle object reference object for the ACC *next* routine.

*int *count* pointer to variable to receive number of objects collected.

Returns a pointer of an array of handles found by an ACC *next* routine. Automatically allocates memory for the array, which can be de-allocated using acc_free().

bool acc_compare_handles(object1, object2)

handle object1 handle for an object.

handle object2 handle for an object.

Returns *true* if two handles refer to the same object; *false* if they do not.

bool acc_configure(config_param, value)

int config_param One of the constants listed in following table

*char *value* configuration value as a character string.

Configures specific ACC routines as shown in the following table. Returns 1 if successful, 0 if an error occurred.

config_param	value	description
accDisplayErrors	"true"	<i>default</i> —print run-time errors generated by ACC routines
	"false"	do not print run-time errors
accDisplayWarnings	"true"	print run-time warnings generated by ACC routines
	"false"	<i>default</i> —do not print run-time warnings
affects all ACC routines		
accMinTypMaxDelays	"false"	<i>default</i> —access typical delays only
	"true"	access min, typ and max delay sets
affects acc_append_delays(), acc_fetch_delays(), acc_replace_delays()		
accPathDelayCount	"1"	access single delay for all path transitions
	"2"	access rise, fall delay for each transition
	"3"	access rise, fall, toZ delays
	"6"	<i>default</i> —access 0->1, 1->0, 0->z, z->1, 1->z, z->0 delays
	"12"	access 0->1, 1->0, 0->z, z->1, 1->z, z->0, 0->x, x->1, 1->x, x->0, x->z, z->x
		affects acc_append_delays(), acc_fetch_delays(), acc_replace_delays(), acc_append_pulsere(), acc_fetch_pulsere(), acc_replace_pulsere()

config_param	value	description
accToHizzDelay	"min" "max" "average" "from_user"	set toZ delay as minimum of rise and fall set toZ delay as maximum of rise and fall set toZ delay as average of rise and fall <i>default</i> —specify toZ delay as input to ACC routine affects acc_append_delays() and acc_replace_delays() when accMinTypMaxDelays is "false"
accMapToMipd	"min" "max" "latest"	map shortest inter-module path delay to module input port delay (mipd) <i>default</i> —map longest inter-module delay to mipd map most recent inter-mod. delay to mipd affects acc_append_delays() and acc_replace_delays()
accPathDelimStr	<i>quoted string</i>	define delimiter between source and destination in path names default is "\$" affects acc_fetch_name(), acc_fetch_fullname(), acc_fetch_attribute(), acc_fetch_attribute_int(), acc_fetch_attribute_str()
accDefaultAttr0	"true" "false"	<i>default</i> return is 0 if attribute not found <i>default</i> —specify default return value as input to ACC routine affects acc_fetch_attribute(), acc_fetch_attribute_int(), acc_fetch_attribute_str()
accEnableArgs		enables certain inputs to specific ACC routines—see the description of the specific routines affected for values affects acc_handle_modpath(), acc_handle_tchk(), acc_set_scope()
accDevelopmentVersion	<i>quoted string</i>	documents version of PLI standard for which an application was developed has no affects

int acc_count(next_routine, object)

handle *next_routine name of any ACC *next* routine except acc_next_topmod()
handle object reference object for the ACC *next* routine.

Returns the number of objects found by an ACC *next* routine.

int acc_fetch_argc()

Returns the number of command line arguments given on the command line used to invoke a Verilog simulator.

char **acc_fetch_argv()

Returns a pointer to an array of character string pointers containing the command line arguments used to invoke a Verilog simulator.

double acc_fetch_attribute(object, attribute, default)

handle object handle for a named object.
*char *attribute* name of the attribute associated with the object.
double default default value to be returned if the attribute does not exist for the object.

Returns the value as a double of an object's attribute parameter or specparam. The default argument can be dropped by setting acc_configure(accDefaultAttr0, "true"). The path name delimiter can be set using acc_configure(accPathDelimStr, ...); the default delimiter is "\$".

int acc_fetch_attribute_int(object, attribute, default)

handle object handle for a named object.
*char *attribute* name of the attribute associated with the object.
int default default value to be returned if the attribute does not exist for the object.

Returns the value as an integer of an object's attribute parameter or specparam. The default argument can be dropped by setting acc_configure(accDefaultAttr0, "true"). The path name delimiter can be set using acc_configure(accPathDelimStr, ...); the default delimiter is "\$".

char *acc_fetch_attribute_str(object, attribute, default)

handle object handle for a named object.
*char *attribute* name of the attribute associated with the object.
double default default value to be returned if the attribute does not exist for the object.

Returns the value as a string of an object's attribute parameter or specparam. The default argument can be dropped by setting acc_configure(accDefaultAttr0, "true"). The path name delimiter can be set using acc_configure(accPathDelimStr, ...); the default delimiter is "\$".

char *acc_fetch_defname(object)

handle object handle for a module or primitive instance.

Returns a pointer to a character string containing the definition name of a module instance or primitive instance.

int acc_fetch_delay_mode(object)

handle object handle for a module instance.

Returns a constant indicating the delay mode of a module. The constant is one of: *accDelayModeNone*, *accDelayModeZero*, *accDelayModeUnit*, *accDelayModePath*, *accDelayModeDistrib* or *accDelayModeMTM*.

bool acc_fetch_delays(object, d1, d2, ... d12)

bool acc_fetch_delays(object, dset_array)

handle object handle for a primitive, module path, timing check, module input port, module input port bit select, or inter-module path.

double *d1...*d12 pointer to variables to receive delay values.

double *dset_array pointer to an array of delay values.

Fetches the delays of an object. How delay values are retrieved is set by *acc_configure(accMinTypmaxDelays,...)*:

- If “*false*” (default), delay values are retrieved to pointers passed as arguments. A single delay value is used for each transition.
- If “*true*”, delay values are retrieved to an array passed as a pointer. Three delay values are used for each transition, representing a min:typ:max delay set.

The number of delays depends on the type of object:

- 2-state primitives have 2 delays or delay sets.
- 3-state primitives, module ports, inter-module paths have 3 delays or delay sets.
- Module paths may have 1, 2, 3, 6 or 12 delays or delay sets, depending on the setting of *acc_configure(accPathDelayCount,...)*.
- Timing checks have one delay or delay set, representing the limit.

Returns 1 if successful, 0 if an error occurred.

int acc_fetch_direction(object)

handle object handle for a port or terminal.

Returns an integer constant representing the direction of a module port or primitive terminal. The constant returned is one of *accInput*, *accOutput*, *accInout*, or *accMixedio*.

int acc_fetch_edge(object)

handle object handle for a module path input or output terminal or a timing check terminal.

Returns an integer value representing the edge specifier of a module path or timing check terminal. The value returned is a mask of one or more of the constants: *accNoedge*, *accPosedge*, *accNegedge*, *accEdge01*, *accEdge10*, *accEdge0x*, *accEdgex1*, *accEdge1x* or *accEdgex0*.

char *acc_fetch_fullname(object)

handle object handle for an object.

Returns a pointer to a string containing the full hierarchical path name of an object.

int acc_fetch_fulltype(object)

handle object handle for an object.

Returns a constant representing the full type of an object. Refer to section 5.1 for a list of full type constants.

int acc_fetch_index(object)

handle object handle for a module port or primitive terminal.

Returns the index number of a module port or primitive terminal, starting with the left-most index as 0. A port index is its position in the module definition. A primitive index is its position in the primitive instance.

void acc_fetch_location(location, object)

p_location location pointer to an application-allocated *s_location* structure to receive the location.

handle object handle for an object

Retrieves the file name and line number in the Verilog source code for an object into an *s_location* structure.

```
typedef struct t_location {
    int line_no;
    char *filename;
} s_location, *p_location;
```

char*acc_fetch_name(object)

handle object handle for an object.

Returns a pointer to a string containing the name of an object.

int acc_fetch_paramtype(object)

handle object handle for a parameter or specparam.

Returns a constant representing the data type of a parameter or specparam. One of: *accIntegerParam*, *accRealParam*, or *accStringParam*.

double acc_fetch_paramval(object)

handle object handle for a parameter or specparam.

Returns the value of a parameter of specparam as a double.

int acc_fetch_polarity(object)

handle object handle for a module path or data path.

Returns an integer constant representing the polarity of a module path or data path. One of: *accPositive*, *accNegative*, or *accUnknown*.

int acc_fetch_precision()

Returns an integer representing the smallest simulation time unit of an instantiated design: 2==100 seconds, 1==1 s, 0==1 ms, -1==100 ms, -2 ==10 ms, -3==1 ms,... -6==1 us,... -9==1 ns,... -12==1 ps,... -15==1 fs.

bool acc_fetch_pulsere(object, r1, e1, ... r12, e12)

handle object handle for a module path.

*double *r1... *r12* pointer to variables to receive reject limit values.

*double *e1... *e12* pointer to variables to receive error limit values.

Retrieves pulse control values of a module path. Number of reject/error value pairs retrieved is set by acc_configure(accPathDelayCount,...). Returns 1 if successful, 0 if an error occurred.

int acc_fetch_range(object, msb, lsb)

handle object handle for a net or register.

*int *msb* pointer to an integer to receive the value of the most-significant (left-most) bit of a vector.

*int *lsb* pointer to an integer to receive the value of the least-significant (right-most) bit of a vector.

Retrieves the most-significant bit and least-significant bit declaration values of a net or register vector.

int acc_fetch_size(object)

handle object handle for a port, net, reg or integer.

Returns the number of bits in a port, net, reg or integer.

double acc_fetch_tfarg(n)***double acc_fetch_itfarg(n, tfinst)***

int n position number of a PLI system task/function arg.

handle tfinst handle for an instance of a PLI system task/function

Returns the value of an argument in the calling, or a specific instance, of a system task/function. Arguments are numbered from left to right beginning with 1.

int acc_fetch_tfarg_int(n)***int acc_fetch_itfarg_int(n, tfinst)***

int n position number of a PLI system task/function arg.

handle tfinst handle for an instance of a PLI system task/function

Returns the value of an argument in the calling, or a specific instance, of a system task/function. Arguments are numbered from left to right beginning with 1.

```
char *acc_fetch_tfarg_str(n)
char *acc_fetch_itfarg_str(n, tfinst)
```

int n position number of a PLI system task/function arg.
handle tfinst handle for an instance of a PLI system task/function

Returns a pointer to a string with the value of an argument in the calling, or a specific instance, of a system task/function. Arguments are numbered from left to right beginning with 1.

```
void acc_fetch_timescale_info(object, timescale)
```

handle object handle for a module instance, module definition, PLI system task/function, or *null*.

p_timescale_info timescale pointer to an application-allocated *s_timescale_info* structure to receive the timescale information.

Retrieves the timescale defined for a Verilog module into an *s_timescale_info* structure. If **object** *null*, the timescale for the active \$timeformat system task is returned. Time is represented as an integer, where: 2==100 seconds, 1==10 s, 0==1s, -1==100 ms, -2==10 ms, -3==1 ms,... -6==1 us,... -9==1 ns,... -12==1 ps,... -15==1 fs.

The timescale unit and precision are retrieved into the structure.

```
typedef struct t_timescale_info {
    short unit;
    short precision;
} s_timescale_info, *p_timescale_info;
```

```
int acc_fetch_type(object)
```

handle object handle for an object.

Returns a constant representing the type of an object. Refer to ACC object diagrams for type constant names for each object.

```
char *acc_fetch_type_str(type)
```

int type type or fulltype constant.

Returns a pointer to a string containing the name of the type or fulltype integer constant returned by *acc_fetch_type()* or *acc_fetch_full_type()*.

char *acc_fetch_value(object, format_str, value)

handle object handle for a net or register.
char *format_str character string controlling the radix of the retrieved value; must be "%b", "%o", "%d", "%h", "%v" or "%%".
p_acc_value value pointer to an application-allocated *s_acc_value* structure to receive the value as an aval/bval pair. Only used if format_str is "%%".

Retrieves the logic value of an object. If the format string is "%b", "%o", "%d", "%h" or "%v", the value is returned as a pointer to a string. If format is "%%", the value is retrieved into an *s_acc_value* structure based on the structure format field.

```
typedef struct t_setval_value {
    int format;      /* accBinStrVal, accOctStrVal,
                      accDecStrVal, accHexStrVal,
                      accScalarVal, accIntVal,
                      accRealVal, accStringVal,
                      accVectorVal */

    union {
        char *str;
        int scalar;
        int integer;
        double real;
        p_acc_vecval vector;
    } value;
} s_setval_value, *p_setval_value,
s_acc_value, *p_acc_value;

typedef struct t_acc_vecval {
    int aval;        /* bit encoding: aval/bval:
                      0/0==0, 1/0==1, 1/1==X, 0/1==Z */
} s_acc_vecval, *p_acc_vecval;
```

void acc_free(array_ptr)

handle *array_ptr pointer to an array of handles.

Frees memory allocated by *acc_collect()*.

handle acc_handle_by_name(obj_name, scope)

char *obj_name name of an object as a string.
handle scope handle for a scope, or *null*.

Returns a handle for a named object based on its name. Searches for object beginning in the scope specified and follows Verilog HDL search rules. If scope is *null*, searches for object beginning in the scope containing the calling system task/function which called the PLI application. Cannot obtain handles for ports, module paths, data paths, or inter-module paths.

handle acc_handle_calling_mod_m

Returns a handle for the module from which the PLI application was called. NOTE: This routine is defined as a macro, not a function. Therefore, it should not be called with parenthesis at the end of the name.

handle acc_handle_condition(object)

handle object handle for a module path, data path or timing check.

Returns a handle for a conditional expression of a module path, data path, or timing check terminal. Returns *null* if there is no condition or an *ifnone* condition.

handle acc_handle_conn(object)

handle object handle for a primitive terminal, path terminal or timing check terminal.

Returns a handle for the net connected to a terminal.

handle acc_handle_datapath(object)

handle object handle for a module path.

Returns a handle for the data path associated with an edge sensitive module path.

handle acc_handle_hiconn(object)

handle object handle for a scalar port or bit select of vector port.

Returns a handle for the net connected externally to a port (hierarchically higher).

handle acc_handle_interactive_scope()

Returns a handle for the interactive debug scope where a Verilog simulator is currently pointing.

handle acc_handle_loconn(object)

handle object handle for a scalar port or bit select of vector port.

Returns a handle for the net connected internally to a port (hierarchically lower).

handle acc_handle_modpath(object, src_name, dest_name, src_handle, dest_handle)

handle object handle for a module.
char * src_name name of net connected to path source, or *null*.
char * dest_name name of net connected to path destination, or *null*.
handle src_handle (optional) handle for net connected to path source.
handle dest_handle (optional) handle for net connected to path destination.

Returns a handle for a module path using either path input/output net names or net handles. *acc_configure(accEnableArgs,...)* controls which arguments are used. If *accEnableArgs* is “no_acc_handle_modpath” (default), then only names are used. If *accEnableArgs* is “acc_handle_modpath”, then names are used if passed in as arguments, and handles are used if the names are *null*.

handle acc_handle_notifier(object)

handle object handle for a timing check.

Returns a handle for the notifier register of a timing check.

handle acc_handle_object(obj_name)

char * obj_name name of an object.

Returns a handle for a named object using either a local name or a hierarchical path name. The search scope is set using *acc_set_scope()*; by default the search scope is the scope from with the PLI application was called. Verilog HDL search rules are followed for local names.

handle acc_handle_parent(object)

handle object handle for any object.

For most objects, returns the handle for the module containing an object. Exceptions: The parent of a primitive terminal is a primitive. The parent of a bit-select or part-select of a port is a port. The parent of a part-select of a vector is a vector (the parent of a bit-select of a vector is a module).

handle acc_handle_path(source, destination)

handle source handle for a scalar output or inout port, or bit-select of a vector output or inout port.
handle destination handle for a scalar input or inout port, or bit-select of a vector input or inout port.

Returns a handle for a module interconnect path (output of one module to the input of another module).

handle acc_handle_pathin(object)

handle object handle for a module path.

Returns a handle for the net connected to the first source in a module path.

handle acc_handle_pathout(object)

handle object handle for a module path.

Returns handle for the net connected to the first destination in module path.

handle acc_handle_port(object, index)

handle object handle for a module.

int index port position index.

Returns a handle for a specific module port based on the port position in the module definition. Ports are numbered from left to right starting with 0.

handle acc_handle_scope(object)

handle object handle for an object.

Returns the handle for the scope containing an object.

handle acc_handle_simulated_net(object)

handle object handle for a net.

Returns the handle for the net being simulated after equivalent nets have been collapsed.

handle acc_handle_tchk(object, type, name1, edge1, name2, edge2, conn1, conn2)

handle	object	handle for a module.
int	type	constant representing timing check type. One of: <i>accHold</i> , <i>accNochange</i> , <i>accPeriod</i> , <i>accRecovery</i> , <i>accSetup</i> , <i>accSkew</i> or <i>accWidth</i> .
char	* name1	name of net connected to the 1st tchk argument, or <i>null</i> .
int	edge1	constant representing the edge of the 1st timing check argument (constant names are listed below).
char	* name2	name of net connected to the 2nd tchk argument, or <i>null</i> .
int	edge2	edge of the 2nd tchk arg (constant names are listed below).
handle	conn1	(optional) handle for net connected to 1st tchk argument.
handle	conn2	(optional) handle for net connected to 2nd tchk argument.

Returns the handle for a timing check using a description of the check. **edge1** and **edge2** identifiers are one of the following:

- One of the constants: *accNoedge*, *accPosedge*, *accNégedge*
- List of constants separated by +: *accEdge01*, *accEdge0x*, *accEdgex1*
- List of constants separated by +: *accEdge10*, *accEdge1x*, *accEdgex0*.

The routine *acc_configure(accEnableArgs,...)* controls which arguments are used. If *accEnableArgs* is “no_acc_handle_tchk” (default), then only names are used. If *accEnableArgs* is “acc_handle_tchk”, then names are used if passed in as arguments, and handles are used if the names are *null*.

handle acc_handle_tchkarg1(object)

handle object handle for a timing check.

Returns the handle for the timing check terminal connected to the first argument of a timing check.

handle acc_handle_tchkarg2(object)

handle object handle for a timing check.

Returns the handle for the timing check terminal connected to the second argument of a timing check.

handle acc_handle_terminal(object, index)

handle object handle for a primitive.

int index primitive terminal position.

Returns the handle for a specific primitive terminal based on the terminal position in the primitive instance. Terminals are numbered from left to right starting with 0.

handle acc_handle_tfarg(n)

handle acc_handle_itfarg(n, tfinst)

int n position number of a PLI system task/function argument.

handle tfinst handle for an instance of a PLI system task/function.

Returns a handle for the object named as an argument in the calling system task/function or a specific instance of a system task/function. Arguments are numbered from left to right beginning with 1. If the argument is the name of a port, a handle for the loconn of the port (the internal signal connected to the port) is returned.

handle acc_handle_tfinst()

Returns a handle for the system task/function which called the PLI application.

bool acc_initialize()

Initializes the ACC environment. Resets configurations to default values. Returns 1 if successful and 0 if an error occurred.

handle acc_next(type_list, scope, prev_object)

int *type_list static array with list of type or fulltype constants.

handle scope handle for the scope in which to scan for objects.

handle prev_object handle for the previous object found; initially *null*.

Returns the handle for the next object of the types listed in an array within the scope specified. Type and fulltype constants must be for modules, primitives, nets, regs, variables or parameters. See the object diagrams for the constant names.

handle acc_next_bit(vector, prev_bit)

handle vector handle for a port, net or path terminal.

handle prev_bit handle for the previous bit found; initially *null*.

Returns the handle for the next bit within a port, net or path terminal. For nets, *acc_object_of_type(accExpandedVector)* must return *true* and *acc_object_of_type(accCollapsedNet)* must return *false*.

handle acc_next_cell(module, prev_cell)

handle module handle for a module.

handle prev_cell handle for the previous cell found; initially *null*.

Returns the handle for the next cell module at or below the specified scope. A cell module is a module which is flagged with the `cell_define compiler directive or was loaded using a library scan option.

handle acc_next_cell_load(net, prev_load)

handle net handle for a scalar net or bit-select of a vector net.
handle prev_load handle for the previous load found; initially *null*.

Returns the handle for the next primitive terminal of a cell load on a net. Loads that are not in cell modules are not returned. Only the first load within a cell will be returned if there are multiple loads within the cell.

handle acc_next_child(module, prev_child)

handle module handle for a module, or *null*.
handle prev_child handle for the previous child found; initially *null*.

Returns the handle for the next module instantiated within the specified module. If the module handle is *null*, the next top-level module is returned.

handle acc_next_driver(net, prev_driver)

handle net handle for a scalar net or bit-select of a vector net.
handle prev_driver handle for the previous driver found; initially *null*.

Returns the handle for the next primitive terminal driver on a net.

handle acc_next_hiconn(port, prev_hiconn)

handle port handle for a port.
handle prev_hiconn handle for the previous connection found; initially *null*.

Returns the handle for the next externally connected net (hierarchically higher) on a module port. Vectored ports are scanned beginning with the most-significant bit.

handle acc_next_input(path, prev_input)

handle path handle for a module path or data path.
handle prev_input handle for the previous input found; initially *null*.

Returns handle for the next input terminal of a mod path or source of a data path.

handle acc_next_load(net, prev_load)

handle net handle for a scalar net or bit-select of a vector net.
handle prev_load handle for the previous load found; initially *null*.

Returns the handle for the next primitive terminal load on a net. If there are multiple loads within the module, all loads will be returned .

handle acc_next_loconn(port, prev_loconn)

handle port handle for a port.

handle prev_loconn handle for the previous connection found; initially *null*.

Returns the handle for the next internally connected net (hierarchically lower) on a module port. Vectored ports are scanned beginning with the most-significant bit.

handle acc_next_modpath(module, prev_path)

handle module handle for the module in which to scan for paths.

handle prev_path handle for the previous path found; initially *null*.

Returns the handle for the next module path specified within a module.

handle acc_next_net(module, prev_net)

handle module handle for a module.

handle prev_net handle for the previous net found; initially *null*.

Returns the handle for the next net within a module. Both explicitly and implicitly declared nets are returned. Vector nets are returned as a handle for the vector.

handle acc_next_output(path, prev_output)

handle path handle for a module path or data path.

handle prev_output handle for the previous output found; initially *null*.

Returns handle for next output terminal of a mod path or destination of a data path.

handle acc_next_parameter(module, prev_param)

handle module handle for a module.

handle prev_param handle for previous parameter found; initially *null*.

Returns the handle for the next parameter within a module.

handle acc_next_port(object, prev_port)

handle object handle for a module, scalar net or bit-select of a vector net.

handle prev_port handle for the previous port found; initially *null*.

If object is a module handle, returns the handle for the next port in the module definition. If object is a net handle, returns the handle for the next port in the module definition connected to the net. (Note: some simulators return a port bit handle instead of a port handle when the object is a net).

handle acc_next_portout(module, prev_port)

handle module handle for a module.

handle prev_port handle for the previous port found; initially *null*.

Returns the handle for the next output or inout port in a module definition.

handle acc_next_primitive(module, prev_prim)

handle module handle for a module.

handle prev_prim handle for previous primitive found; initially *null*.

Returns the handle for the next primitive within a module.

handle acc_next_scope(ref_scope, prev_scope)

handle ref_scope handle for a scope.

handle prev_scope handle for the previous scope found; initially *null*.

Returns the handle for the next scope within a scope.

handle acc_next_specparam(module, prev_sparam)

handle module handle for a module.

handle prev_sparam handle for prev. specparam found; initially *null*.

Returns the handle for the next specparam within a module.

handle acc_next_tchk(module, prev_tchk)

handle module handle for a module.

handle prev_tchk handle for the previous timing check found; initially *null*.

Returns the handle for the next timing check task within a module.

handle acc_next_terminal(gate, prev_term)

handle primitive handle for a primitive.

handle prev_term handle for previous terminal found; initially *null*.

Returns the handle for the next terminal of a primitive. Terminals are returned in the order of the primitive instance, starting with terminal 0.

handle acc_next_topmod(prev_topmod)

handle prev_topmod handle for the previous top level module found; initially *null*.

Returns the handle for the next top level module within an instantiated design. `acc_next_topmod()` cannot be used as an argument to `acc_collect()` or `acc_count()`—use `acc_next_child(null)` instead.

bool acc_object_of_type(object, type)

handle object handle for an object.

int type type, fulltype or special-type property constant.

Returns *true* if an object matches a type, fulltype or special-type property, returns *false* if it does not match. Refer the ACC object diagrams for names type and fulltype constants. special-type property constants are *accScalar*, *accVector*, *accCollapsedNet*, *accExpandedVector*, *accUnexpandedVector*, *accScope* and *accModPathHasNone*.

bool acc_object_in_typelist(object, type_list)

handle object handle for an object.

*int *type_list* static array of type, fulltype and special-type property constants.

Returns *true* if object matches any of a list of type, fulltype or special-type properties, *false* if it does not match. Refer to ACC object diagrams for type and fulltype constants, and *acc_object_of_type()* for special-type constants.

int acc_product_type()

Returns a constant representing the type of simulator that called the PLI application. Product types are *accSimulator*, *accTimingAnalyzer*, *accFaultSimulator* and *accOther*.

char *acc_product_version()

Returns a string with the version of the simulator that called the PLI application.

void acc_release_object(object)

handle object handle for a module path or data path terminal.

Frees memory allocated by calls to *acc_next_input()* and *acc_next_output()* for the module path or data path.

bool acc_replace_delays(object, d1, d2,...d12)

bool acc_replace_delays(object, dset_array)

handle object handle for a primitive, module path, timing check, module input port, module input port bit select, or inter-module path.

double d1...d12 delay values.

double *dset_array pointer to an array of delay values.

Replaces the delays of an object. How delay values are passed is set by *acc_configure(accMinTypmaxDelays,...)*:

- If “*false*” (default), delay values are passed as arguments. A single delay value is used for each transition.
- If “*true*”, delay values are passed as a pointer to an array. Three delay values are used for each transition, representing a min:typ:max delay set.

The number of delays depends on the type of object:

- 2-state primitives may have 1 or 2 delays or delay sets.
- 3-state primitives, module ports and inter-module paths may have 1 or 2 delays or delay sets if *acc_configure(accToHZDelay,...)* is “*average*”, “*min*” or “*max*”, and must have 3 delays or delay sets if it is “*from_user*”.
- Module paths may have 1, 2, 3, 6 or 12 delays or delay sets, depending on the setting of *acc_configure(accPathDelayCount,...)*.
- Timing checks have one delay or delay set; the timing check limit.

Returns 1 if successful, 0 if an error occurred.

bool acc_replace_pulsere(object, r1, e1,... r12, e12)

handle object handle for a module path.

double r1... r12 pulse reject limit values.

double e1... e12 pulse error limit values.

Replaces pulse control values of a module path. The number of reject and error value pairs depends on the setting of *acc_configure(accPathDelayCount,...)*. Returns 1 if successful, 0 if an error occurred.

void acc_reset_buffer()

Resets the pointer to the ACC string buffer to its beginning. The string buffer is used by many of the ACC routines that return pointers to strings.

handle acc_set_interactive_scope()

Sets the PLI scope to the scope of the interactive debug mode of a simulator. Returns a handle for the interactive scope.

```
void acc_set_pulsere(object, r_percent, e_percent)
```

handle object handle for a module path.

double r_percent pulse reject limit.

double e_percent pulse error limit.

Sets pulse control values of a path as a percentage of the path delays.

```
char *acc_set_scope(module, name)
```

handle module handle for a module.

*char *name* (optional) name of a module.

Sets the scope and search rules for acc_handle_object(), based on the configuration of acc_configure(*accEnableArgs*,...):

- If “*no_acc_set_scope*” (default), and **module** is a valid module handle, then the PLI scope is set to the level of the module handle.
- If “*no_acc_set_scope*” (default), and **module** is *null*, then the PLI scope is set to the first top-level module.
- If “*acc_set_scope*”, and **module** is a valid module handle, then the PLI scope is set to the level of the module handle.
- If “*acc_set_scope*”, and **module** is *null*, then the PLI scope is set to the level of the module name.
- If “*acc_set_scope*”, and **module** is *null* and module name is *null*, then the PLI scope is set to the first top-level module.

Returns a pointer to a string containing the full hierarchical path name of the new PLI scope, or *null* if an error occurred.

```
int acc_set_value(object, value, delay)
```

handle object handle for a net, register or sequential UDP.
p_setval_value value pointer to an application-allocated *s_setval_value* structure containing the value to be set.
p_setval_delay delay pointer to an application-allocated *s_setval_delay* structure containing a propagation delay value.

Deposits and propagates a value on a register or sequential UDP. Forces a value on a net or register. The value and delay information are stored in structures. Returns 0 if successful and non-zero if an error occurred. The *s_setval_value* structure definition is:

```
typedef struct t_setval_value {
    int format;          /* accBinStrVal, accOctStrVal,
                           accDecStrVal, accHexStrVal,
                           accScalarVal, accIntVal,
                           accRealVal,   accStringVal,
                           accVectorVal */

    union {
        char *str;
        int scalar;
        int integer;
        double real;
        p_acc_vecval vector;
    } value;
} s_setval_value, *p_setval_value,
s_acc_value, *p_acc_value;

typedef struct t_acc_vecval {
    int aval;
    int bval;
} s_acc_vecval, *p_acc_vecval;

typedef struct t_setval_delay {
    s_acc_time time;
    int model;
} s_setval_delay, *p_setval_delay;
```

void acc_vcl_add(object, consumer, user_data, vcl_flag)

handle object handle for a net, register, event, port, primitive output terminal or primitive inout terminal.
int * consumer unquoted name of a C consumer routine.
char * user_data user-defined data value.
int vcl_flag constant: *vcl_verilog_logic*, *vcl_verilog_strength*.

Adds a value change link monitor to an object. Each time the monitored object changes value, the consumer routine is called and passed a pointer to a vc_record structure, which contains information about the change.

```

typedef struct t_vc_record {
    int vc_reason;      /* one of: logic_value_change,
                           strength_value_change,
                           vector_value_change,
                           sregister_value_change,
                           vregister_value_change,
                           integer_value_change,
                           real_value_change,
                           time_value_change,
                           event_value_change */
    int vc_hightime; /* upper 32-bits of sim. time */
    int vc_lowtime; /* lower 32-bits of sim. time */
    char *user_data; /* value passed to acc_vcl_add() */
    union {
        unsigned char logic_value; /* for logic_value_change;
                                      one of: vc10, vc11,
                                      vc1Z, vc1X */
        double real_value; /*for real_value_change */
        handle vector_handle; /*object that changed for
                               vector_value_change,
                               vregister_value_change,
                               integer_value_change,
                               time_value_change */
        s_strengths strengths_s; /* strength_value_change */
    } out_value;
} s_vc_record, *p_vc_record;

typedef struct t_strengths {
    unsigned char logic_value; /* one of: vc10, vc11,
                                vc1Z, vc1Z */
    unsigned char strength1; /* one of: vclSupply, */
    unsigned char strength2; /* vclStrong, vclPull,
                            vclLarge, vclWeak,
                            vclMedium, vclSmall,
                            vclHighZ */
} s_strengths, *p_strengths;

```

void acc_vcl_delete(*object*, *consumer*, *user_data*, *vcl_flag*)

handle object handle for an object with a VCL flag.
*int * consumer* unquoted name of a C consumer function.
*char * user_data* user-defined data value.
int vcl_flag constant: *vcl_verilog*.

Removes a value change link monitor on an object. (Note: some simulators allow the *vcl_flag* constant to be *vcl_verilog_logic* or *vcl_verilog_strength*).

char *acc_version()

Returns a pointer to a character string with the version of ACC routines.

Index

Symbols

`$<example_name>`, see Examples
\$hold **508**
\$setup **508**
\$timeformat **565**
'celldefine **499**
'endcelldefine **499**
'timescale **129, 317, 342, 563**

Numerics

64-bit math operations **323**
64-bit time operations **320**
64-bit value conversions **324**

A

ACC constants
acc0 **541**
acc1 **541**
accAndGate **747**
accAssignFlag **555**
accBinStrVal **544, 546, 757, 769**
accBitSelectPort **496, 739**
accBufGate **747**
accBufif0Gate **747**
accBufif1Gate **747**
accCellInstance **737–738**
accCmosGate **747**
accCollapsedNet **529, 741, 762, 766**
accCombPrim **747**
accConcatPort **496, 739**

accConstant **738**
accDataPath **745**
accDeassignFlag **555**
accDecStrVal **544, 546, 757, 769**
accDefaultAttr0 **477, 751**
accDelayModeDistrib **737, 752**
accDelayModeMTM **737, 752**
accDelayModeNone **737, 752**
accDelayModePath **737, 752**
accDelayModeUnit **737, 752**
accDelayModeZero **737, 752**
accDevelopmentVersion **477, 751**
accDisplayErrors **477, 750**
accDisplayWarnings **477, 750**
accEdge01 **510, 745–746, 753, 761**
accEdge0x **510, 745–746, 753, 761**
accEdge10 **510, 745–746, 753, 761**
accEdge1x **510, 745–746, 753, 761**
accEdgex0 **510, 745–746, 753, 761**
accEdgex1 **510, 745–746, 753, 761**
accEnableArgs **477, 507, 510, 520, 751, 759, 761, 768**
accExpandedVector **529, 741, 762, 766**
accFaultSimulator **766**
accForceFlag **555**
accFunction **738**
accHexStrVal **544, 546, 757, 769**
accHold **510, 746, 761**
accInertialDelay **554, 558**
accInout **739, 747, 753**
accInoutTerminal **747**
accInput **739, 747, 753**

accInputTerminal **747**
accIntegerParam **738, 744, 754**
accIntegerVar **738, 743**
accIntermodPath **746**
accIntVal **545–547, 757, 769**
accMapToMipd **478, 751**
accMinTypMaxDelays **477, 569, 576–577, 750**
accMinTypmaxDelays **749, 753, 767**
accMixedlo **739, 747, 753**
accModPath **745**
accModPathHasIfnone **529, 766**
accModule **737–738**
accModuleInstance **737–738**
accNamedBeginStat **738**
accNamedEvent **743**
accNamedForkStat **738**
accNandGate **747**
accNegative **745, 754**
accNegedge **510, 745–746, 753, 761**
accNet **738, 741**
accNmosGate **747**
accNochange **510, 746, 761**
accNoDelay **554, 558–559**
accNoedge **510, 745–746, 753, 761**
accNorGate **747**
accNotGate **747**
accNotif0Gate **747**
accNotif1Gate **747**
accOctStrVal **544, 546, 757, 769**
accOrGate **747**
accOther **766**
accOutput **739, 747, 753**
accOutputTerminal **747**
accParameter **738, 744**
accPartSelect **738**
accPartSelectPort **496, 739**
accPath **745**
accPathDelayCount **477, 569, 576, 587–588, 749–750, 753, 767**
accPathDelimStr **477, 751**
accPathInput **745**
accPathOutput **745**
accPathTerminal **745**
accPeriod **510, 746, 761**
accPmosGate **747**
accPort **739**
accPortBit **739**
accPosedge **510, 745–746, 753, 761**
accPositive **745, 754**
accPrimitive **747**
accPulldownGate **747**
accPullupGate **747**
accPureTransportDelay **554, 558**
accRcmosGate **747**
accRealParam **738, 744, 754**
accRealTime **555**
accRealVal **545, 547, 757, 769**
accRealVar **738, 743**
accRecovery **510, 746, 761**
accReg **738, 742**
accRegBit **738, 742**
accRegister **738**
accReleaseFlag **555**
accRnmosGate **747**
accRpmosGate **747**
accRtranGate **747**
accRtranif0Gate **747**
accRtranif1Gate **747**
accScalar **529, 741–743, 766**
accScalarPort **495, 739**
accScalarVal **545, 757, 769**
accScope **529, 738, 766**
accSeqPrim **747**
accSetup **510, 746, 761**
accSetuphold **746**
accSimTime **556**
accSimulator **766**
accSkew **510, 746, 761**
accSpecparam **744**
accStatement **738, 744**
accStringParam **738, 744, 754**
accStringVal **545, 548, 757, 769**
accSupply0 **738, 741**
accSupply1 **738, 741**
accTask **738**
accTchk **746**
accTchkTerminal **746**
accTerminal **747**
accTime **556**
accTimeVar **738, 743**
accTimingAnalyzer **766**
accToHiZDelay **478, 576–577, 749, 751, 767**
accTopModule **737–738**
accTranGate **747**
accTranif0Gate **747**
accTranif1Gate **747**
accTransportDelay **554, 558**
accTri **738, 741**
accTri0 **738, 741**
accTri1 **738, 741**
accTriand **738, 741**
accTrior **738, 741**
accTrireg **738, 741**
accUnExpandedVector **529**

accUnexpandedVector 766
 accUnknown 745, 754
 accUserFunction 748
 accUserRealFunction 748
 accUserTask 748
 accVector 529, 741–743, 766
 accVectorPort 495, 739
 accVectorVal 545, 548, 551, 757, 769
 accWand 738, 741
 accWidth 510, 746, 761
 accWire 738, 741
 accWirePath 746
 accWor 738, 741
 accX 541
 accXnorGate 747
 accXorGate 747
 accZ 541
 event_value_change 596
 integer_value_change 596
 logic_value_change 596
 real_value_change 596
 sregister_value_change 596
 strength_value_change 596
 time_value_change 596
 vcl_verilog 595
 vcl_verilog_logic 592–593, 770
 vcl_verilog_strength 592–593, 770
 vcl0 597
 vcl1 597
 vclHighZ 598
 vclLarge 598
 vclMedium 598
 vclPull 598
 vclSmall 598
 vclStrong 598
 vclSupply 598
 vclWeak 598
 vclX 597
 vclZ 597
 vector_value_change 596
 vregister_value_change 596
 ACC fulltype properties 526
 ACC handle routines 484
 ACC next routines 485
 ACC routines, definition 4, 30, 276, 446
 ACC special-type properties 526
 ACC type properties 526
 acc_append_delays() 477–478, 575–580, 749
 acc_append_pulsere() 587, 749
 acc_close() 475, 749
 acc_collect() 513, 750
 acc_compare_handles() 520, 750
 acc_configure() 476, 507, 510, 519, 569,
 576–578, 587, 750
 acc_count() 511, 751
 acc_error_flag 478, 538
 acc_fetch_argc() 532, 751
 acc_fetch_argv() 532, 751
 acc_fetch_attribute() 477, 584, 752
 acc_fetch_attribute_int() 584, 752
 acc_fetch_defname() 457, 752
 acc_fetch_delay_mode() 752
 acc_fetch_delays() 477–478, 568–574, 753
 acc_fetch_direction() 753
 acc_fetch_edge() 753
 acc_fetch_fullname() 457, 477, 753
 acc_fetch_fulltype() 455, 527, 754
 acc_fetch_index() 754
 acc_fetch_itfarg() 539, 755
 acc_fetch_itfarg_int() 539, 755
 acc_fetch_itfarg_str() 539, 756
 acc_fetch_location() 530, 754
 acc_fetch_name() 457, 477, 754
 acc_fetch_paramtype() 581, 754
 acc_fetch_paramval() 581, 754
 acc_fetch_polarity() 754
 acc_fetch_precision() 566, 755
 acc_fetch_pulsere() 586, 755
 acc_fetch_range() 755
 acc_fetch_size() 755
 acc_fetch_tfarg() 537, 755
 acc_fetch_tfarg_int() 537, 755
 acc_fetch_tfarg_str() 537, 756
 acc_fetch_timescale_info() 564, 756
 acc_fetch_type() 454, 526, 756
 acc_fetch_type_str() 455, 527, 756
 acc_fetch_value() 458, 541, 757
 acc_free() 514, 757
 acc_handle_by_name() 484, 516, 757
 acc_handle_calling_mod_m 468, 484, 758
 acc_handle_condition() 484, 506, 758
 acc_handle_conn() 484, 490, 506, 509, 758
 acc_handle_datapath() 484, 758
 acc_handle_hiconn() 484, 758
 acc_handle_interactive_scope() 484, 758
 acc_handle_itfarg() 484, 539, 762
 acc_handle_loconn() 484, 758
 acc_handle_modpath() 477, 484, 506–507,
 759
 acc_handle_notifier() 484, 759
 acc_handle_object() 484, 517, 759
 acc_handle_parent() 484, 759
 acc_handle_path() 484, 508, 759
 acc_handle_pathin() 484, 506, 760
 acc_handle_pathout() 484, 489, 506, 760
 acc_handle_port() 484, 760

- acc_handle_scope() **477, 484, 760**
 acc_handle_simulated_net() **484, 530, 760**
 acc_handle_tchck() **477**
 acc_handle_tchk **510, 761**
 acc_handle_tchk() **484, 509–510, 761**
 acc_handle_tchkarg1() **484, 509, 761**
 acc_handle_tchkarg2() **484, 509, 761**
 acc_handle_terminal() **484, 761**
 acc_handle_tfarg() **452, 484, 762**
 acc_handle_tfinst() **484, 539, 559, 762**
 acc_initialize() **475, 762**
 acc_next() **462, 762**
 acc_next_bit() **762**
 acc_next_cell() **500, 762**
 acc_next_cell_load() **502, 763**
 acc_next_child() **500, 763**
 acc_next_driver() **502, 763**
 acc_next_hiconn() **763**
 acc_next_input() **506–507, 763**
 acc_next_load() **502, 763**
 acc_next_loconn() **764**
 acc_next_modpath() **488, 506, 764**
 acc_next_net() **453, 764**
 acc_next_output() **489, 506–507, 764**
 acc_next_parameter() **764**
 acc_next_port() **491, 764**
 acc_next_portout() **764**
 acc_next_primitive() **765**
 acc_next_scope() **765**
 acc_next_specparam() **765**
 acc_next_tchk() **509, 765**
 acc_next_terminal() **765**
 acc_next_topmod() **765**
 acc_object_in_typelist() **528, 766**
 acc_object_of_type() **528, 766**
 acc_product_type() **766**
 acc_product_version() **766**
 acc_release_object() **507, 766**
 acc_replace_delays() **477–478, 575–580, 767**
 acc_replace_pulsere() **587, 767**
 acc_reset_buffer() **767**
 acc_set_interactive_scope() **767**
 acc_set_pulsere() **587, 768**
 acc_set_scope **520, 768**
 acc_set_scope() **517, 519, 768**
 acc_set_value() **553, 559, 769**
 acc_user.h, including **447**
 acc_vcl_add() **592, 770**
 acc_vcl_delete() **595, 771**
 acc_version() **771**
 accMinTypMaxDelays **578**
 assign **555**
 attributes **168, 583**
- Author's web page **xvii, 7**
 aval/bval logic encoding **150, 349, 372, 548–552**
 average **577**
- C**
- calltf routines, definition for TF/ACC **283**
 calltf routines, definition for VPI **37**
 cbAfterDelay **715**
 cbAtStartOfSimTime **715**
 cbEndOfCompile **715**
 cbEndOfRestart **715**
 cbEndOfSave **715**
 cbEndOfSimulation **715**
 cbEnterInteractive **715**
 cbError **715**
 cbExitInteractive **715**
 cbForce **715**
 cbInteractiveScopeChange **715**
 cbNextSimTime **715**
 cbReadOnlySynch **715**
 cbReadWriteSynch **715**
 cbRelease **715**
 cbStartOfRestart **715**
 cbStartOfSave **715**
 cbStartOfSimulation **715**
 cbStmt **715**
 cbTchkViolation **715**
 cbUnresolvedSystf **715**
 cbValueChange **715**
 cells **499**
 checktf routines, definition for TF/ACC **284**
 checktf routines, limitations **285**
 comparing handles **137**
 compiletf routines, definition for VPI **38**
 compiletf routines, limitations **39**
 configuring ACC routines **476–478**
 consumer routine, definition **591**
 cr_vlog **646**
- D**
- deassign **555**
 drivers **502**
- E**
- error handling, ACC routines **478**
 error handling, VPI routines **100**
 error limit **585**
 event_value_change **596**
 example source code, where to obtain **7**
 Examples

-
- ACC**
- \$count_loads **511**
 - \$display_all_nets **514**
 - \$func_72bit **560**
 - \$hello **266**
 - \$list_cells **501**
 - \$list_nets **543**
 - \$list_params **582**
 - \$list_path_delays **574**
 - \$list_pathout_ports **492**
 - \$list_prim_delays **570**
 - \$mipd_delays **580**
 - \$my_monitor **599**
 - \$port_info **497**
 - \$pow **291**
 - \$print_invoke_commands **534**
 - \$read_vector_value **552**
 - \$show_all_nets **459–461**
 - \$show_all_signals, version 1 **464**
 - \$show_all_signals, version 2 **469**
 - \$show_all_signals, version 3 **471**
 - \$show_value **270, 296**
 - C model, combinational logic **607–608, 611**
 - C model, latched logic **622**
 - C model, sequential logic **615**
 - C model, synchronized logic **618**
 - TF
 - \$dump_mem_ascii **375**
 - \$dump_mem_bin **384, 386**
 - \$dump_mem_hex **375**
 - \$exprinfo_test **352**
 - \$fill_mem **380**
 - \$my_monitor1 **404**
 - \$my_monitor2 **406**
 - \$my_strobe **399**
 - \$nodeinfo_test **365**
 - \$propagatep_test **357**
 - \$read_4state_value **354**
 - \$read_stimulus **402**
 - \$read_stimulus_ba **411**
 - \$read_test_vector **326**
 - \$realpow **336**
 - C model, combinational logic **423–424, 427, 430**
 - C model, latched logic **438, 440**
 - C model, sequential logic **432**
 - C model, synchronized logic **435**
- C model, with delays** **442**
- VPI**
- \$count_all_prims **126**
 - \$list_pathout_ports **121**
 - \$mipd_delays **180**
 - \$my_monitor **223**
 - \$my_strobe **209**
 - \$port_info **123**
 - \$pow **44**
 - \$read_delays **177**
 - \$read_stimulus_ba **214**
 - \$read_strengthval **157**
 - \$read_test_vector **200**
 - \$read_timeval **159**
 - \$read_vecval **154**
 - \$realpow **167**
 - \$show_all_nets **73**
 - \$show_all_signals, version 1 **80**
 - \$show_all_signals, version 2 **87**
 - \$show_all_signals, version 3 **92**
 - \$show_value **23, 25**
 - C model, combinational logic **233–234, 237**
 - C model, latched logic **251**
 - C model, sequential logic **242**
 - C model, synchronized logic **246**
 - C model, with delays **256**
 - PLIbook_get_vpiworkarea() **187**
 - PLIbook_getarg_intval_vpi() **146**
 - PLIbook_getarg_realval_vpi() **147**
 - PLIbook_getarg_stringval_vpi() **148**
 - PLIbook_numargs_vpi() **114**
 - PLIbook_set_vpiworkarea() **187**
 - exception return values, VPI routines **100**
 - exception values **479**
 - expanded vectors **496**
- F**
- f invocation option **134, 533**
 - FinSim **674**
 - force **162, 555**
 - forwref flag **295**
 - from_user **577**
 - fulltype properties, ACC **526**
- G**
- glitches **585**

H

handle routines, ACC **484**
 handles, ACC **450, 484**
 handles, VPI **60**
 hiconn (ACC) **494**
 hierarchy scope, definition **85, 466**
 highconn (VPI) **122**

I

IEEE 1364
 PLI standard **5**
 Verilog Language Reference Manual **8**
 IEEE 1364-1999, proposed enhancements
 file I/O **128**
 multi-dimensional arrays **172, 367**
 parameter changed to argument **304**
 system function types **50, 166, 634**
 VPI simulation control **46, 73, 204, 302**
 VPI type strings **68, 125**
 VPI work area **186**
 ifnone **758**
 inertial delay **162, 341, 554**
 instance specific ACC routines **538**
 instance specific TF routines **304**
 instances
 of modules **33, 279**
 of system tasks/functions **34, 112, 280, 304,**
538
 instantiated Verilog design **279**
 integer_value_change **596**
 inter-connect path **508**
 interface mechanism
 TF/ACC interface **293, 636**
 VPI interface **48, 632**
 inter-module path **109**
 invocation options **133, 322**
 io_mcdprintf() **308, 718**
 io_printf() **307–308, 718**
 iterator **63–64, 106–108**

L

Language Reference Manual, see IEEE 1364
 loads **502**
 loconn (ACC) **493**
 logic_value_change **596**
 lowconn (VPI) **122**
 LRM, see IEEE 1364

M

max **577**

mc_scan_plusargs() **322, 718**
 memory arrays **171, 367–388**
 min **577**
 miscf routines, definition for TF/ACC **287**
 ModelSim **654**
 modified transport delay **162, 341, 554**
 module path delays **504**

N

NC-Verilog **649**
 next routines, ACC **485**
 no_acc_handle_modpath **507**
 no_acc_handle_tchk **510, 761**
 no_acc_set_scope **520, 768**

O

object diagrams, ACC, how to use **480**
 object diagrams, VPI, how to use **102**
 object names, VPI **68**
 object relationships, ACC **483**
 Object relationships, VPI **61**
 object type properties, VPI **67**
 objects, VPI **57, 59**
 Open Verilog International **3**
 overloading **32, 278**
 OVI, see Open Verilog International

P

parameter **581**
 paramvc input, for miscf routines **390**
 path delays **504**
 PATHPULSE\$ **585**
 performance considerations **98, 117, 131,**
141, 149, 303, 311, 338, 392, 418,
448, 458, 473, 475, 516
 pin-to-pin delays **504**
 PLI
 definition **1**
 origin **2**
 purpose **13**
 PLI 1.0, definition **4**
 PLI 2.0, definition **4**
 PLI applications, definition **30**
 PLI applications, deifnition **276**
 PLI routines, definition **30, 276**
 Polaris **658**
 printf() **308**
 printing messages **307**
 Programming Language Interface, see PLI
 pulse **585**
 PVC flags **404**

R

real_value_change 596
reason input, for misctf routines 390
reason_paramdrc 718
reason_paramvc 718
reason_reactivate 729
reason_restart 728
reason_rosynch 722, 728
reason_save 734
reason_synch 733
 registering PLI applications
 in Avant! Polaris 658
 in Cadence NC-Verilog 649
 in Cadence Verilog-XL 641
 in Fintronic FinSim 674
 in Model Technology ModelSim 654
 in Simucad Silos III 661
 in Synopsys VCS 667
 in VeriBest VeriBest 674
 using the TF/ACC interface 293, 636
 using the VPI interface 48, 50, 633
reject limit 585
release 162, 555

S

s_acc_time structure 555, 557
s_acc_value structure 541, 544, 546, 556–557
s_acc_vecval structure 551
s_cb_data structure 193–194
s_location structure 530
s_nodeinfo structure 361
s_setval_delay structure 554–555, 557
s_setval_value structure 554
s_strengths structure 598
s_strengthal structure 363
s_tfcell structure 294, 637
s_tfexprinfo structure 347–348
s_tfnodeinfo structure 360
s_timescale_info structure 564
s_vc_record structure 595
s_vecval structure 351
s_vpi_delay structure 173–174
s_vpi_error_info structure 101
s_vpi_strengthal structure 155
s_vpi_systf_data structure 48, 633
s_vpi_time structure 71, 130, 158
s_vpi_value structure 70, 141–142, 161
s_vpi_vecval structure 152–153
s_vpi_vlog_info structure 133
scope, definition 85, 466, 519
shell module 233, 256, 419, 442, 607, 626
Silos III 661

simulation callback routines, definition 41, 184

simulation time step activity 208
sizetf routines, definition for TF/ACC 286
sizetf routines, definition for VPI 40
sizetf routines, limitations 41
sizetf routines, limitations 287
special-type properties, ACC 526
specparam 168, 581, 583
sregister_value_change 596
strength_value_change 596
string buffer, VPI 99
system task/function
 definition 16, 30–31, 264, 276–277
 multiple instances 34, 112, 280, 304, 536
 obtaining ACC instance handle 539
 obtaining TF instance pointer 305
 obtaining VPI instance handle 113
 overloading 32, 278
 returning function value 165, 331, 335, 343, 559
 types of functions 48–50, 165
system task/function arguments
 accessing with ACC routines 536–540
 accessing with VPI routines 65, 91, 111–112
 verifying with TF routines 310

T

TF constants
ERR_ERROR 309
ERR_INTERNAL 309
ERR_MESSAGE 309
ERR_SYSTEM 309
ERR_WARNING 309
REASON_DISABLE 391
REASON_ENDOFCOMPILE 313, 391, 394–395
REASON_ENDOFRESET 392
REASON_FINISH 391, 393–394
REASON_FORCE 391
REASON_INTERACTIVE 391, 393–394
REASON_PARAMDRC 391
REASON_PARAMVC 391, 403
REASON.REACTIVATE 391, 400
REASON_RELEASE 391
REASON_RESET 392
REASON.RESTART 391, 394
REASON_ROSYNCH 318, 391
REASON_SAVE 391, 394
REASON_SCOPE 391
REASON_STARTOFSAVE 391

REASON_SYNCH 391, 397
tf_integer_node 726
tf_memory_node 726
tf_netscalar_node 726
tf_netvector_node 726
tf_null_node 726
TF_NULLPARAM 312, 348
tf_nullparam 720, 727, 733
TF_READONLY 312, 348
tf_readonly 720, 727, 733
TF_READONLYREAL 312, 348
tf_READONLYREAL 720, 727, 733
TF_READWRITE 312, 348, 356
tf_readwrite 720, 727, 733
TF_READWRITEREAL 312, 348, 356
tf_readwritereal 720, 727, 733
tf_real_node 726
tf_reg_node 726
TF_RWBITLELECT 348, 356
tf_rwbitselect 720, 727
TF_RWMEMSELECT 348, 356
tf_rwmemselect 720, 727
TF_RWPARTSELECT 348, 356
tf_rwpartsselect 720, 727
TF_STRING 312, 348
tf_string 720, 727, 733
tf_time_node 726
userfunction 295, 638
userrealfunction 295, 638
usertask 295, 638
TF routines, definition 4, 30, 276, 301
TF work area 313–316
tf_add_long() 323, 718
tf_asynchoff() 403, 718
tf_asynchon() 403, 719
tf_clearalldelays() 400, 719
tf_compare_long() 719
tf_copypvc_flag() 405, 719
tf_divide_long() 324, 719
tf_dofinish() 307, 719
tf_dostop() 307, 720
tf_error() 308, 720
tf_evaluatep() 346, 353, 720
tf_exprinfo() 346, 720
tf_getcstringp() 345, 721
tf_getinstance() 305, 721
tf_getlongp() 332, 721
tf_getlongtime() 721
tf_getnextlongtime() 318, 722
tf_getp() 332, 722
tf_getpchange() 406, 722
tf_getreapl() 333, 722
tf_getrealtime() 318, 723
tf_gettime() 317–318, 723
tf_gettimeprecision() 319, 723
tf_gettimeunit() 319, 723
tf_getworkarea() 313, 723
tf_iasynchoff() 718
tf_iasynchon() 719
tf_iclearalldelays() 719
tf_icopypvc_flag() 719
tf_ievaluatep() 720
tf_iexprinfo() 720
tf_igetcstringp() 721
tf_igetlongp() 721
tf_igetlongtime() 721
tf_igetlp() 722
tf_igetpchange() 722
tf_igetreapl() 722
tf_igetrealtime() 723
tf_igettime() 723
tf_igettimelprecision() 723
tf_igettimeunit() 319, 723
tf_igetworkarea() 723
tf_imipname() 724
tf_imovepvc_flag() 725
tf_inodeinfo() 725
tf_inump() 726
tf_ipropagtep() 727
tf_iputlongp() 727
tf_iputp() 728
tf_iputreapl() 728
tf_irosynchronize() 728
tf_isetdelay() 729
tf_isetlongdelay() 729
tf_isetrealdelay() 730
tf_isetworkarea() 730
tf_isizep() 730
tf_ispname() 730
tf_istrdelpup() 731
tf_istrgetp() 731
tf_istrongdelputp() 732
tf_istrrealdelputp() 732
tf_isynchronize() 733
tf_itestpvc_flag() 733
tf_itypep() 733
tf_long_to_real() 324, 724
tf_longtime_tostr() 320, 724
tf_message() 309, 724
tf_mipname() 325, 724
tf_movepvc_flag() 405, 725
tf_multiply_long() 324, 725
tf_nodeinfo() 359, 367, 725
tf_nodeinfo(), portability issues 360
tf_nump() 311, 726
tf_propagtep() 346, 356, 379, 727

tf_putstrp() 334, 344, 727
 tf_putstrp() 330, 334, 344, 728
 tf_putreapl() 334, 344, 728
 tf_read_restart() 408, 728
 tf_real_to_long() 324, 728
 tf_rosynchronize() 397, 728
 tf_scale_longdelay() 321, 729
 tf_scale_realdelay() 321, 729
 tf_setdelay() 400, 729
 tf_setlongdelay() 400, 729
 tf_setrealdelay() 400, 730
 tf_setnetworkarea() 313, 730
 tf_sizep() 312, 730
 tf_spname() 325, 730
 tf_strdelputp() 339–344, 731
 tf_strgetp() 338, 731
 tf_strgettime() 731
 tf_strlongdelputp() 340–344, 732
 tf_strrealdelputp() 340–344, 732
 tf_subtract_long() 323, 733
 tf_synchronize() 397, 733
 tf_testpvc_flag() 405, 733
 tf_text() 310, 733
 tf_typep() 312, 331, 733
 tf_unscale_longdelay() 321, 734
 tf_unscale_realdelay() 321, 734
 tf_warning() 308, 734
 tf_write_save() 407, 734
 tferrmessage 296
 tfveritool 296
 time scaling 128–132, 317, 319–320,
 342–343
 time step activity 208
 time_value_change 596
 transport delay 162, 341, 554
 traversing hierarchy, using VPI 125
 type properties, ACC 526

U

unexpanded vectors 496
 user_data
 in ACC VCL routines 593
 in TF/ACC applications 297, 636
 in VPI applications 52, 633
 in VPI simulation callback
 routines 193–195
 user-defined invocation options 133, 322
 user-defined system tasks 30, 264, 276
 user-defined system tasks/functions 16
 userfunction 295
 userrealfunction 295
 usertask 295

V

Value Change Link 591
 variable array 367
 variable arrays 171
 variables, definition 77
 VCL 591
 vcl_verilog 595
 vcl_verilog_logic 592–593
 vcl_verilog_strength 592–593
 vcl0 597
 vcl1 597
 vclHighZ 598
 vclLarge 598
 vclMedium 598
 vclPull 598
 vclSmall 598
 vclStrong 598
 vclSupply 598
 vclWeak 598
 vclX 597
 vclZ 597
 vconfig 644
 VCS 667
 vector_value_change 596
 VeriBest 674
 Verilog-XL 641
 veriuser.c 643
 veriuser.h, including 302
 veriusertfs array 293, 636
 vlog_startup_routines 51, 635
 vlog_startup_routines, limitations 52
 VPI constants
 cbAfterDelay 204, 211, 715
 cbAssign 220, 222–223
 cbAtStartOfSimTime 204, 211, 715
 cbDeassign 220, 223
 cbDisable 220
 cbEndOfCompile 196, 199, 715
 cbEndOfRestart 197, 715
 cbEndOfSave 197, 715
 cbEndOfSimulation 197, 715
 cbEnterInteractive 197, 715
 cbError 197, 715
 cbExitInteractive 197, 715
 cbForce 220, 222–223, 715
 cbInteractiveScopeChange 197, 715
 cbNextSimTime 204, 211, 715
 cbPLIError 197
 cbReadOnlySynch 204, 206–207, 209, 715
 cbReadWriteSynch 204, 206–207, 209, 715
 cbRelease 220, 223, 715
 cbStartOfRestart 197, 715

cbStartOfSave 197, 715
cbStartOfSimulation 196, 199, 715
cbStmt 220, 223, 715
cbTchkViolation 197, 715
cbUnresolvedSystf 197, 715
cbValueChange 220, 222–223, 715
compiletf routine 199
vlog_startup_routines 199
vpi0 141, 143, 155, 710, 714
vpi1 141, 143, 155, 710, 714
vpiAddOp 686
vpiAlways 698
vpiAndPrim 695
vpiAnyEdge 704
vpiArray 171–172, 692
vpiAssignment 698–699
vpiAssignStmt 698, 702
vpiBegin 698
vpiBinaryConst 685, 694, 703
vpiBinStrVal 70, 143, 149, 221, 710, 714
vpiBitAndOp 686
vpiBitNegOp 686
vpiBitOrOp 686
vpiBitXNorOp 686
vpiBitXorOp 686
vpiBlocking 699
vpiBufif0Prim 695
vpiBufif1Prim 695
vpiBufPrim 695
vpiCallback 705
vpiCancelEvent 162
vpiCase 698, 701
vpiCaseEqOp 686
vpiCaseExact 701
vpiCaseItem 701
vpiCaseNeqOp 686
vpiCaseType 701
vpiCaseX 701
vpiCaseZ 701
vpiCellInstance 683
vpiChargeStrength 689
vpiCmosPrim 695
vpiCombPrim 695
vpiCompile 101, 707
vpiConcatOp 686
vpiConditionOp 686
vpiConnByName 687
vpiConstant 685
vpiConstType 685, 694, 703
vpiContAssign 693
vpiDataPolarity 703
vpiDeassign 698, 702
vpiDecConst 685, 694, 703
vpiDecStrVal 143, 149, 221, 710, 714
vpiDefDecayTime 684
vpiDefDelayMode 684
vpiDefFile 684
vpiDefLineNo 684
vpiDefName 68, 684, 695
vpiDefNetType 683
vpiDefParam 694
vpiDelayControl 698–699
vpiDelayModeDistrib 684
vpiDelayModeMTM 684
vpiDelayModeNone 684
vpiDelayModePath 684
vpiDelayModeUnit 684
vpiDelayModeZero 684
vpiDirection 687–688, 695, 703
vpiDisable 698, 702
vpiDivOp 686
vpiDontCare 141, 143, 710, 714
vpiEdge 704
vpiEdge01 704
vpiEdge0x 704
vpiEdge10 704
vpiEdge1x 704
vpiEdgex0 704
vpiEdgex1 704
vpiEqOp 686
vpiError 101, 707
vpiEventControl 698–699
vpiEventOrOp 686
vpiEventStmt 698
vpiExpanded 689
vpiExplicitName 687
vpiExplicitScalared 689
vpiExplicitVectored 689
vpiFile 687–694, 696–704
vpiFor 698, 700
vpiForce 698, 702
vpiForceFlag 162
vpiForever 698, 700
vpiFork 698
vpiFullName 68, 684, 689–695, 703
vpiFuncCall 685, 697–698
vpiFuncInt 709, 716
vpiFuncReal 709, 716
vpiFuncSized 709, 716
vpiFuncTime 709, 716
vpiFunction 684, 696
vpiGate 695
vpiGeOp 686
vpiGtOp 686
vpiH 141, 143, 710, 714
vpiHexConst 685, 694, 703

vpiHexStrVal **70, 143, 149, 221, 710, 714**
 vpiHighZ **683**
 vpiHiZ **710, 714**
 vpiHold **704**
 vpiIf **698, 701**
 vpiIfElse **698, 701**
 vpiImplicitDecl **689**
 vpiIndex **171–172, 695**
 vpiInertialDelay **162, 165, 713**
 vpiInitial **698**
 vpiInout **687–688, 695, 703**
 vpiInput **687–688, 695, 703**
 vpiIntegerVar **685, 691–692**
 vpiInterModPath **109, 704, 711**
 vpiInternal **101, 707**
 vpiInertialDelay **164**
 vpiIntFunc **50, 166, 634**
 vpiIntVal **143, 145, 222, 710, 714**
 vpiIODecl **688**
 vpiIterator **712**
 vpiL **141, 143, 710, 714**
 vpiLargeCharge **710, 714**
 vpiLeftRange **171–172, 692**
 vpiLeOp **686**
 vpiLhs **694**
 vpiLineNo **687–694, 696–704**
 vpiListOp **686**
 vpiLoad **690**
 vpiLogAndOp **686**
 vpiLogOrOp **686**
 vpiLShiftOp **686**
 vpiLtOp **686**
 vpiMediumCharge **710, 714**
 vpiMemory **685, 691–692**
 vpiMemoryWord **685, 691**
 vpiMinTypMaxOp **686**
 vpiMinusOp **686**
 vpiMixedIO **687–688, 695, 703**
 vpiModOp **686**
 vpiModule **683–684**
 vpiMultiConcatOp **686**
 vpiMultOp **686**
 vpiName **68, 684, 687–695, 697, 703**
 vpiNamedBegin **684, 698**
 vpiNamedEvent **693**
 vpiNamedFork **684, 698**
 vpiNandPrim **695**
 vpiNegative **703**
 vpiNegedge **704**
 vpiNegedgeOp **686**
 vpiNeqOp **686**
 vpiNet **685, 689**
 vpiNetDeclAssign **689, 693**
 vpiNetType **689**
 vpiNmosPrim **695**
 vpiNoChange **704**
 vpiNoDelay **162, 164–165, 713**
 vpiNoDirection **687–688, 695, 703**
 vpiNoEdge **704**
 vpiNorPrim **695**
 vpiNotice **101, 707**
 vpiNotif0Prim **695**
 vpiNotif1Prim **695**
 vpiNotOp **686**
 vpiNotPrim **695**
 vpiNullOp **686**
 vpiNullStmt **698**
 vpiObjTypeVal **143, 160, 222, 710, 714**
 vpiOctConst **685, 694, 703**
 vpiOctStrVal **143, 149, 221, 710, 714**
 vpiOperation **685**
 vpiOpType **685**
 vpiOrPrim **695**
 vpiOutput **687–688, 695, 703**
 vpiParamAssign **694**
 vpiParameter **685, 694**
 vpiPartSelect **685**
 vpiPathFull **703**
 vpiPathParallel **703**
 vpiPathTerm **703**
 vpiPathType **703**
 vpiPeriod **704**
 vpiPLI **101, 707**
 vpiPlusOp **686**
 vpiPmosPrim **695**
 vpiPolarity **703**
 vpiPort **687, 690**
 vpiPortBit **687**
 vpiPortIndex **687**
 vpiPortInst **690**
 vpiPosedge **704**
 vpiPosedgeOp **686**
 vpiPositive **703**
 vpiPrimTerm **695**
 vpiPrimType **695**
 vpiProtected **683, 695**
 vpiPull1 **683**
 vpiPulldownPrim **695**
 vpiPullDrive **710, 714**
 vpiPullupPrim **695**
 vpiPureTransportDelay **162, 164, 713**
 vpiRcmosPrim **695**
 vpiRealConst **685, 694, 703**
 vpiRealFunc **50, 166, 634**
 vpiRealVal **143, 146, 222, 710, 714**
 vpiRealVar **685, 692**

vpiRecovery **704**
vpiReg **685, 690**
vpiRegBit **690**
vpiRelease **698, 702**
vpiReleaseFlag **162**
vpiRepeat **698, 700**
vpiRepeatControl **699**
vpiReturnEvent **162, 714**
vpiRightRange **171–172, 691–692**
vpiRnmosPrim **695**
vpiRpmosPrim **695**
vpiRShiftOp **686**
vpiRtranif0Prim **695**
vpiRtranif1Prim **695**
vpiRtranPrim **695**
vpiRun **101, 707**
vpiScalar **687–690**
vpiScalarVal **143, 222, 710, 714**
vpiScaledRealDelay **165**
vpiScaledRealTime **71, 130, 163, 174, 205, 221, 708–710, 713–715**
vpiSchedEvent **714**
vpiSeqPrim **695**
vpiSetup **704**
vpiSetupHold **704**
vpiSimTime **71–72, 130, 163, 174, 205, 221, 708–710, 713–715**
vpiSize **171–172, 685, 687–692, 695**
vpiSizedFunc **50, 166, 634**
vpiSkew **704**
vpiSmallCharge **710, 714**
vpiSpecParam **703**
vpiStrength0 **689, 693, 695**
vpiStrength1 **689, 693, 695**
vpiStrengthVal **143, 155, 222, 710, 714**
vpiStringConst **685, 694, 703**
vpiStringVal **143, 148, 222, 710, 714**
vpiStrongDrive **710, 714**
vpiSubOp **686**
vpiSupply0 **683, 689**
vpiSupply1 **683, 689**
vpiSupplyDrive **710, 714**
vpiSuppressTime **205, 221, 708–710, 713–715**
vpiSuppressVal **221–222, 710, 714**
vpiSwitch **695**
vpiSysFunc **49, 633, 709, 716**
vpiSysFuncCall **685, 697**
vpiSysFuncInt **49, 166, 633–634, 697**
vpiSysFuncReal **49–50, 166, 633–634, 697**
vpiSysFuncSized **49–50, 166, 633–634, 697**
vpiSysFuncTime **49–50, 166, 633–634, 697**
vpiSysFunction **166**
vpiSysFuncType **697**
vpiSysTask **49, 633, 709, 716**
vpiSysTaskCall **697**
vpiSystem **101, 707**
vpiSysTfCall **697**
vpiTableEntry **695**
vpiTask **684, 696**
vpiTaskCall **696–698**
vpiTchk **704**
vpiTchkDataTerm **704**
vpiTchkRefTerm **704**
vpiTchkTerm **704**
vpiTchkType **704**
vpiTimeFunc **50, 166, 634**
vpiTimePrecision **131, 683**
vpiTimeQueue **705–706**
vpiTimeUnit **131, 683, 707**
vpiTimeVal **143, 222, 710, 714**
vpiTimeVar **685, 692**
vpiTopModule **683**
vpiTranif0Prim **695**
vpiTranif1Prim **695**
vpiTranPrim **695**
vpiTransportDelay **162, 164, 713**
vpiTri **683, 689**
vpiTri0 **683, 689**
vpiTri1 **683, 689**
vpiTriAnd **683, 689**
vpiTriOr **683, 689**
vpiTriReg **683, 689**
vpiType **683–685, 687–694, 696–706**
vpiUDP **695**
vpiUDPDefn **695**
vpiUnaryAndOp **686**
vpiUnaryNandOp **686**
vpiUnaryNorOp **686**
vpiUnaryOrOp **686**
vpiUnaryXNorOp **686**
vpiUnaryXorOp **686**
vpiUnconnDrive **683**
vpiUnknown **703**
vpiUse **686**
vpiUserDefn **697**
vpiUserSystf **696**
vpiVarSelect **685**
vpiVector **687–690**
vpiVectorVal **143, 150, 222, 710, 714**
vpiWait **698, 700**
vpiWand **683, 689**
vpiWarning **101, 707**
vpiWeakDrive **710, 714**
vpiWhile **698, 700–701**

vpiWidth **704**
 vpiWire **683, 689**
 vpiWor **683, 689**
 vpiX **141, 143, 155, 710, 714**
 vpiXnorPrim **695**
 vpiXorPrim **695**
 vpiZ **141, 143, 155, 710, 714**
 VPI error handling **100**
 VPI Interface Mechanism **632**
 VPI object diagram legend **682**
 VPI object diagrams, how to use **102**
 VPI routines, advantages **57**
 VPI routines, definition **5, 30, 276**
 VPI string buffer **99**
 VPI work area **186**
 vpi_chk_error() **100, 707**
 vpi_compare_objects() **137, 707**
 vpi_free_object() **66, 108, 707**
 vpi_get() **67, 707**
 vpi_get_cb_info() **708**
 vpi_get_delays() **173, 708**
 vpi_get_str() **67, 709**
 vpi_get_systf_info() **709**
 vpi_get_time() **71, 130, 709**
 vpi_get_value() **69, 141, 710**
 vpi_get_vlog_info() **133, 711**
 vpi_handle() **62, 106, 711**
 vpi_handle_by_index() **110, 711**
 vpi_handle_by_name() **109–110, 711**
 vpi_handle_multi() **109, 711**
 vpi_iterate() **63–64, 106, 712**
 vpi_mcd_close() **128, 712**
 vpi_mcd_name() **128, 712**
 vpi_mcd_open() **127, 712**
 vpi_mcd_printf() **128, 712**
 vpi_printf() **66, 712**
 vpi_put_delays() **179, 713**
 vpi_put_value() **161, 713**
 vpi_register_cb() **192, 715**
 vpi_register_systf() **716**
 vpi_remove_cb() **196, 716**
 vpi_scan() **63–64, 107, 716**
 vpi_user.h, including **57**
 vpiHandle data type **60**
 vregister_value_change **596**

W

work area
 for TF/ACC applications **313–316**
 for VPI applications **186**